

PERFORMANCE MATTERS

Emery Berger

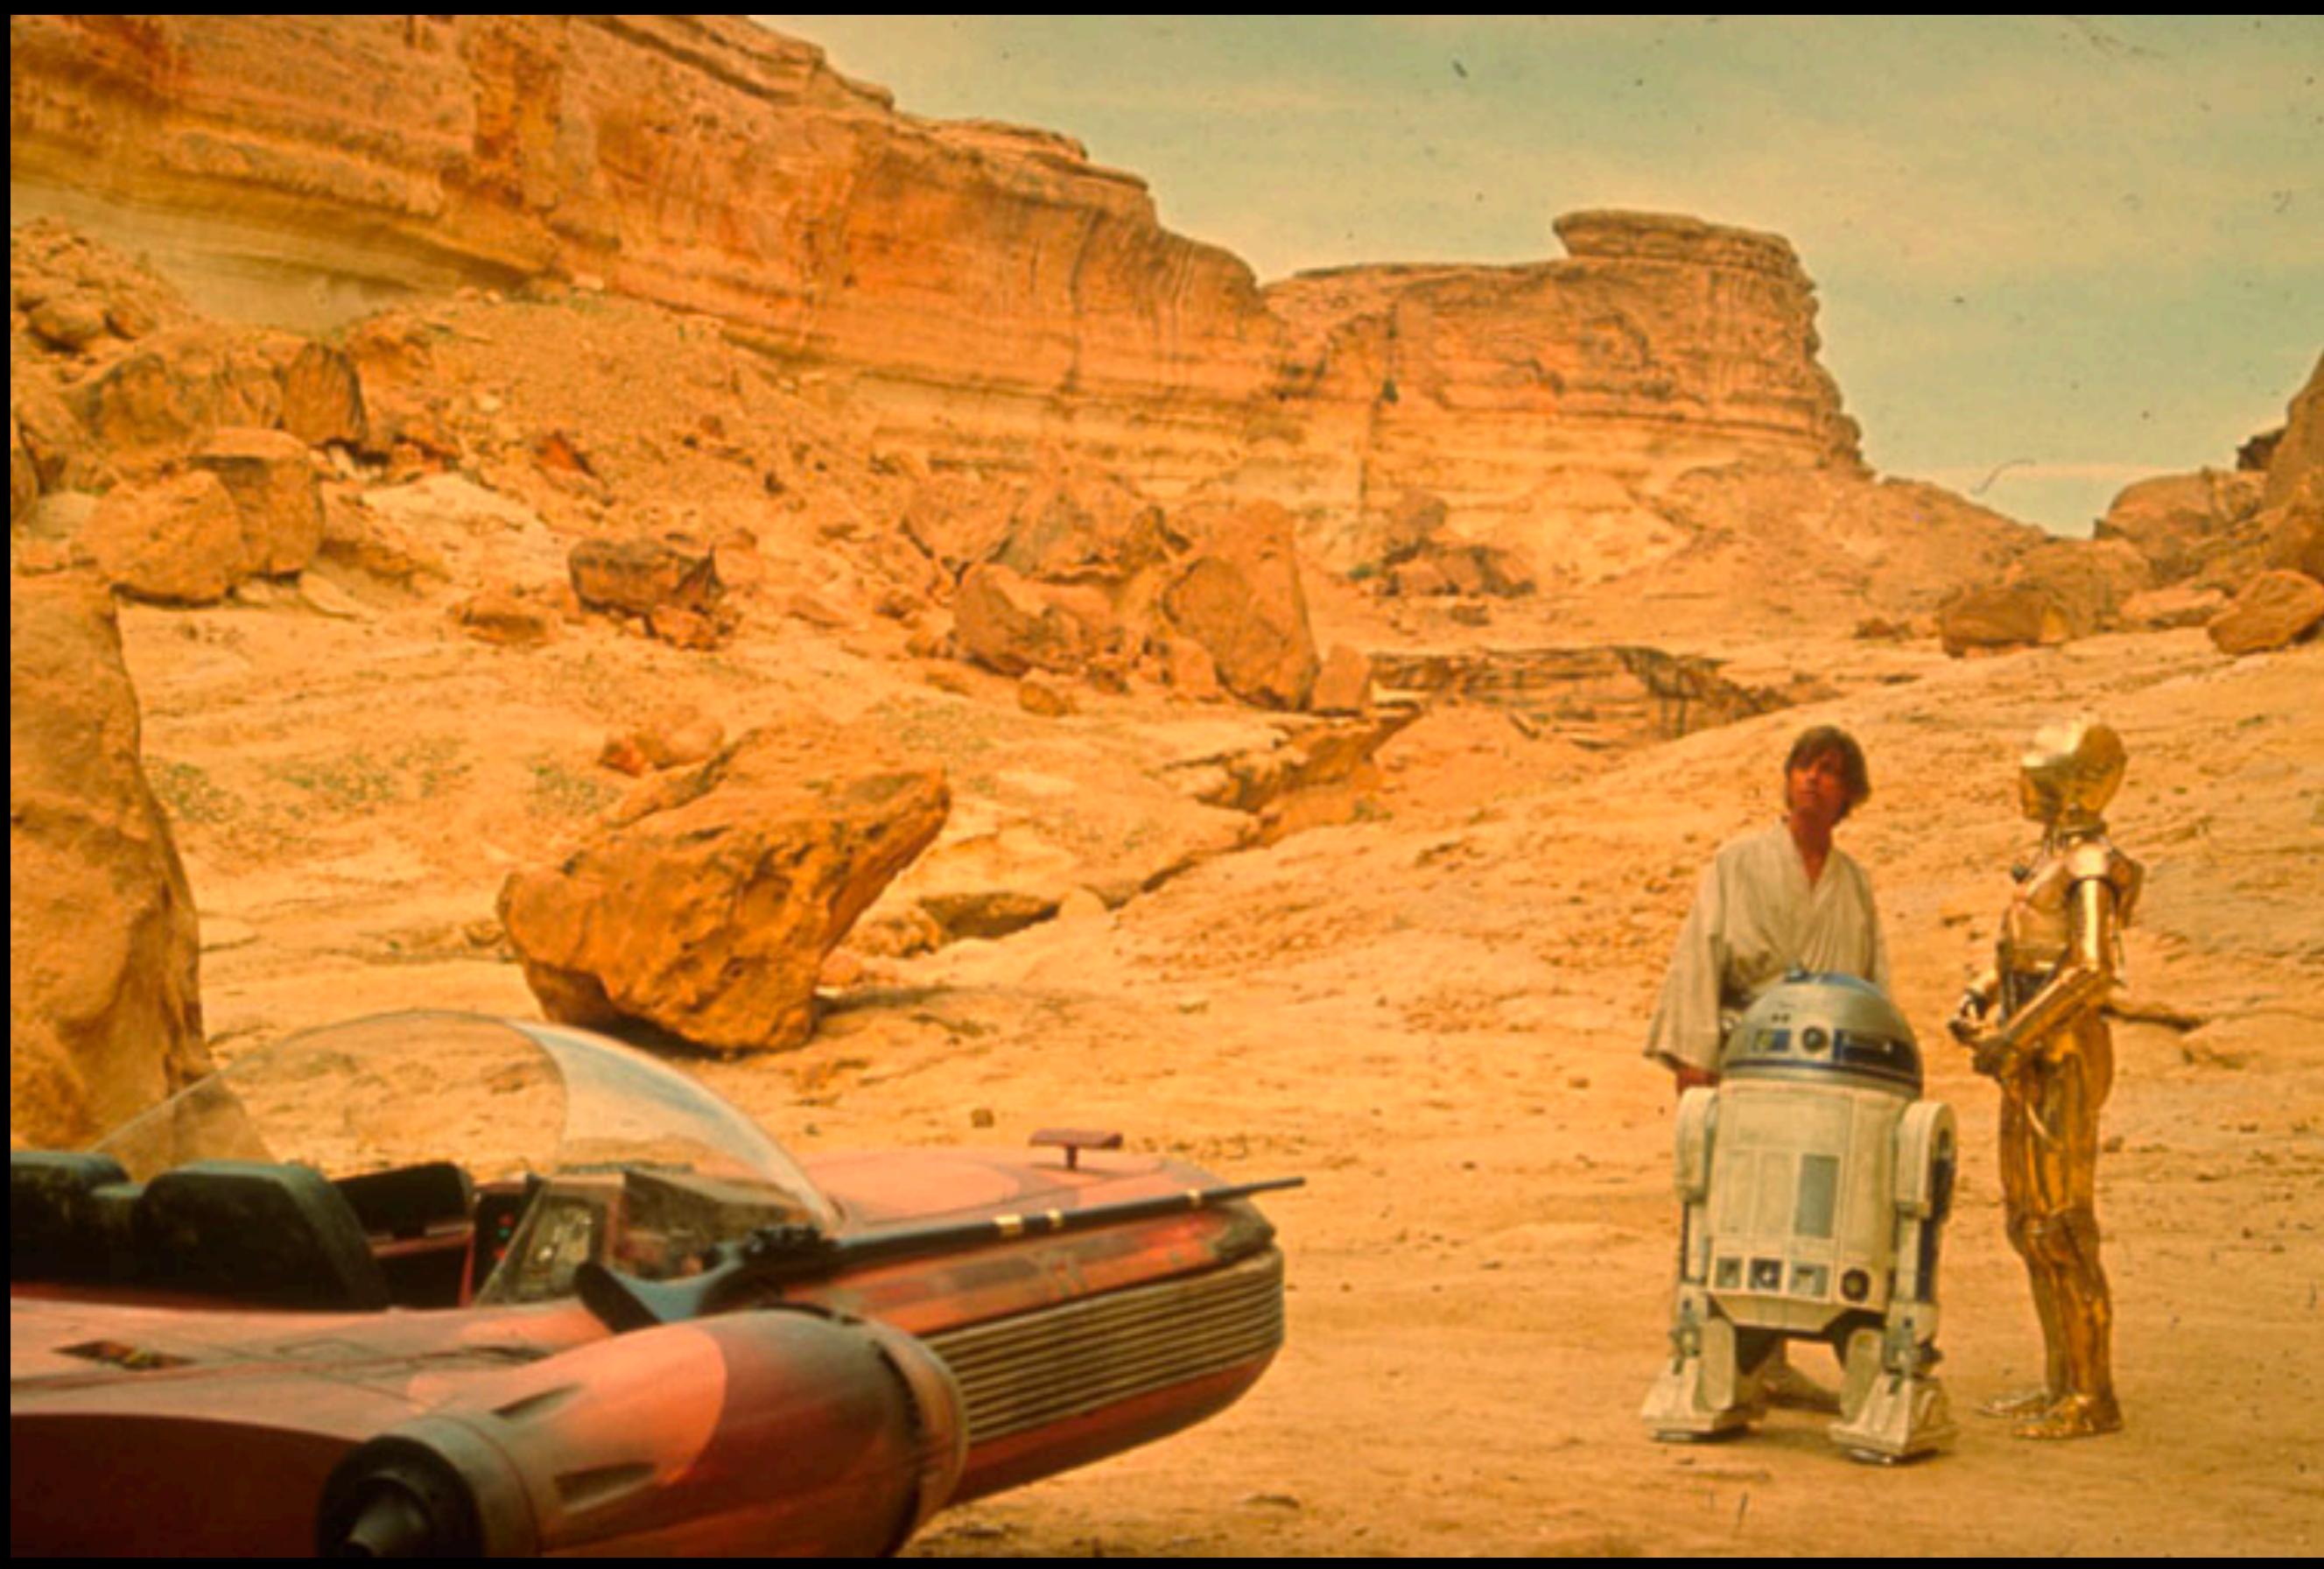
*College of Information
and Computer Sciences*

UMASS AMHERST

(joint work with Charlie Curtsinger, Grinnell College)

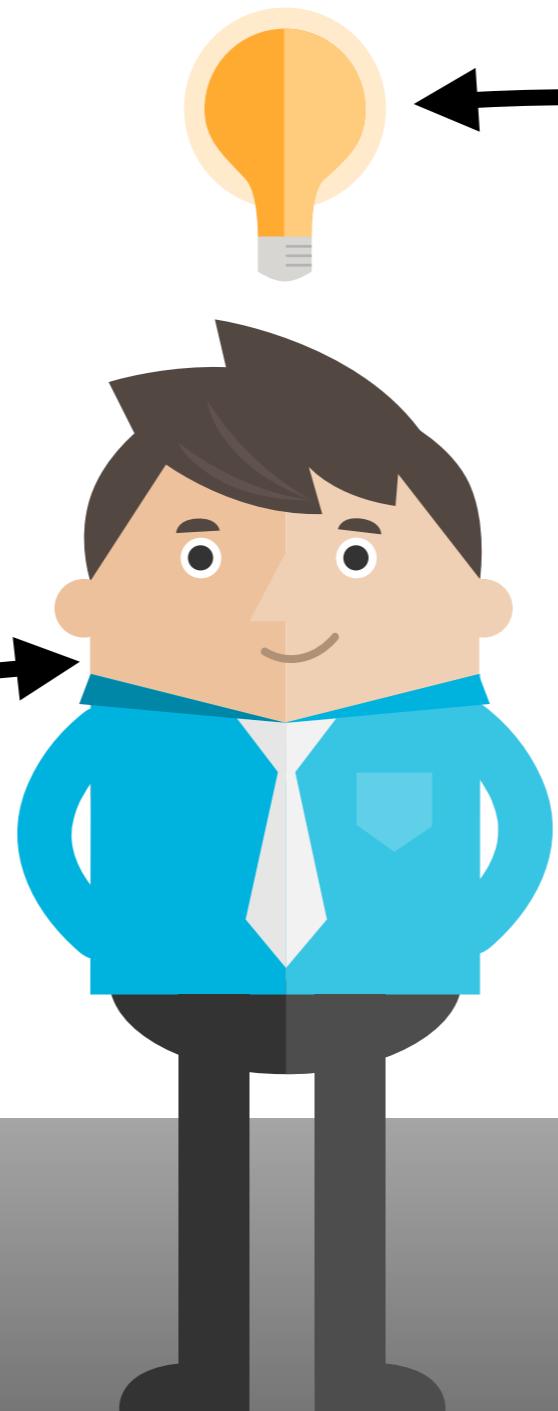
emeryberger.com, @emeryberger

A short time ago in a
valley far, far away...

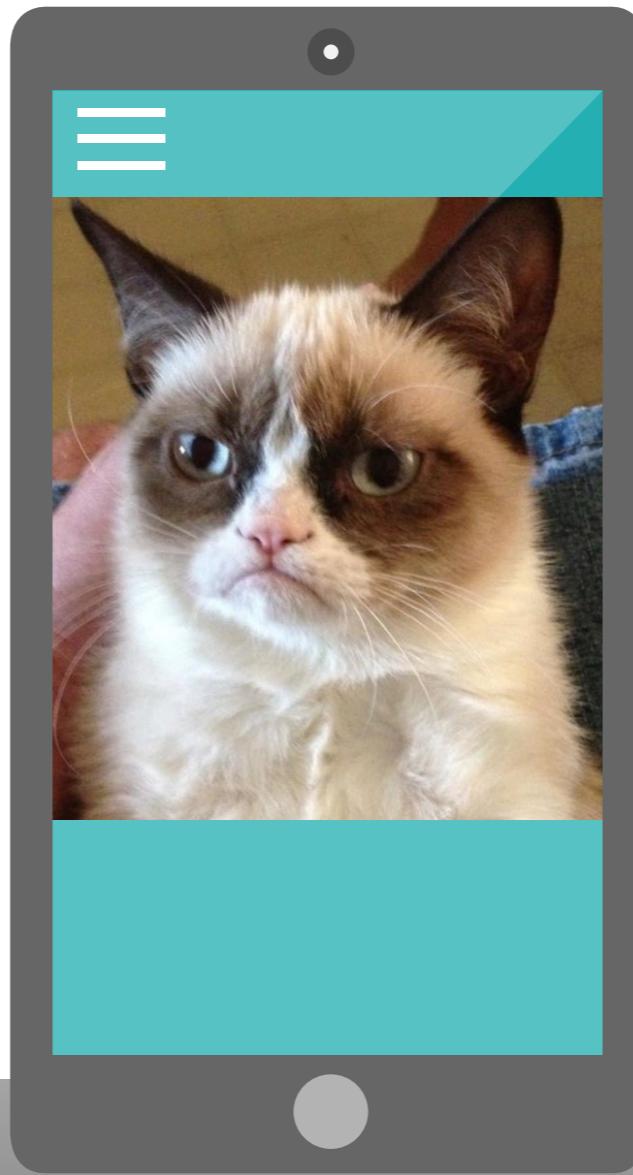




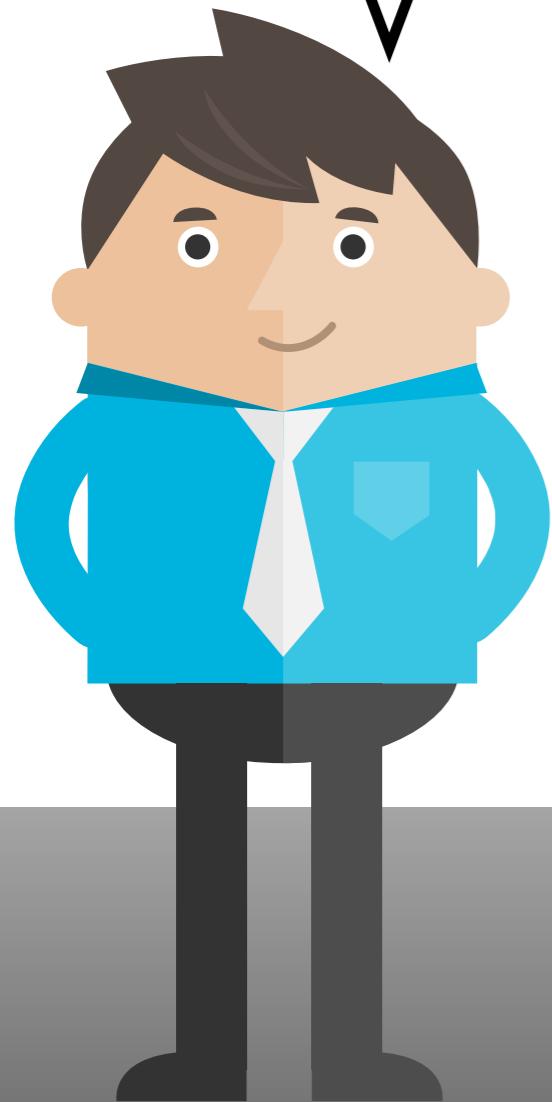
Luke
(our hero)



Luke has an idea
for an app.

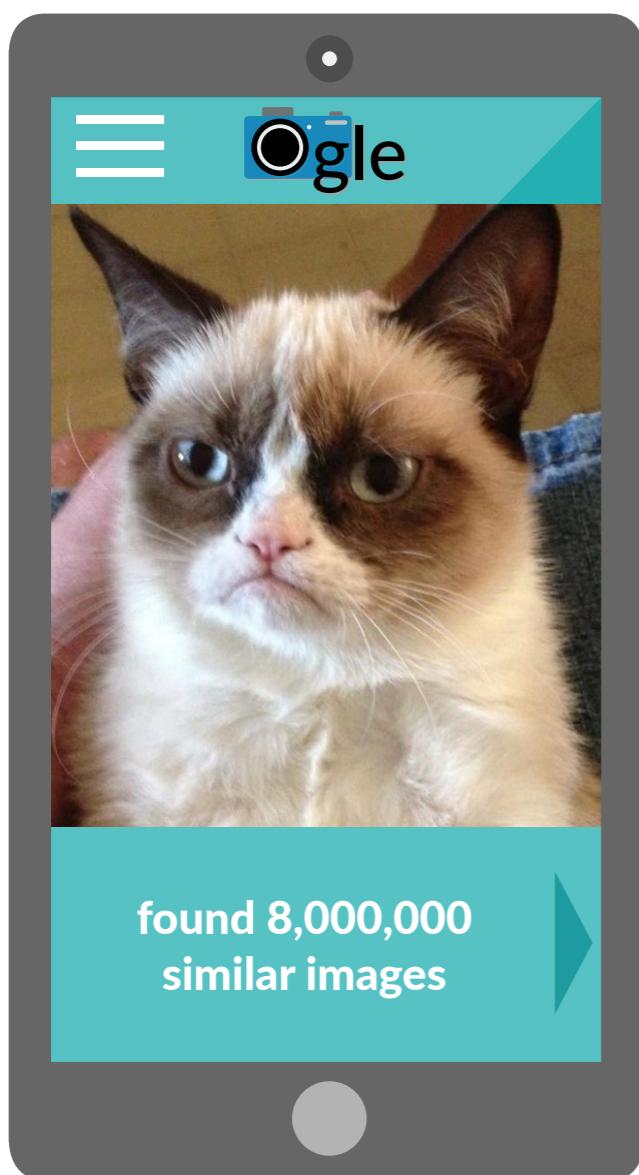


It's going to
totally disrupt
image search.

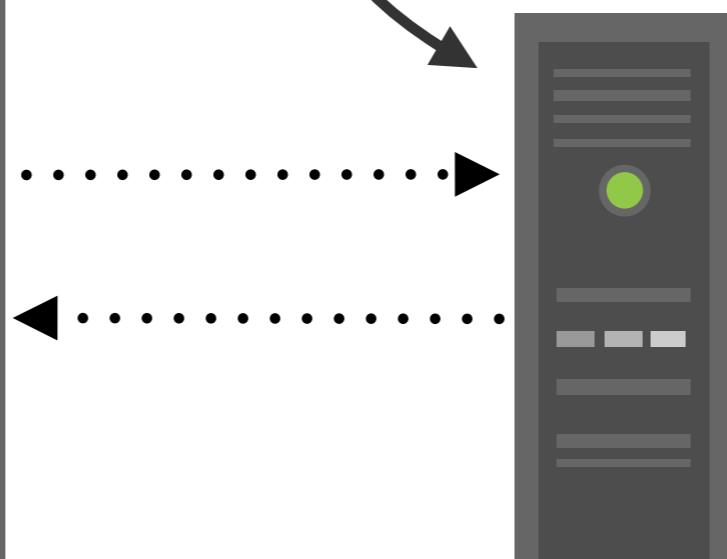


The Prototype Ogle

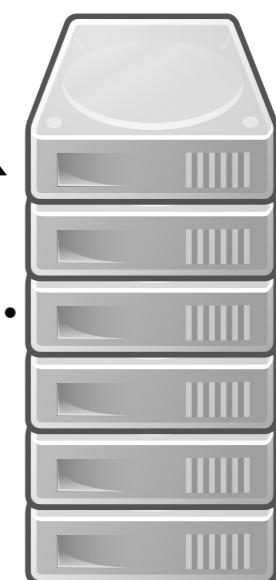
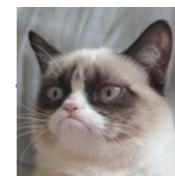
Take a picture



Send it to Ogle

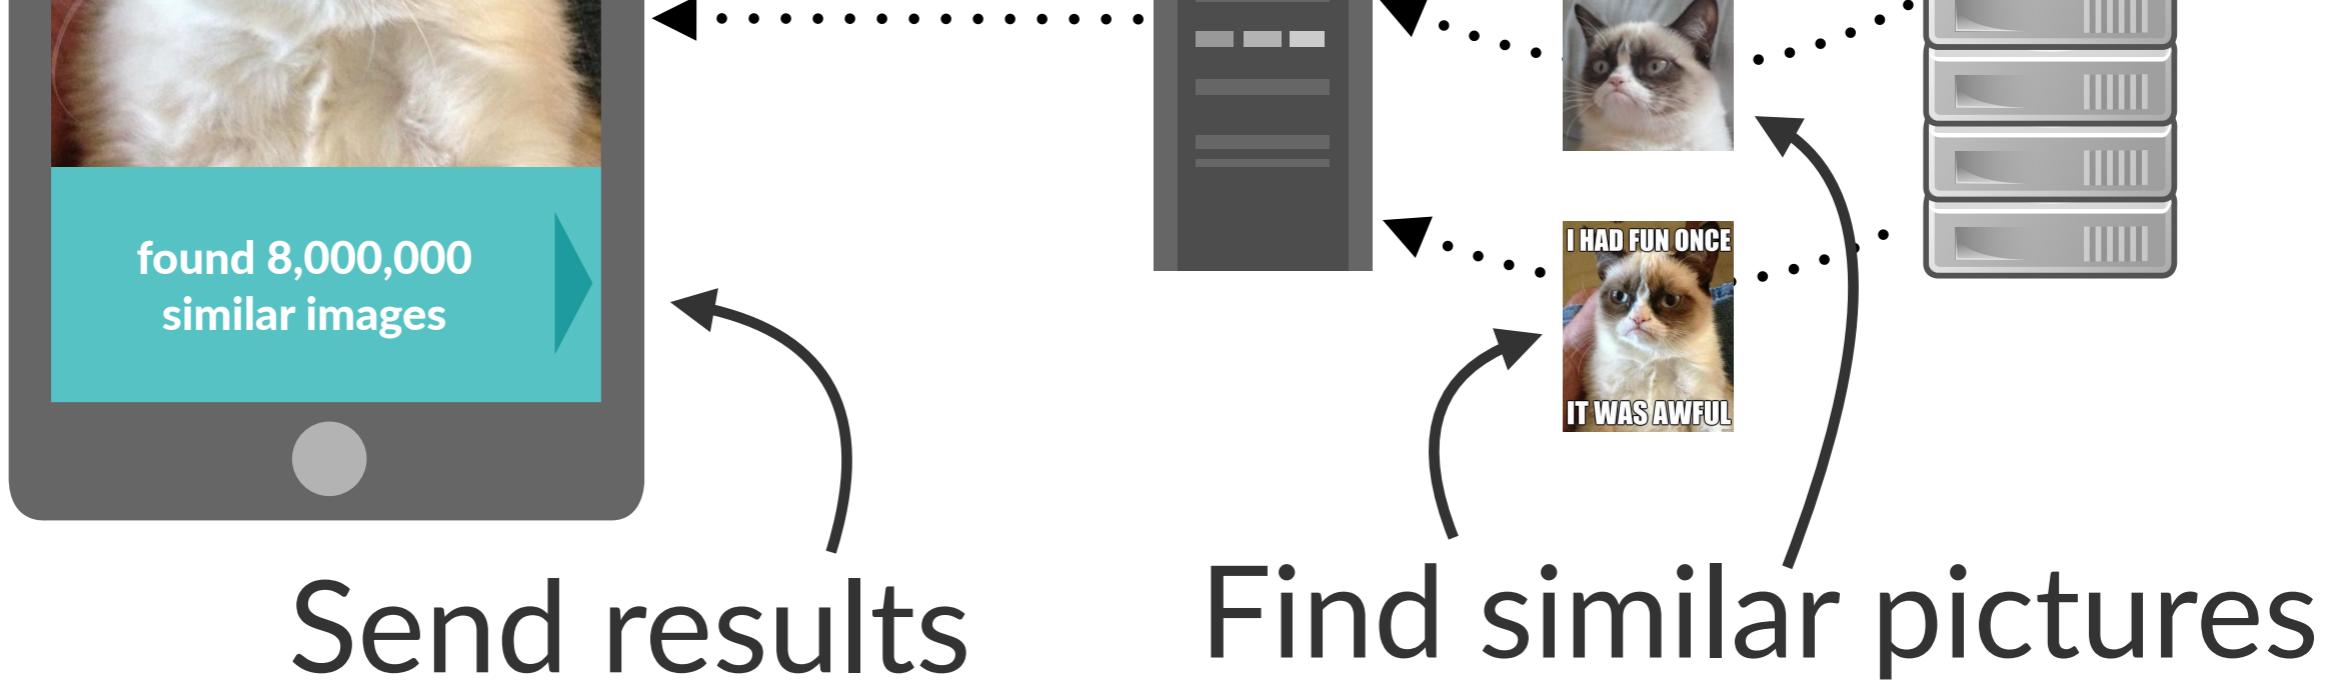


Add it to
the database

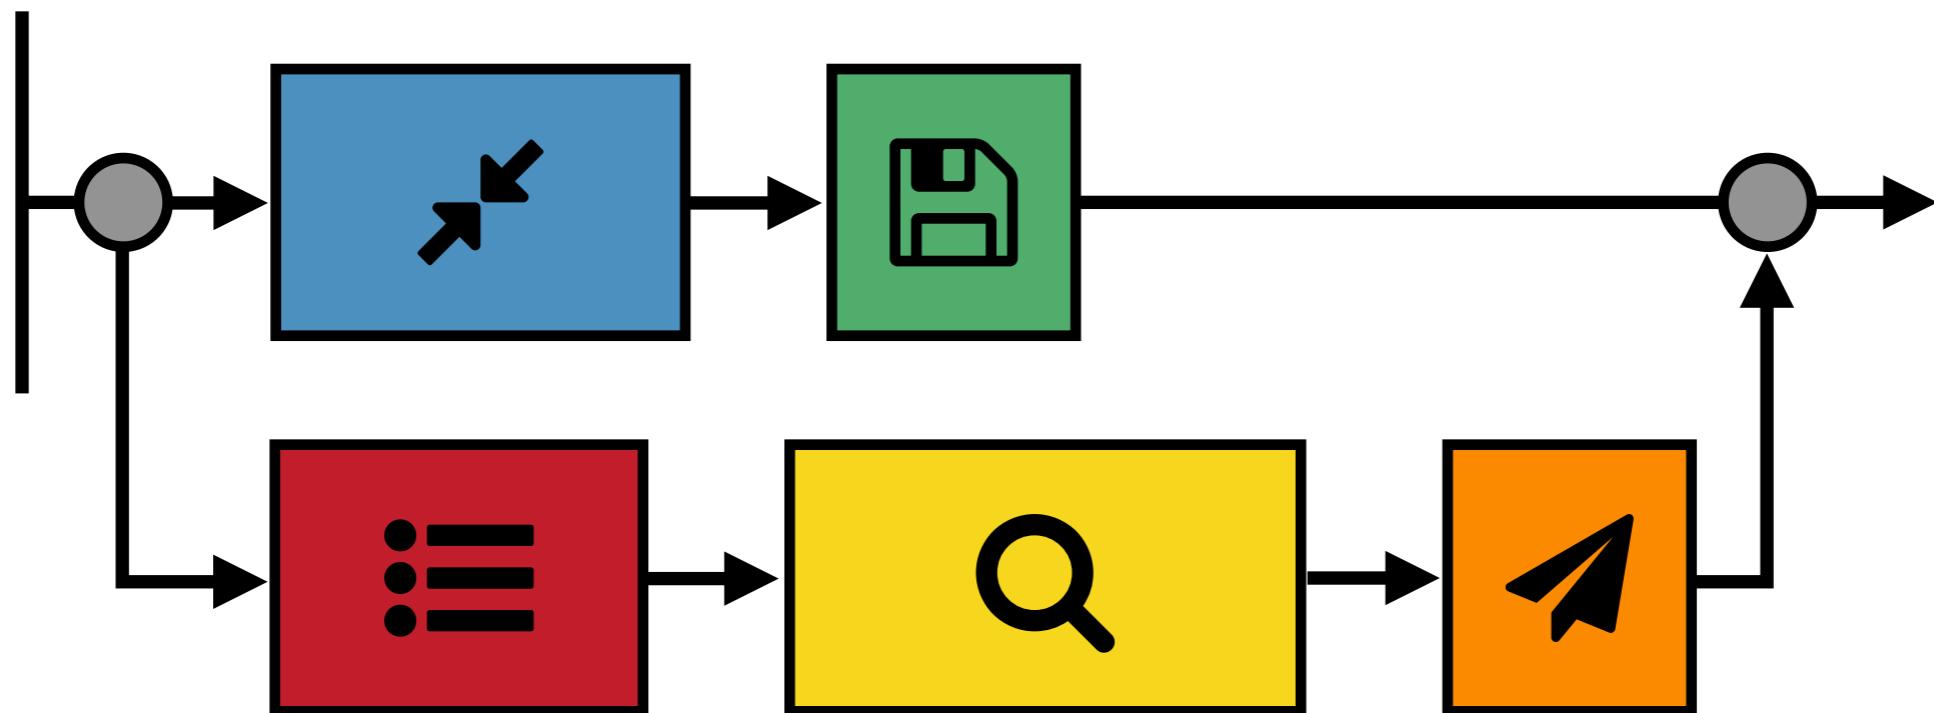


Find similar pictures

Send results



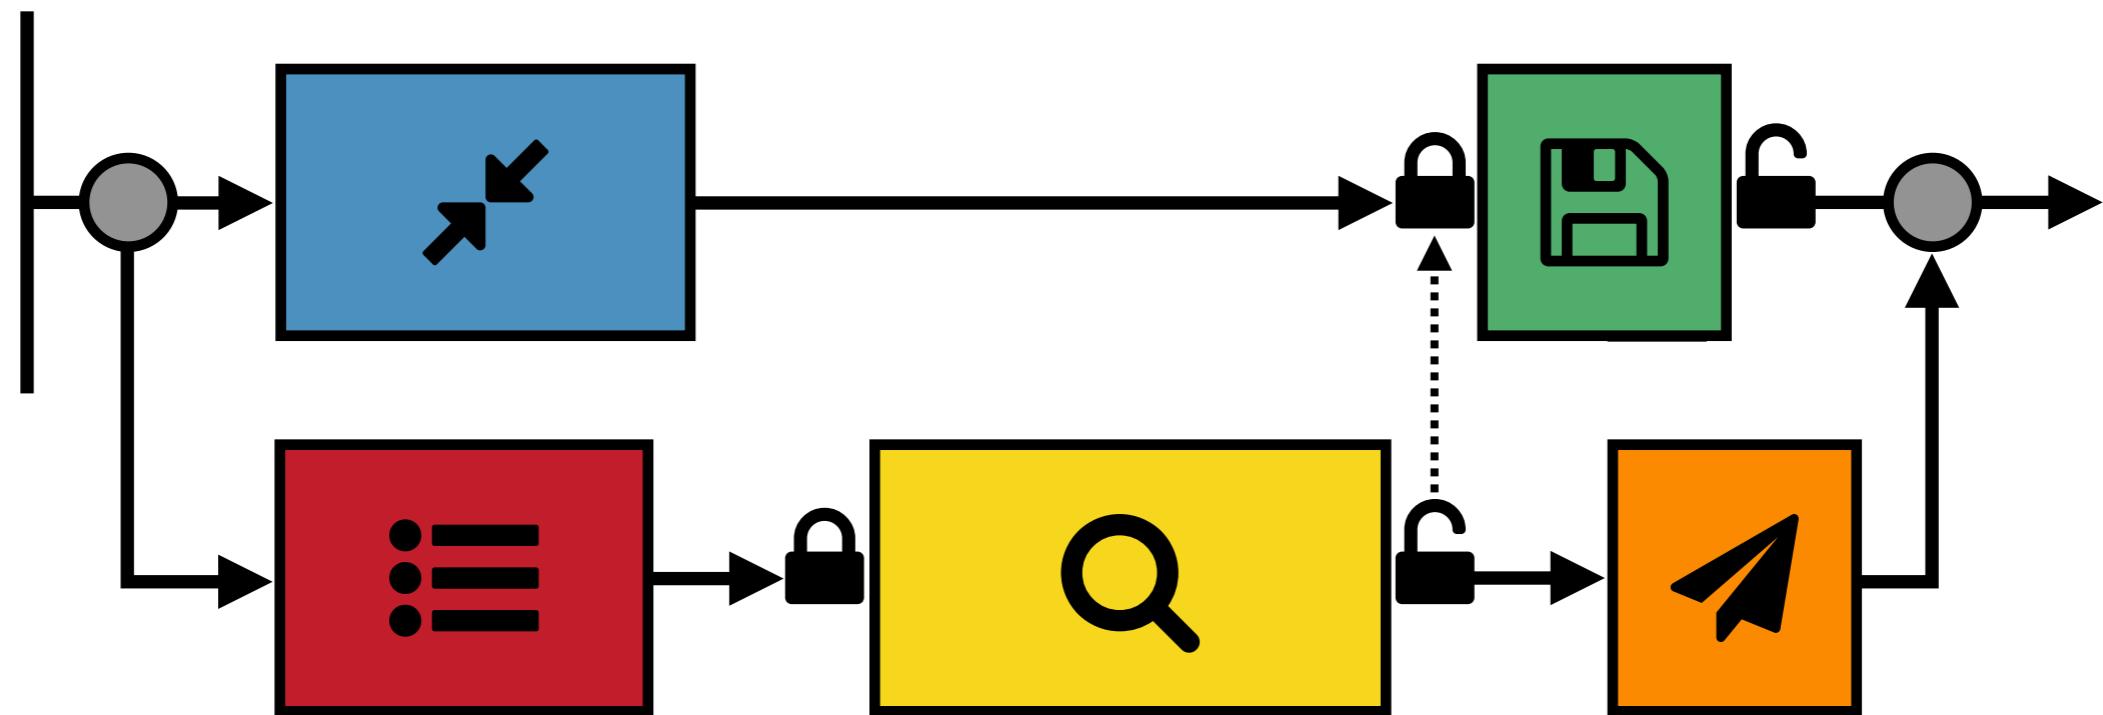
Find similar pictures



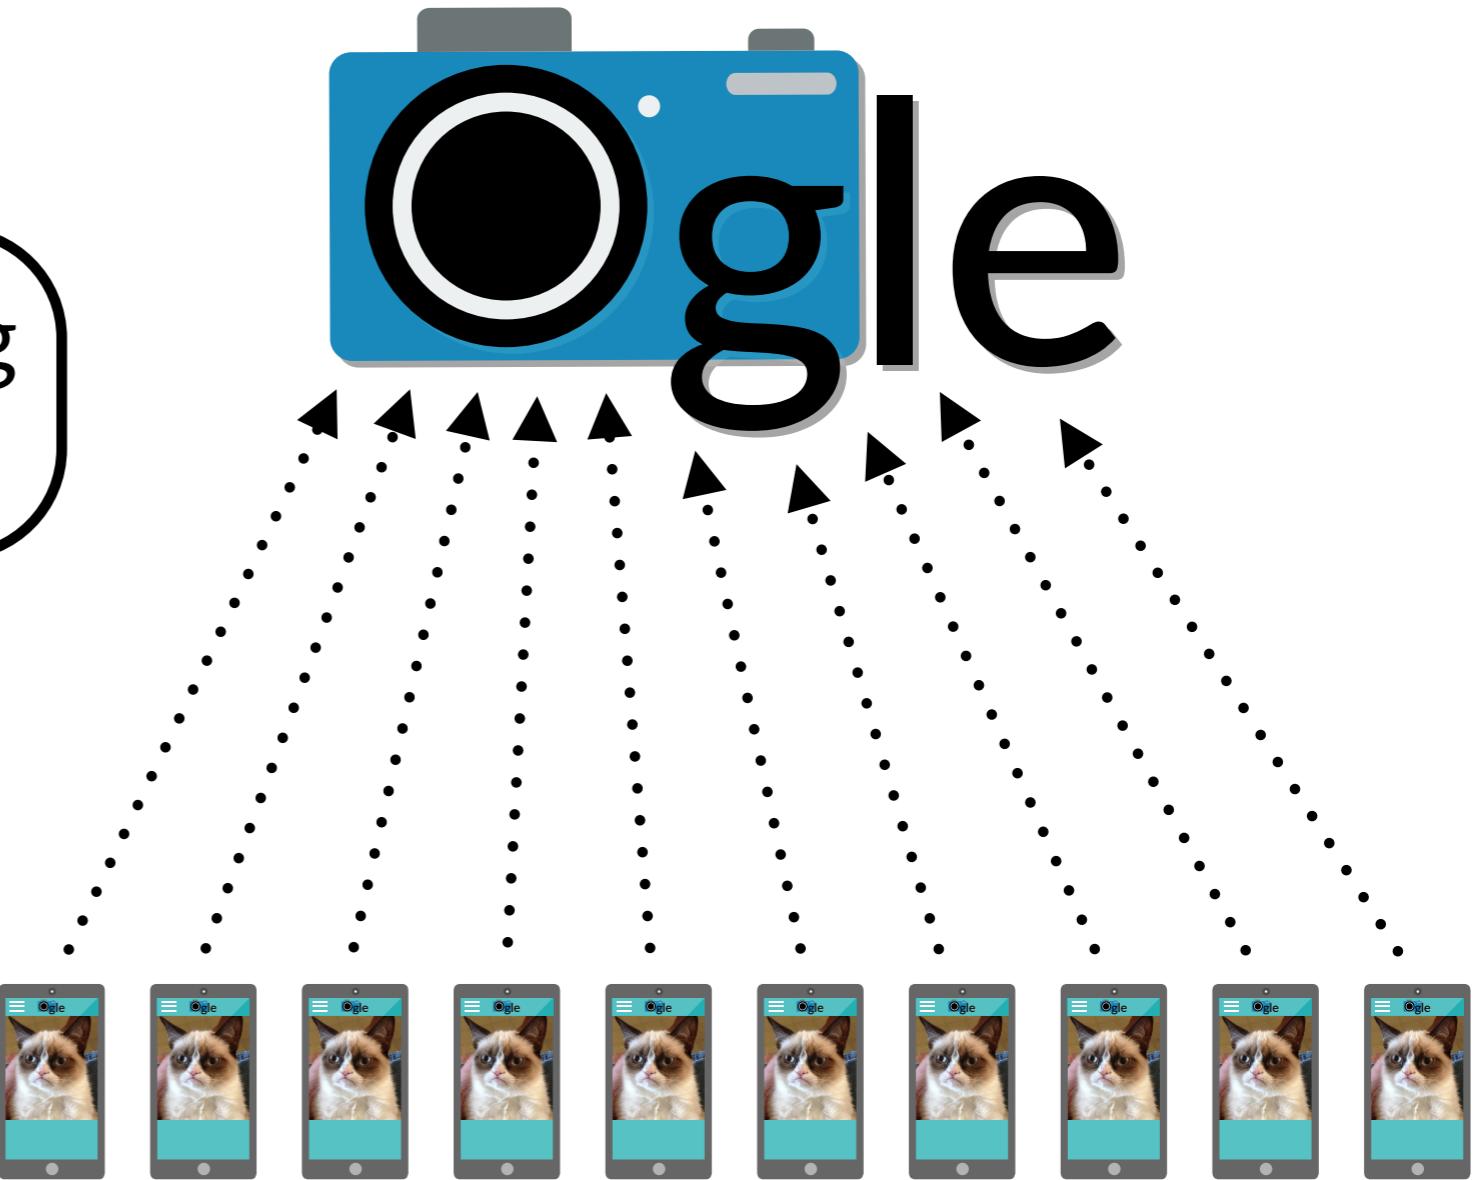
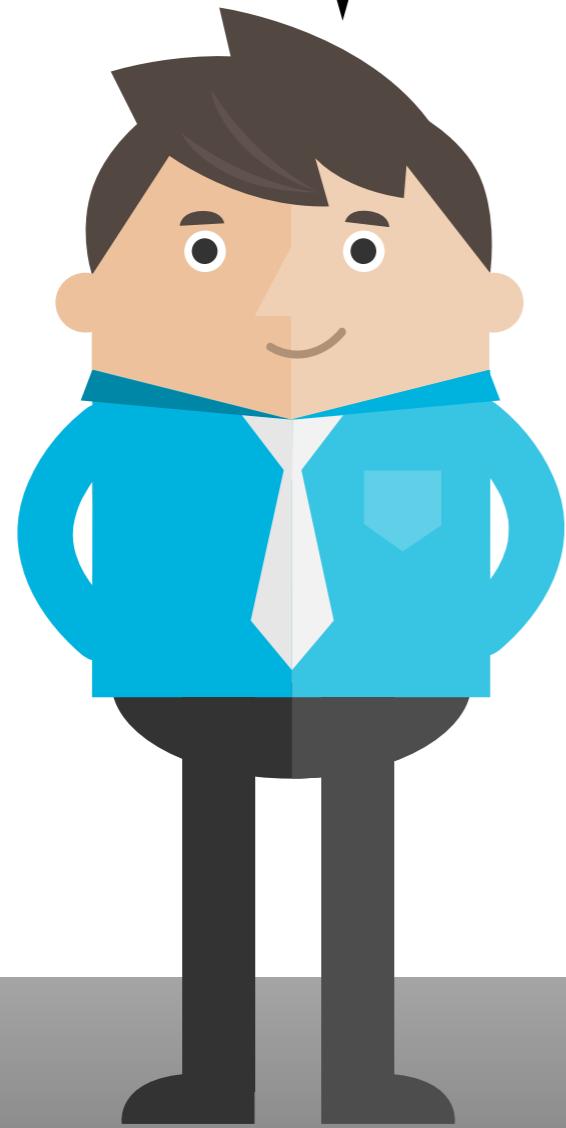


Send results

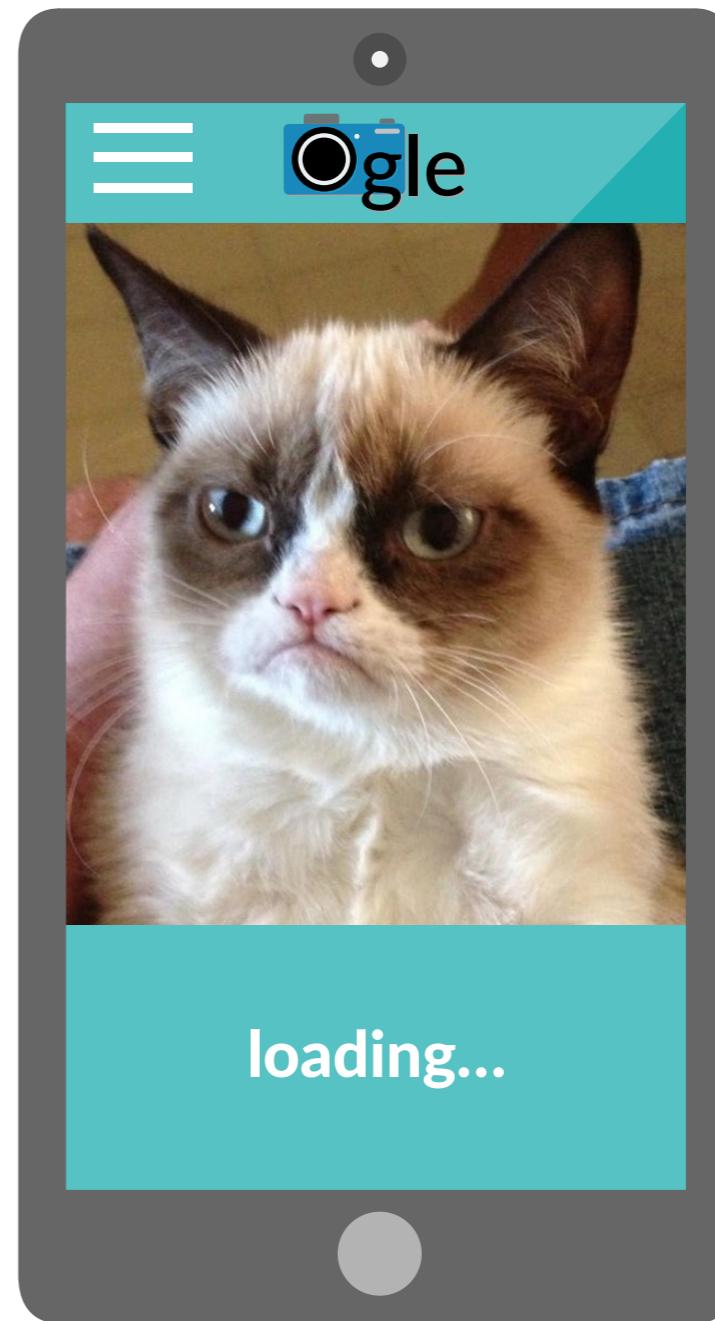
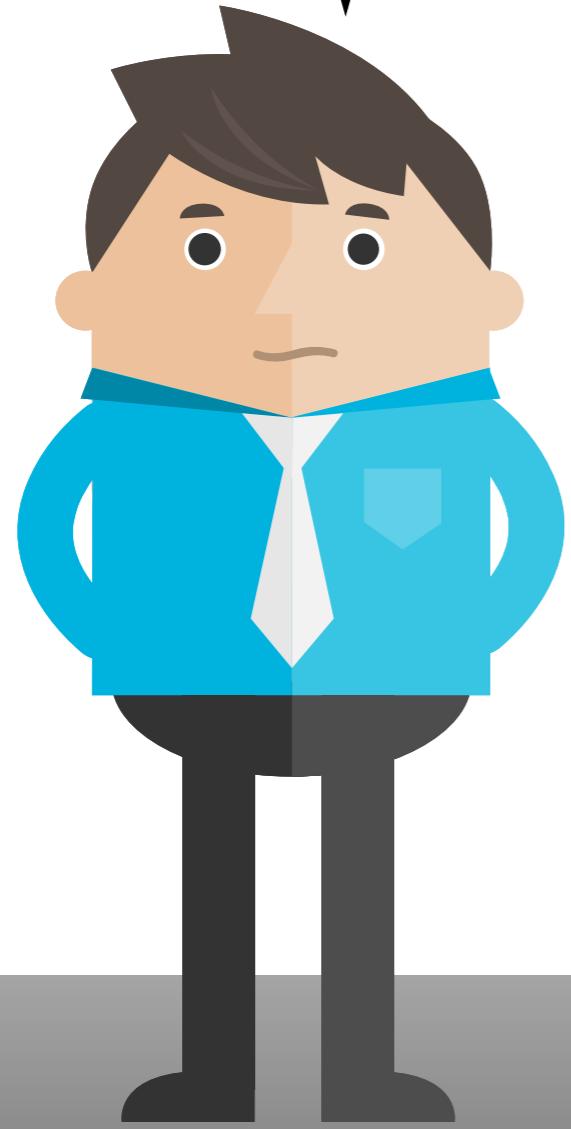
Find similar pictures



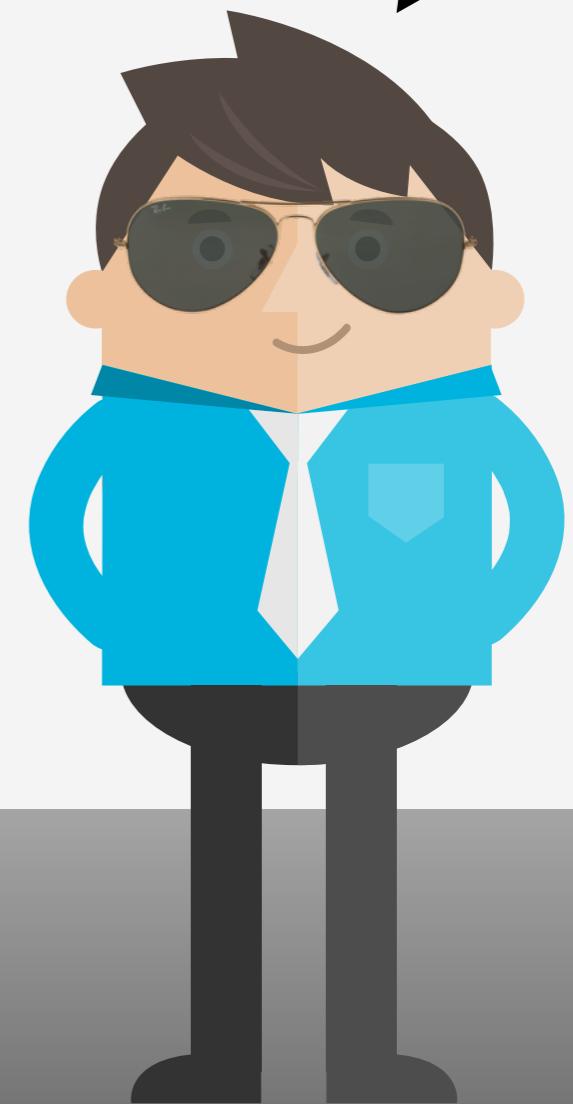
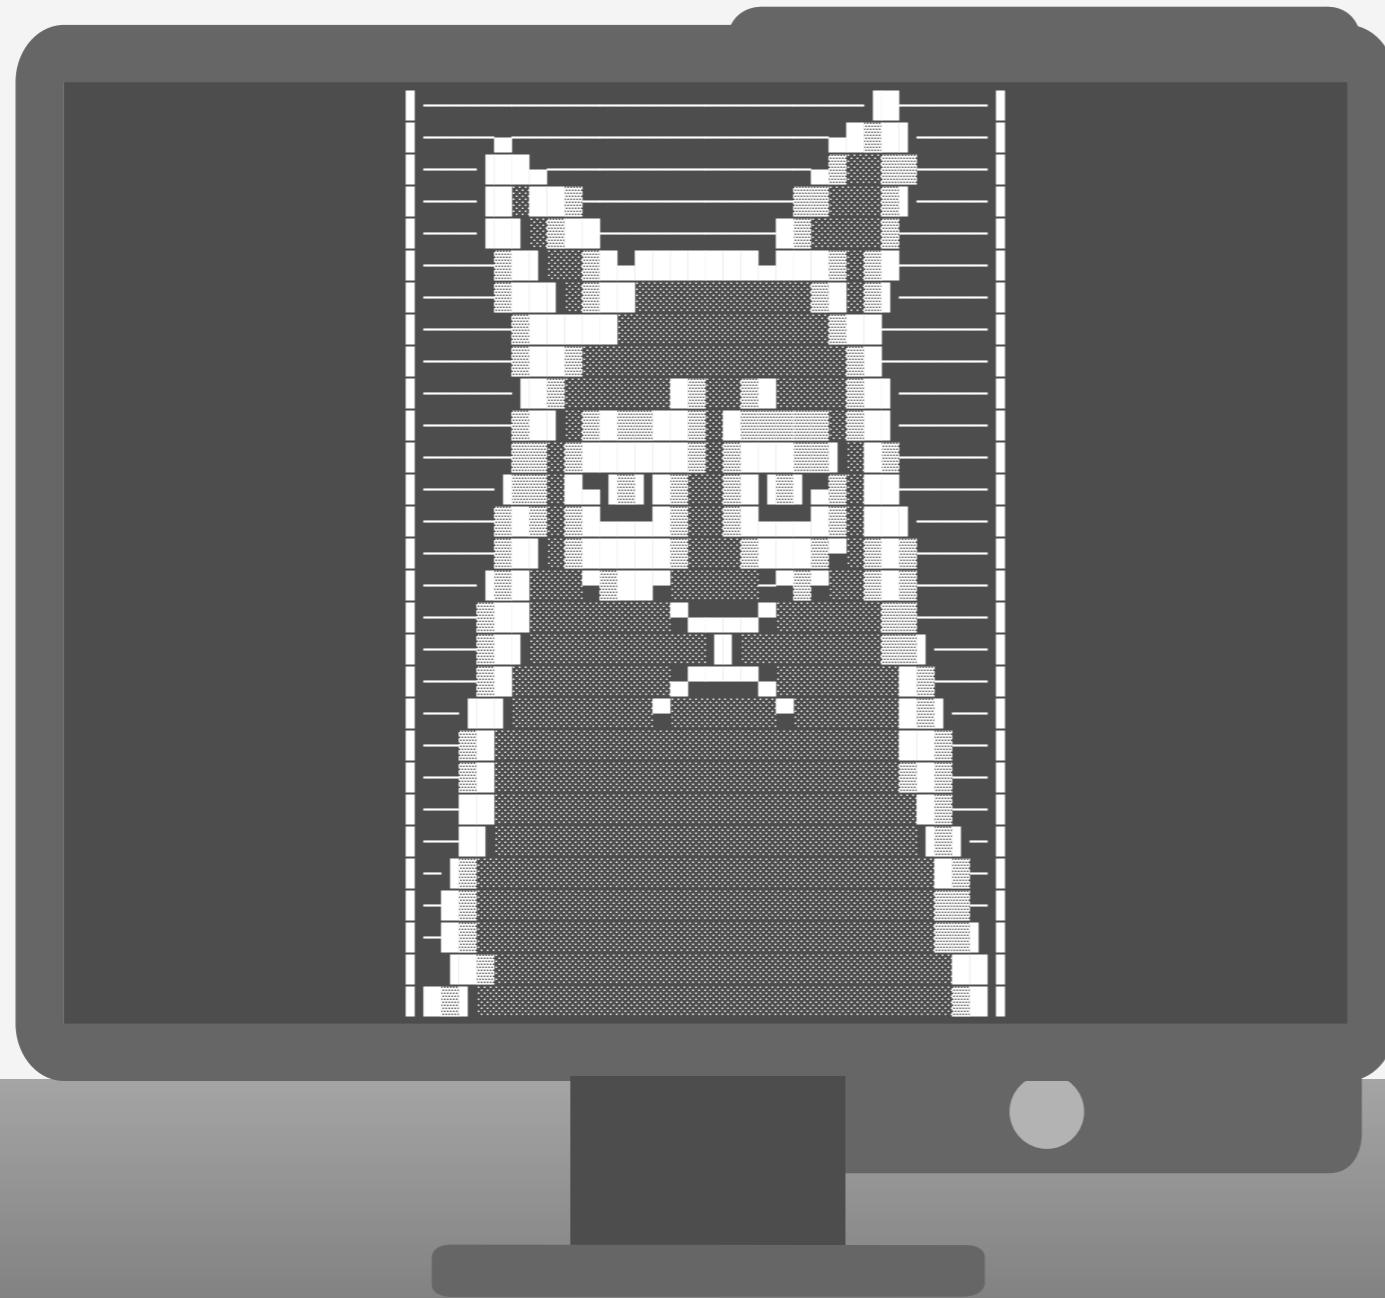
Ogle is totally disrupting
image search.



Ogle is too slow!



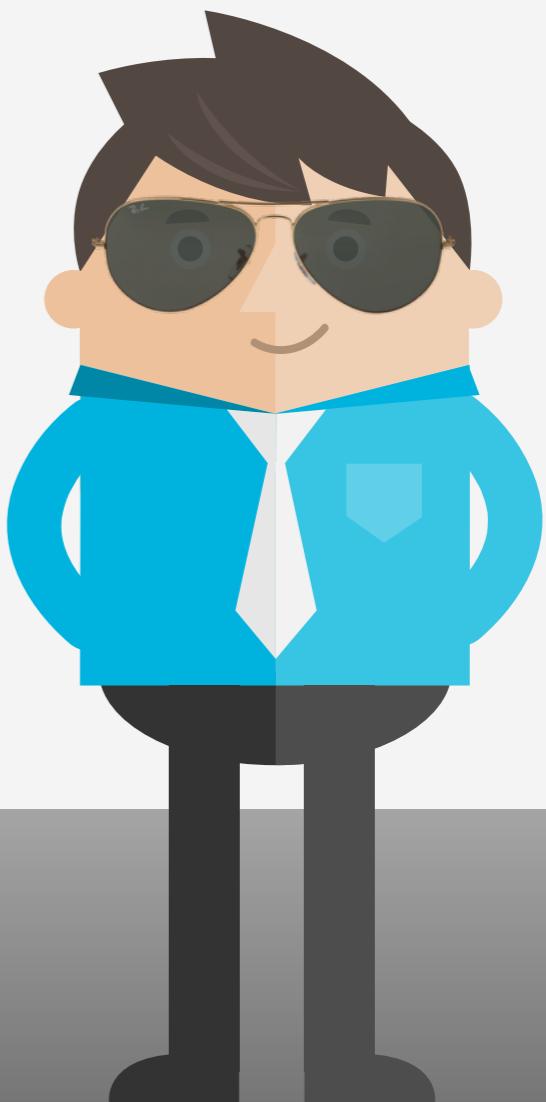
This is also Luke.
in the '80s



OGLE '84 1.0

Searching database...

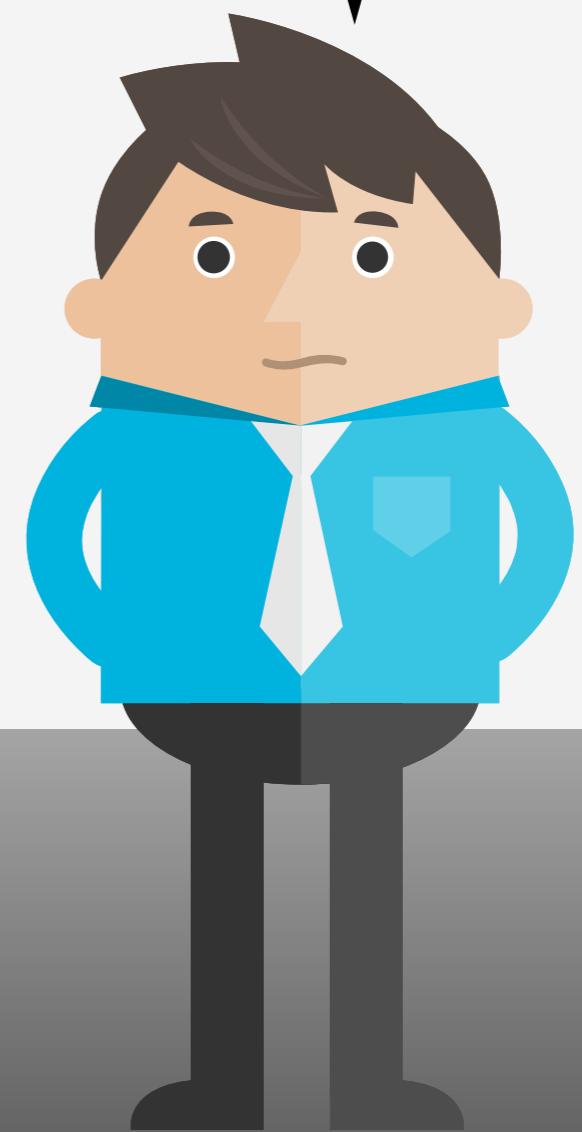
found 4 results:
catisgumpy.bmp
funisawful.bmp
catdoesnotlikefun.bmp
from_dad.bmp



OGLE'84 is too slow!

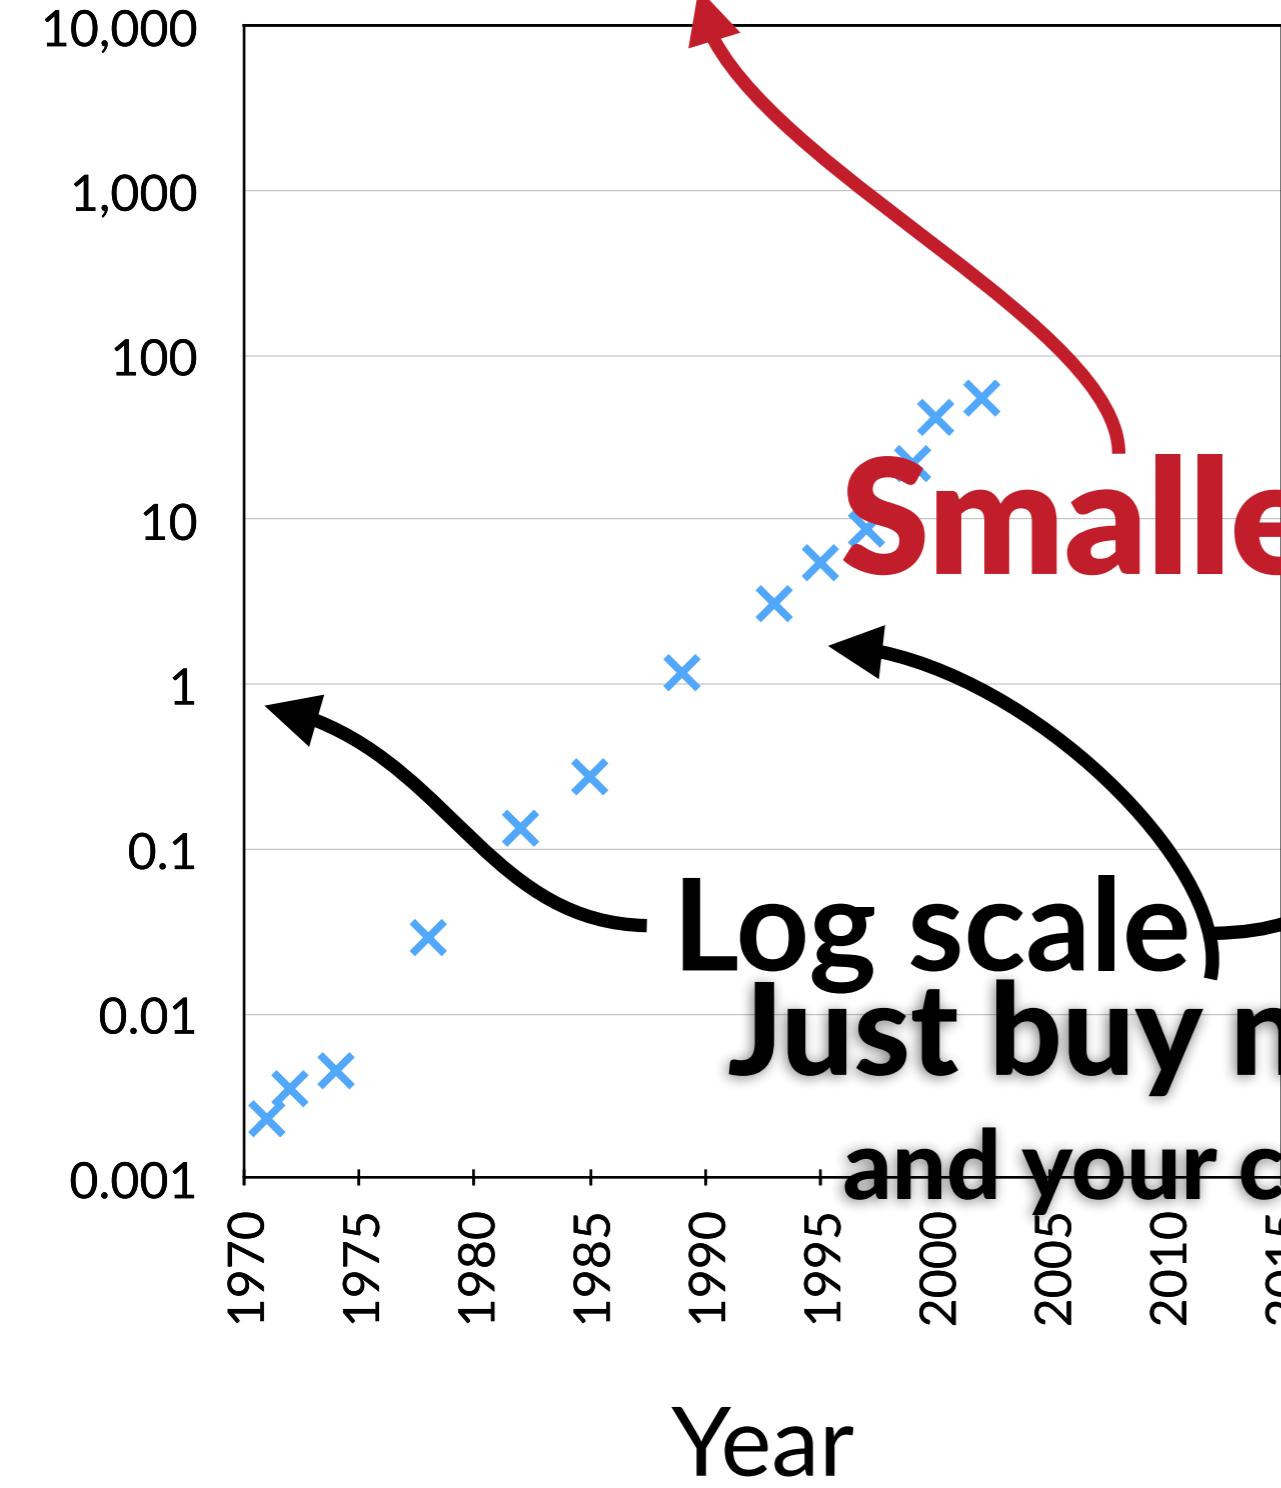


Ogle is too slow!

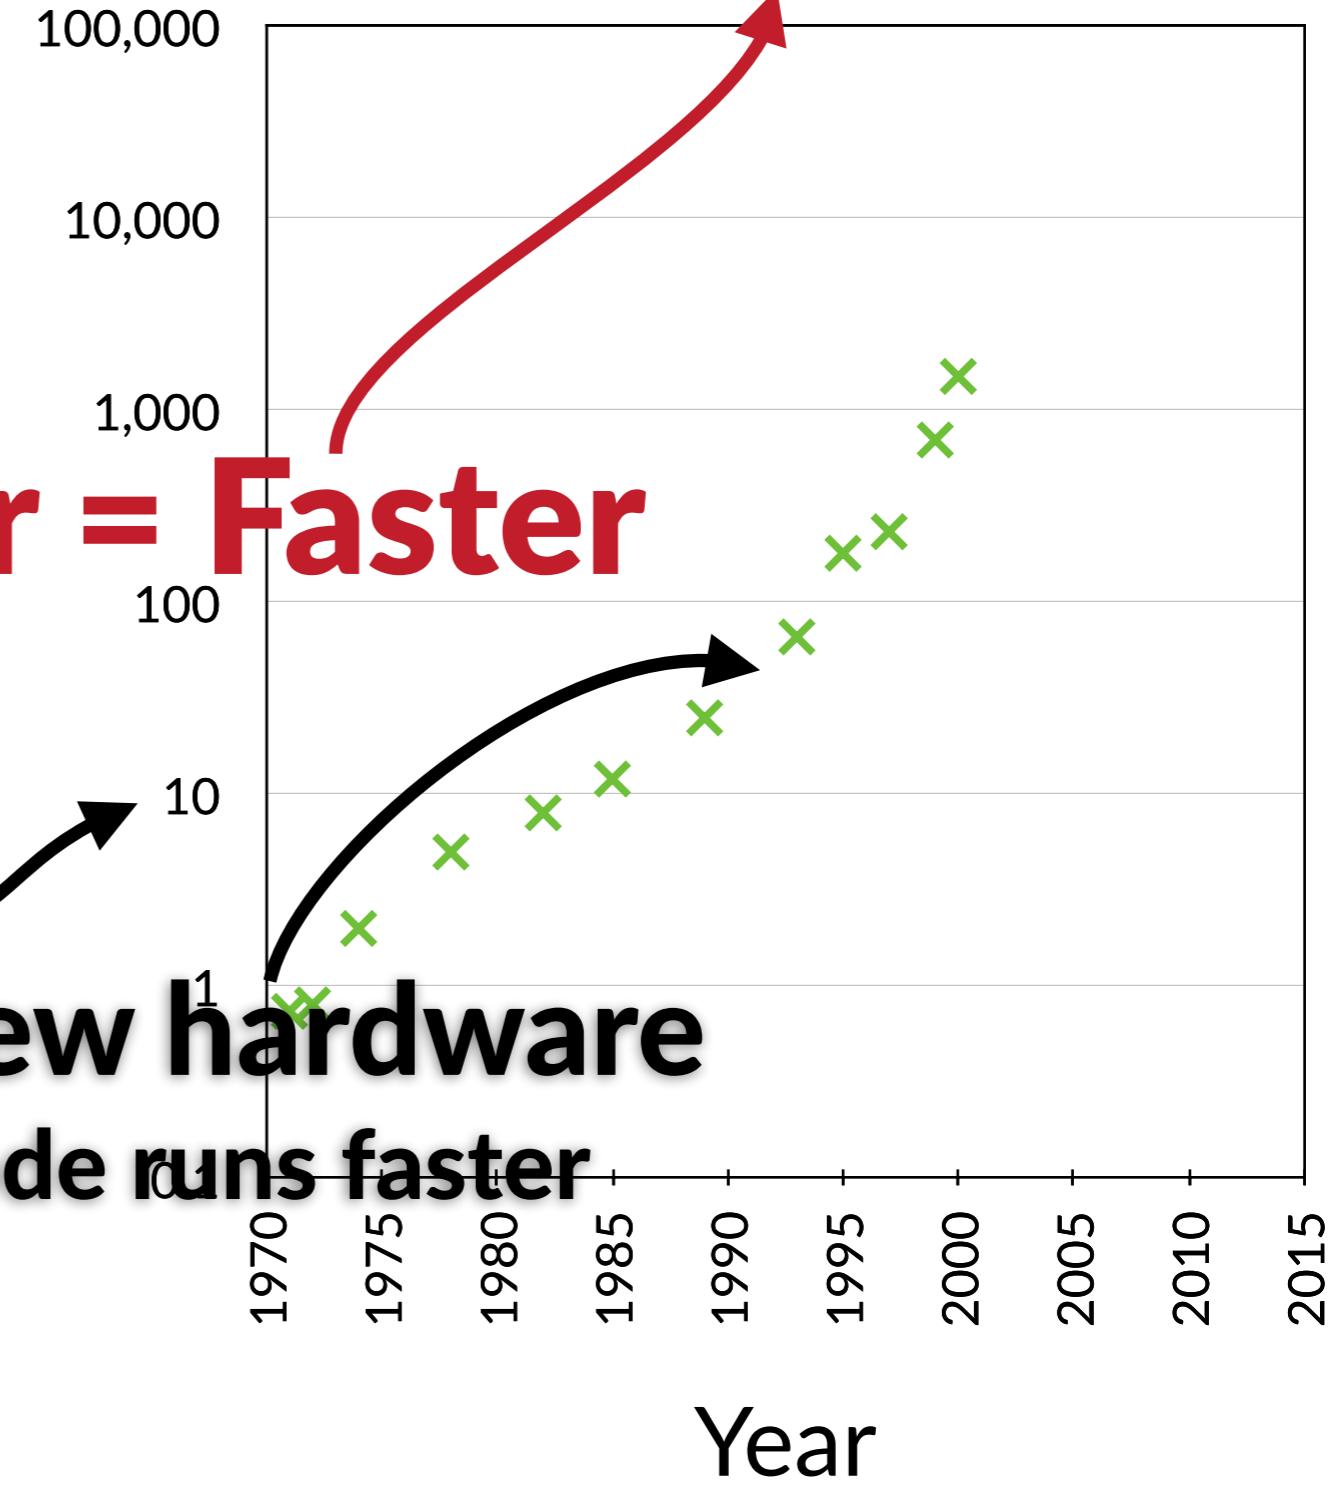


Performance used to be easy

Transistors (millions)



Clock Speed (MHz)



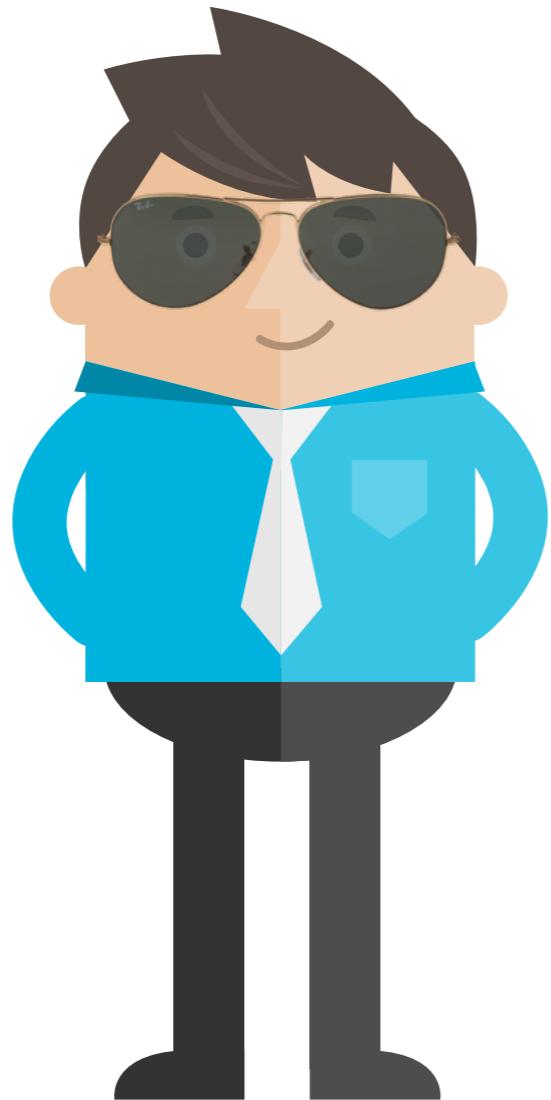
Smaller = Faster

Log scale

Just buy new hardware

and your code runs faster

Performance improvement in the '80s



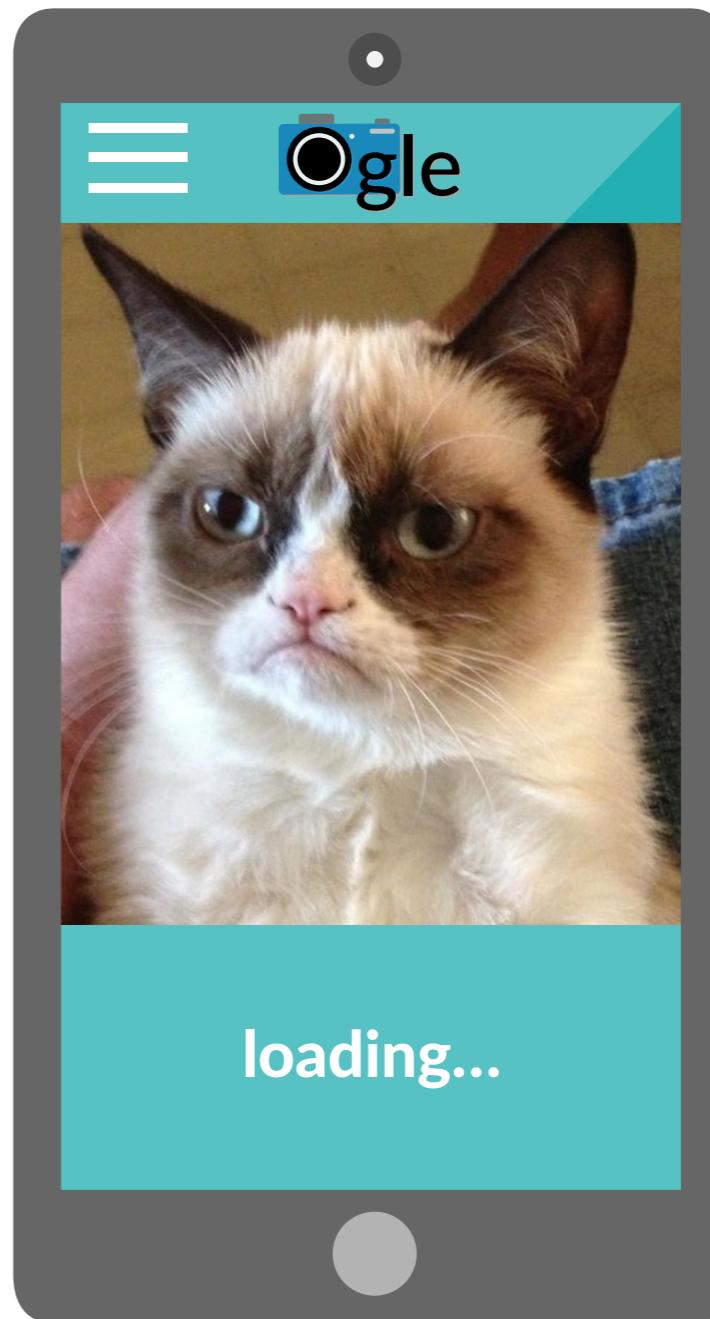
Performance improvement in the '80s



just chill and wait for faster hardware!

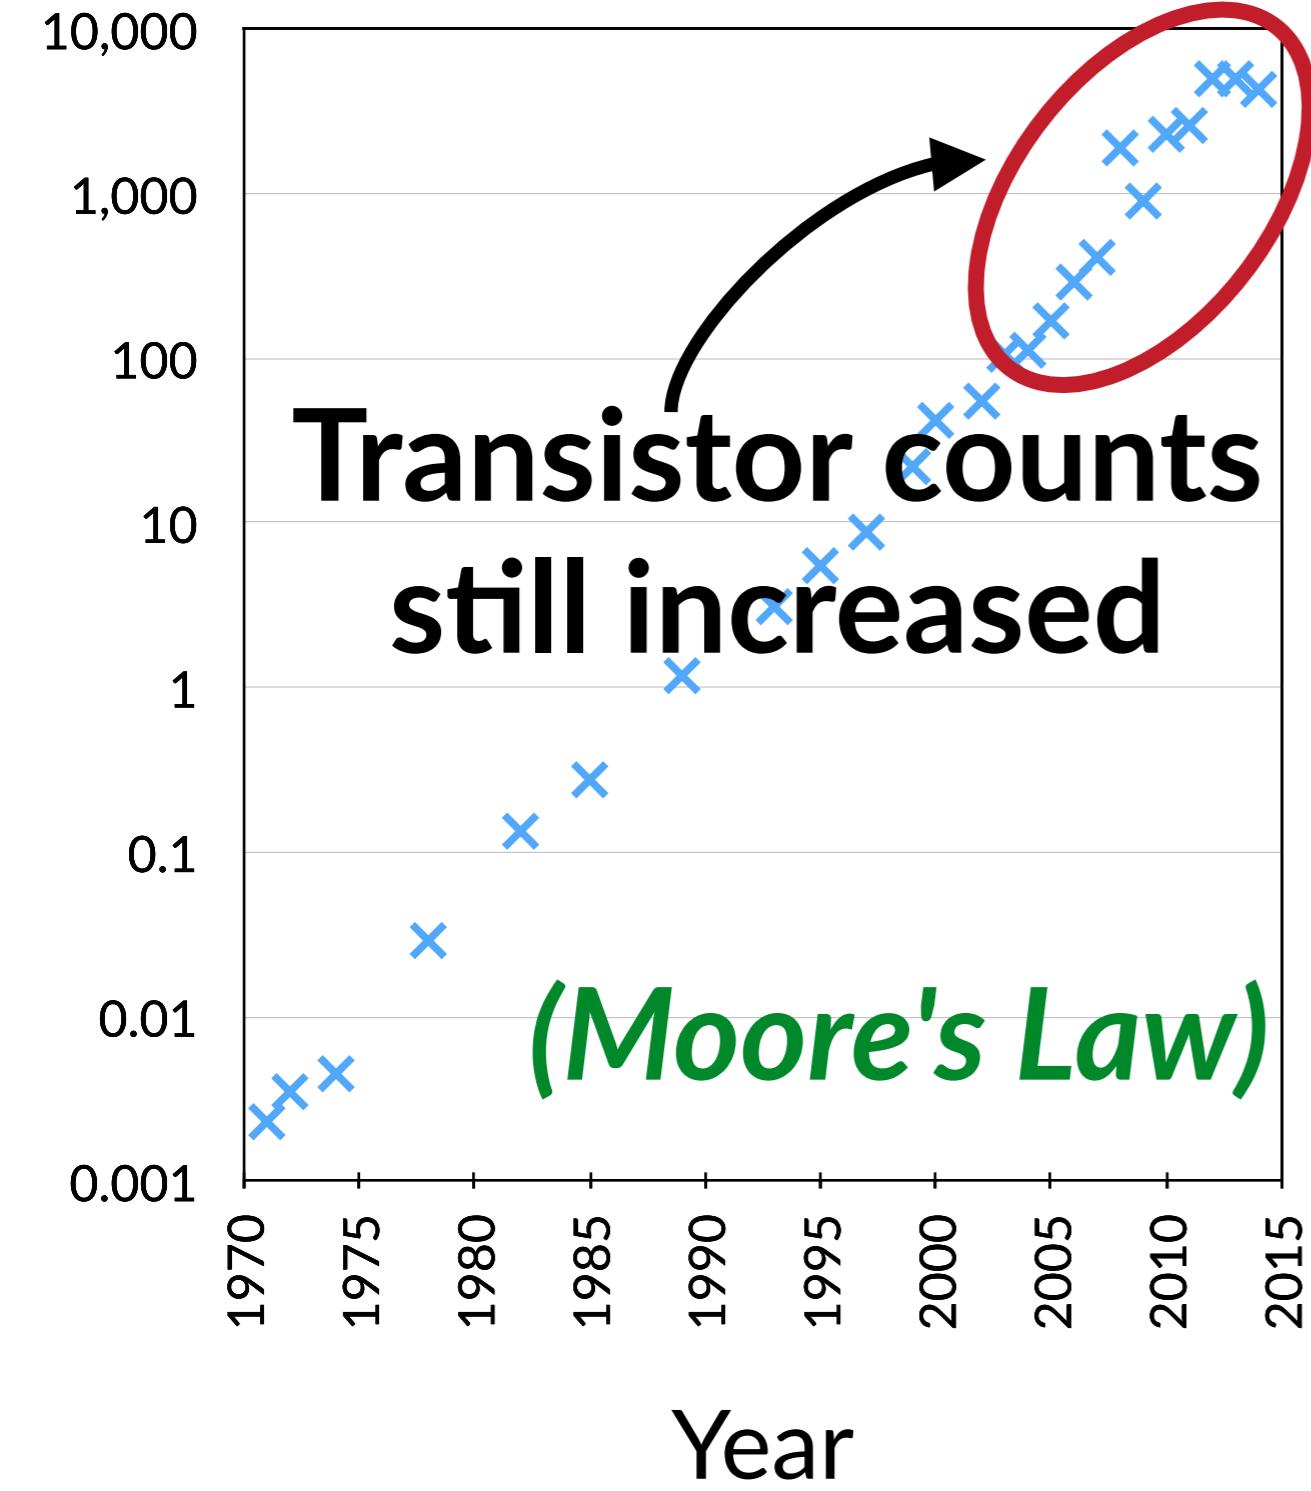
Back to the present...

No mojitos for me...

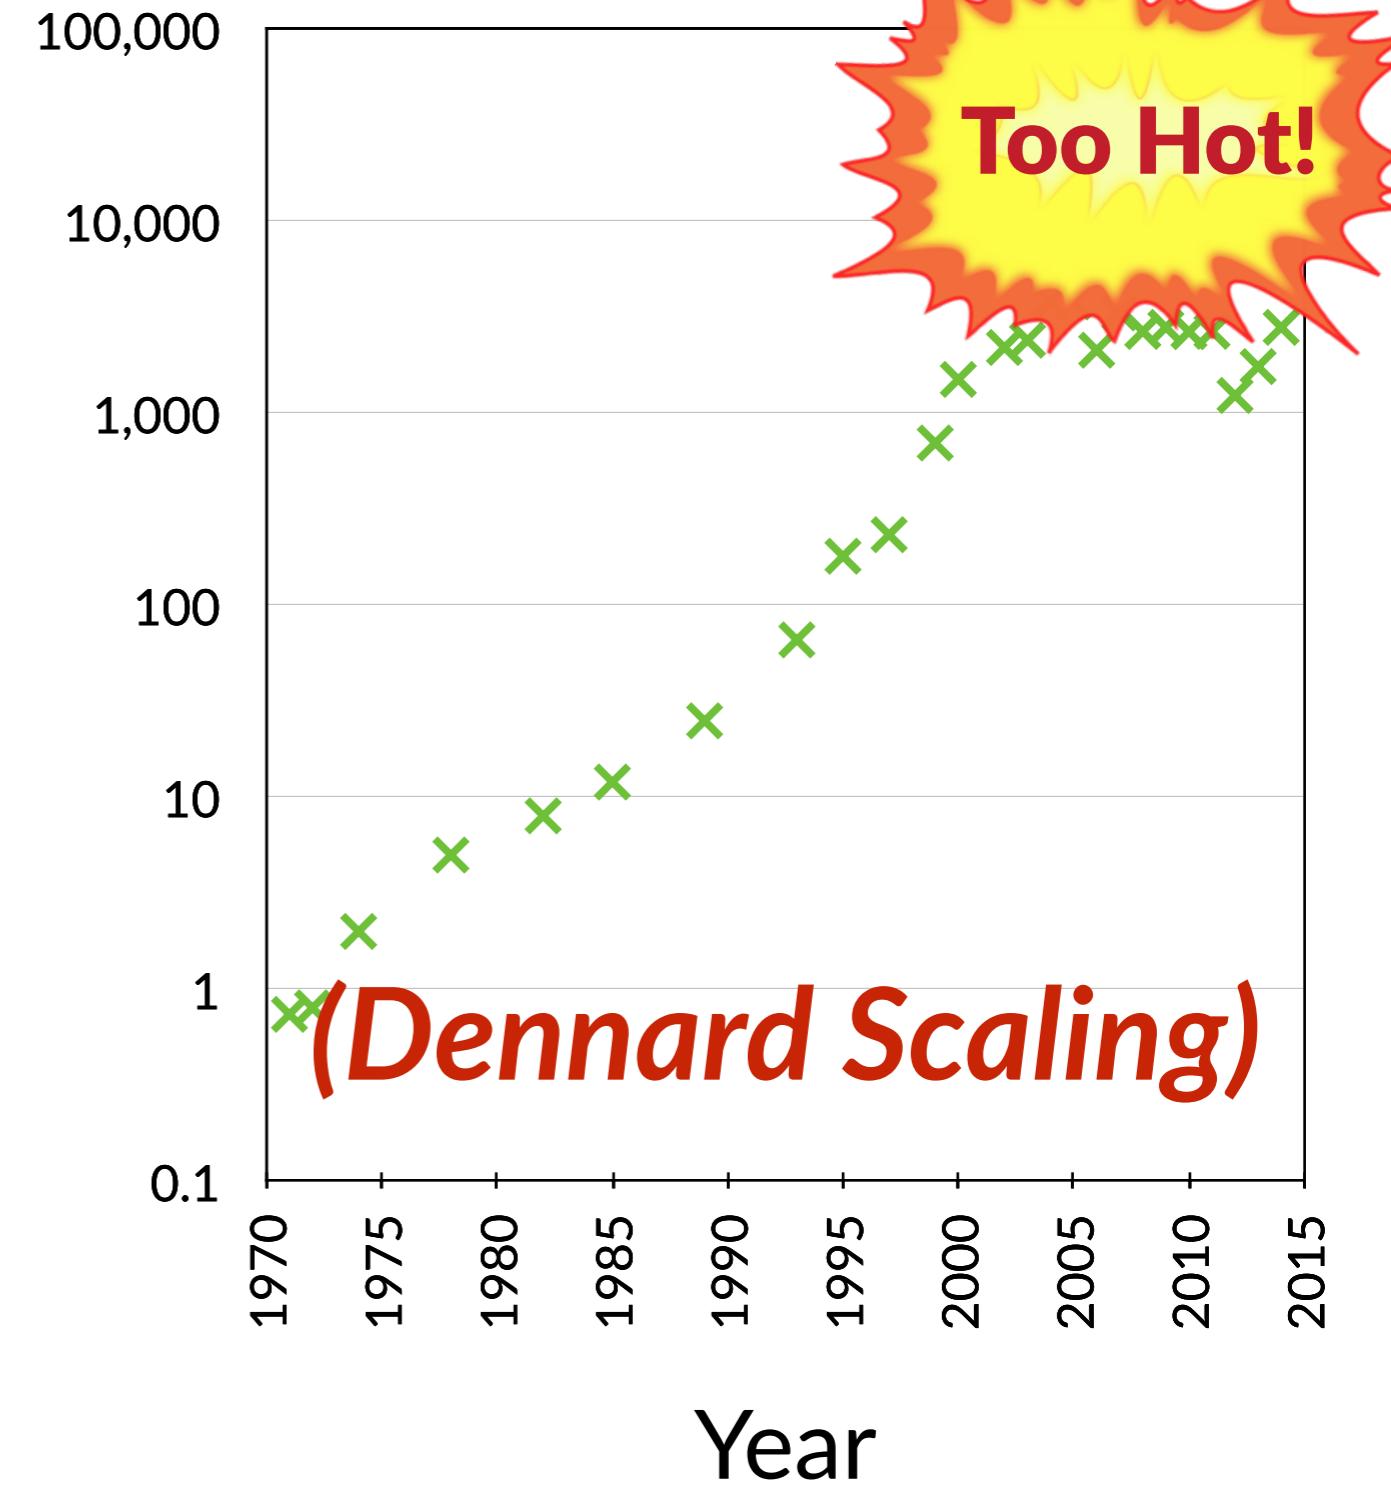


Performance not easy anymore

Transistors (millions)



Clock Speed (MHz)



8:39 ↗

8:39 ↗

8:40 ↗

8:40 ↗

.

Updates
Dropbox
Today

What's new:

- Under-the-hood updates for performance

We release updates regularly, and we're always looking for ways to make things better. If you have any feedback or run into issues, please let us know in our forums. We're happy to help!


MyFitnessPal
Today

Whether you want to log your last 10k, your first run in weeks, c


Google Maps
GPS Navigation
Today

Thanks for using Google Maps! We've brought bug fixes and awesome



Today



Games



Apps



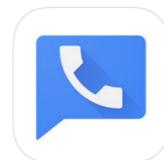
Today



Games



Apps

Updates
Google Voice
Mar 17, 2018

- Support for Smart Invert Colors
- Bug fixes and performance improvements

Note: This release no longer sup


Signal - Private Messenger
Mar 17, 2018

- You can enable an optional Recovery PIN that will be required in order


JetBlue
Mar 17, 2018

- Updated Flight Tracker
- Updated flight search



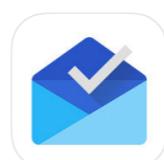
Today



Games



Apps

Updates
Inbox by Gmail
Mar 17, 2018

- Bug fixes and performance improvements


Gmail - Email by Google
Mar 17, 2018

- You can now change your profile picture and more! Just go to Account in settings.


Expedia: Hotels, Flights & Car
Mar 17, 2018

Paper is so 2017. Ditch the physical paperless with the Expedia App. Get all your travel details in one convenient place.


Bank of America Mobile Banking

Today

Games



Apps


Yelp: Discover Local Favorites
Mar 15, 2018
OPEN

We've decided to start referring to "performance improvements" as "app engine calibrations" because it sounds cooler. Therefore: For this version, we've further calibrated the app engine. See what we mean? Also, we fixed a bunch of bugs.


Facebook
Mar 15, 2018
OPEN

Thanks for using Facebook! To make our app better for you, we bring updates to the App Store regularly.

Every update of our Facebook app includes improvements for speed and reliability. As new features become available, we'll highlight those for you in the app.


WhatsApp Messenger
OPEN

Today

Games



Apps

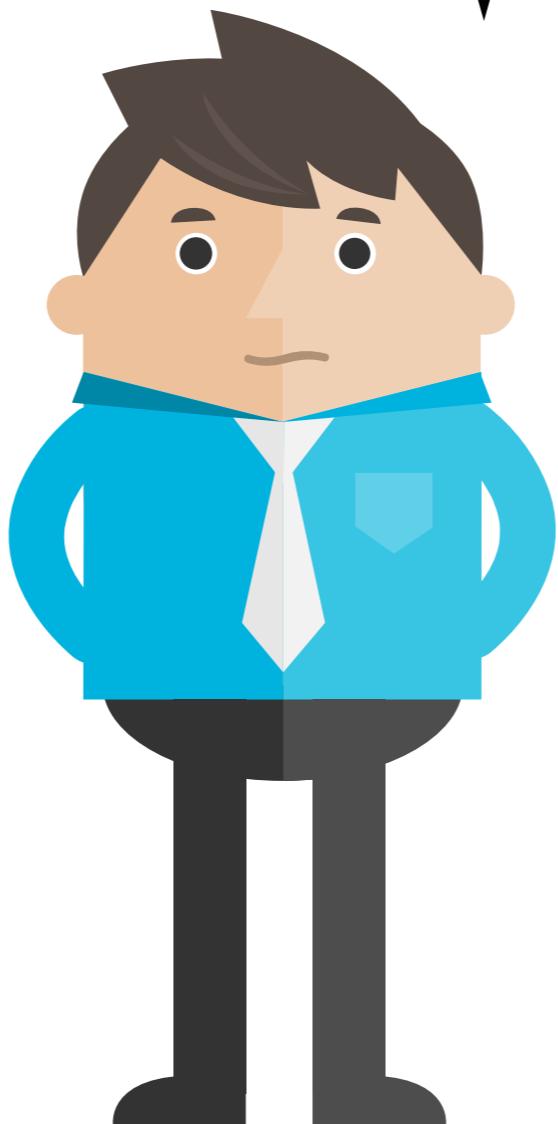


Updates

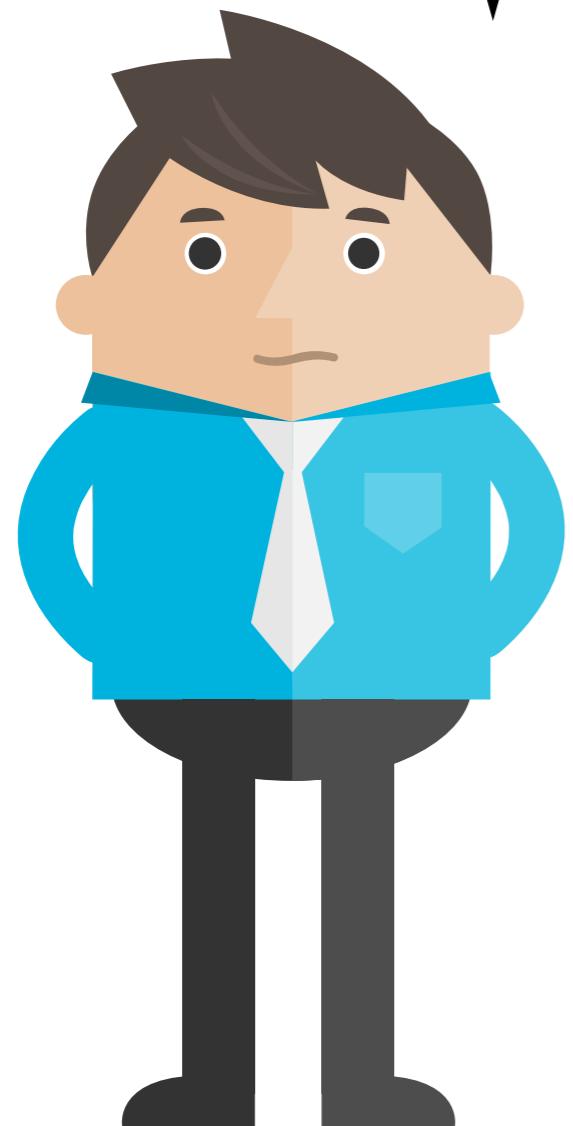


Search

Why is this so hard?



Why is this so hard?



Performance Analysis

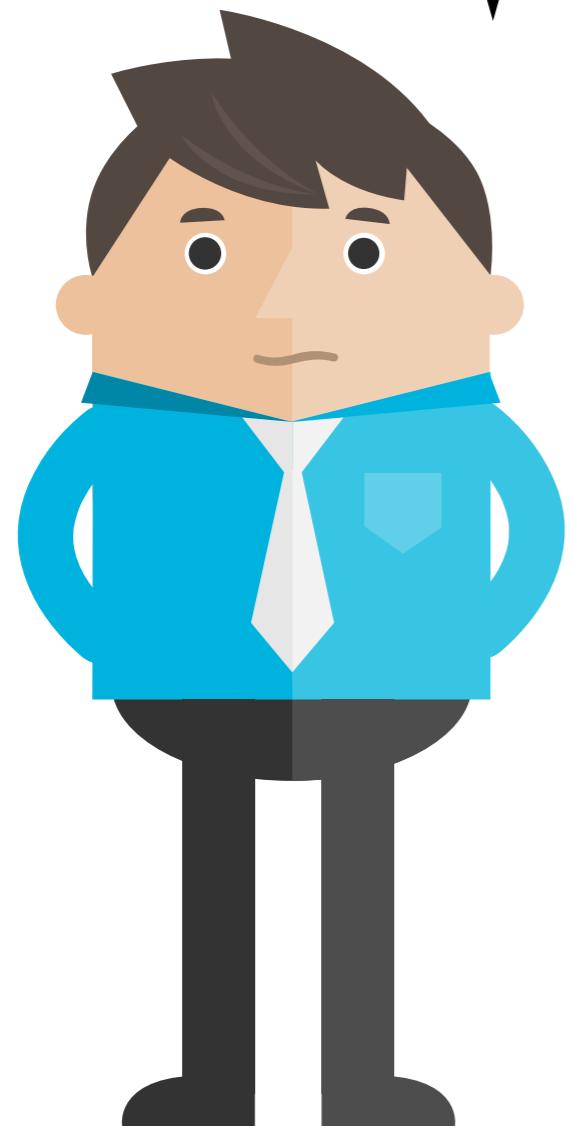
how to do it right

Performance Profiling

how to do it better



Why is this so hard?

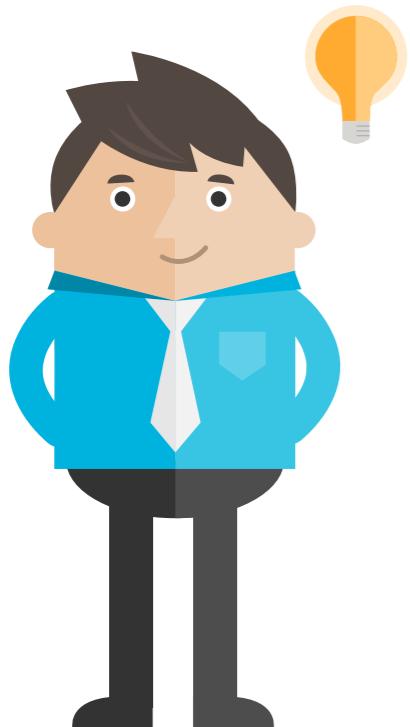
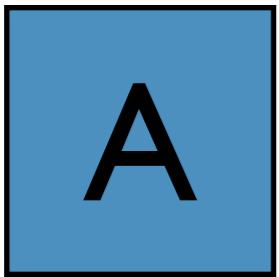


Performance Analysis

how to do it right



Typical performance evaluation



```
int main(int argc, char **argv) {
    topFrame = (void**)__builtin_frame_address(0);
    setHandler(Trap::TrapSignal, onTrap);
    setHandler(SIGALRM, onTimer);
    setHandler(SIGSEGV, onFault);
    for(Function* f: functions) {
        f->setTrap();
    }
    setTimer(interval);
    int r = stabilizer_main(argc, argv);
    return r;
}

void setTimer(int msec) {
    struct itimerval timer;

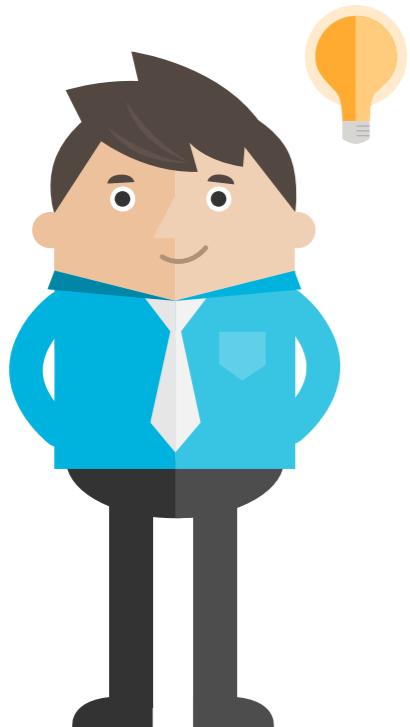
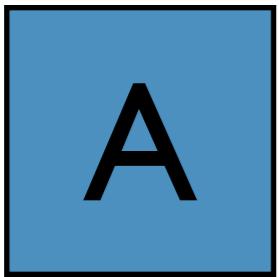
    timer.it_value.tv_sec = (msec - msec % 1000) / 1000;
    timer.it_value.tv_usec = 1000 * (msec % 1000);
    timer.it_interval.tv_sec = 0;
    timer.it_interval.tv_usec = 0;

    setitimer(ITIMER_REAL, &timer, 0);
}
```

```
static void flush_icache(void* begin, size_t size) {
    uintptr_t p = (uintptr_t)begin & ~15UL;
    for (size_t i = 0; i < size; i += 16) {
        asm("icbi 0,%0" : : "r"(p));
        p += 16;
    }
    asm("isync");
}
```

```
DataHeapType* getDataHeap() {
    static char buf[sizeof(DataHeapType)];
    static DataHeapType* _theDataHeap = new (buf) DataHeapType;
    return _theDataHeap;
}
```

Typical performance evaluation



```
int main(int argc, char **argv) {
    topFrame = (void**)__builtin_frame_address(0);
    setHandler(Trap::TrapSignal, onTrap);
    setHandler(SIGALRM, onTimer);
    setHandler(SIGSEGV, onFault);
    for(Function* f: functions) {
        f->setTrap();
    }
    setTimer(interval);
    int r = stabilizer_main(argc, argv);
    return r;
}

void setTimer(int msec) {
    struct itimerval timer;

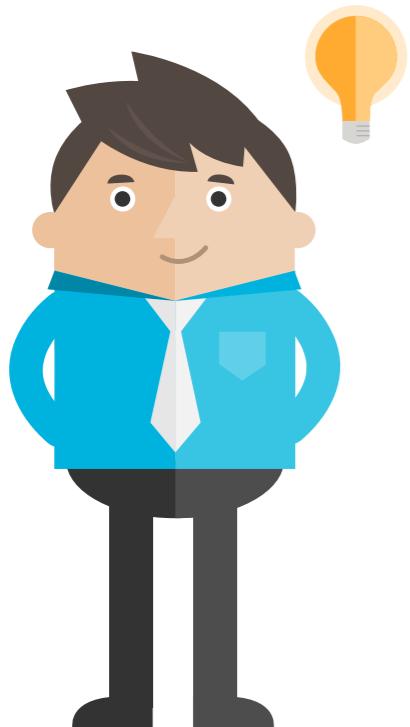
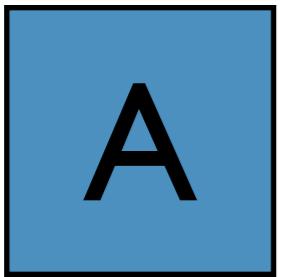
    timer.it_value.tv_sec = (msec - msec % 1000) / 1000;
    timer.it_value.tv_usec = 1000 * (msec % 1000);
    timer.it_interval.tv_sec = 0;
    timer.it_interval.tv_usec = 0;

    setitimer(ITIMER_REAL, &timer, 0);
}
```

```
static void flush_icache(void* begin, size_t size) {
    uintptr_t p = (uintptr_t)begin & ~15UL;
    size_t meaning_of_life=42;
    for (size_t i = 0; i < size; i += 32) {
        asm("icbi 0,%0" : : "r"(p));
        p += 32;
    }
    for (size_t i = 16; i < size; i += 32) {
        asm("icbi 0,%0" : : "r"(p));
        p += 32;
    }
    asm("isync");
}
```

```
DataHeapType* getDataHeap() {
    static char buf[sizeof(DataHeapType)];
    static DataHeapType* _theDataHeap = new (buf) DataHeapType;
    return _theDataHeap;
}
```

Typical performance evaluation



```
int main(int argc, char **argv) {
    topFrame = (void**)__builtin_frame_address(0);
    setHandler(Trap::TrapSignal, onTrap);
    setHandler(SIGALRM, onTimer);
    setHandler(SIGSEGV, onFault);
    for(Function* f: functions) {
        f->setTrap();
    }
    setTimer(interval);
    int r = stabilizer_main(argc, argv);
    return r;
}

void setTimer(int msec) {
    struct itimerval timer;

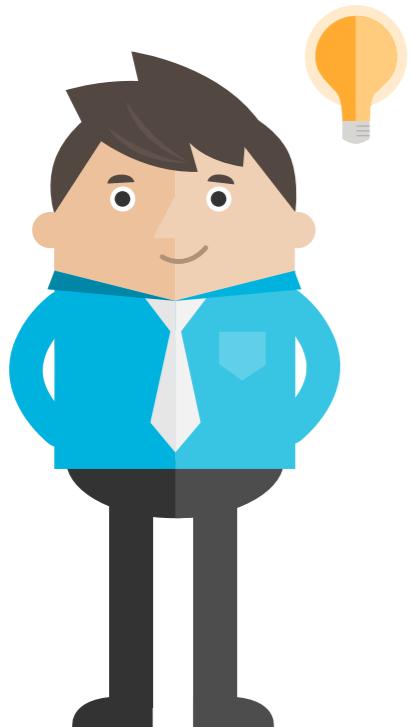
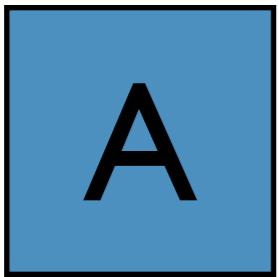
    timer.it_value.tv_sec = (msec - msec % 1000) / 1000;
    timer.it_value.tv_usec = 1000 * (msec % 1000);
    timer.it_interval.tv_sec = 0;
    timer.it_interval.tv_usec = 0;

    setitimer(ITIMER_REAL, &timer, 0);
}
```

```
static void flush_icache(void* begin, size_t size) {
    uintptr_t p = (uintptr_t)begin & ~15UL;
    size_t meaning_of_life=42;
    for (size_t i = 0; i < size; i += 32) {
        asm("icbi 0,%0" : : "r"(p));
        p += 32;
    }
    for (size_t i = 16; i < size; i += 32) {
        asm("icbi 0,%0" : : "r"(p));
        p += 32;
    }
    asm("isync");
}
```

```
DataHeapType* getDataHeap() {
    static char buf[sizeof(DataHeapType)];
    static DataHeapType* _theDataHeap = new (buf) DataHeapType;
    return _theDataHeap;
}
```

Typical performance evaluation



```
DataHeapType* getDataHeap() {
    static char buf[sizeof(DataHeapType)];
    static DataHeapType* _theDataHeap = new (buf) DataHeapType;
    return _theDataHeap;
}
```

```
int main(int argc, char **argv) {
    topFrame = (void**)__builtin_frame_address(0);
    setHandler(Trap::TrapSignal, onTrap);
    setHandler(SIGALRM, onTimer);
    setHandler(SIGSEGV, onFault);
    for(Function* f: functions) {
        f->setTrap();
    }
    setTimer(interval);
    int r = stabilizer_main(argc, argv);
    return r;
}

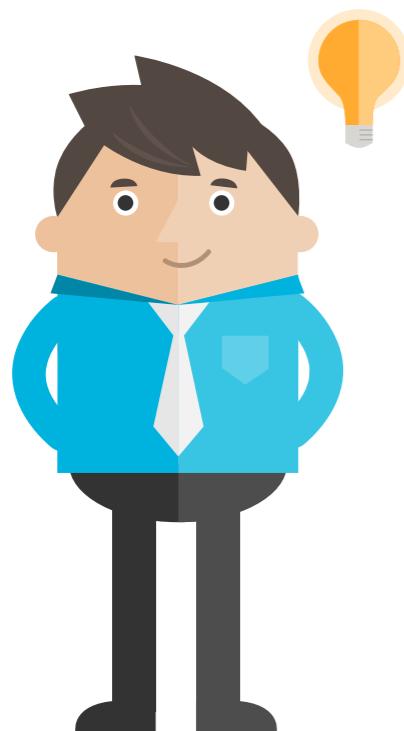
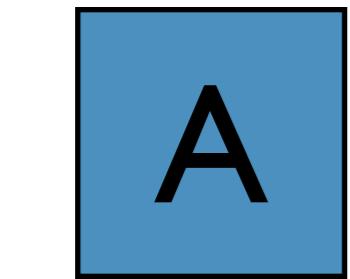
void setTimer(int msec) {
    struct itimerval timer;

    timer.it_value.tv_sec = (msec - msec % 1000) / 1000;
    timer.it_value.tv_usec = 1000 * (msec % 1000);
    timer.it_interval.tv_sec = 0;
    timer.it_interval.tv_usec = 0;

    setitimer(ITIMER_REAL, &timer, 0);
}
```

```
static void flush_icache(void* begin, size_t size) {
    uintptr_t p = (uintptr_t)begin & ~15UL;
    size_t meaning_of_life=42;
    for (size_t i = 0; i < size; i += 32) {
        asm("icbi 0,%0" : : "r"(p));
        p += 32;
    }
    for (size_t i = 16; i < size; i += 32) {
        asm("icbi 0,%0" : : "r"(p));
        p += 32;
    }
    asm("isync");
}
```

Typical performance evaluation



```
DataHeapType* getDataHeap() {
    static char buf[sizeof(DataHeapType)];
    static DataHeapType* _theDataHeap = new (buf) DataHeapType;
    return _theDataHeap;
}
```

```
int main(int argc, char **argv) {
    topFrame = (void**)__builtin_frame_address(0);
    setHandler(Trap::TrapSignal, onTrap);
    setHandler(SIGALRM, onTimer);
    setHandler(SIGSEGV, onFault);
    for(Function* f: functions) {
        f->setTrap();
    }
    setTimer(interval);
    int r = stabilizer_main(argc, argv);
    return r;
}
```

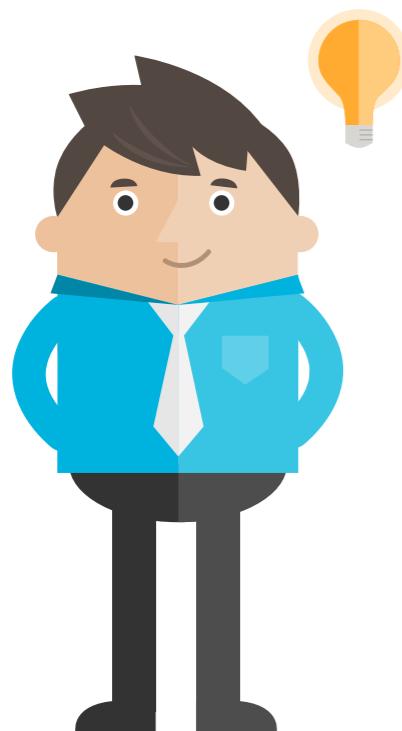
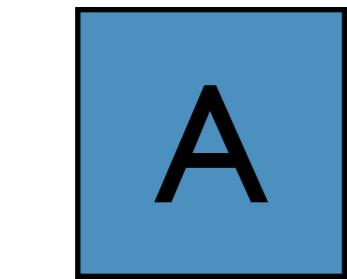
```
static void flush_icache(void* begin, size_t size) {
    uintptr_t p = (uintptr_t)begin & ~15UL;
    size_t meaning_of_life=42;
    for (size_t i = 0; i < size; i += 32) {
        asm("icbi 0,%0" : : "r"(p));
        p += 32;
    }
    for (size_t i = 16; i < size; i += 32) {
        asm("icbi 0,%0" : : "r"(p));
        p += 32;
    }
    asm("isync");
}
```

```
void setTimer(int msec) {
    struct itimerval timer;

    timer.it_value.tv_sec = (msec - msec % 1000) / 1000;
    timer.it_value.tv_usec = 1000 * (msec % 1000);
    timer.it_interval.tv_sec = 0;
    timer.it_interval.tv_usec = 0;

    setitimer(ITIMER_REAL, &timer, 0);
}
```

Typical performance evaluation



```
DataHeapType* getDataHeap() {
    static char buf[sizeof(DataHeapType)];
    static DataHeapType* _theDataHeap = new (buf) DataHeapType;
    return _theDataHeap;
}
```

```
int main(int argc, char **argv) {
    topFrame = (void**)__builtin_frame_address(0);
    setHandler(Trap::TrapSignal, onTrap);
    setHandler(SIGALRM, onTimer);
    setHandler(SIGSEGV, onFault);
    for(Function* f: functions) {
        f->setTrap();
    }
    setTimer(interval);
    int r = stabilizer_main(argc, argv);
    return r;
}
```

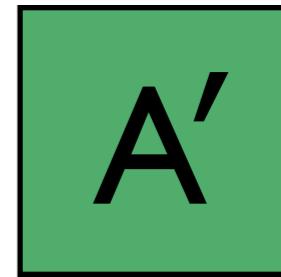
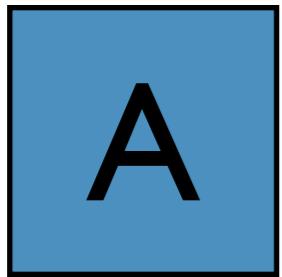
```
static void flush_icache(void* begin, size_t size) {
    uintptr_t p = (uintptr_t)begin & ~15UL;
    size_t meaning_of_life=42;
    for (size_t i = 0; i < size; i += 32) {
        asm("icbi 0,%0" : : "r"(p));
        p += 32;
    }
    for (size_t i = 16; i < size; i += 32) {
        asm("icbi 0,%0" : : "r"(p));
        p += 32;
    }
    asm("isync");
}
```

```
void setTimer(int msec) {
    struct itimerval timer;

    timer.it_value.tv_sec = (msec - msec % 1000) / 1000;
    timer.it_value.tv_usec = 1000 * (msec % 1000);
    timer.it_interval.tv_sec = 0;
    timer.it_interval.tv_usec = 0;

    setitimer(ITIMER_REAL, &timer, 0);
}
```

Typical performance evaluation

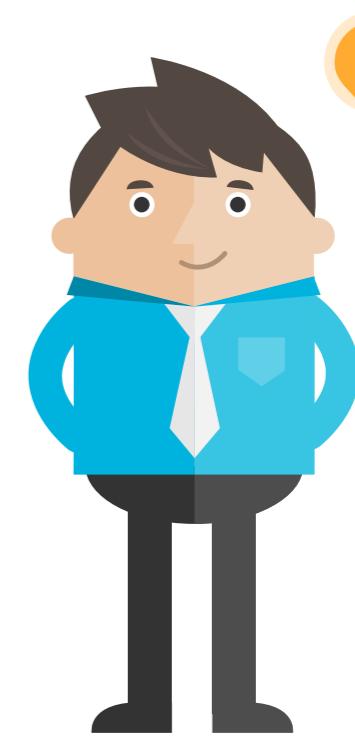


```
DataHeapType* getDataHeap() {  
    static char buf[sizeof(DataHeapType)];  
    static DataHeapType* _theDataHeap = new (buf) DataHeapType;  
    return _theDataHeap;  
}
```

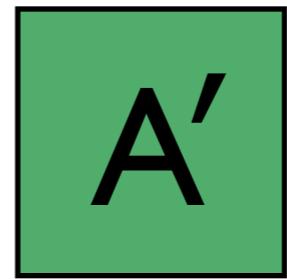
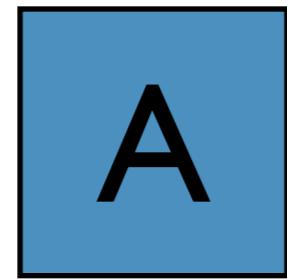
```
int main(int argc, char **argv) {  
    topFrame = (void**)__builtin_frame_address(0);  
    setHandler(Trap::TrapSignal, onTrap);  
    setHandler(SIGALRM, onTimer);  
    setHandler(SIGSEGV, onFault);  
    for(Function* f: functions) {  
        f->setTrap();  
    }  
    setTimer(interval);  
    int r = stabilizer_main(argc, argv);  
    return r;  
}
```

```
static void flush_icache(void* begin, size_t size) {  
    uintptr_t p = (uintptr_t)begin & ~15UL;  
    size_t meaning_of_life=42;  
    for (size_t i = 0; i < size; i += 32) {  
        asm("icbi 0,%0" : : "r"(p));  
        p += 32;  
    }  
    for (size_t i = 16; i < size; i += 32) {  
        asm("icbi 0,%0" : : "r"(p));  
        p += 32;  
    }  
    asm("isync");  
}
```

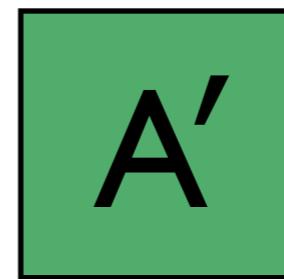
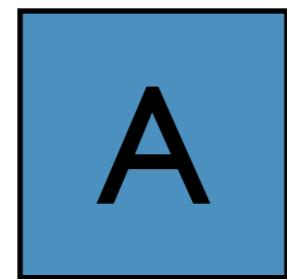
```
void setTimer(int msec) {  
    struct itimerval timer;  
  
    timer.it_value.tv_sec = (msec - msec % 1000) / 1000;  
    timer.it_value.tv_usec = 1000 * (msec % 1000);  
    timer.it_interval.tv_sec = 0;  
    timer.it_interval.tv_usec = 0;  
  
    setitimer(ITIMER_REAL, &timer, 0);  
}
```



Which is faster?

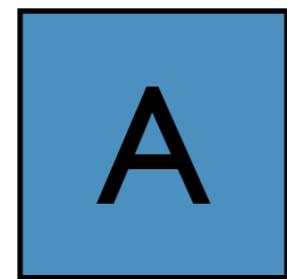


Which is faster?

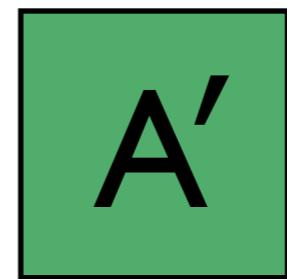


90s

Which is faster?



90s



87.5s

Which is faster?

A

90s

A'

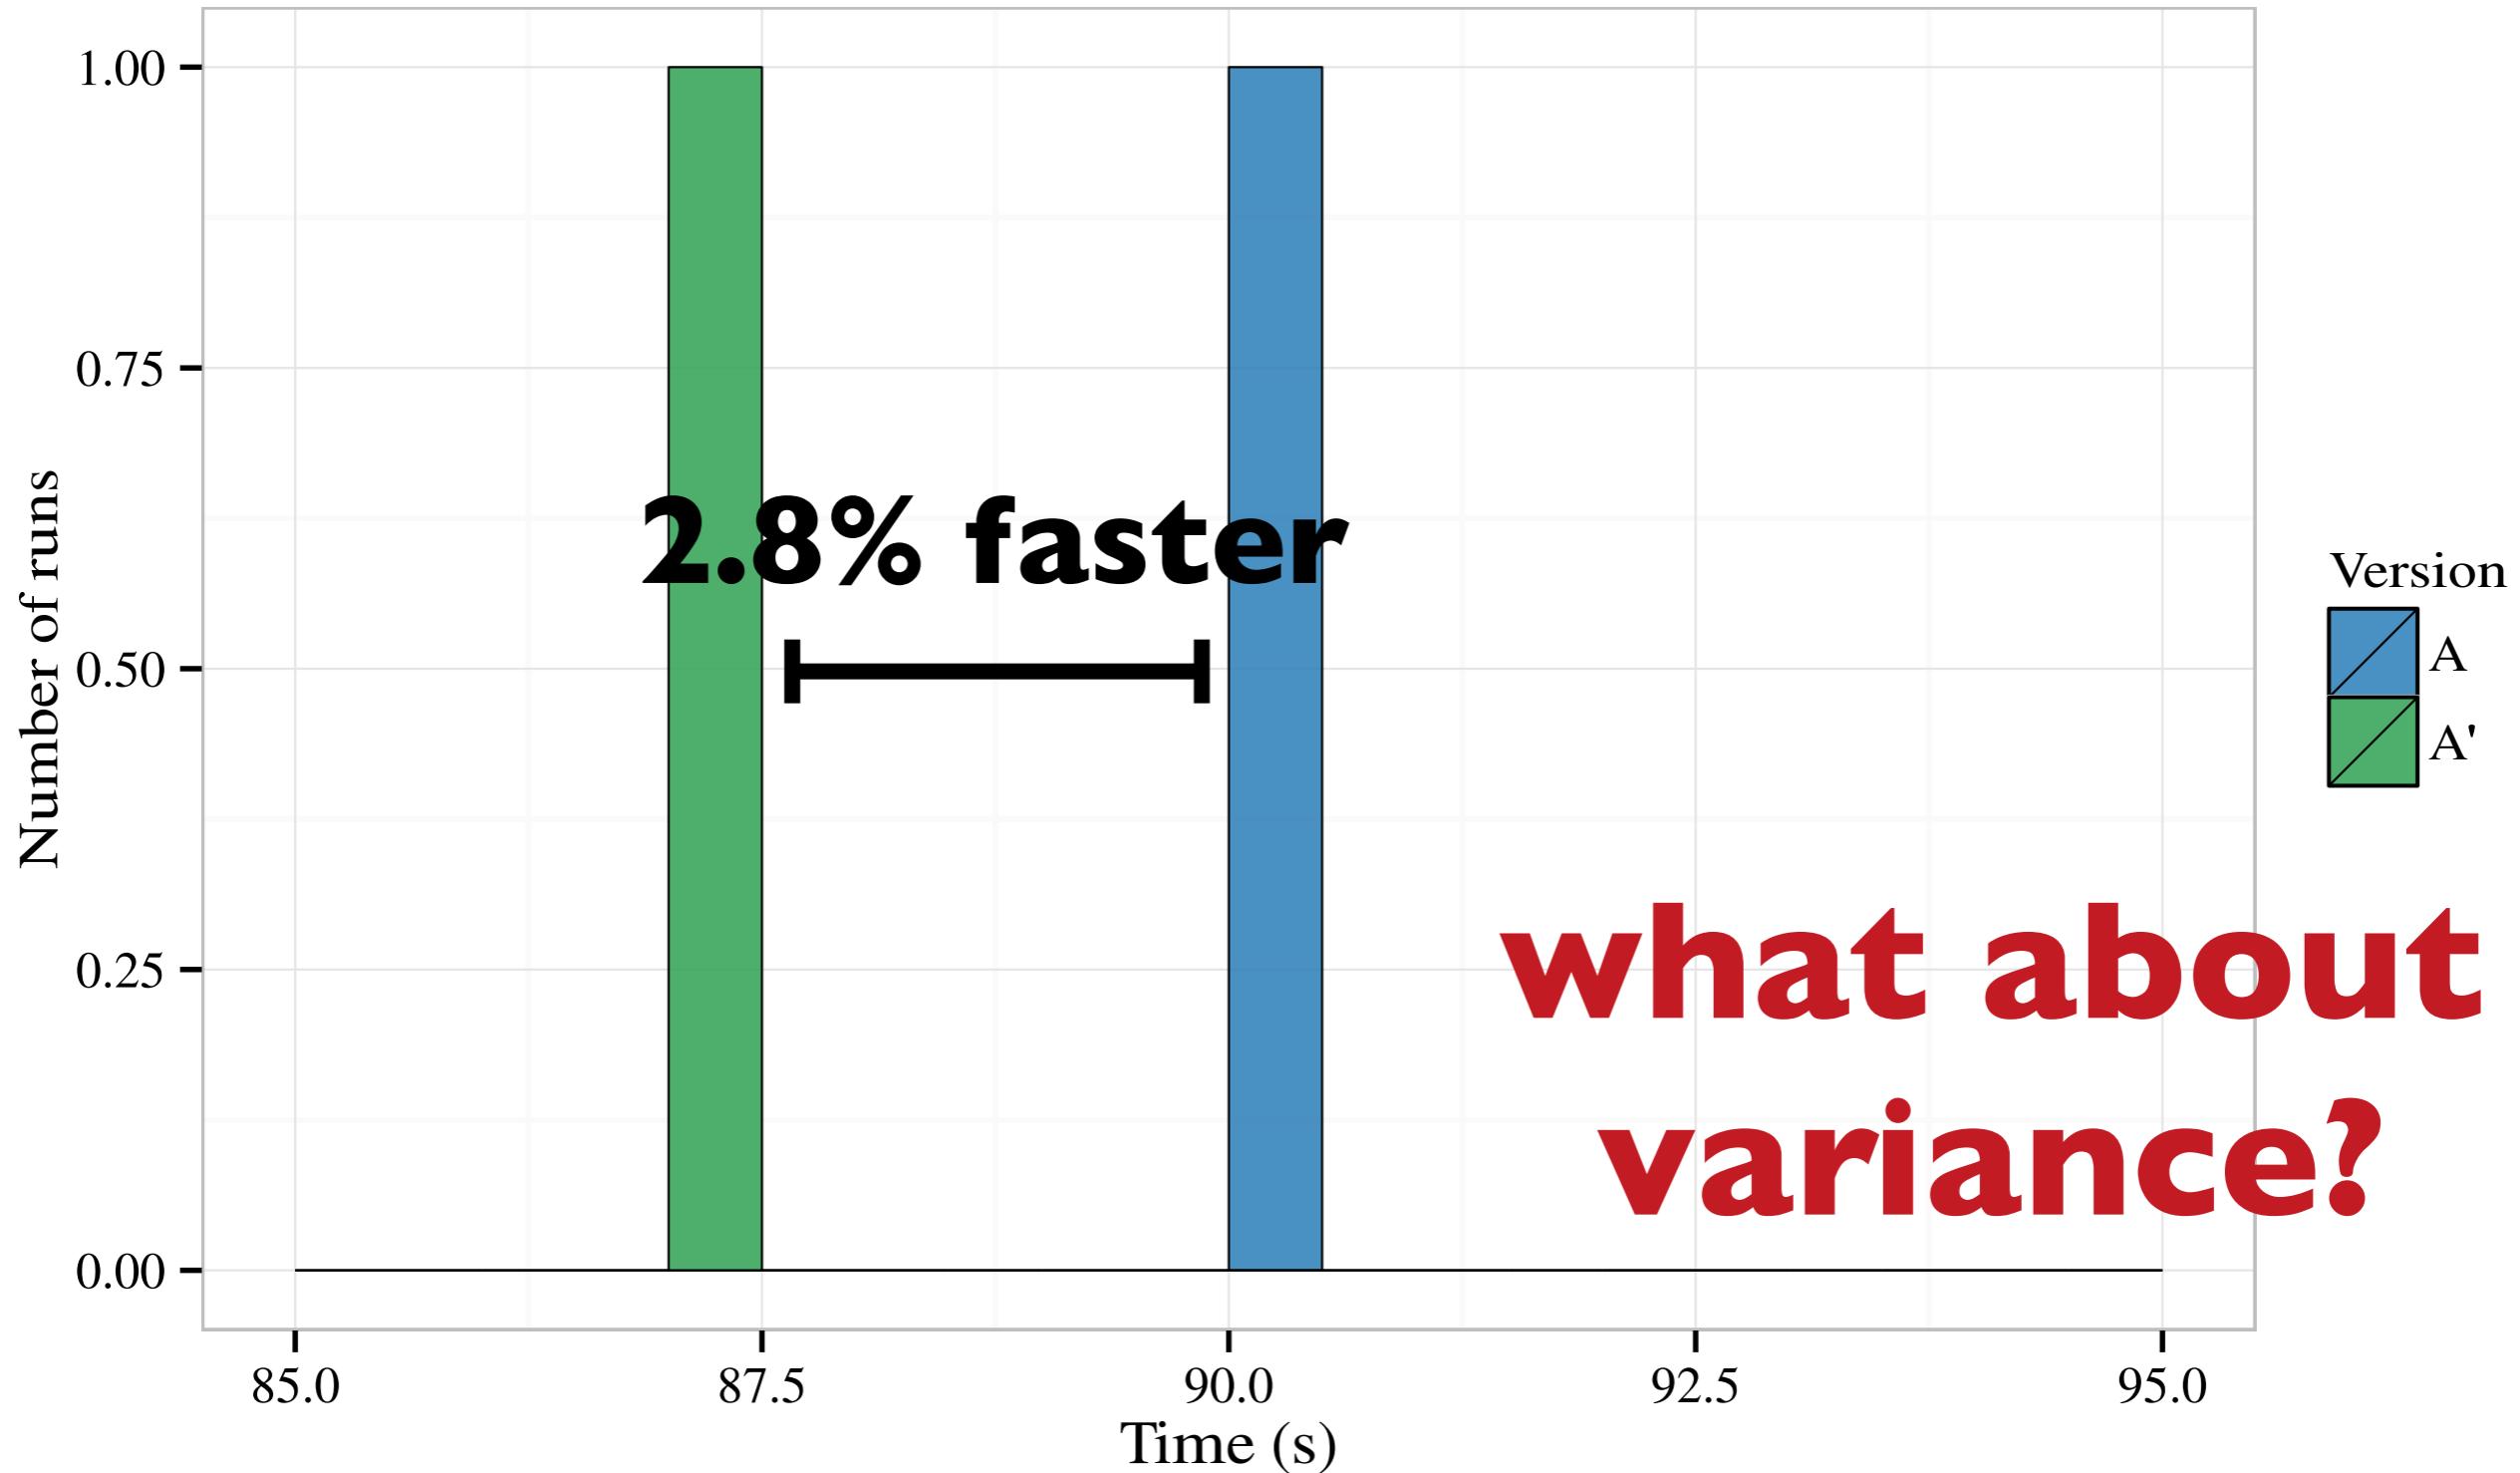
87.5s

A'

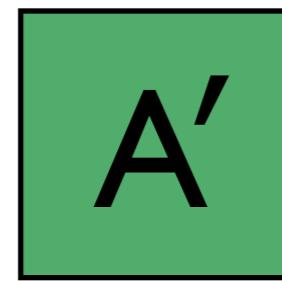
2.8% faster!



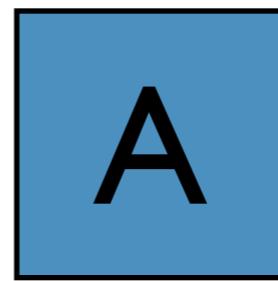
Is A' faster than A ?



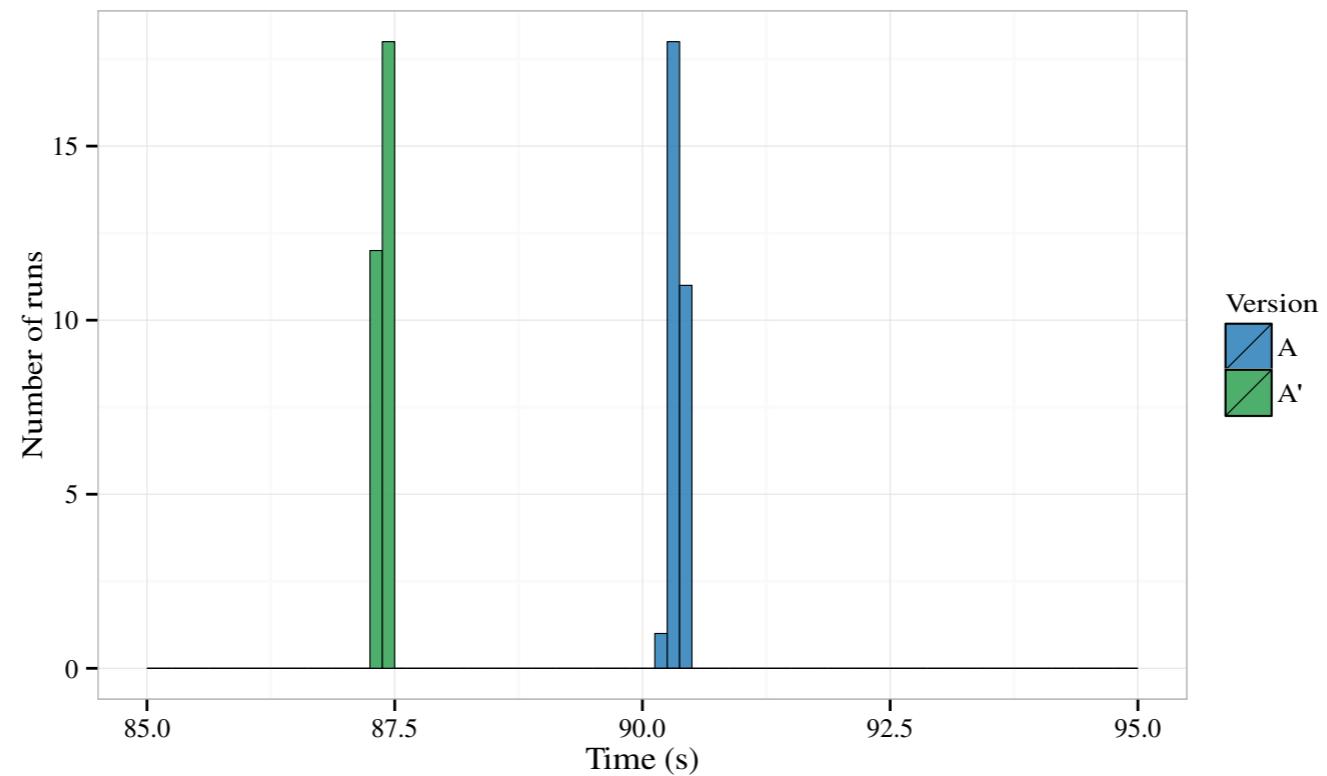
Which is faster?



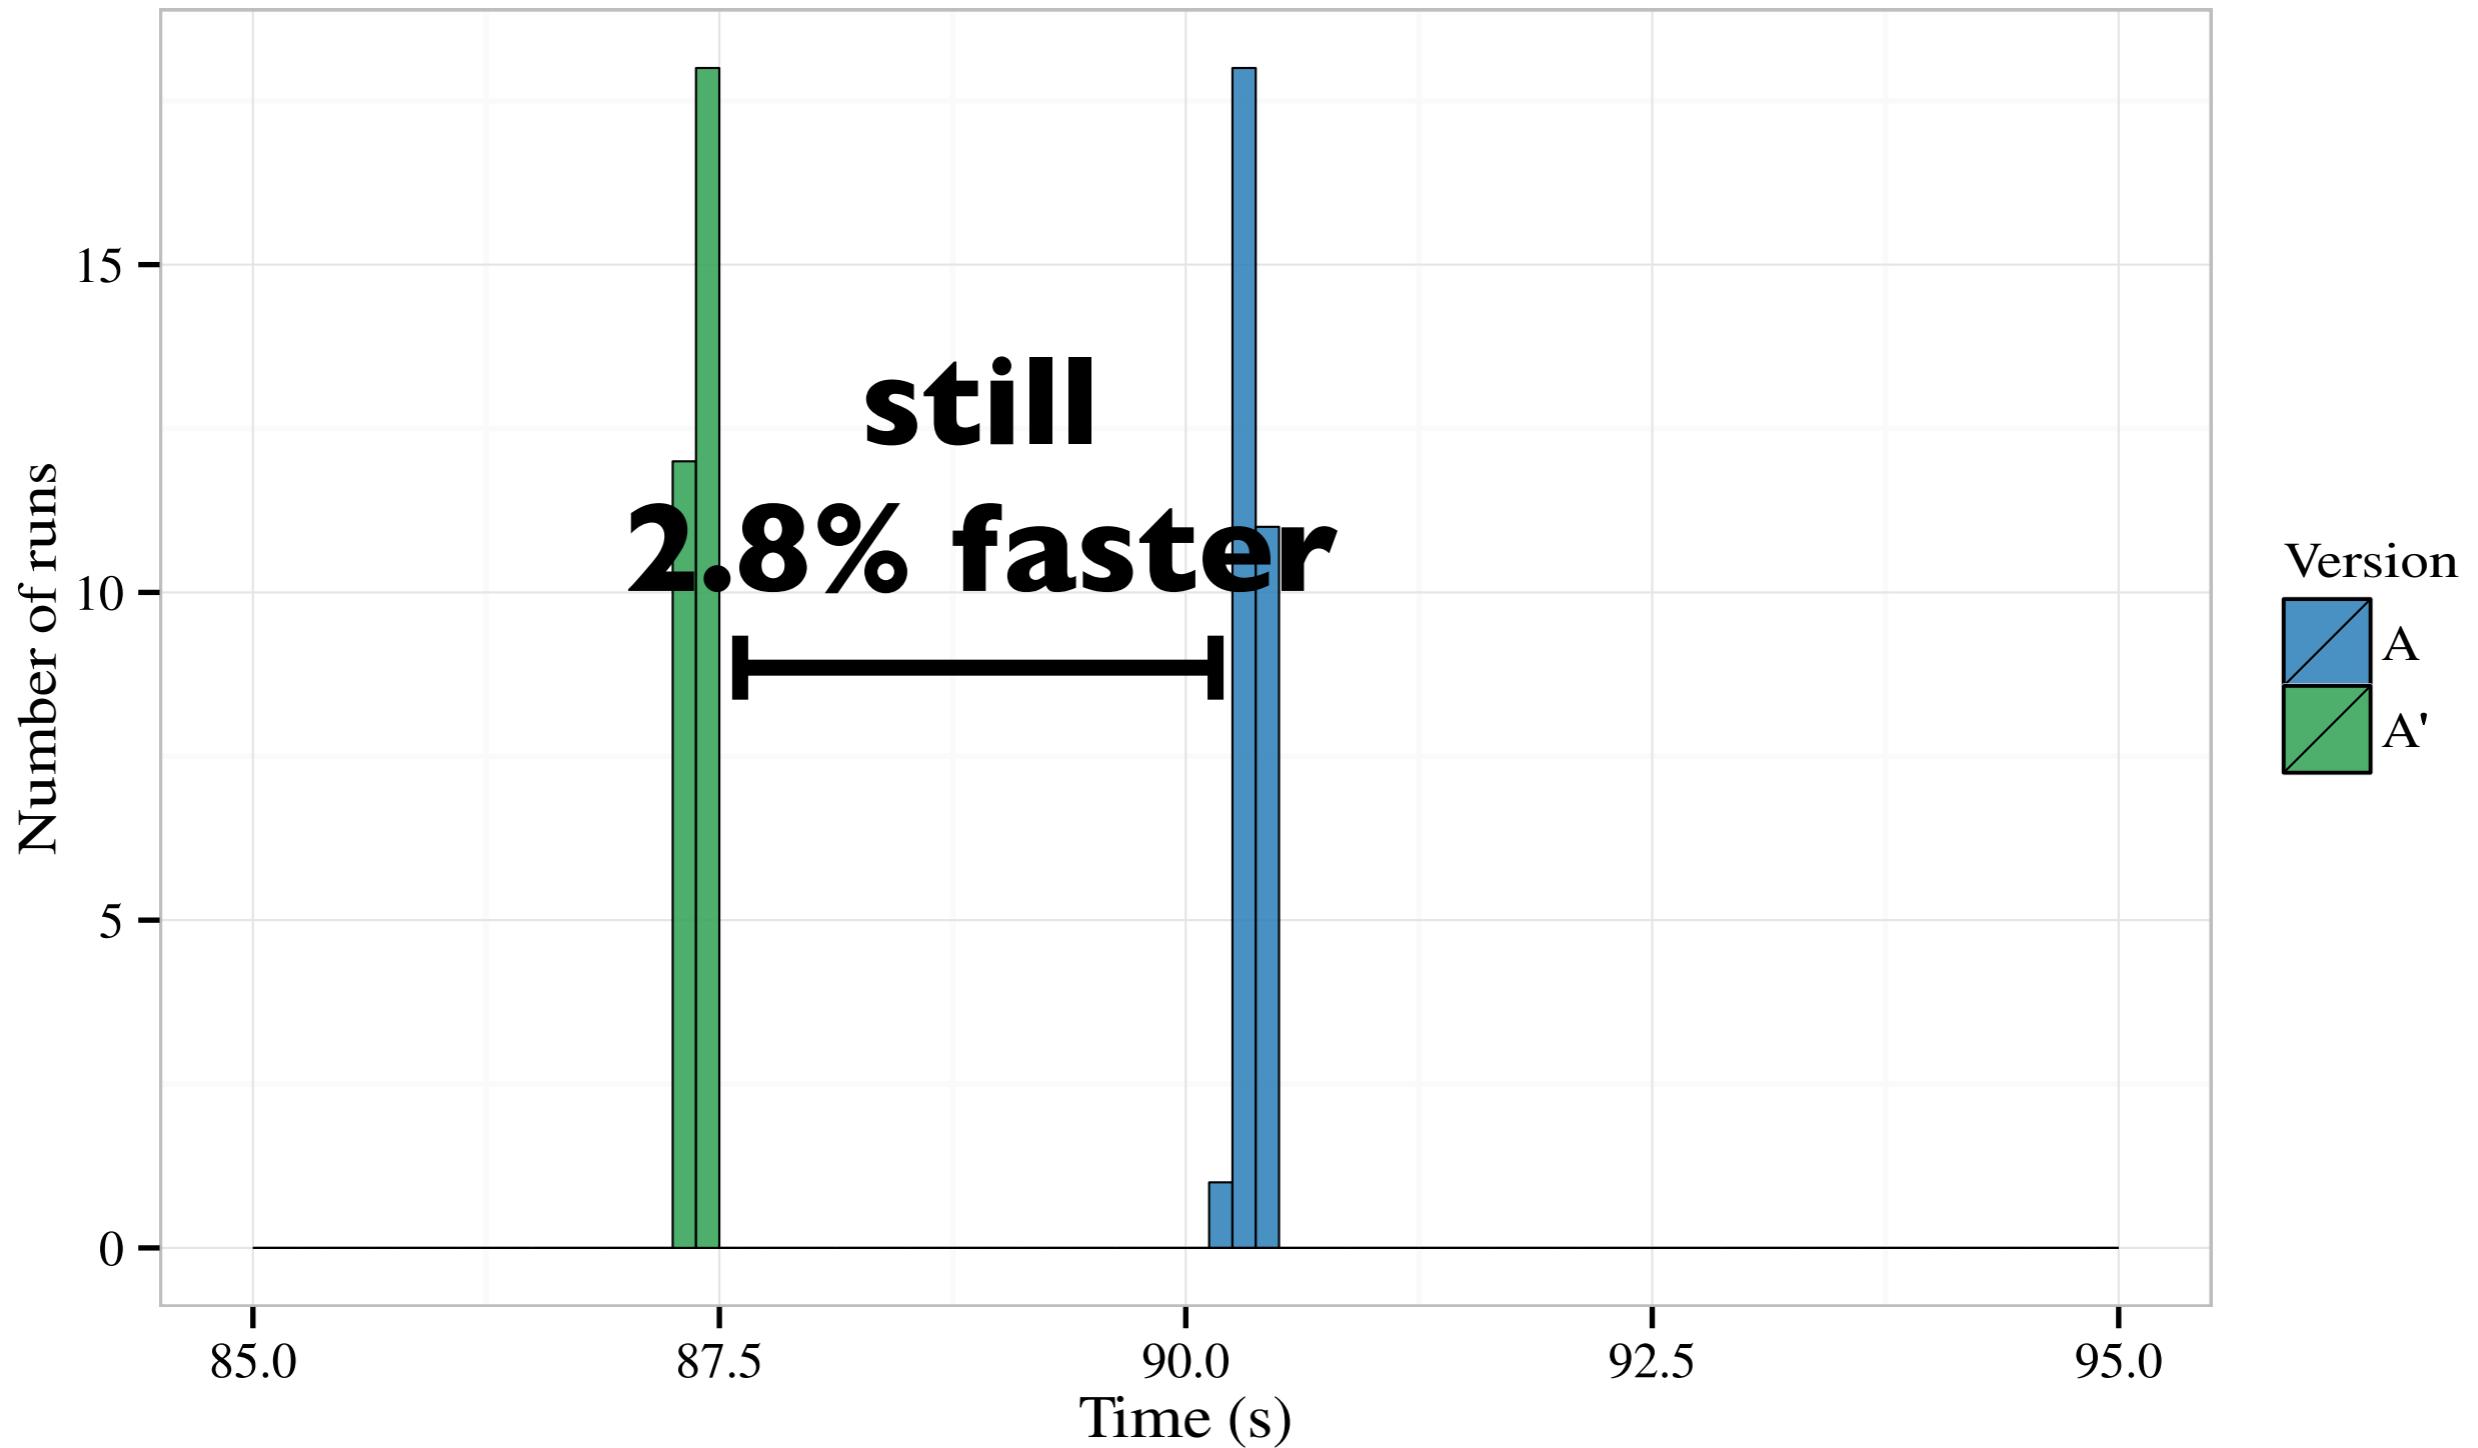
$\times 30$



$\times 30$



Is A' faster than A ?



Why is A' faster than A ?

Why is A' faster than A?

```
DataHeapType* getDataHeap() {
    static char buf[sizeof(DataHeapType)];
    static DataHeapType* _theDataHeap = new (buf) DataHeapType;
    return _theDataHeap;
}
```

```
int main(int argc, char **argv) {
    topFrame = (void**)__builtin_frame_address(0);
    setHandler(Trap::TrapSignal, onTrap);
    setHandler(SIGALRM, onTimer);
    setHandler(SIGSEGV, onFault);
    for(Function* f: functions) {
        f->setTrap();
    }
    setTimer(interval);
    int r = stabilizer_main(argc, argv);
    return r;
}
```

```
static void flush_icache(void* begin, size_t size) {
    uintptr_t p = (uintptr_t)begin & ~15UL;
    for (size_t i = 0; i < size; i += 16) {
        asm("icbi 0,%0" : : "r"(p));
        p += 16;
    }
}
```

```
void setTimer(int msec) {
    struct itimerval timer;

    timer.it_value.tv_sec = (msec - msec % 1000) / 1000;
    timer.it_value.tv_usec = 1000 * (msec % 1000);
    timer.it_interval.tv_sec = 0;
    timer.it_interval.tv_usec = 0;
    setitimer(ITIMER_REAL, &timer, 0);
}
```

```
int main(int argc, char **argv) {
    topFrame = (void**)__builtin_frame_address(0);
    setHandler(Trap::TrapSignal, onTrap);
    setHandler(SIGALRM, onTimer);
    setHandler(SIGSEGV, onFault);
    for(Function* f: functions) {
        f->setTrap();
    }
    setTimer(interval);
    int r = stabilizer_main(argc, argv);
    return r;
}

void setTimer(int msec) {
    struct itimerval timer;

    timer.it_value.tv_sec = (msec - msec % 1000) / 1000;
    timer.it_value.tv_usec = 1000 * (msec % 1000);
    timer.it_interval.tv_sec = 0;
    timer.it_interval.tv_usec = 0;

    setitimer(ITIMER_REAL, &timer, 0);
}
```

```
static void flush_icache(void* begin, size_t size) {
    uintptr_t p = (uintptr_t)begin & ~15UL;
    for (size_t i = 0; i < size; i += 16) {
        asm("icbi 0,%0" : : "r"(p));
        p += 16;
    }
    asm("isync");
}
```

```
DataHeapType* getDataHeap() {
    static char buf[sizeof(DataHeapType)];
    static DataHeapType* _theDataHeap = new (buf) DataHeapType;
    return _theDataHeap;
}
```

Why is A' faster than A?

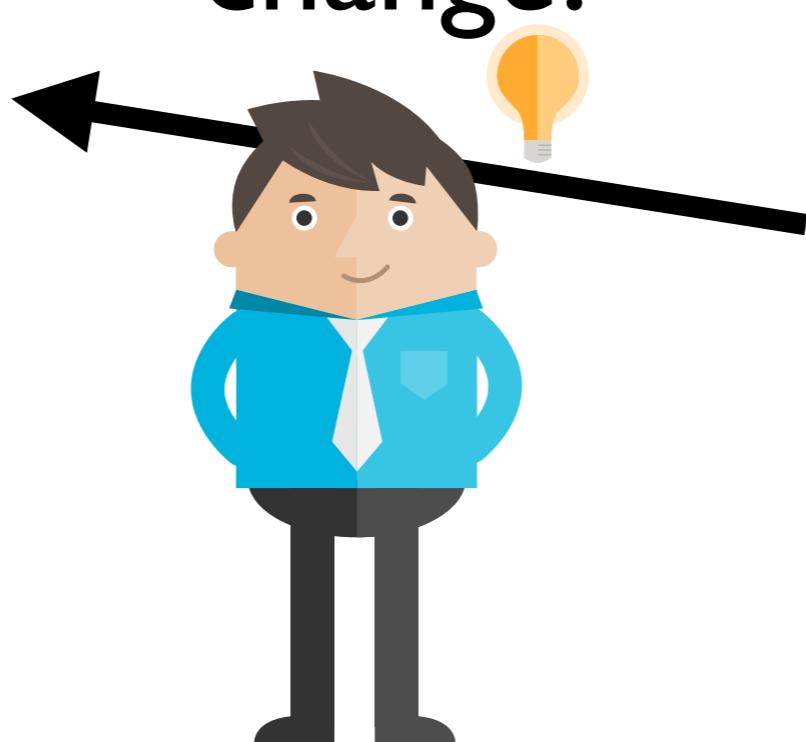
```
DataHeapType* getDataHeap() {  
    static char buf[sizeof(DataHeapType)];  
    static DataHeapType* _theDataHeap = new (buf) DataHeapType;  
    return _theDataHeap;  
}
```

```
int main(int argc, char **argv) {  
    topFrame = (void**)__builtin_frame_address(0);  
    setHandler(Trap::TrapSignal, onTrap);  
    setHandler(SIGALRM, onTimer);  
    setHandler(SIGSEGV, onFault);  
    for(Function* f: functions) {  
        f->setTrap();  
    }  
    setTimer(interval);  
    int r = stabilizer_main(argc, argv);  
    return r;  
}
```

```
static void flush_icache(void* begin, size_t size) {  
    uintptr_t p = (uintptr_t)begin & ~15UL;  
    for (size_t i = 0; i < size; i += 16) {  
        asm("icbi 0,%0" : : "r"(p));  
        p += 16;  
    }  
}
```

```
void setTimer(int msec) {  
    struct itimerval timer;  
  
    timer.it_value.tv_sec = (msec - msec % 1000) / 1000;  
    timer.it_value.tv_usec = 1000 * (msec % 1000);  
    timer.it_interval.tv_sec = 0;  
    timer.it_interval.tv_usec = 0;  
  
    setitimer(ITIMER_REAL, &timer, 0);  
}
```

The code change!



```
int main(int argc, char **argv) {  
    topFrame = (void**)__builtin_frame_address(0);  
    setHandler(Trap::TrapSignal, onTrap);  
    setHandler(SIGALRM, onTimer);  
    setHandler(SIGSEGV, onFault);  
    for(Function* f: functions) {  
        f->setTrap();  
    }  
    setTimer(interval);  
    int r = stabilizer_main(argc, argv);  
    return r;  
}  
  
void setTimer(int msec) {  
    struct itimerval timer;  
  
    timer.it_value.tv_sec = (msec - msec % 1000) / 1000;  
    timer.it_value.tv_usec = 1000 * (msec % 1000);  
    timer.it_interval.tv_sec = 0;  
    timer.it_interval.tv_usec = 0;  
  
    setitimer(ITIMER_REAL, &timer, 0);  
}
```

```
static void flush_icache(void* begin, size_t size) {  
    uintptr_t p = (uintptr_t)begin & ~15UL;  
    for (size_t i = 0; i < size; i += 16) {  
        asm("icbi 0,%0" : : "r"(p));  
        p += 16;  
    }  
    asm("isync");  
}
```

```
DataHeapType* getDataHeap() {  
    static char buf[sizeof(DataHeapType)];  
    static DataHeapType* _theDataHeap = new (buf) DataHeapType;  
    return _theDataHeap;  
}
```

Why is A' faster than A?

Or was it

an accident?

```
DataHeapType* getDataHeap() {
    static char buf[sizeof(DataHeapType)];
    static DataHeapType* _theDataHeap = new (buf) DataHeapType;
    return _theDataHeap;
}
```

```
int main(int argc, char **argv) {
    topFrame = (void**)__builtin_frame_address(0);
    setHandler(Trap::TrapSignal, onTrap);
    setHandler(SIGALRM, onTimer);
    setHandler(SIGSEGV, onFault);
    for(Function* f: functions) {
        f->setTrap();
    }
    setTimer(interval);
    int r = stabilizer_main(argc, argv);
    return r;
}
```

```
static void flush_icache(void* begin, size_t size) {
    uintptr_t p = (uintptr_t)begin & ~15UL;
    for (size_t i = 0; i < size; i += 16) {
        asm("icbi 0,%0" : : "r"(p));
        p += 16;
    }
}
```

```
void setTimer(int msec) {
    struct itimerval timer;

    timer.it_value.tv_sec = (msec - msec % 1000) / 1000;
    timer.it_value.tv_usec = 1000 * (msec % 1000);
    timer.it_interval.tv_sec = 0;
    timer.it_interval.tv_usec = 0;

    setitimer(ITIMER_REAL, &timer, 0);
}
```

```
int main(int argc, char **argv) {
    topFrame = (void**)__builtin_frame_address(0);
    setHandler(Trap::TrapSignal, onTrap);
    setHandler(SIGALRM, onTimer);
    setHandler(SIGSEGV, onFault);
    for(Function* f: functions) {
        f->setTrap();
    }
    setTimer(interval);
    int r = stabilizer_main(argc, argv);
    return r;
}

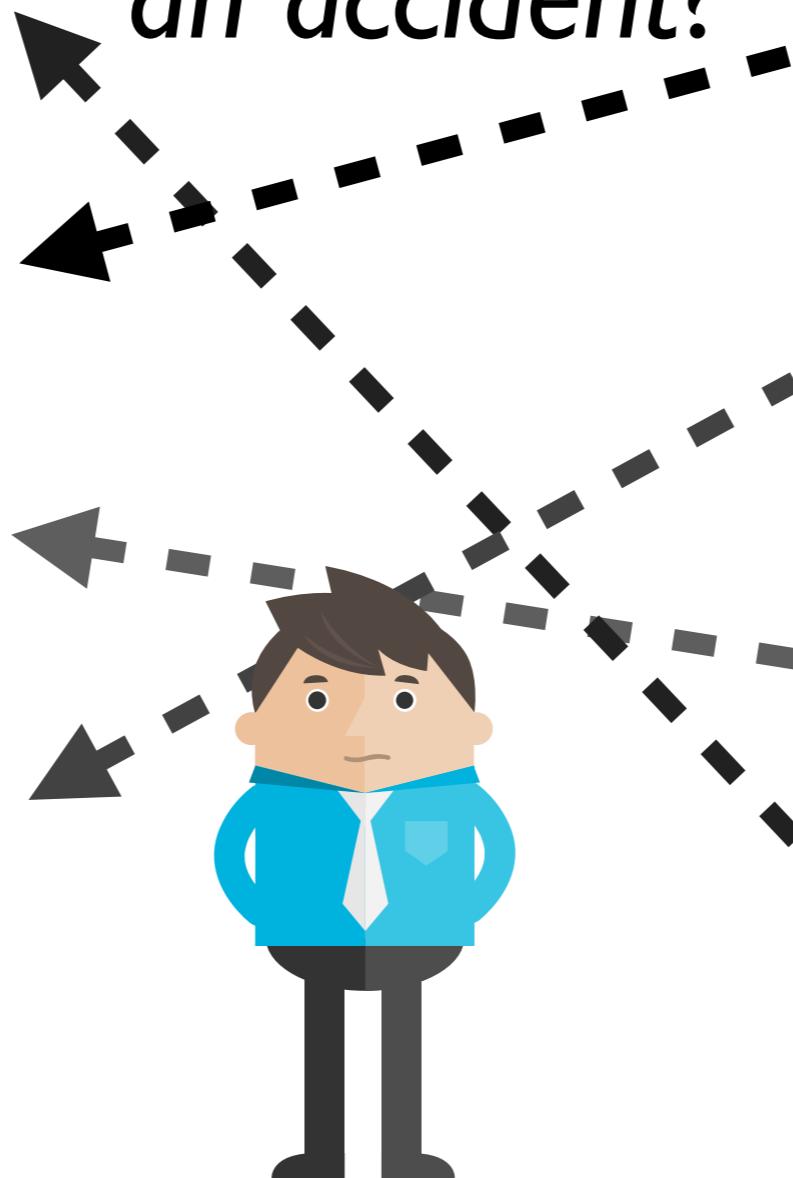
void setTimer(int msec) {
    struct itimerval timer;

    timer.it_value.tv_sec = (msec - msec % 1000) / 1000;
    timer.it_value.tv_usec = 1000 * (msec % 1000);
    timer.it_interval.tv_sec = 0;
    timer.it_interval.tv_usec = 0;

    setitimer(ITIMER_REAL, &timer, 0);
}
```

```
static void flush_icache(void* begin, size_t size) {
    uintptr_t p = (uintptr_t)begin & ~15UL;
    for (size_t i = 0; i < size; i += 16) {
        asm("icbi 0,%0" : : "r"(p));
        p += 16;
    }
    asm("isync");
}
```

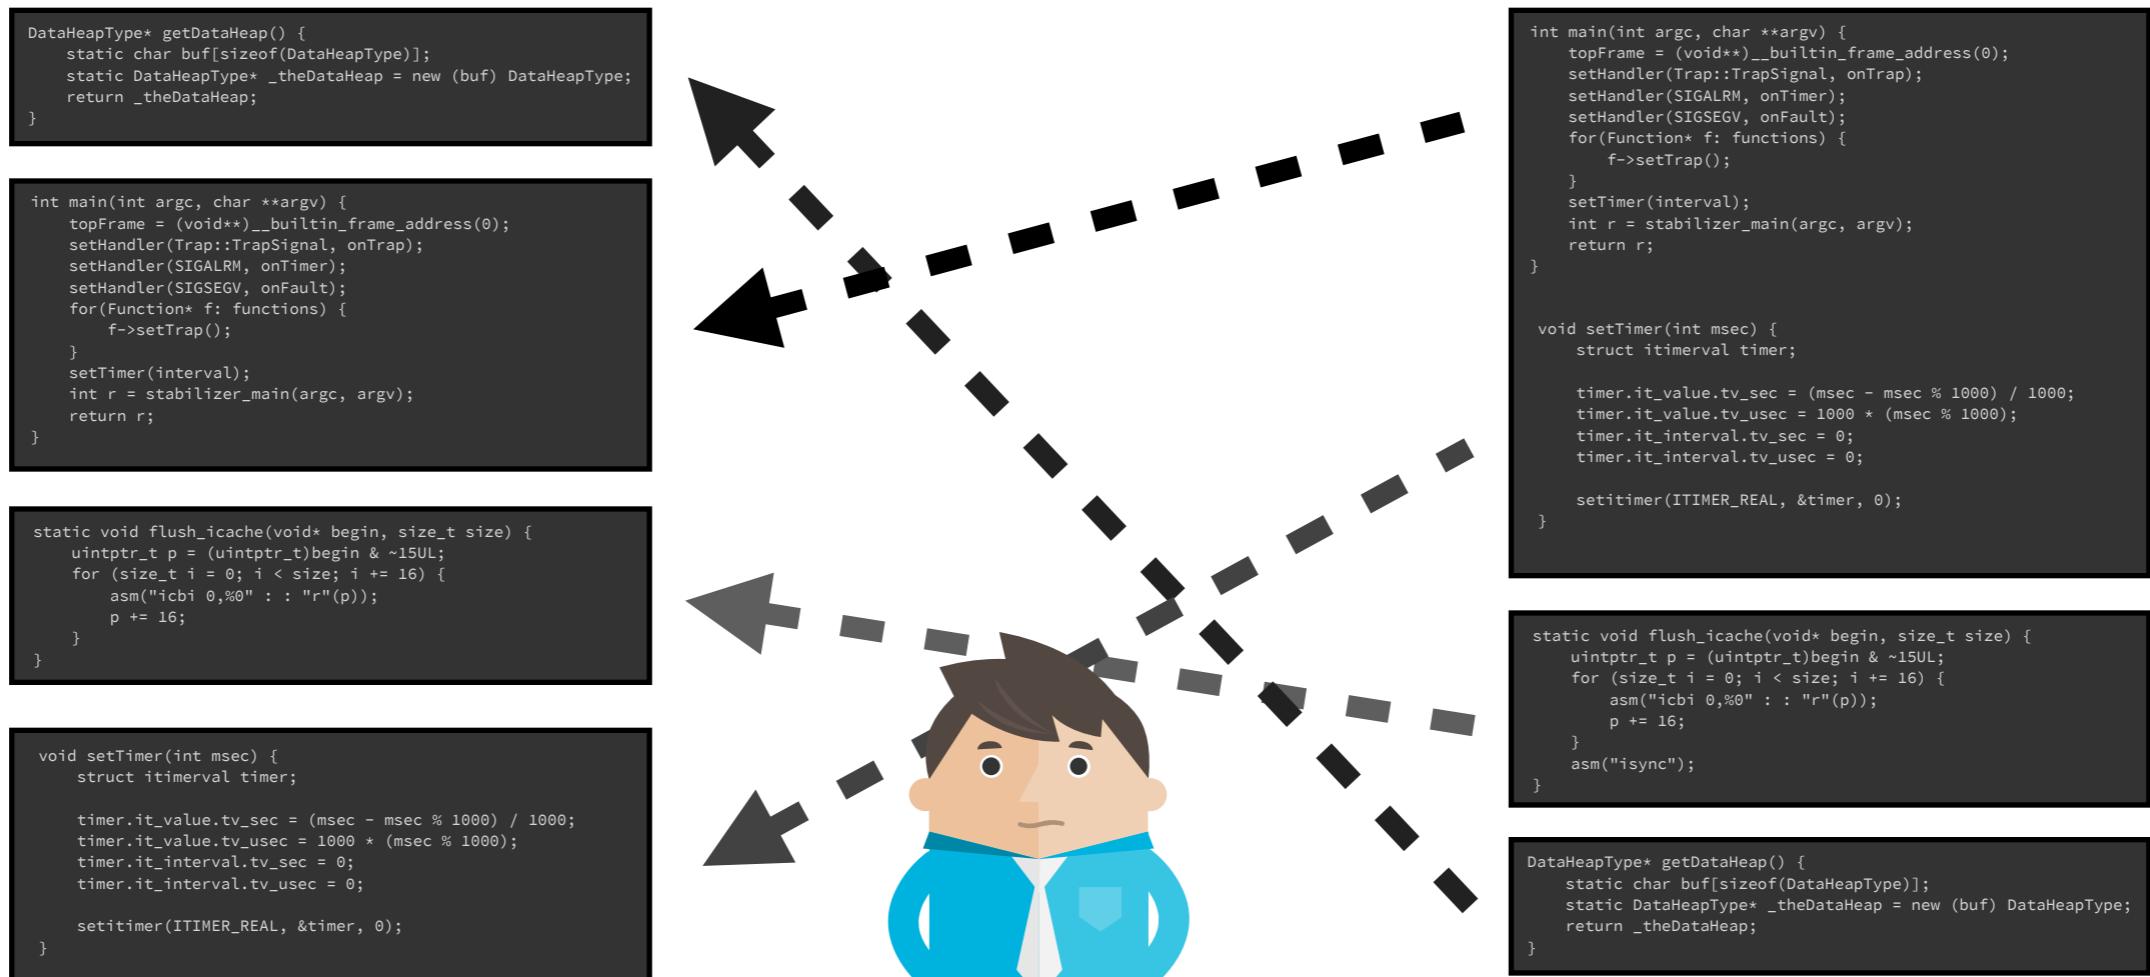
```
DataHeapType* getDataHeap() {
    static char buf[sizeof(DataHeapType)];
    static DataHeapType* _theDataHeap = new (buf) DataHeapType;
    return _theDataHeap;
}
```



Why is A' faster than A ?

Layout biases measurement

Mytkowicz et al. (ASPLOS'09)



Layout biases measurement

Mytkowicz et al. (ASPLOS'09)

Producing Wrong Data Without Doing Anything Obviously Wrong!

Todd Mytkowicz Amer Diwan

Department of Computer Science
University of Colorado
Boulder, CO, USA

{mytkowit,diwan}@colorado.edu

Matthias Hauswirth

Faculty of Informatics
University of Lugano
Lugano, CH

Matthias.Hauswirth@unisi.ch

Peter F. Sweeney

IBM Research
Hawthorne, NY, USA
pfs@us.ibm.com

Link Order

Changes function addresses

Environment

Variable Size

Moves the program stack

Layout biases measurement

Mytkowicz et al. (ASPLOS'09)

Link Order

Changes function addresses

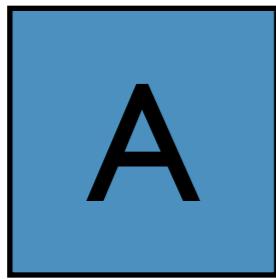
**Larger than the impact
of -O3! ($\pm 40\%$)**

Environment

Variable Size

Moves the program stack

Why is A' faster than A?



```
int main(int argc, char **argv) {
    topFrame = (void**)__builtin_frame_address(0);
    setHandler(Trap::TrapSignal, onTrap);
    setHandler(SIGALRM, onTimer);
    setHandler(SIGSEGV, onFault);
    for(Function* f: functions) {
        f->setTrap();
    }
    setTimer(interval);
    int r = stabilizer_main(argc, argv);
    return r;
}

void setTimer(int msec) {
    struct itimerval timer;

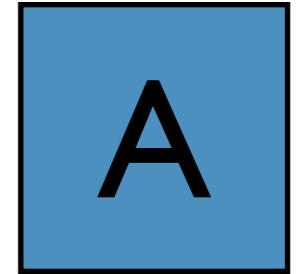
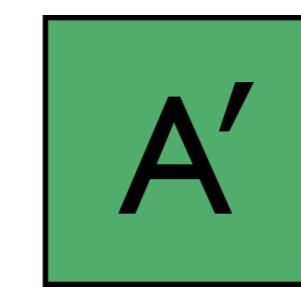
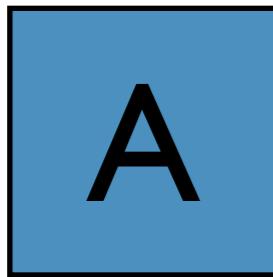
    timer.it_value.tv_sec = (msec - msec % 1000) / 1000;
    timer.it_value.tv_usec = 1000 * (msec % 1000);
    timer.it_interval.tv_sec = 0;
    timer.it_interval.tv_usec = 0;

    setitimer(ITIMER_REAL, &timer, 0);
}
```

```
static void flush_icache(void* begin, size_t size) {
    uintptr_t p = (uintptr_t)begin & ~15UL;
    for (size_t i = 0; i < size; i += 16) {
        asm("icbi 0,%0" : : "r"(p));
        p += 16;
    }
    asm("isync");
}
```

```
DataHeapType* getDataHeap() {
    static char buf[sizeof(DataHeapType)];
    static DataHeapType* _theDataHeap = new (buf) DataHeapType;
    return _theDataHeap;
}
```

Why is A' faster than A?



```
int main(int argc, char **argv) {
    topFrame = (void**)__builtin_frame_address(0);
    setHandler(Trap::TrapSignal, onTrap);
    setHandler(SIGALRM, onTimer);
    setHandler(SIGSEGV, onFault);
    for(Function* f: functions) {
        f->setTrap();
    }
    setTimer(interval);
    int r = stabilizer_main(argc, argv);
    return r;
}
```

```
void setTimer(int msec) {
    struct itimerval timer;

    timer.it_value.tv_sec = (msec - msec % 1000) / 1000;
    timer.it_value.tv_usec = 1000 * (msec % 1000);
    timer.it_interval.tv_sec = 0;
    timer.it_interval.tv_usec = 0;

    setitimer(ITIMER_REAL, &timer, 0);
}
```

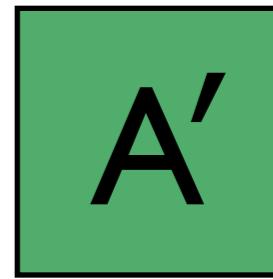
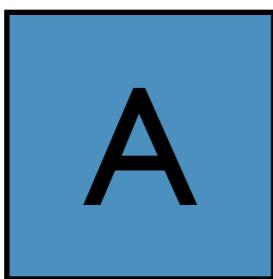
```
static void flush_icache(void* begin, size_t size) {
    uintptr_t p = (uintptr_t)begin & ~15UL;
    for (size_t i = 0; i < size; i += 16) {
        asm("icbi 0,%0" : : "r"(p));
        p += 16;
    }
    asm("isync");
}
```

```
DataHeapType* getDataHeap() {
    static char buf[sizeof(DataHeapType)];
    static DataHeapType* _theDataHeap = new (buf) DataHeapType;
    return _theDataHeap;
}
```

map to same
cache set
conflict



Why is A' faster than A?



```
int main(int argc, char **argv) {
    topFrame = (void**)__builtin_frame_address(0);
    setHandler(Trap::TrapSignal, onTrap);
    setHandler(SIGALRM, onTimer);
    setHandler(SIGSEGV, onFault);
    for(Function* f: functions) {
        f->setTrap();
    }
    setTimer(interval);
    int r = stabilizer_main(argc, argv);
    return r;
}

void setTimer(int msec) {
    struct itimerval timer;

    timer.it_value.tv_sec = (msec - msec % 1000) / 1000;
    timer.it_value.tv_usec = 1000 * (msec % 1000);
    timer.it_interval.tv_sec = 0;
    timer.it_interval.tv_usec = 0;

    setitimer(ITIMER_REAL, &timer, 0);
}
```



```
static void flush_icache(void* begin, size_t size) {
    uintptr_t p = (uintptr_t)begin & ~15UL;
    for (size_t i = 0; i < size; i += 16) {
        asm("icbi 0,%0" : : "r"(p));
        p += 16;
    }
    asm("isync");
}
```

```
DataHeapType* getDataHeap() {
    static char buf[sizeof(DataHeapType)];
    static DataHeapType* _theDataHeap = new (buf) DataHeapType;
    return _theDataHeap;
}
```

```
DataHeapType* getDataHeap() {
    static char buf[sizeof(DataHeapType)];
    static DataHeapType* _theDataHeap = new (buf) DataHeapType;
    return _theDataHeap;
}
```

```
int main(int argc, char **argv) {
    topFrame = (void**)__builtin_frame_address(0);
    setHandler(Trap::TrapSignal, onTrap);
    setHandler(SIGALRM, onTimer);
    setHandler(SIGSEGV, onFault);
    for(Function* f: functions) {
        f->setTrap();
    }
    setTimer(interval);
    int r = stabilizer_main(argc, argv);
    return r;
}
```

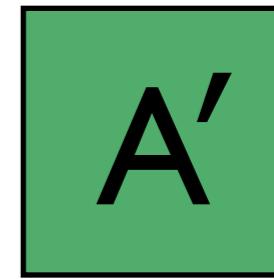
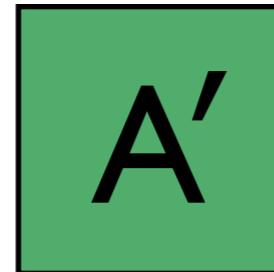
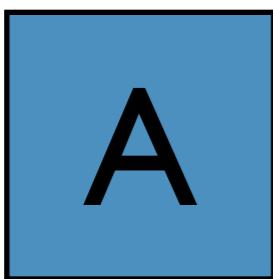
```
static void flush_icache(void* begin, size_t size) {
    uintptr_t p = (uintptr_t)begin & ~15UL;
    for (size_t i = 0; i < size; i += 16) {
        asm("icbi 0,%0" : : "r"(p));
        p += 16;
    }
}
```

```
void setTimer(int msec) {
    struct itimerval timer;

    timer.it_value.tv_sec = (msec - msec % 1000) / 1000;
    timer.it_value.tv_usec = 1000 * (msec % 1000);
    timer.it_interval.tv_sec = 0;
    timer.it_interval.tv_usec = 0;

    setitimer(ITIMER_REAL, &timer, 0);
}
```

Why is A' faster than A?



?

```
int main(int argc, char **argv) {
    topFrame = (void**)__builtin_frame_address(0);
    setHandler(Trap::TrapSignal, onTrap);
    setHandler(SIGALRM, onTimer);
    setHandler(SIGSEGV, onFault);
    for(Function* f: functions) {
        f->setTrap();
    }
    setTimer(interval);
    int r = stabilizer_main(argc, argv);
    return r;
}

void setTimer(int msec) {
    struct itimerval timer;

    timer.it_value.tv_sec = (msec - msec % 1000) / 1000;
    timer.it_value.tv_usec = 1000 * (msec % 1000);
    timer.it_interval.tv_sec = 0;
    timer.it_interval.tv_usec = 0;

    setitimer(ITIMER_REAL, &timer, 0);
}
```

```
static void flush_icache(void* begin, size_t size) {
    uintptr_t p = (uintptr_t)begin & ~15UL;
    for (size_t i = 0; i < size; i += 16) {
        asm("icbi 0,%0" : : "r"(p));
        p += 16;
    }
    asm("isync");
}
```

```
DataHeapType* getDataHeap() {
    static char buf[sizeof(DataHeapType)];
    static DataHeapType* _theDataHeap = new (buf) DataHeapType;
    return _theDataHeap;
}
```

map to
same set

Nothing here →

no conflict

```
DataHeapType* getDataHeap() {
    static char buf[sizeof(DataHeapType)];
    static DataHeapType* _theDataHeap = new (buf) DataHeapType;
    return _theDataHeap;
}
```

```
int main(int argc, char **argv) {
    topFrame = (void**)__builtin_frame_address(0);
    setHandler(Trap::TrapSignal, onTrap);
    setHandler(SIGALRM, onTimer);
    setHandler(SIGSEGV, onFault);
    for(Function* f: functions) {
        f->setTrap();
    }
    setTimer(interval);
    int r = stabilizer_main(argc, argv);
    return r;
}
```

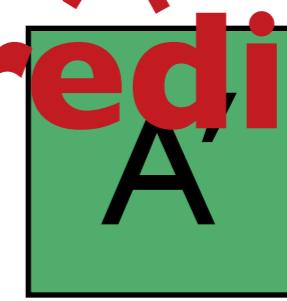
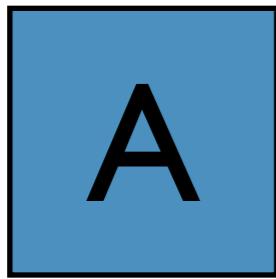
```
static void flush_icache(void* begin, size_t size) {
    uintptr_t p = (uintptr_t)begin & ~15UL;
    for (size_t i = 0; i < size; i += 16) {
        asm("icbi 0,%0" : : "r"(p));
        p += 16;
    }
}
```

```
void setTimer(int msec) {
    struct itimerval timer;

    timer.it_value.tv_sec = (msec - msec % 1000) / 1000;
    timer.it_value.tv_usec = 1000 * (msec % 1000);
    timer.it_interval.tv_sec = 0;
    timer.it_interval.tv_usec = 0;

    setitimer(ITIMER_REAL, &timer, 0);
}
```

It's the cache or branch predictor or branch



or TLB

```
int main(int argc, char **argv) {
    topFrame = (void**)__builtin_frame_address(0);
    setHandler(Trap::TrapSignal, onTrap);
    setHandler(SIGALRM, onTimer);
    setHandler(SIGSEGV, onFault);
    for(Function* f: functions) {
        f->setTrap();
    }
    setTimer(interval);
    int r = stabilizer_main(argc, argv);
    return r;
}

void setTimer(int msec) {
    struct itimerval timer;

    timer.it_value.tv_sec = (msec - msec % 1000) / 1000;
    timer.it_value.tv_usec = 1000 * (msec % 1000);
    timer.it_interval.tv_sec = 0;
    timer.it_interval.tv_usec = 0;

    setitimer(ITIMER_REAL, &timer, 0);
}
```

```
static void flush_icache(void* begin, size_t size) {
    uintptr_t p = (uintptr_t)begin & ~15UL;
    for (size_t i = 0; i < size; i += 16) {
        asm("icbi 0,%0" : : "r"(p));
        p += 16;
    }
    asm("isync");
}
```

```
DataHeapType* getDataHeap() {
    static char buf[sizeof(DataHeapType)];
    static DataHeapType* _theDataHeap = new (buf) DataHeapType;
    return _theDataHeap;
}
```

```
DataHeapType* getDataHeap() {
    static char buf[sizeof(DataHeapType)];
    static DataHeapType* _theDataHeap = new (buf) DataHeapType;
    return _theDataHeap;
}
```

```
int main(int argc, char **argv) {
    topFrame = (void**)__builtin_frame_address(0);
    setHandler(Trap::TrapSignal, onTrap);
    setHandler(SIGALRM, onTimer);
    setHandler(SIGSEGV, onFault);
    for(Function* f: functions) {
        f->setTrap();
    }
    setTimer(interval);
    int r = stabilizer_main(argc, argv);
    return r;
}
```

```
static void flush_icache(void* begin, size_t size) {
    uintptr_t p = (uintptr_t)begin & ~15UL;
    for (size_t i = 0; i < size; i += 16) {
        asm("icbi 0,%0" : : "r"(p));
        p += 16;
    }
}
```

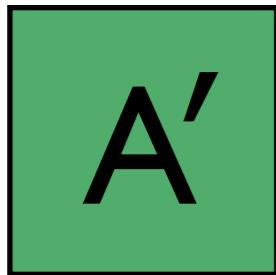
```
void setTimer(int msec) {
    struct itimerval timer;

    timer.it_value.tv_sec = (msec - msec % 1000) / 1000;
    timer.it_value.tv_usec = 1000 * (msec % 1000);
    timer.it_interval.tv_sec = 0;
    timer.it_interval.tv_usec = 0;

    setitimer(ITIMER_REAL, &timer, 0);
}
```

target predictor or prefetcher

Might be faster today...



```
DataHeapType* getDataHeap() {
    static char buf[sizeof(DataHeapType)];
    static DataHeapType* _theDataHeap = new (buf) DataHeapType;
    return _theDataHeap;
}
```

```
int main(int argc, char **argv) {
    topFrame = (void**)__builtin_frame_address(0);
    setHandler(Trap::TrapSignal, onTrap);
    setHandler(SIGALRM, onTimer);
    setHandler(SIGSEGV, onFault);
    for(Function* f: functions) {
        f->setTrap();
    }
    setTimer(interval);
    int r = stabilizer_main(argc, argv);
    return r;
}
```

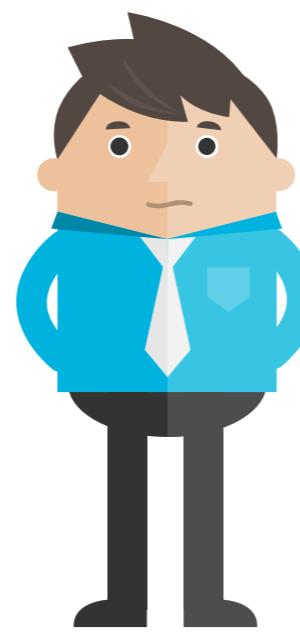
```
static void flush_icache(void* begin, size_t size) {
    uintptr_t p = (uintptr_t)begin & ~15UL;
    for (size_t i = 0; i < size; i += 16) {
        asm("icbi 0,%0" : : "r"(p));
        p += 16;
    }
}
```

```
void setTimer(int msec) {
    struct itimerval timer;

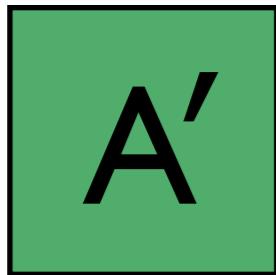
    timer.it_value.tv_sec = (msec - msec % 1000) / 1000;
    timer.it_value.tv_usec = 1000 * (msec % 1000);
    timer.it_interval.tv_sec = 0;
    timer.it_interval.tv_usec = 0;

    setitimer(ITIMER_REAL, &timer, 0);
}
```

One more new
Changes layout



Might be faster *today*...



```
DataHeapType* getDataHeap() {
    static char buf[sizeof(DataHeapType)];
    static DataHeapType* _theDataHeap = new (buf) DataHeapType;
    return _theDataHeap;
}
```

```
int main(int argc, char **argv) {
    topFrame = (void**)__builtin_frame_address(0);
    setHandler(Trap::TrapSignal, onTrap);
    setHandler(SIGALRM, onTimer);
    setHandler(SIGSEGV, onFault);
    for(Function* f: functions) {
        f->setTrap();
    }
    setTimer(interval);
    int r = stabilizer_main(argc, argv);
    return r;
}
```

```
static void flush_icache(void* begin, size_t size) {
    uintptr_t p = (uintptr_t)begin & ~15UL;
    for (size_t i = 0; i < size; i += 16) {
        asm("icbi 0,%0" : : "r"(p));
        p += 16;
    }
}
```

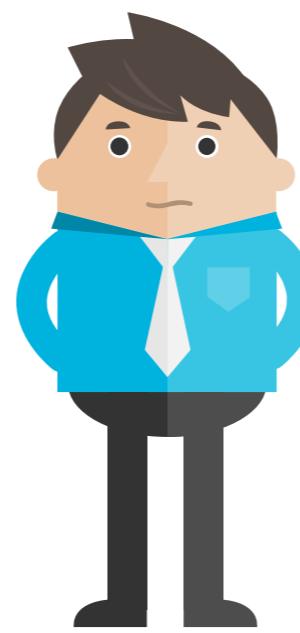
```
void setTimer(int msec) {
    struct itimerval timer;

    timer.it_value.tv_sec = (msec - msec % 1000) / 1000;
    timer.it_value.tv_usec = 1000 * (msec % 1000);
    timer.it_interval.tv_sec = 0;
    timer.it_interval.tv_usec = 0;

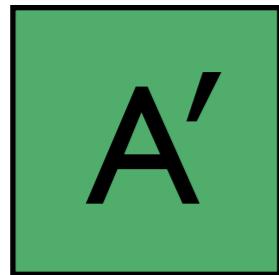
    setitimer(ITIMER_REAL, &timer, 0);
}
```

Change one class

Changes layout



Might be faster today...



```
DataHeapType* getDataHeap() {
    static char buf[sizeof(DataHeapType)];
    static DataHeapType* _theDataHeap = new (buf) DataHeapType;
    return _theDataHeap;
}
```

```
int main(int argc, char **argv) {
    topFrame = (void**)__builtin_frame_address(0);
    setHandler(Trap::TrapSignal, onTrap);
    setHandler(SIGALRM, onTimer);
    setHandler(SIGSEGV, onFault);
    for(Function* f: functions) {
        f->setTrap();
    }
    setTimer(interval);
    int r = stabilizer_main(argc, argv);
    return r;
}
```

```
static void flush_icache(void* begin, size_t size) {
    uintptr_t p = (uintptr_t)begin & ~15UL;
    for (size_t i = 0; i < size; i += 16) {
        asm("icbi 0,%0" : : "r"(p));
        p += 16;
    }
}
```

```
void setTimer(int msec) {
    struct itimerval timer;

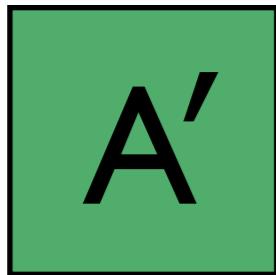
    timer.it_value.tv_sec = (msec - msec % 1000) / 1000;
    timer.it_value.tv_usec = 1000 * (msec % 1000);
    timer.it_interval.tv_sec = 0;
    timer.it_interval.tv_usec = 0;

    setitimer(ITIMER_REAL, &timer, 0);
}
```

Upgrade libc Changes layout



Might be faster *today*...



```
DataHeapType* getDataHeap() {
    static char buf[sizeof(DataHeapType)];
    static DataHeapType* _theDataHeap = new (buf) DataHeapType;
    return _theDataHeap;
}
```

```
int main(int argc, char **argv) {
    topFrame = (void**)__builtin_frame_address(0);
    setHandler(Trap::TrapSignal, onTrap);
    setHandler(SIGALRM, onTimer);
    setHandler(SIGSEGV, onFault);
    for(Function* f: functions) {
        f->setTrap();
    }
    setTimer(interval);
    int r = stabilizer_main(argc, argv);
    return r;
}
```

```
static void flush_icache(void* begin, size_t size) {
    uintptr_t p = (uintptr_t)begin & ~15UL;
    for (size_t i = 0; i < size; i += 16) {
        asm("icbi 0,%0" : : "r"(p));
        p += 16;
    }
}
```

```
void setTimer(int msec) {
    struct itimerval timer;

    timer.it_value.tv_sec = (msec - msec % 1000) / 1000;
    timer.it_value.tv_usec = 1000 * (msec % 1000);
    timer.it_interval.tv_sec = 0;
    timer.it_interval.tv_usec = 0;

    setitimer(ITIMER_REAL, &timer, 0);
}
```

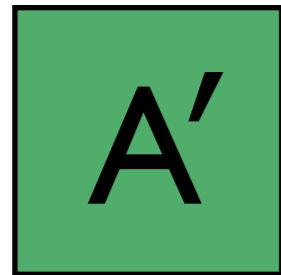
**Run in a new
directory**

Changes layout



% pwd
/users/vader
% cd /users/luke

Might be faster *today*...



**same issues for
perf regression!**

```
DataHeapType* getDataHeap() {
    static char buf[sizeof(DataHeapType)];
    static DataHeapType* _theDataHeap = new (buf) DataHeapType;
    return _theDataHeap;
}
```

```
int main(int argc, char **argv) {
    topFrame = (void**)__builtin_frame_address(0);
    setHandler(Trap::TrapSignal, onTrap);
    setHandler(SIGALRM, onTimer);
    setHandler(SIGSEGV, onFault);
    for(Function* f: functions) {
        f->setTrap();
    }
    setTimer(interval);
    int r = stabilizer_main(argc, argv);
    return r;
}
```

```
static void flush_icache(void* begin, size_t size) {
    uintptr_t p = (uintptr_t)begin & ~15UL;
    for (size_t i = 0; i < size; i += 16) {
        asm("icbi 0,%0" : : "r"(p));
        p += 16;
    }
}
```

```
void setTimer(int msec) {
    struct itimerval timer;

    timer.it_value.tv_sec = (msec - msec % 1000) / 1000;
    timer.it_value.tv_usec = 1000 * (msec % 1000);
    timer.it_interval.tv_sec = 0;
    timer.it_interval.tv_usec = 0;

    setitimer(ITIMER_REAL, &timer, 0);
}
```

Change Username

Changes layout



% sudo -u luke

Layout is Brittle

Layout biases measurement

Mytkowicz et al. (ASPLOS'09)

Can we eliminate the effect of layout?

Layout biases measurement

**Can we eliminate the
effect of layout?**

YES

with *randomization*

STABILIZER

Memory layout affects performance
makes performance evaluation difficult

Use randomization to eliminate the effect of layout
enables sound performance evaluation

Case Studies
evaluation of LLVM's optimizations with STABILIZER

STABILIZER

Memory layout affects performance
makes performance evaluation difficult

Use randomization to eliminate the effect of layout
enables sound performance evaluation

Case Studies
evaluation of LLVM's optimizations with STABILIZER

STABILIZER

randomizes layout

function addresses

stack frame sizes

heap allocations

Layout biases measurement

STABILIZER during
repeatedly randomizes layout[✓]

function addresses

stack frame sizes

heap allocations

Layout biases measurement

STABILIZER

repeatedly randomizes layout

function addresses

stack frame sizes

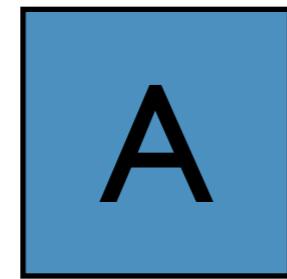
heap allocations

~~Layout biases measurement~~

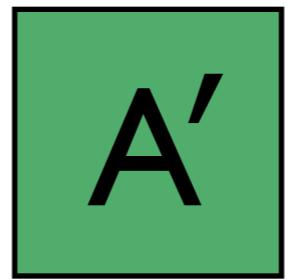
completely random layout

cannot bias results

Sound Performance Evaluation

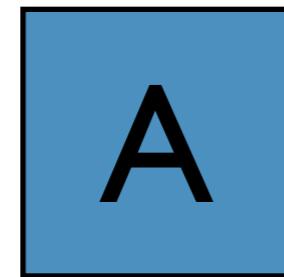


×30

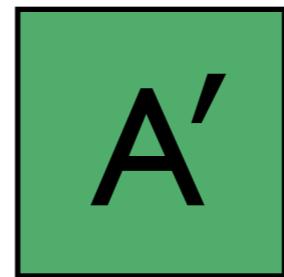


×30

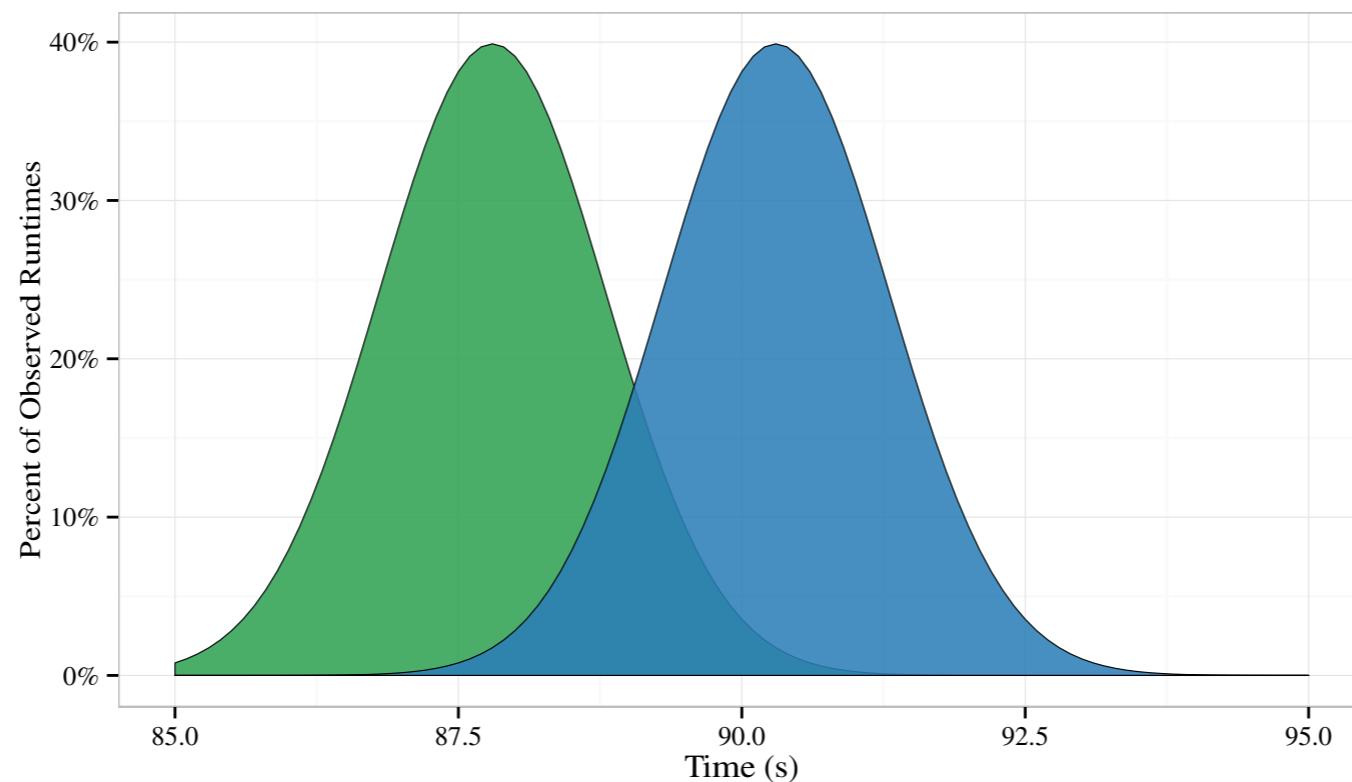
Sound Performance Evaluation

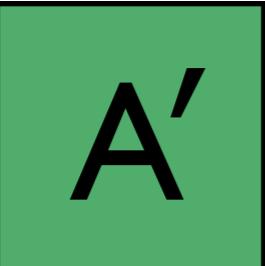
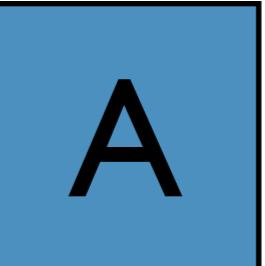


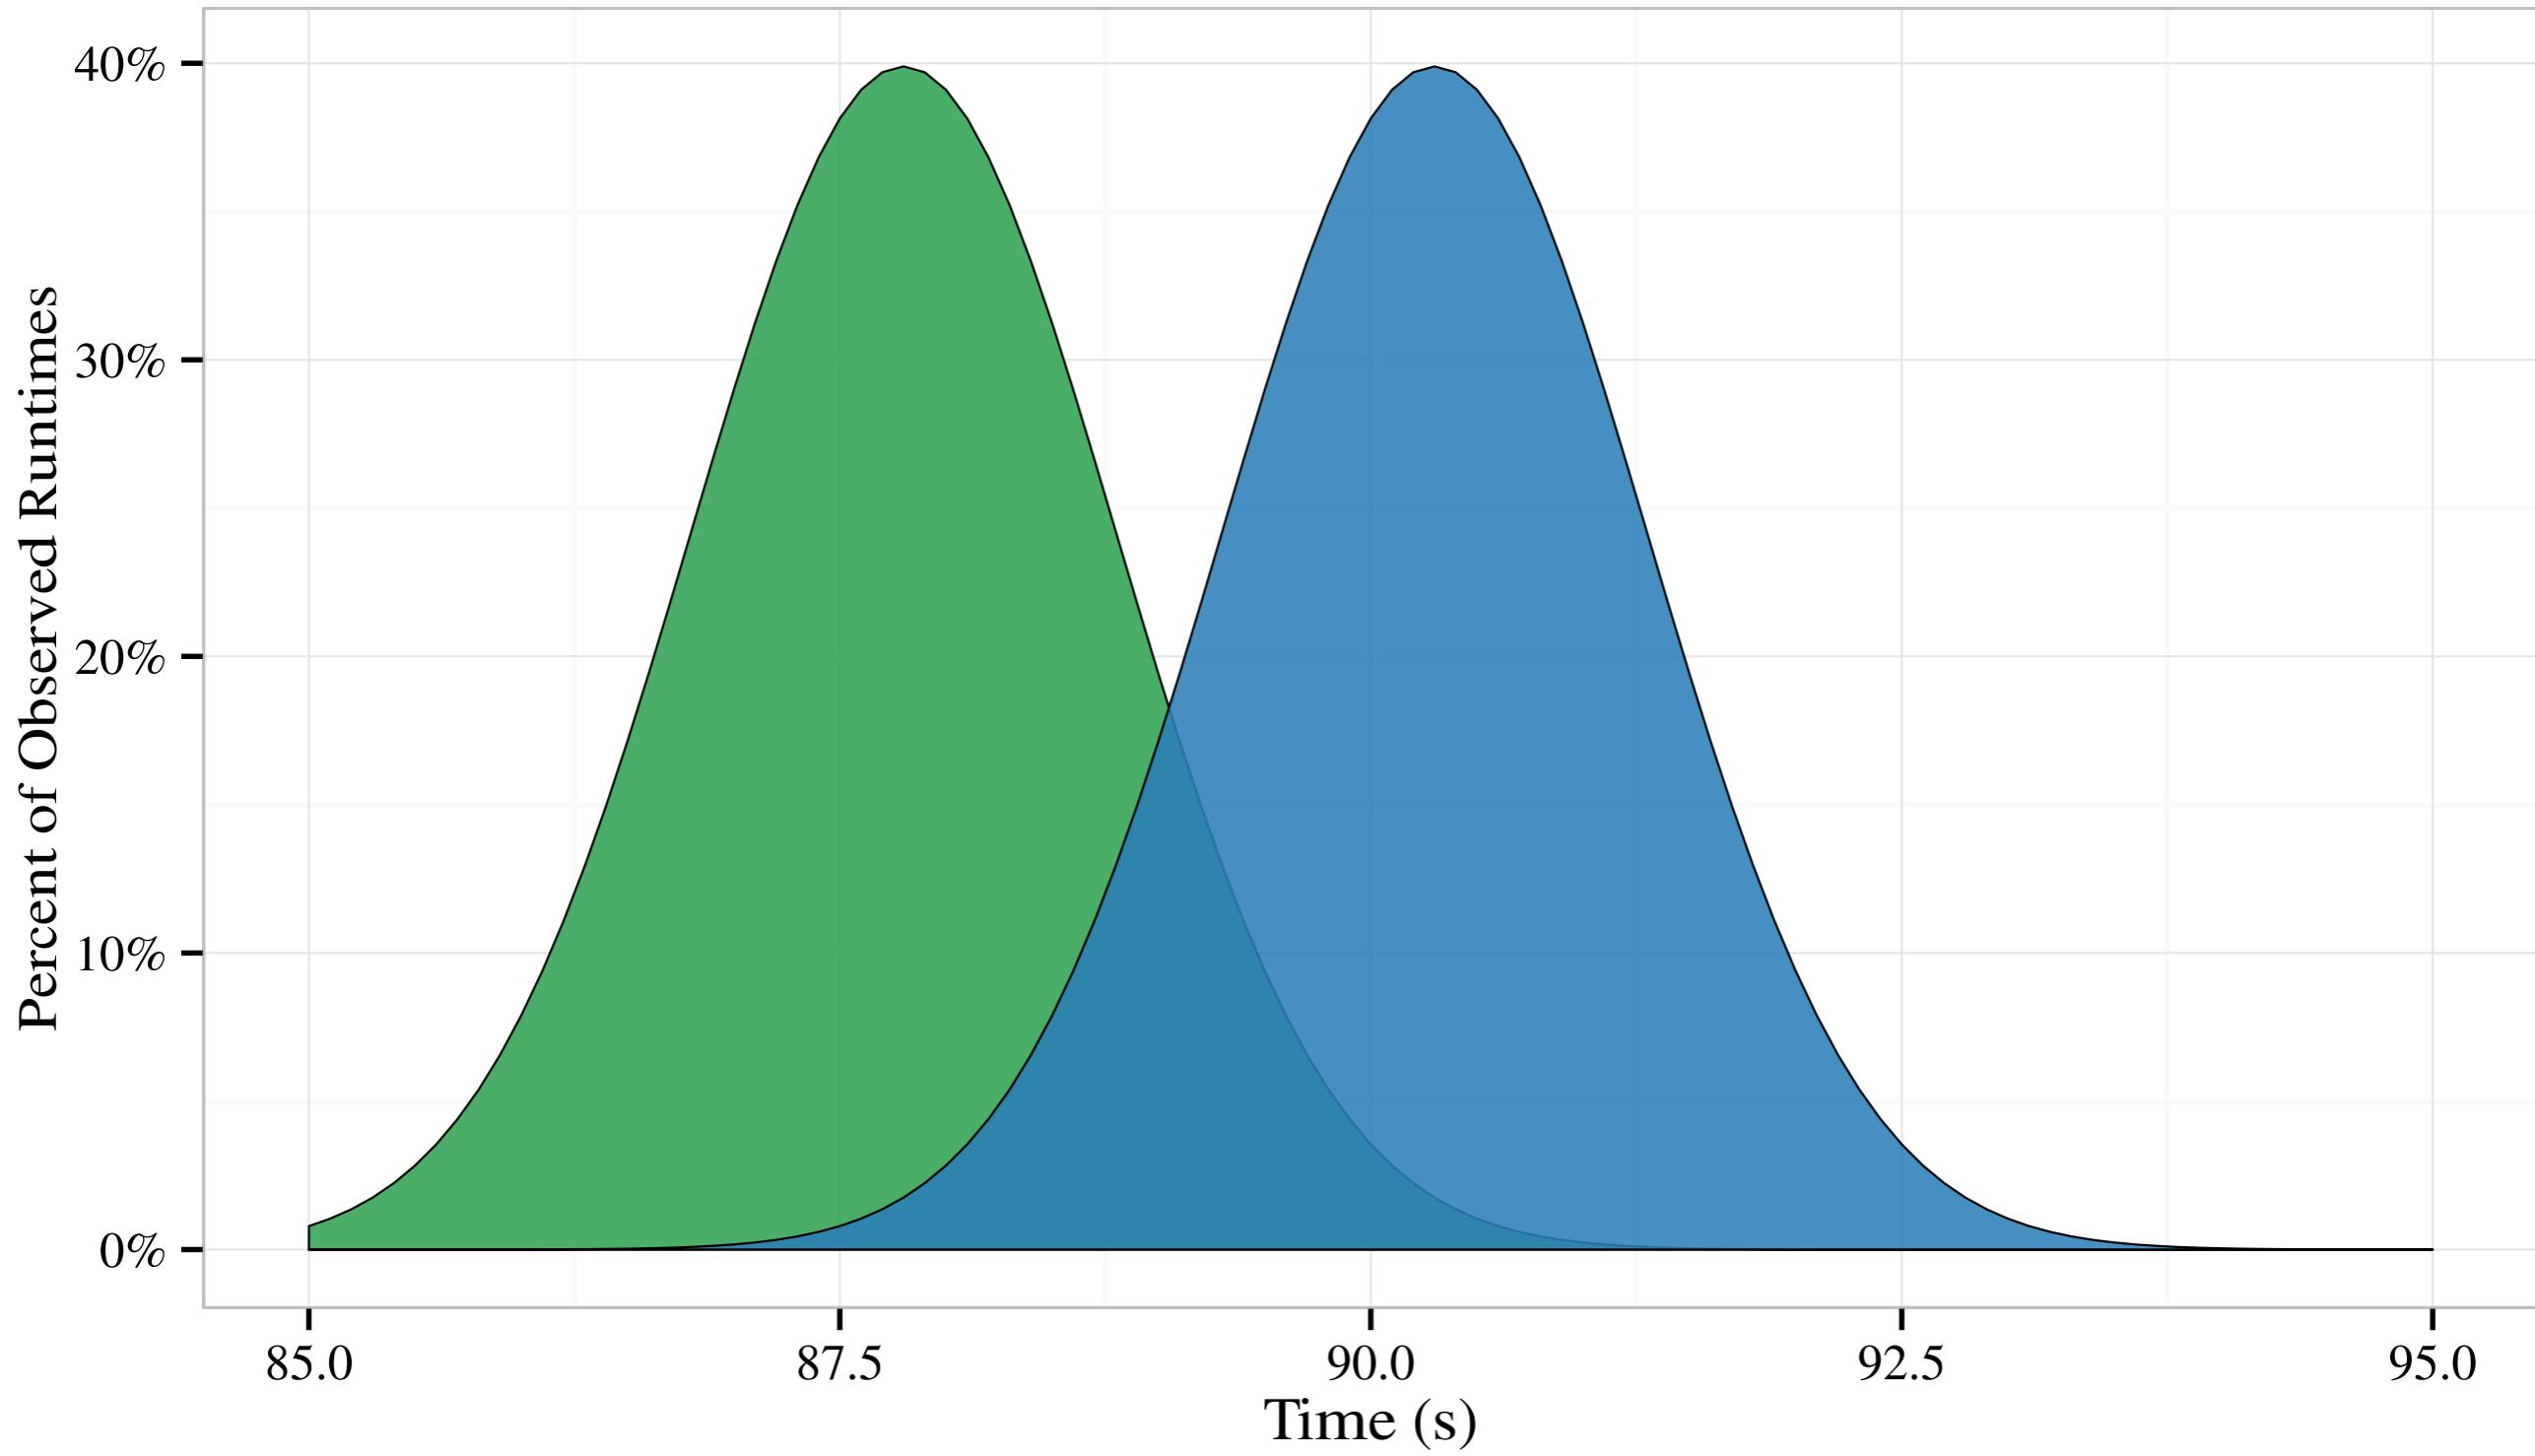
$\times 30$

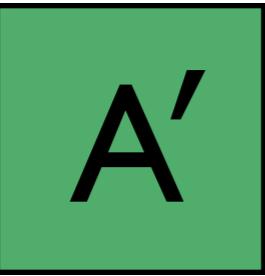
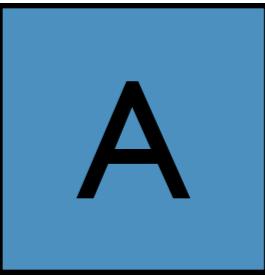


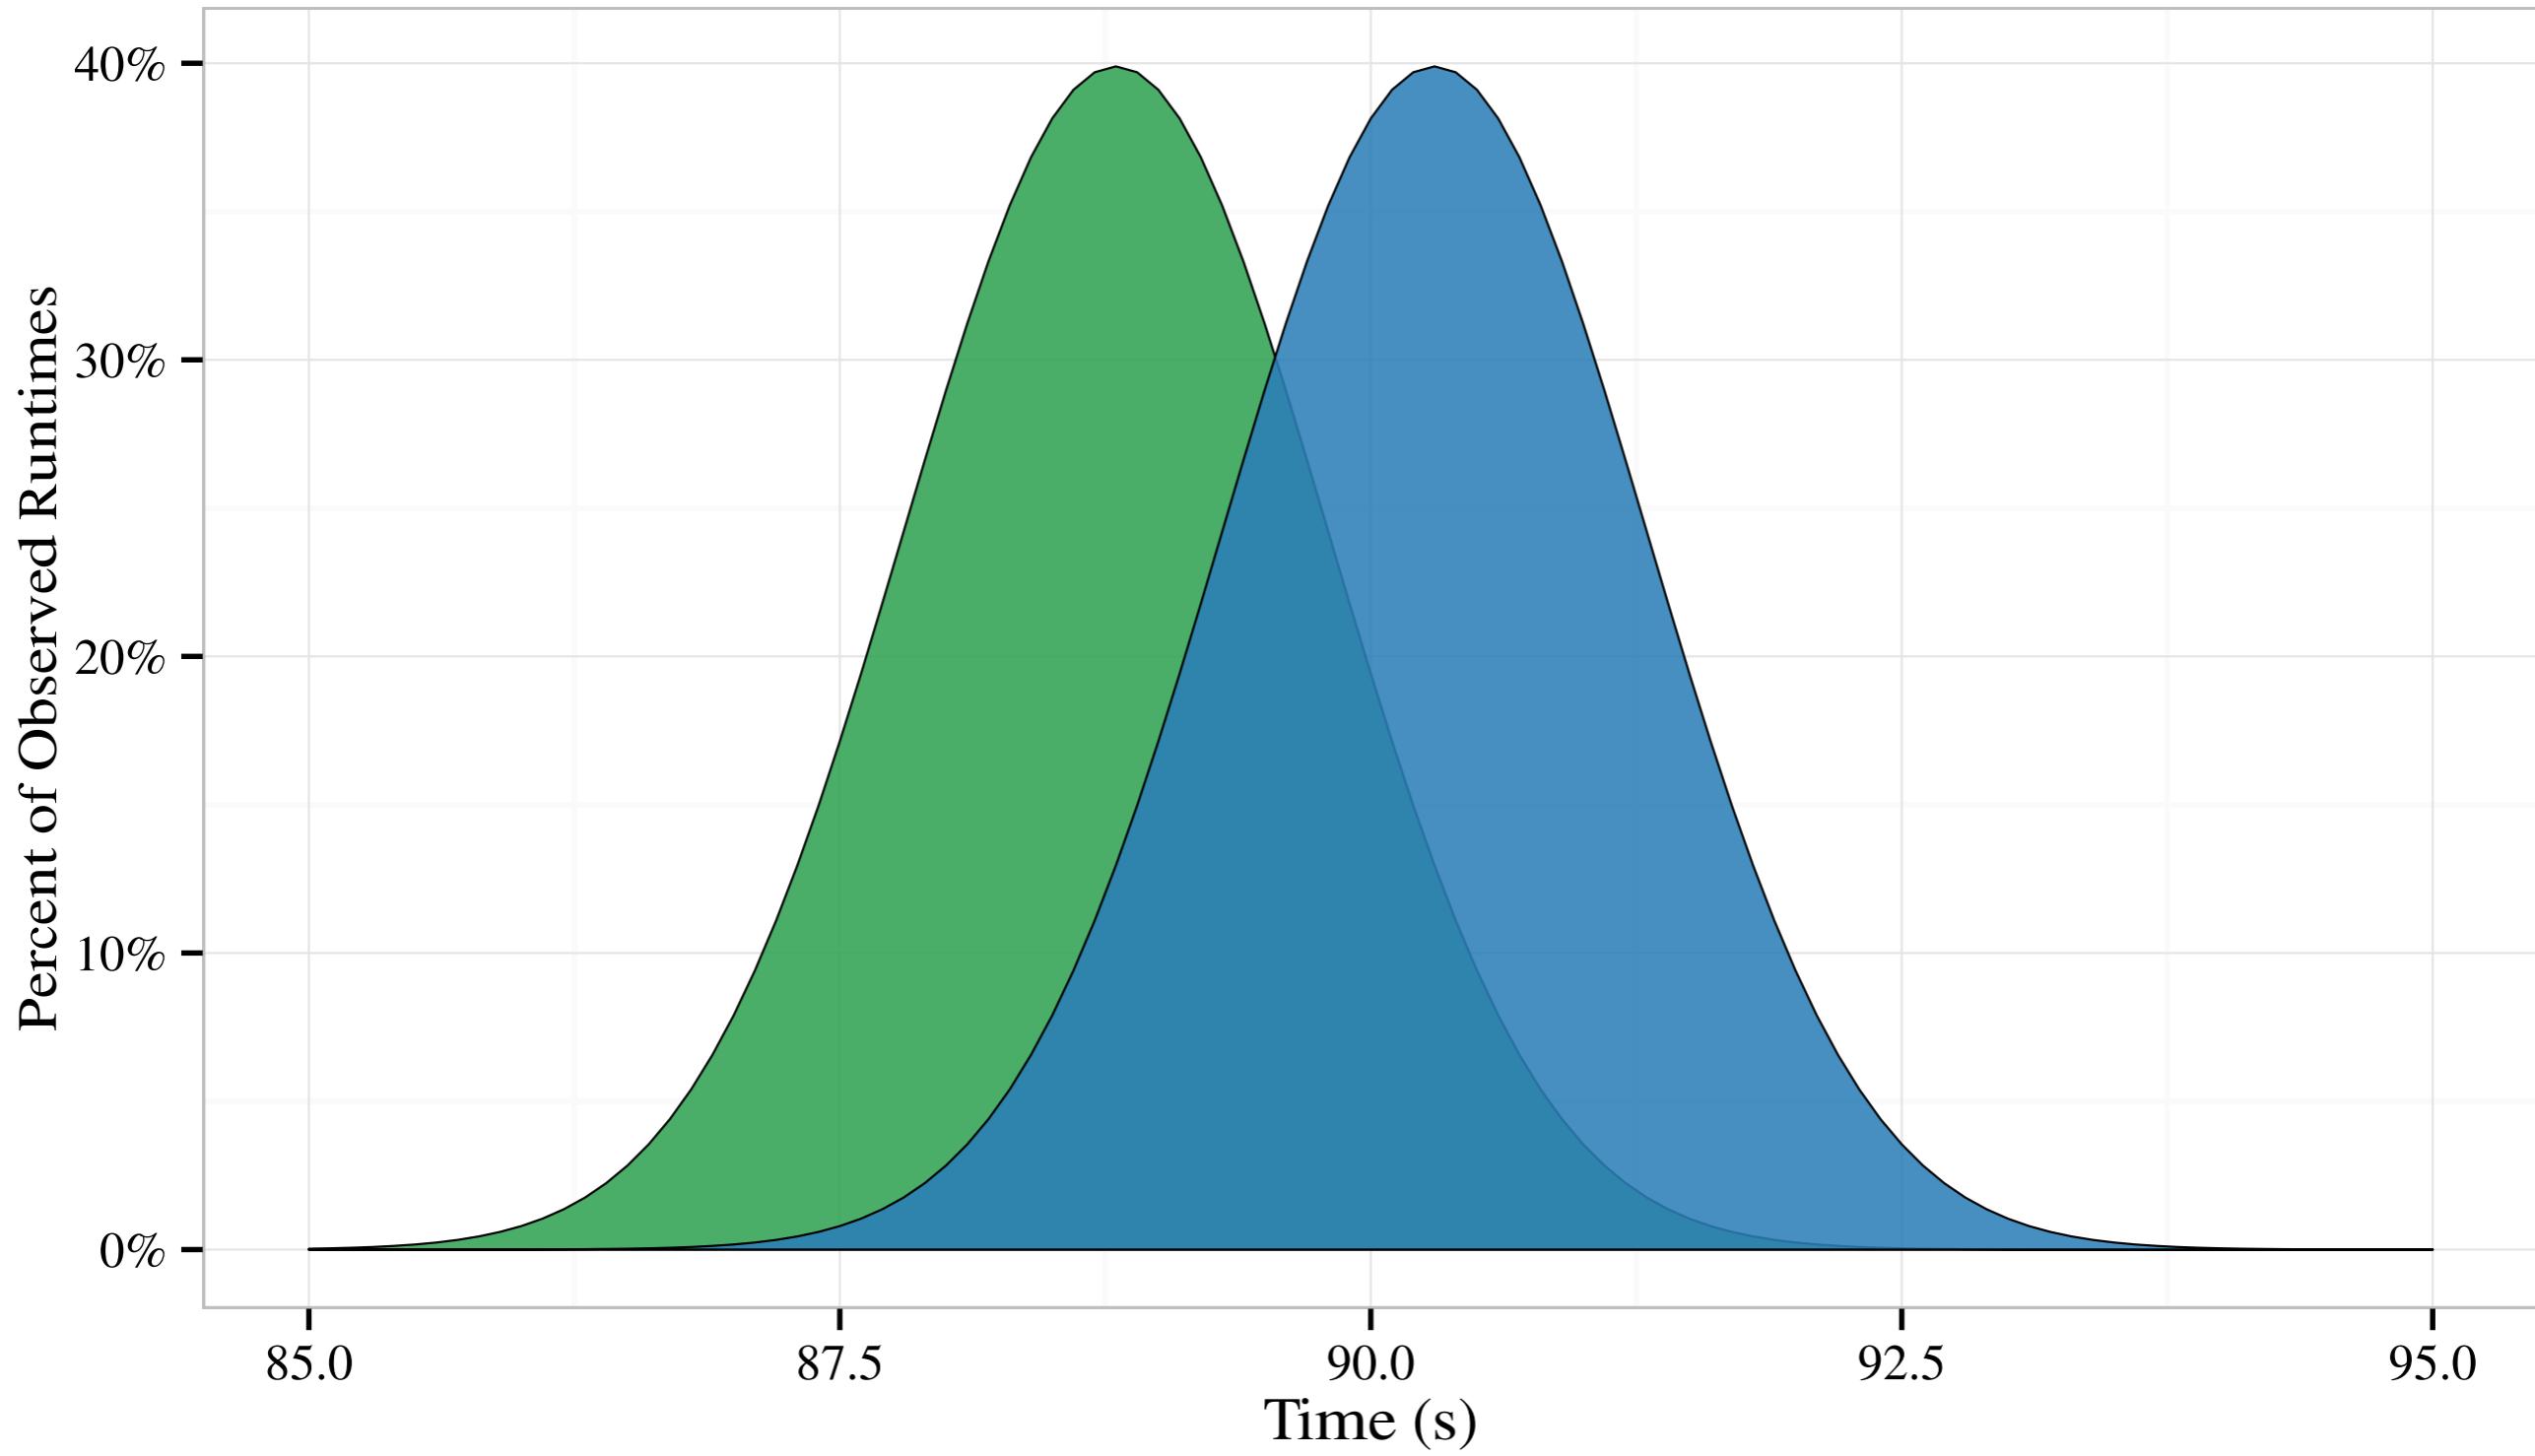
$\times 30$

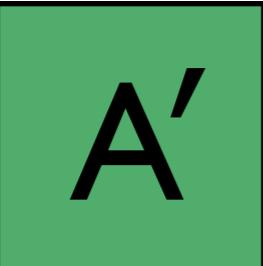
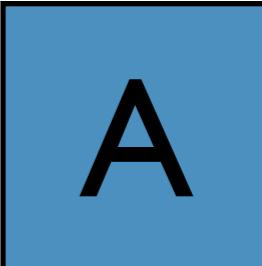


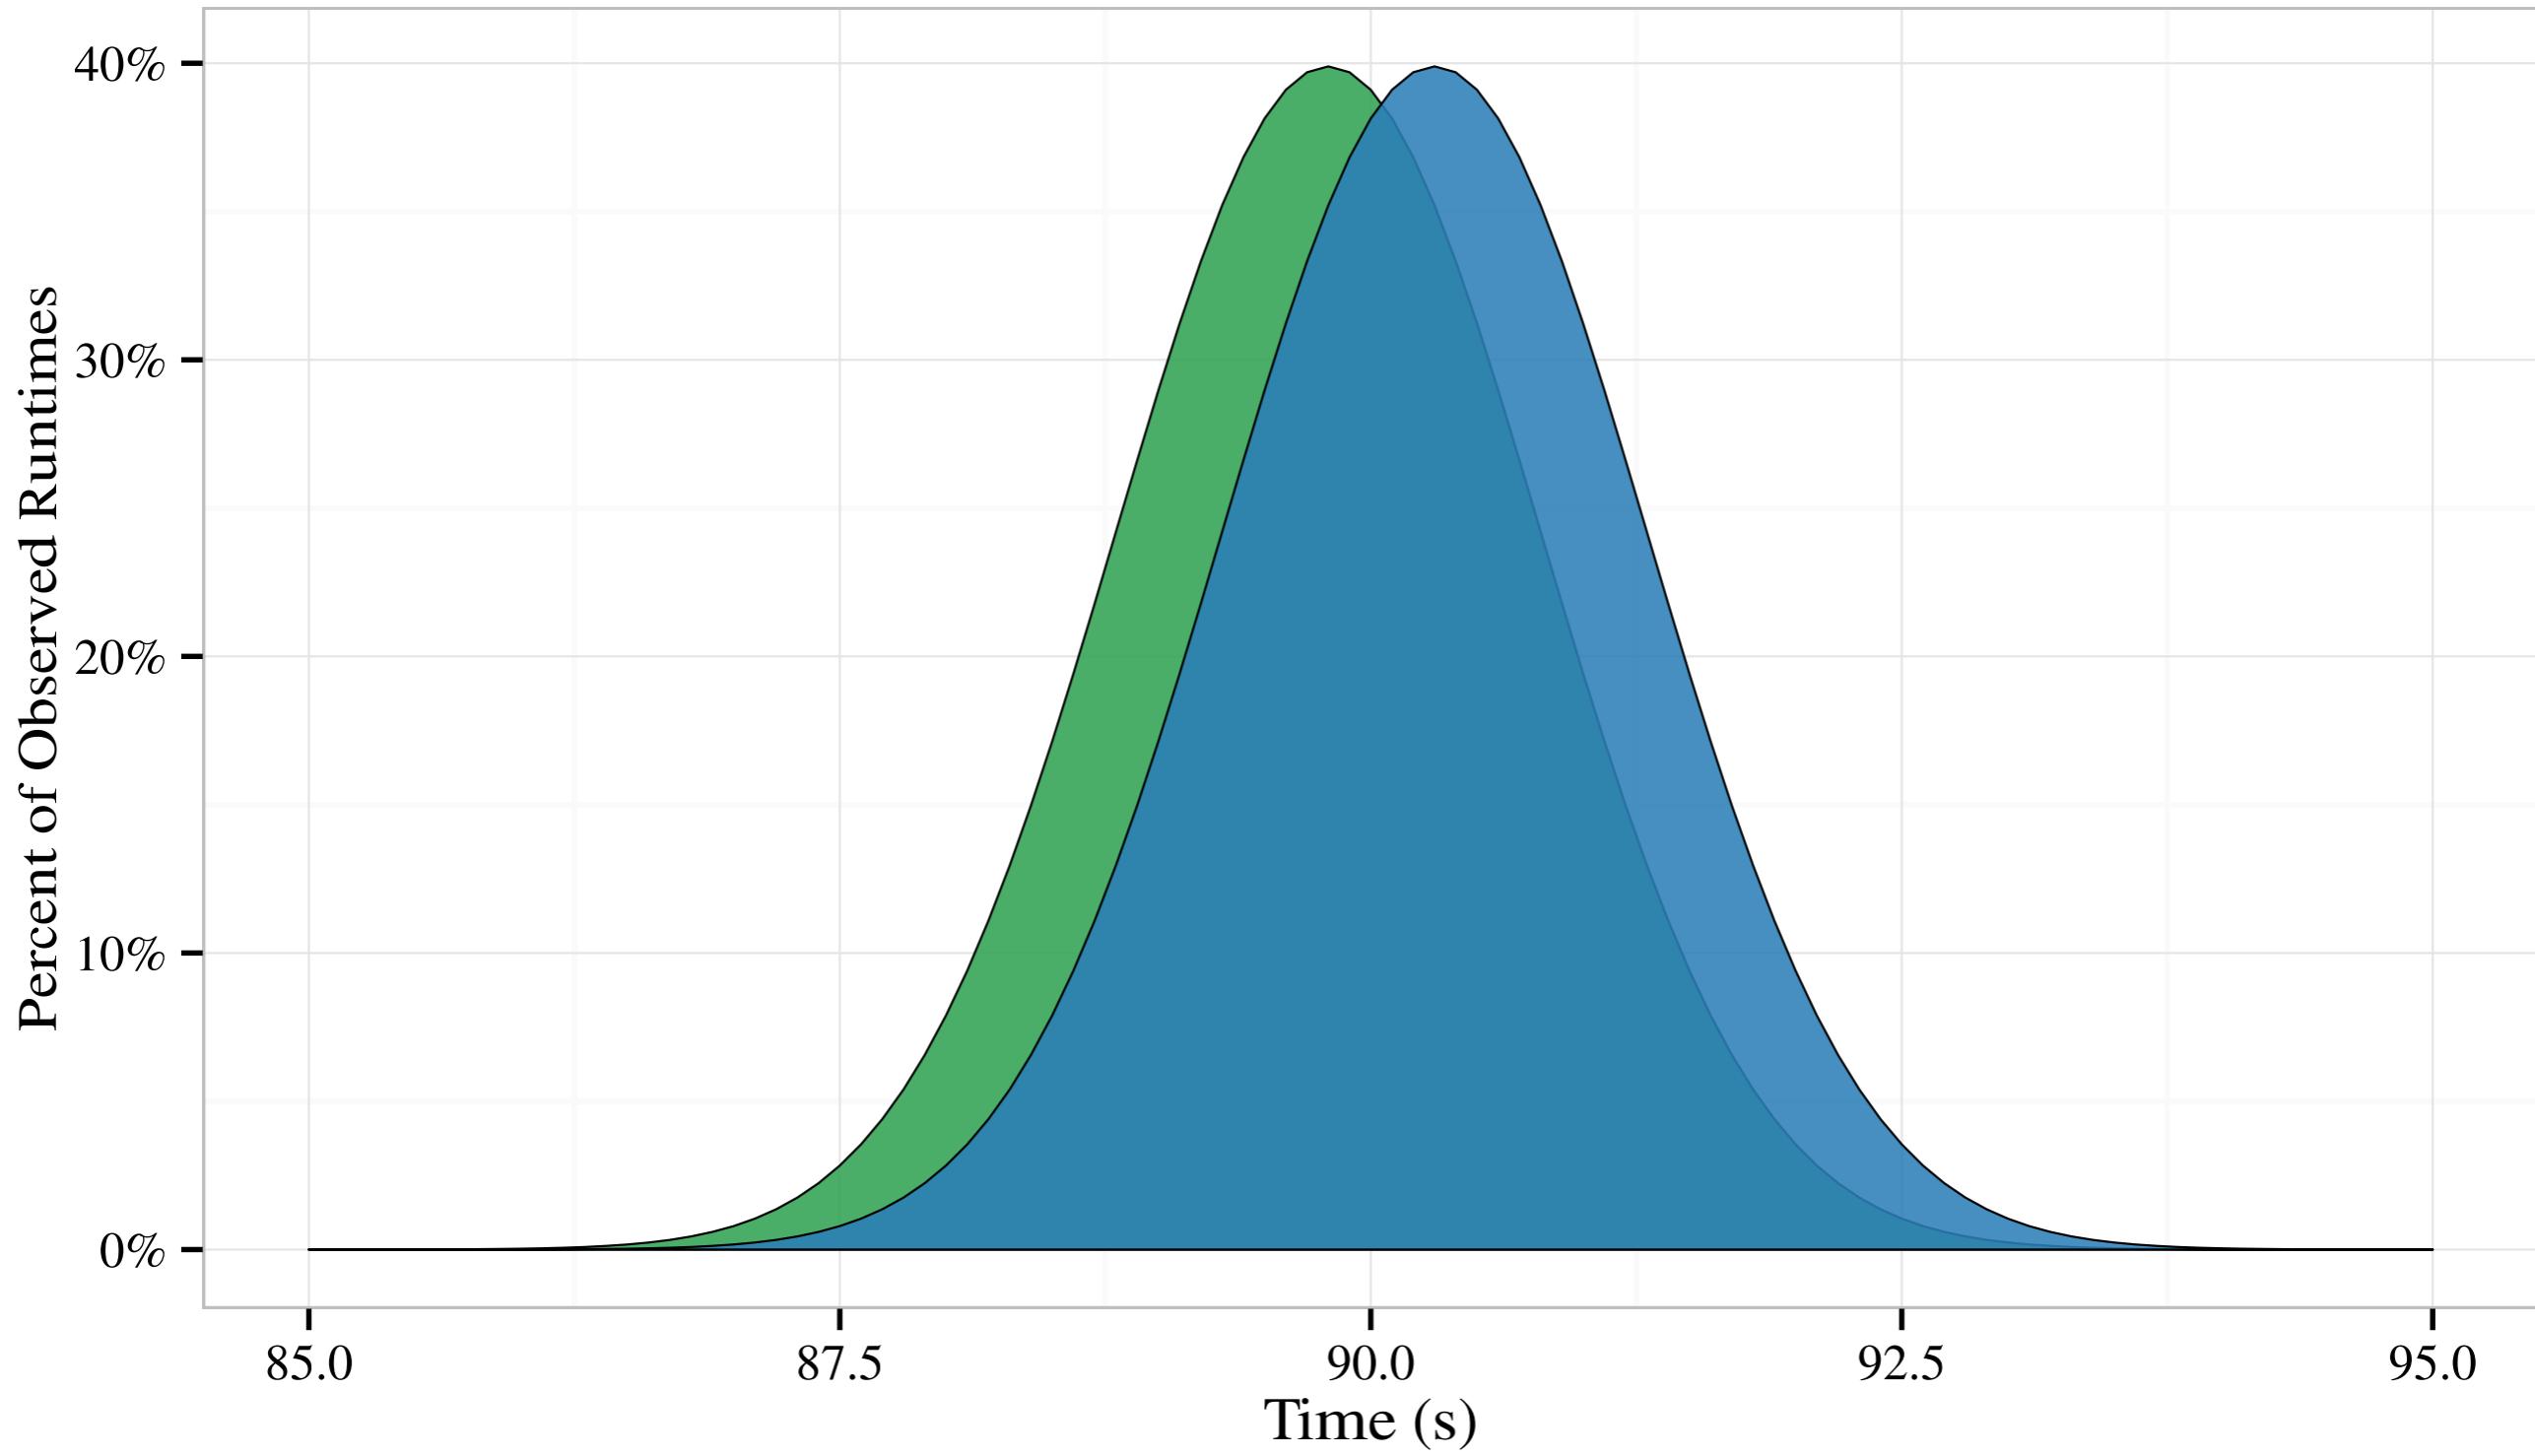
Is  faster than ?

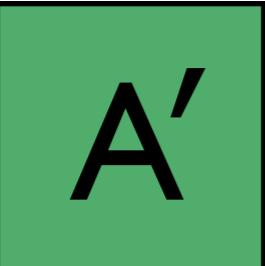
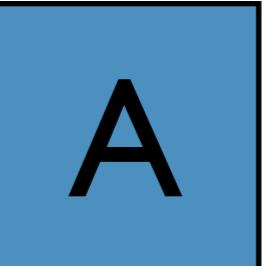


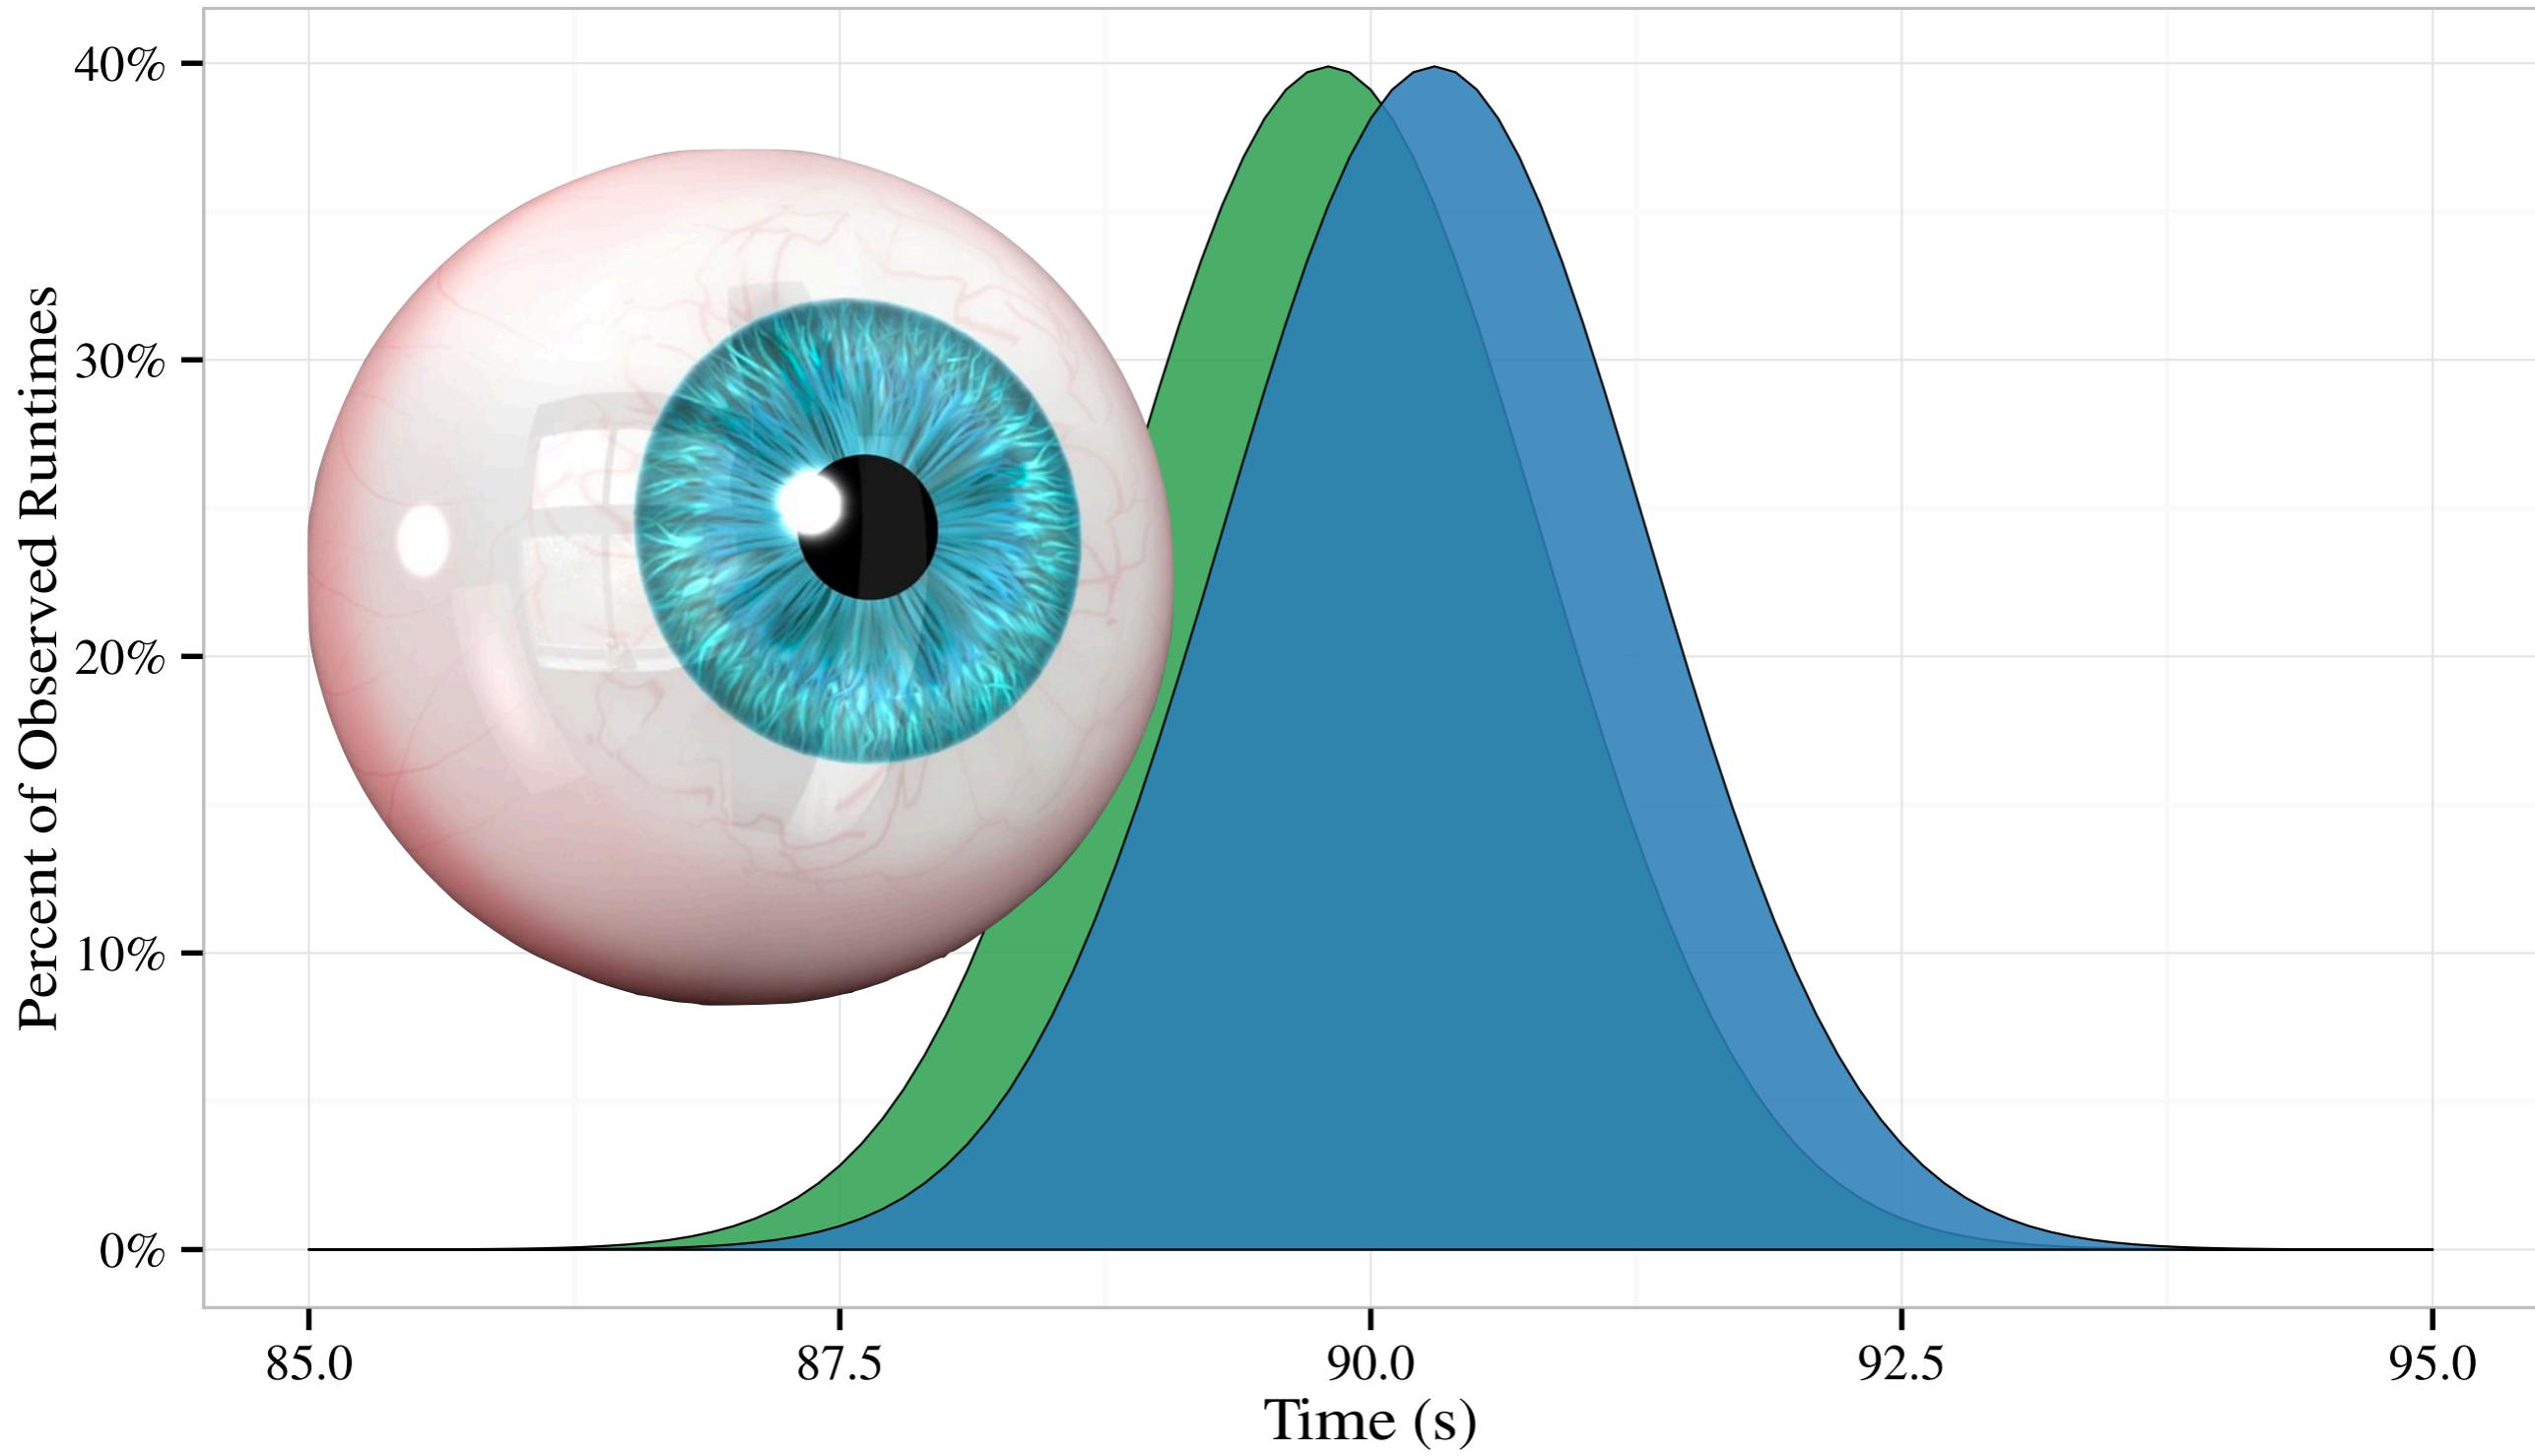
Is  faster than ?



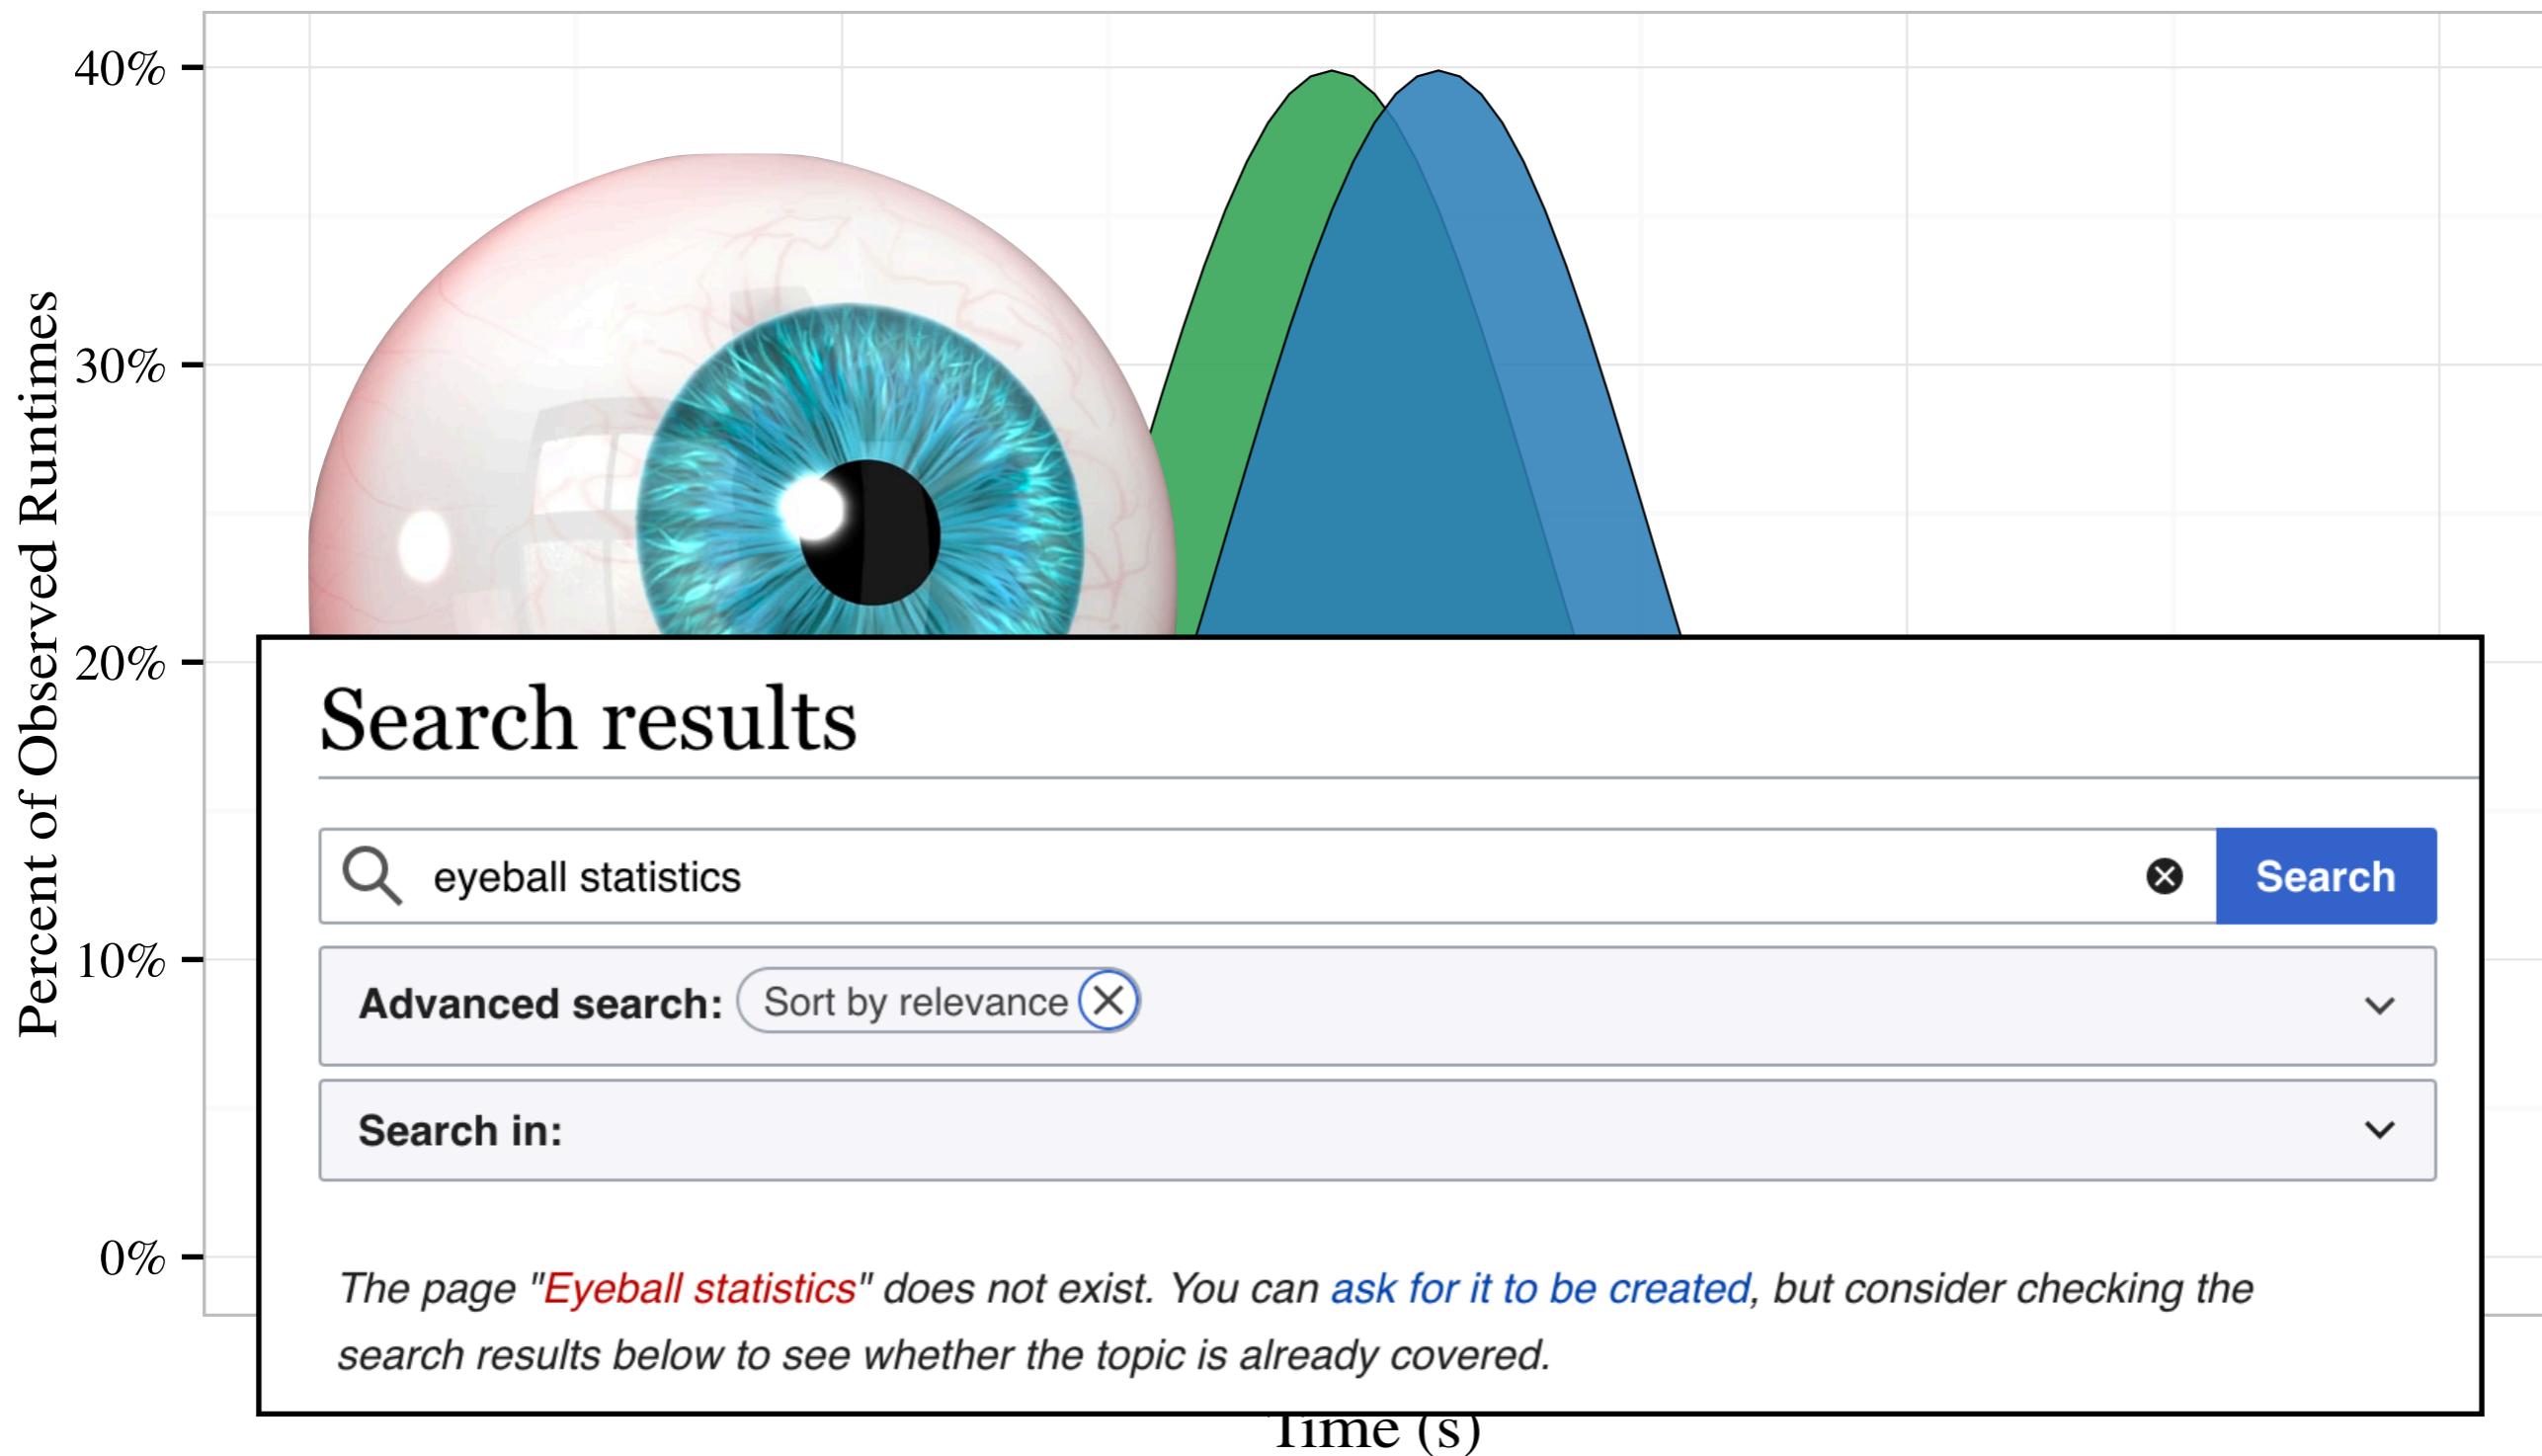
Is  faster than ?



Is  faster than ?



Is A' faster than A?

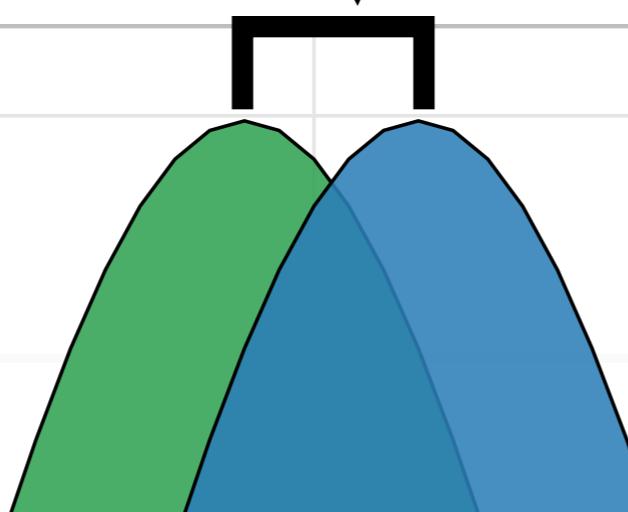


The Statistical Approach

null hypothesis significance testing

If $A' = A$

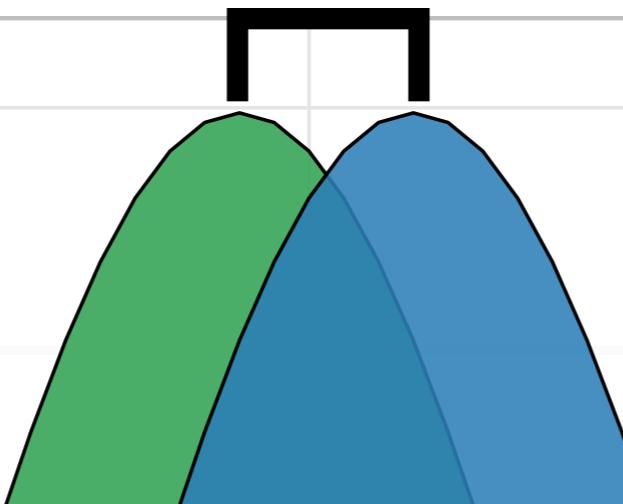
what is the probability of measuring
a speedup this large by chance?



If $A' = A$

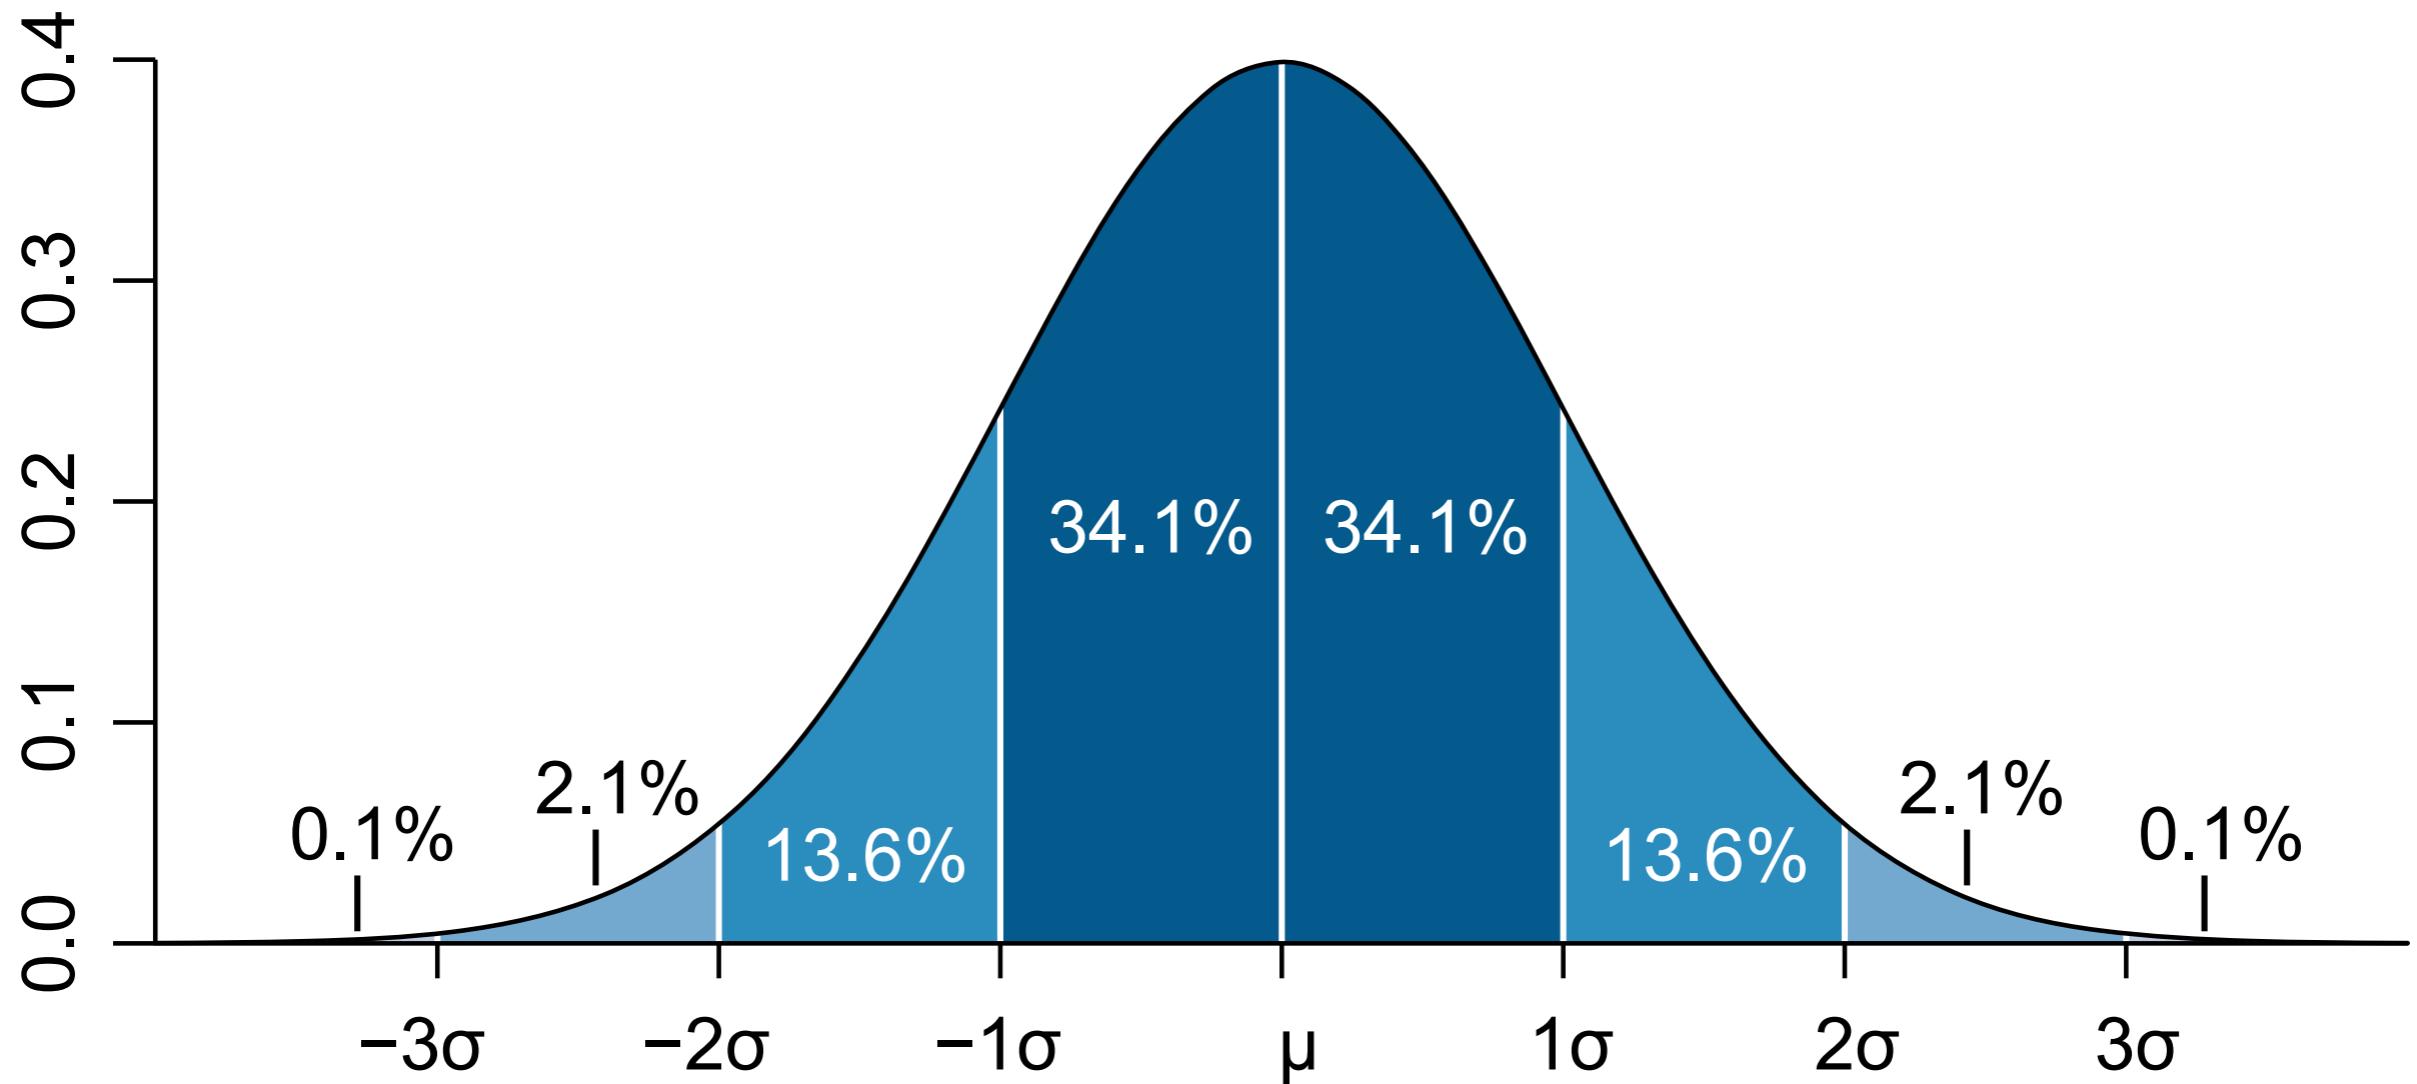
**easy to compute for
the normal distribution**

what is the probability of measuring
a speedup this large by chance?



If $A' = A$

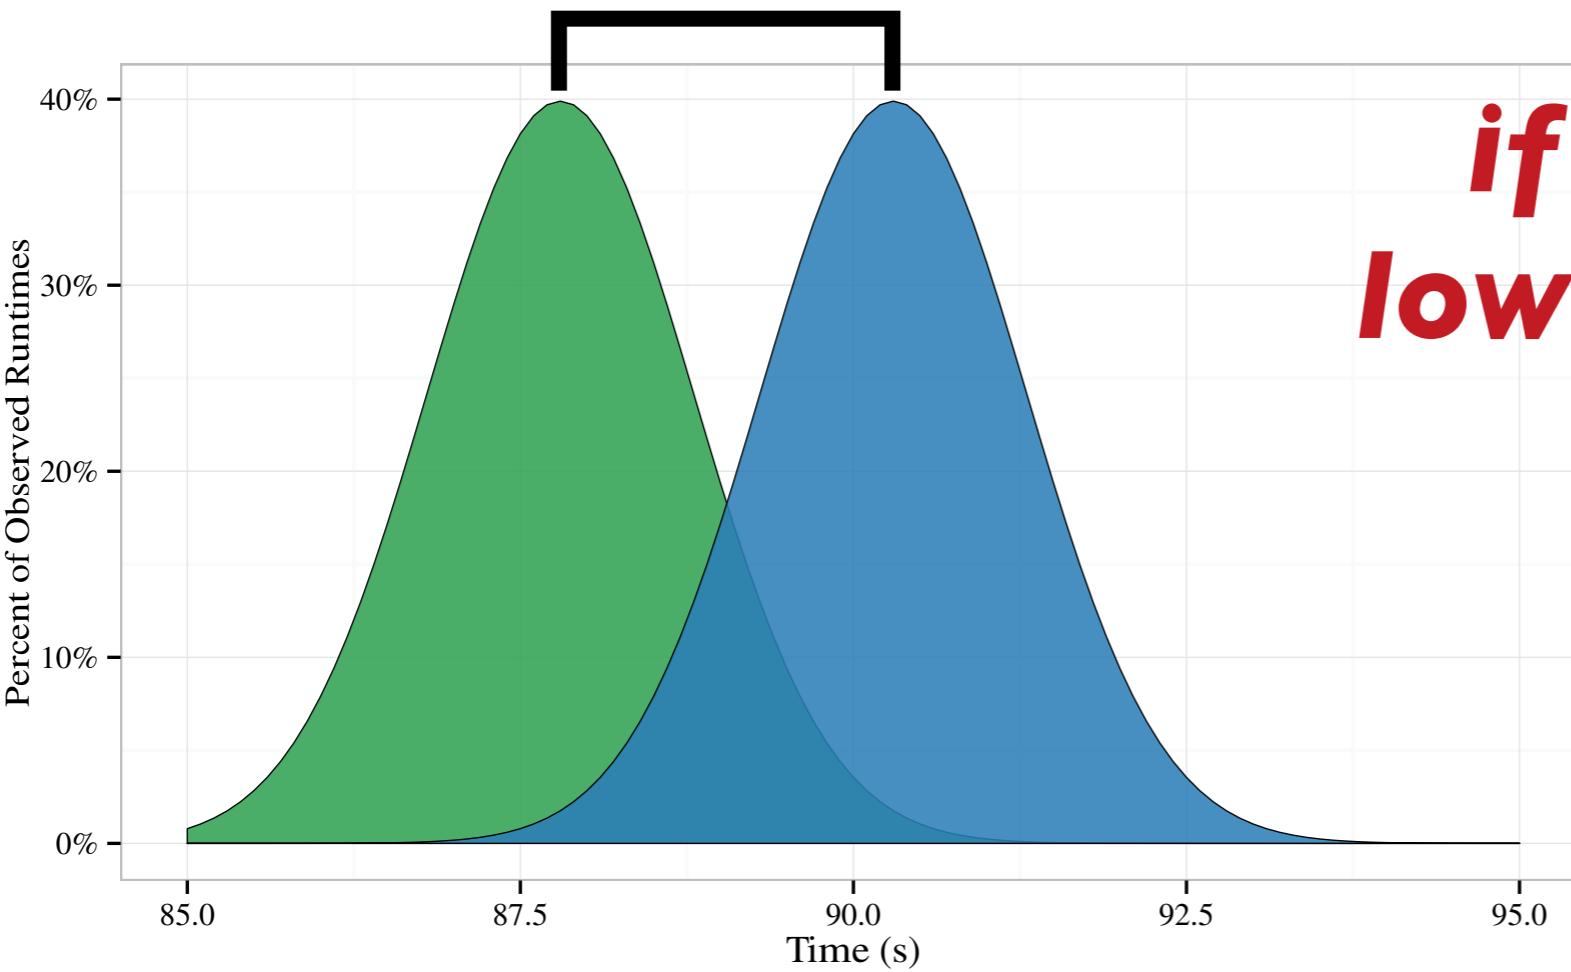
**easy to compute for
the normal distribution**



STABILIZER

repeatedly randomizes layout

what is the probability of measuring
a speedup this large by chance?



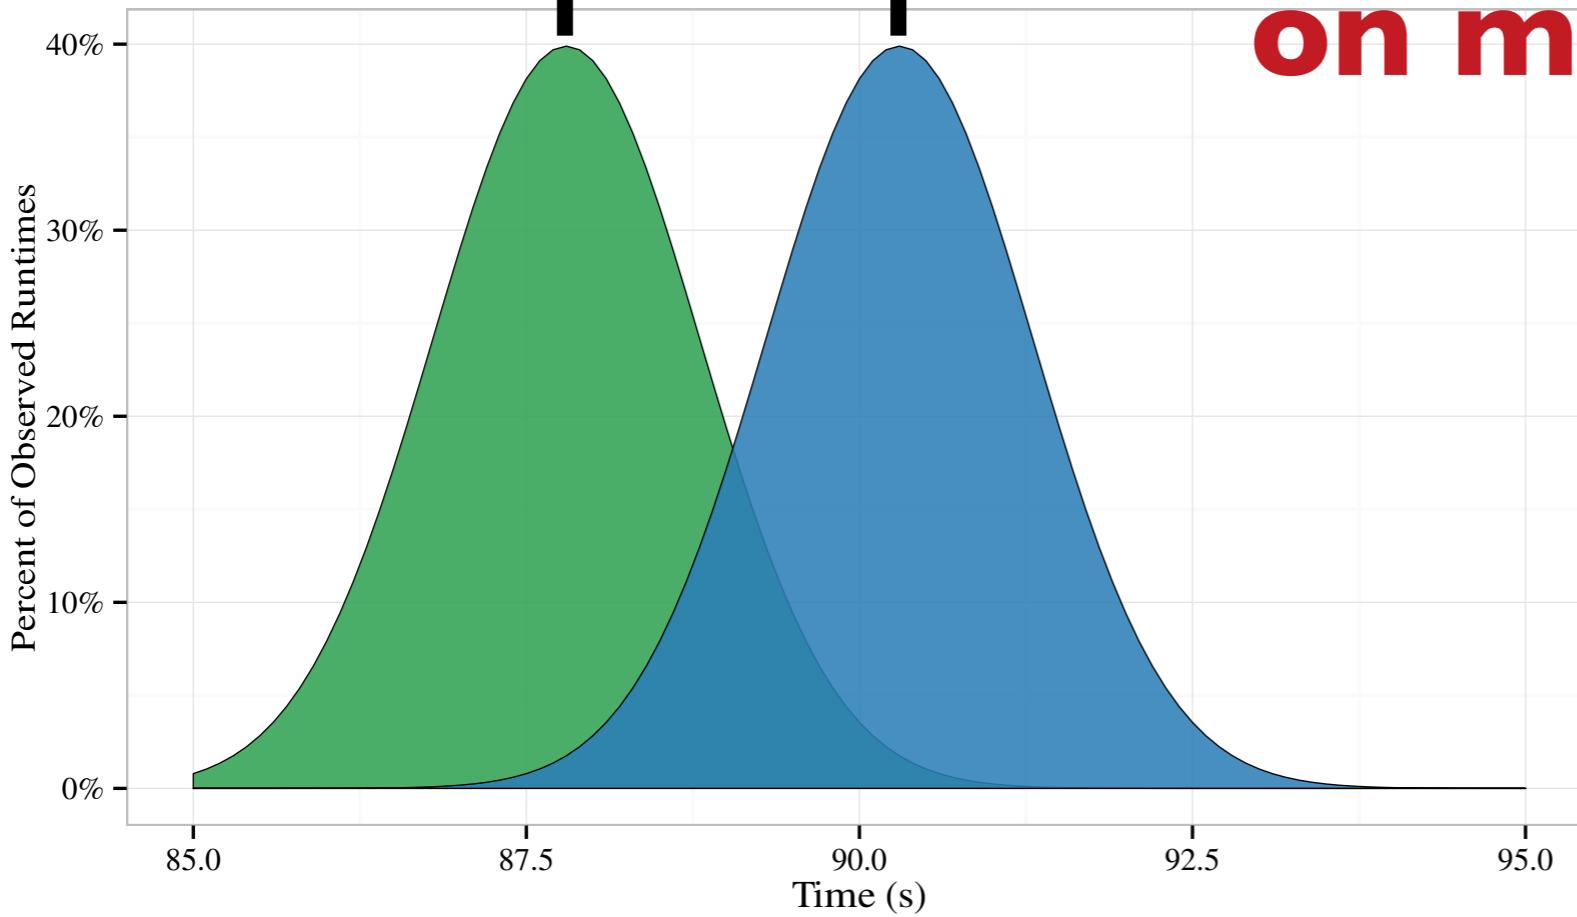
*if there is a
low probability*

STABILIZER

repeatedly randomizes layout

this speedup is real

**not due to the effect
on memory layout**



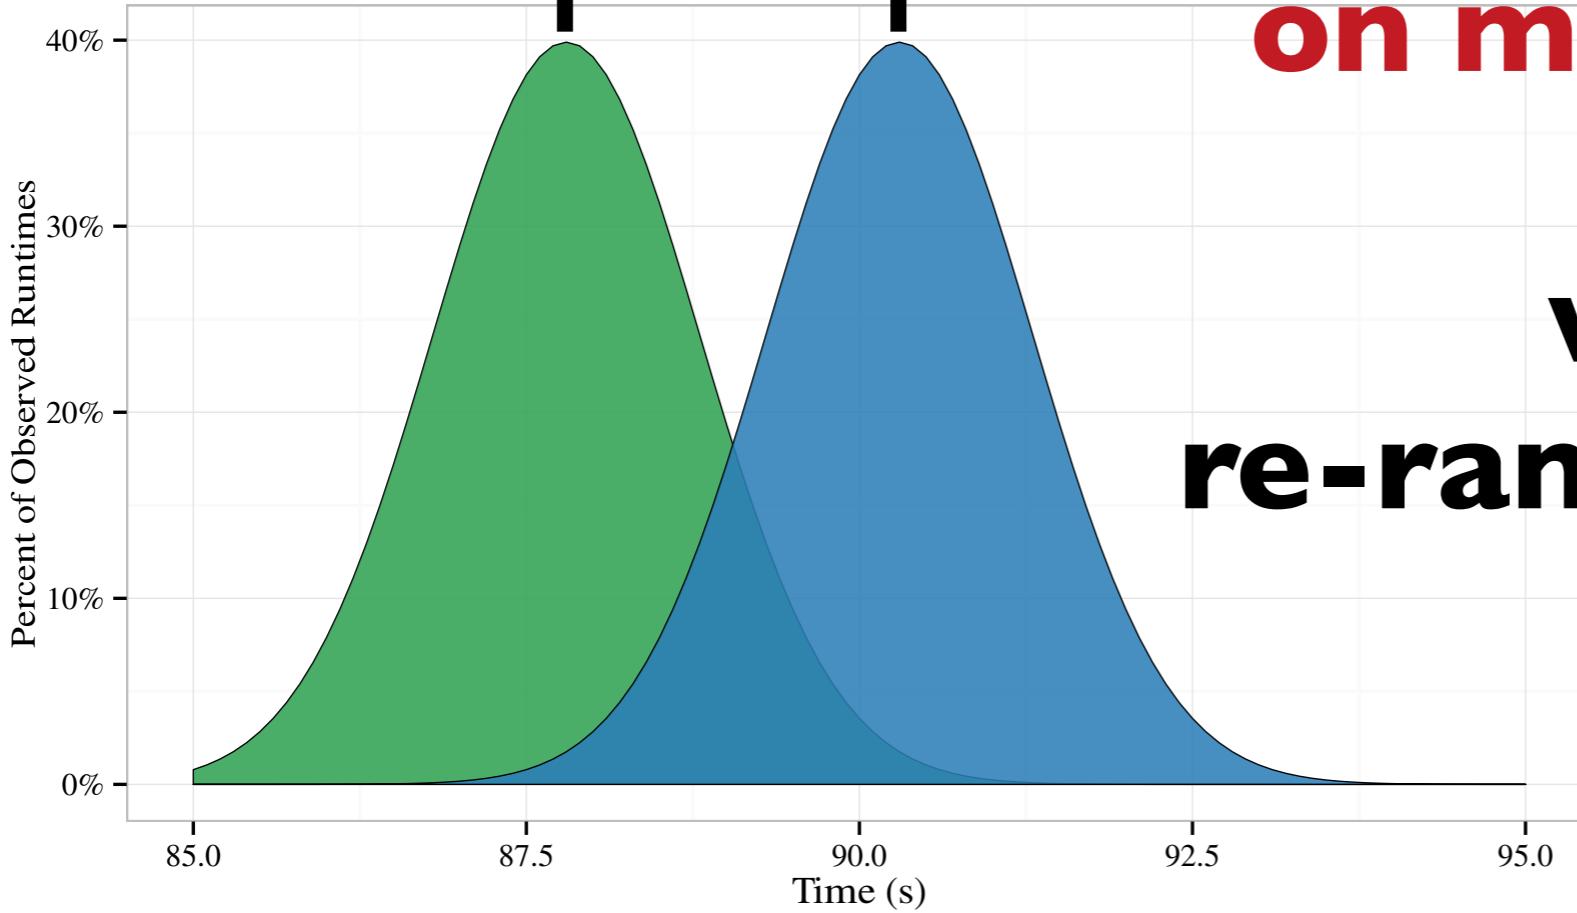
STABILIZER

repeatedly randomizes layout

this speedup is real

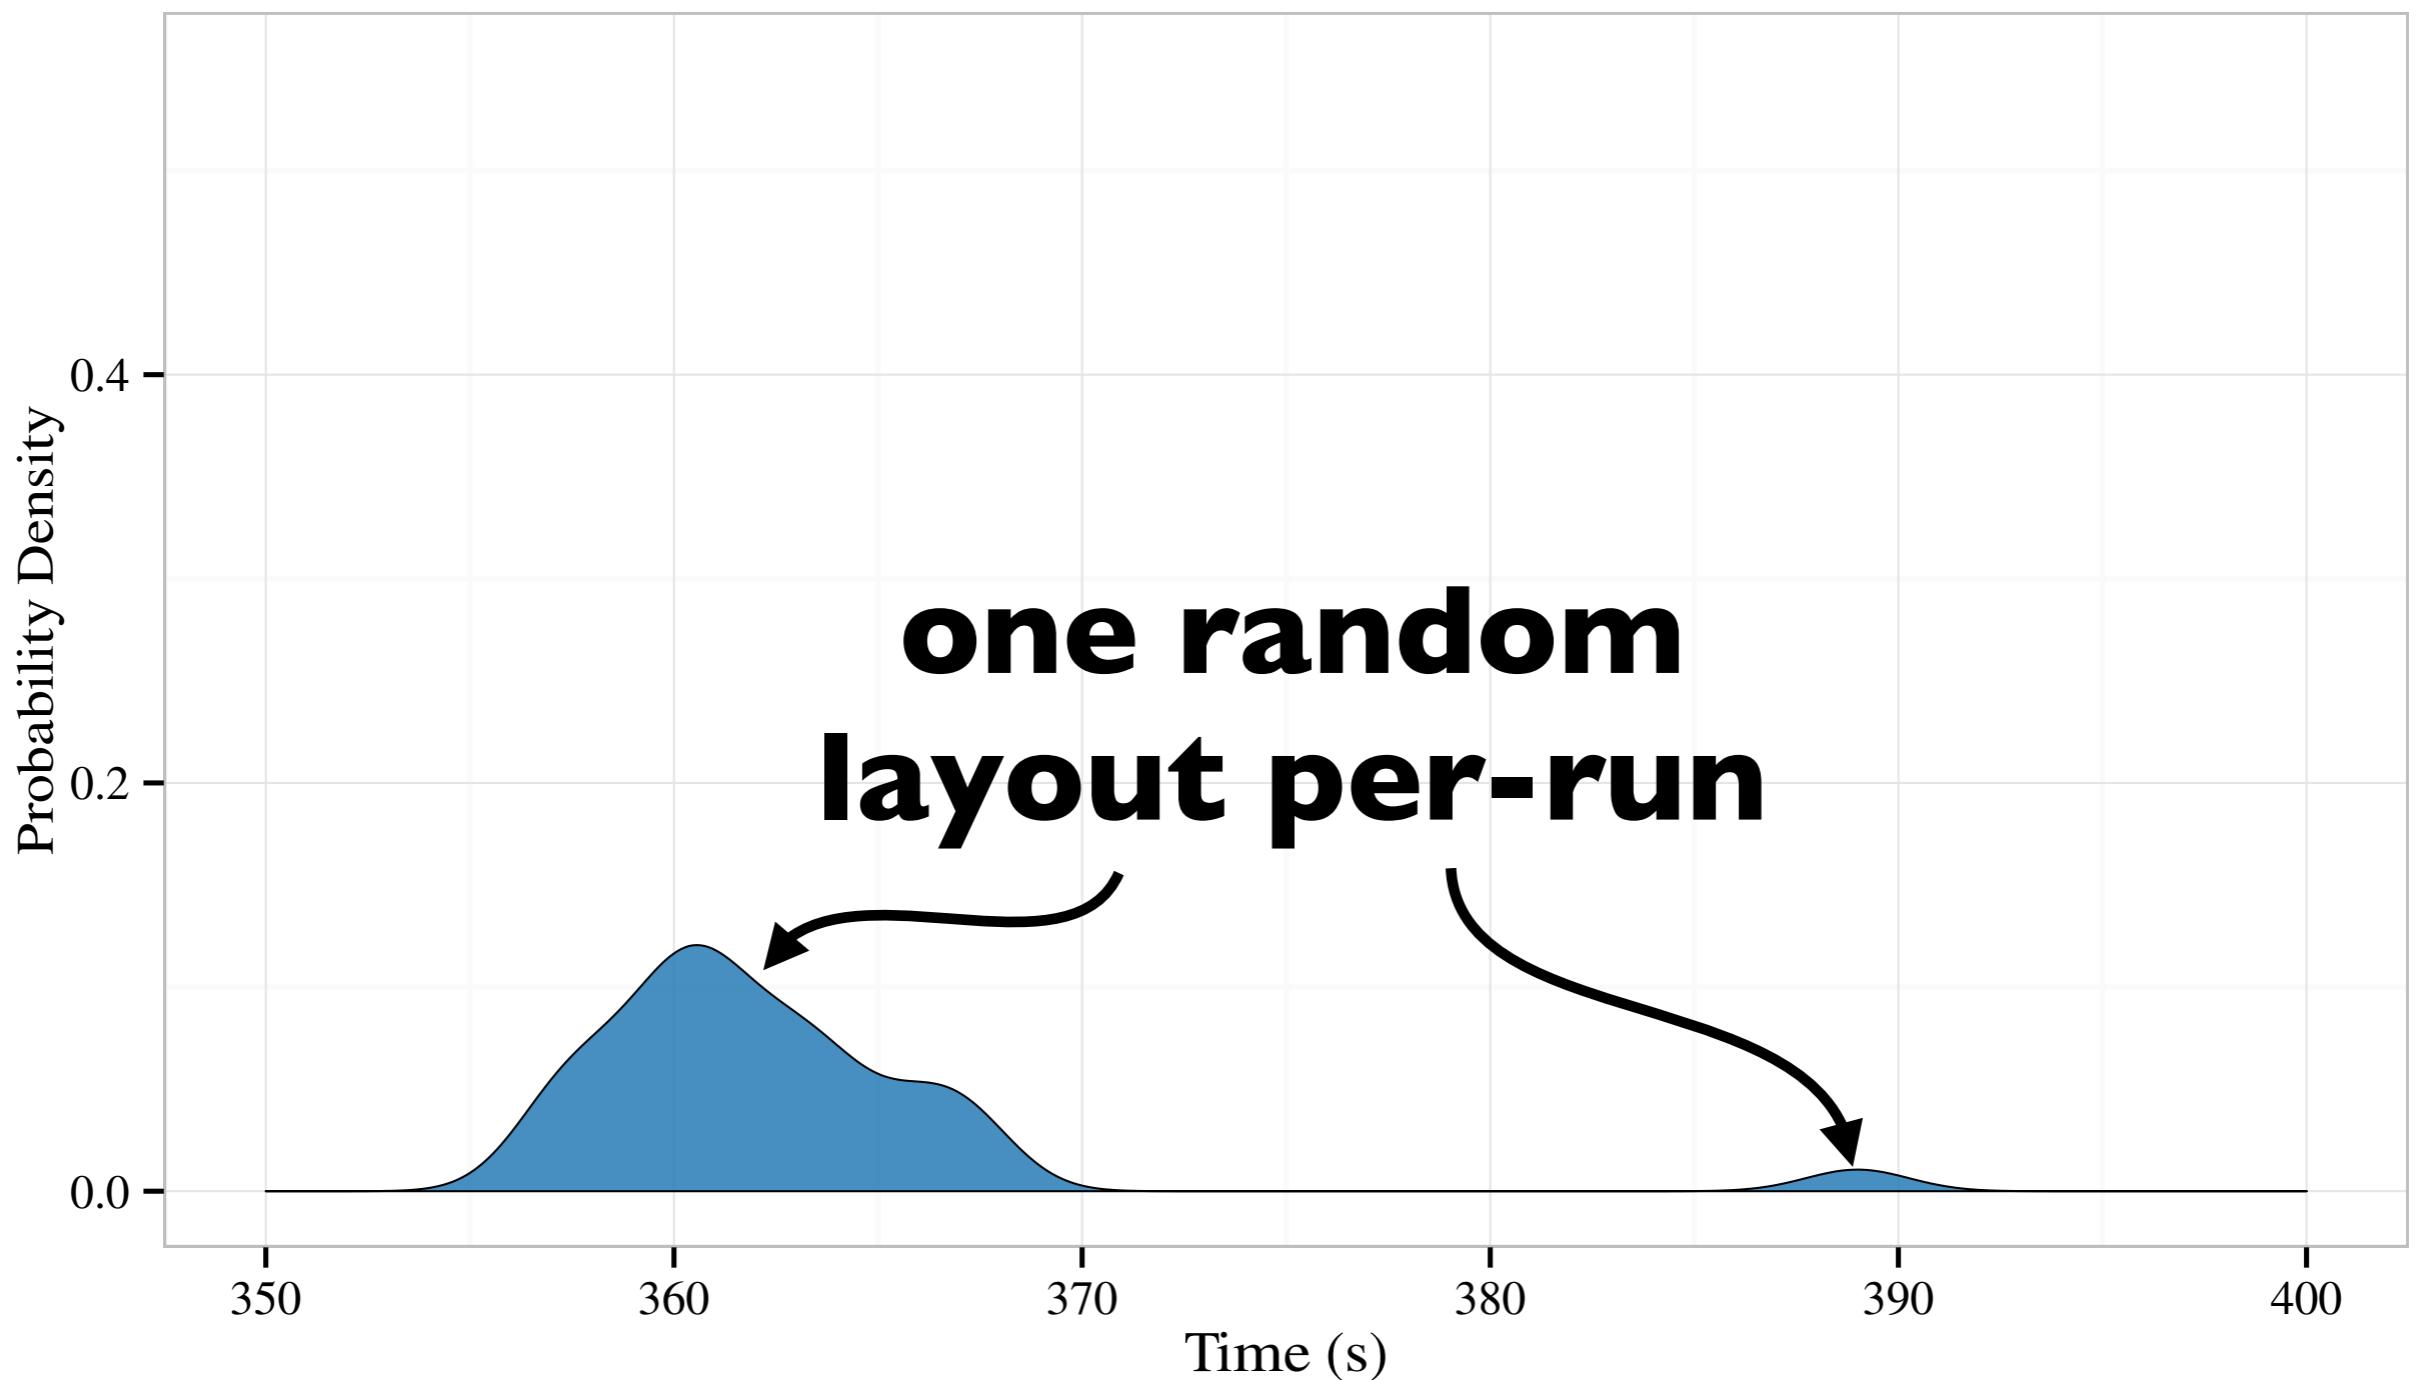
**not due to the effect
on memory layout**

**what does
re-randomization do?**



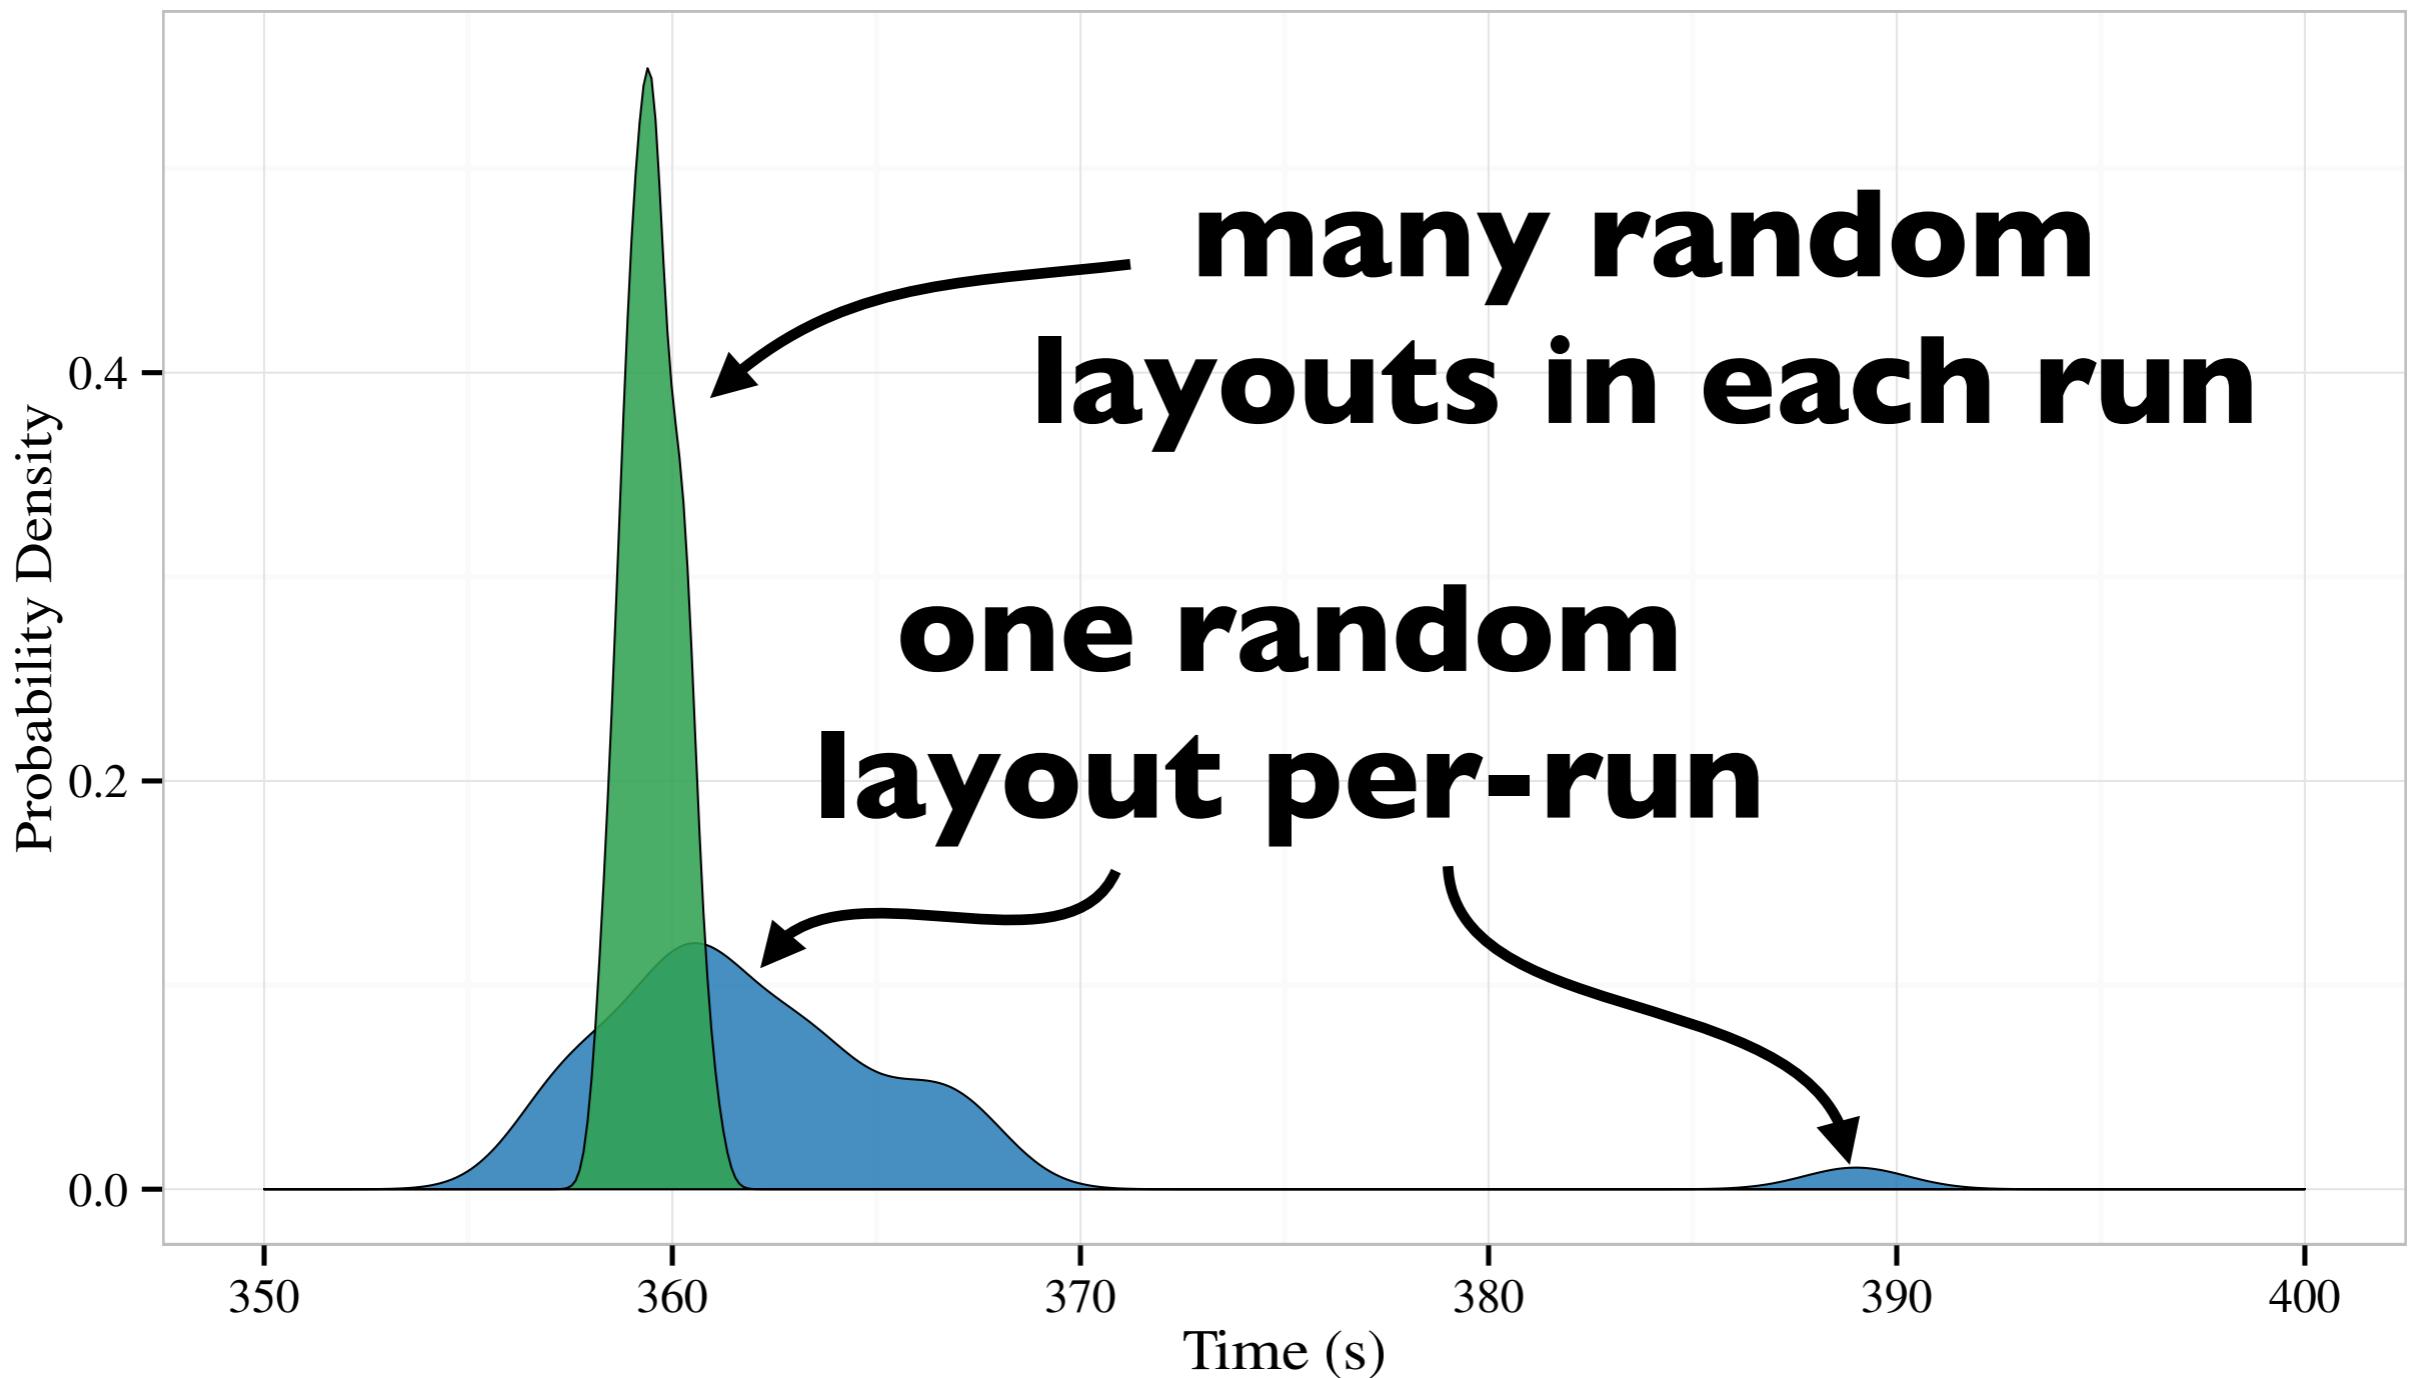
STABILIZER

repeatedly randomizes layout



STABILIZER

repeatedly randomizes layout



STABILIZER

repeatedly randomizes layout

**STABILIZER generates a new
random layout every $\frac{1}{2}$ second**

**Total execution time is
the sum of all periods**

STABILIZER

repeatedly randomizes layout

**STABILIZER generates a new
random layout every $\frac{1}{2}$ second**

**Total execution time is
the sum of all periods**

The sum

STABILIZER

repeatedly randomizes layout

**STABILIZER generates a new
random layout every $\frac{1}{2}$ second**

**Total execution time is
the sum of all periods**

The sum of a sufficient number

STABILIZER

repeatedly randomizes layout

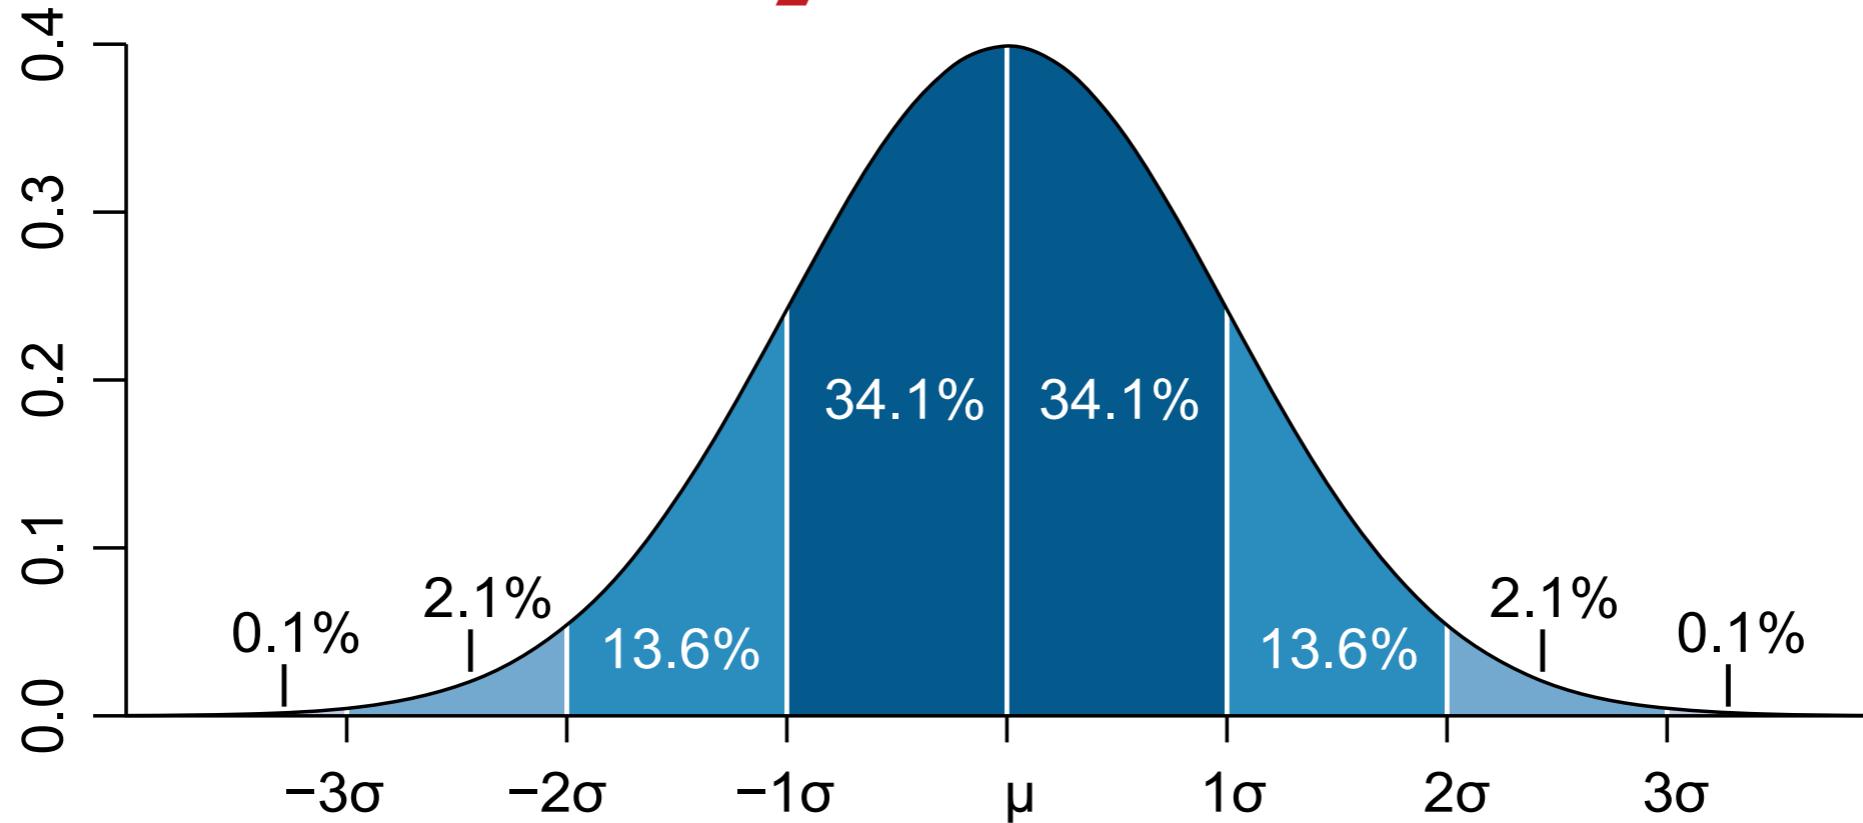
**STABILIZER generates a new
random layout every $\frac{1}{2}$ second**

**Total execution time is
the sum of all periods**

*The sum of a sufficient number of
independent, identically distributed random
variables*

Central Limit Theorem

**execution times are
normally distributed**



The sum of a sufficient number of independent, identically distributed random variables is approximately normally distributed.

STABILIZER

Memory layout affects performance
makes performance evaluation difficult

Use randomization to eliminate the effect of layout
enables sound performance evaluation

Case Studies
evaluation of LLVM's optimizations with STABILIZER

STABILIZER

Memory layout affects performance
makes performance evaluation difficult

Use randomization to eliminate the effect of layout
enables sound performance evaluation

Case Studies

evaluation of LLVM's optimizations with STABILIZER

Case Studies

evaluation of LLVM's optimizations with STABILIZER

on each benchmark

**across the whole
benchmark suite**

first, build benchmarks with STABILIZER

Build programs with **STABILIZER**

```
> szc main.c
```

Build programs with **STABILIZER**

```
> szc main.c
```

Build programs with **STABILIZER**

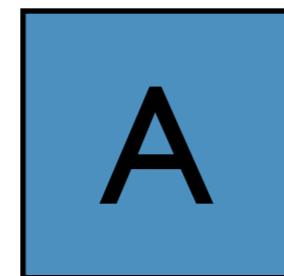
```
> szc -Rcode main.c
```

Build programs with STABILIZER

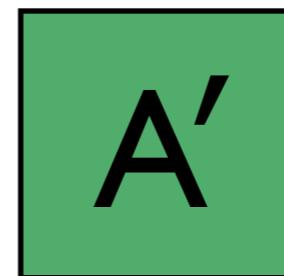
```
> szc -Rcode -Rheap -Rstack main.c
```

now run the benchmarks

Run benchmarks as usual

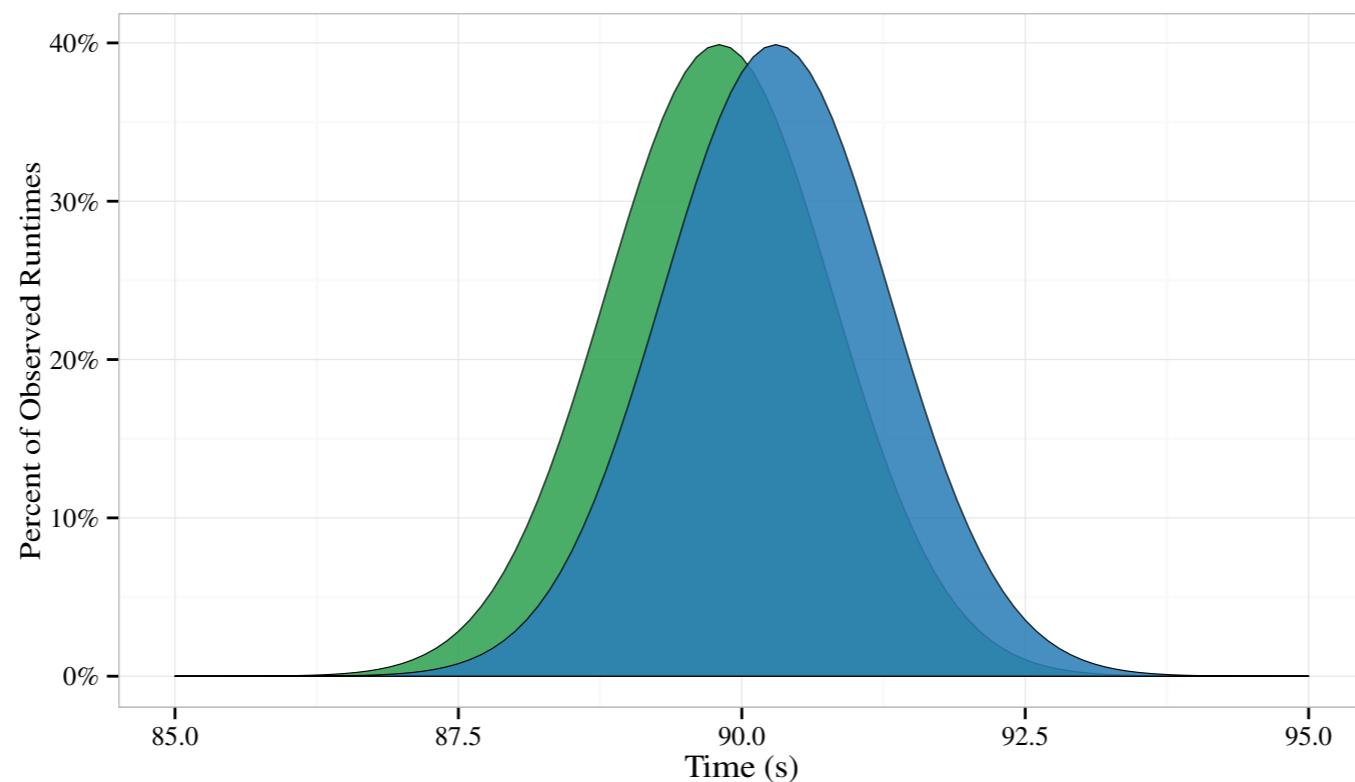


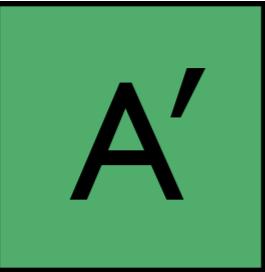
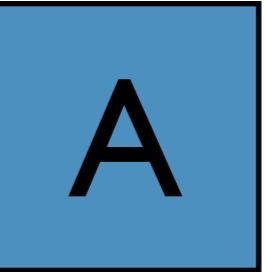
$\times 30$

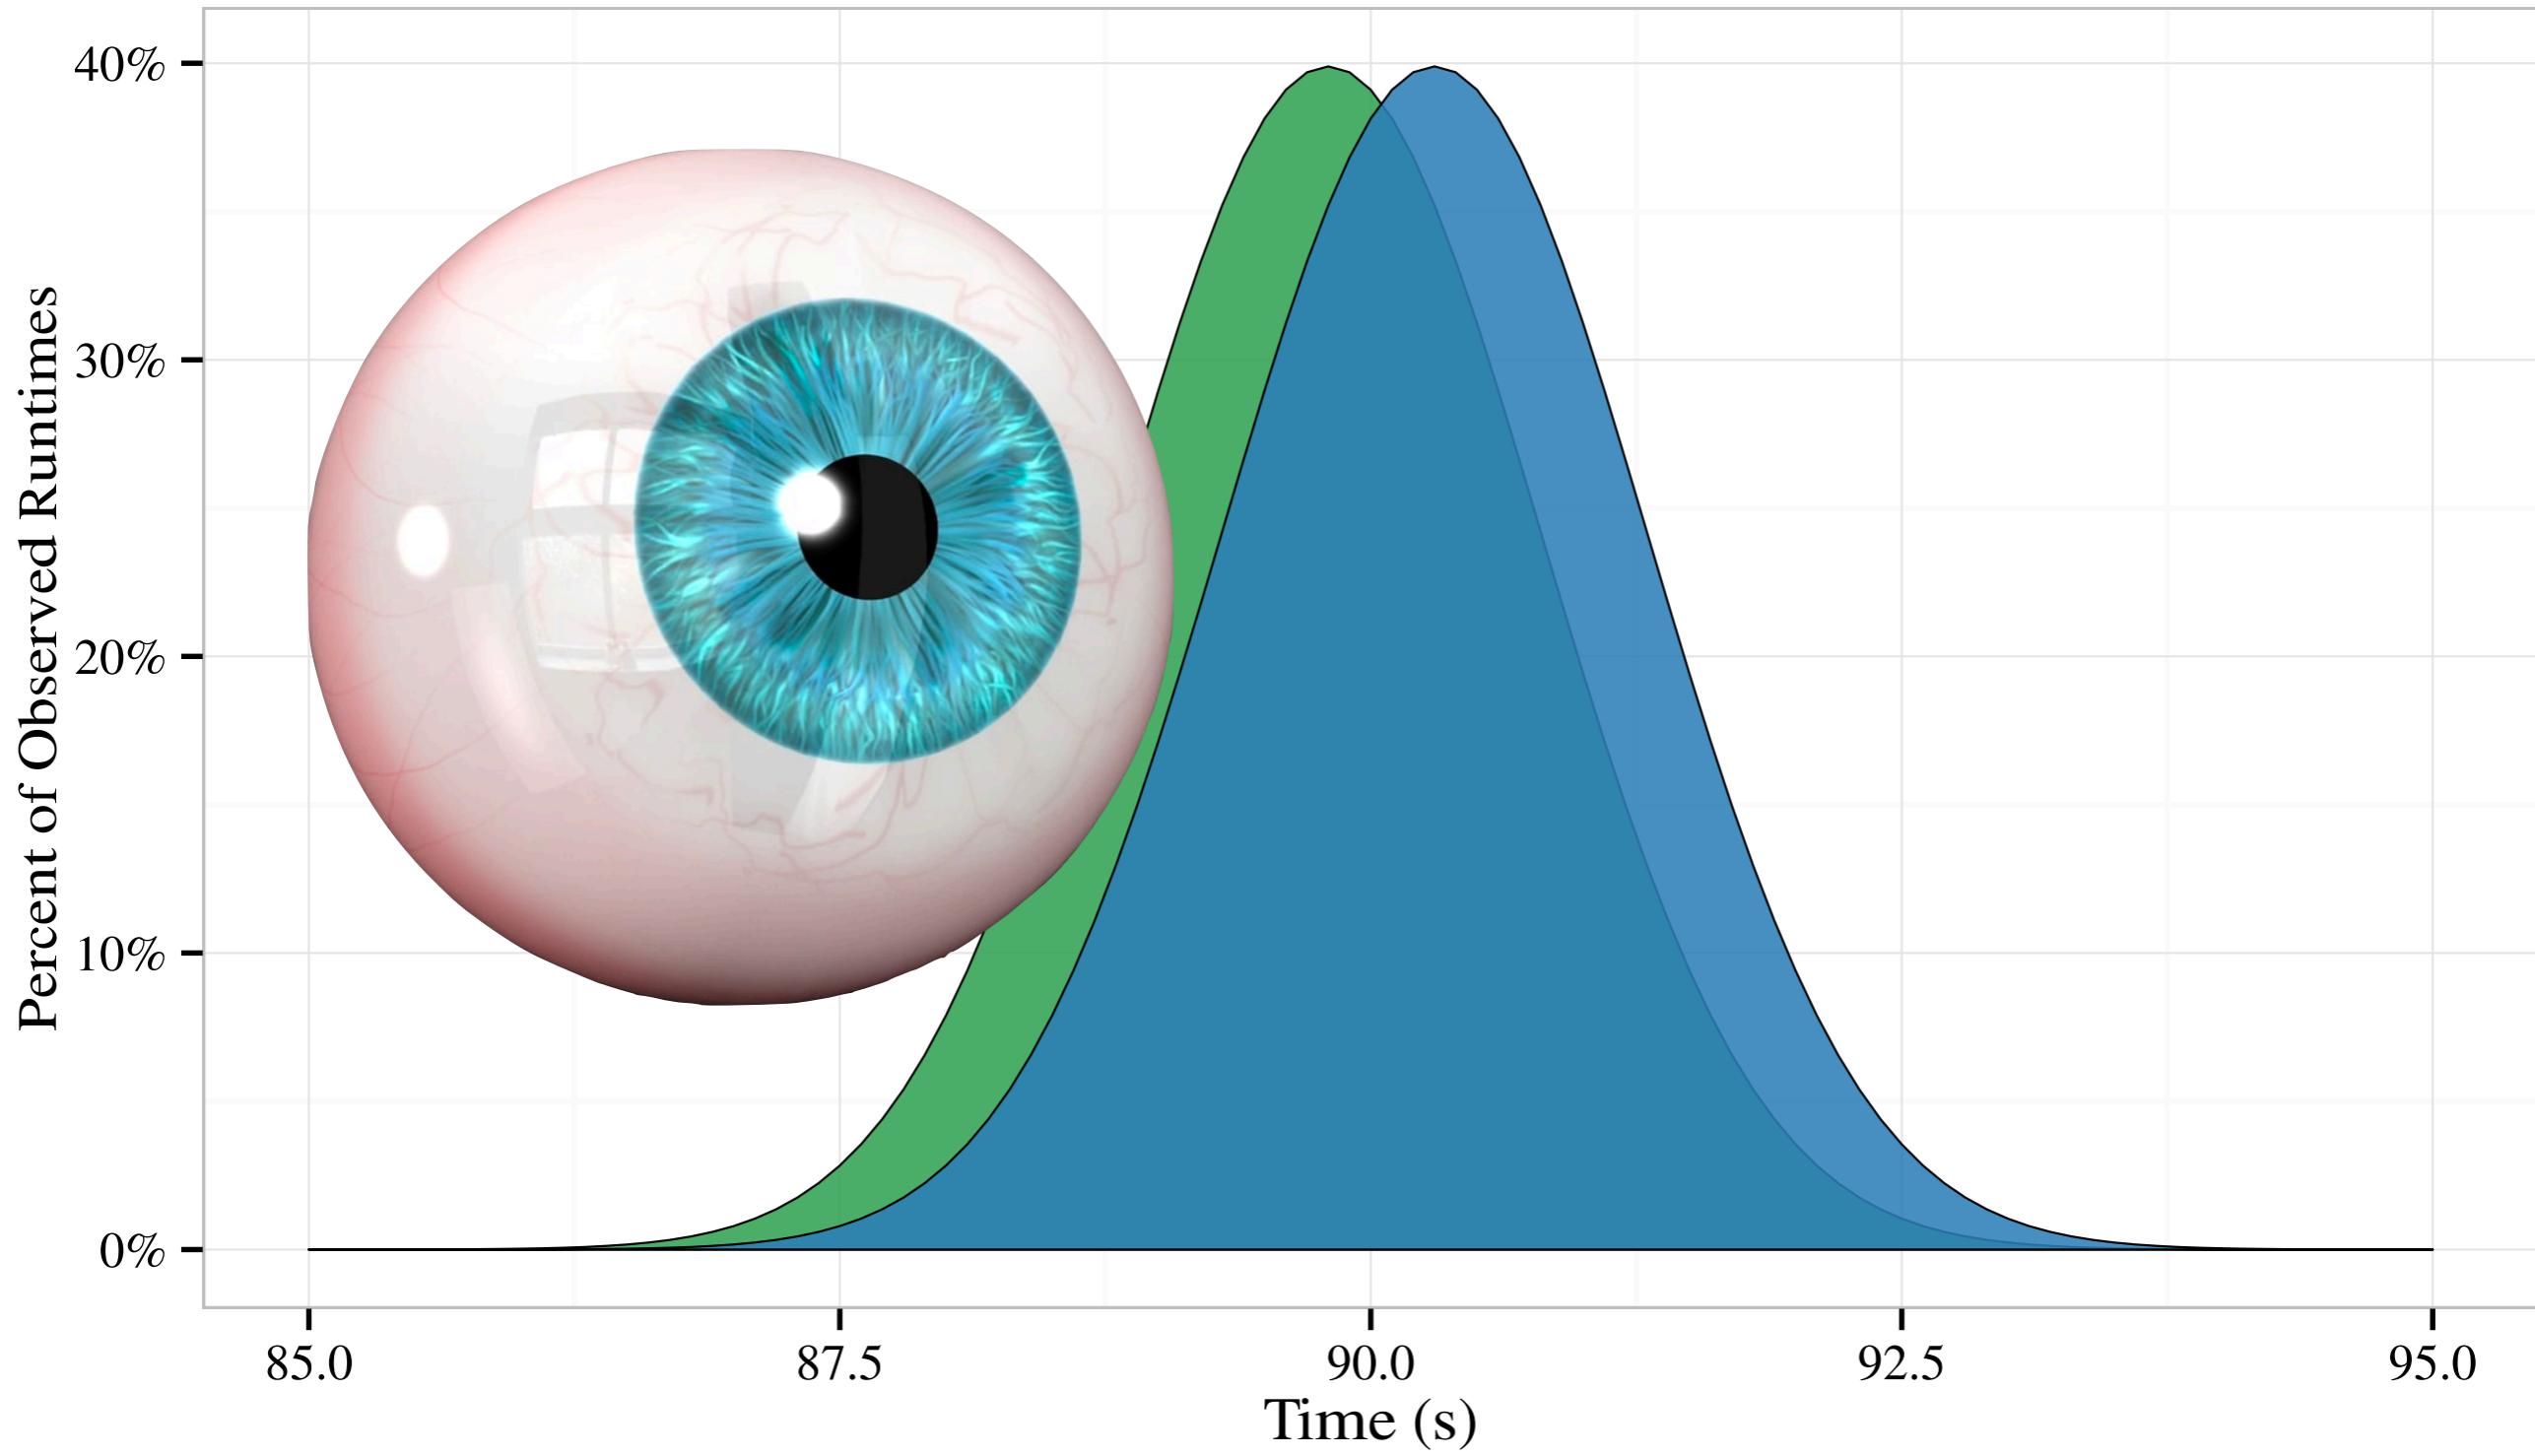


$\times 30$

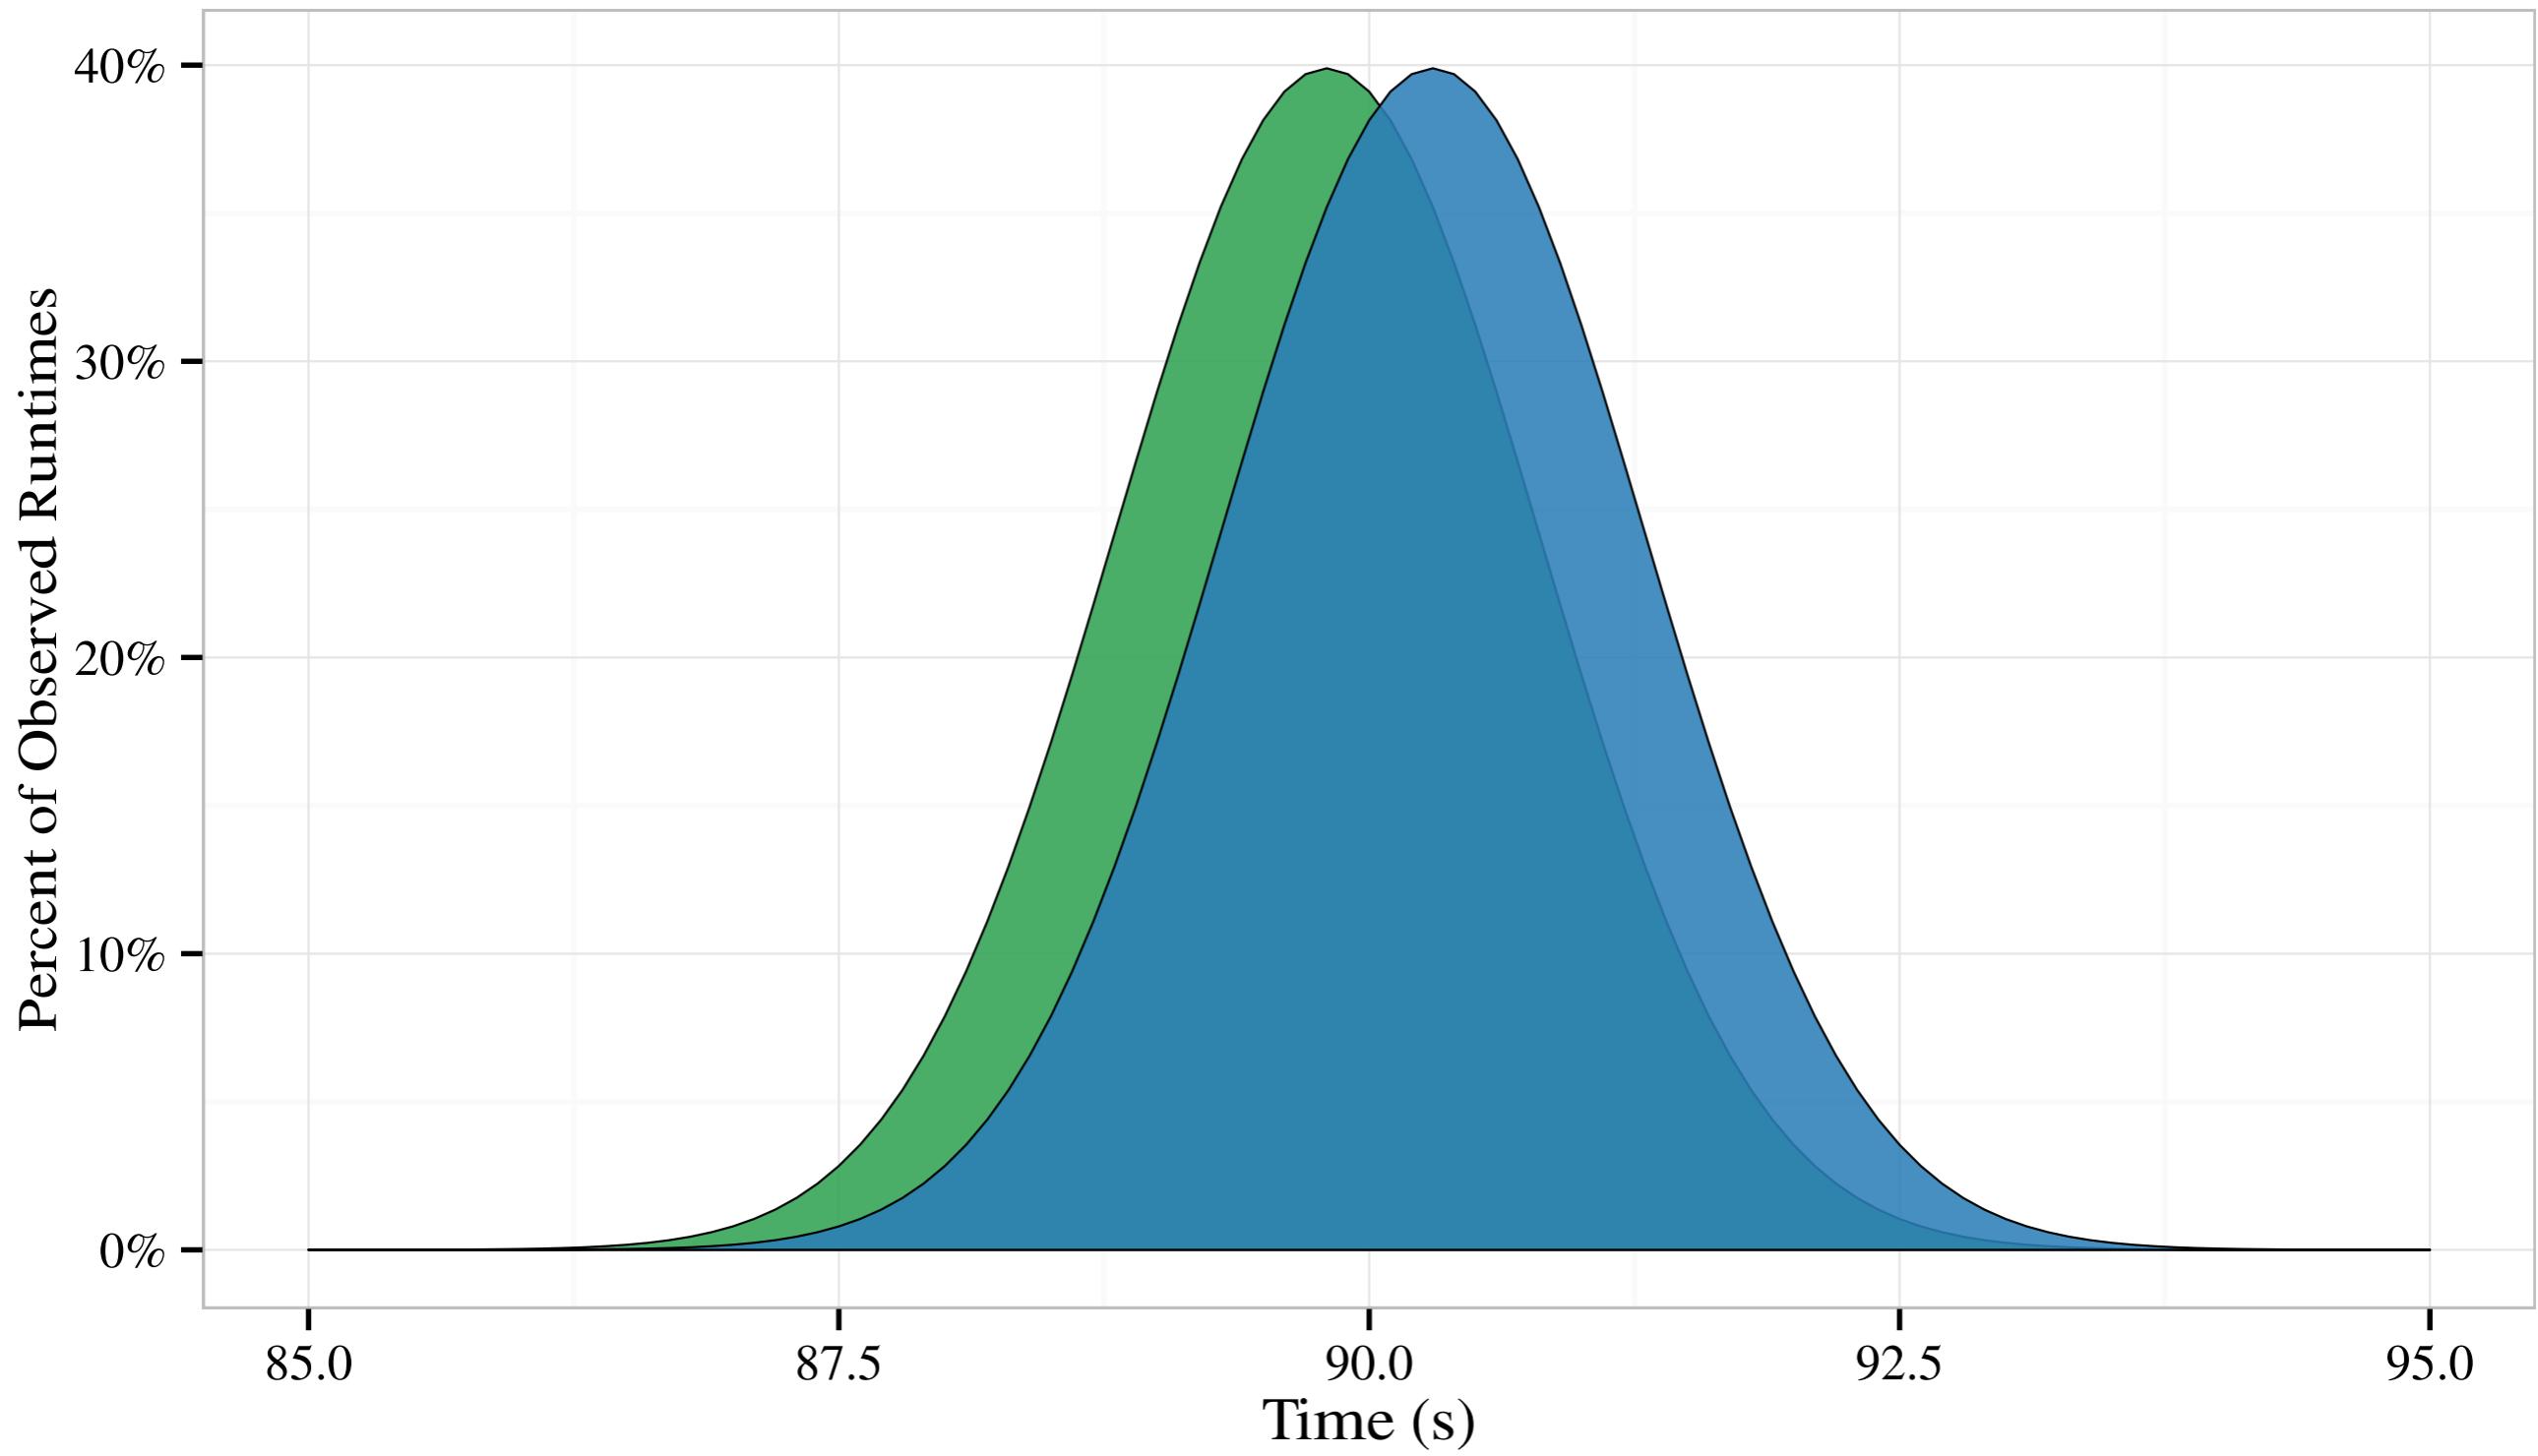
drop the results into R



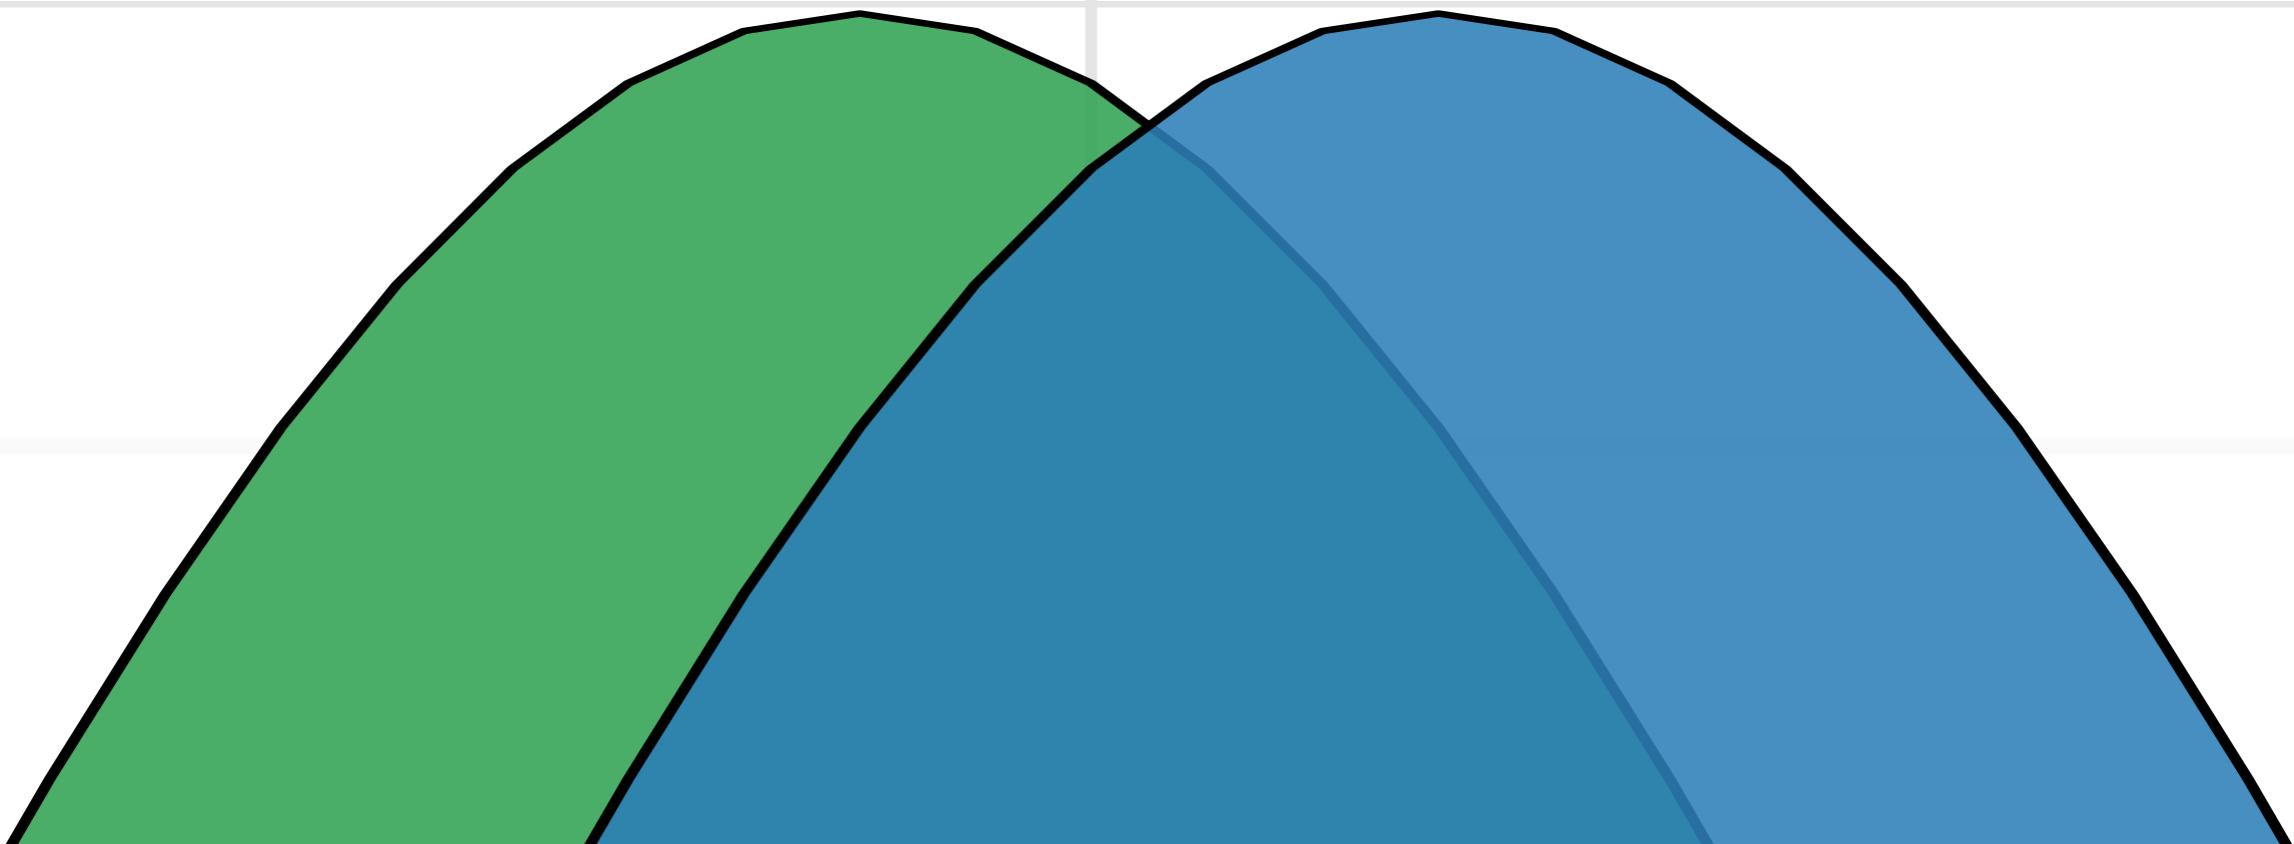
Is  faster than ?



If $A' = A$

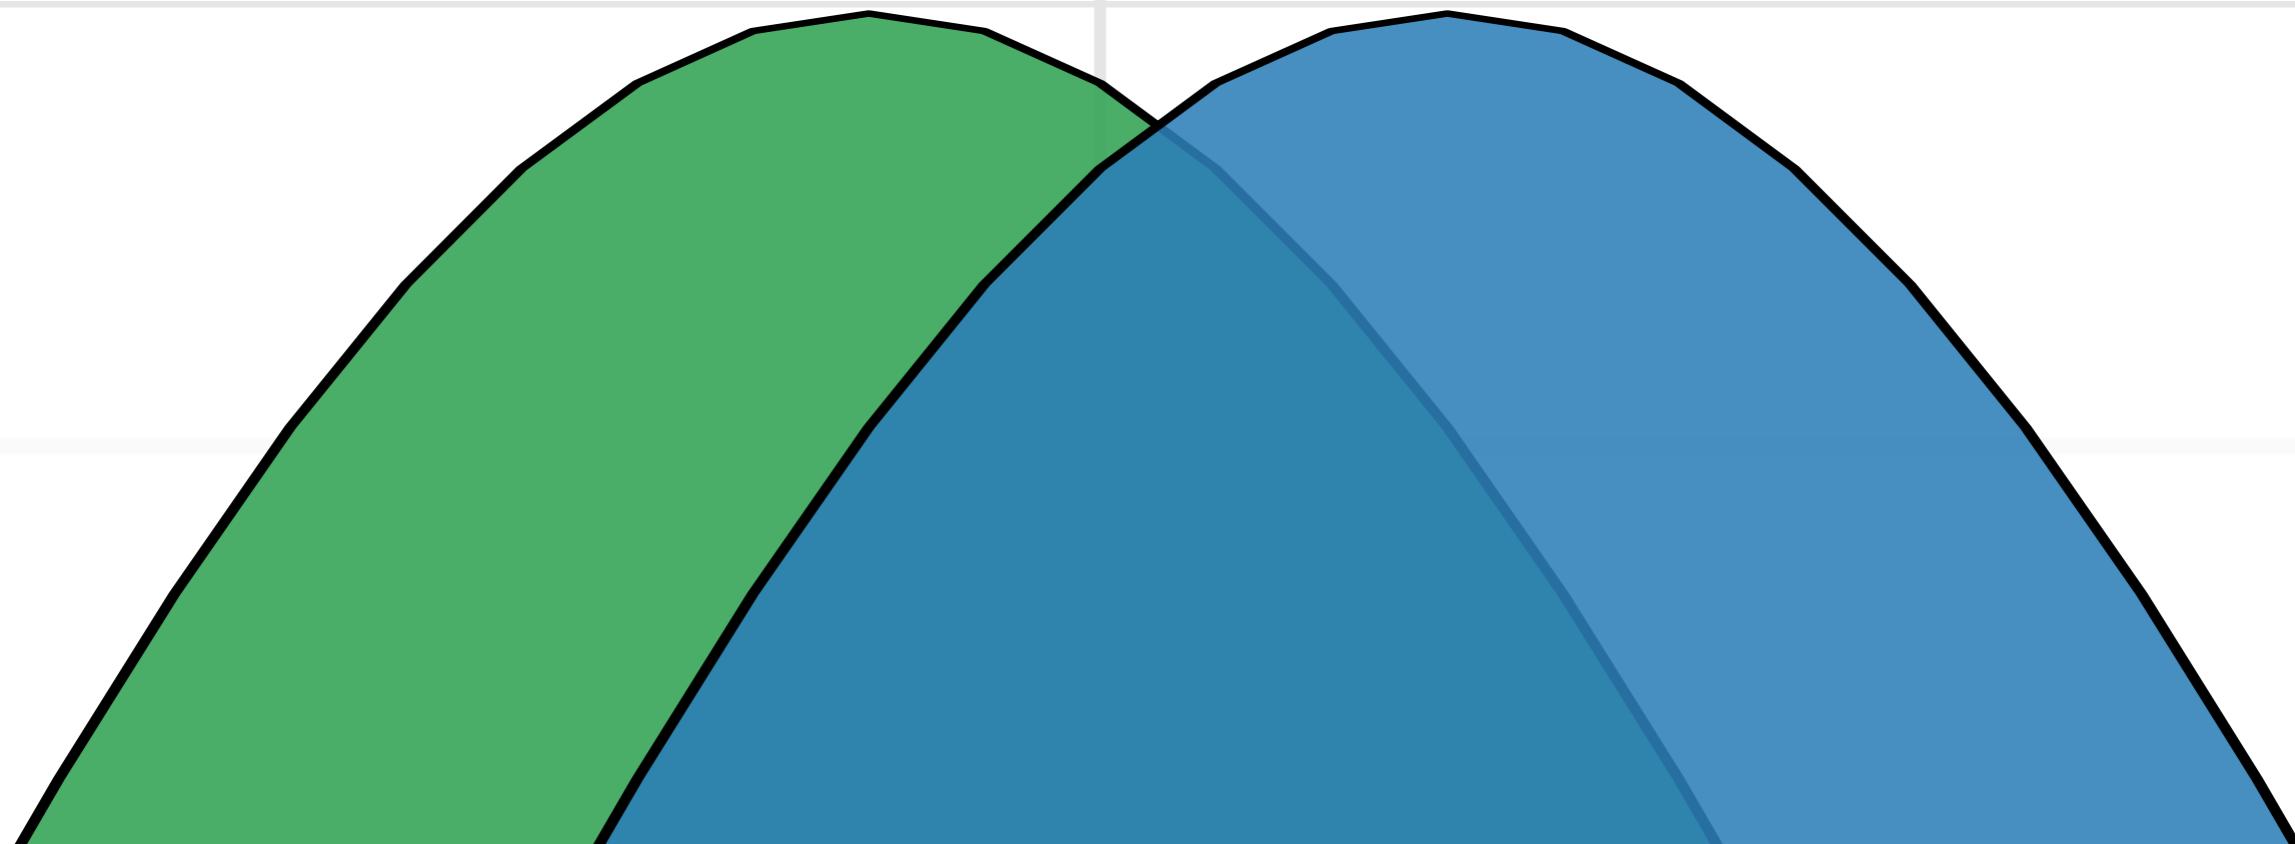
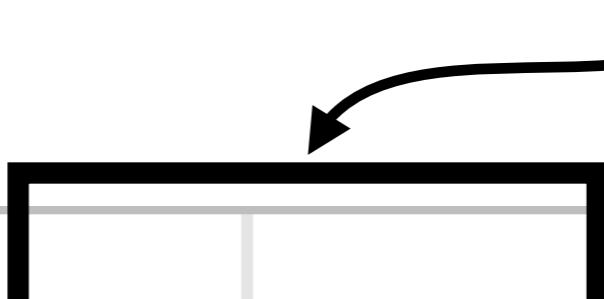


If $A' = A$



If $A' = A$

what is the *probability* of measuring a difference at least this large?



The Student's t-test

`t.test(times.A', times.A)`

If A' = A

what is the *probability* of measuring a difference at least this large?



The Student's t-test

If p-value

If $A' = A$

what is the probability of measuring a difference at least this large?



The Student's t-test

If p-value \leq **5%**

If **A'** = **A**

what is the *probability* of measuring a difference at least this large?



The Student's t-test

If p-value $\leq 5\%$
we reject the null hypothesis

$$\text{If } A' = A$$

what is the *probability* of measuring a
difference at least this large?



The Student's t-test

If p-value $\leq 5\%$
we *reject the null hypothesis*

$$A' \neq A$$

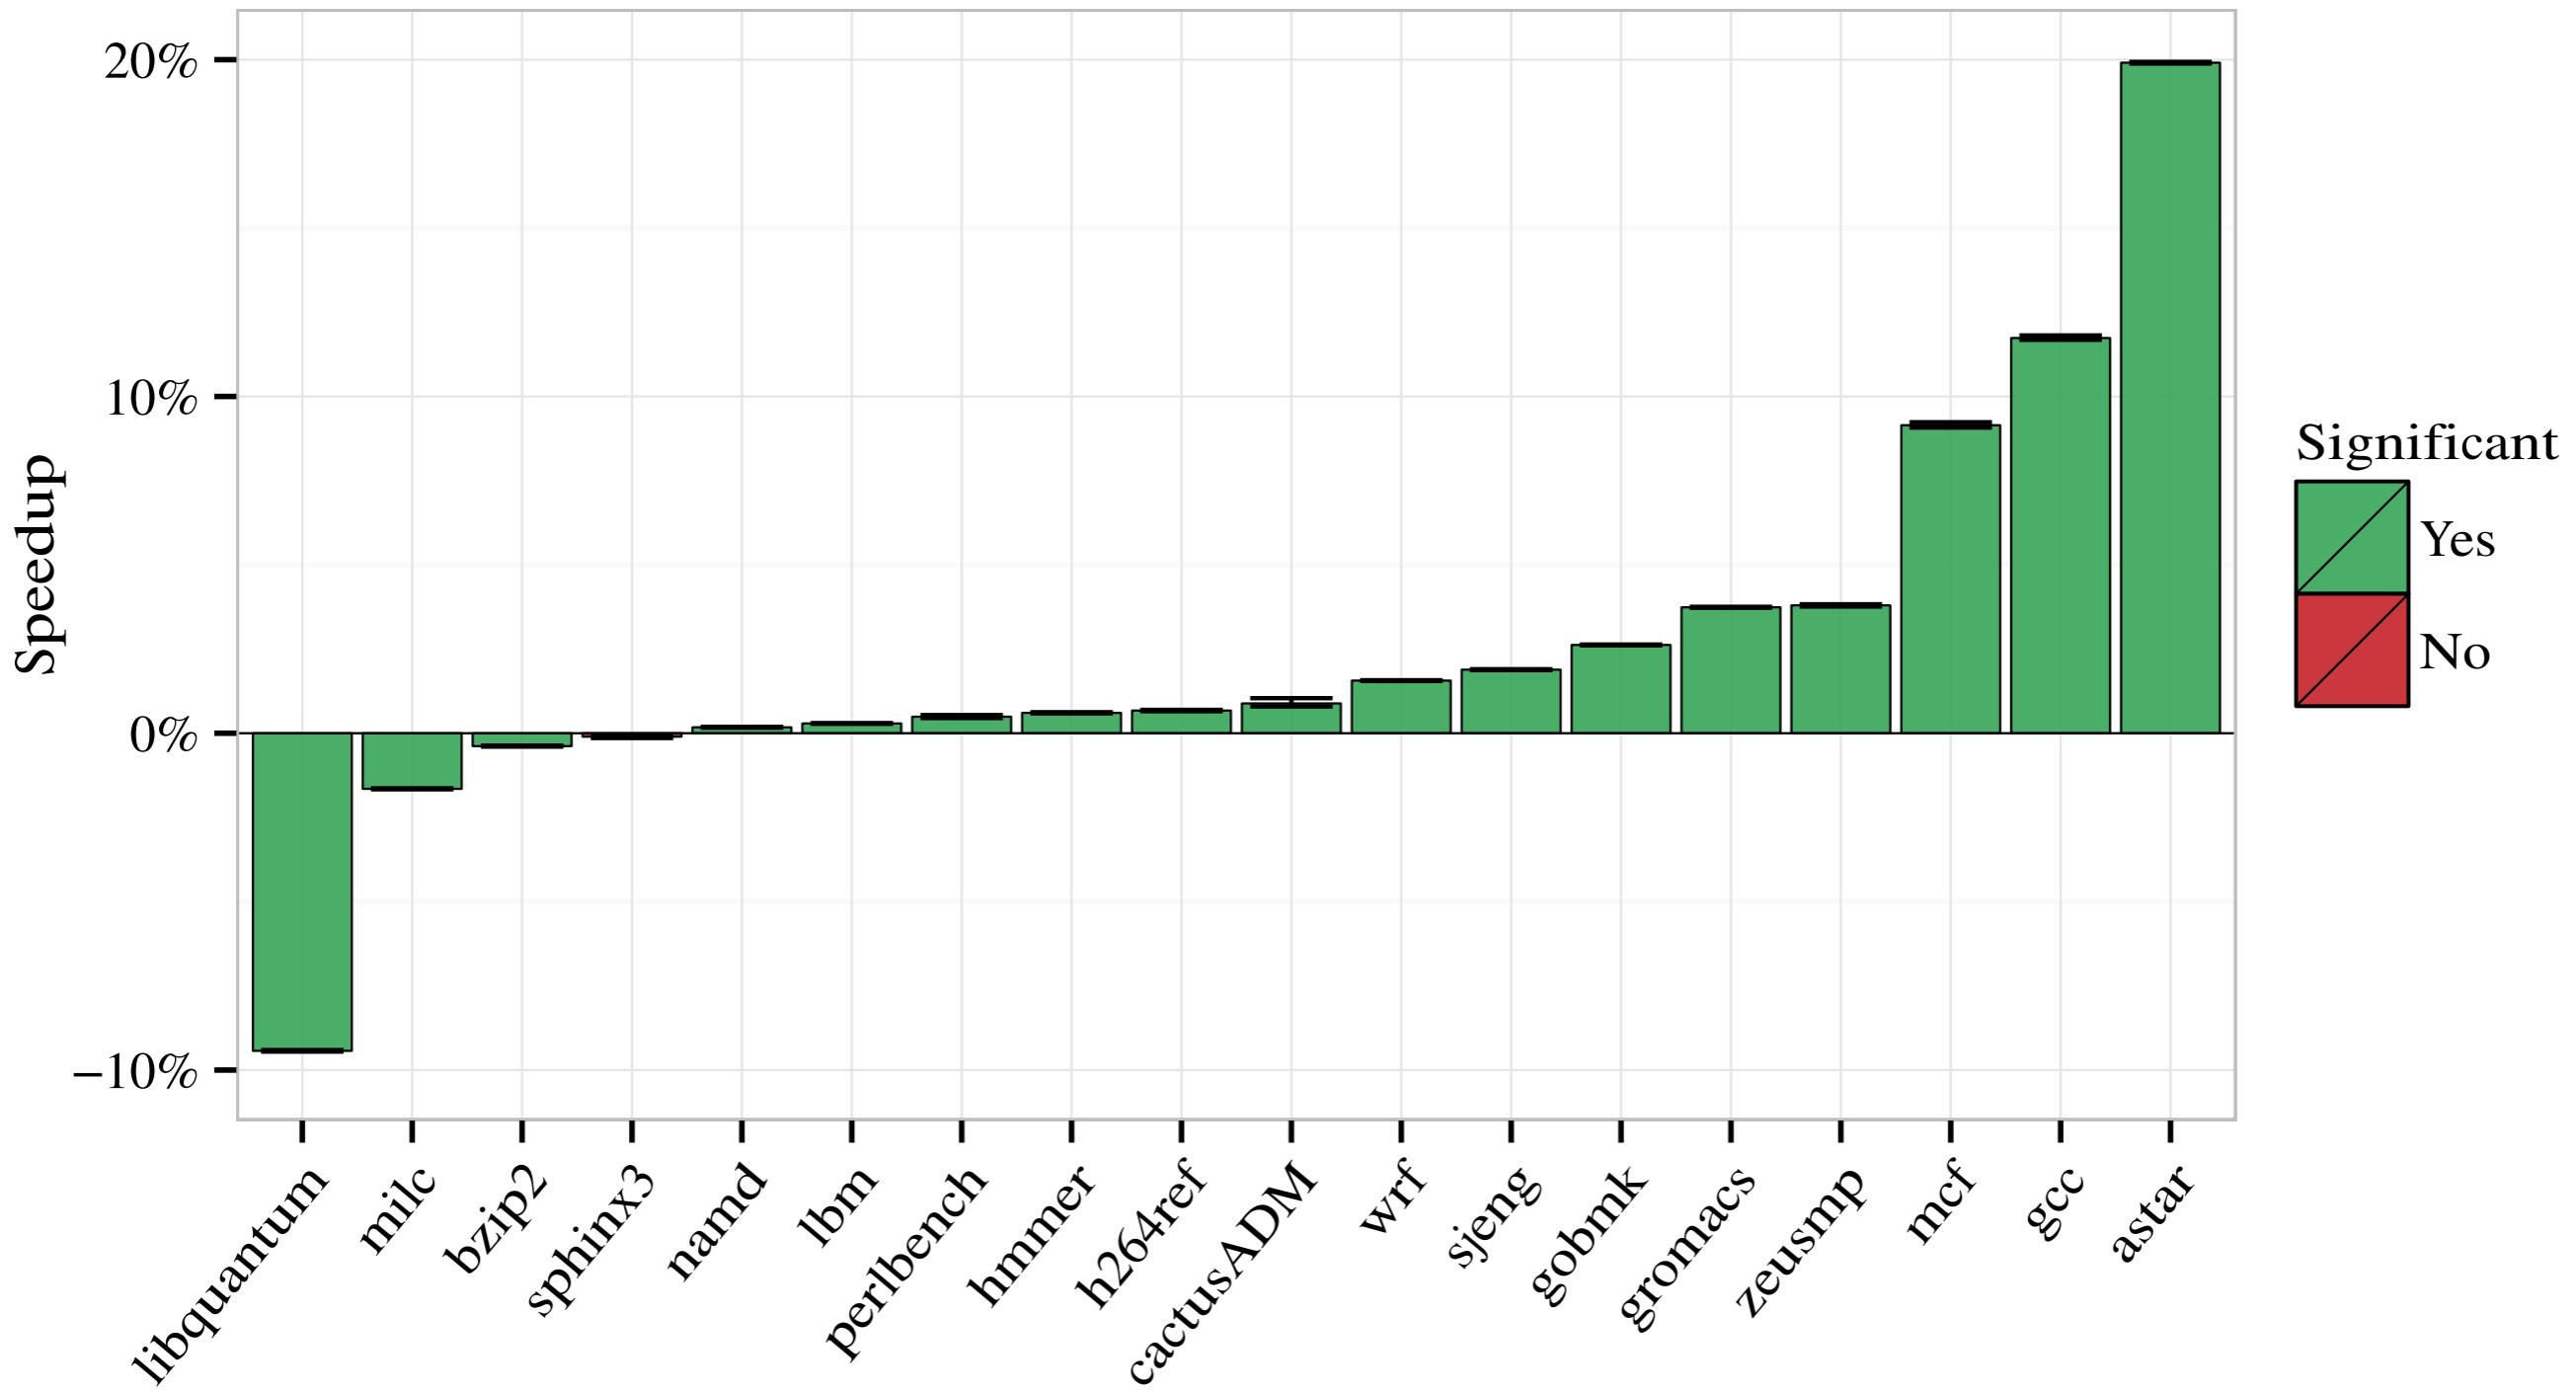
Random chance not responsible for
the measured difference



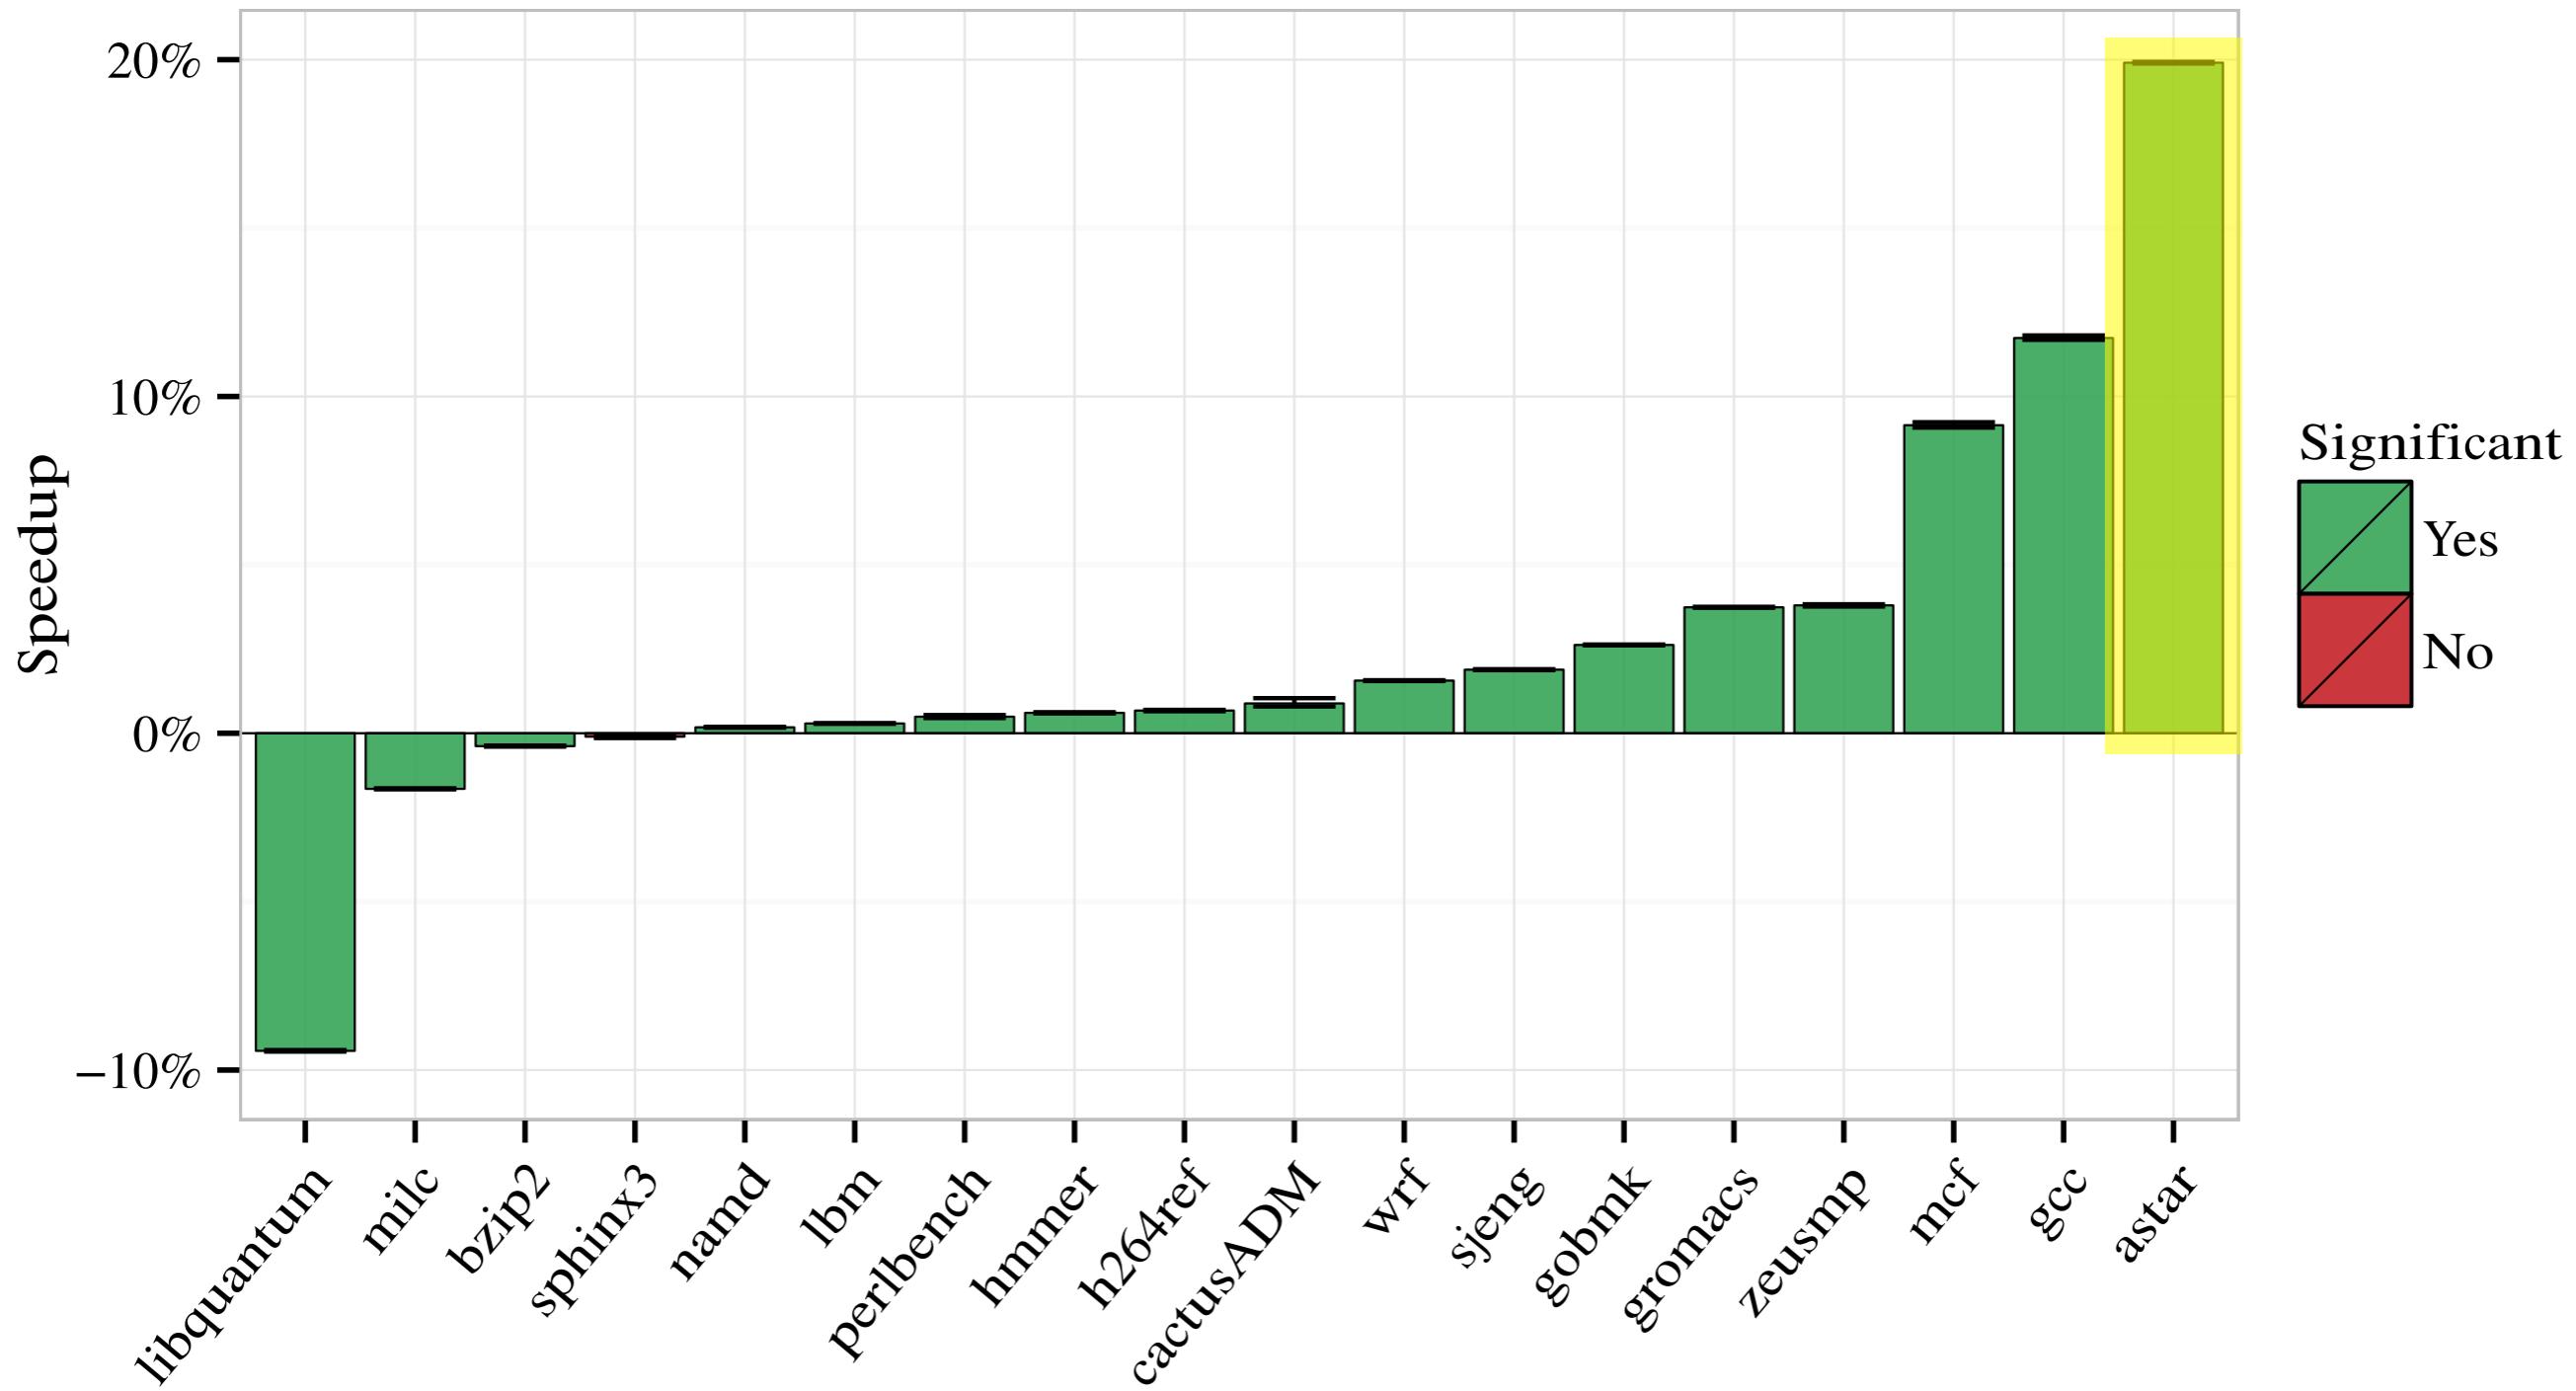
The difference is real

$$A' \neq A$$

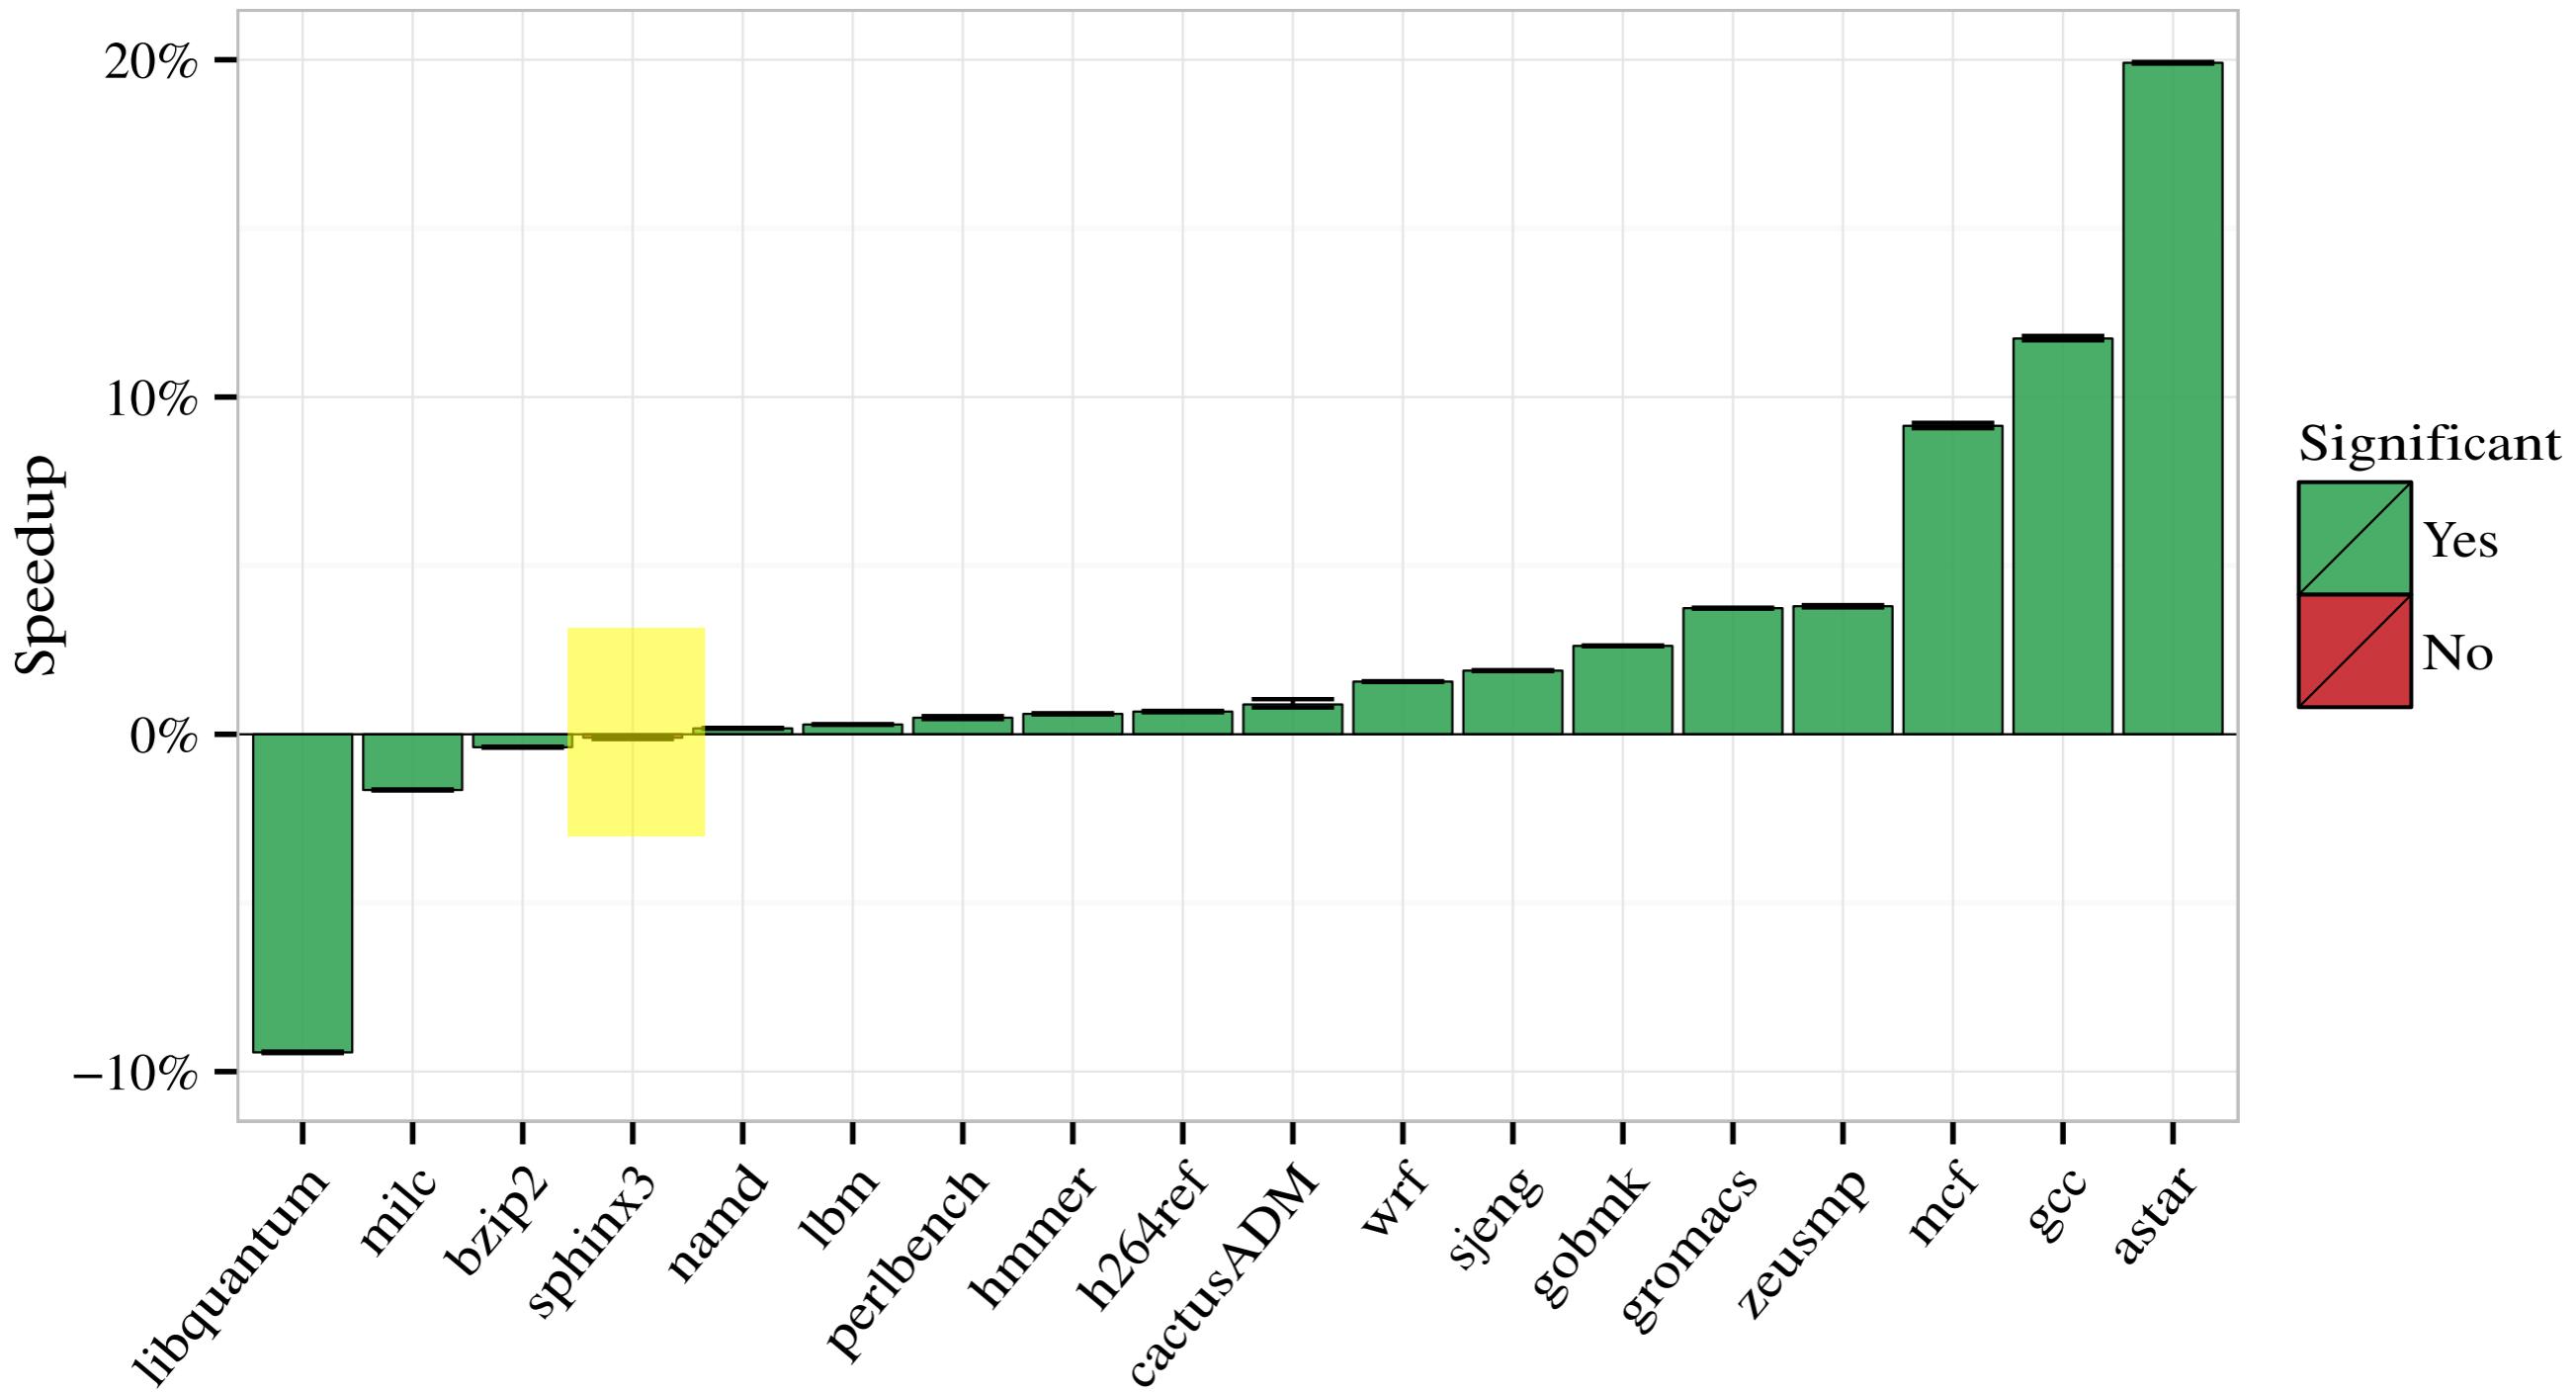
Speedup of -02 over -01



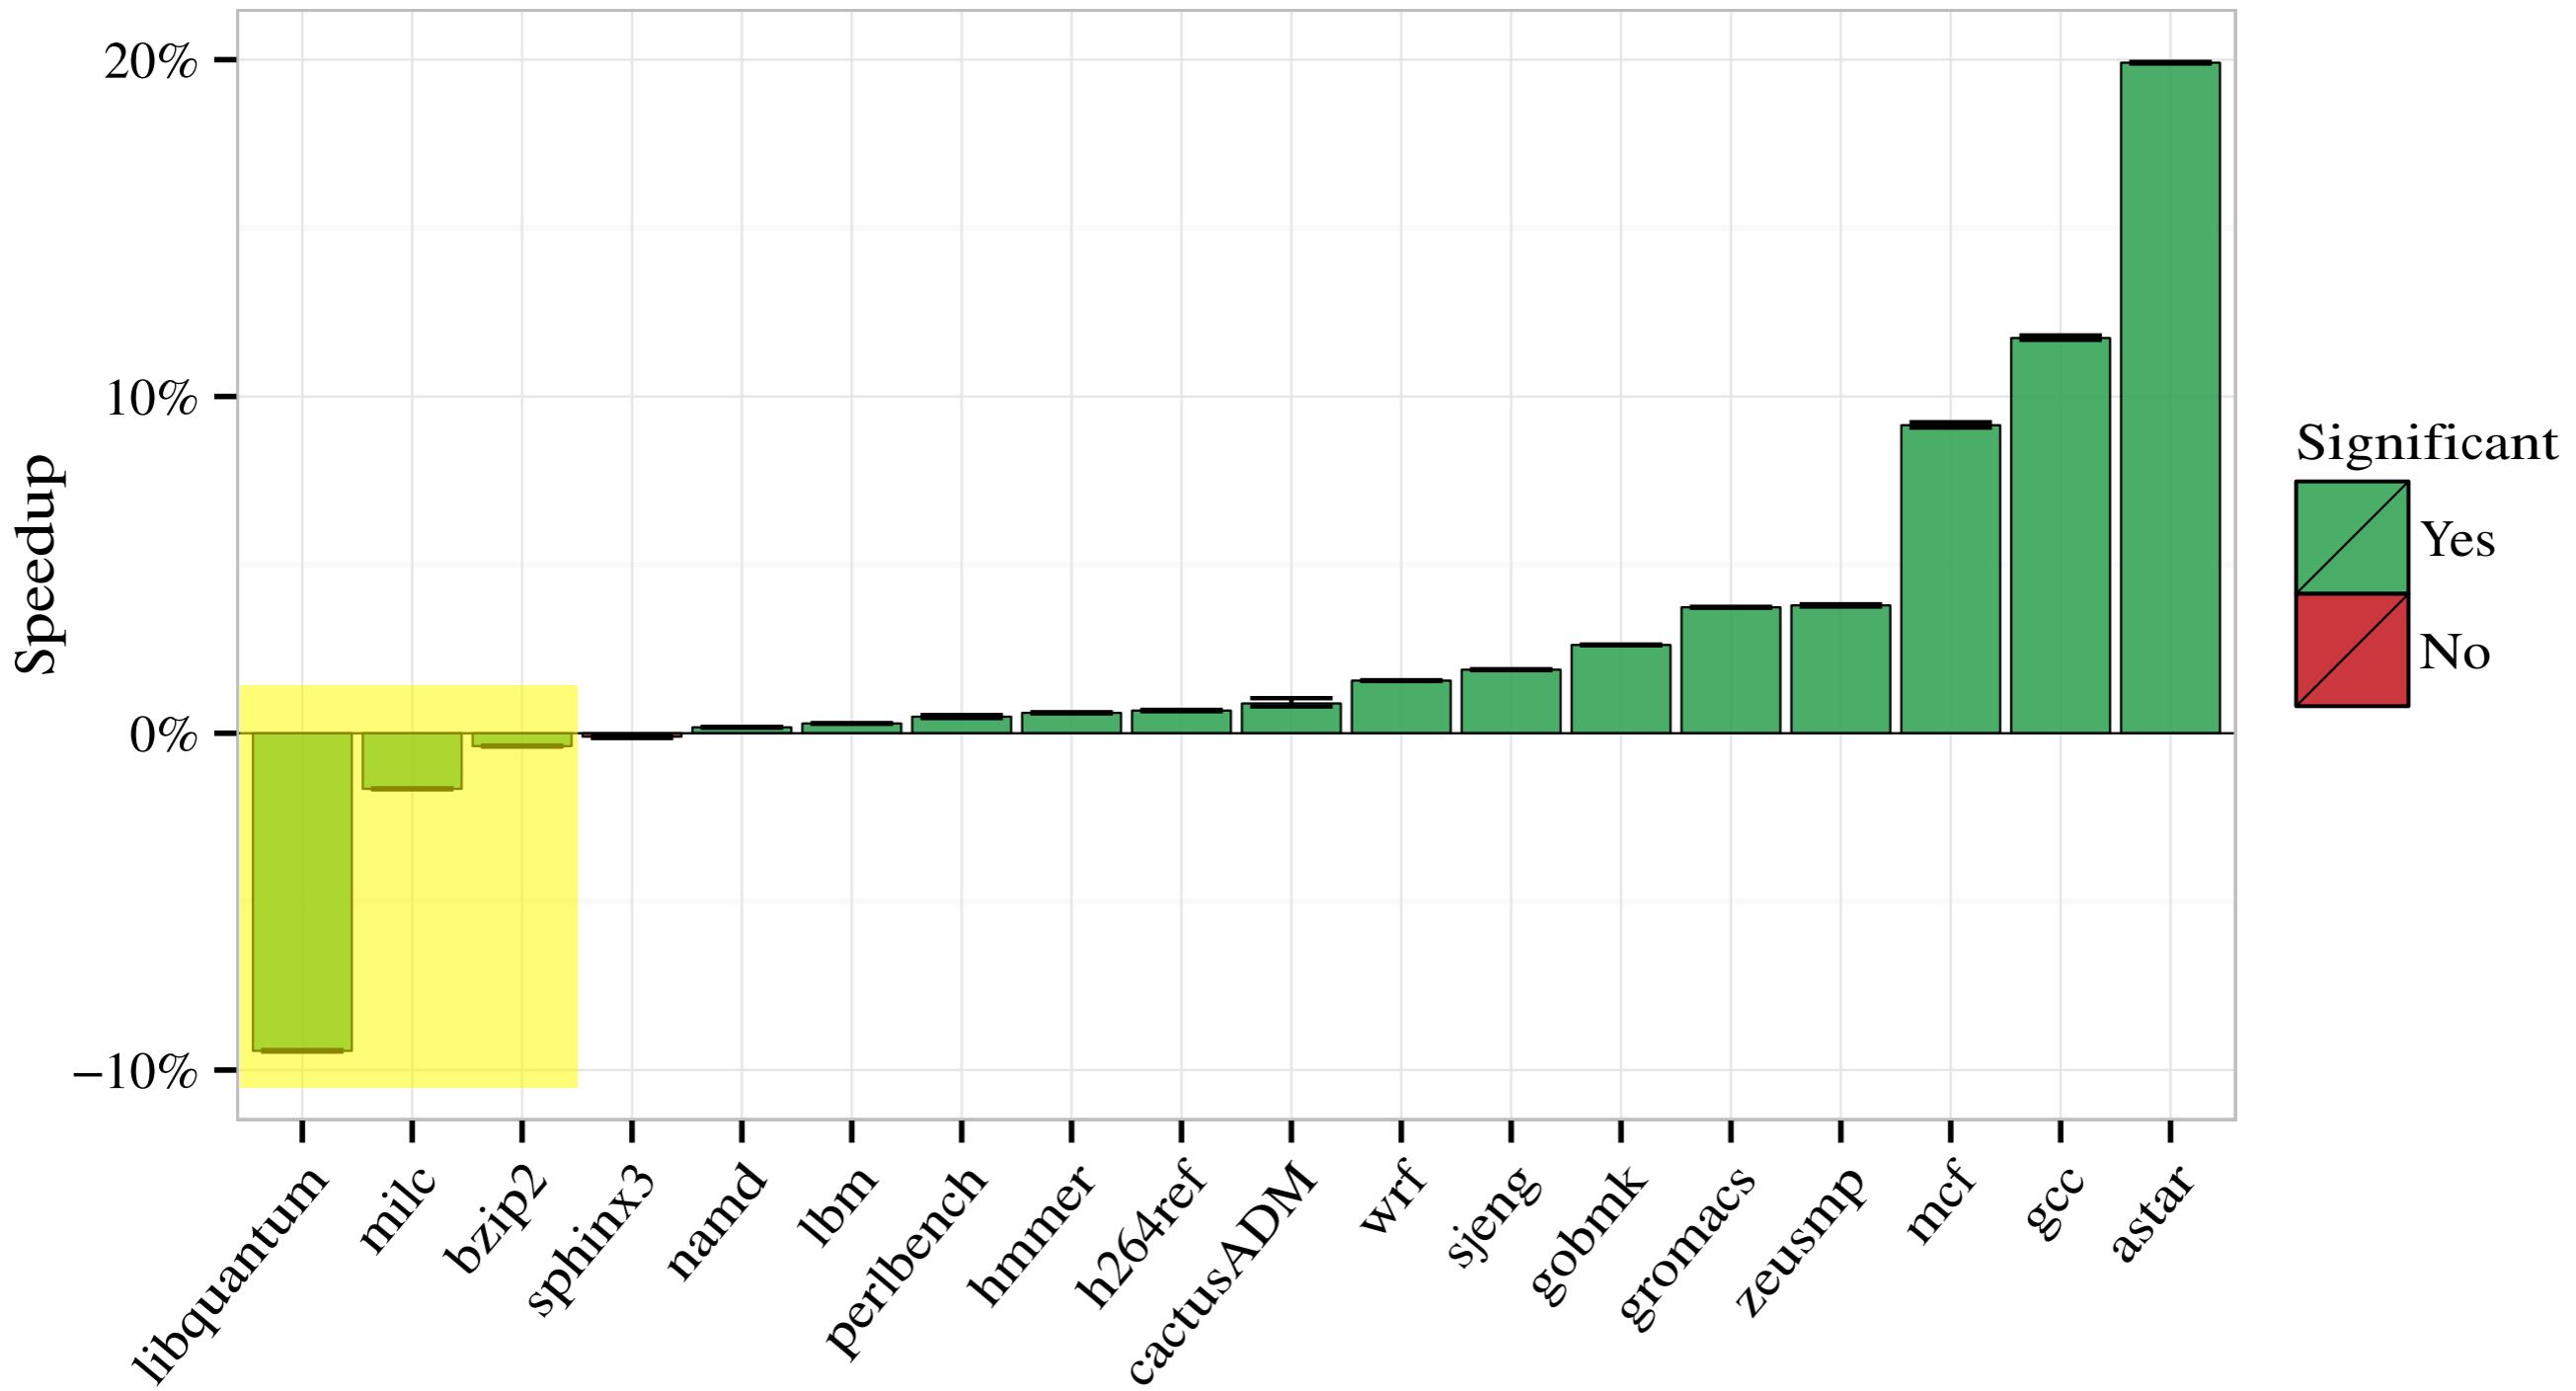
Speedup of -02 over -01



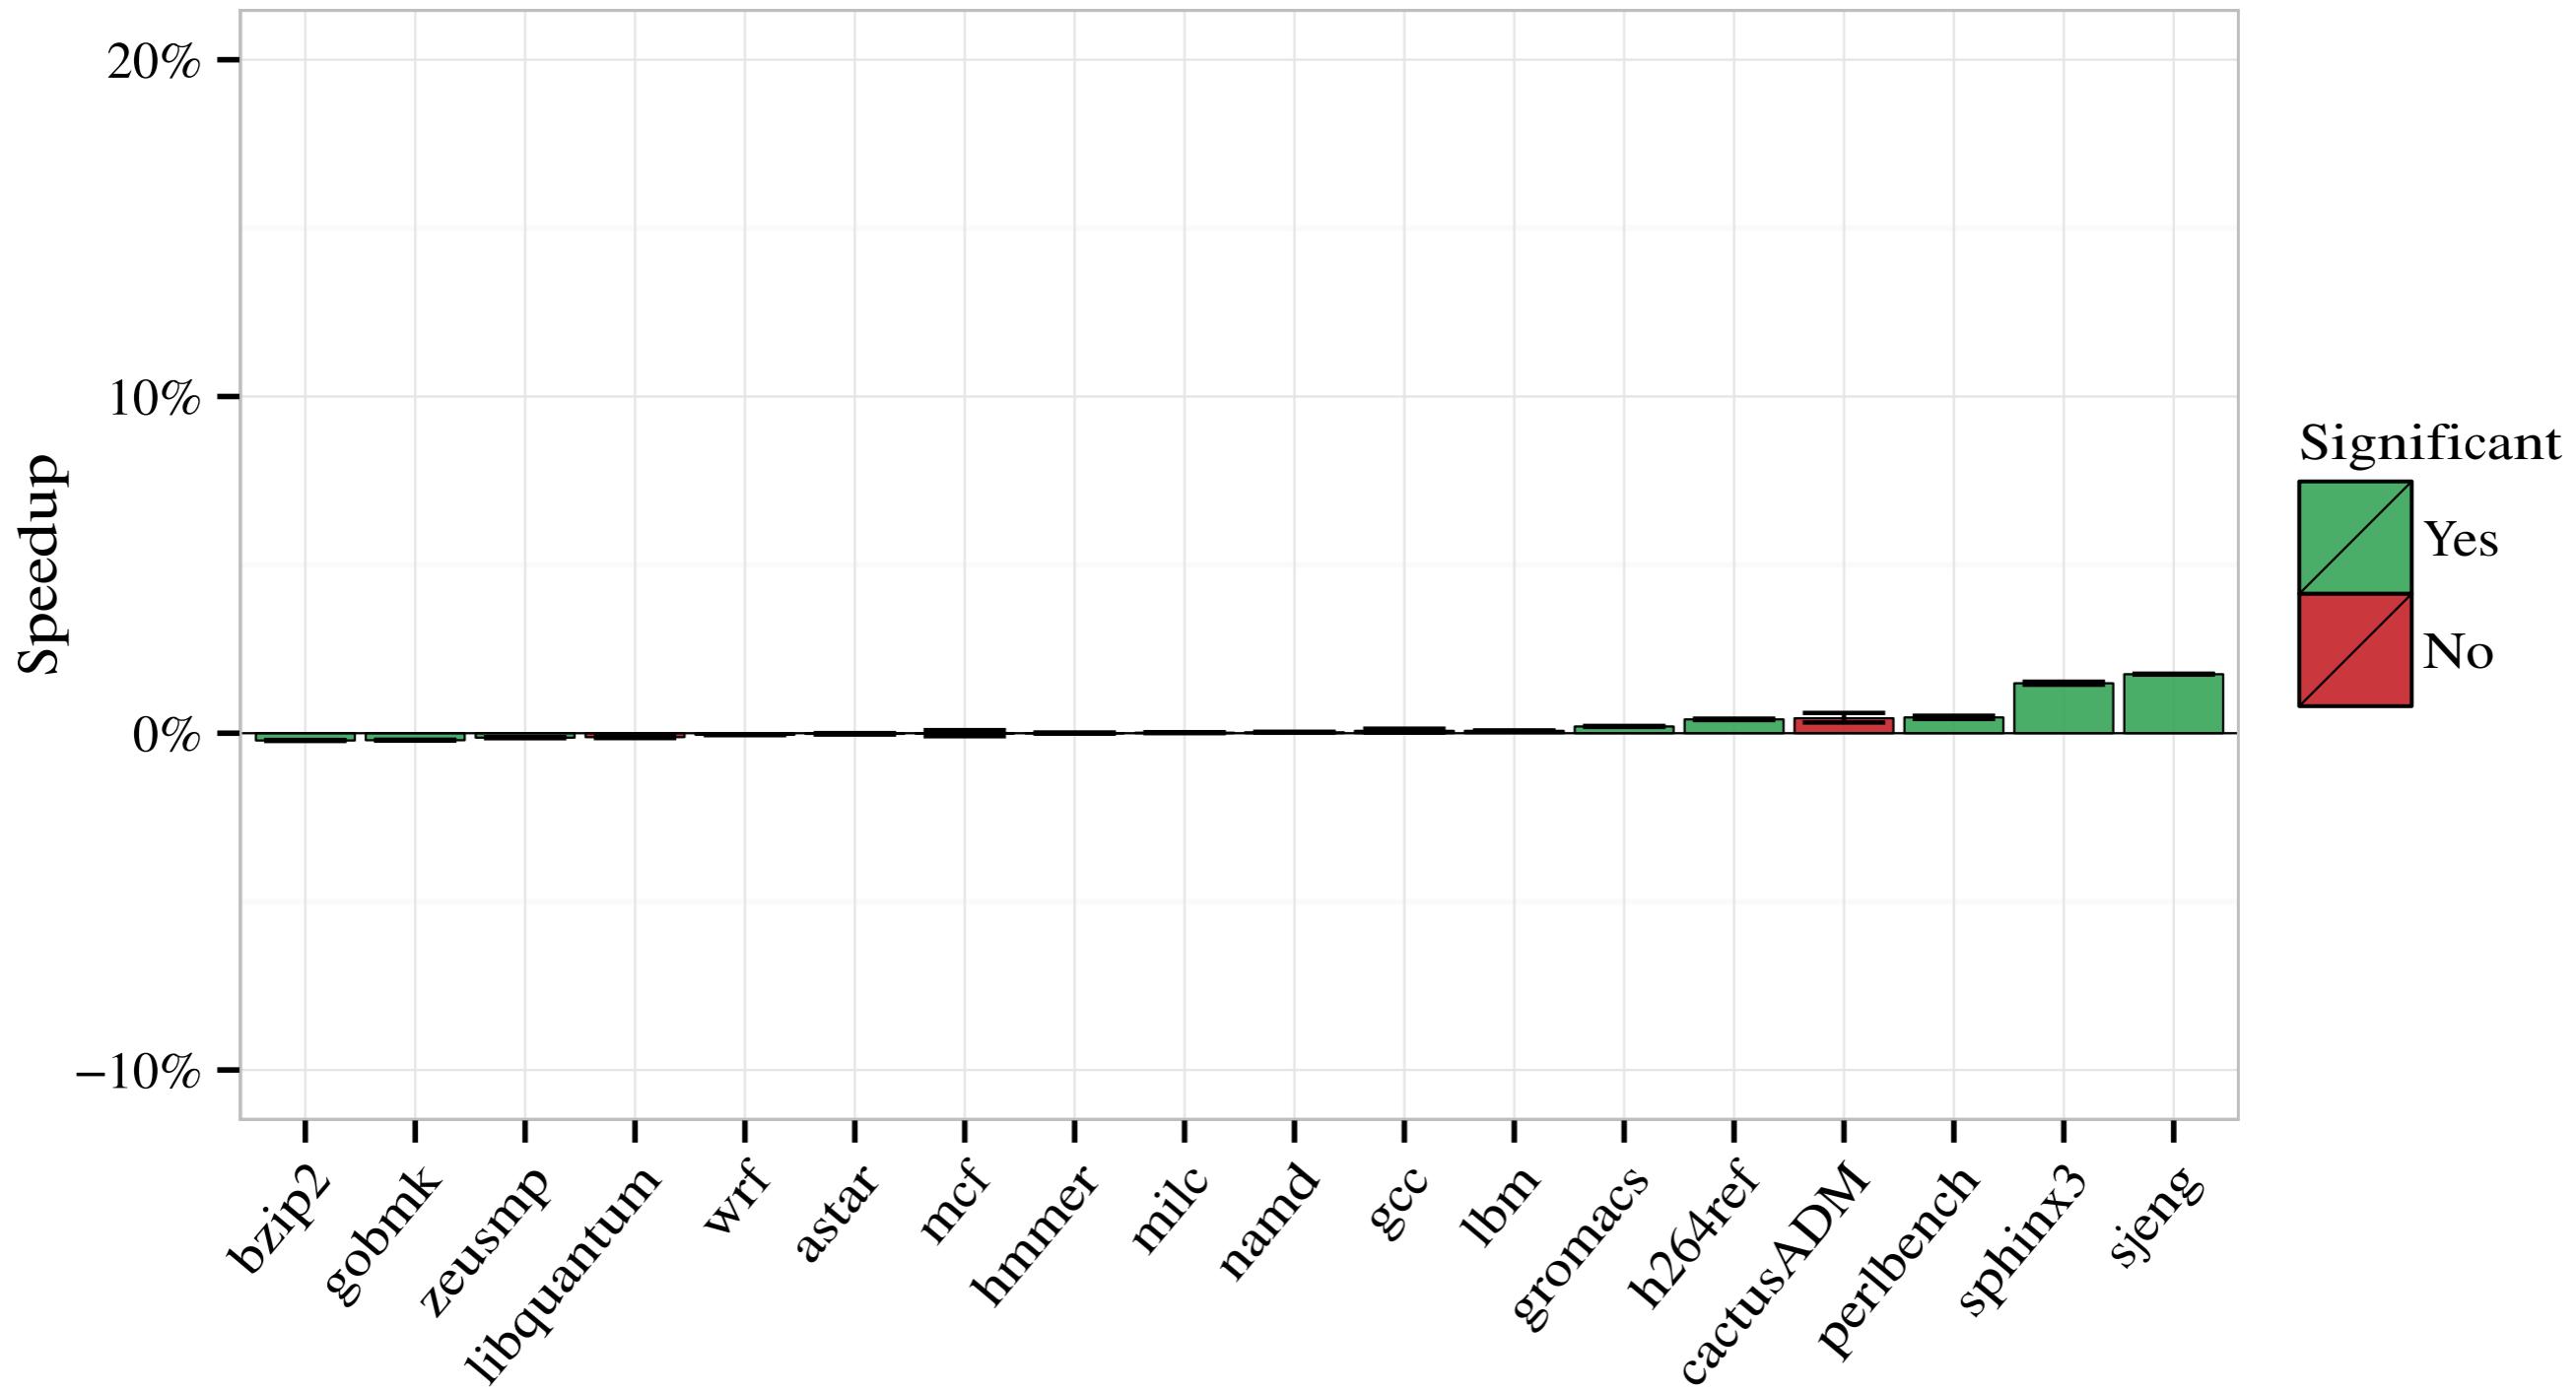
Speedup of -02 over -01



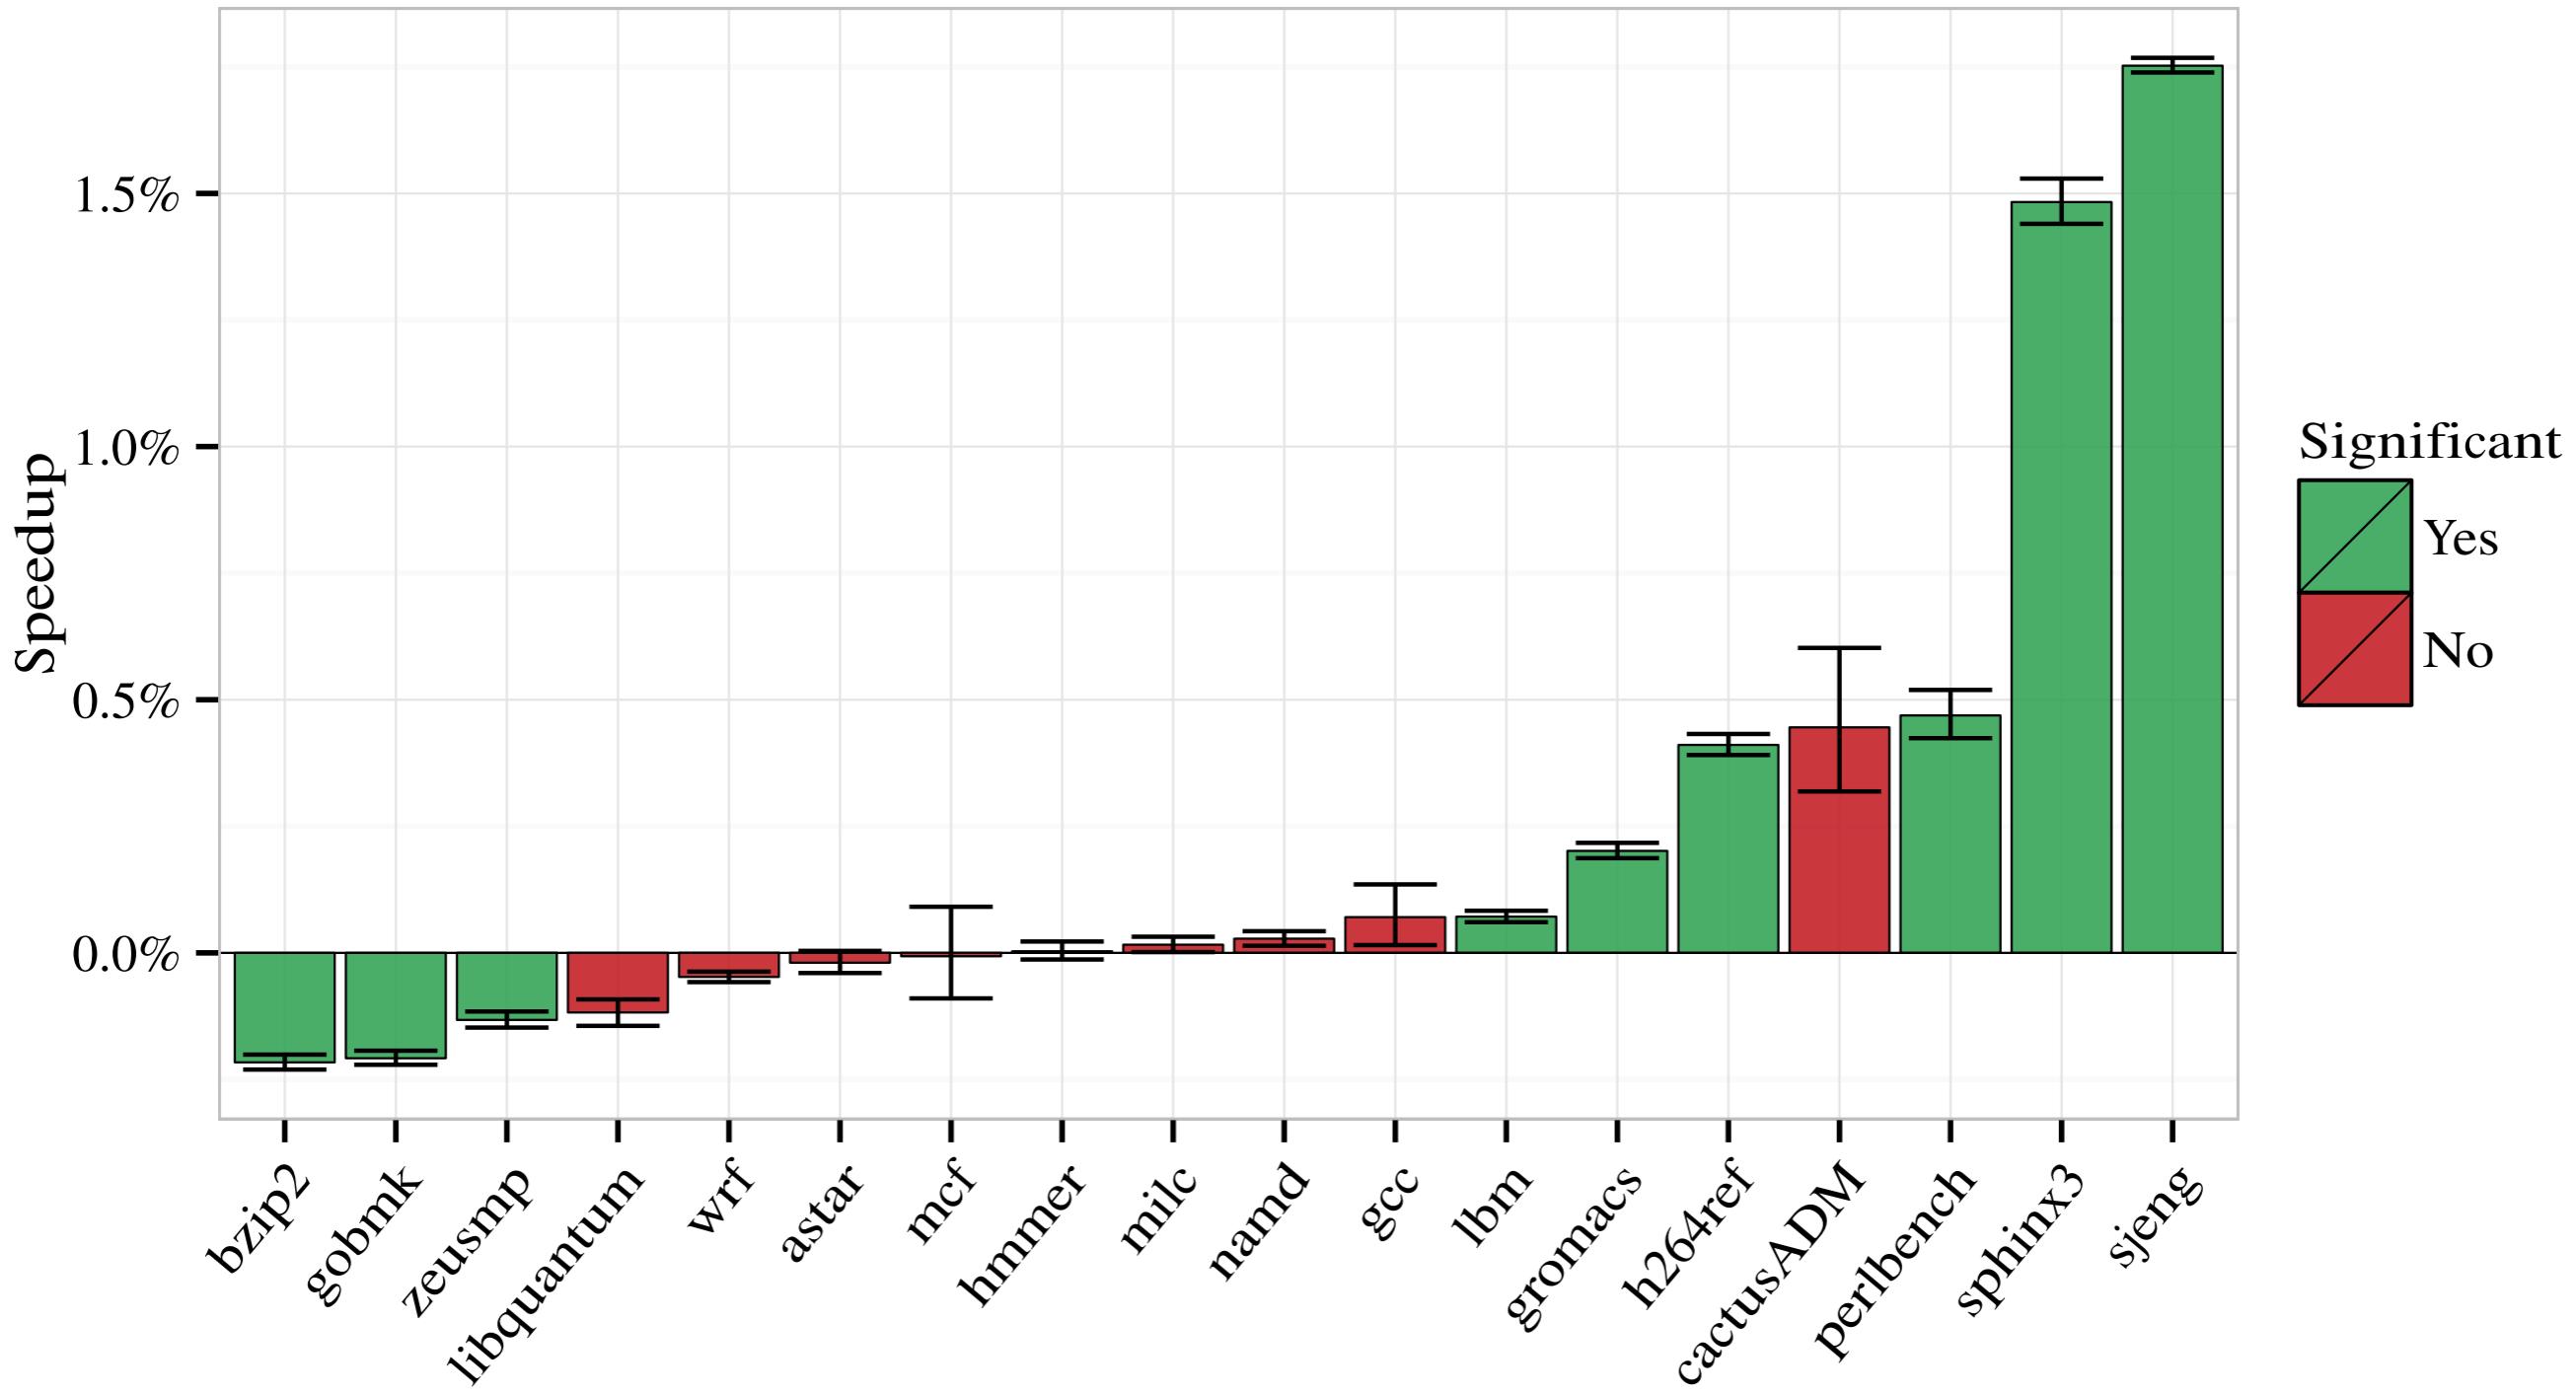
Speedup of -02 over -01



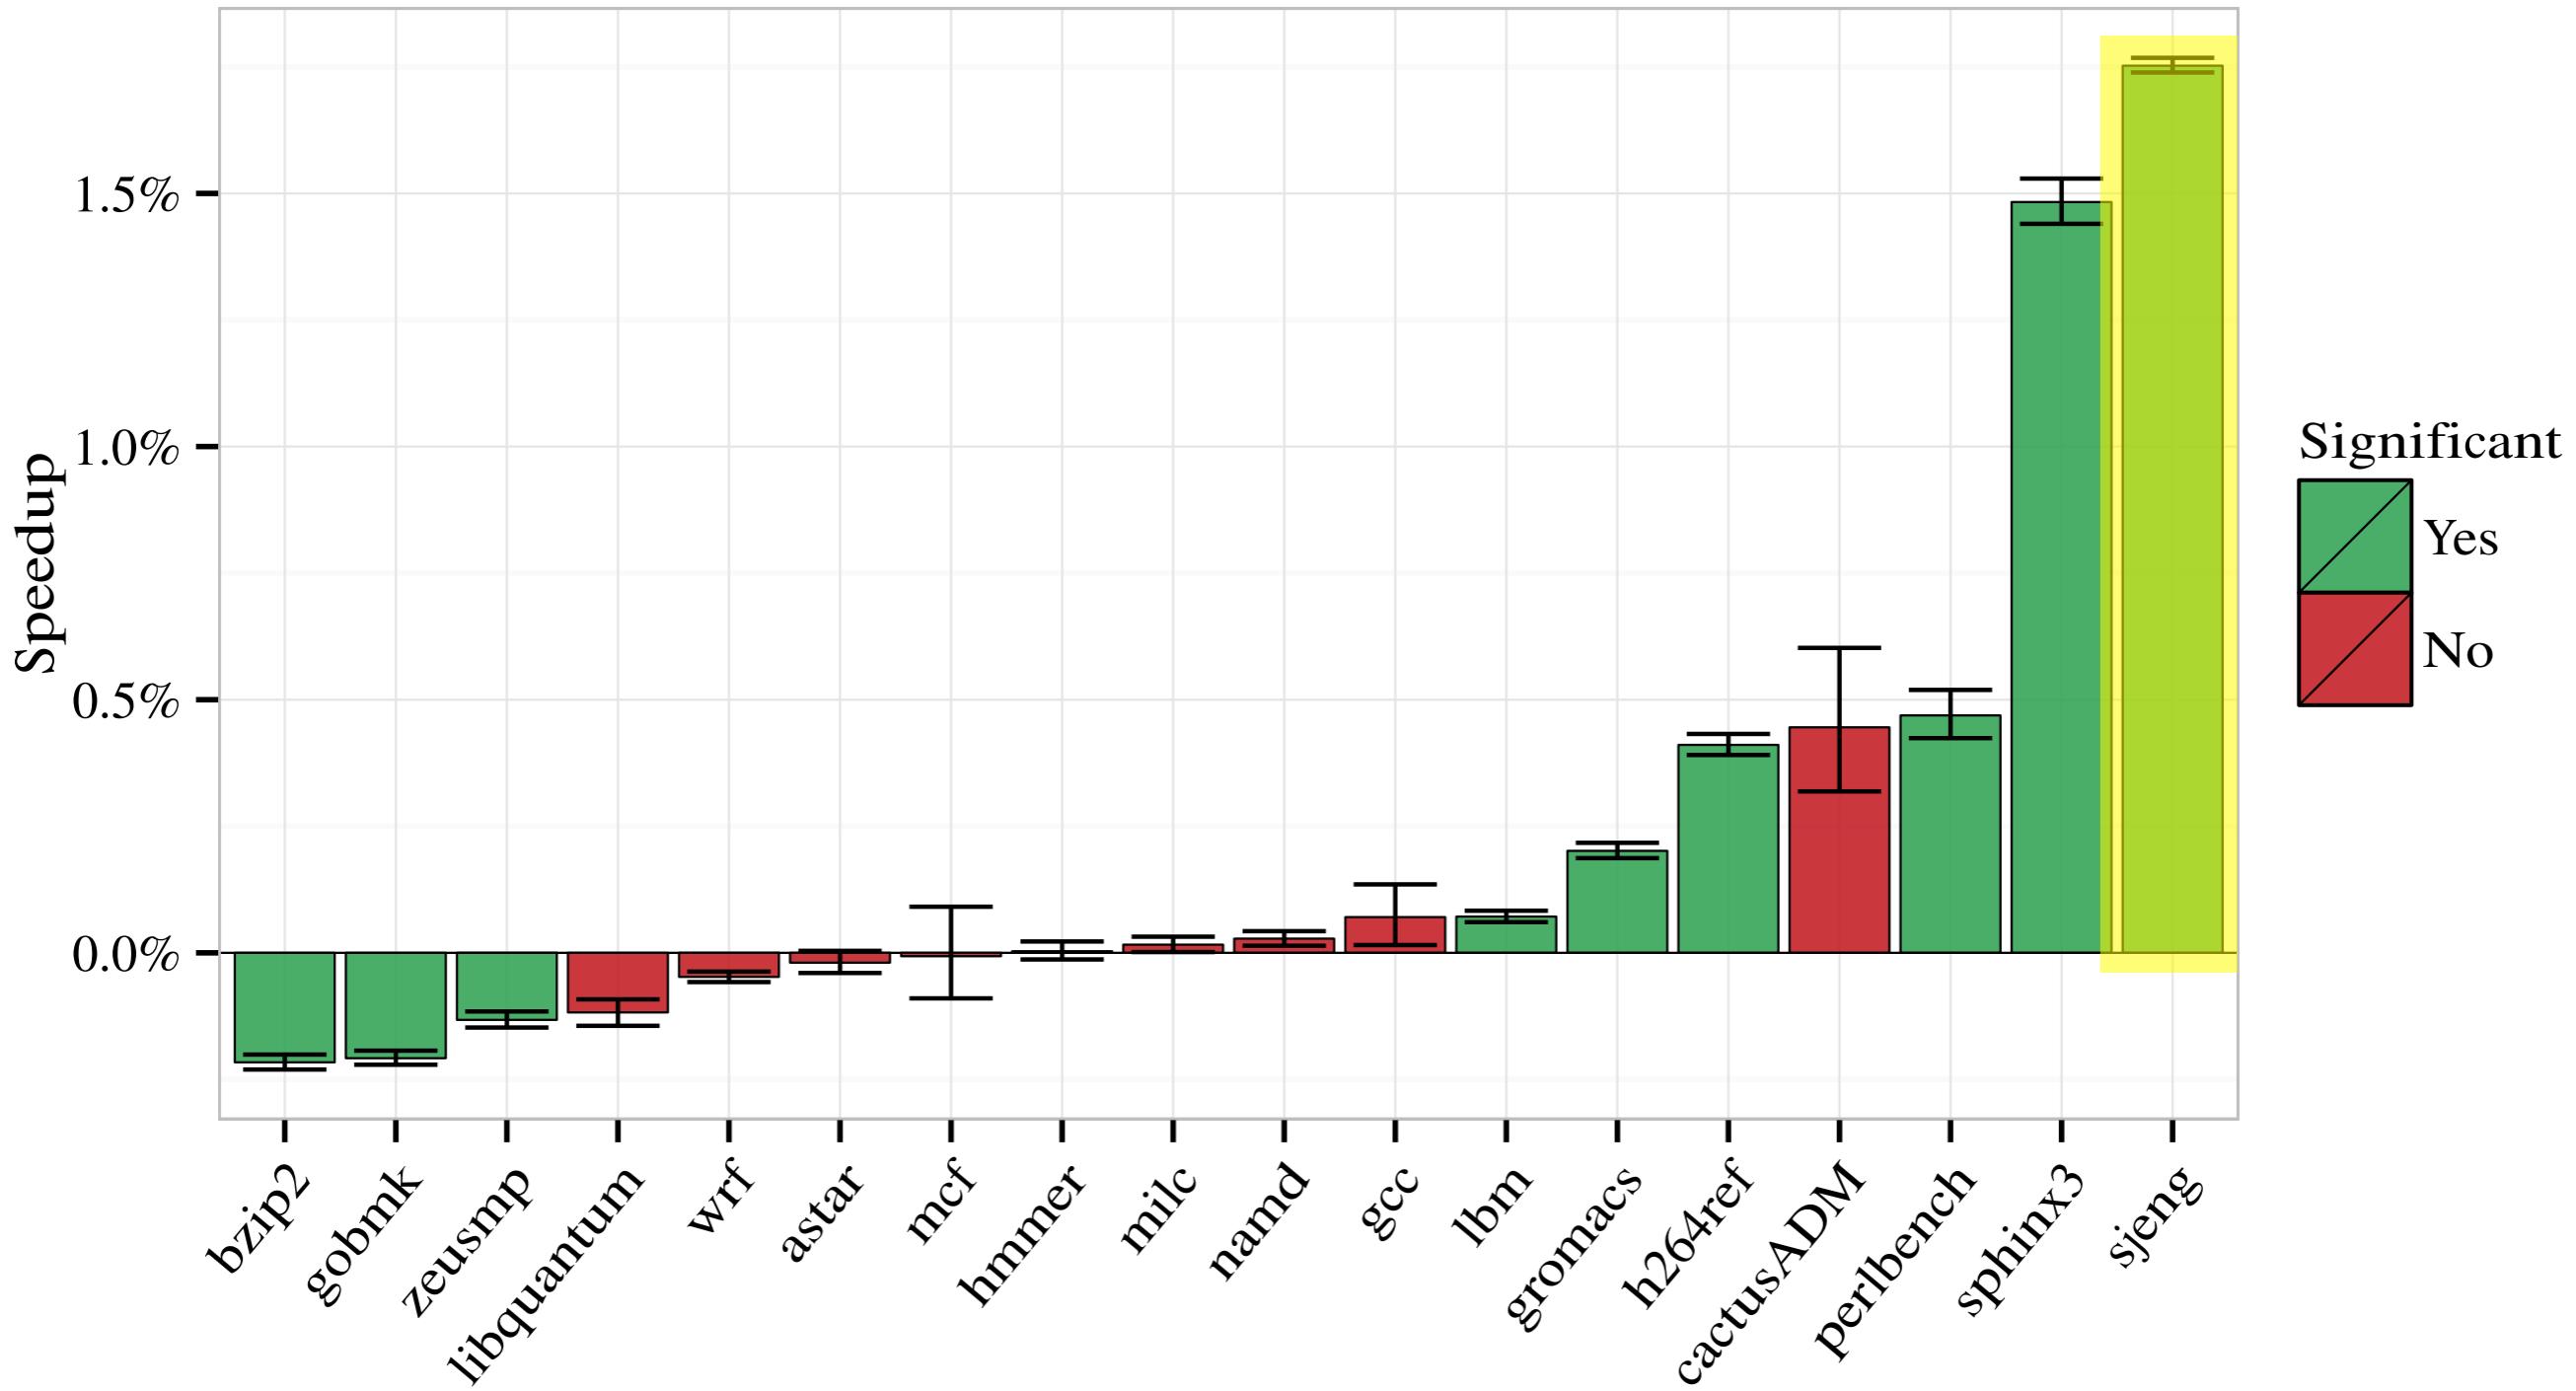
Speedup of -03 over -02



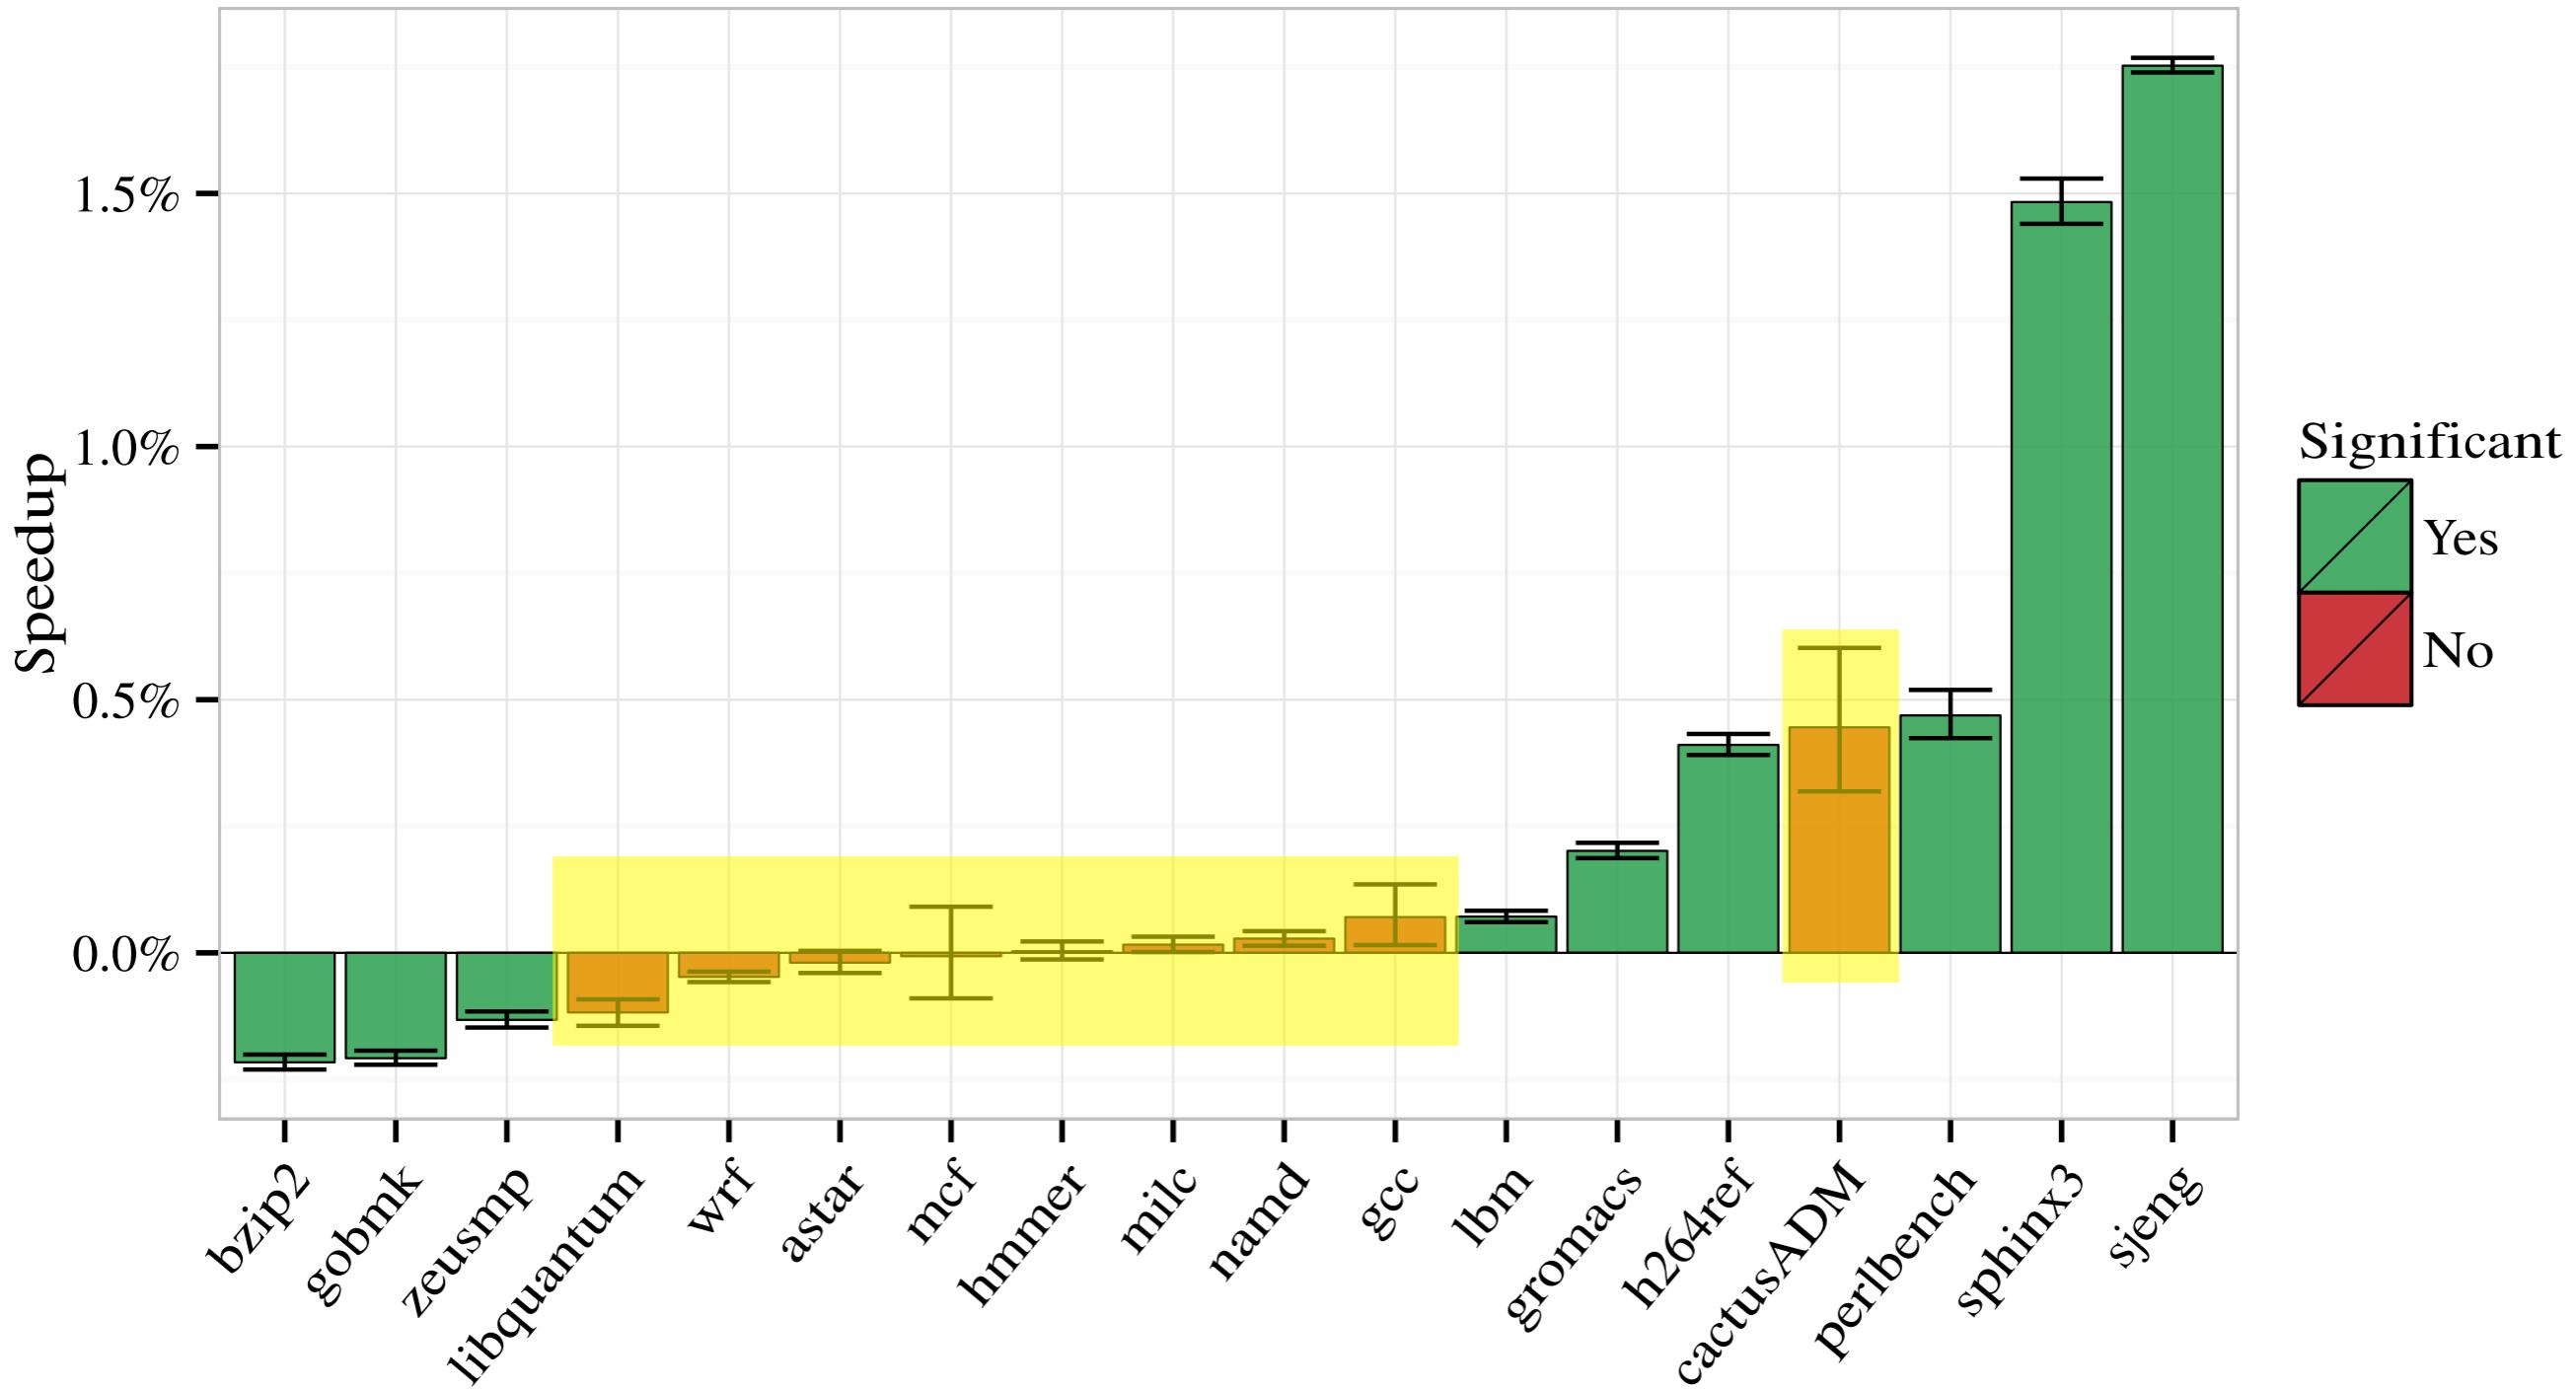
Speedup of -03 over -02



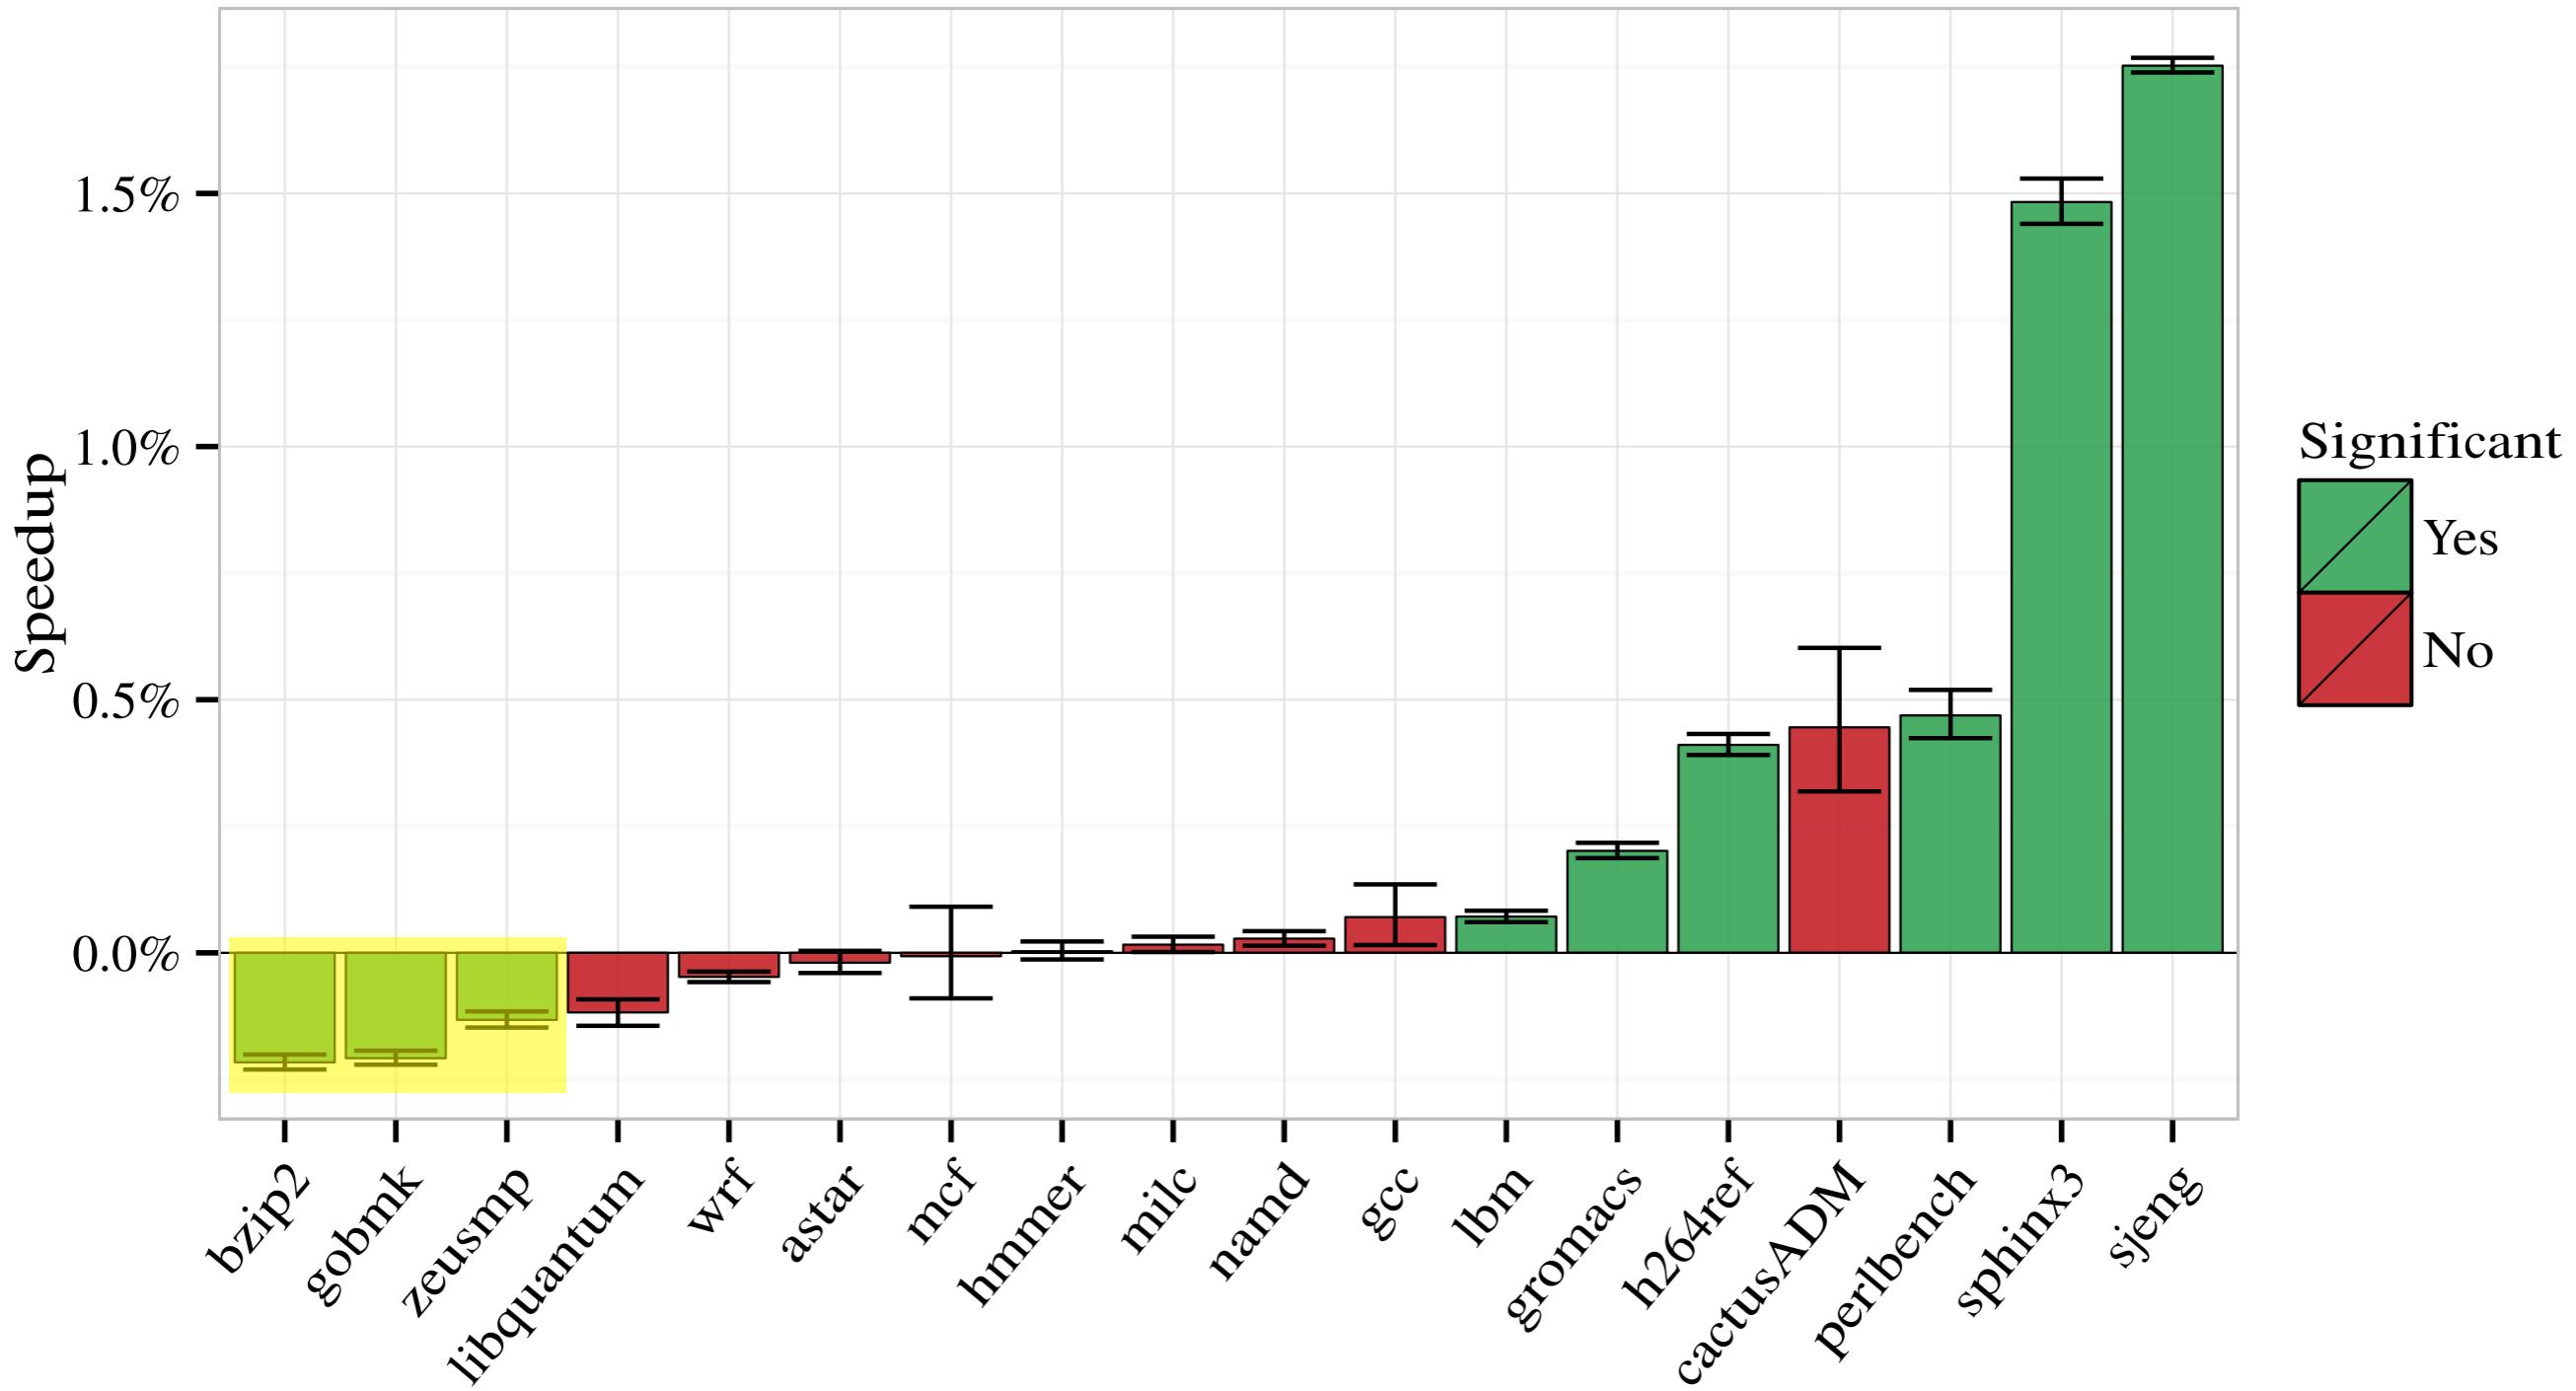
Speedup of -03 over -02



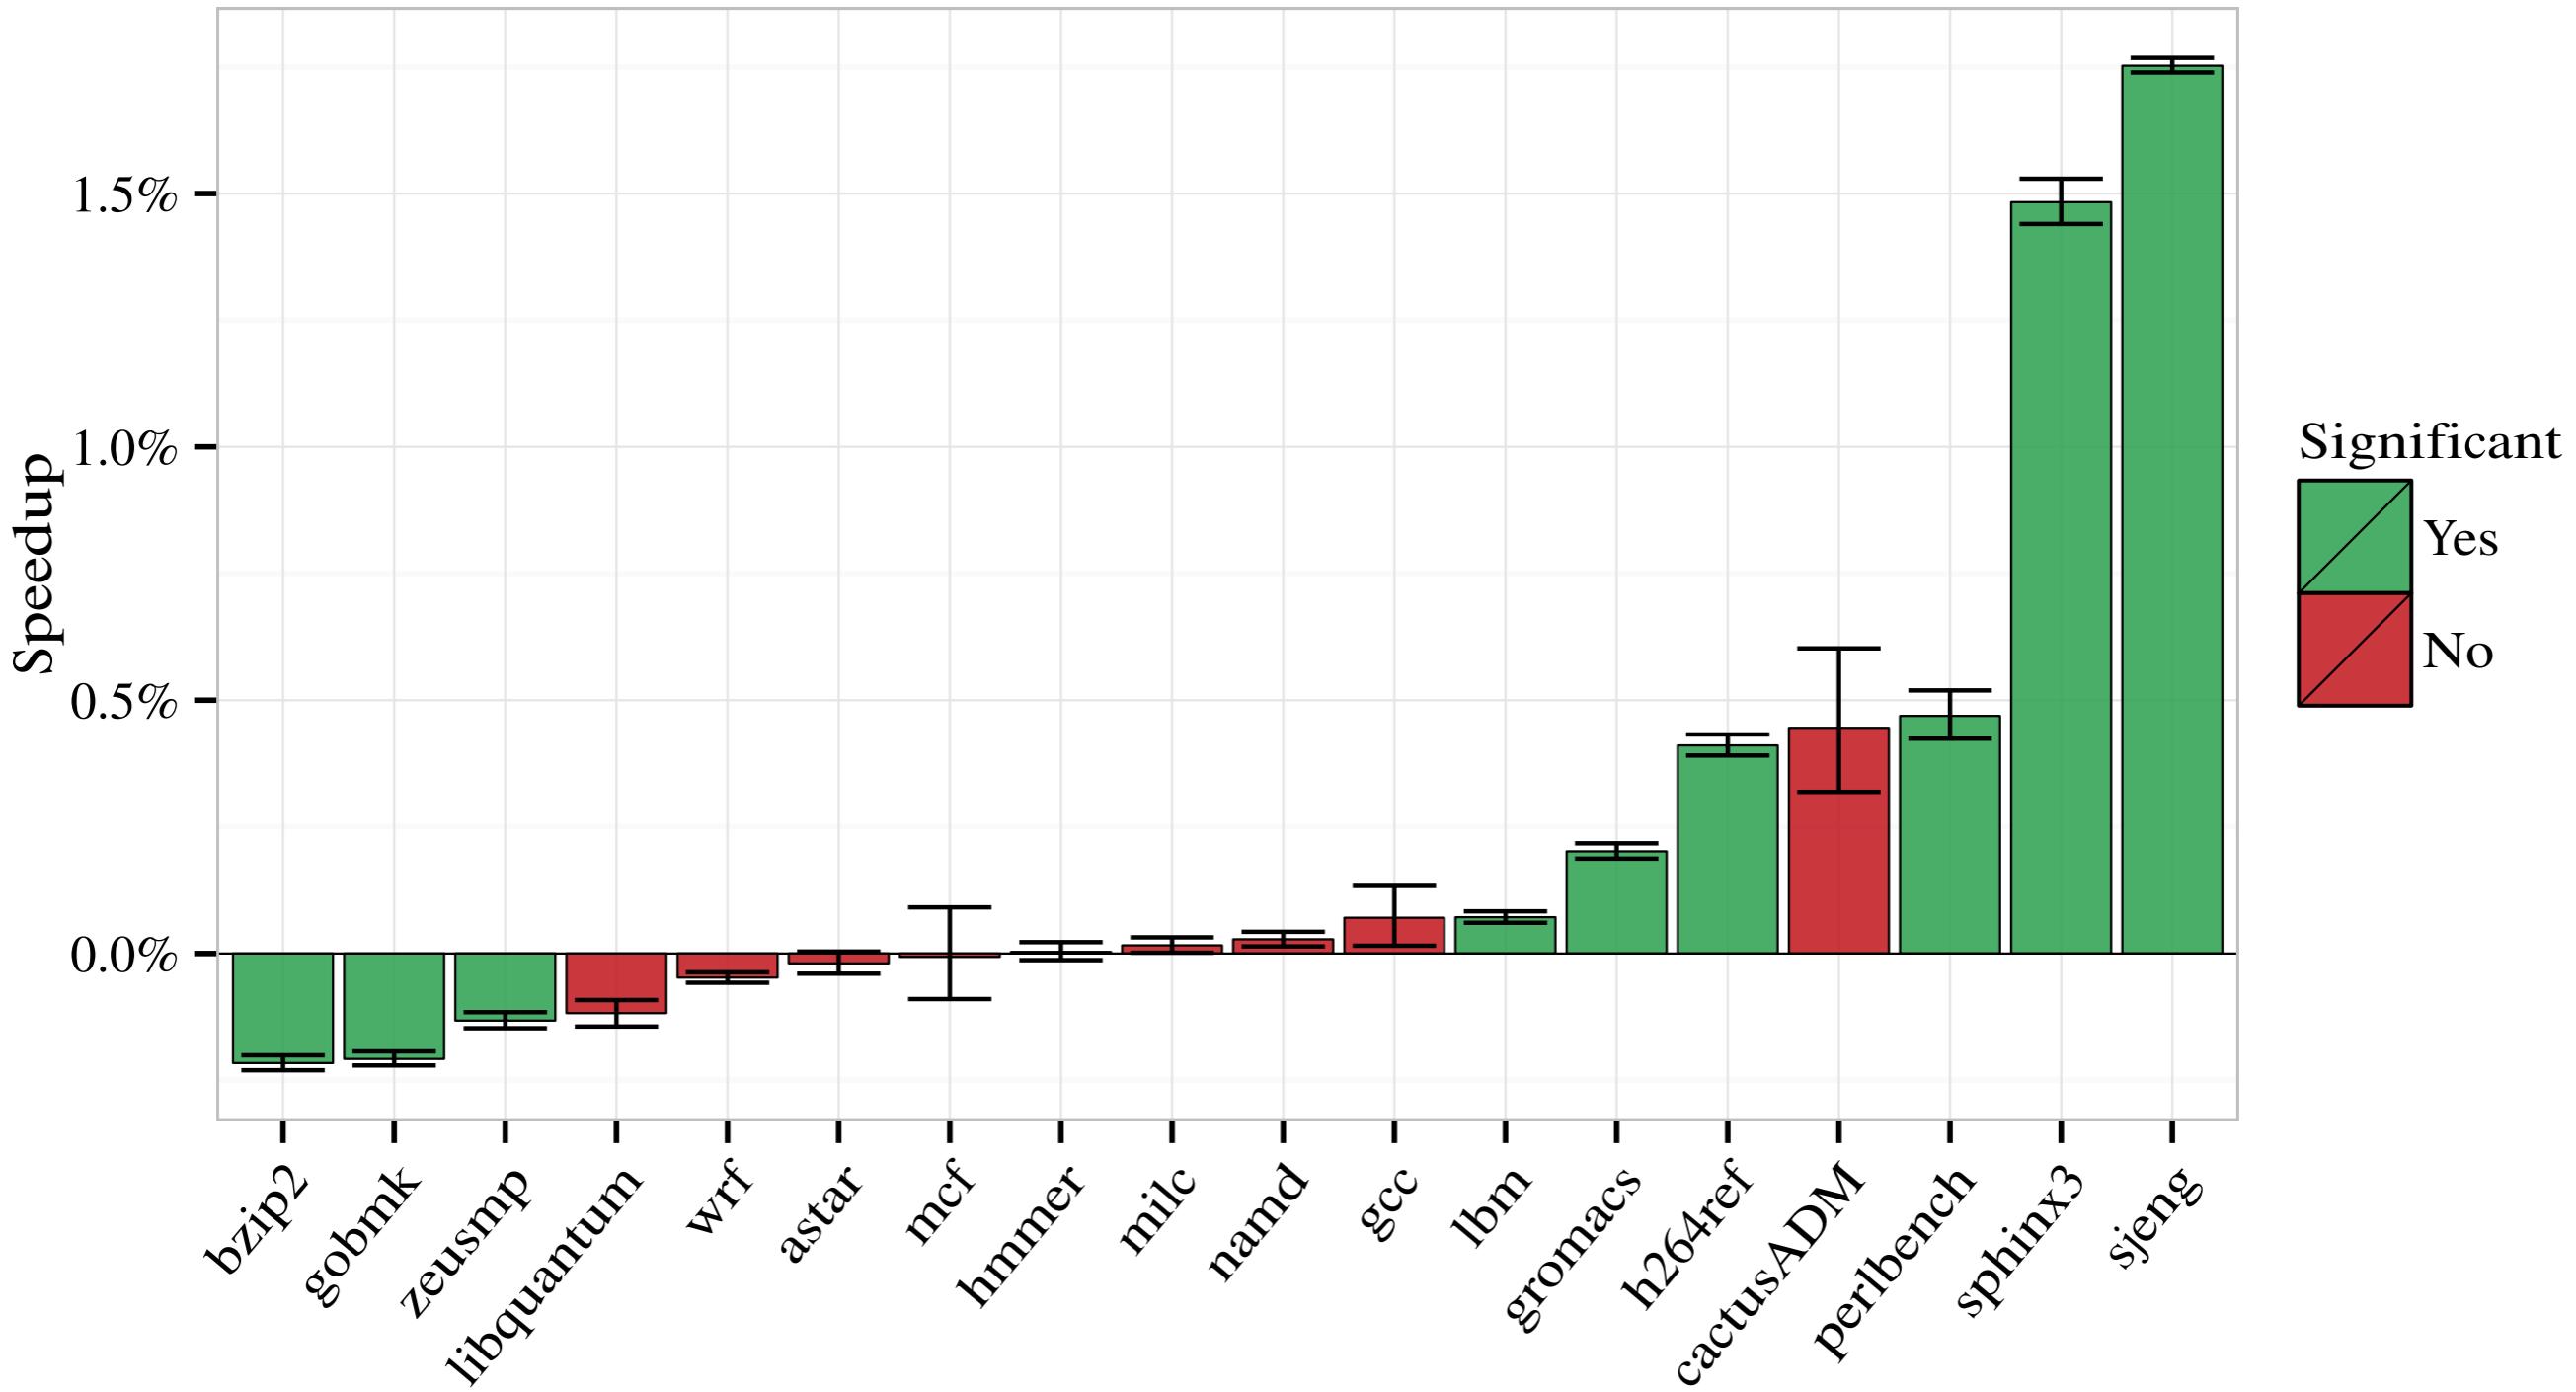
Speedup of -03 over -02



Speedup of -03 over -02



What do the results mean?



Comparing optimizations

-O2 ×30

-O3 ×30

Comparing optimizations

-O2 ×30

Ibm ×30

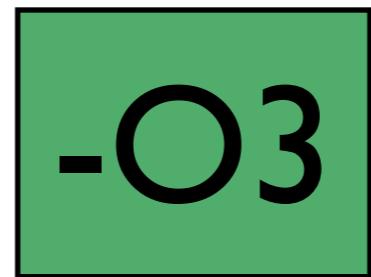
-O3 ×30

Ibm ×30

Comparing optimizations



×30



×30

lbm ×30

lbm ×30

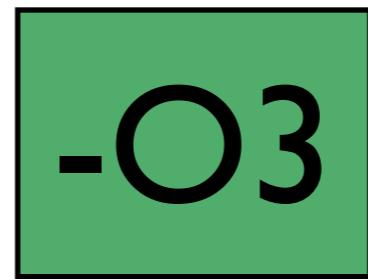
astar ×30

astar ×30

Comparing optimizations



×30



×30

lbm ×30

lbm ×30

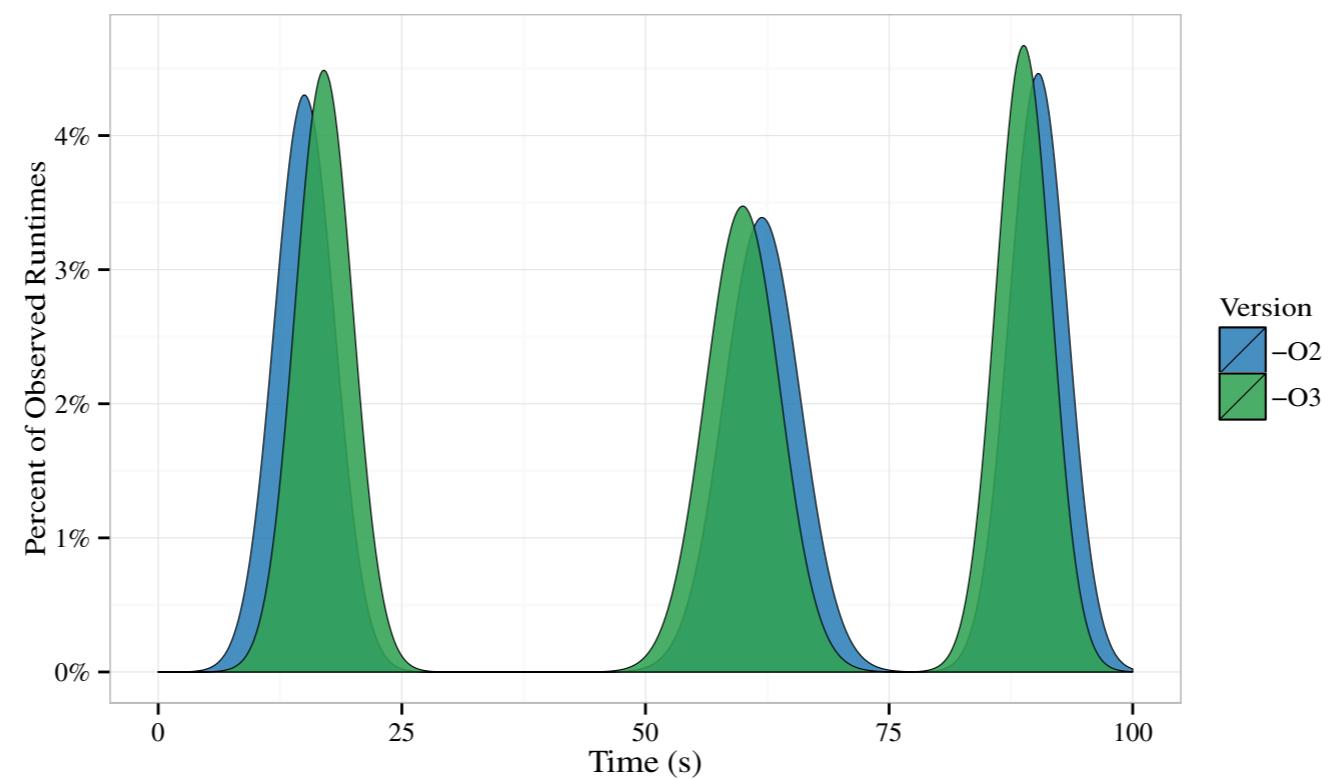
astar ×30

astar ×30

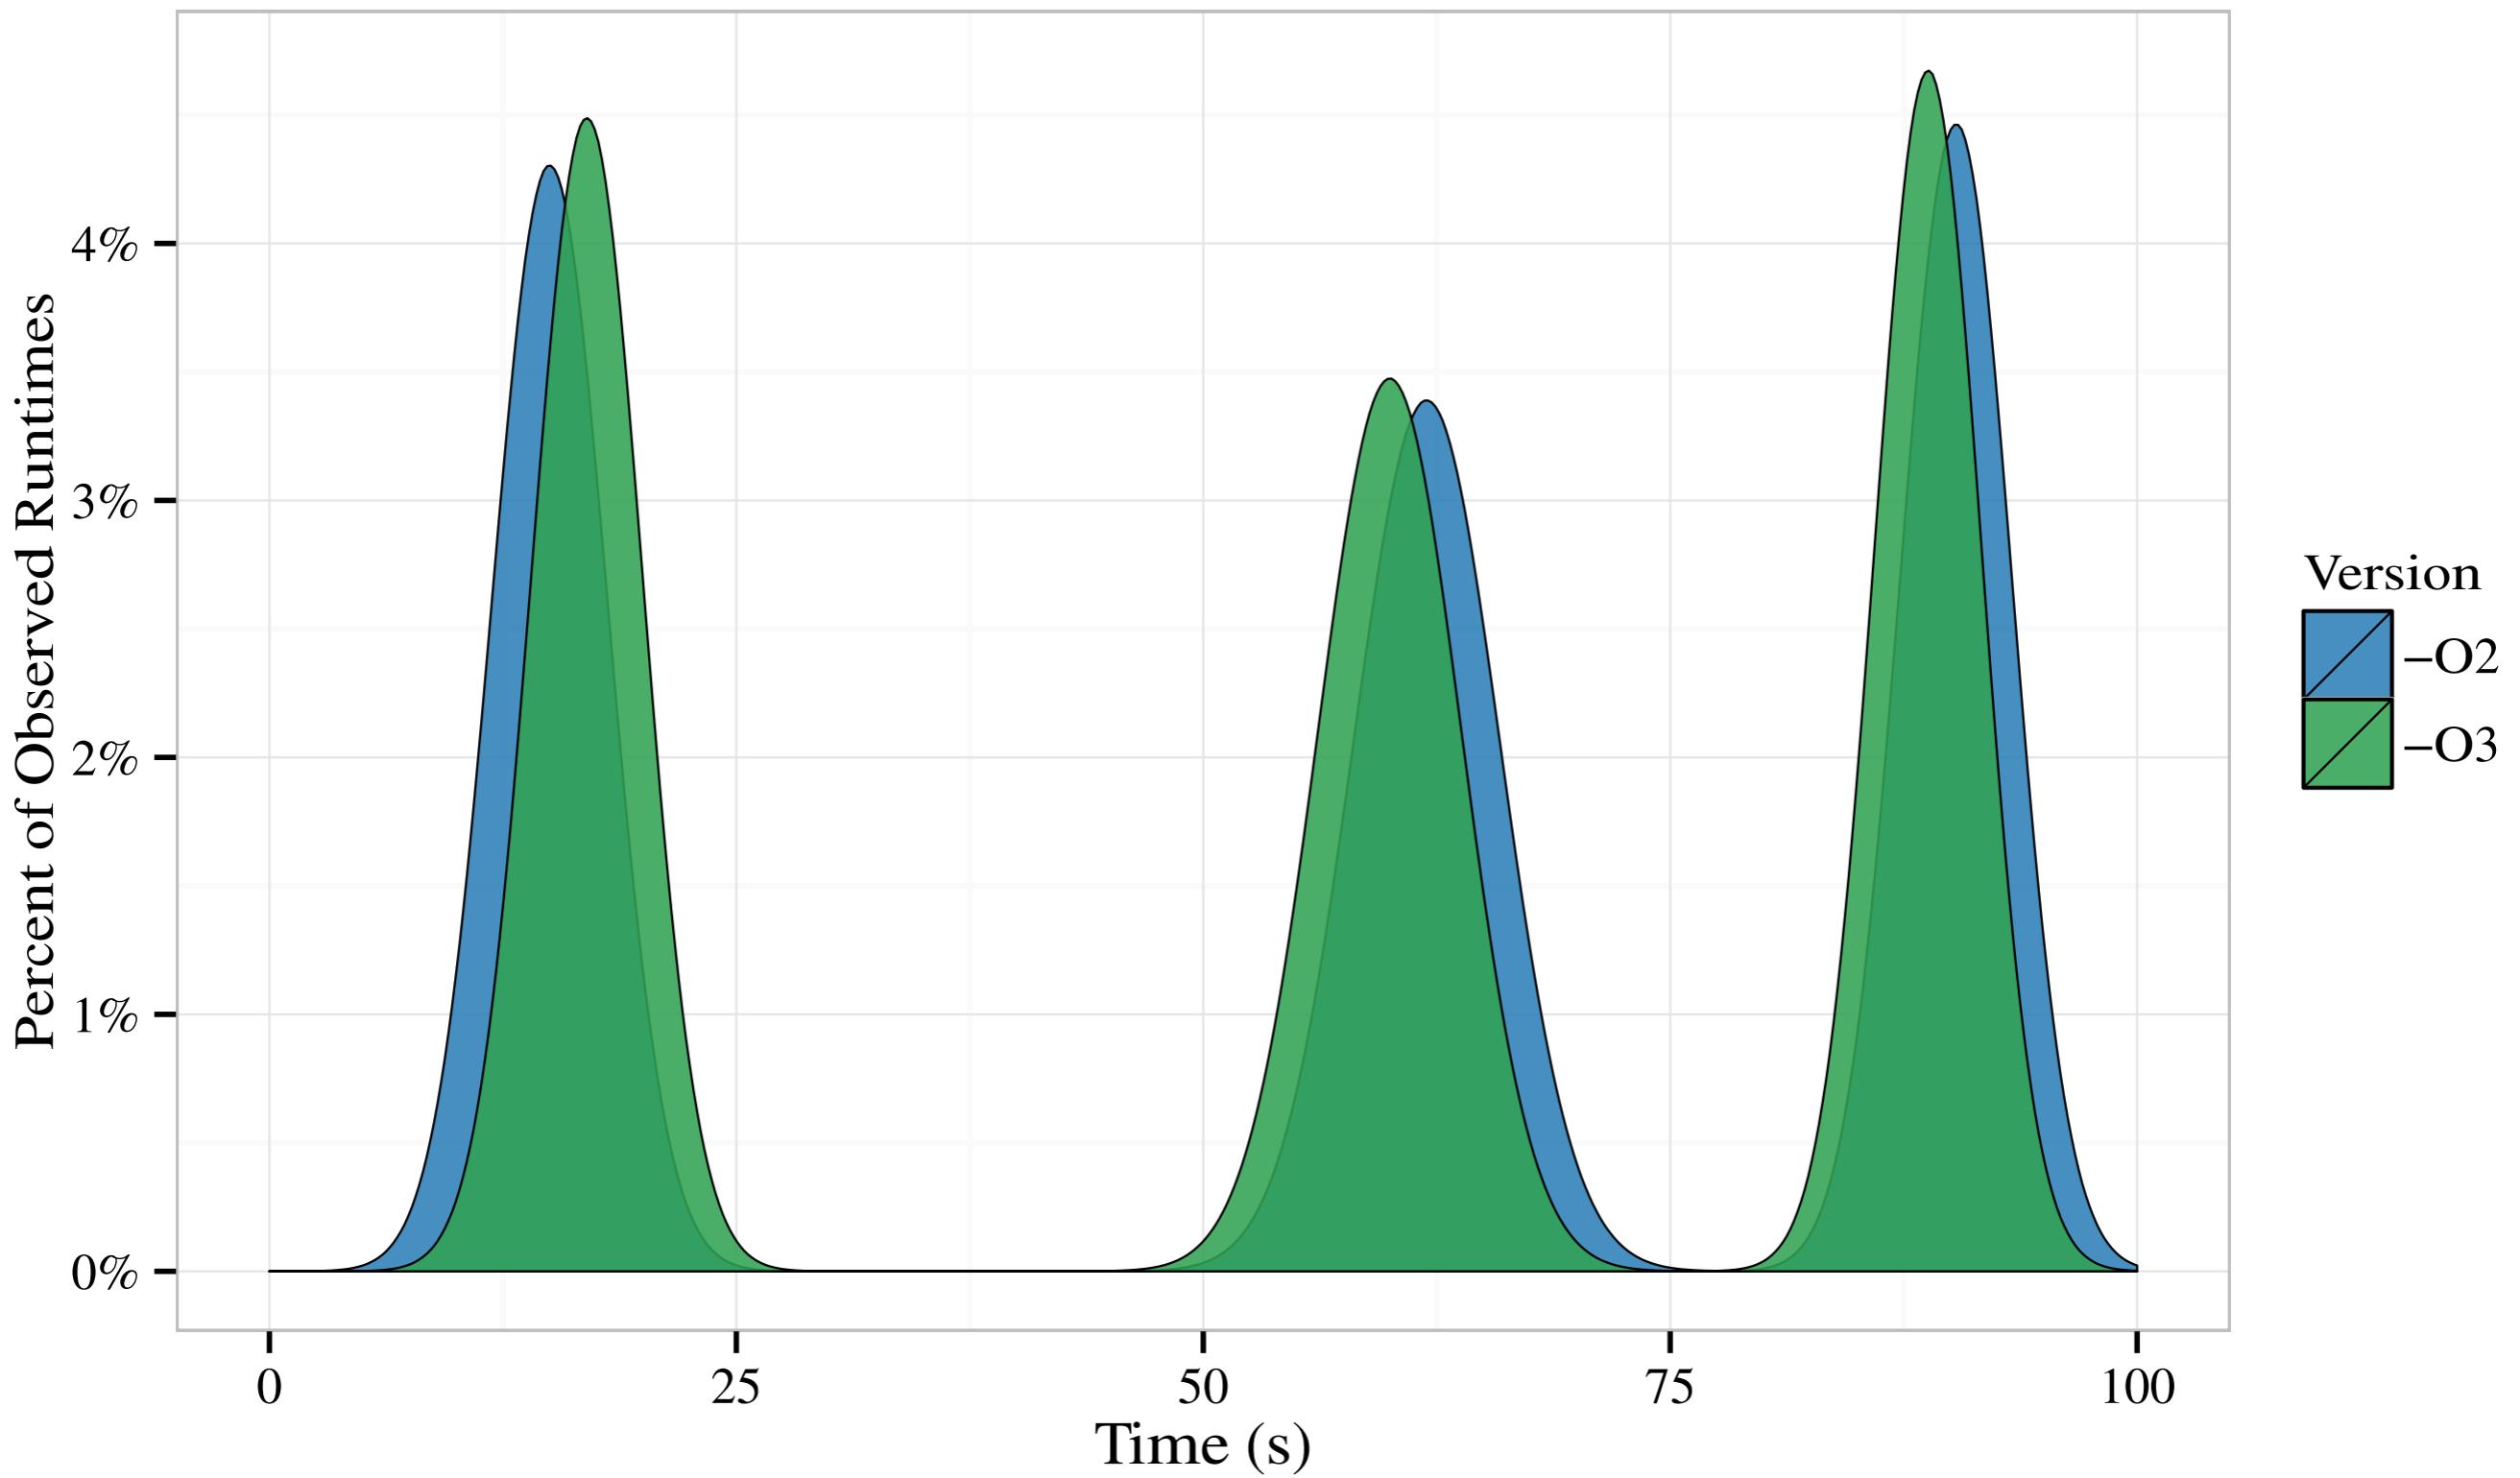
•••

Comparing optimizations

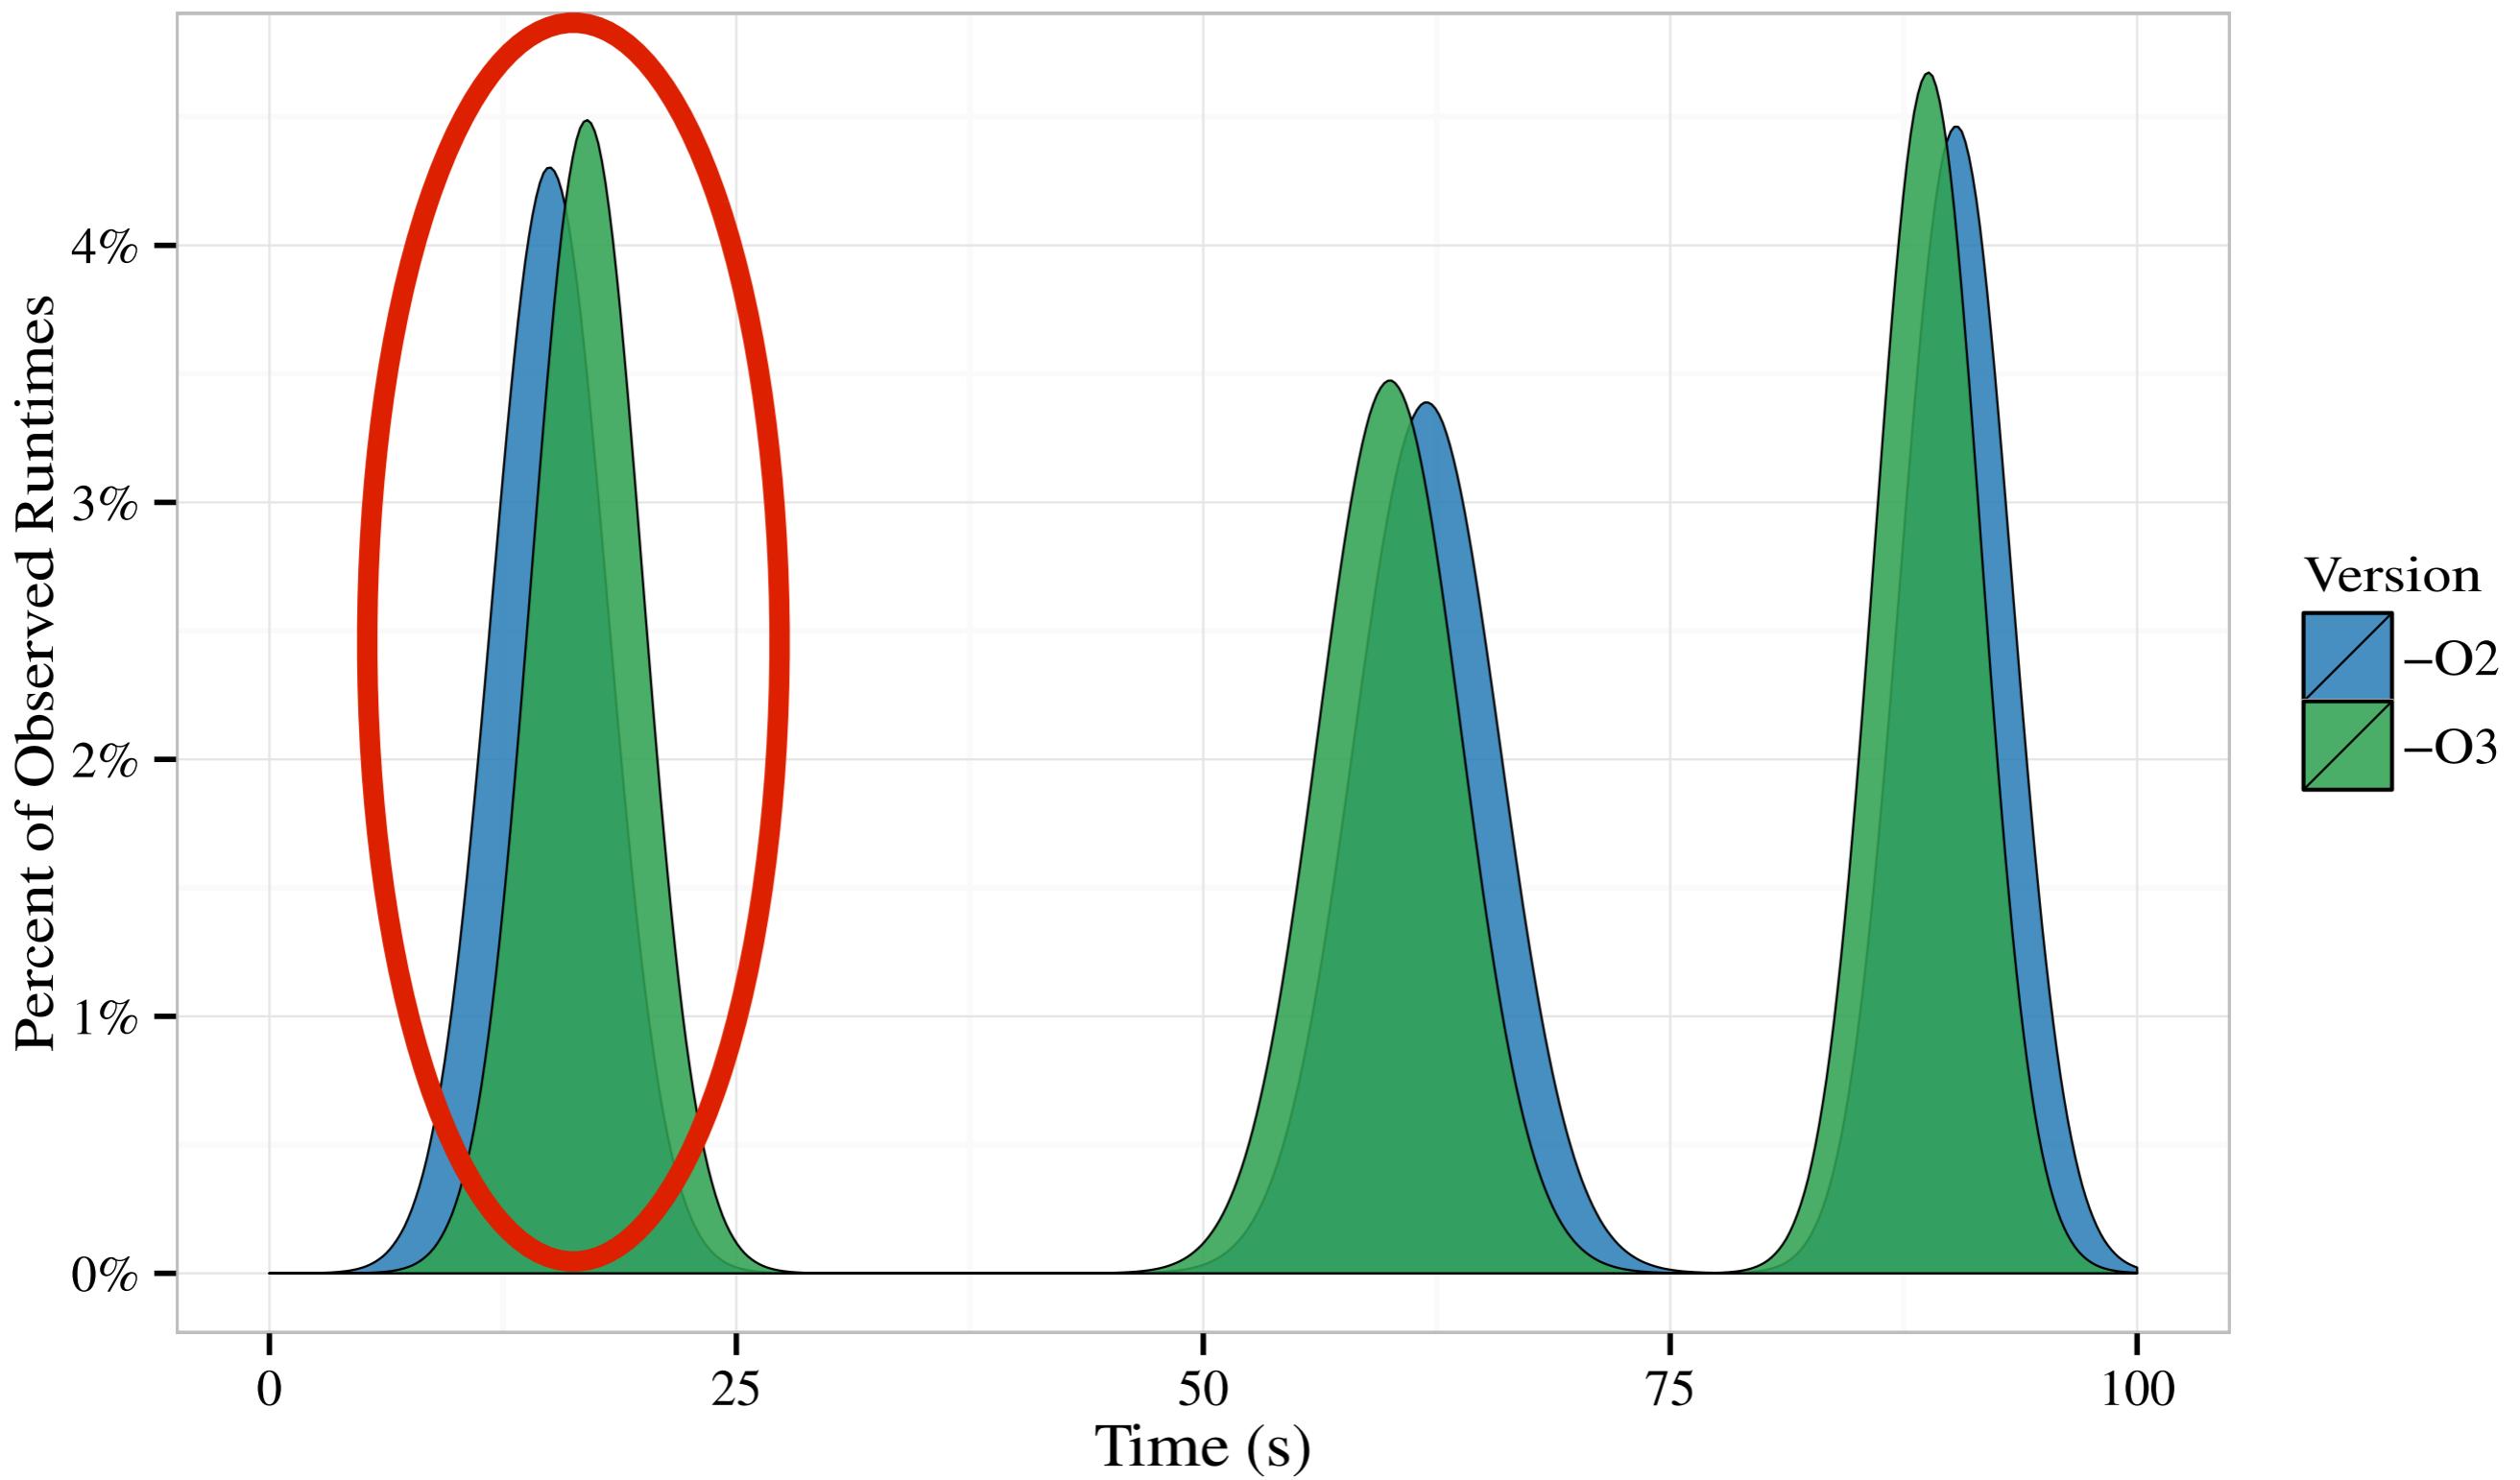
-O2 ×30 **-O3 ×30**



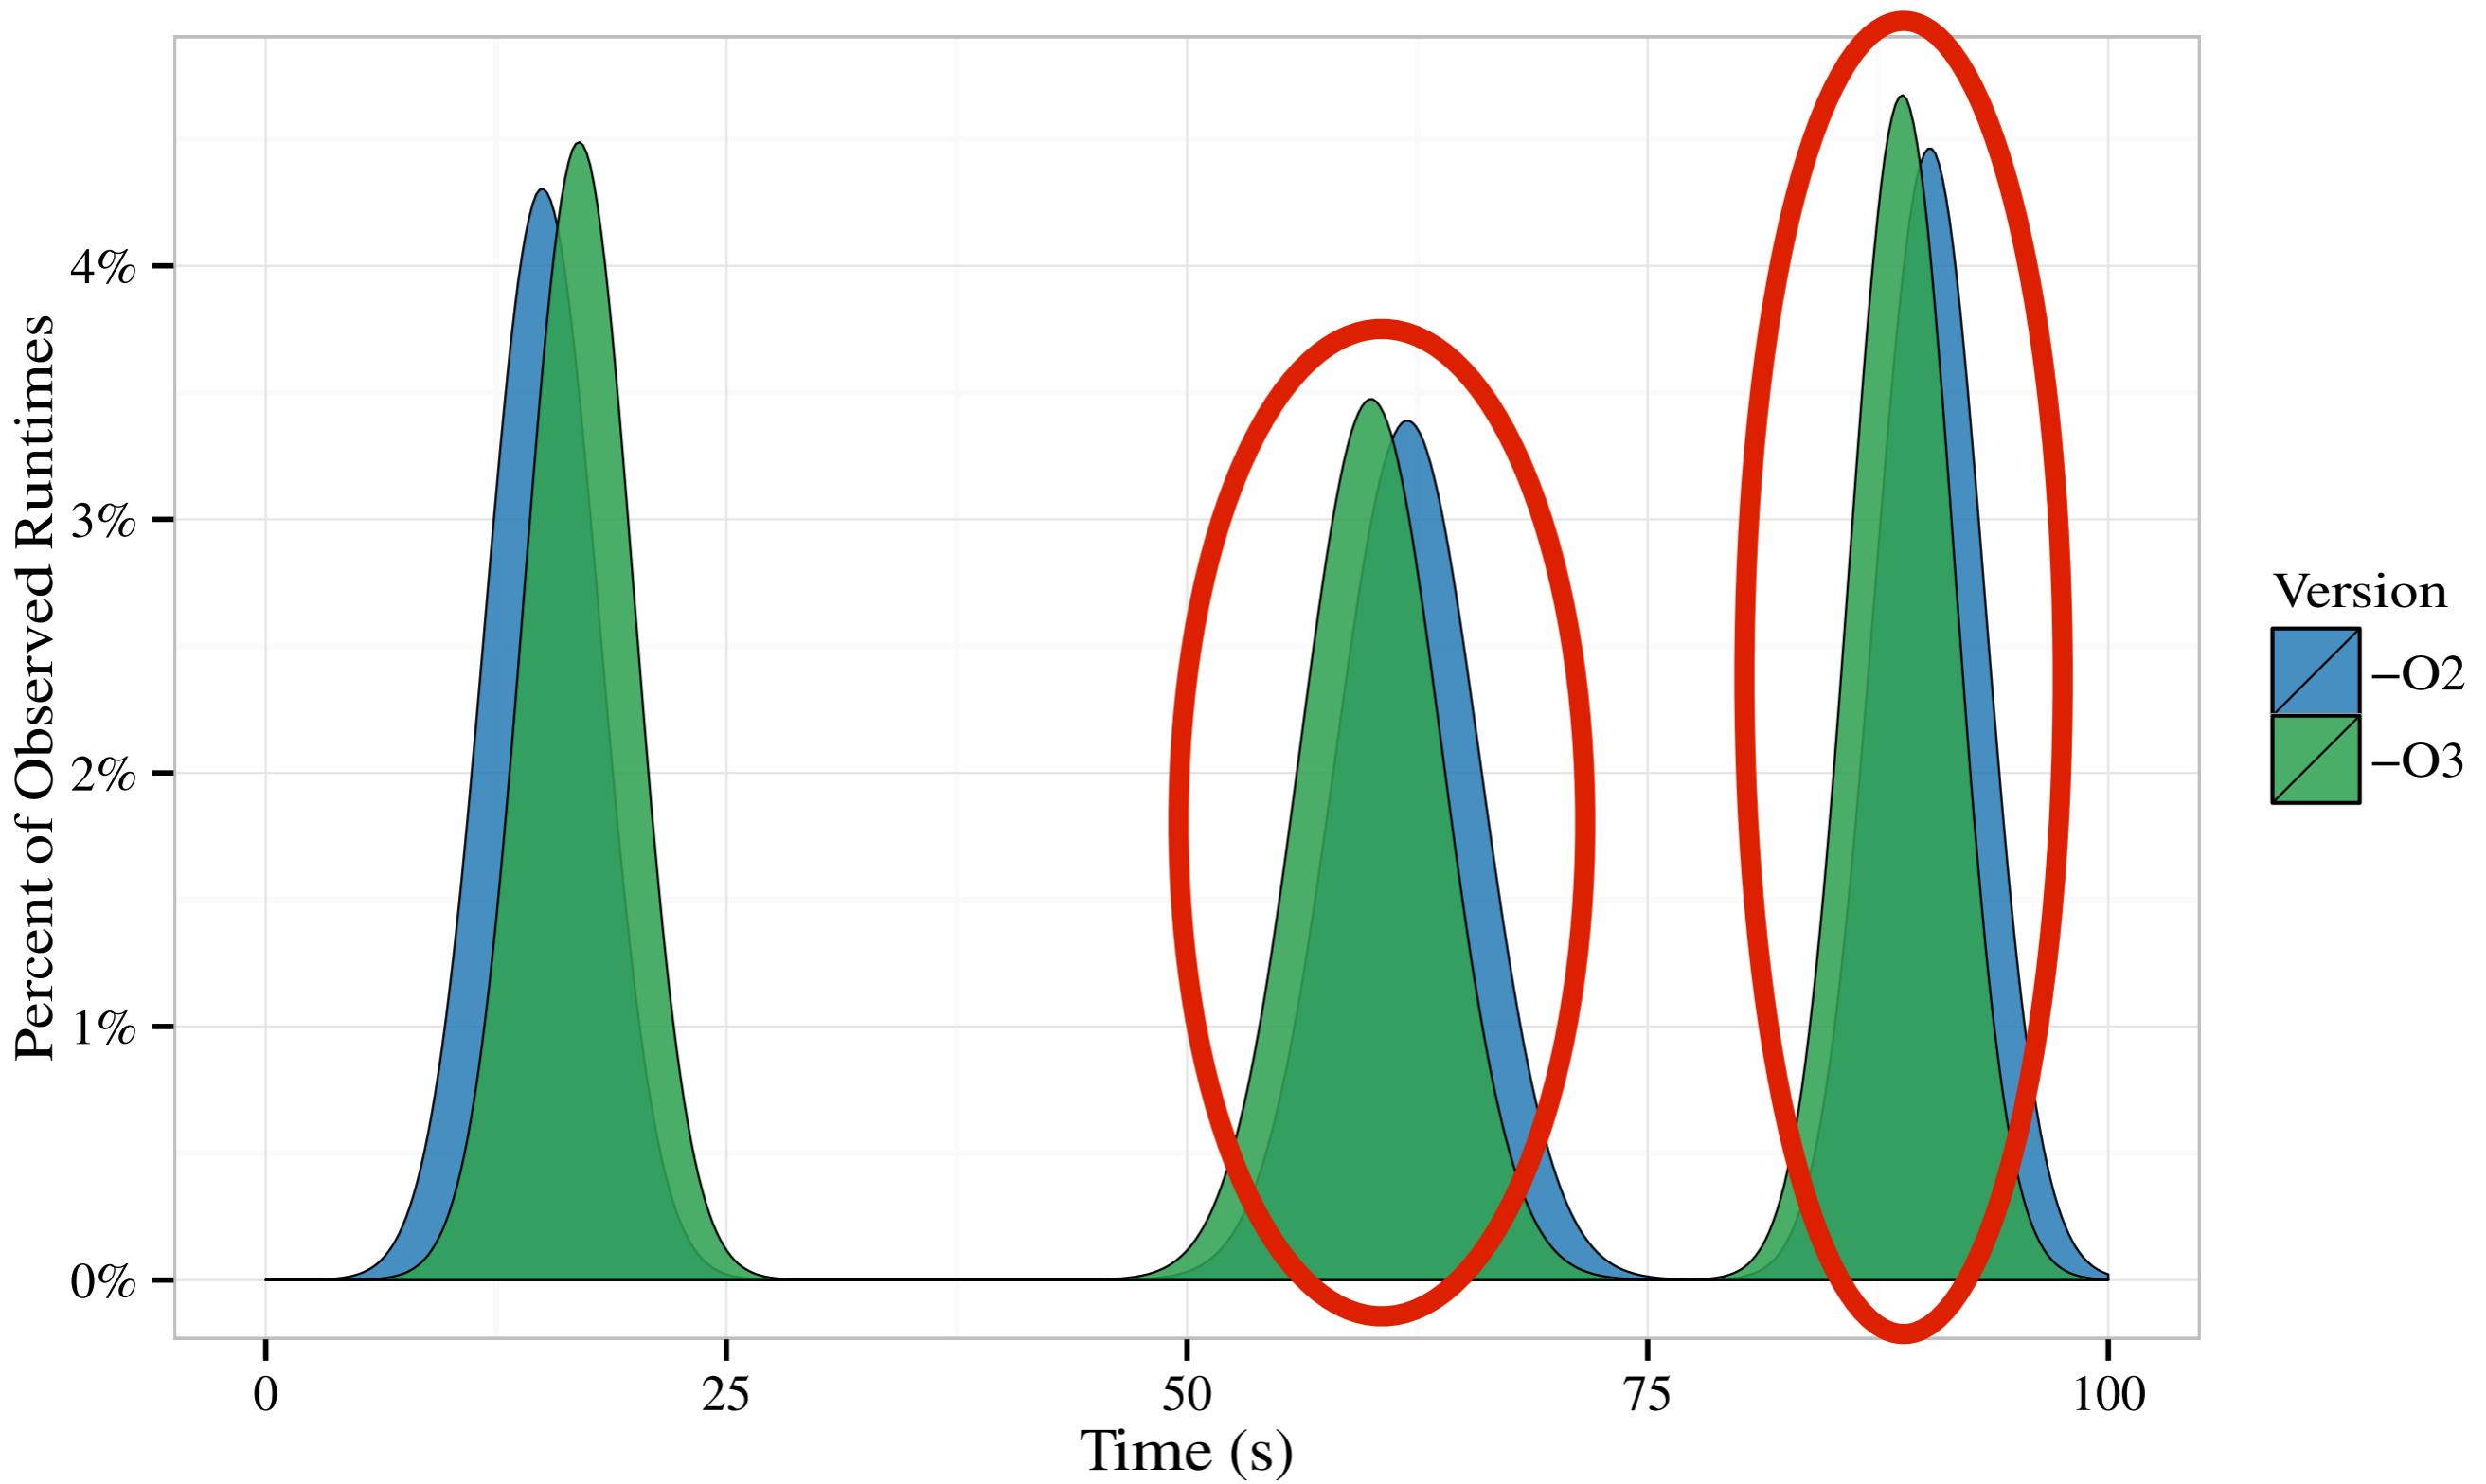
Is **-O3** faster than **-O2**?



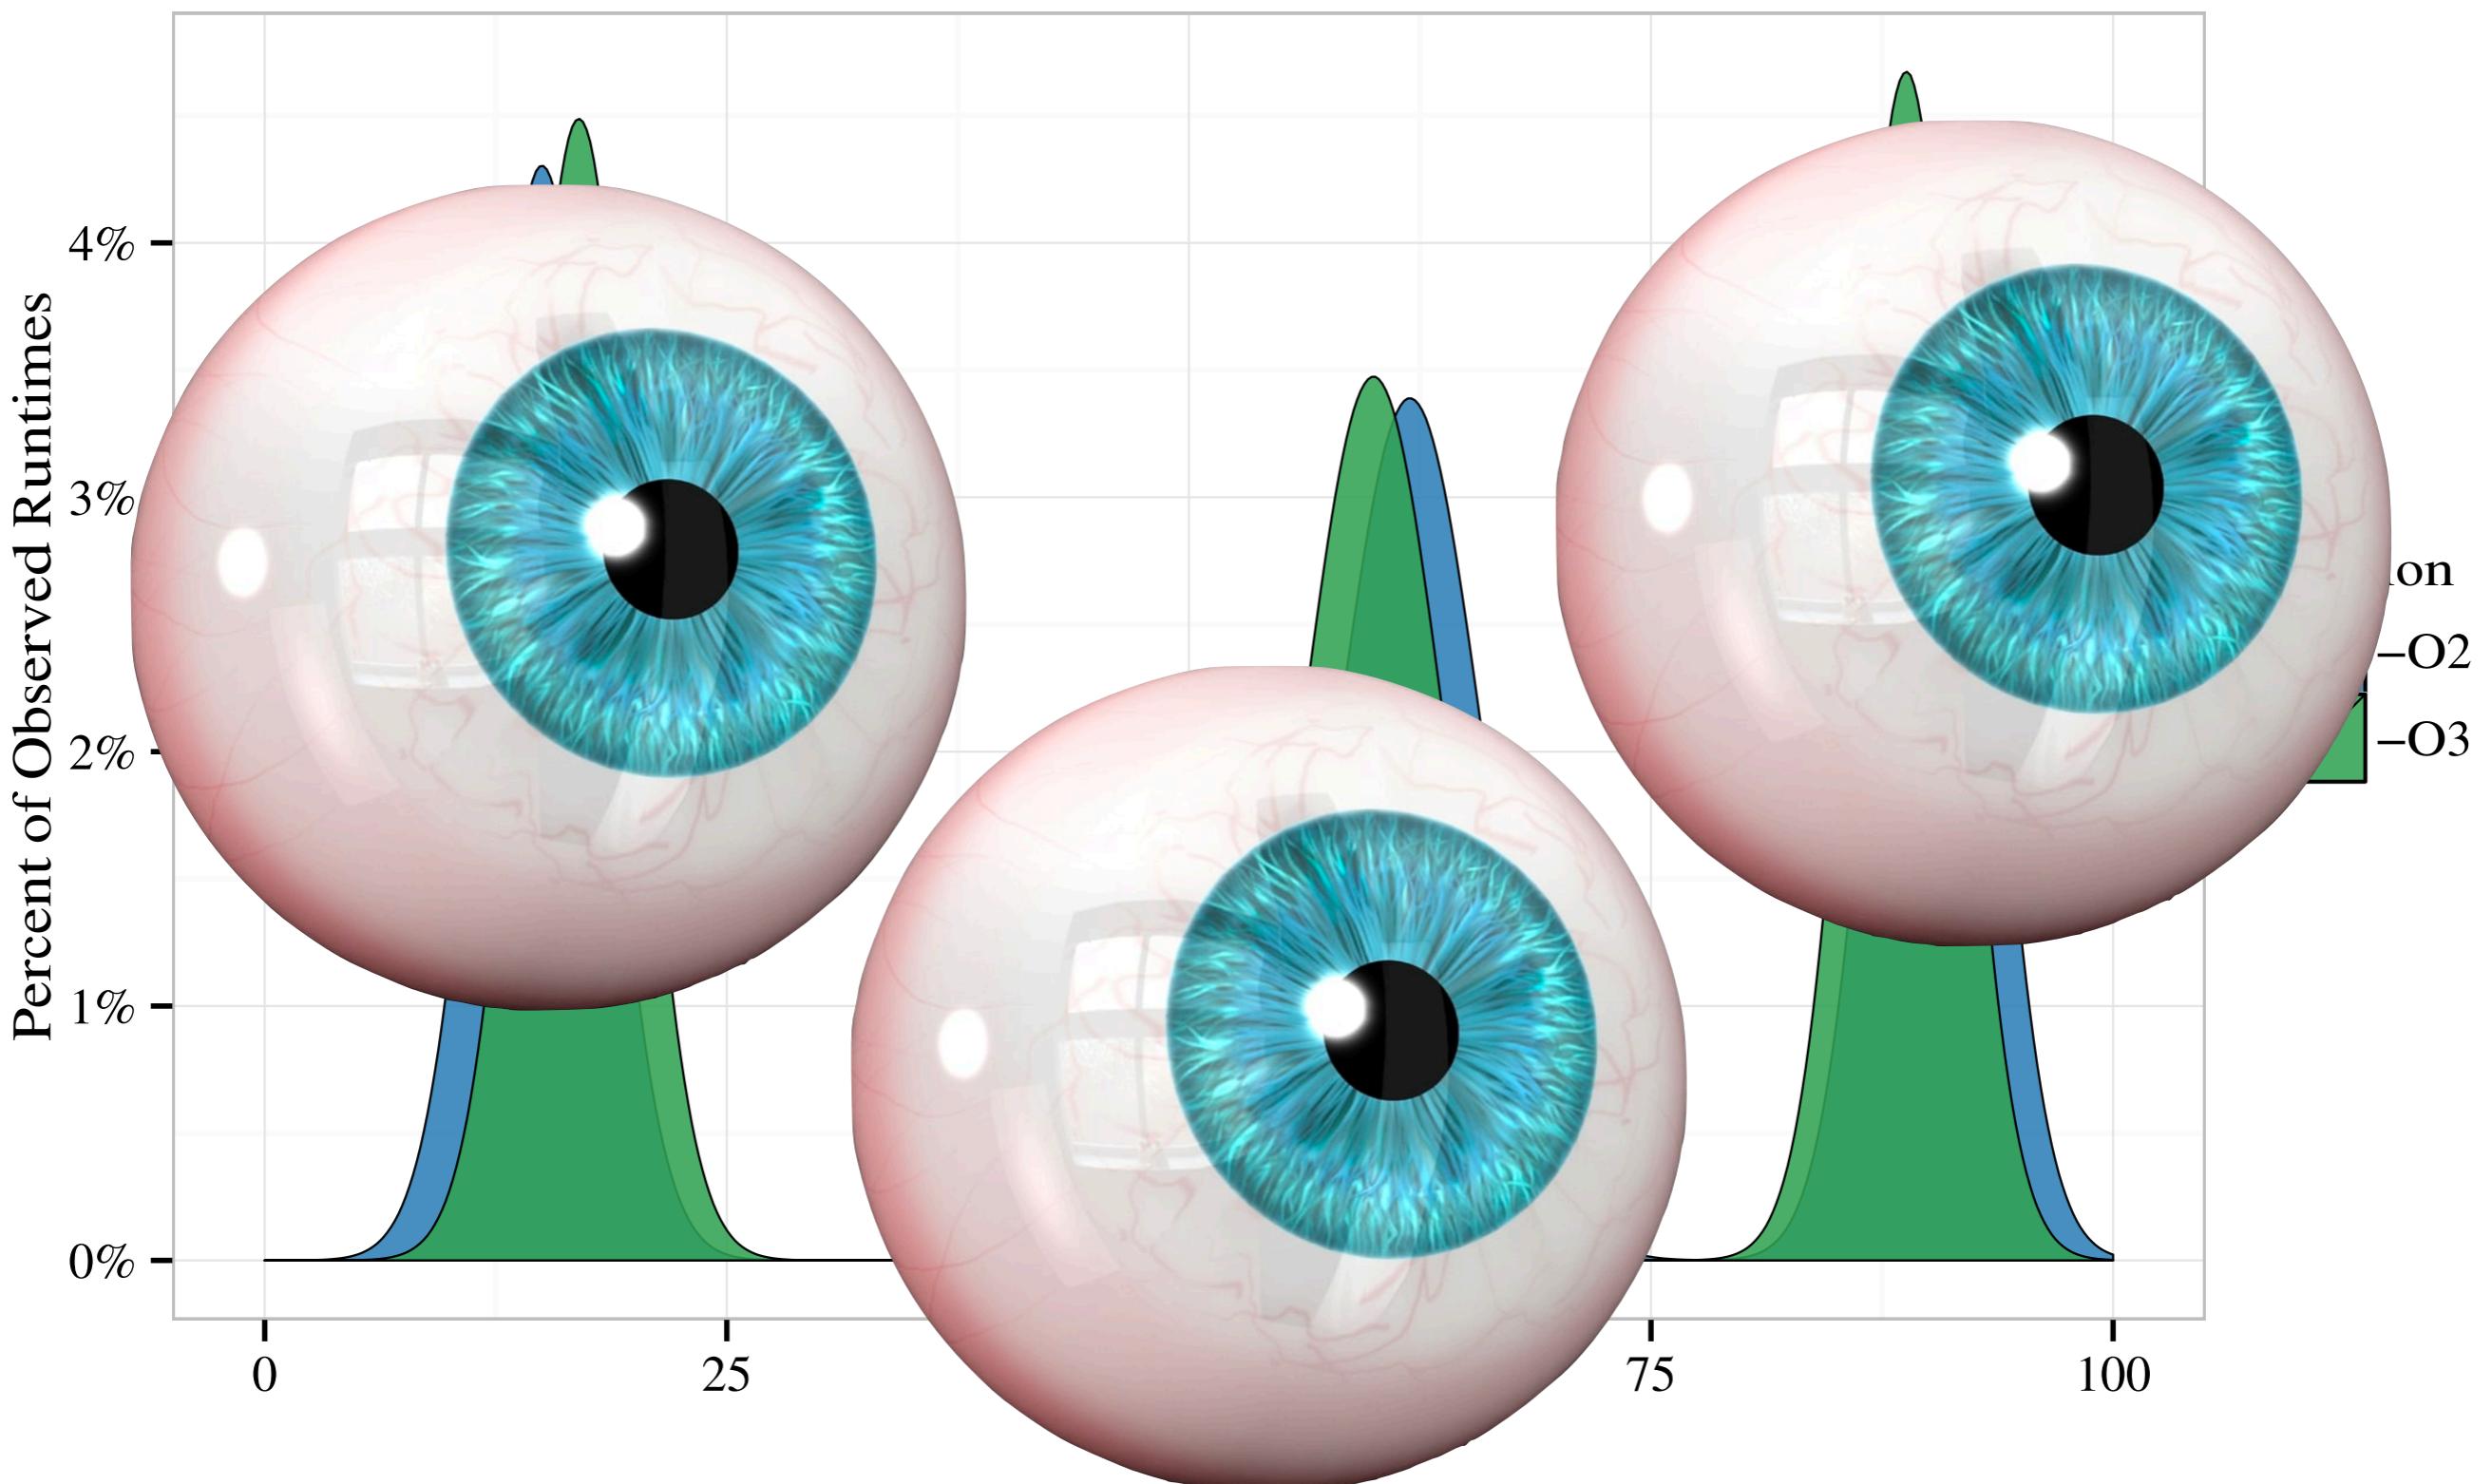
Is **-O3** faster than **-O2**?



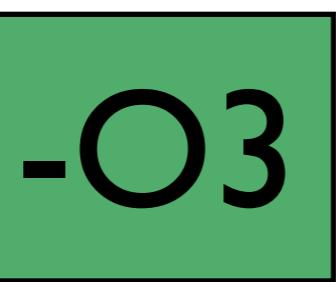
Is **-O3** faster than **-O2**?



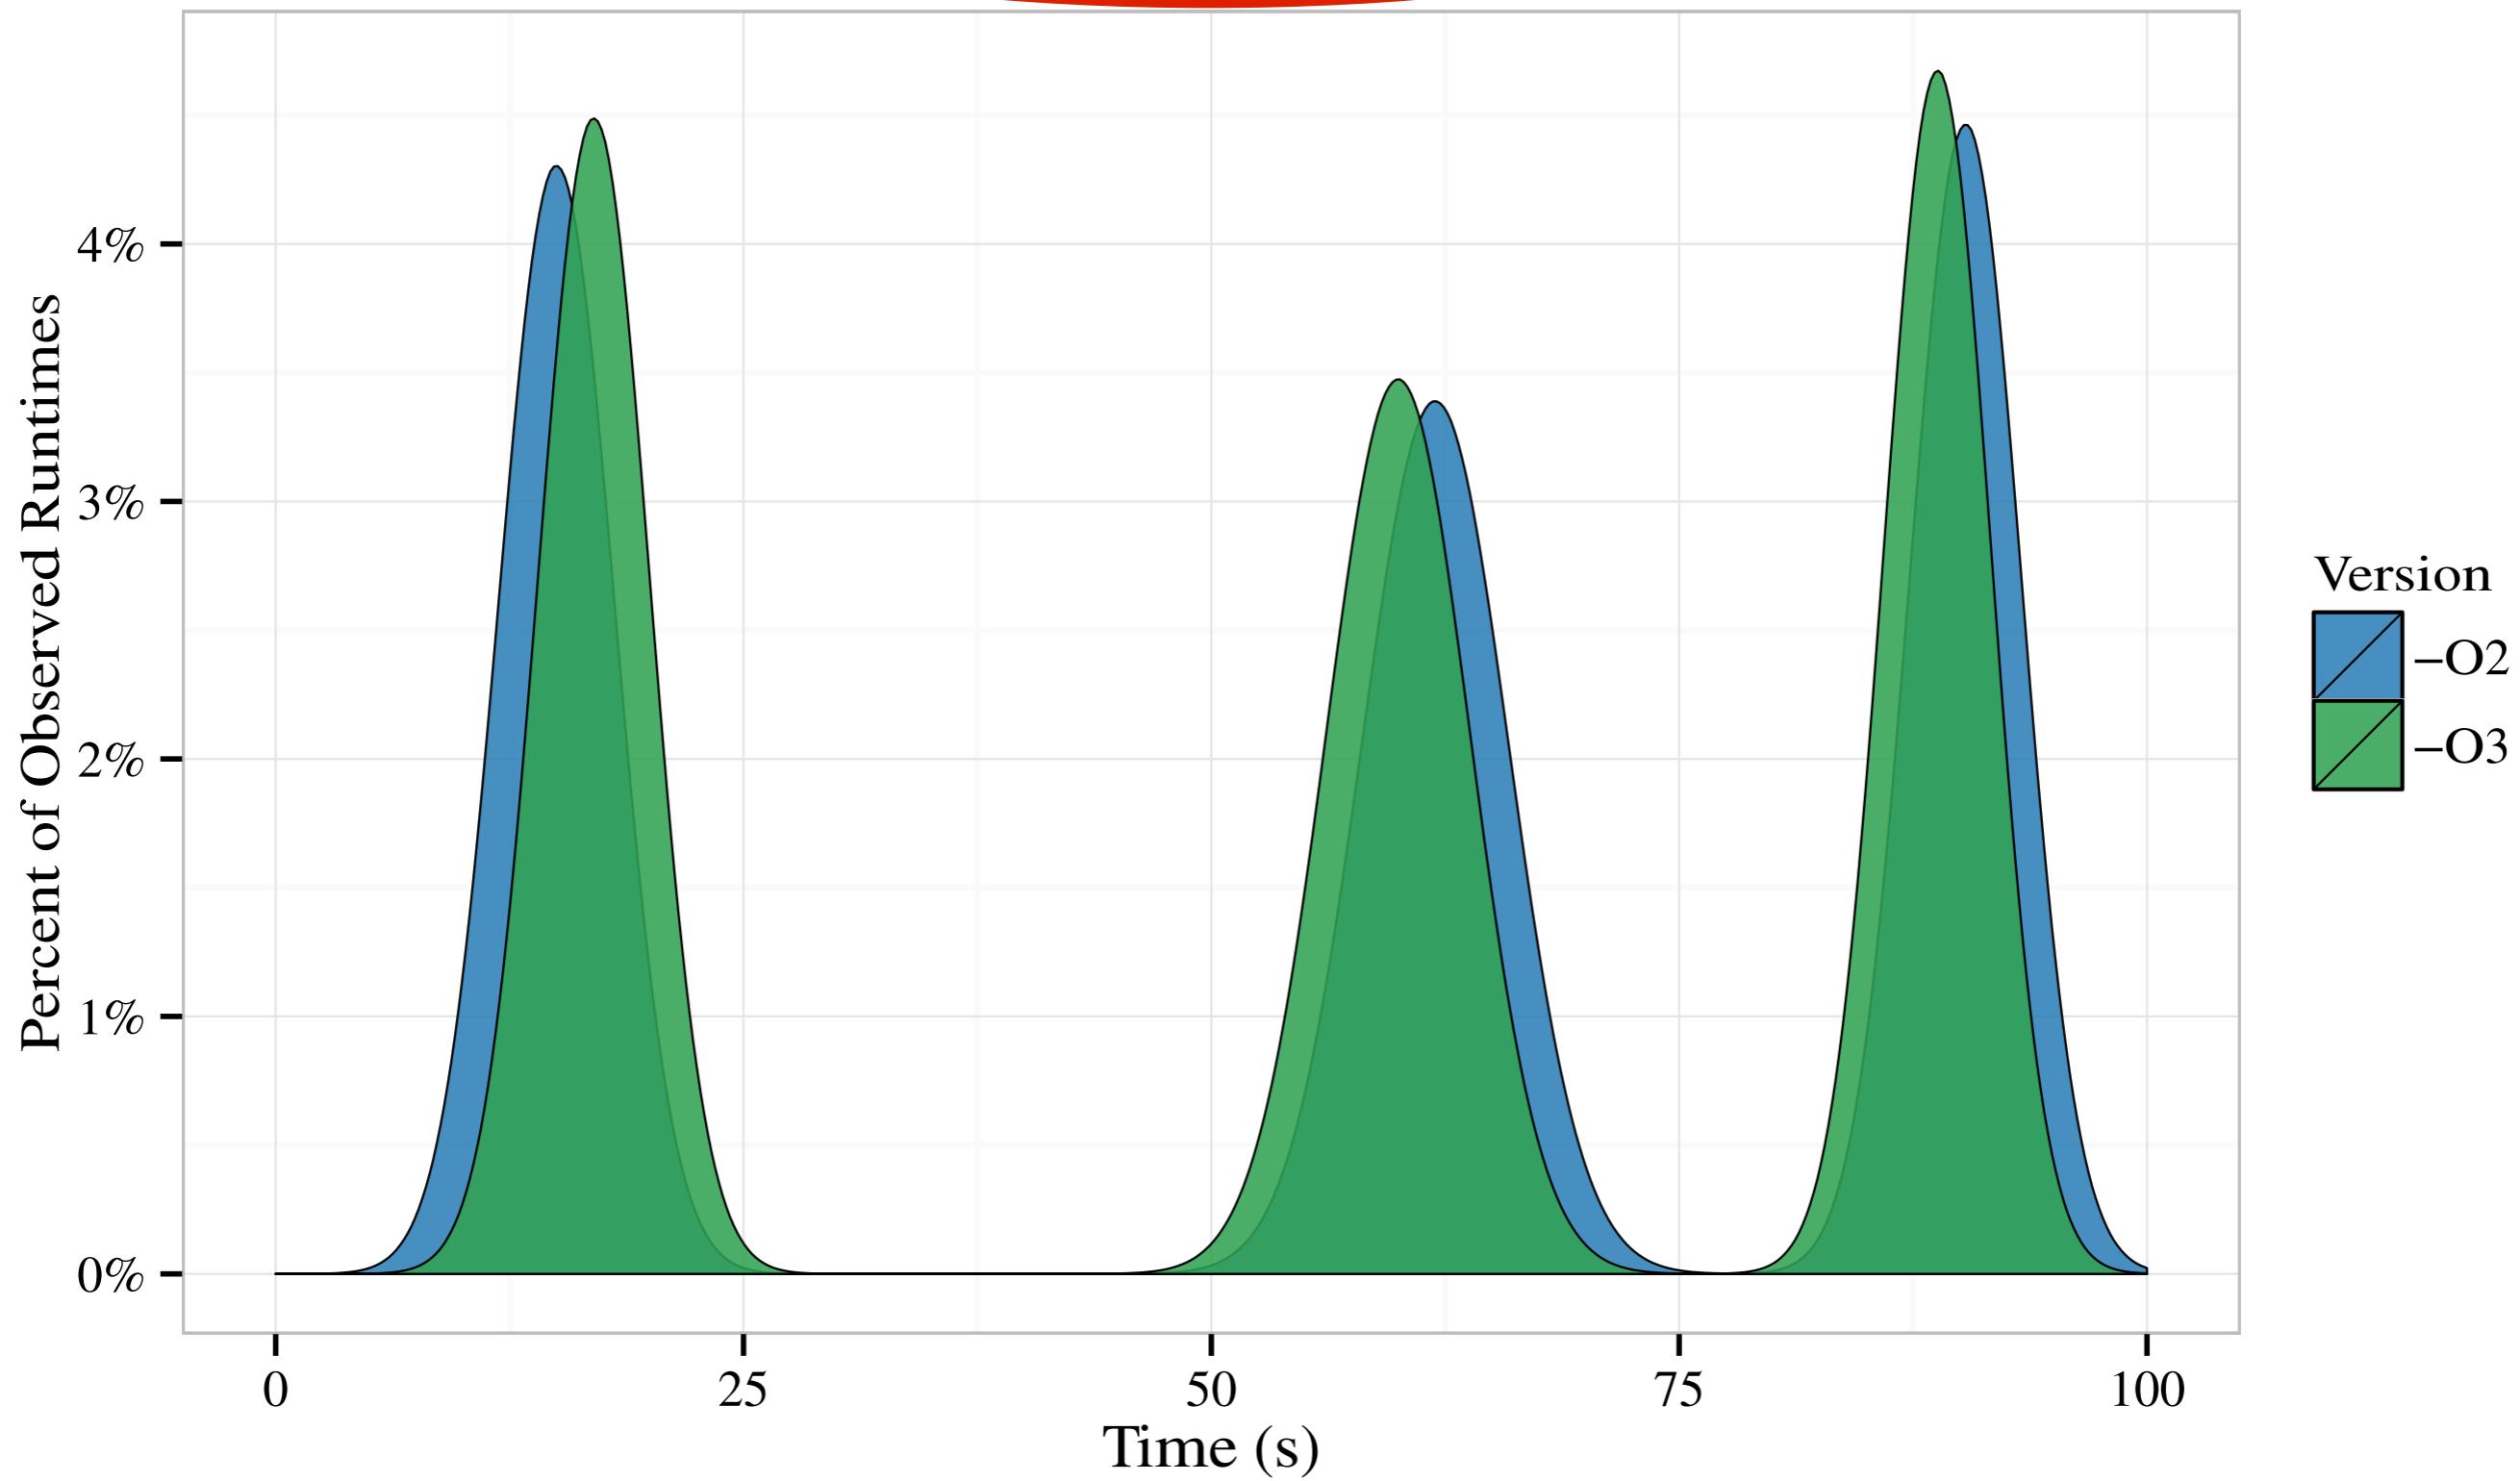
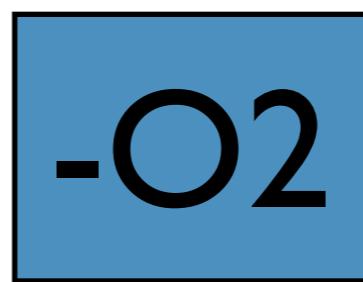
Is **-O3** faster than **-O2**?



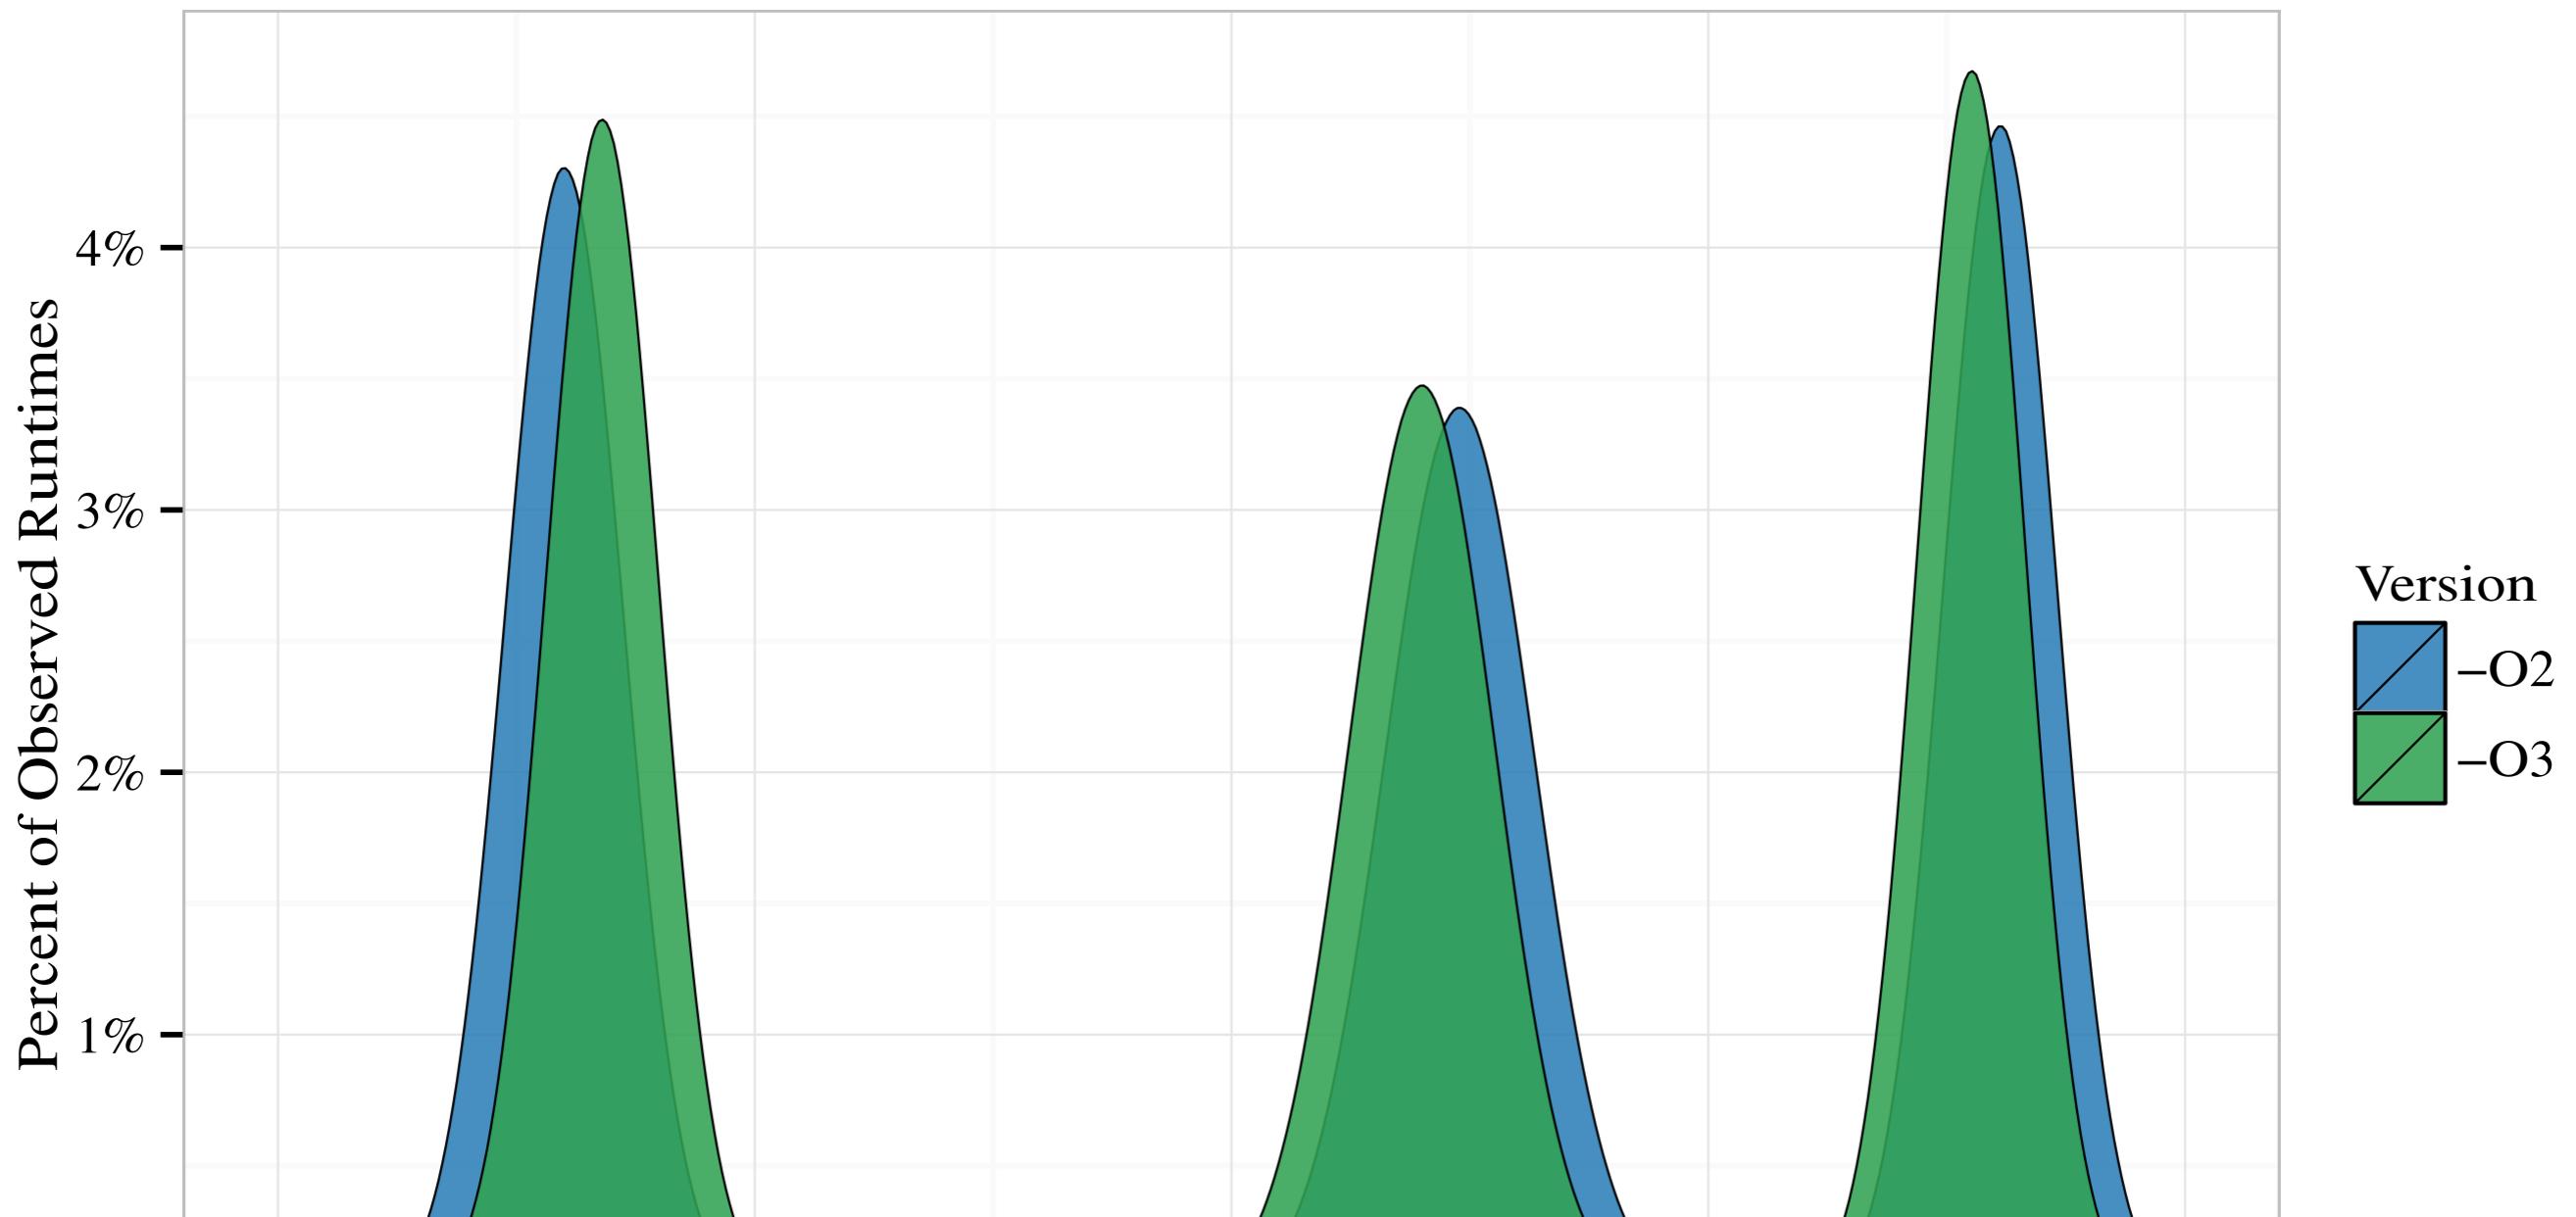
If



=

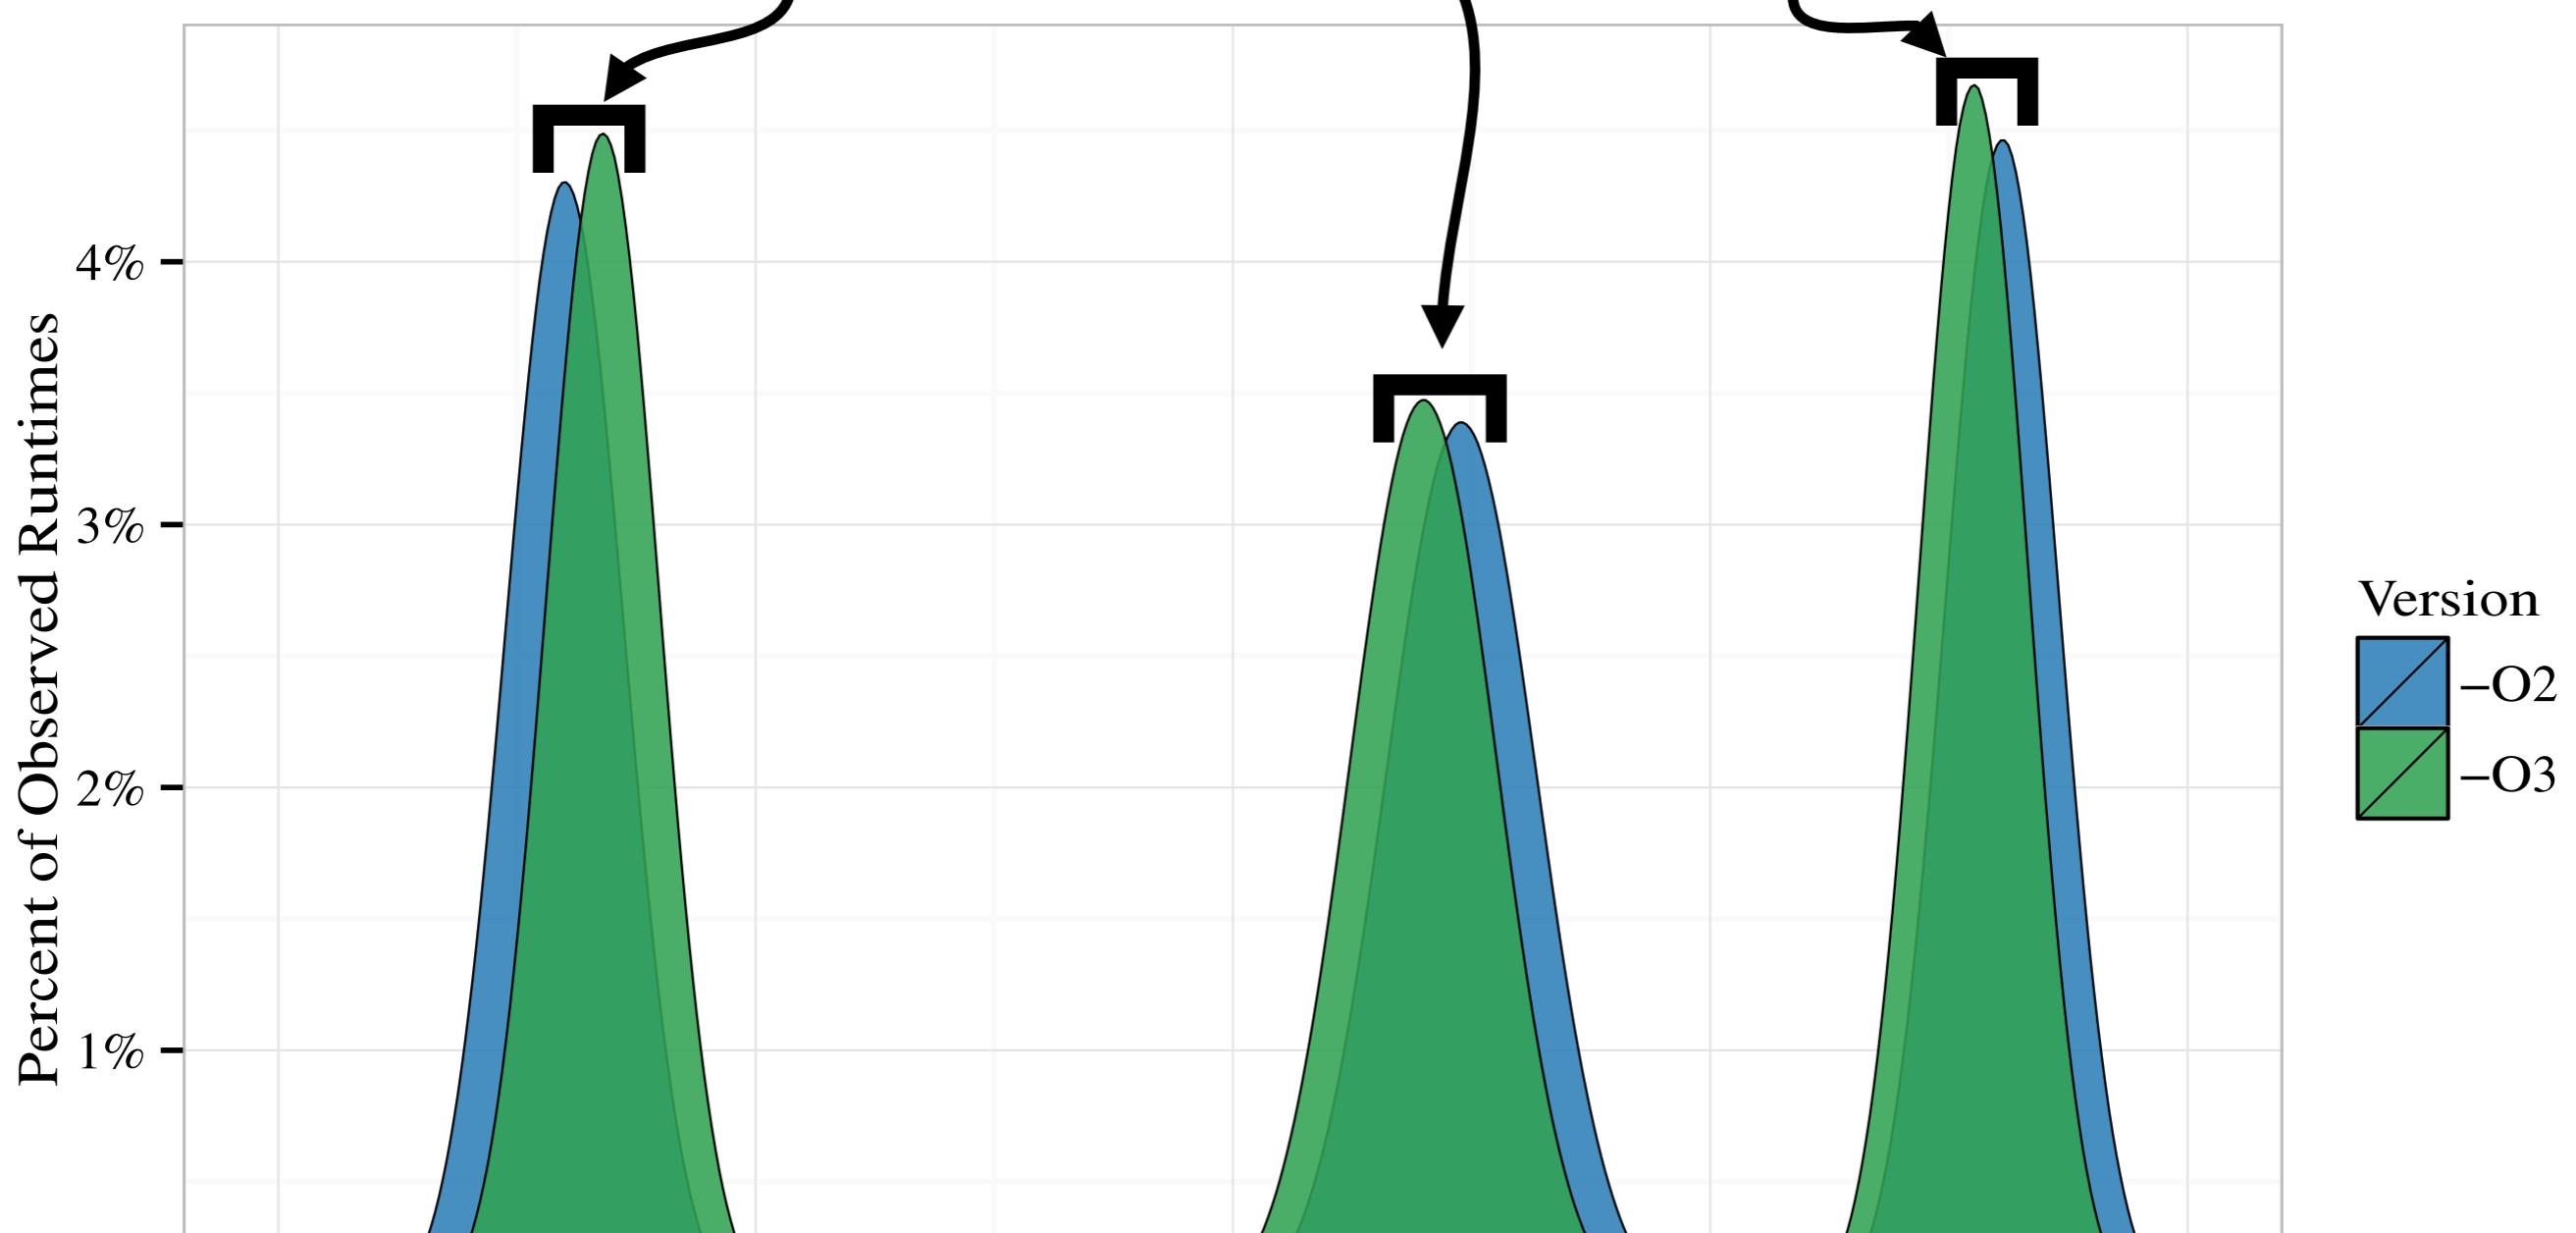


If **-O3** = **-O2**



If **-O3** = **-O2**

what is the probability of measuring
these differences?



Analysis of Variance

```
aov(time~opt+Error(benchmark/opt), times)
```

If $-O_3 = -O_2$

what is the probability of measuring
these differences?



Analysis of Variance

If p-value

If $-O_3 = -O_2$

what is the probability of measuring
these differences?

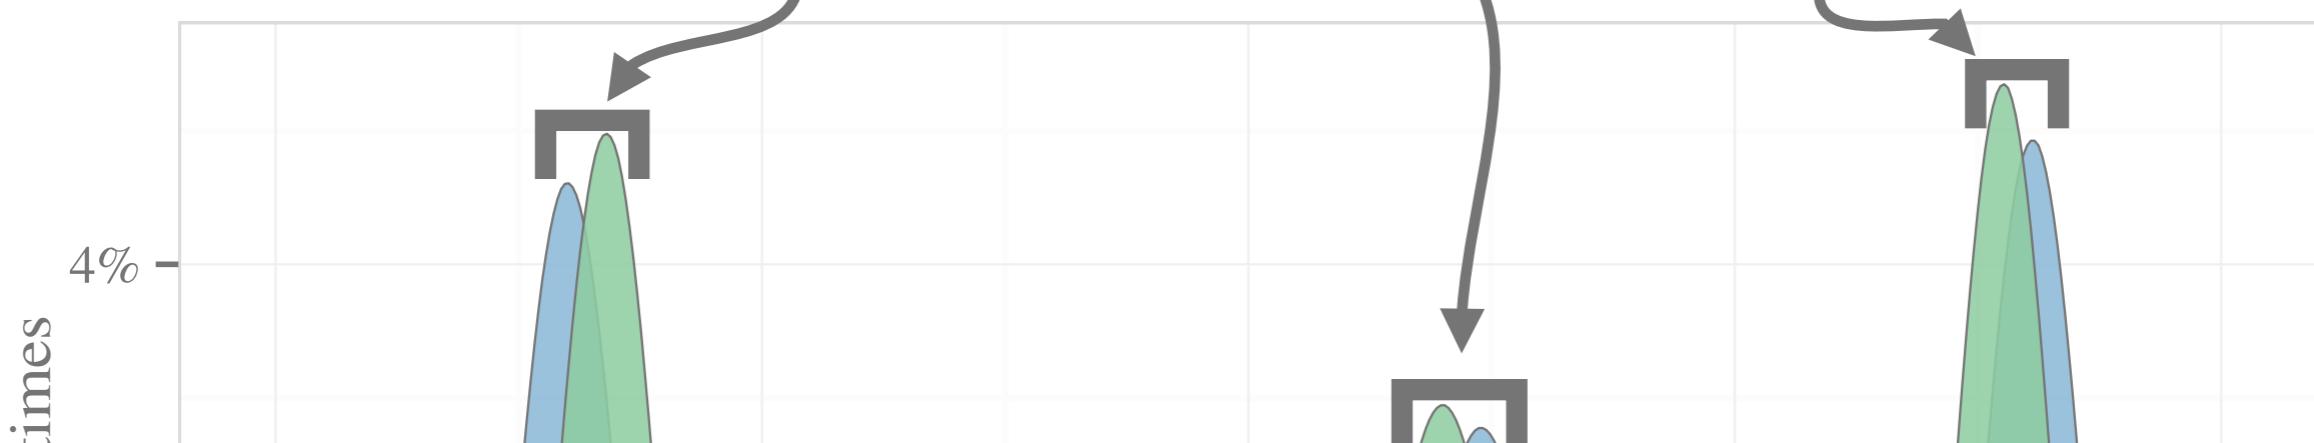


Analysis of Variance

If p-value \leq 5%

If $-O_3 = -O_2$

what is the probability of measuring
these differences?

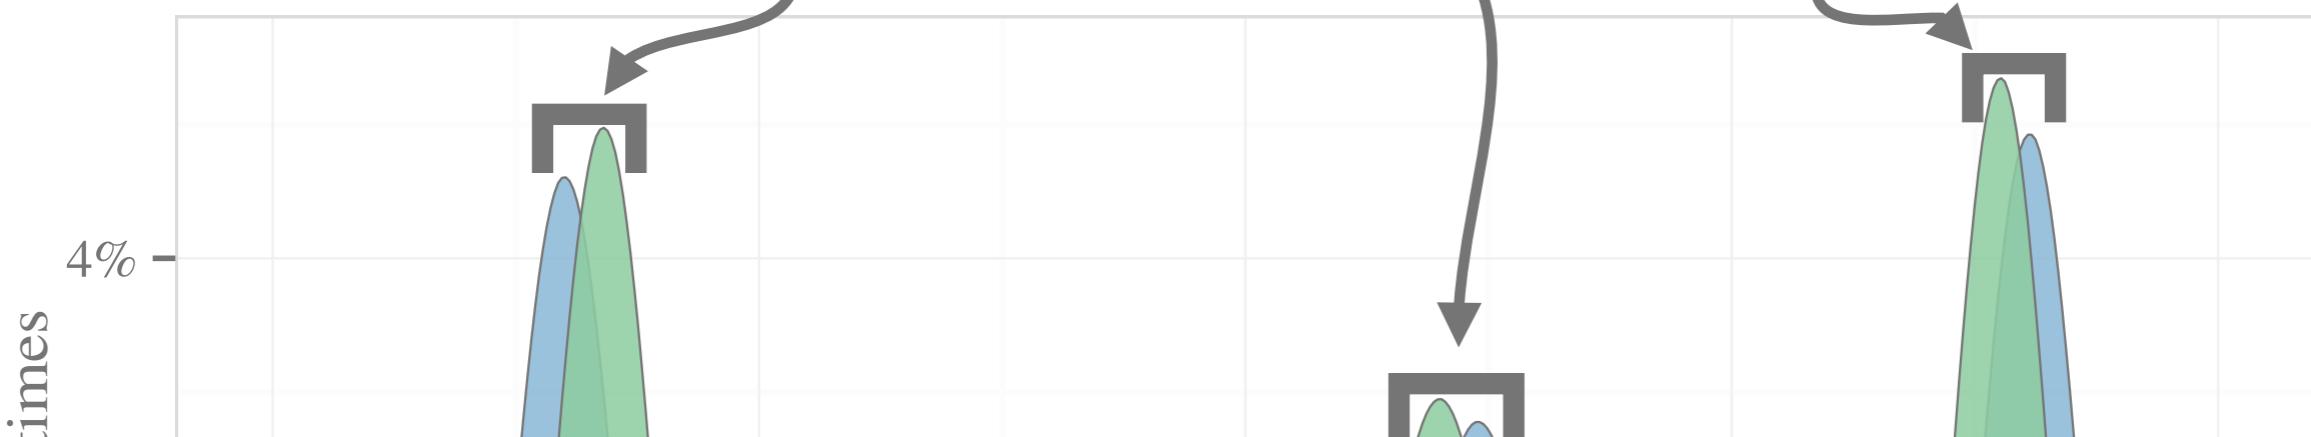


Analysis of Variance

If p-value $\leq 5\%$
we reject the null hypothesis

If $[-O_3] = [-O_2]$

what is the probability of measuring
these differences?



Analysis of Variance

If p-value $\leq 5\%$
we *reject the null hypothesis*

-O3 vs -O2

p-value = 26.4%

one in four experiments will show an effect that does not exist!

Analysis of Variance

If p-value $\leq 5\%$
we *reject the null hypothesis*

-O3 vs -O2

p-value = 26.4%

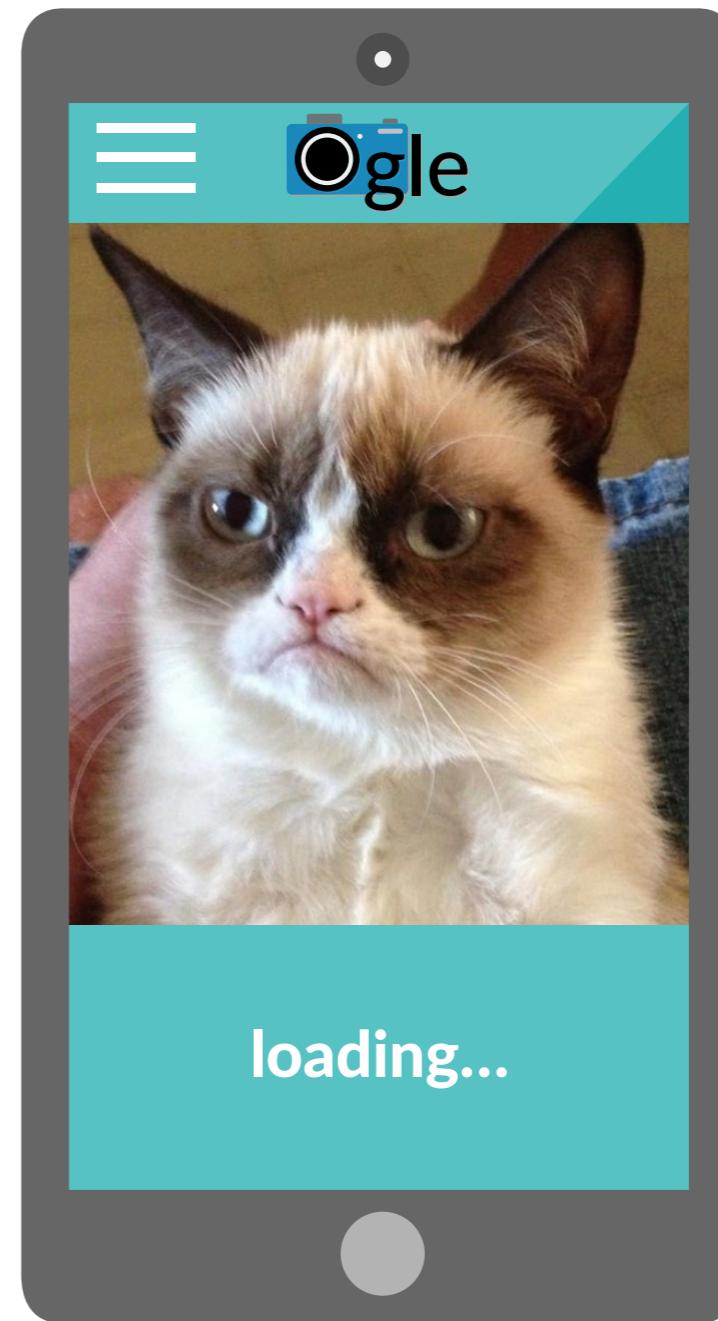
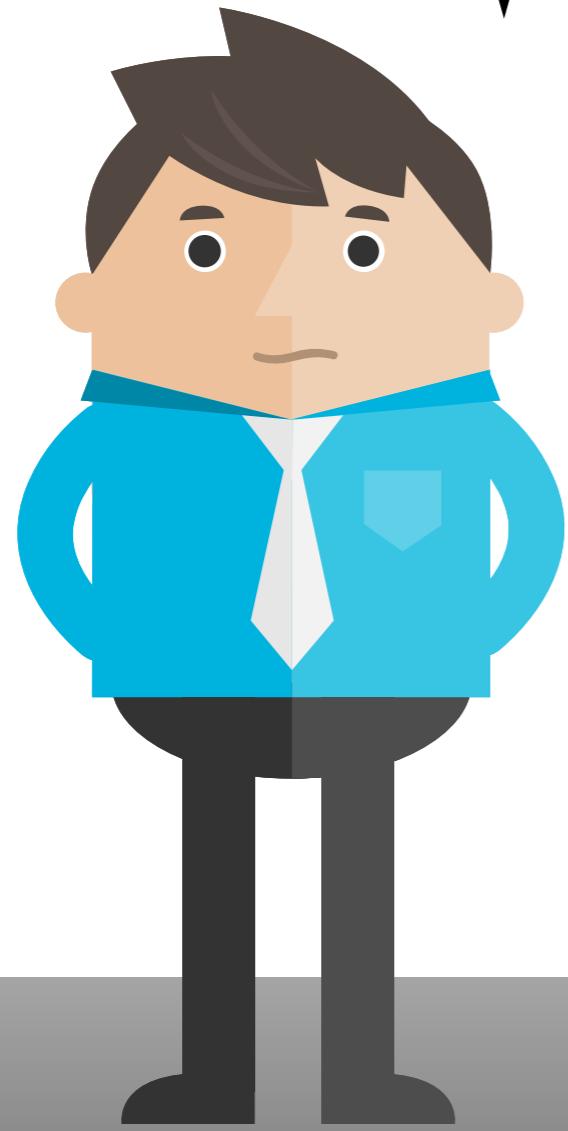
fail to reject the null hypothesis

Analysis of Variance

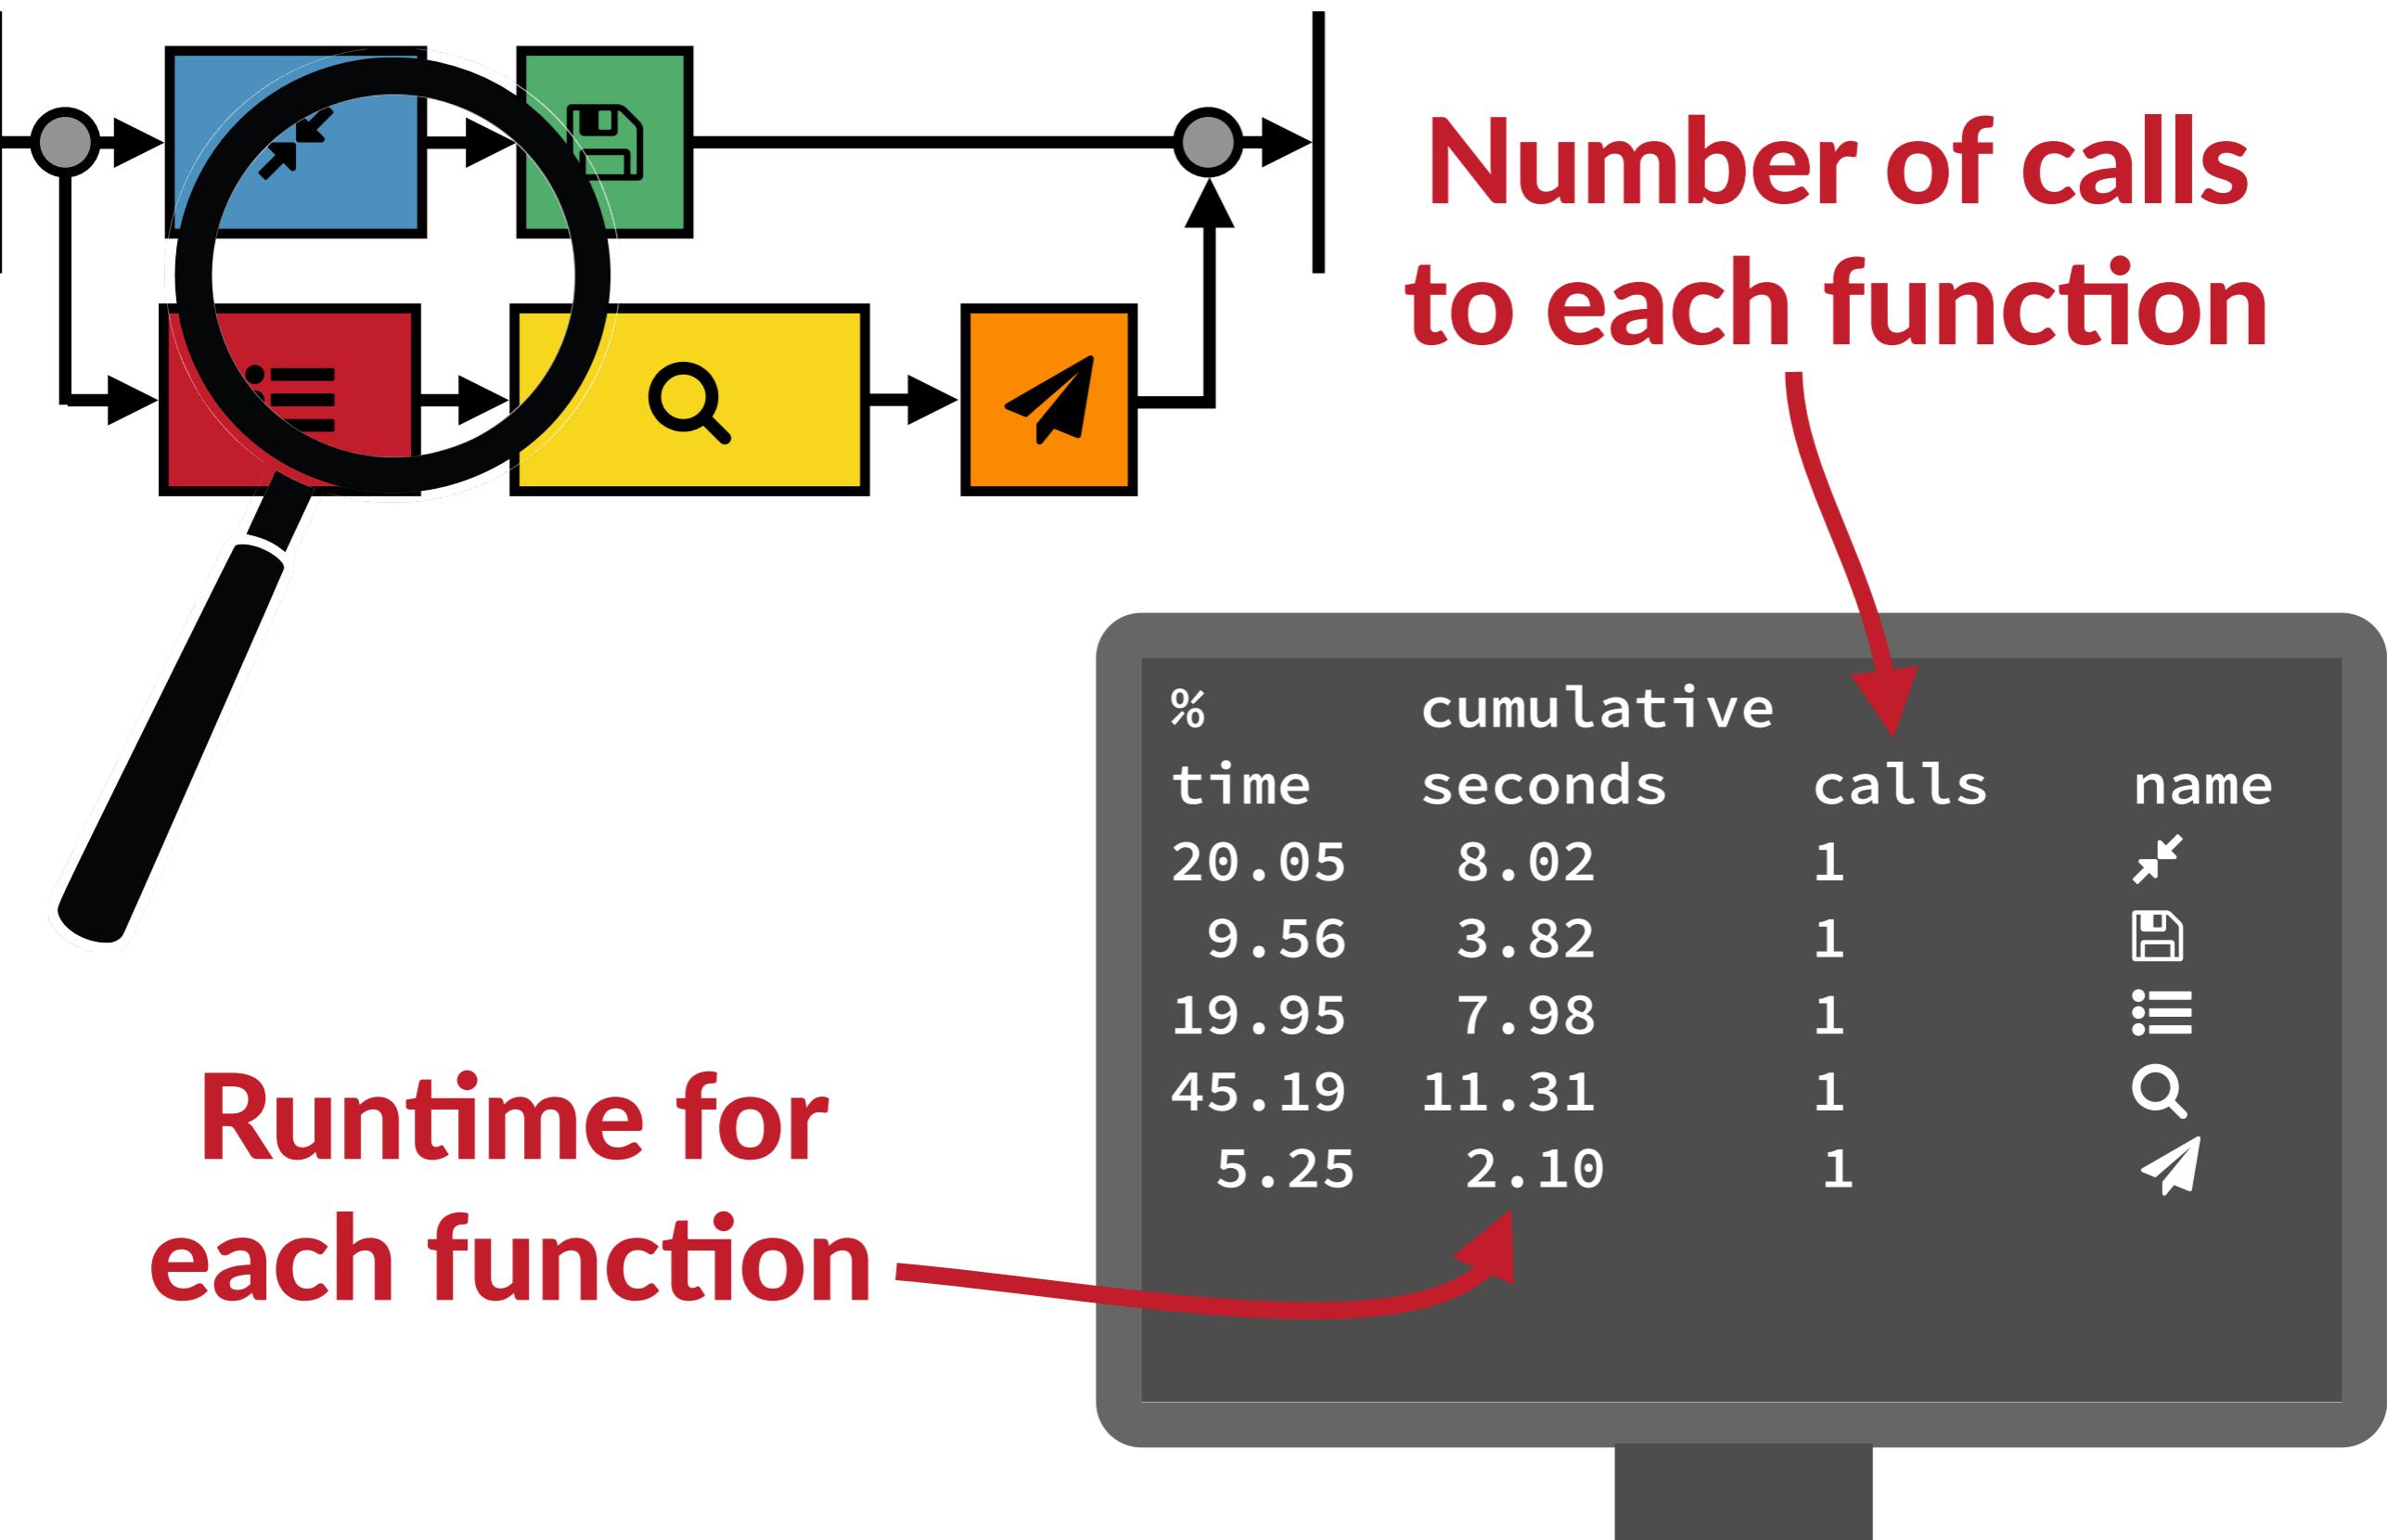
If p-value $\leq 5\%$
we *reject the null hypothesis*

The effect of **-O3** over **-O2** is
indistinguishable from noise

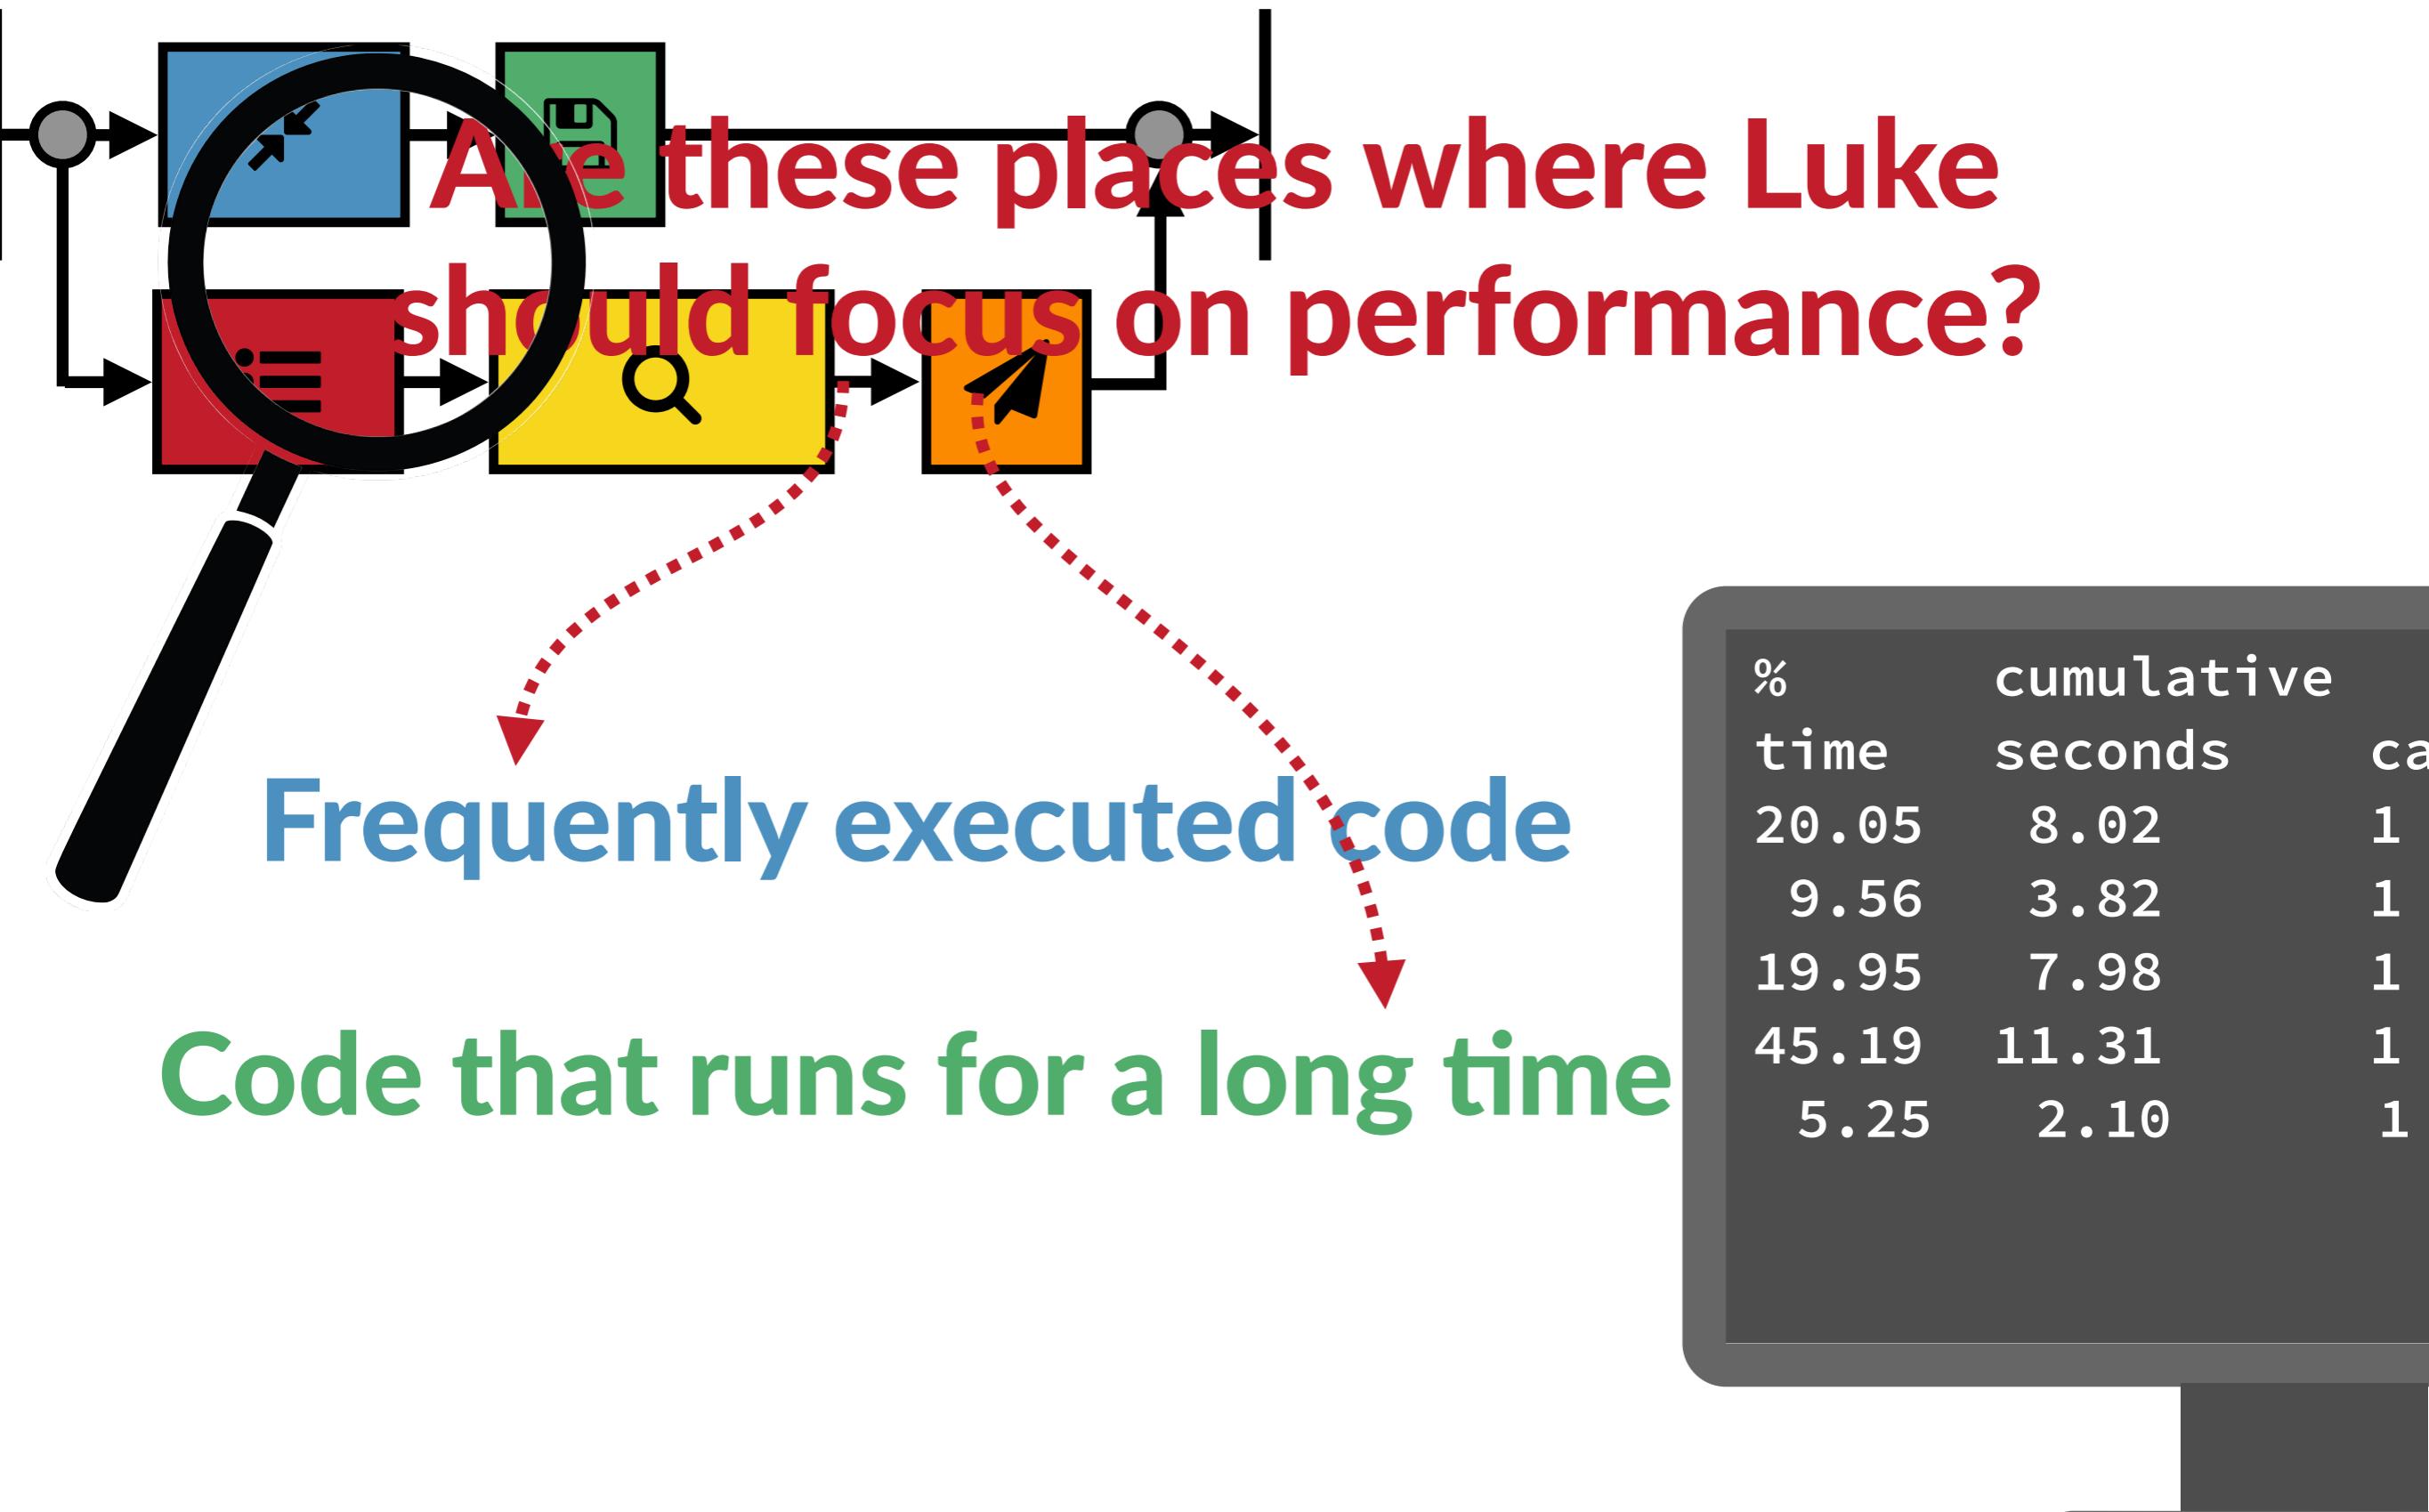
...did you try -O9?



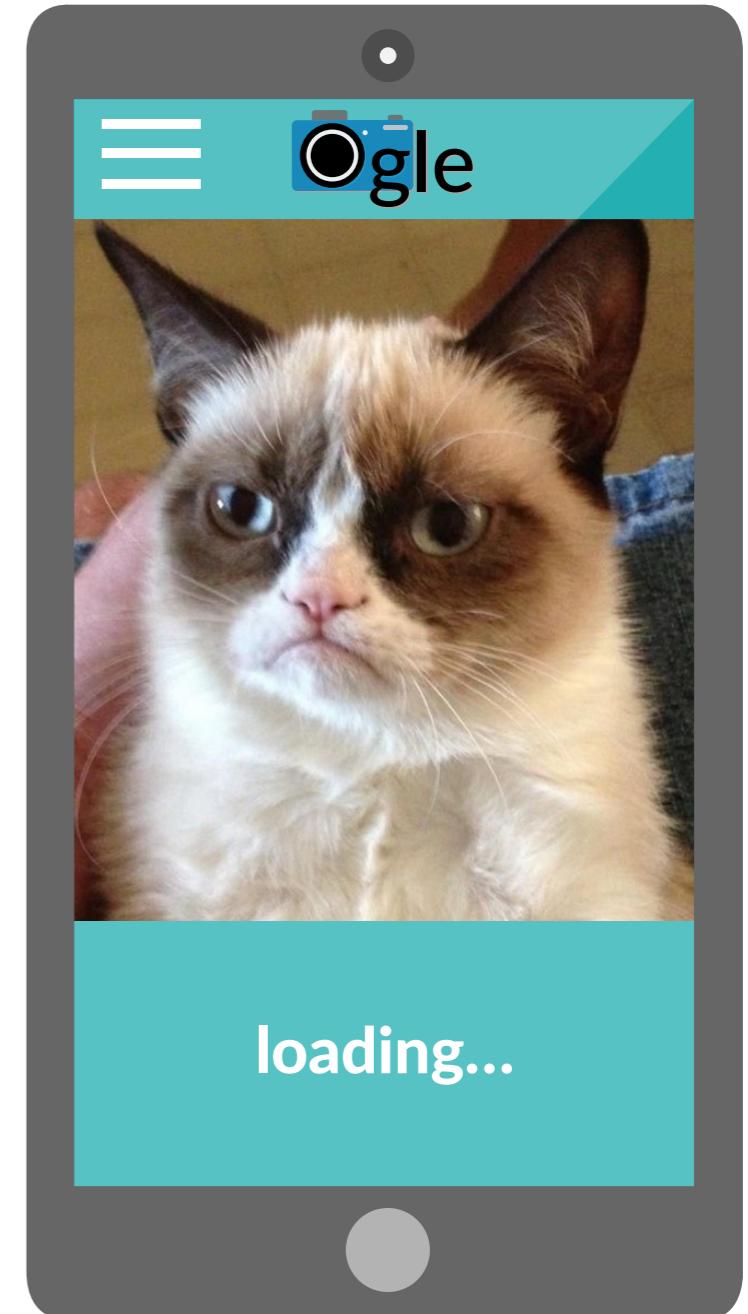
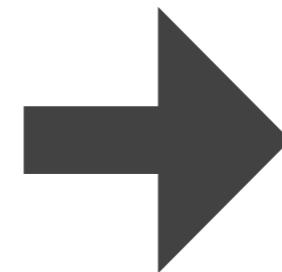
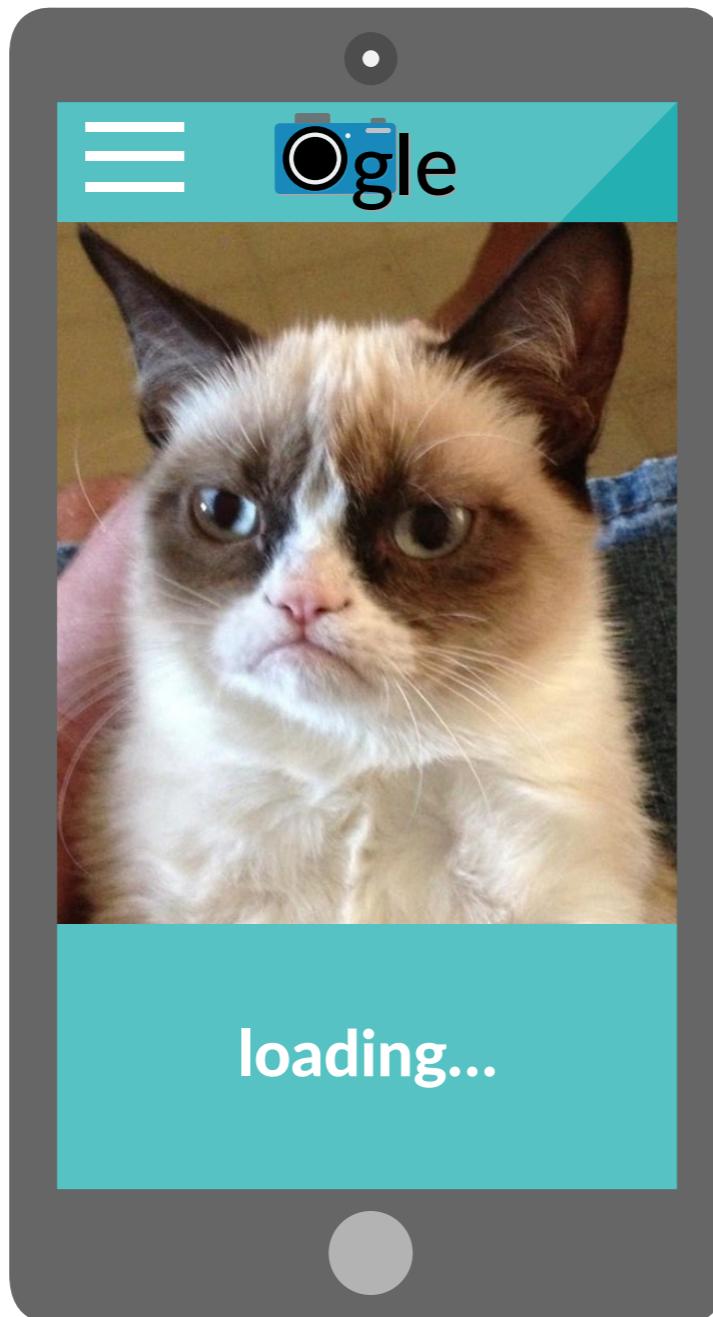
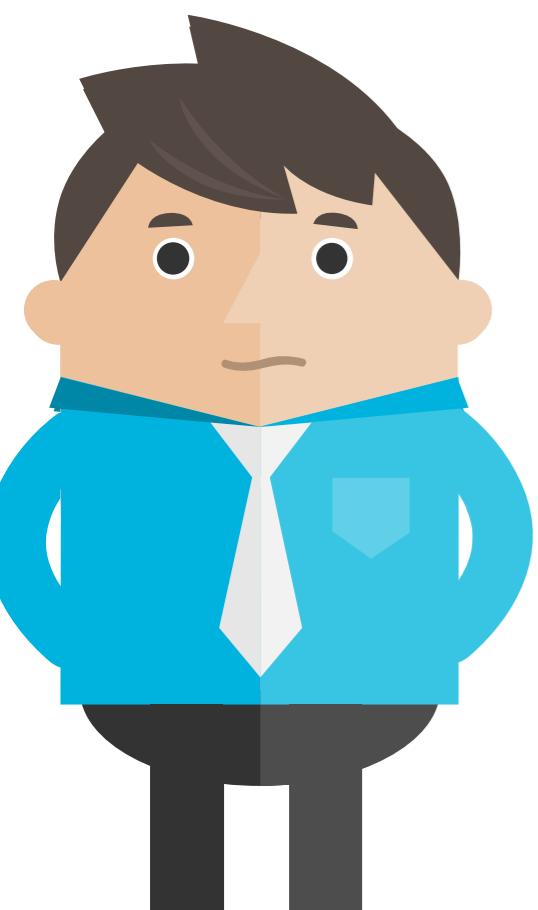
Software Profilers



Software Profilers



Would this speed up Ogle?

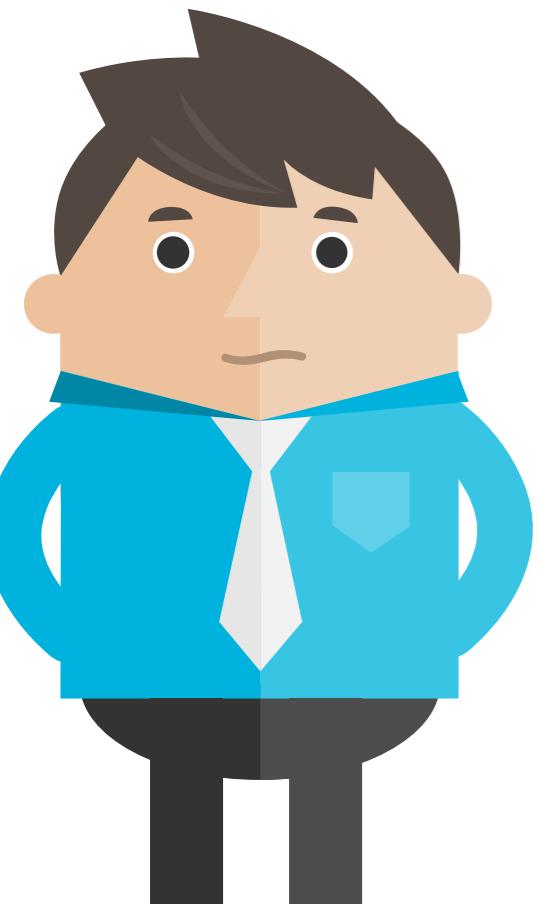


Would this speed up Ogle?

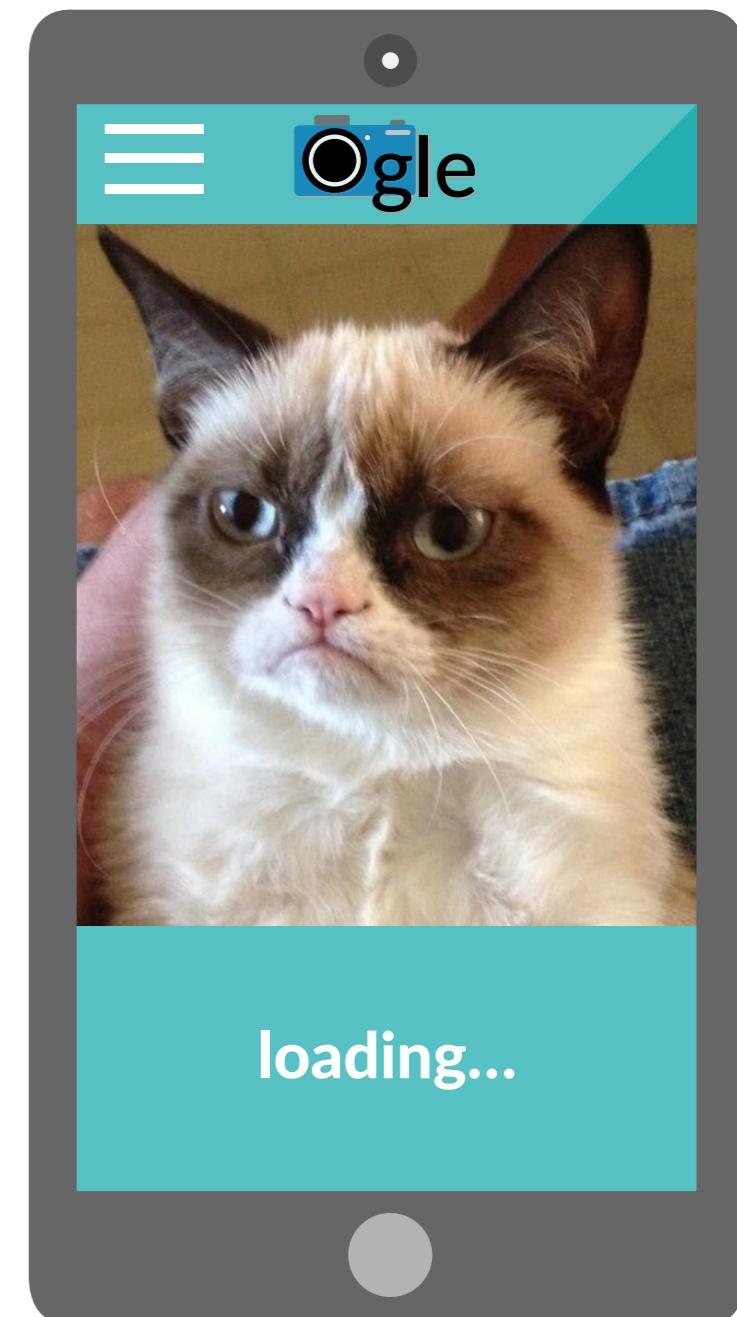
Frequently executed code

Code that runs for a long time

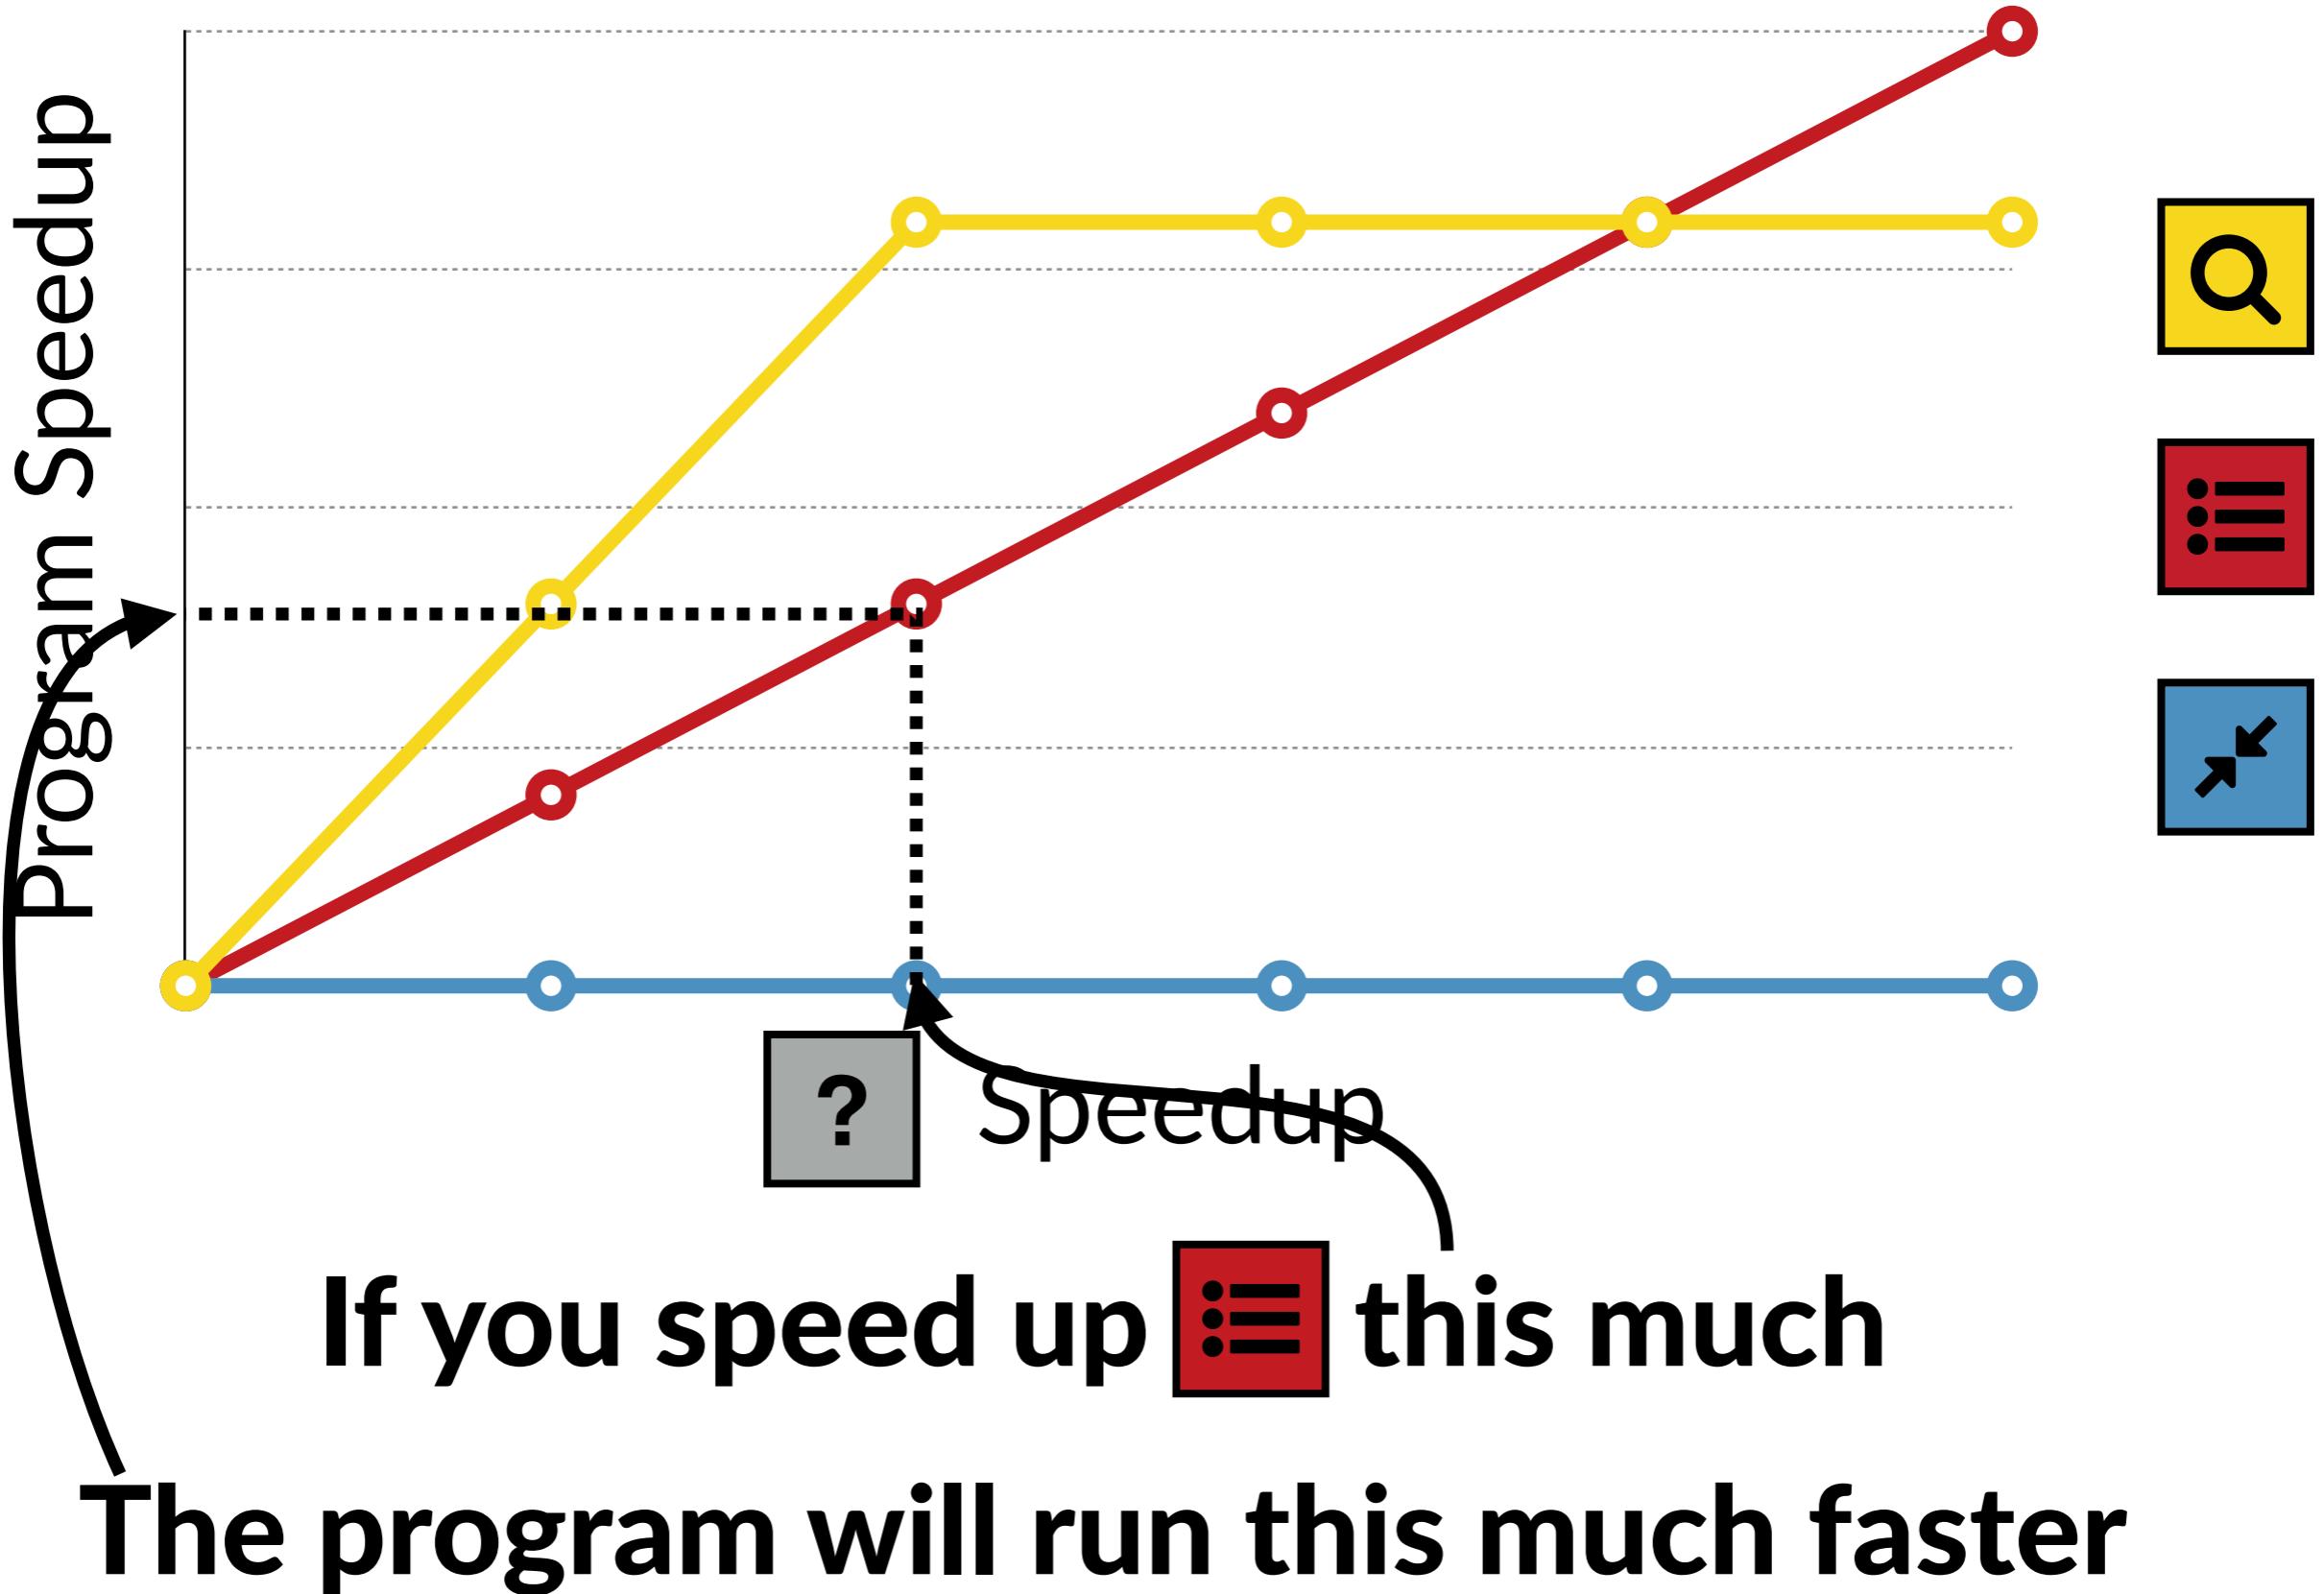
Profilers do a bad job finding important code in parallel / asynchronous / concurrent programs.



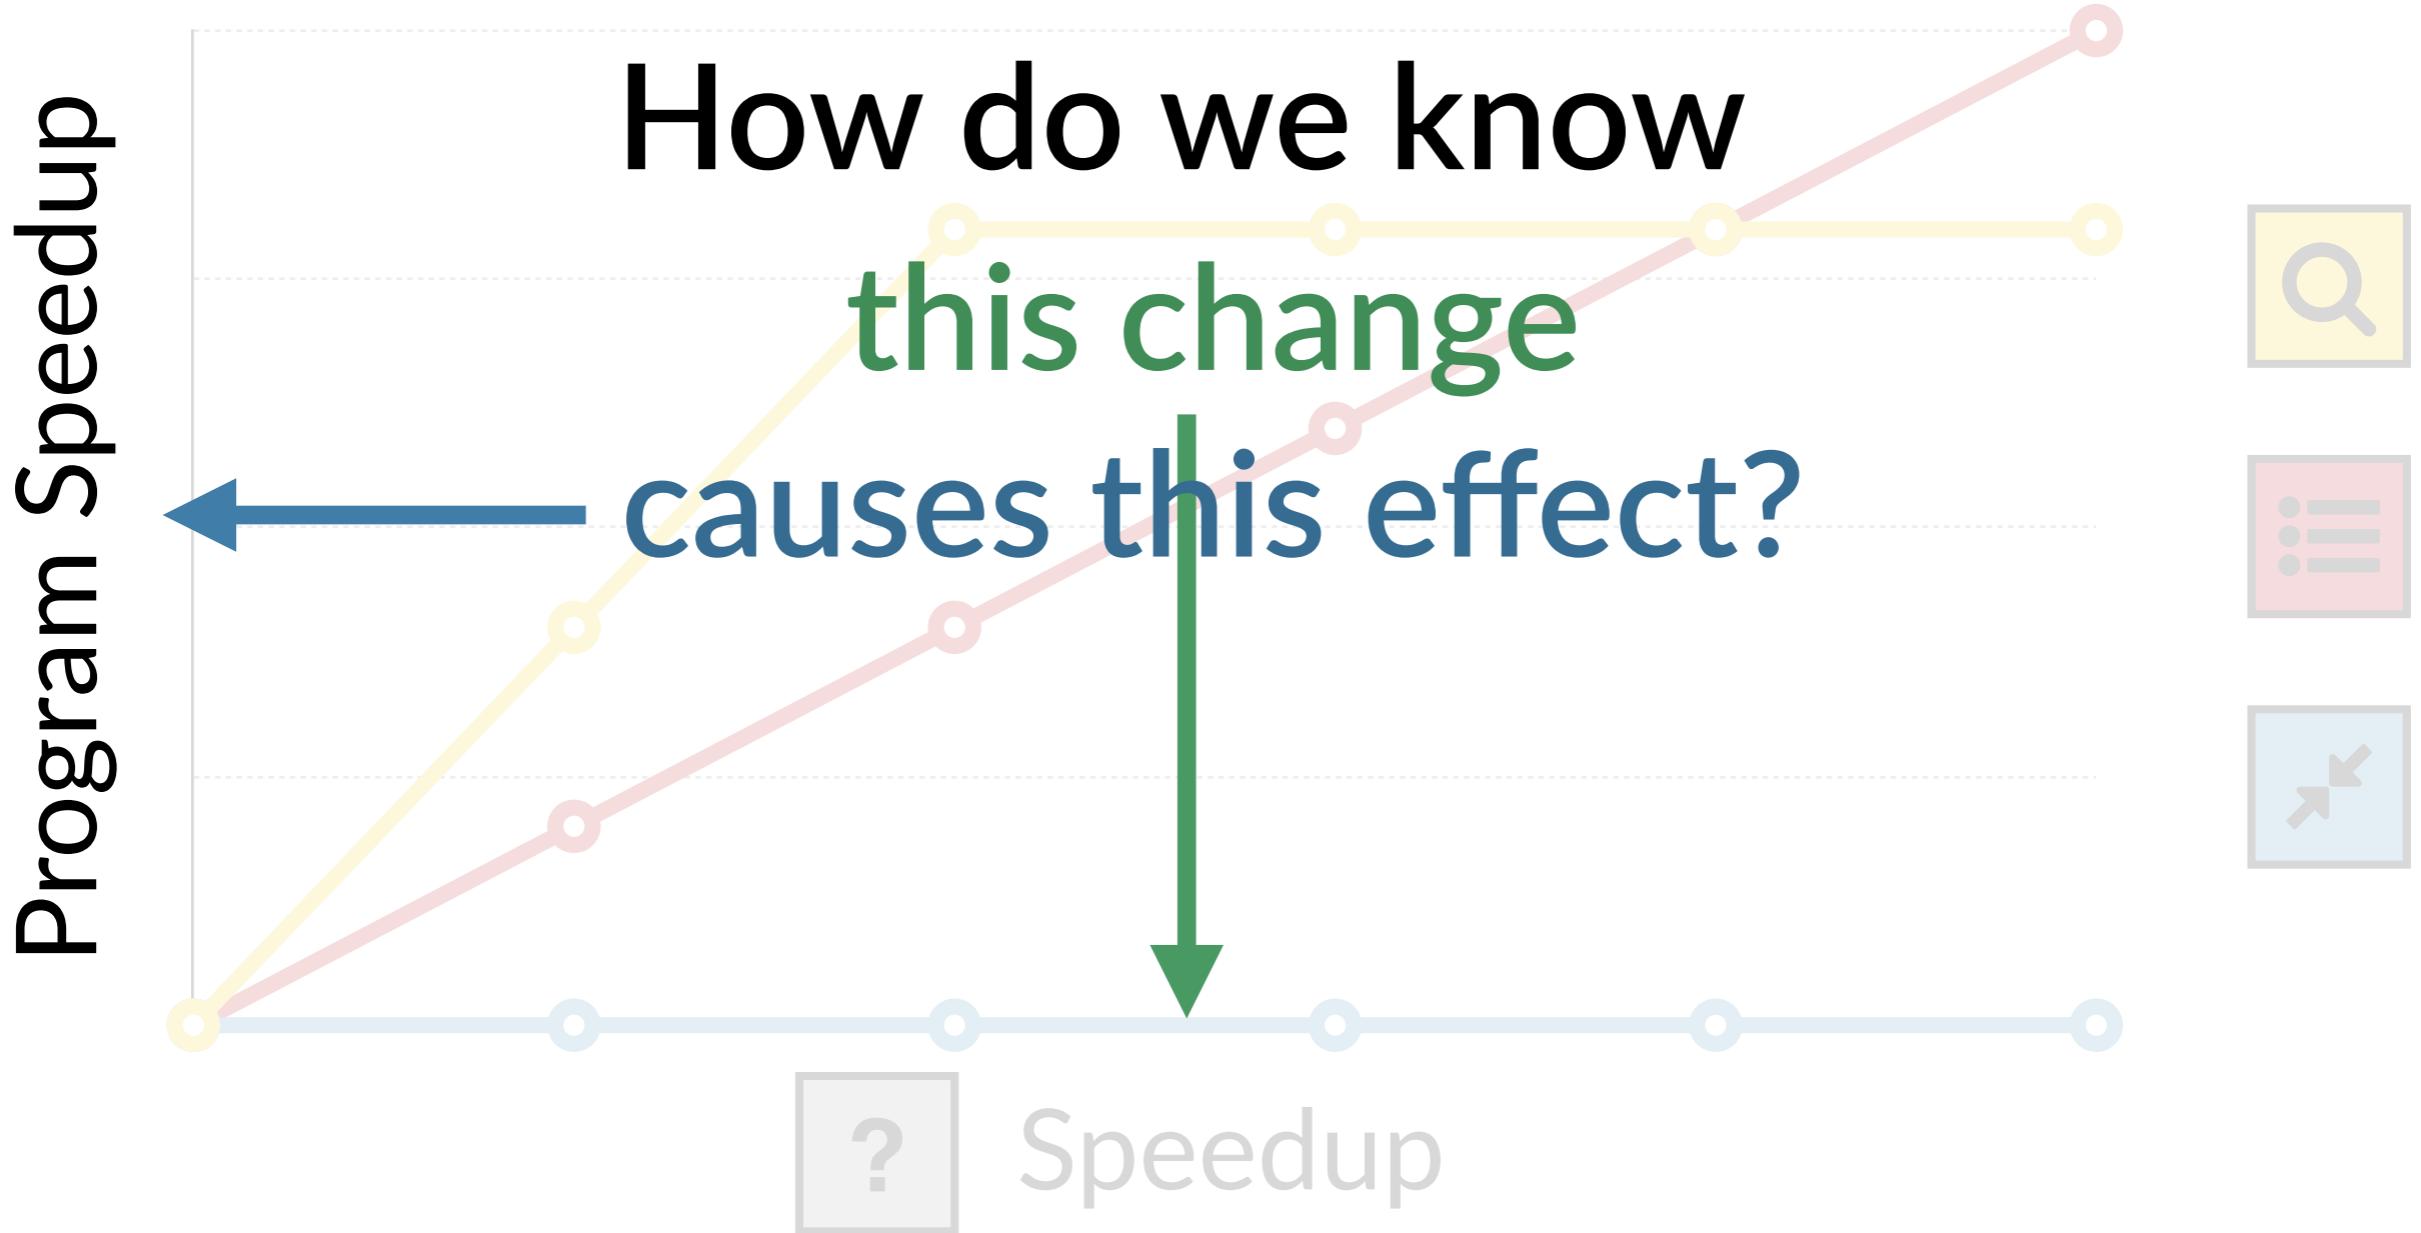
We need to do better.



Causal Profile



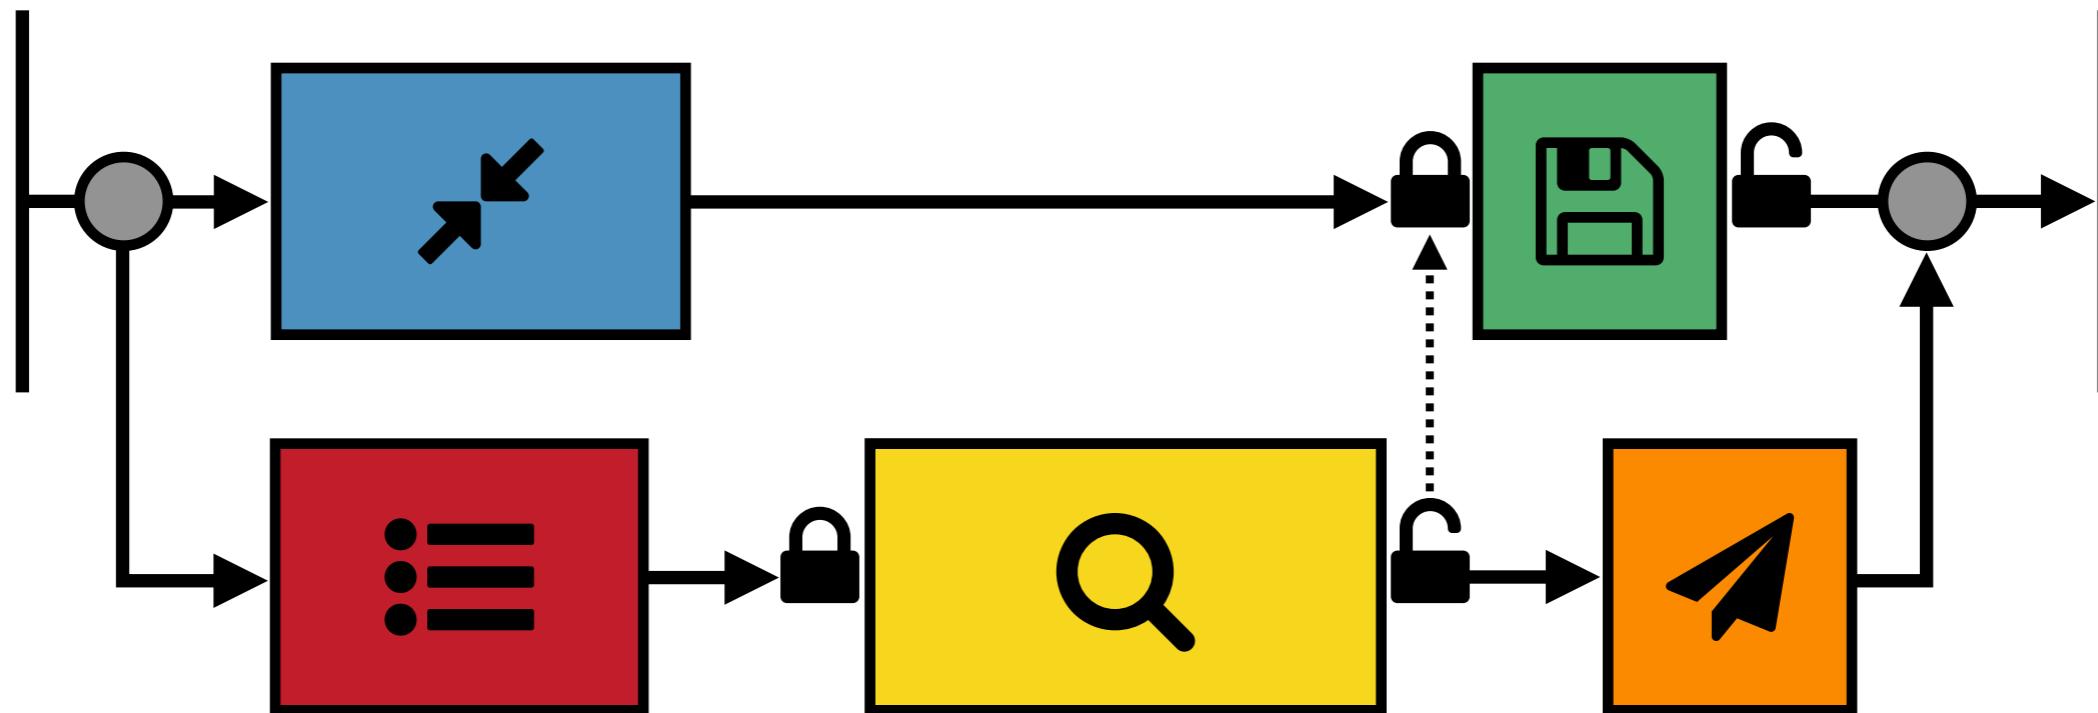
Causal Profile



Run an experiment

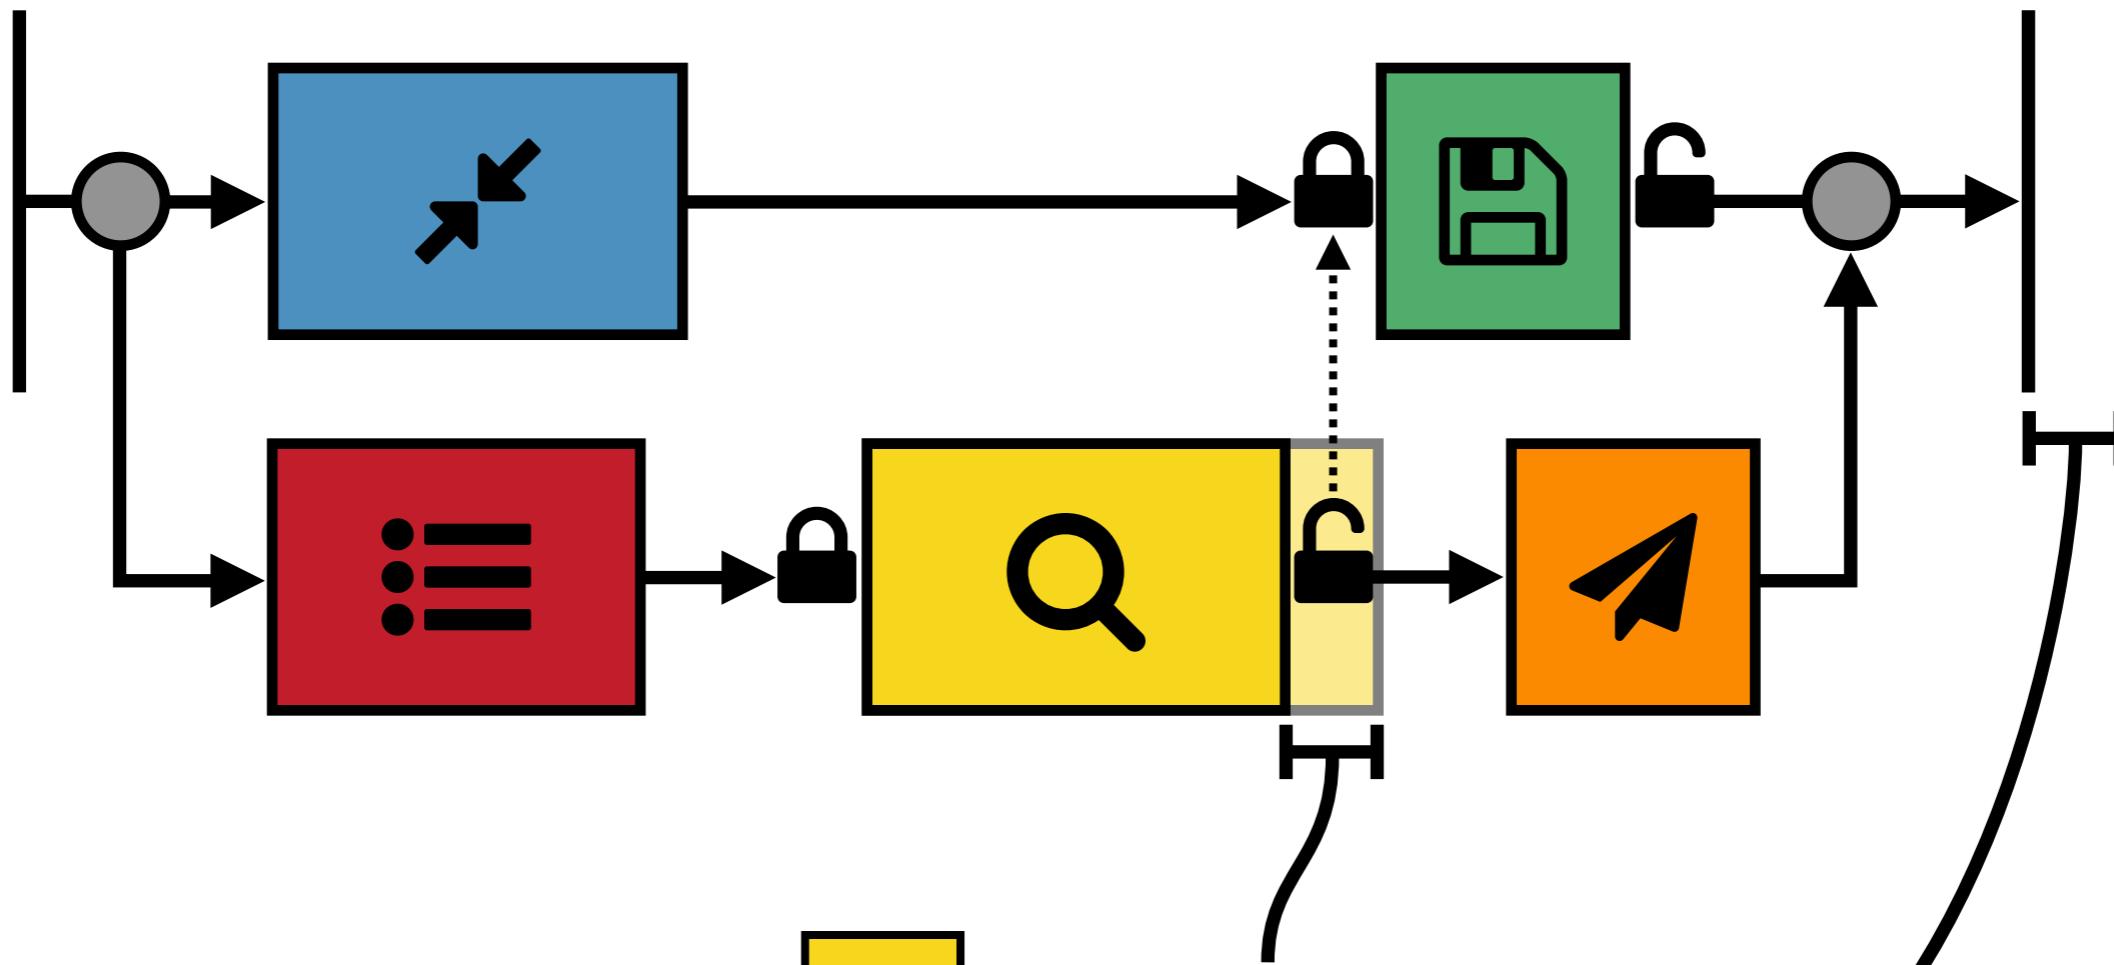
Performance Experiments

If we could magically speed up  ...



Performance Experiments

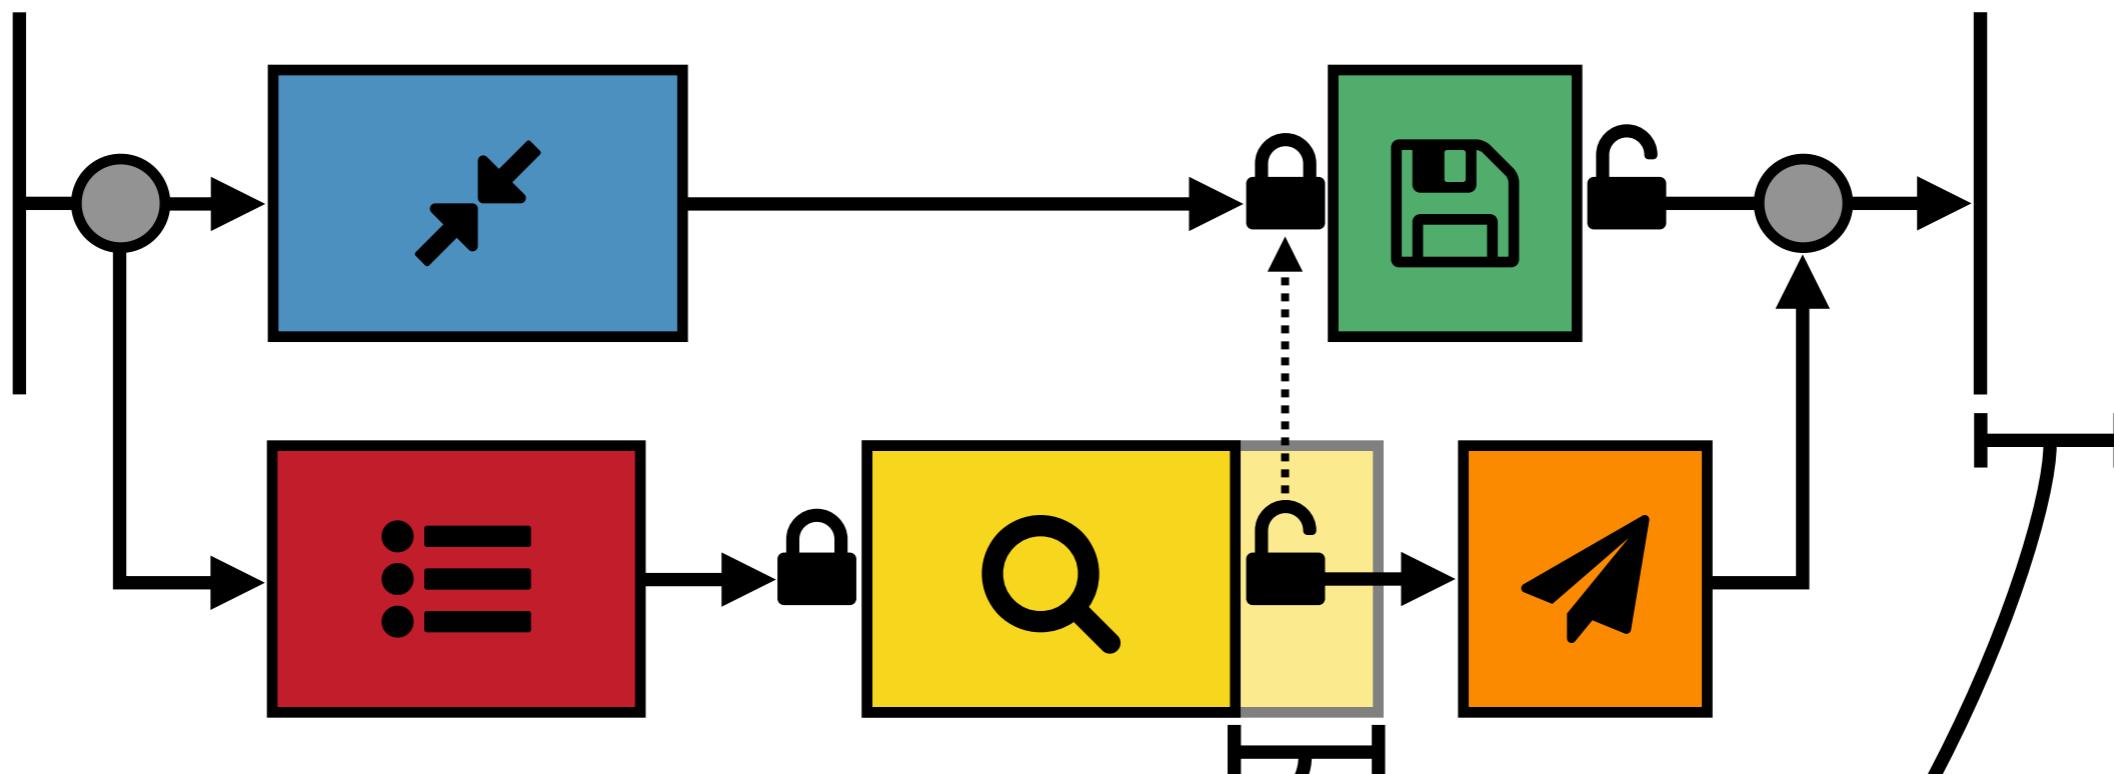
If we could magically speed up  ...



Speeding up  by this much...
speeds up the program by this much.

Performance Experiments

If we could magically speed up  ...

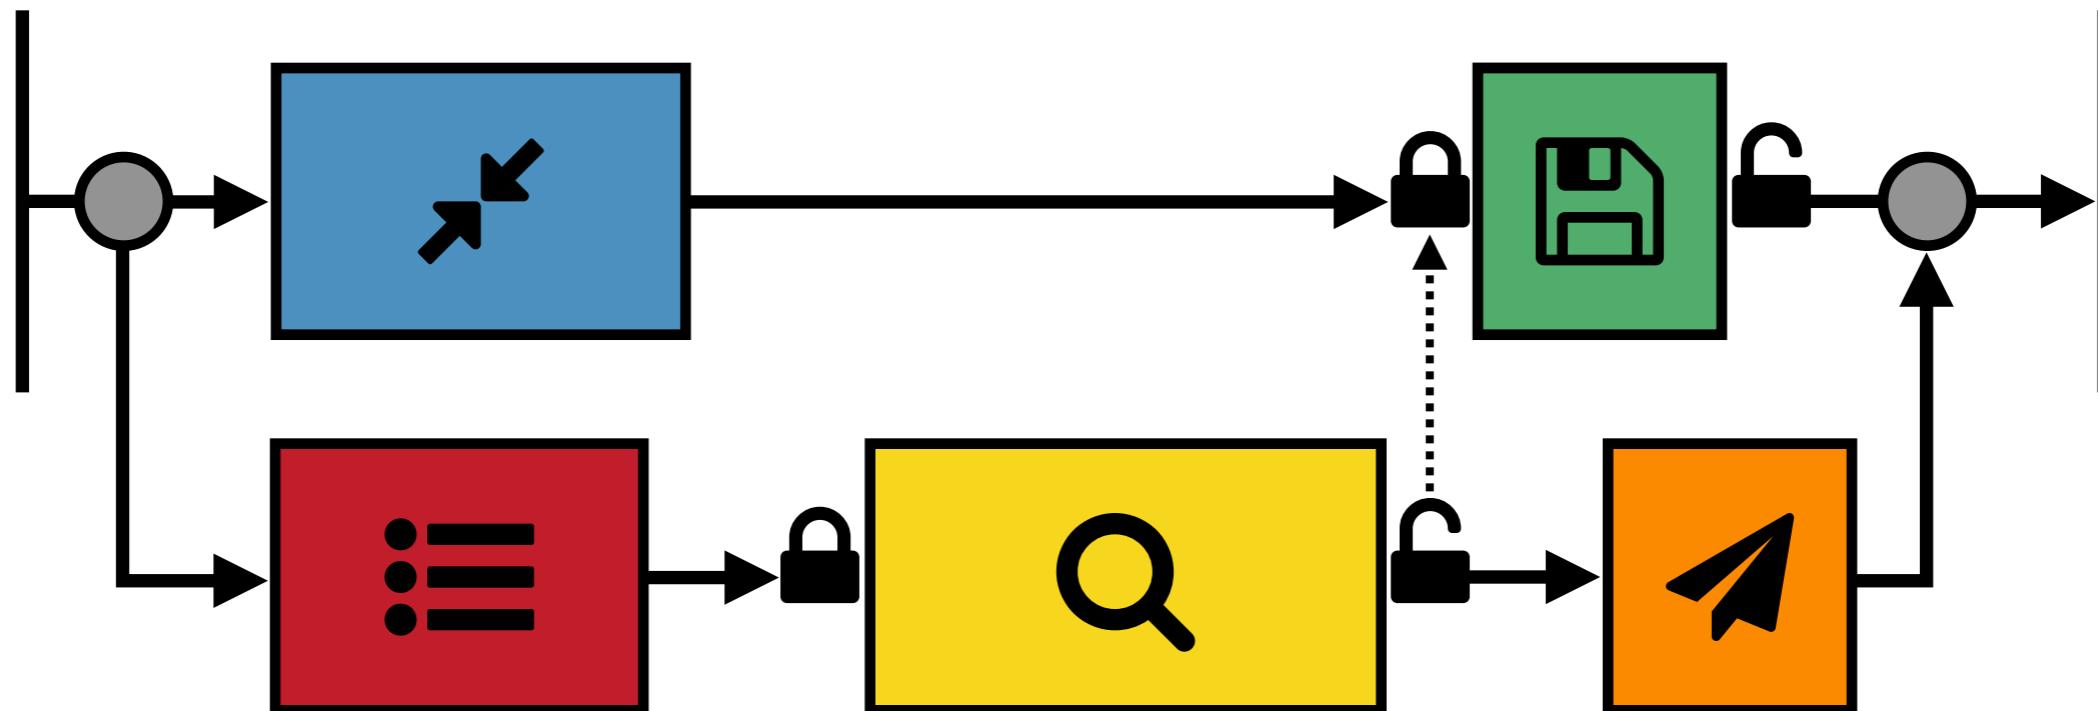


More speedup in  ...

leads to a larger program speedup.

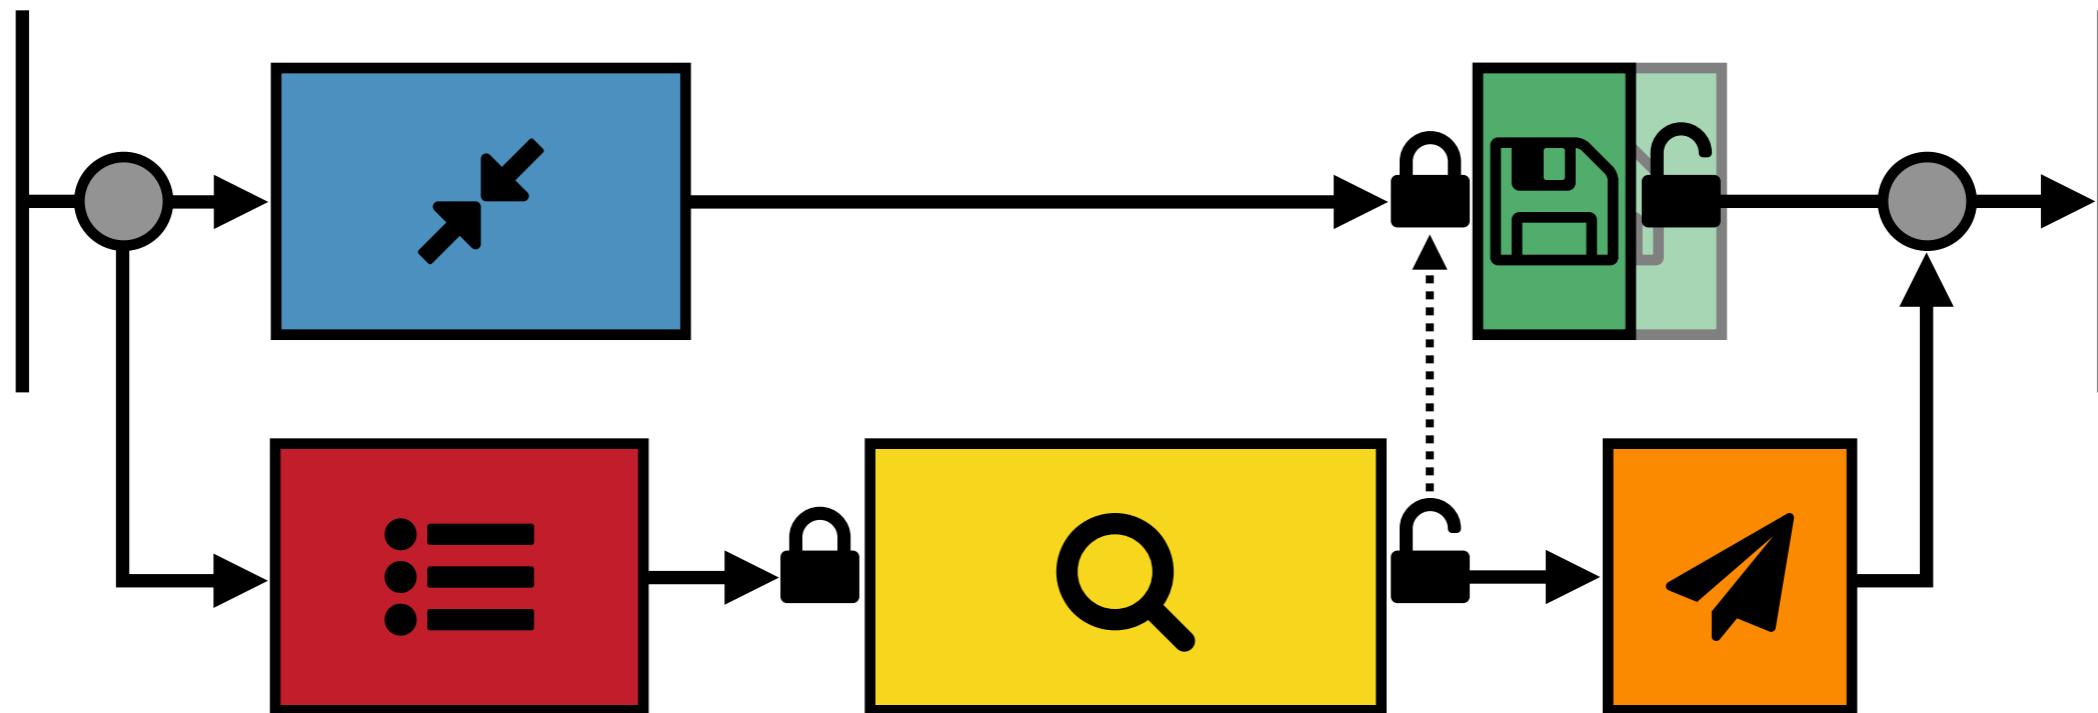
Performance Experiments

If we could magically speed up  ...



Performance Experiments

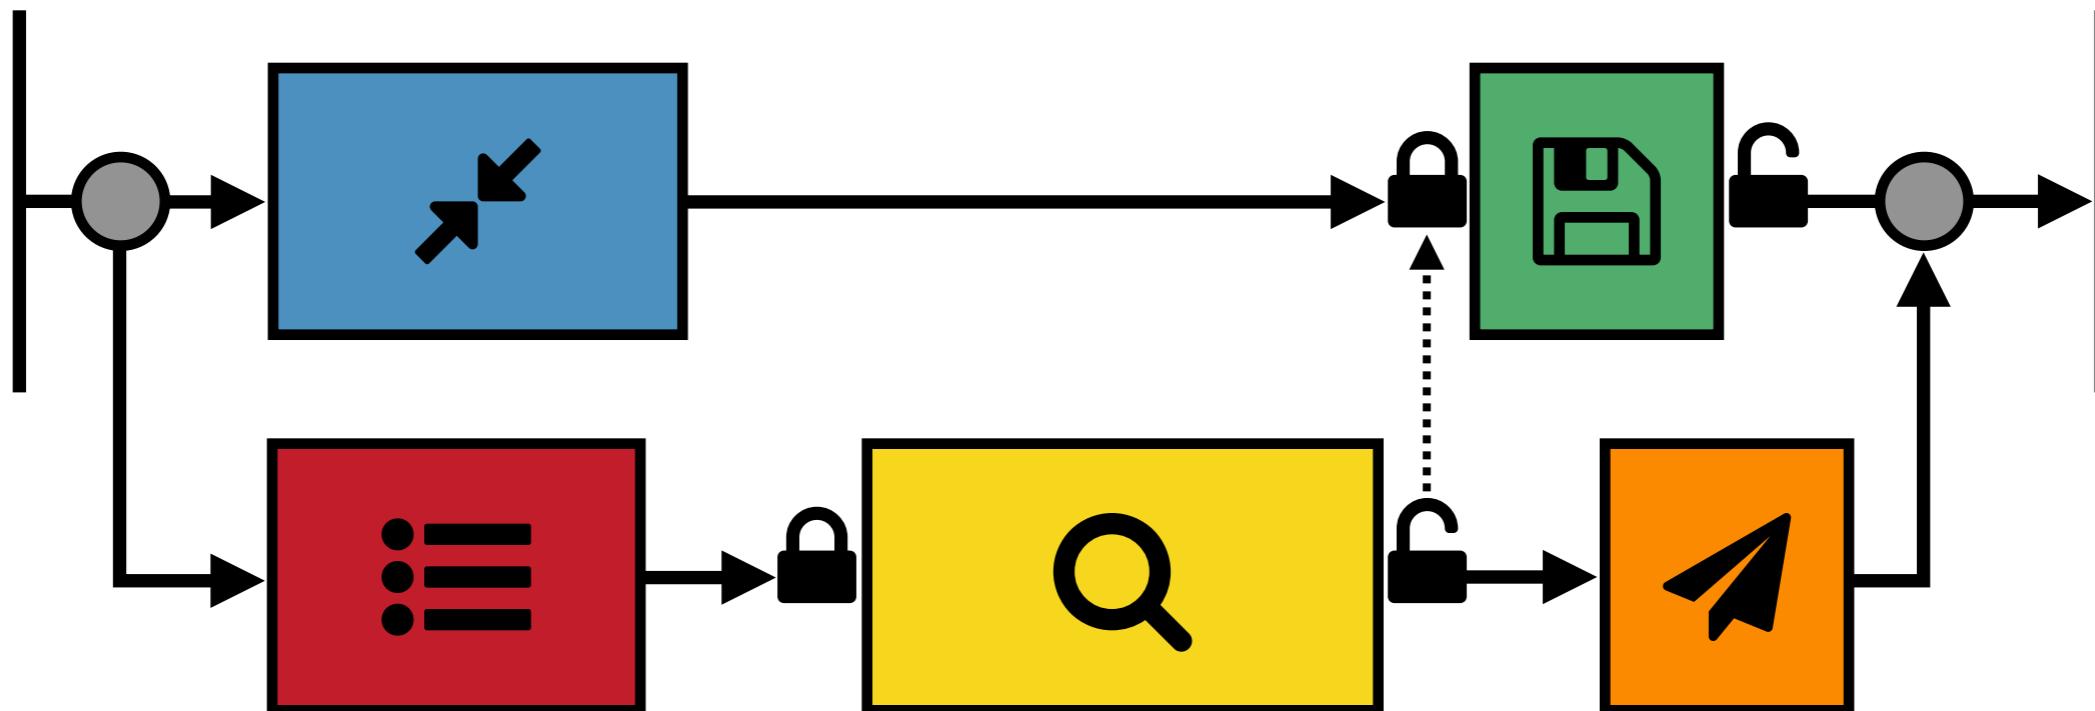
If we could magically speed up  ...



No program speedup

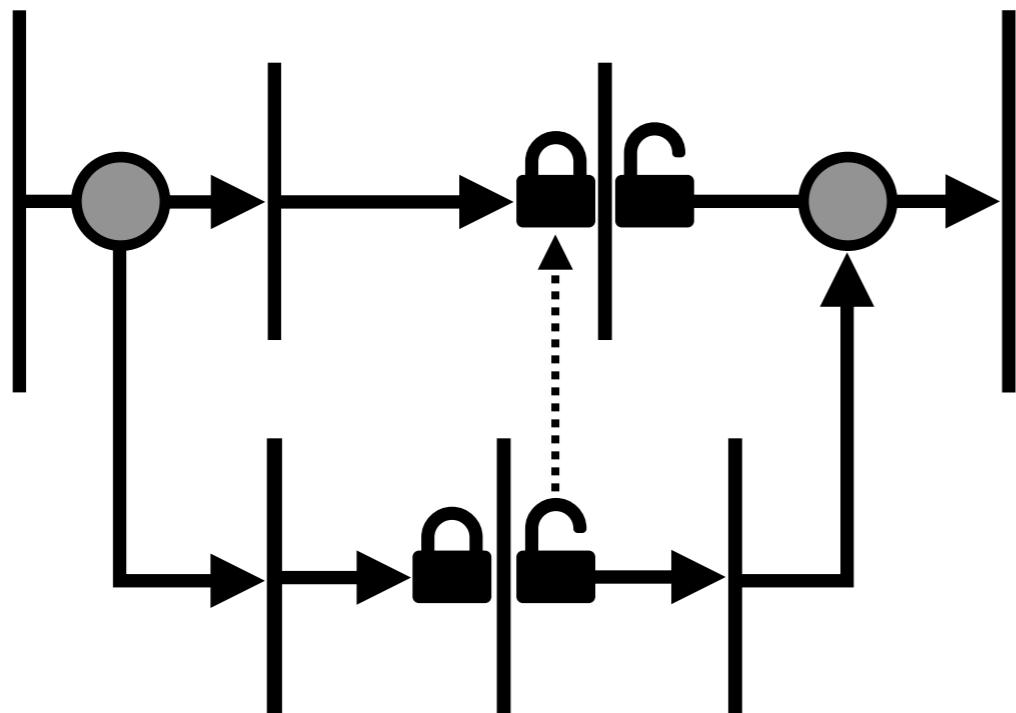
Performance Experiments

We're going to have to do this without magic.



Performance Experiments

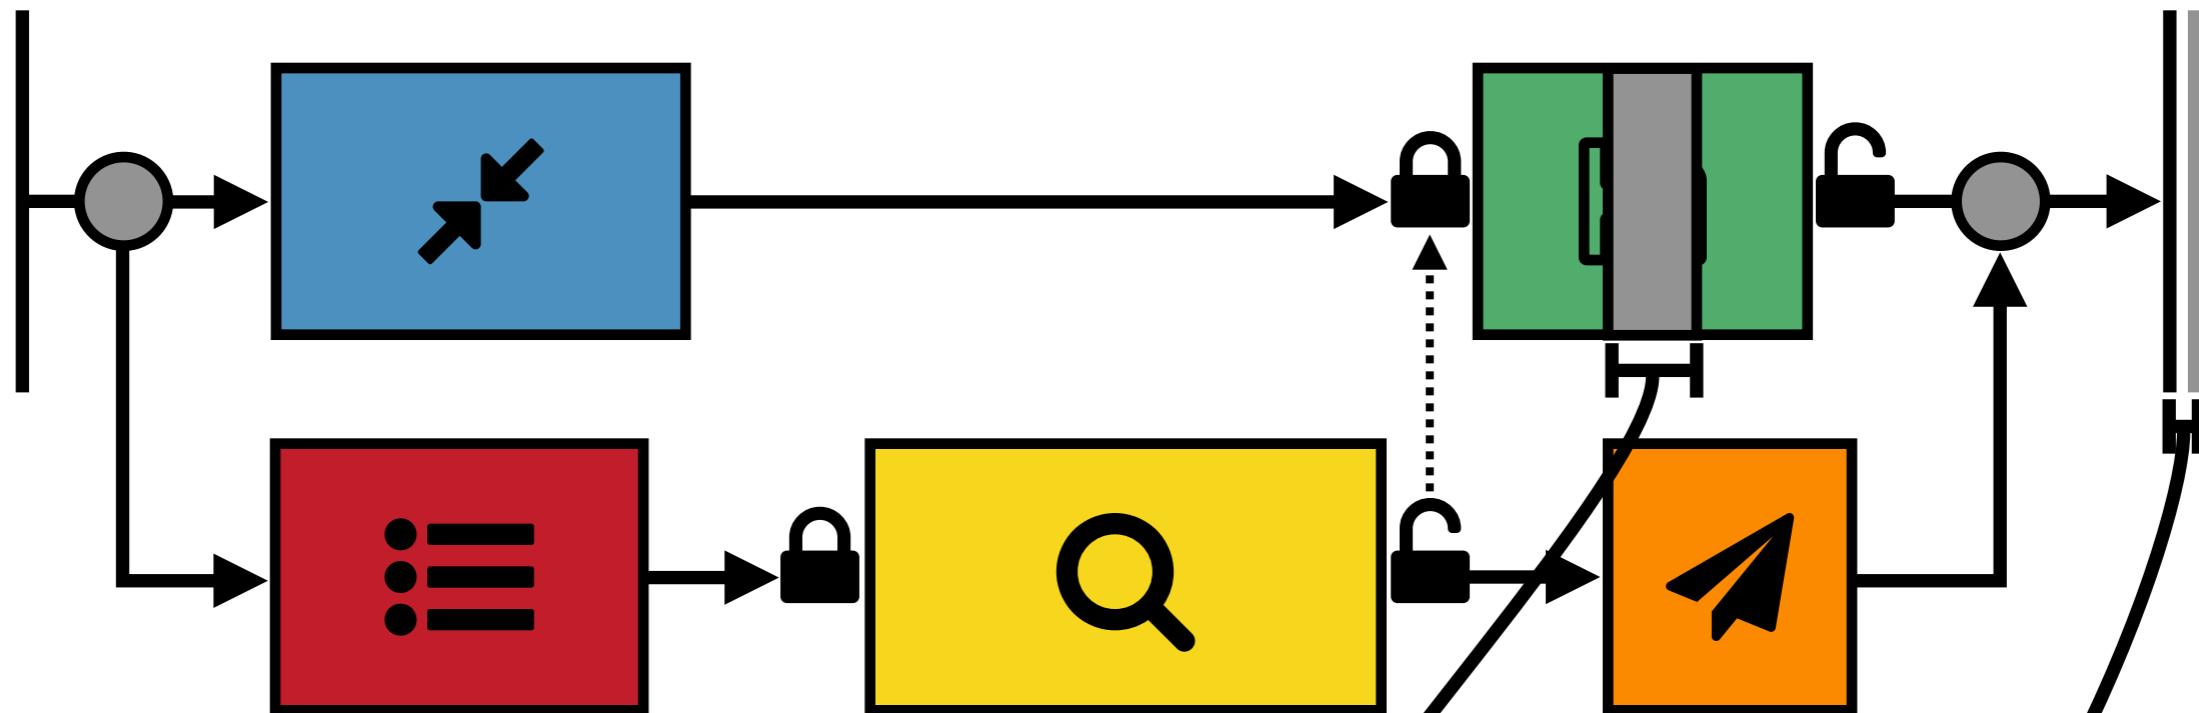
We're going to have to do this *without magic...*



Otherwise, we'd just do this...

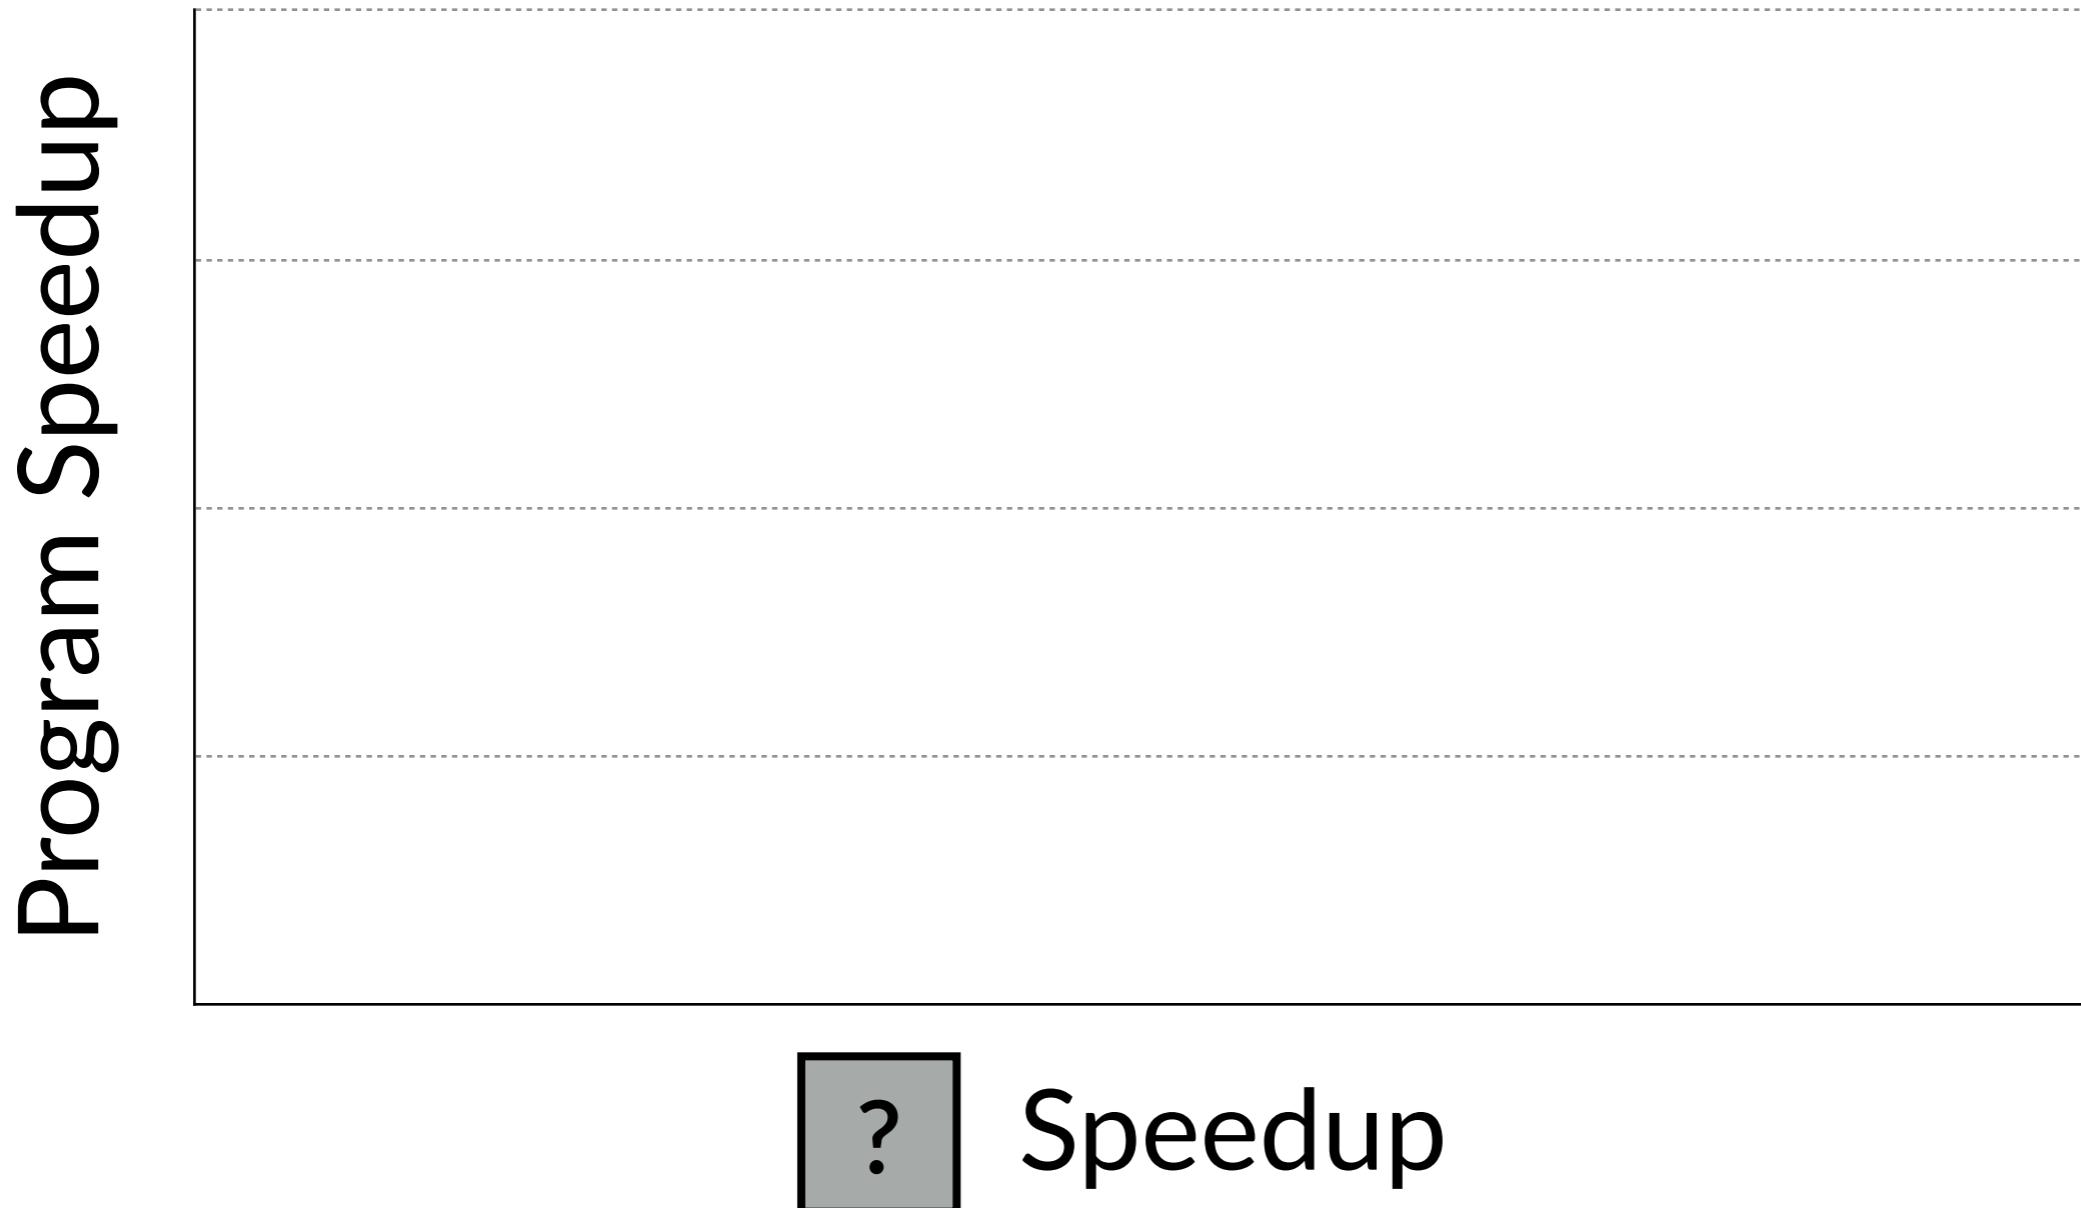
Virtual Speedup

“Speed up”  by slowing everything else down.

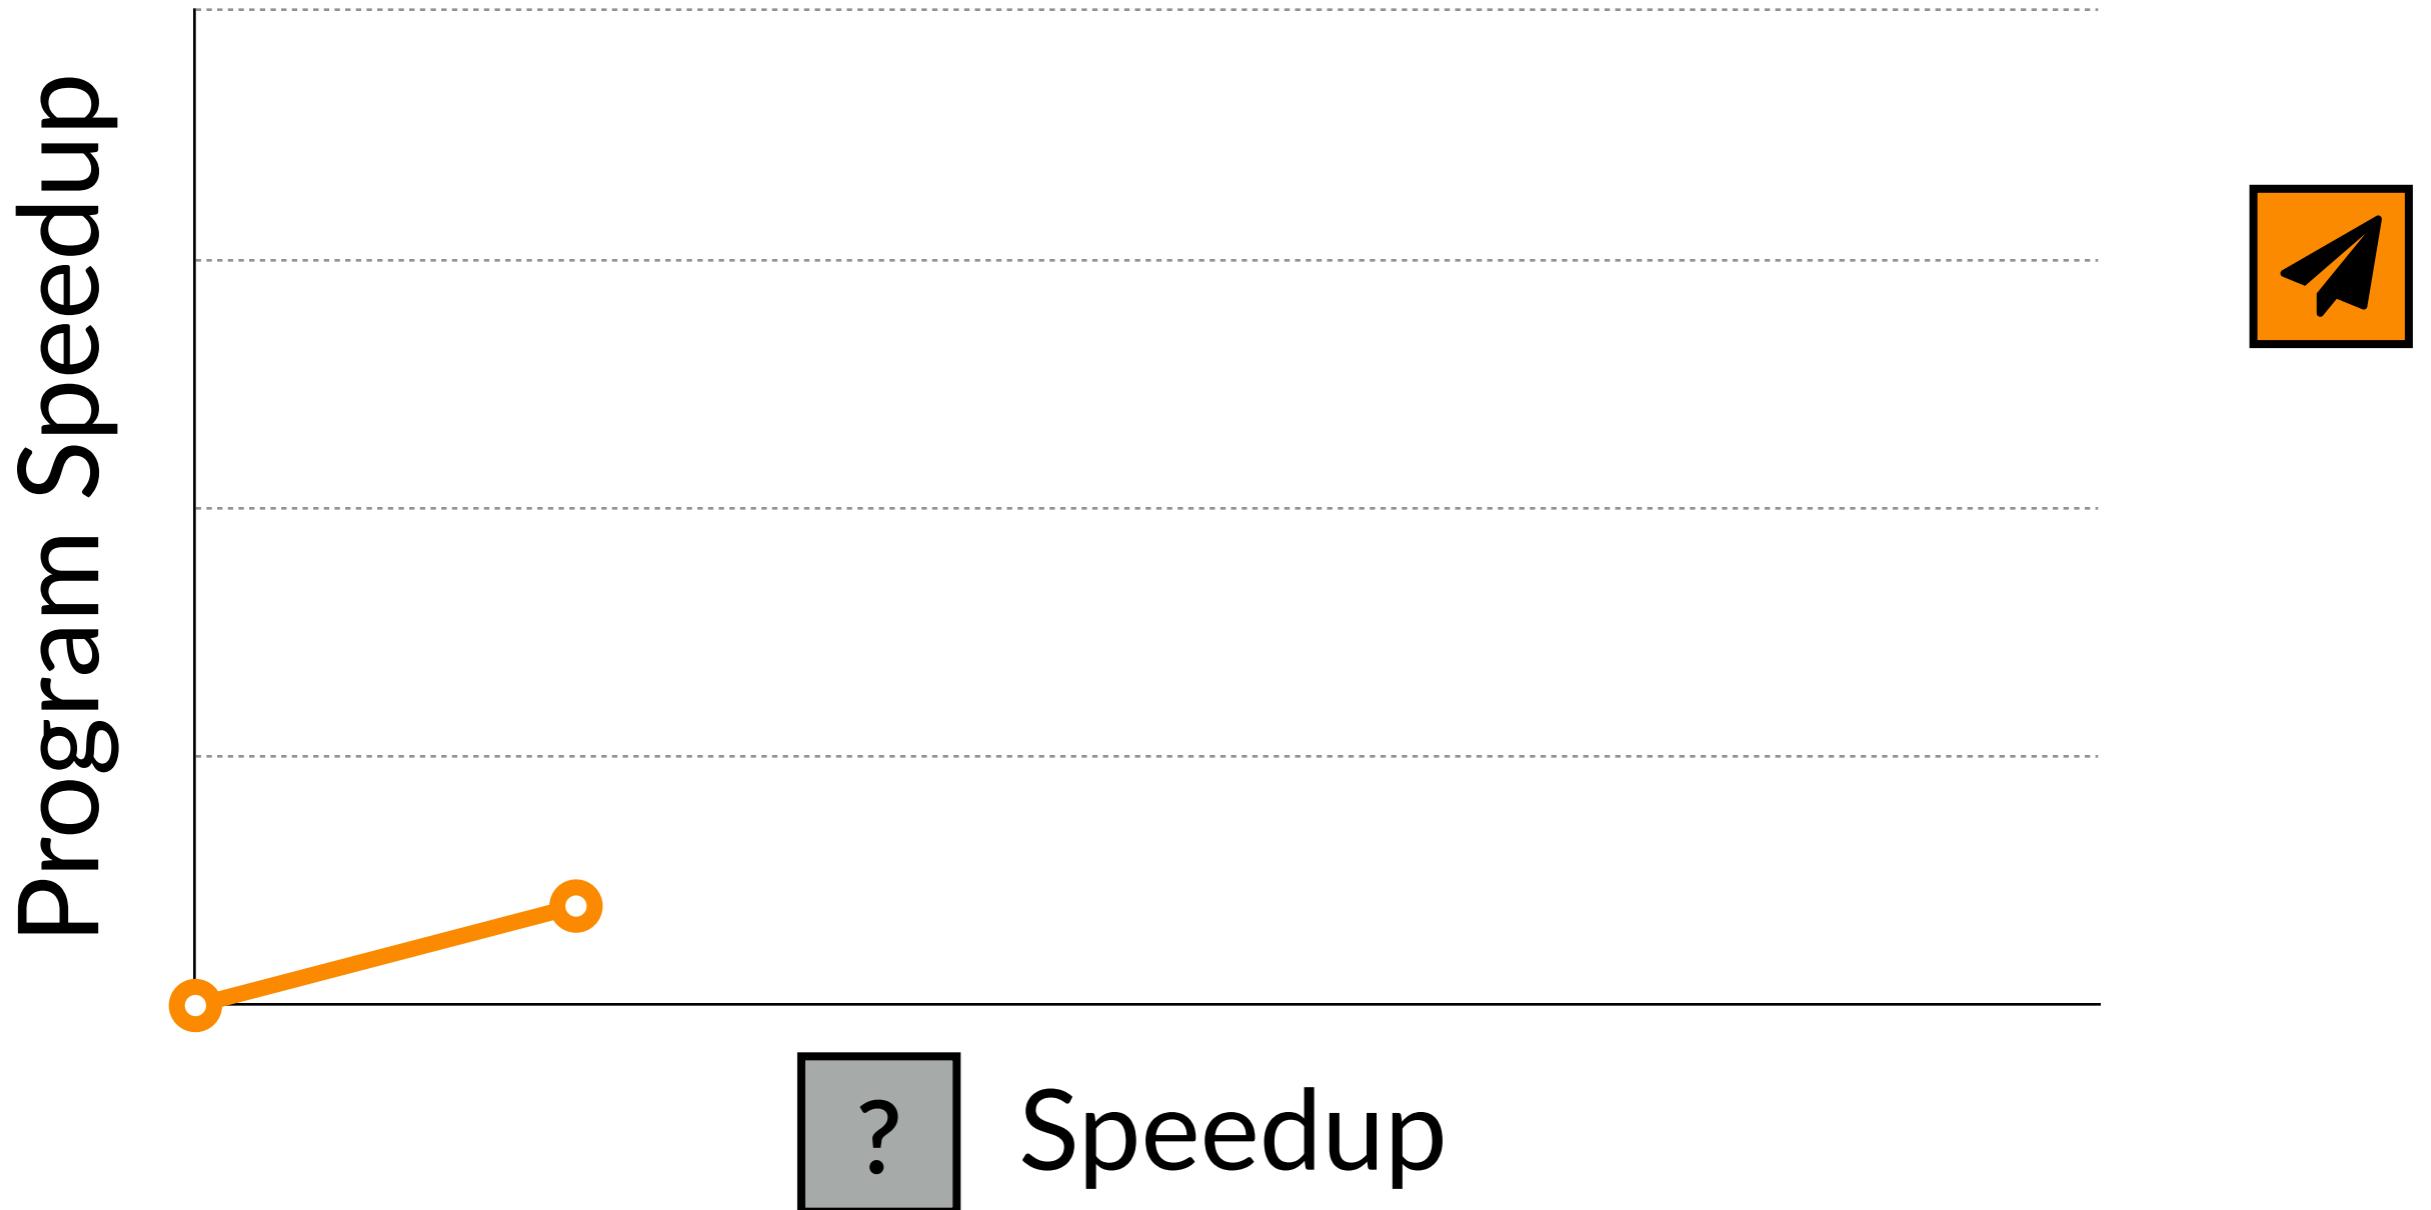


Speeding up  by this much...
speeds up the program by this much.

Speedup Results

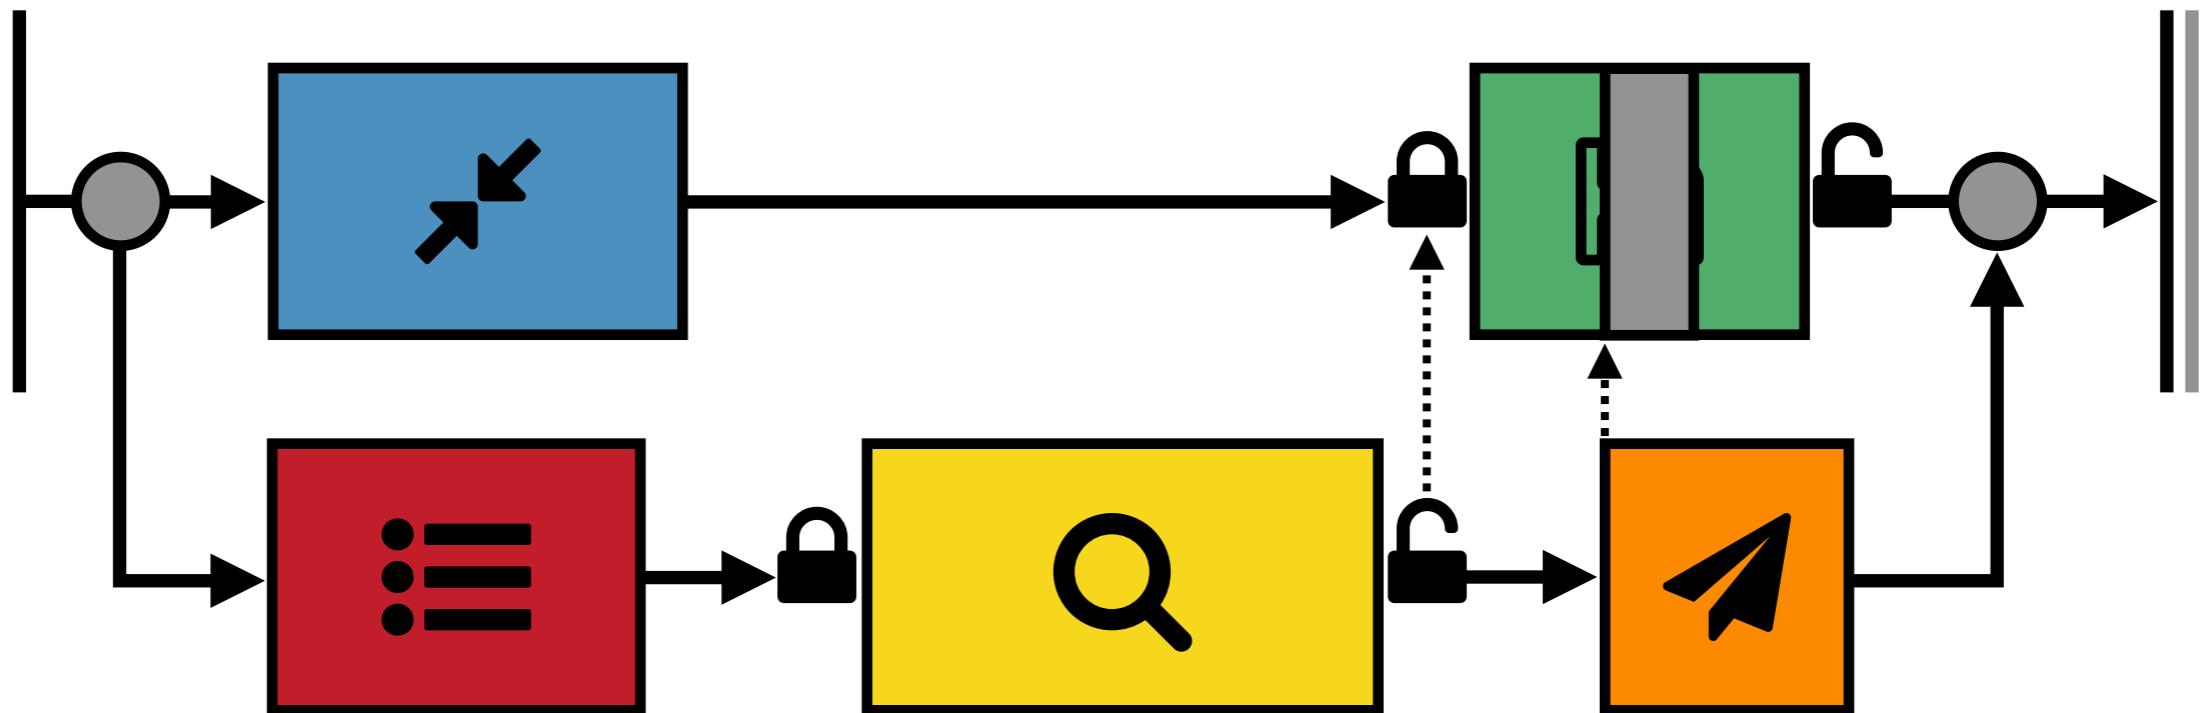


Speedup Results



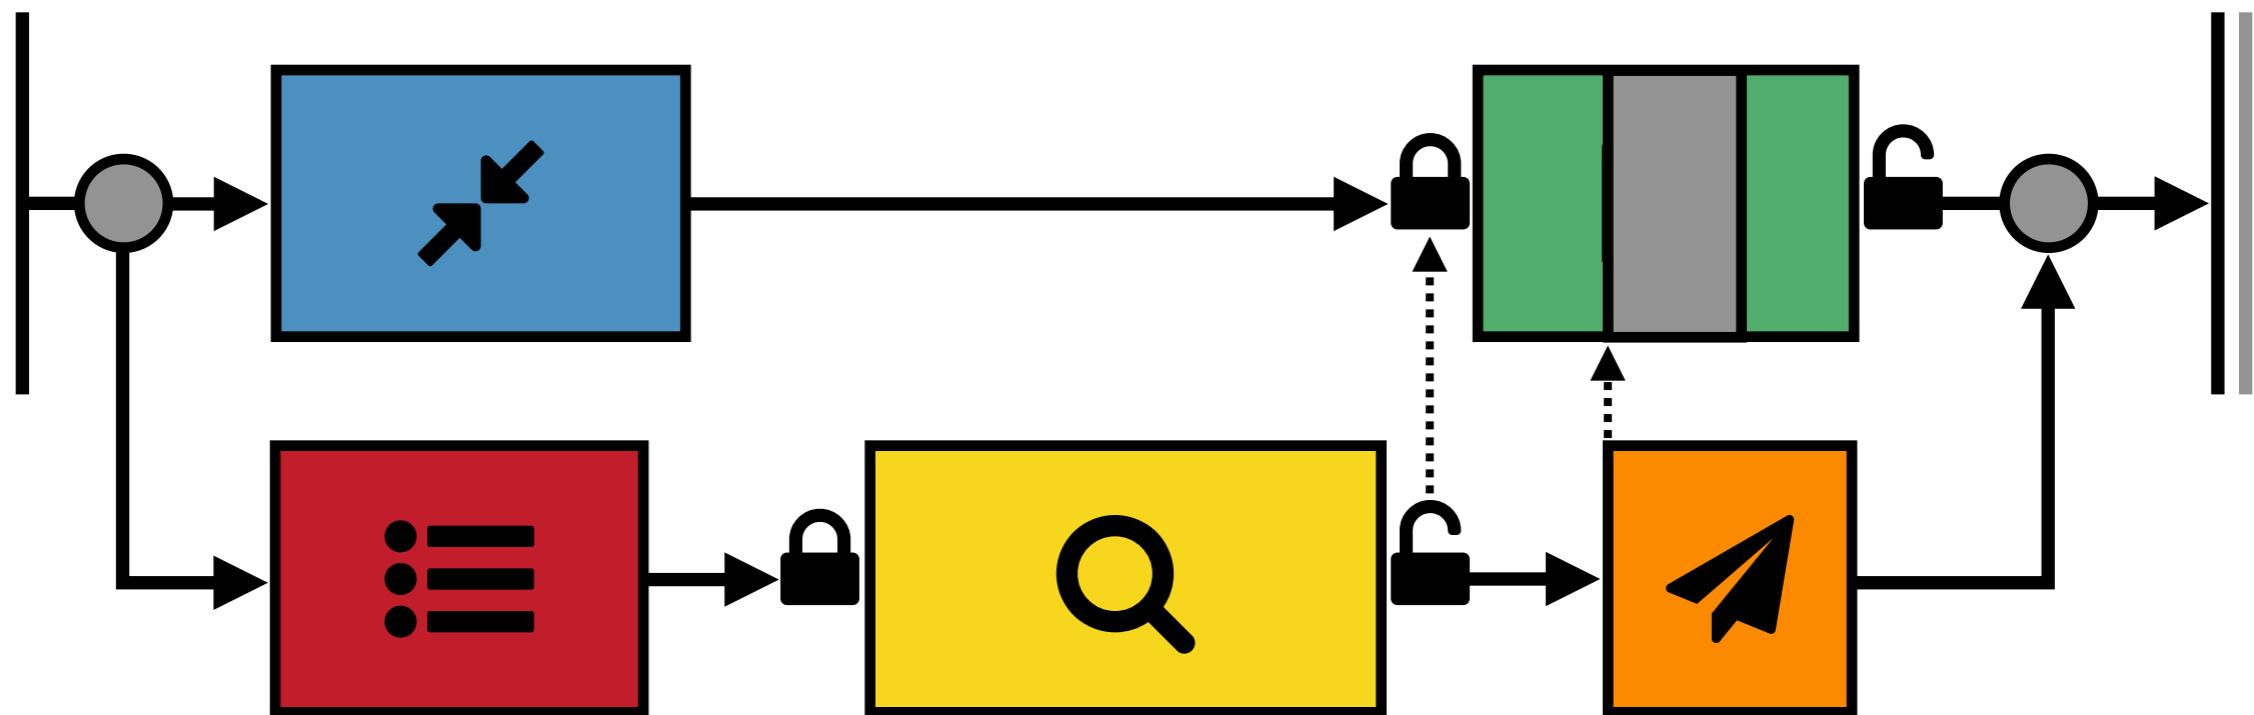
Virtual Speedup

“Speed up”  by slowing everything else down.



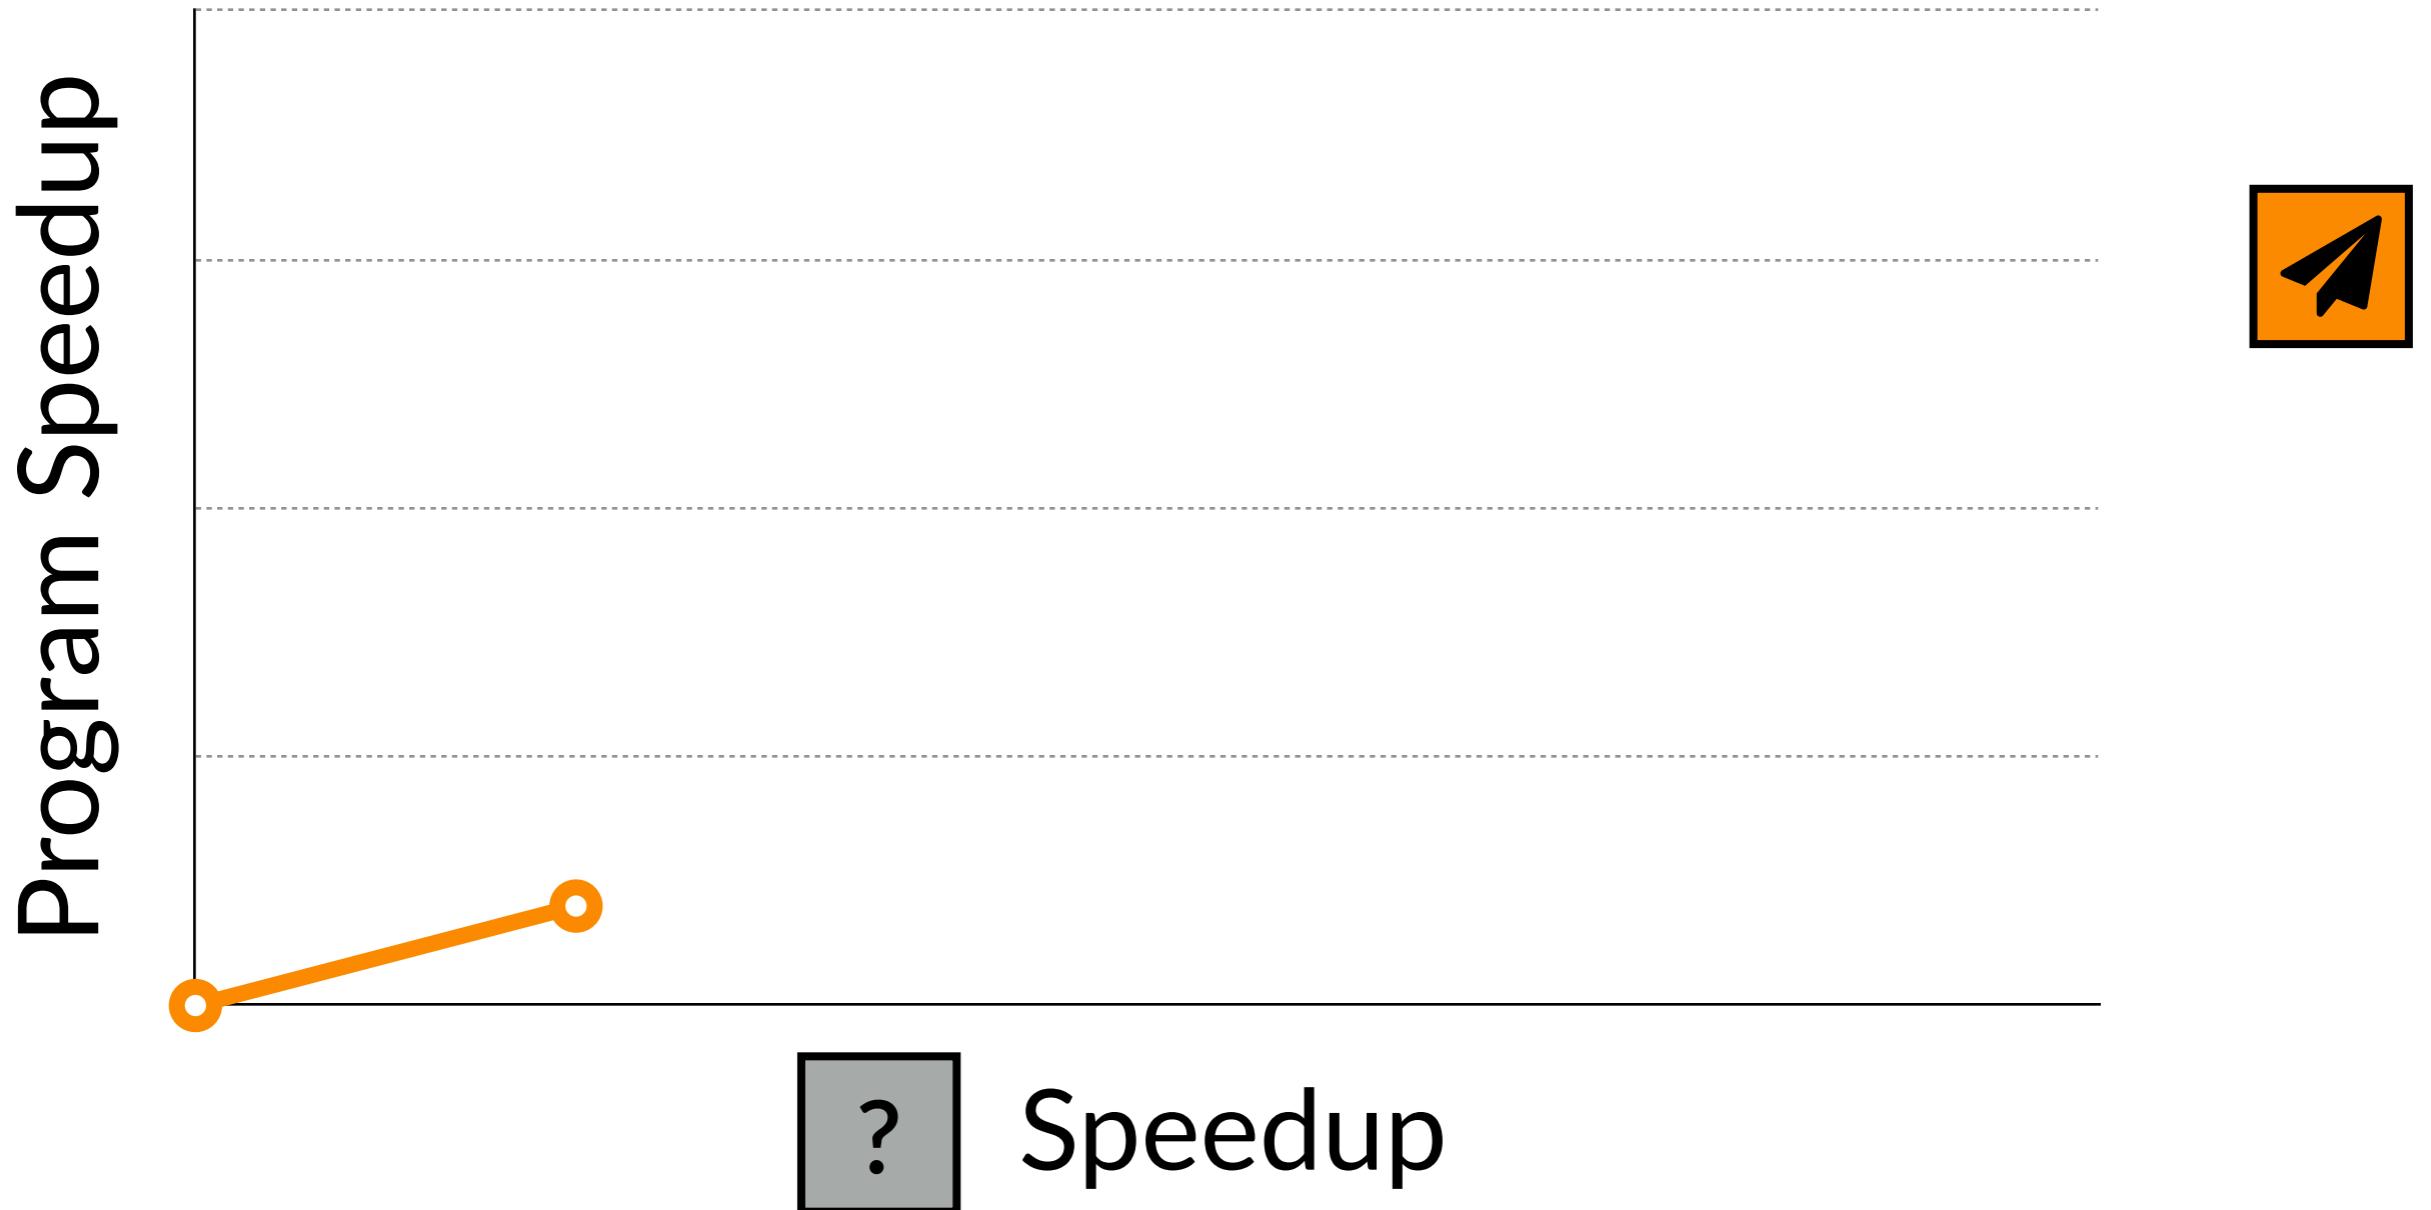
Virtual Speedup

“Speed up”  by slowing everything else down.

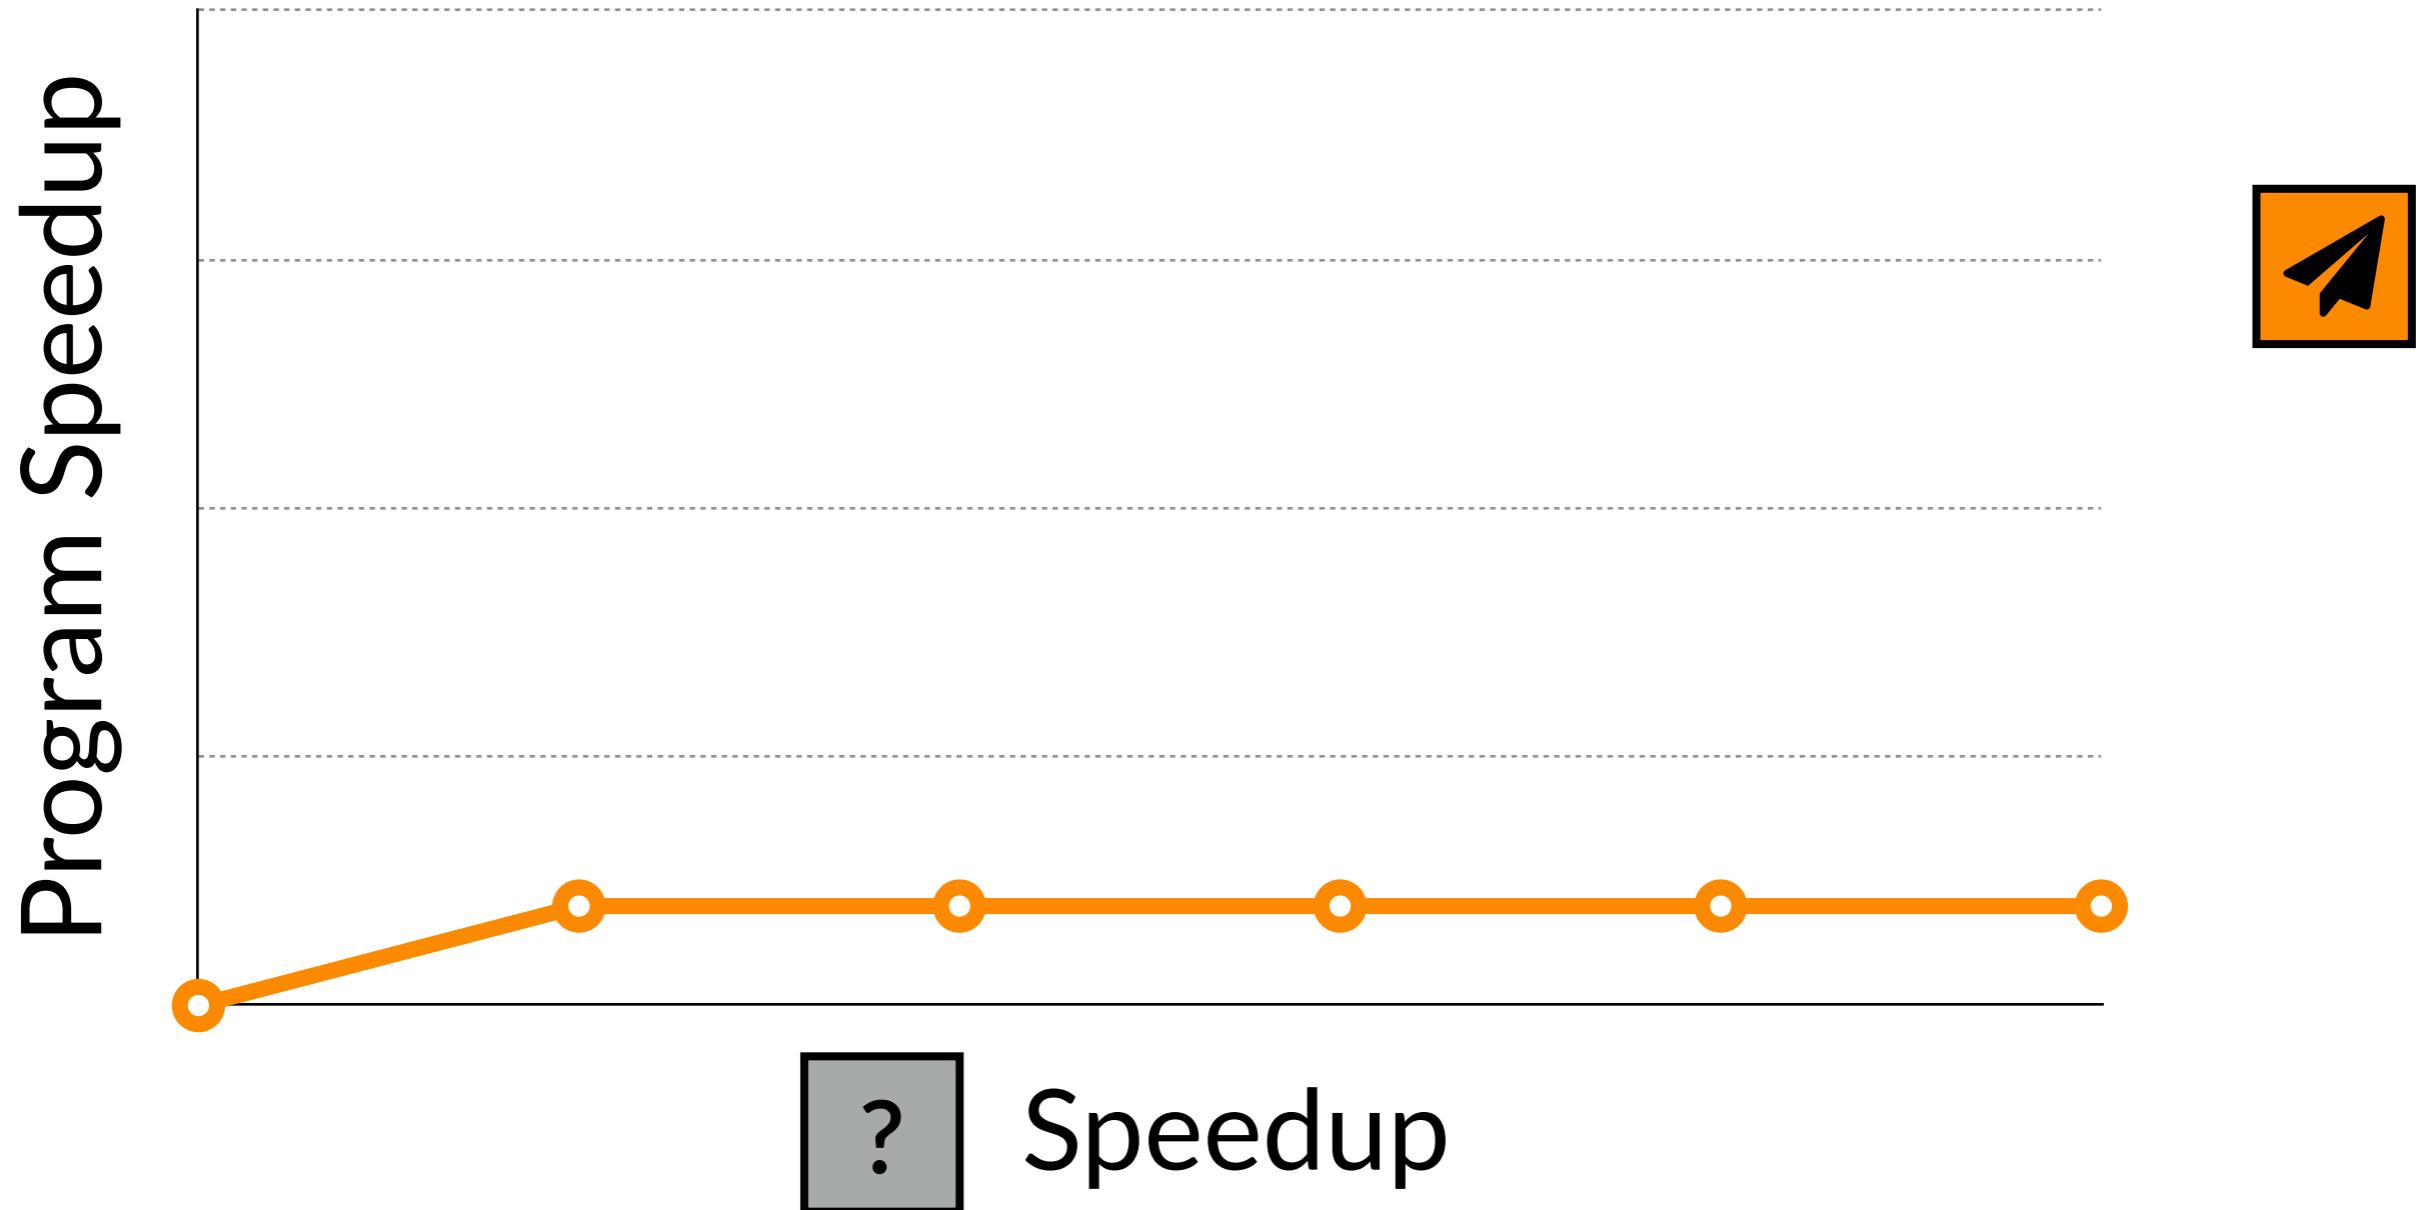


A larger speedup has no additional effect

Speedup Results

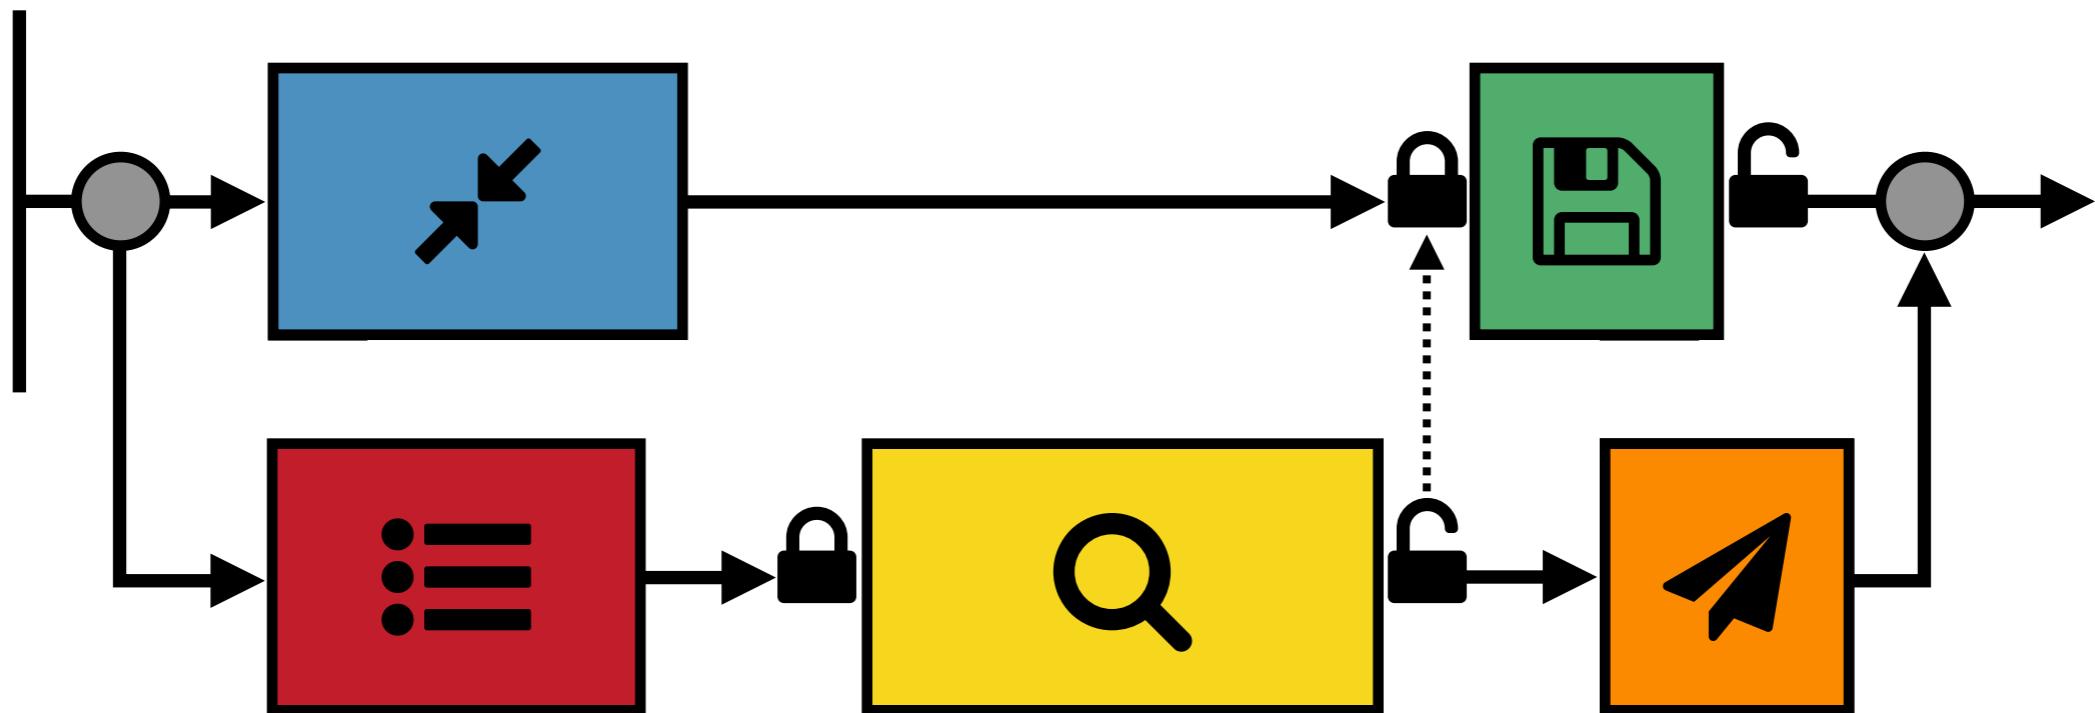


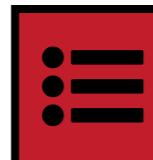
Speedup Results



Virtual Speedup

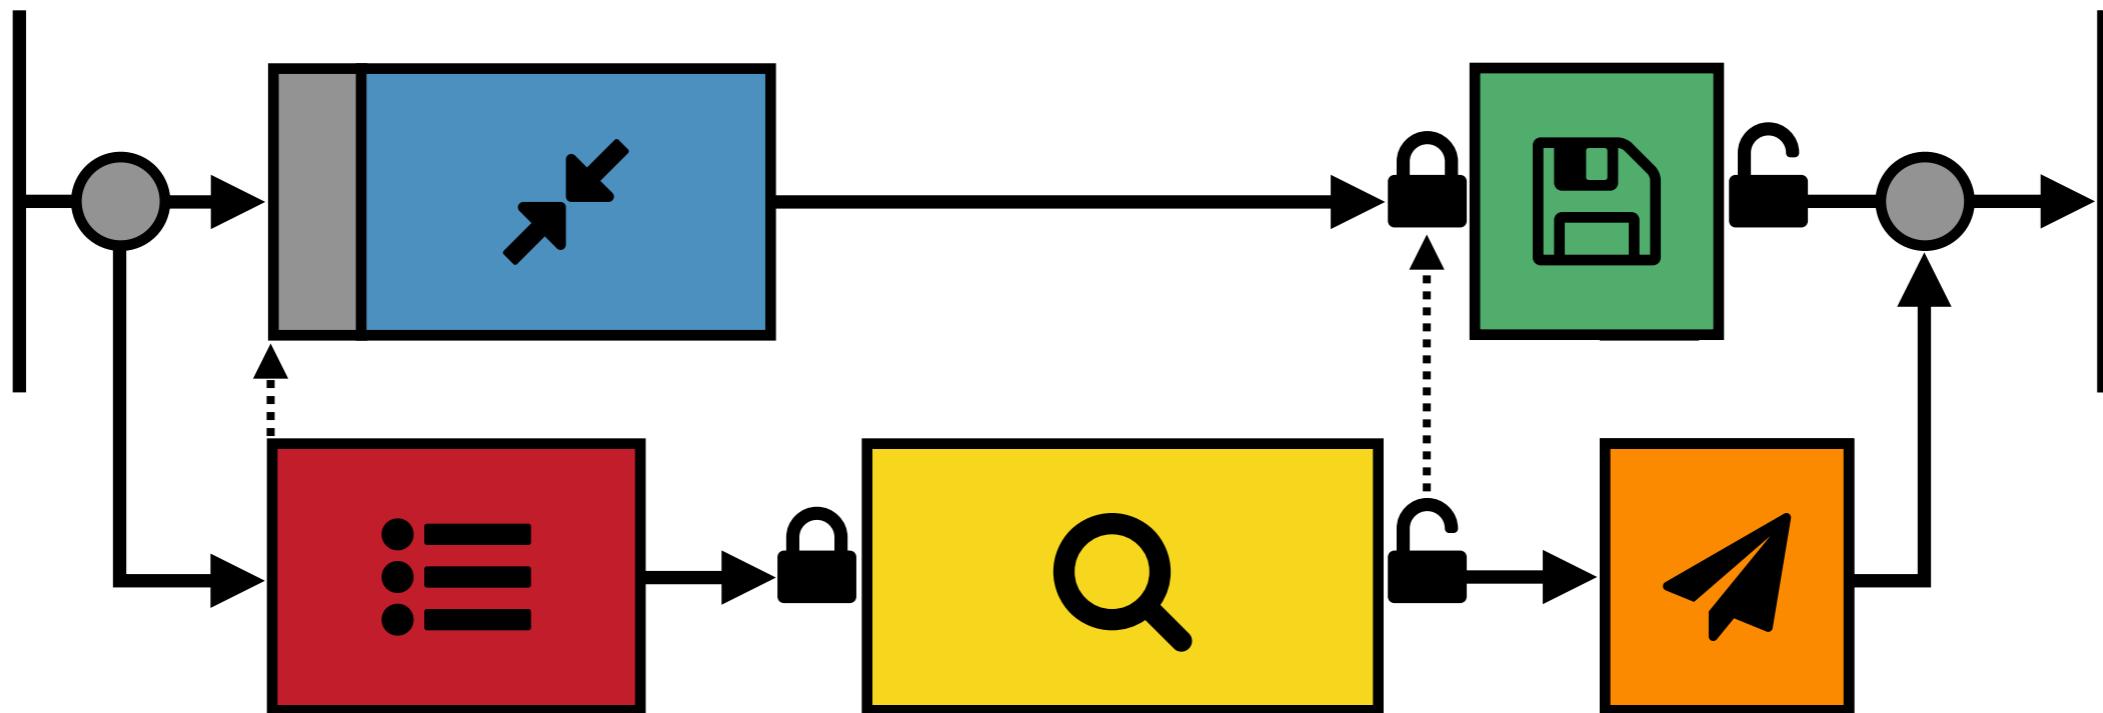
“Speed up”  by slowing everything else down.

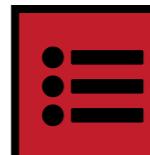


Each time  runs, pause all other threads.

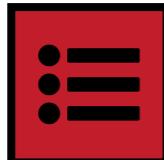
Virtual Speedup

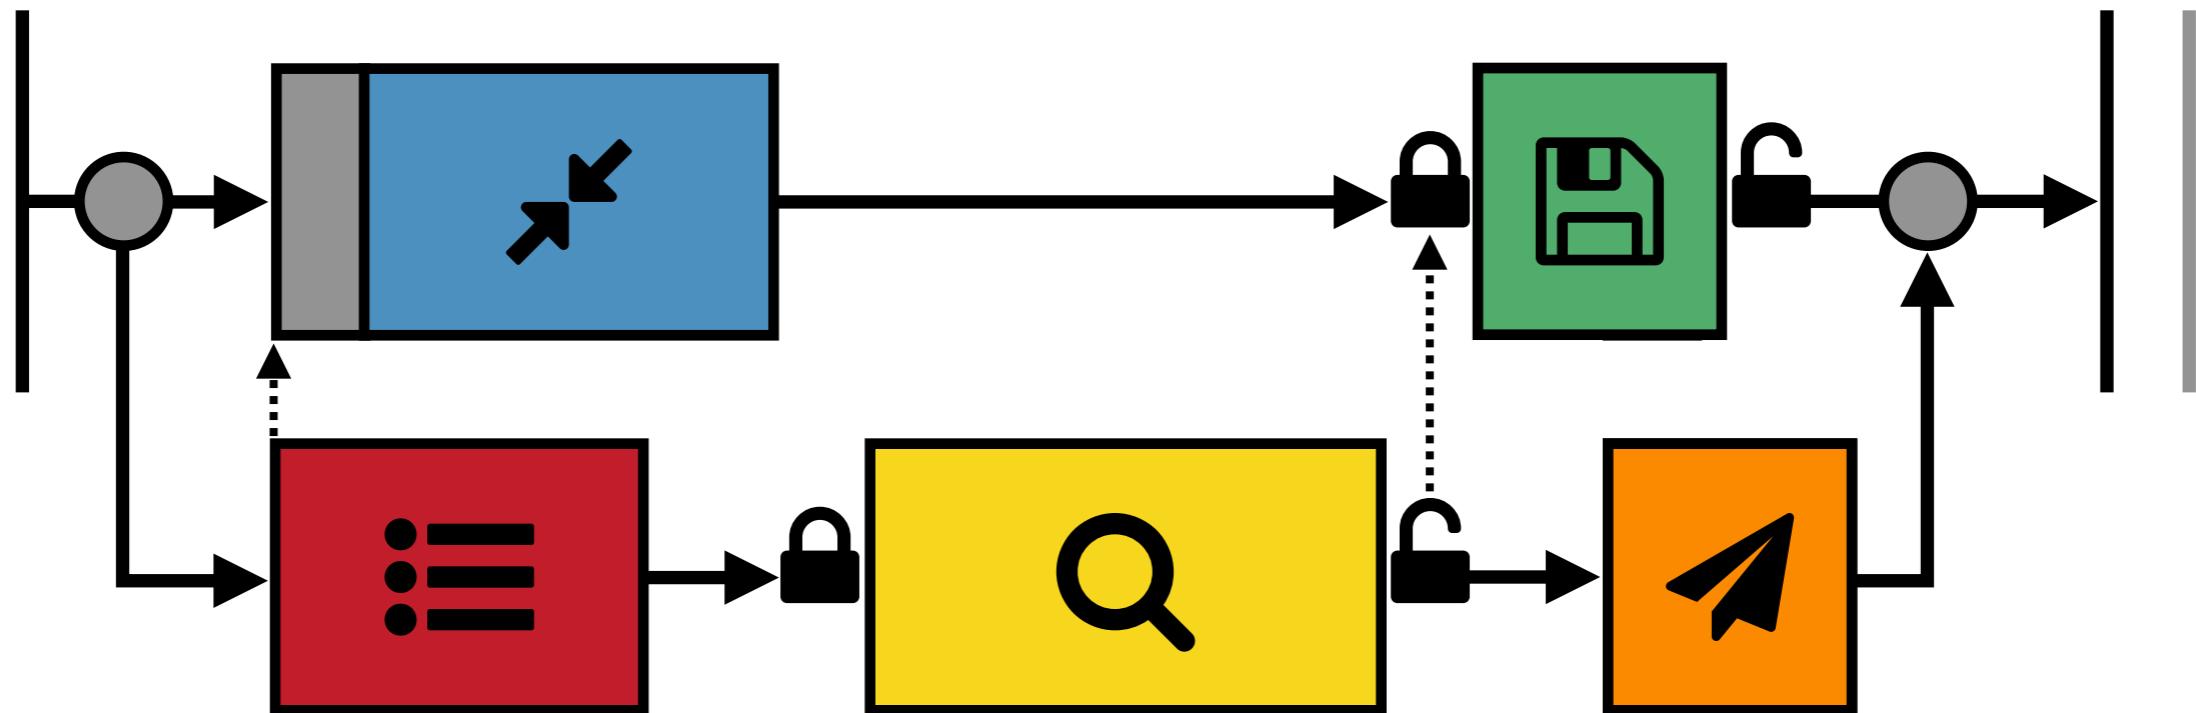
“Speed up”  by slowing everything else down.

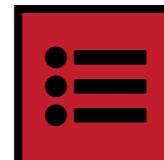


Each time  runs, pause all other threads.

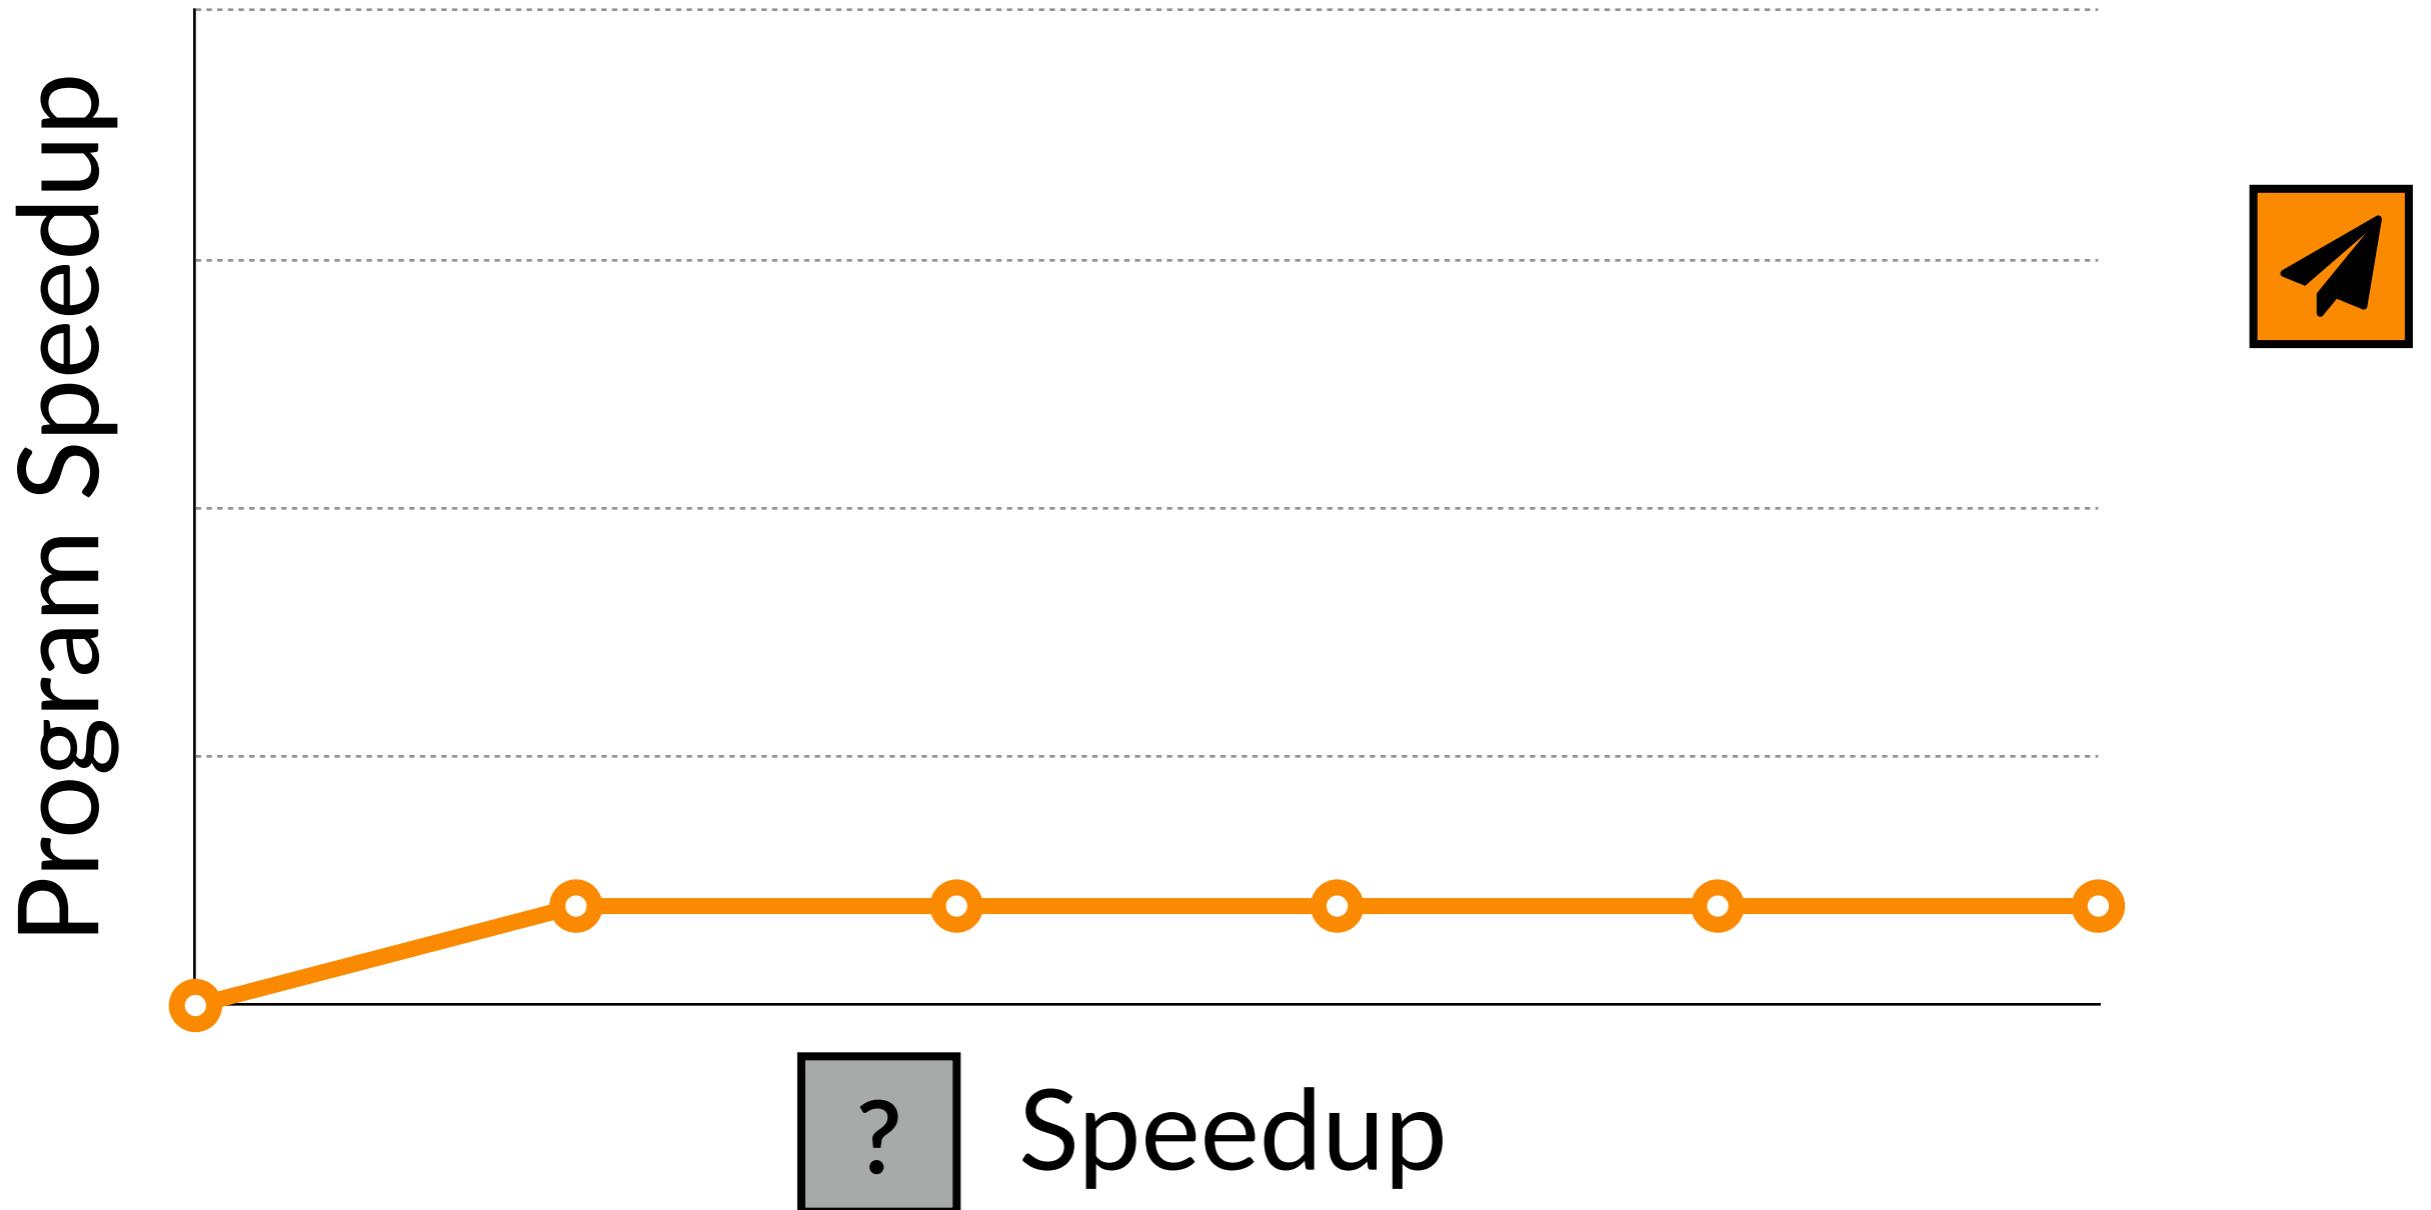
Virtual Speedup

“Speed up”  by slowing everything else down.

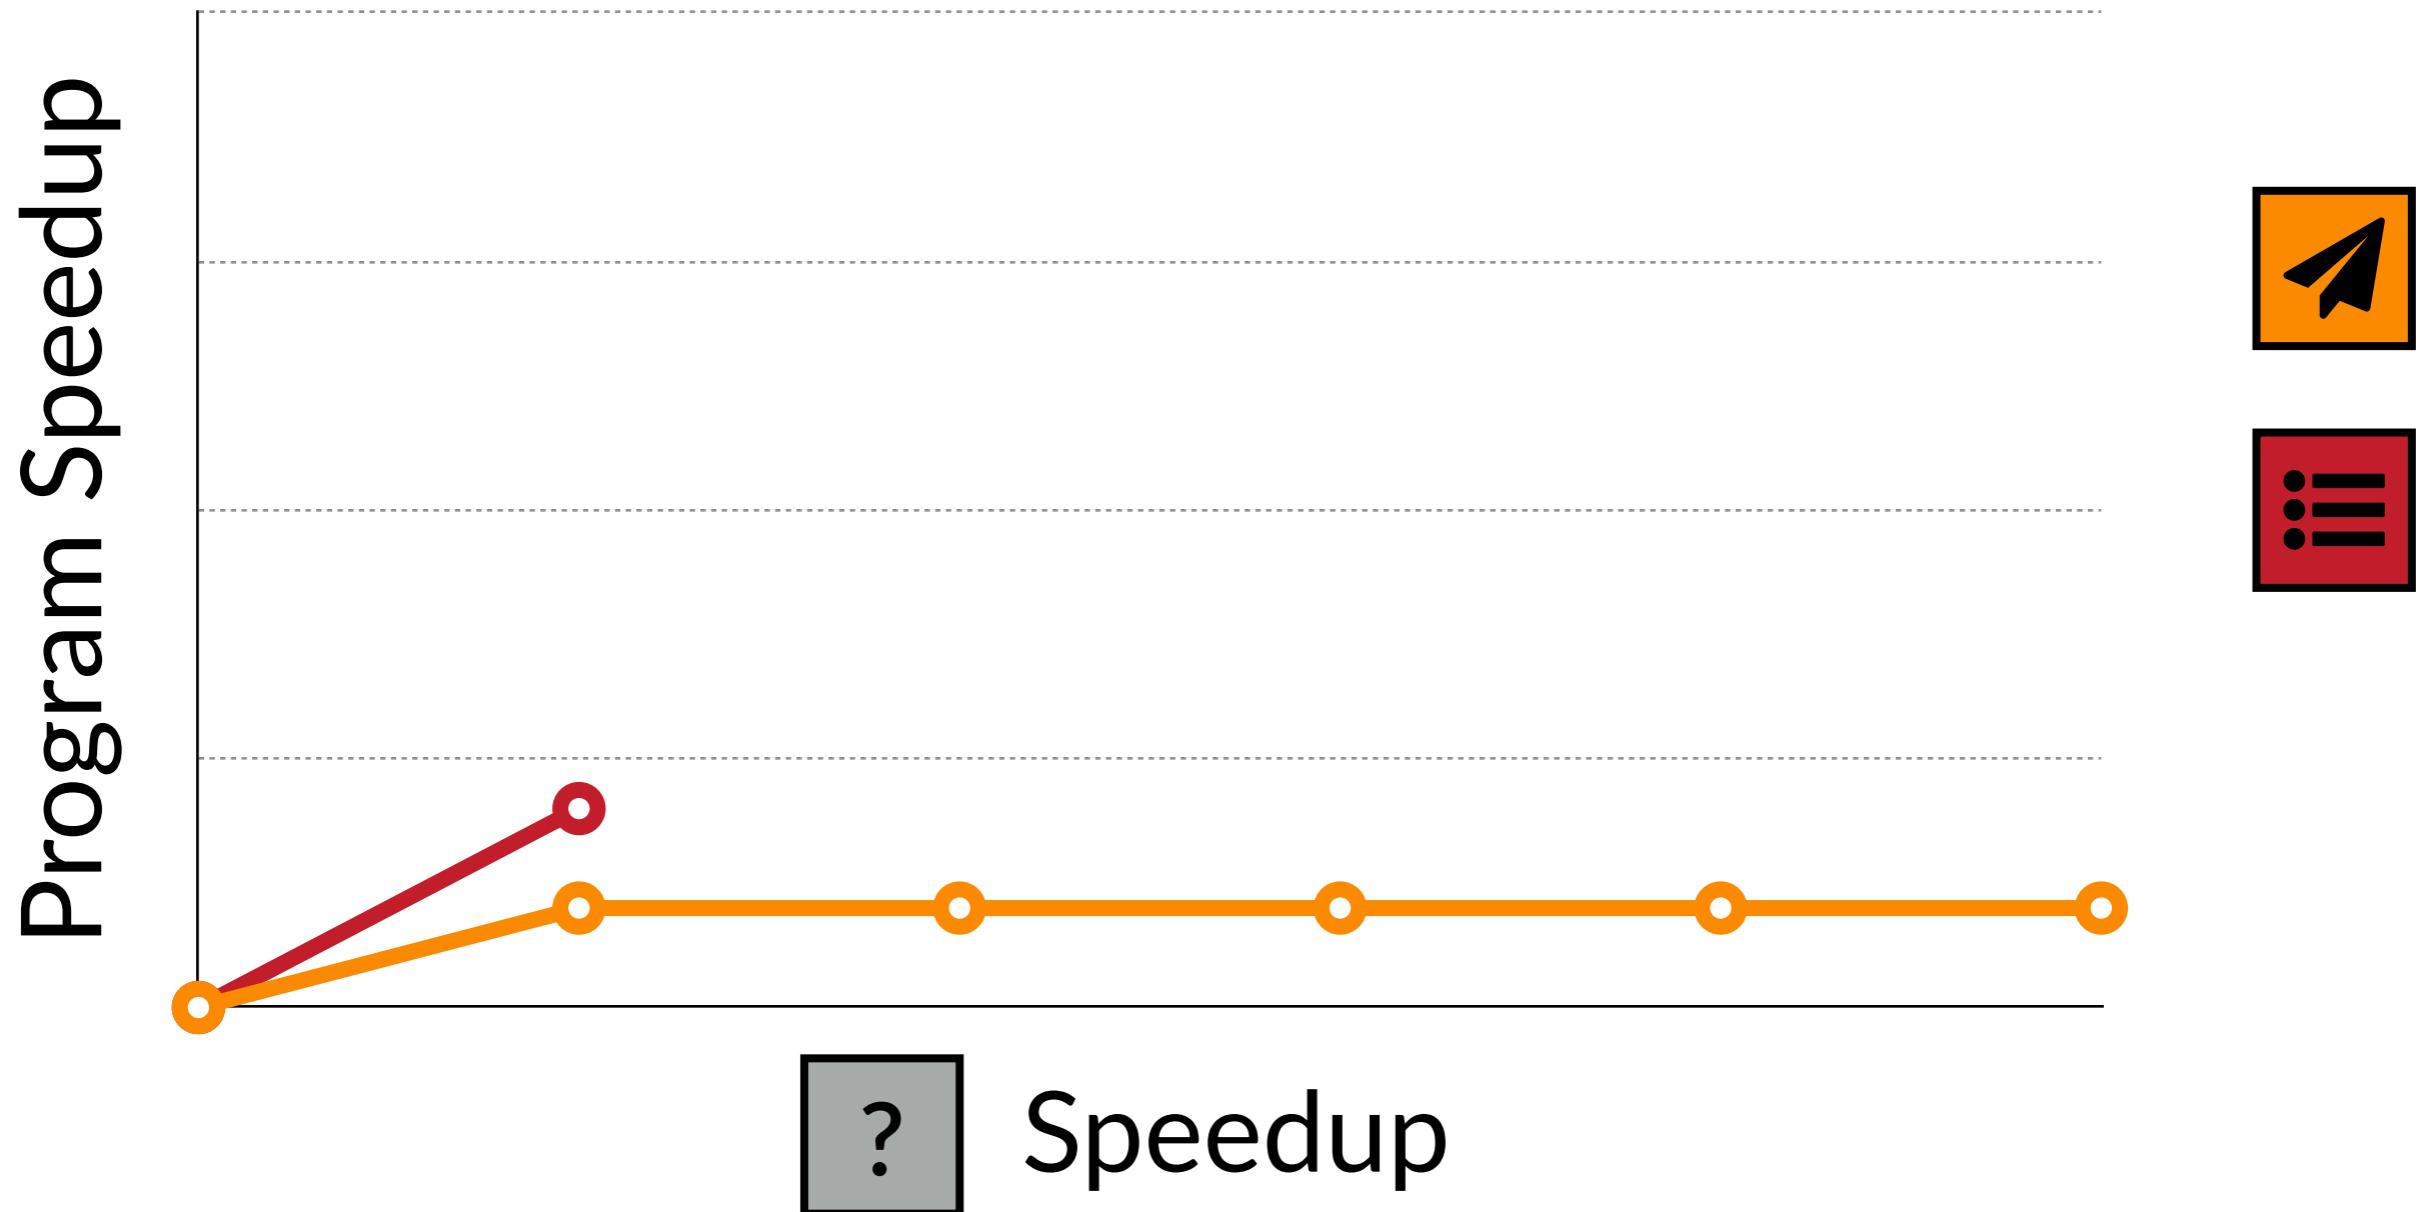


Each time  runs, pause all other threads.

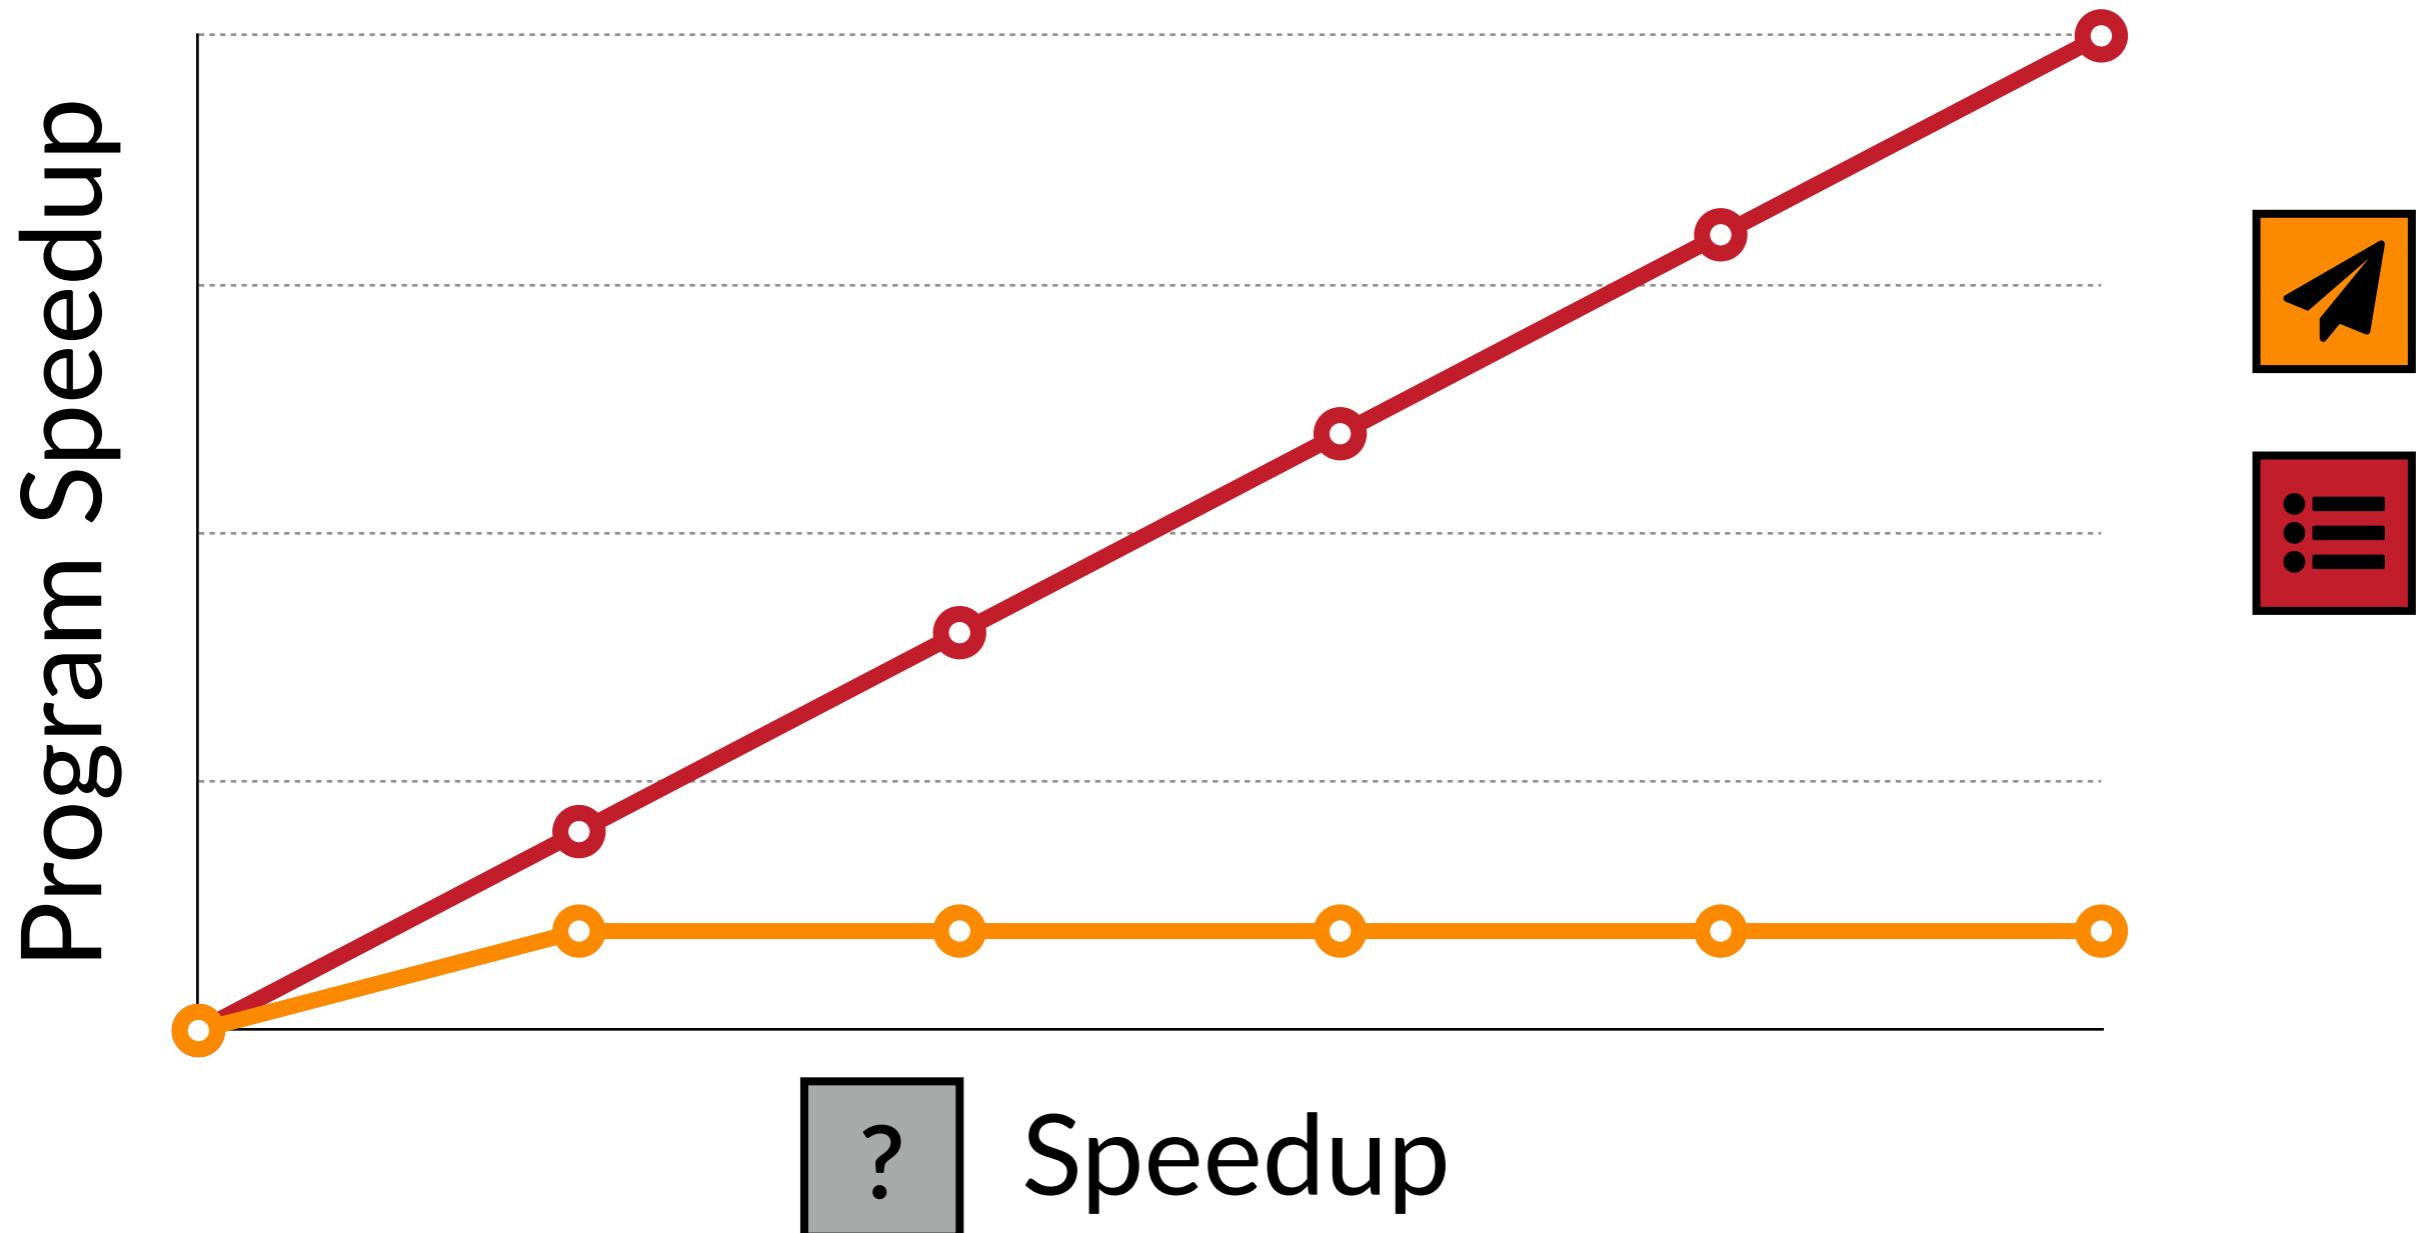
Speedup Results



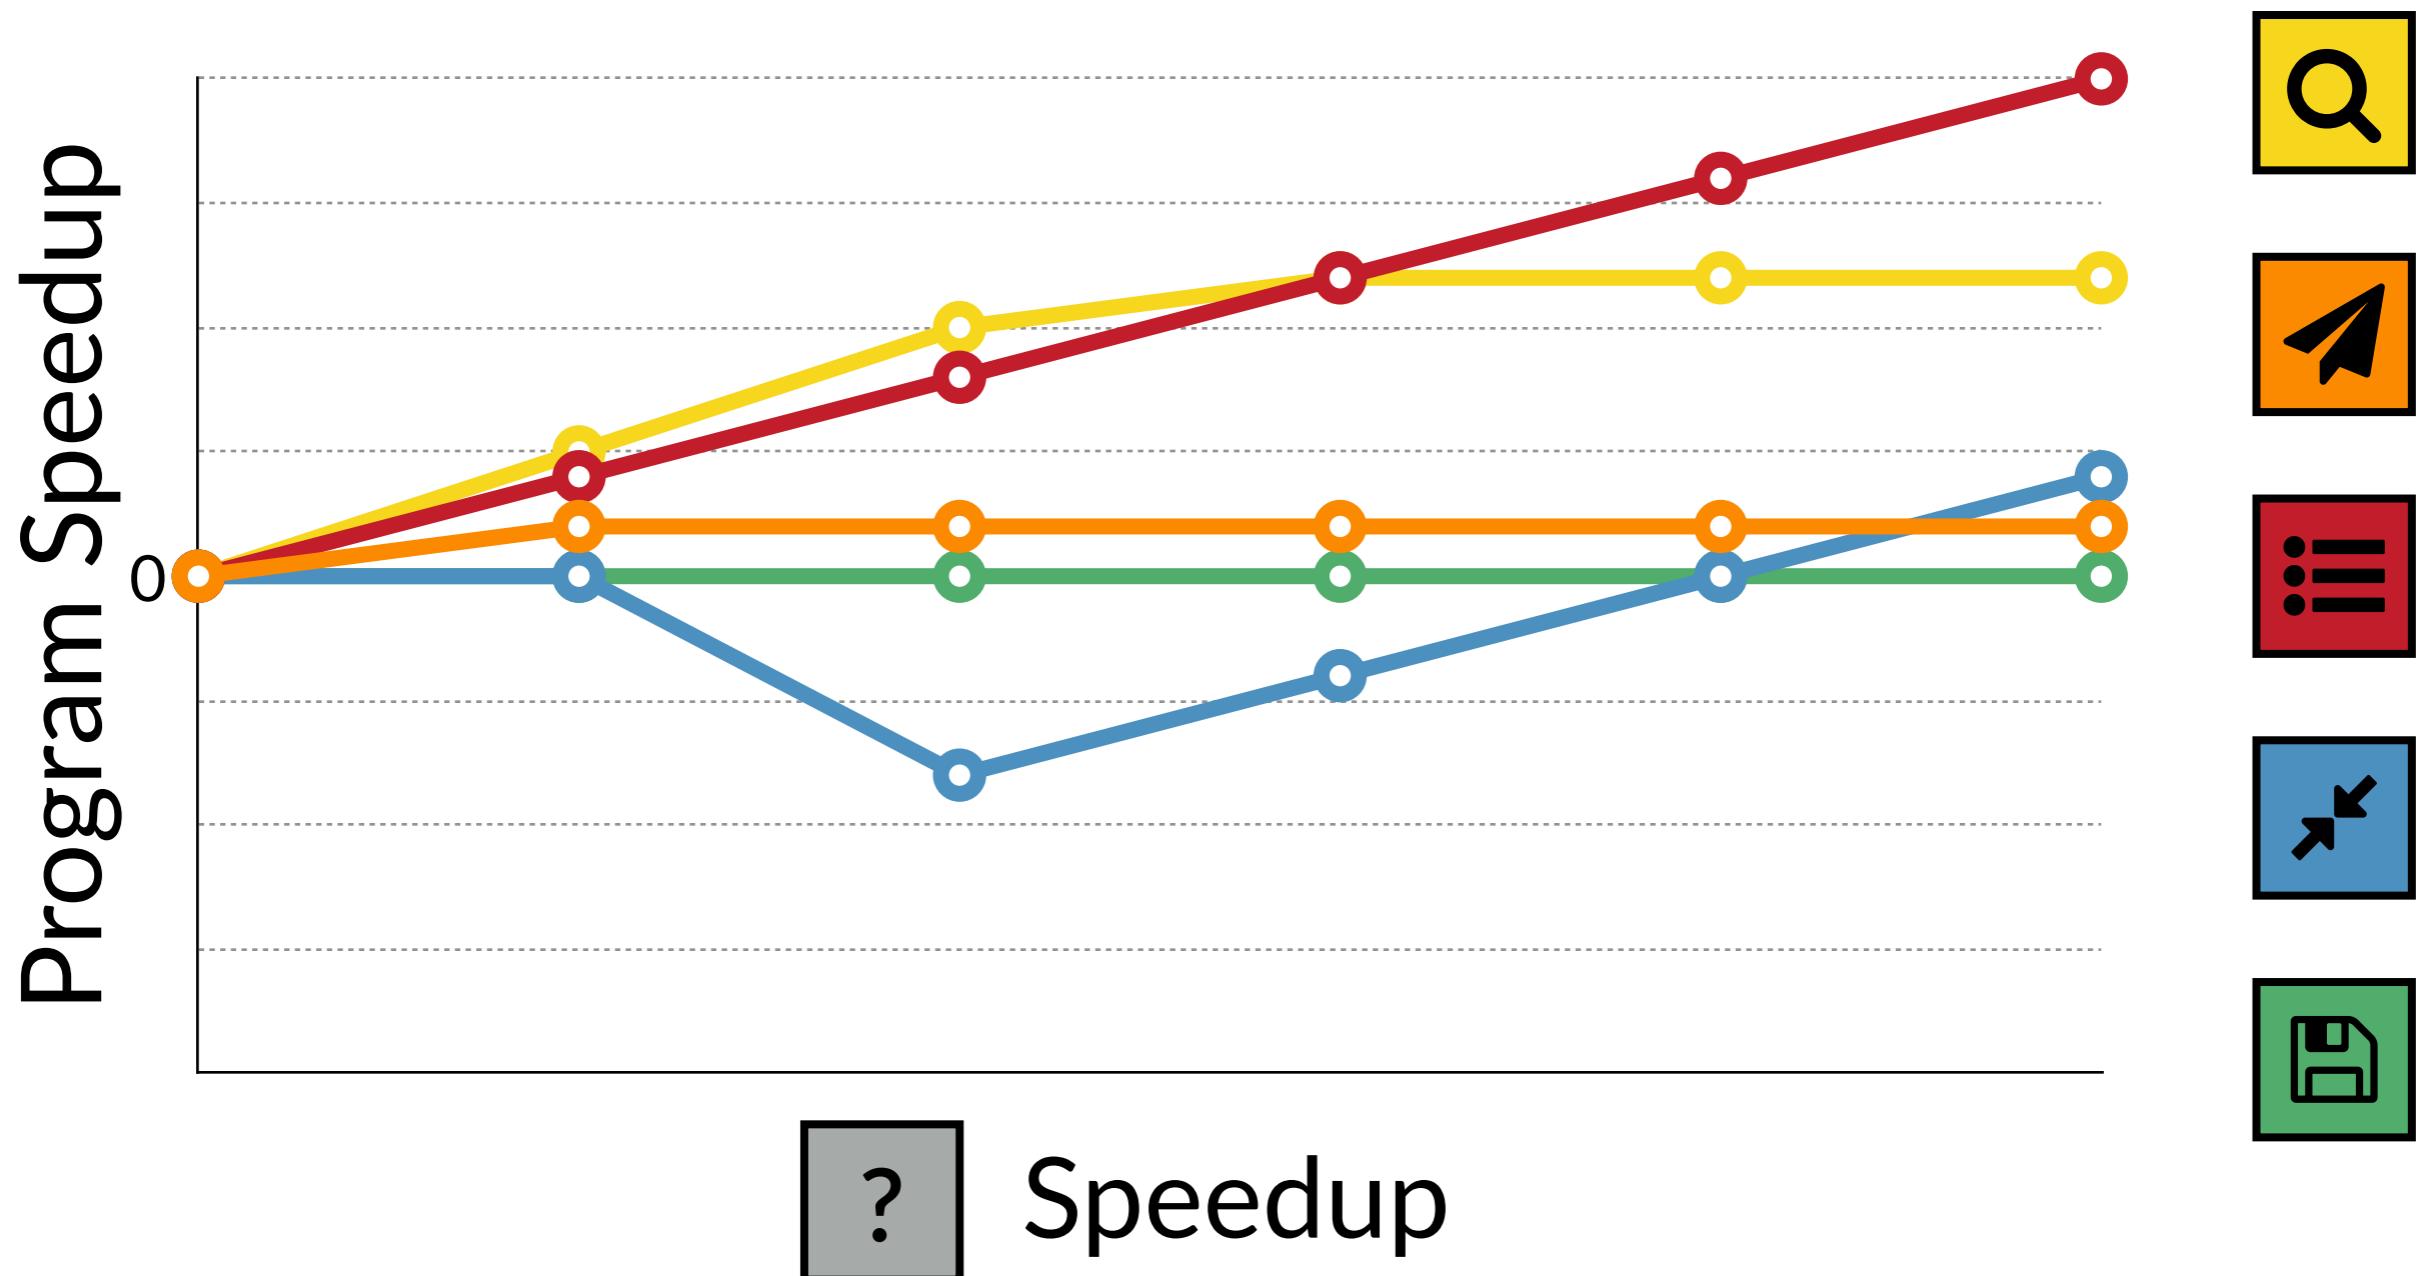
Speedup Results



Speedup Results

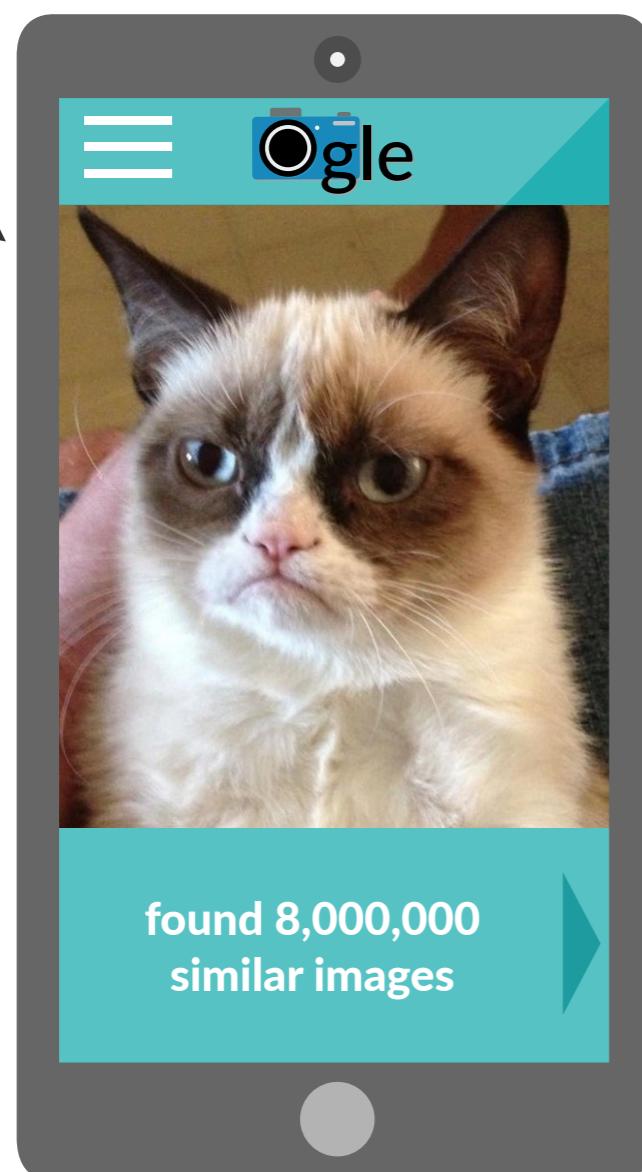


Speedup Results

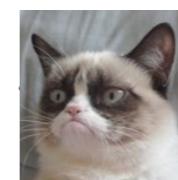
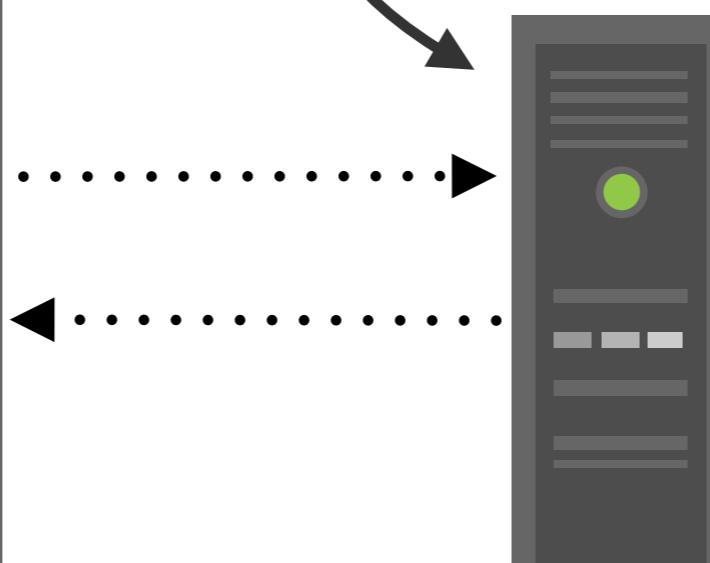




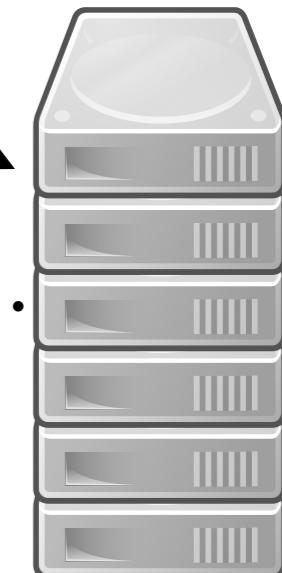
Take a picture



Send it to Ogle



Add it to
the database



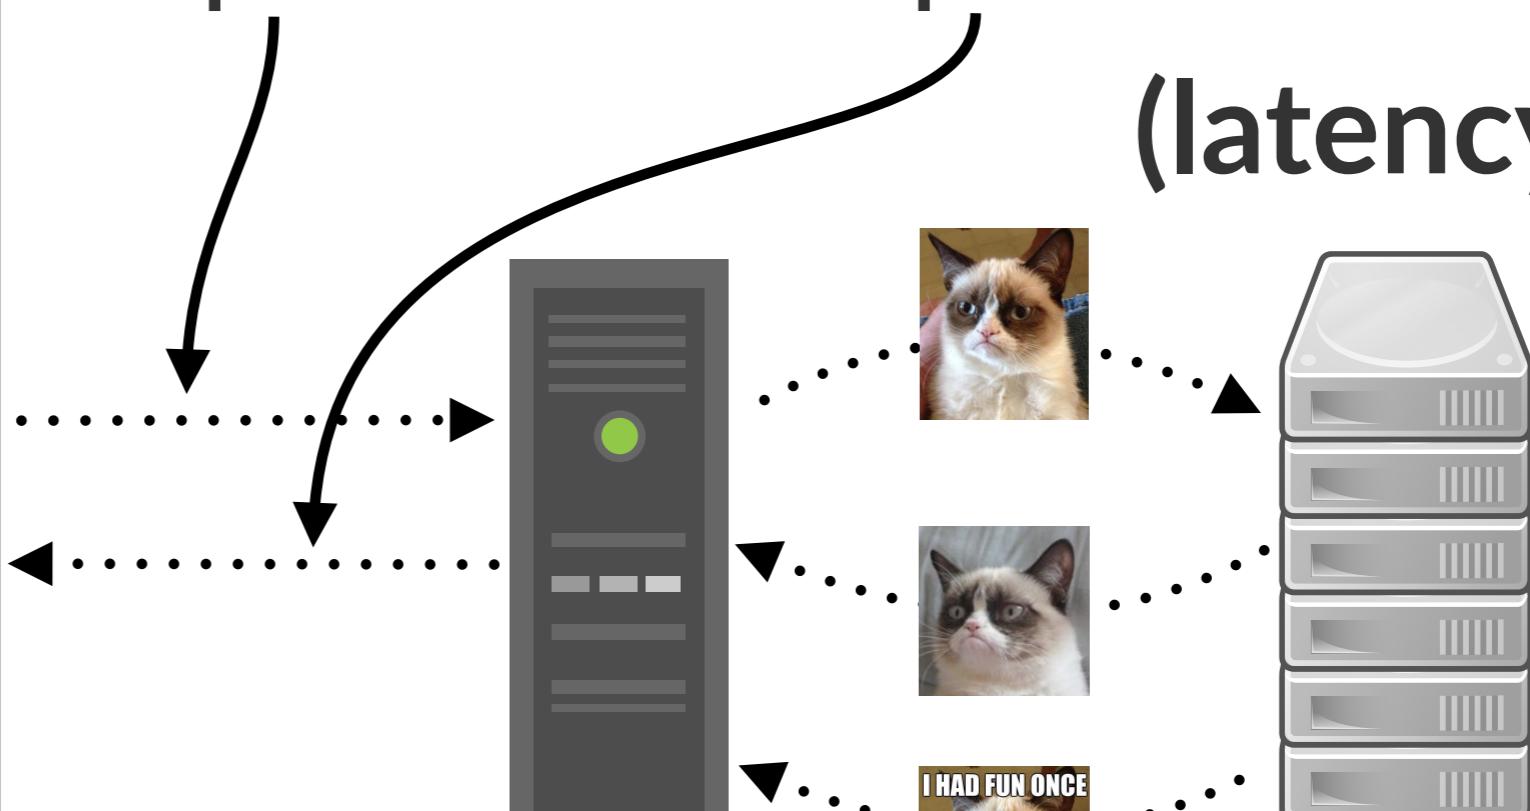
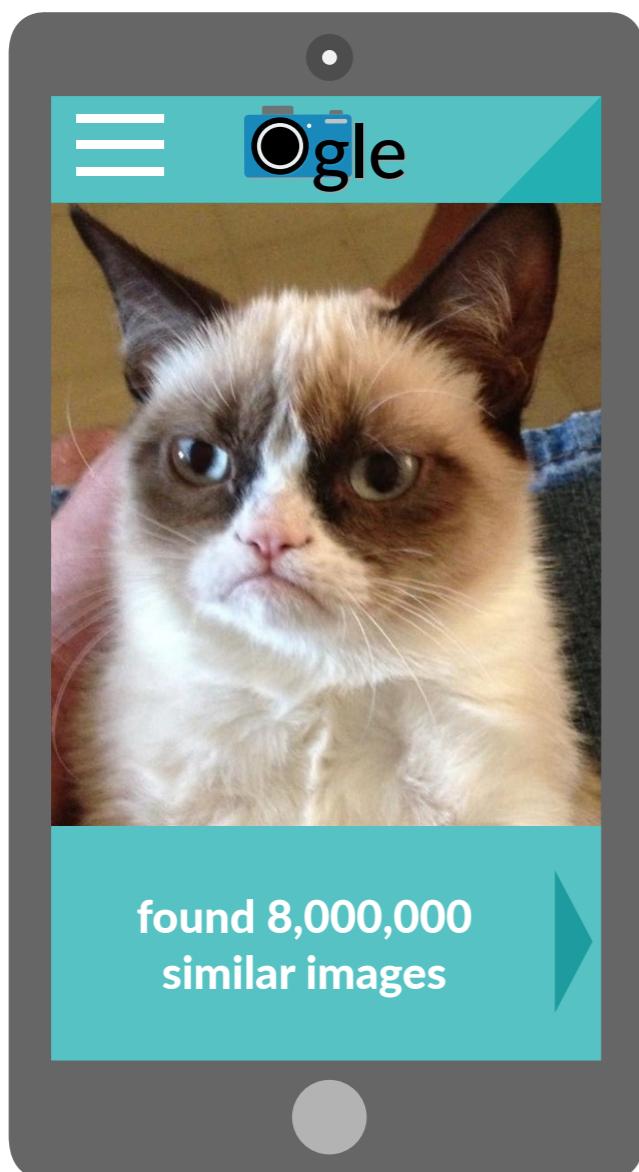
Send results

Find similar pictures



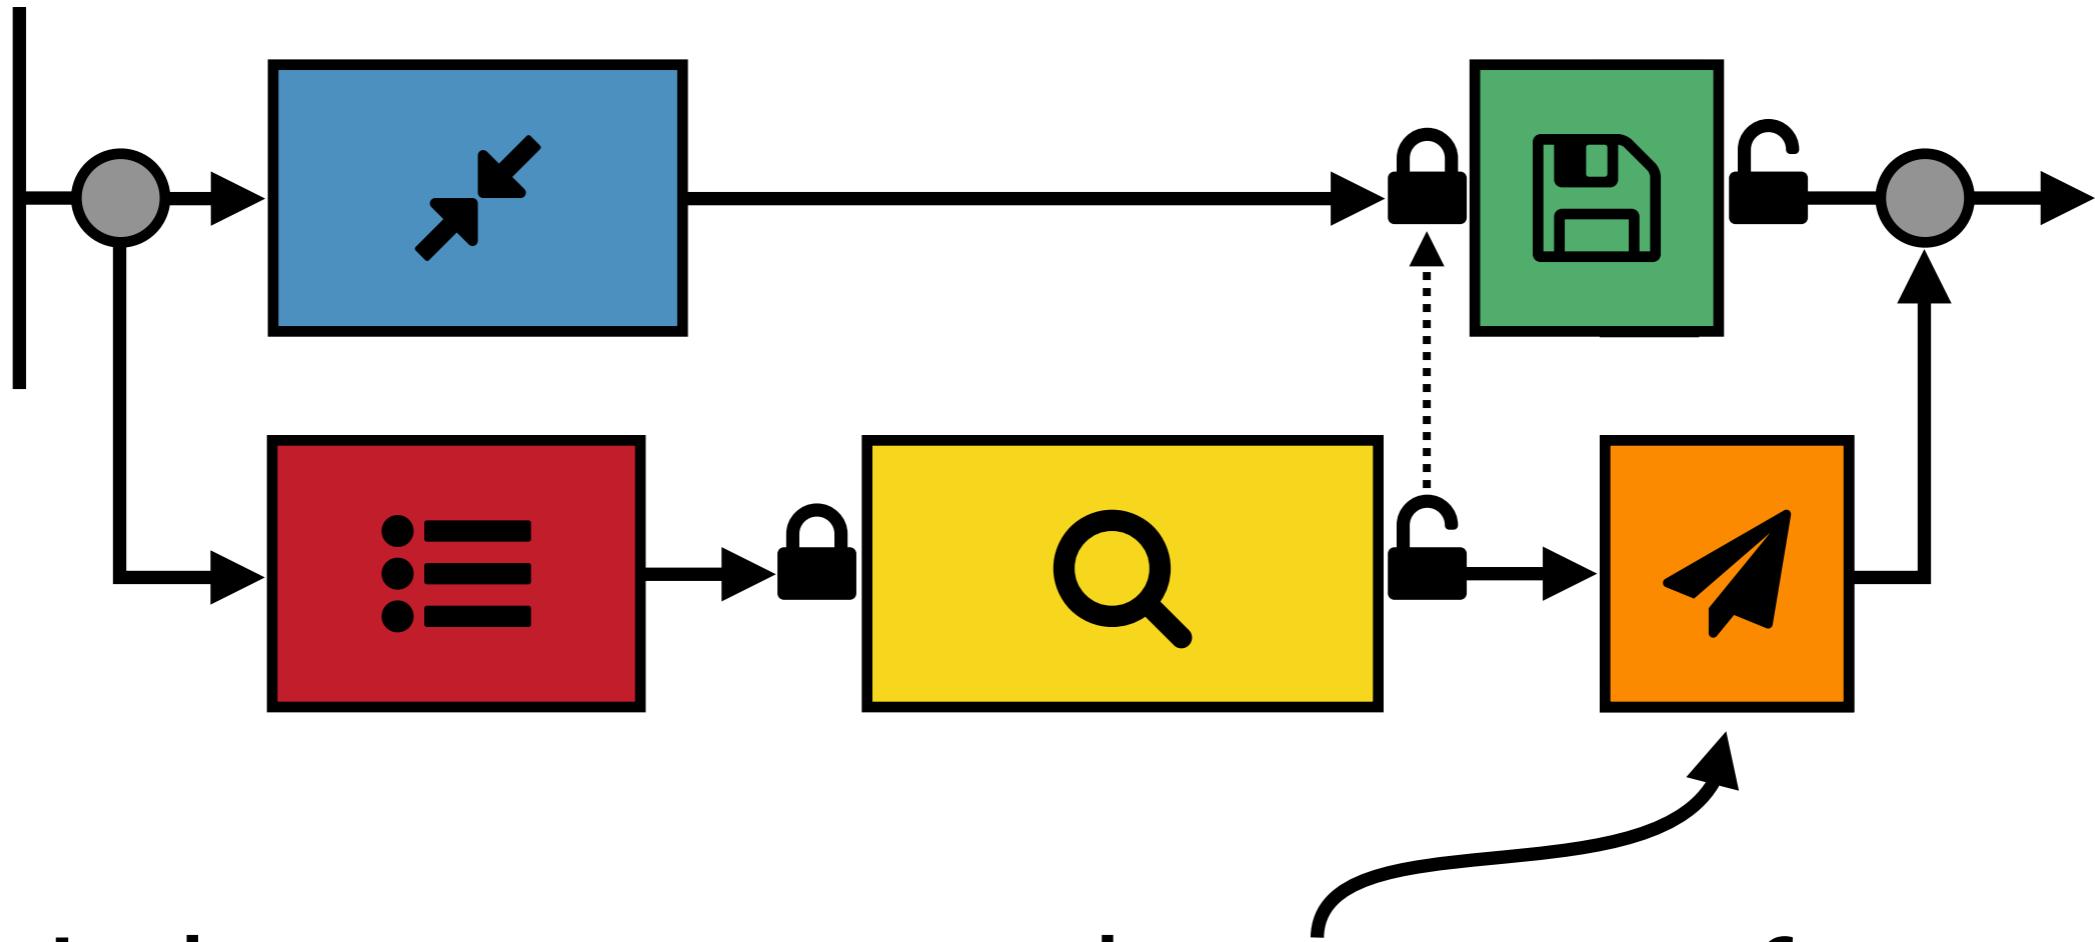
How long between
request and response?

(latency)



How fast do results come back?

Progress Points



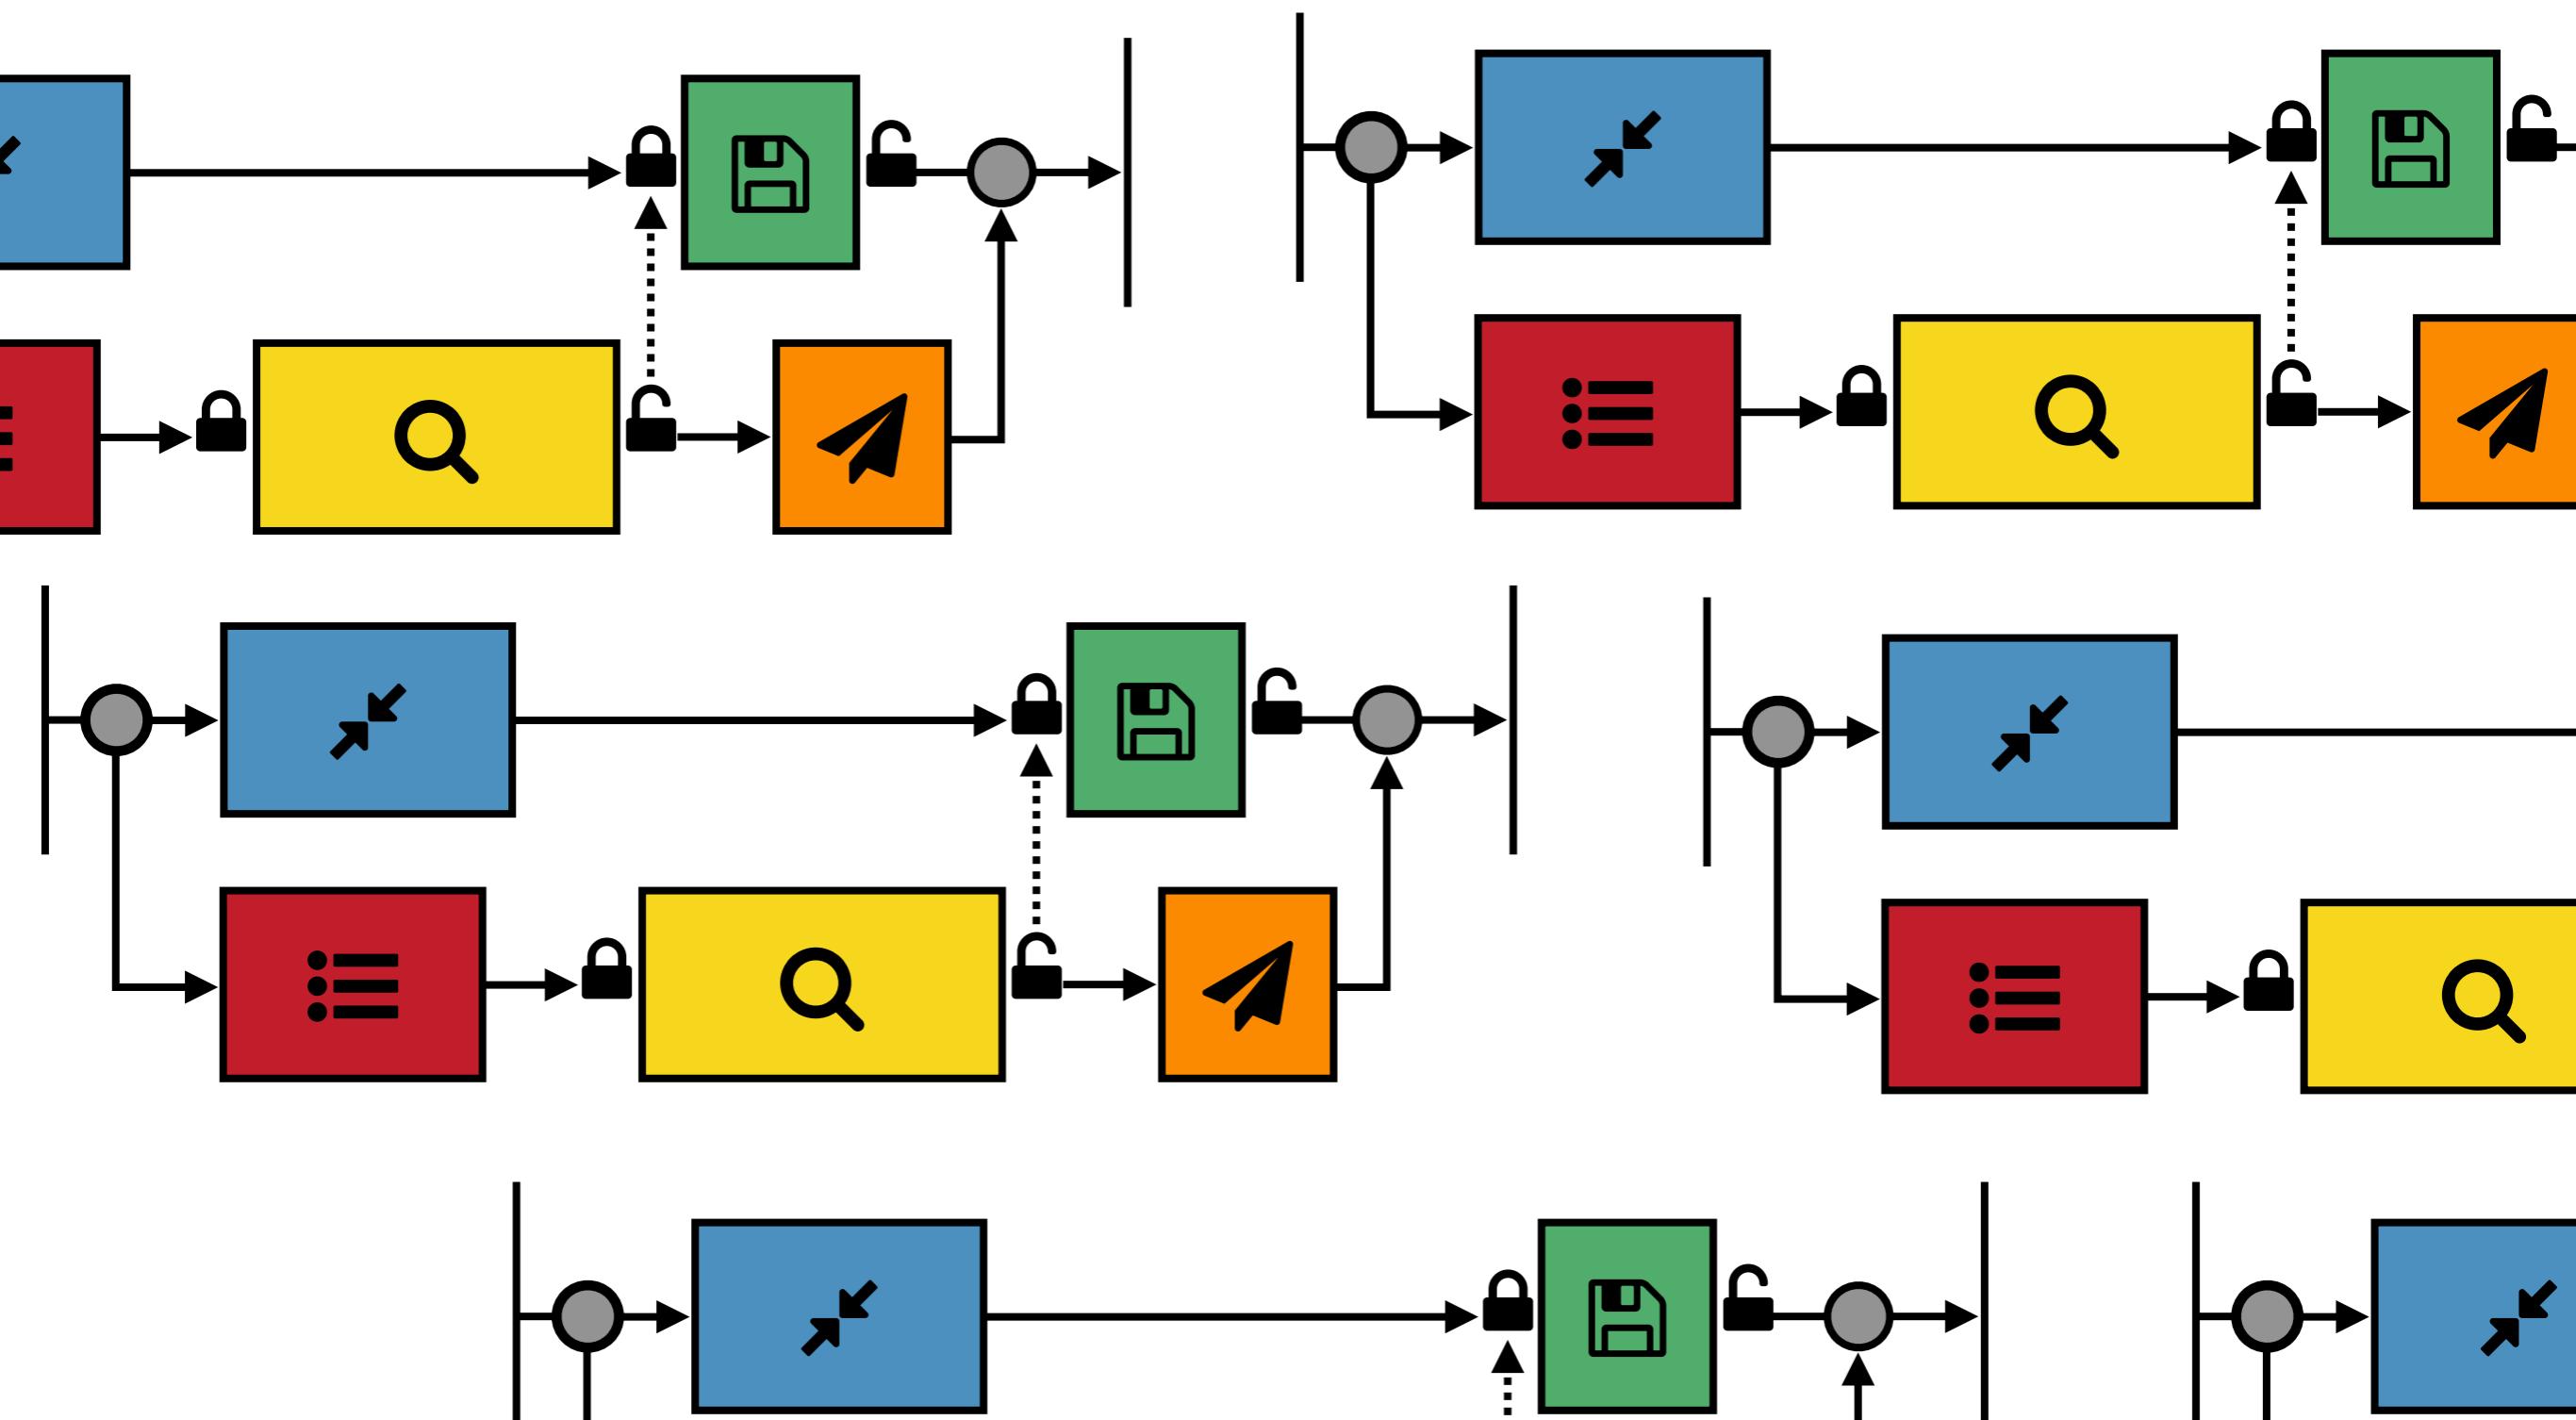
Luke wants to send responses faster.

He marks  as a *progress point*.

Each time this code runs...

Progress Points

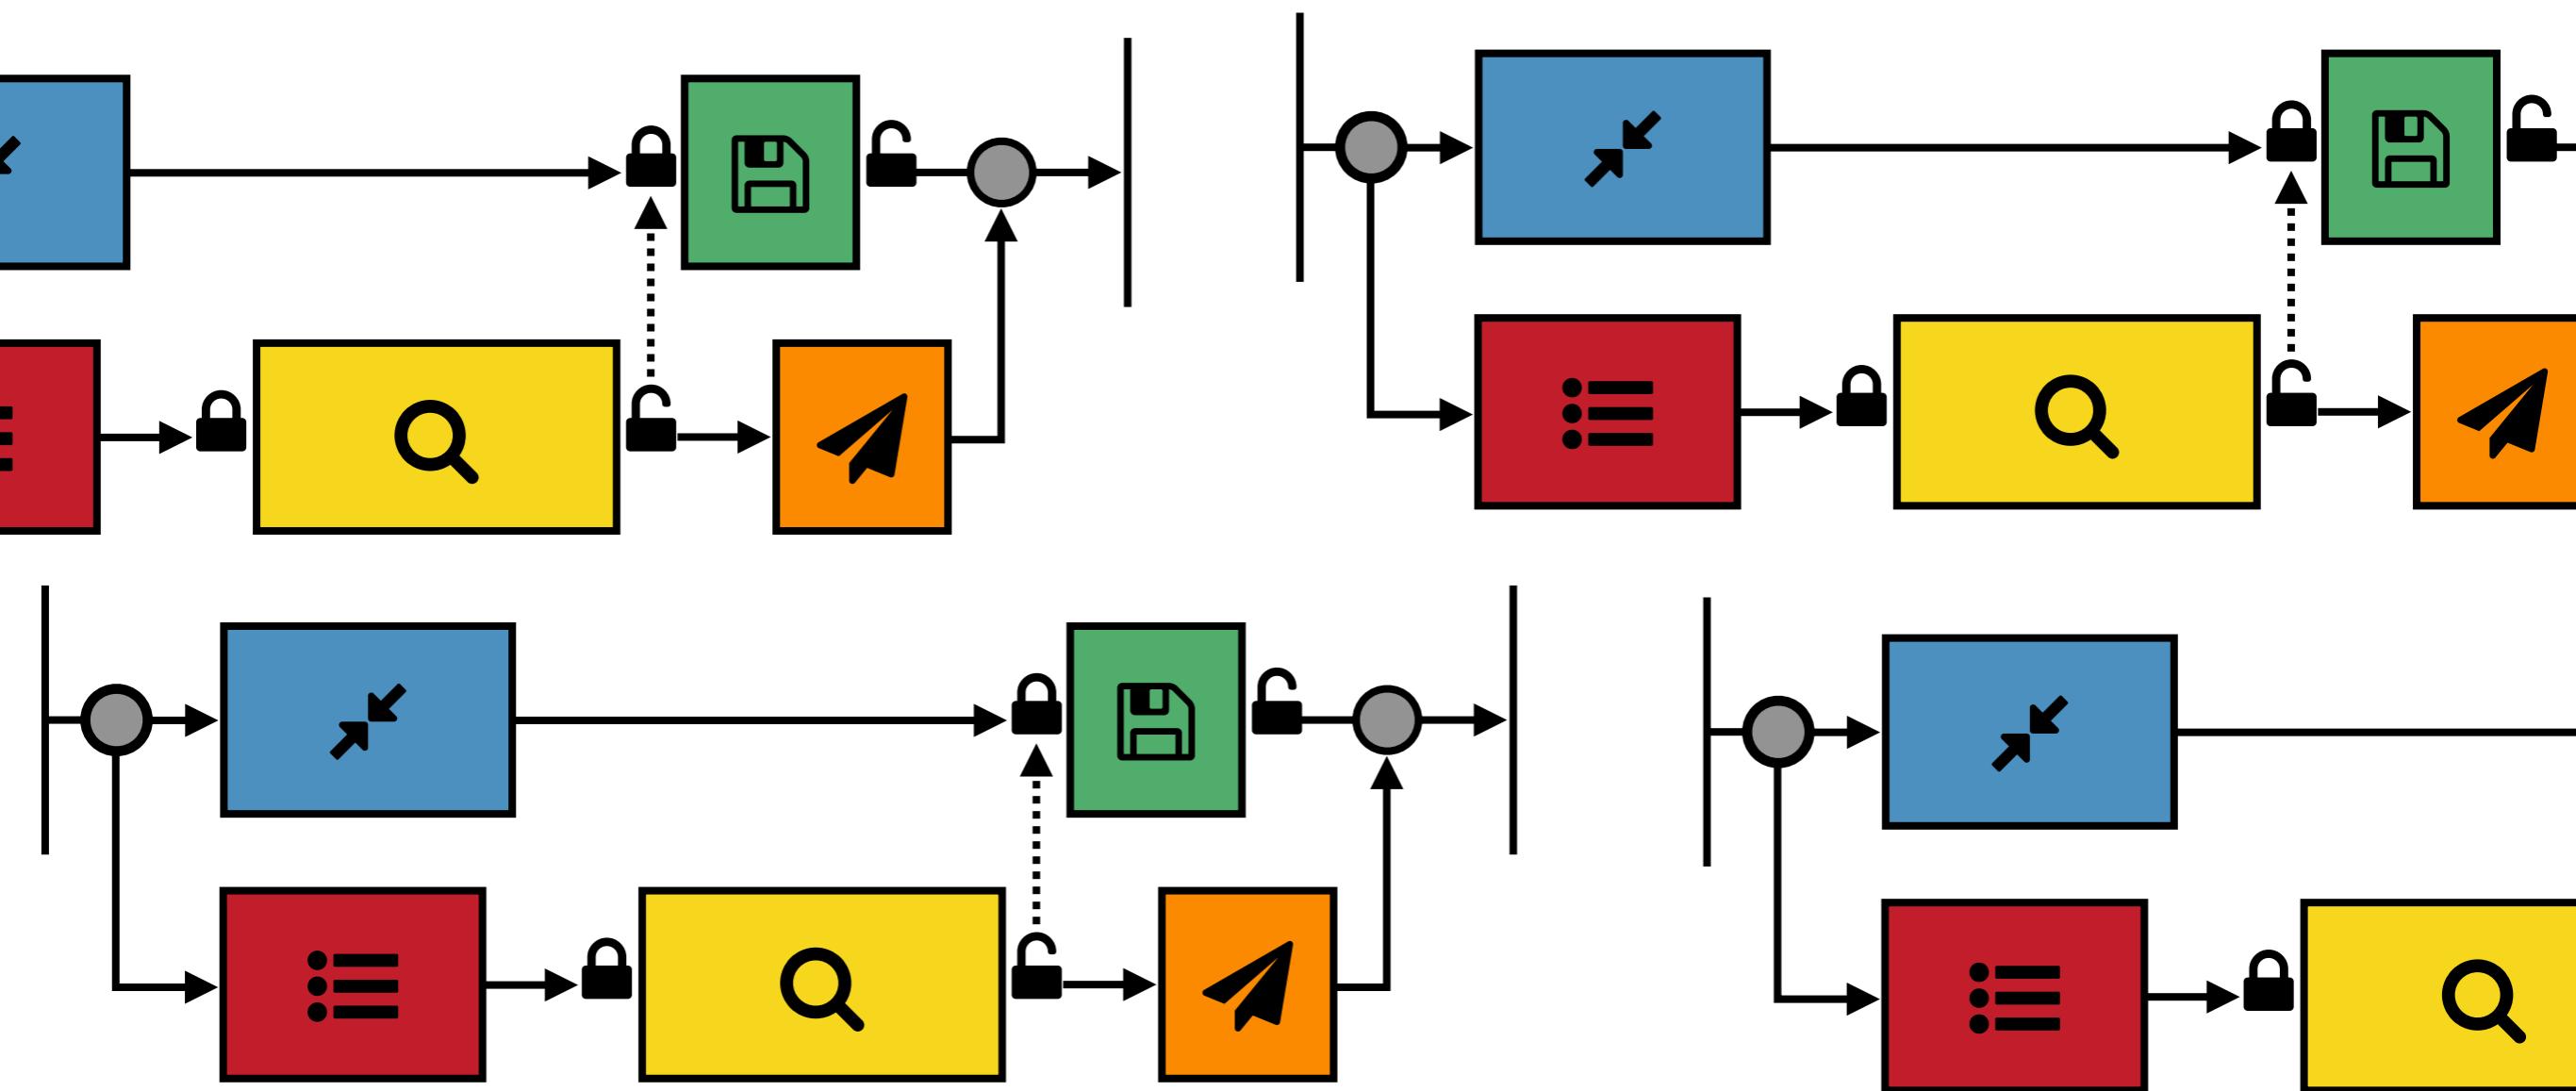
Many requests running for many users.



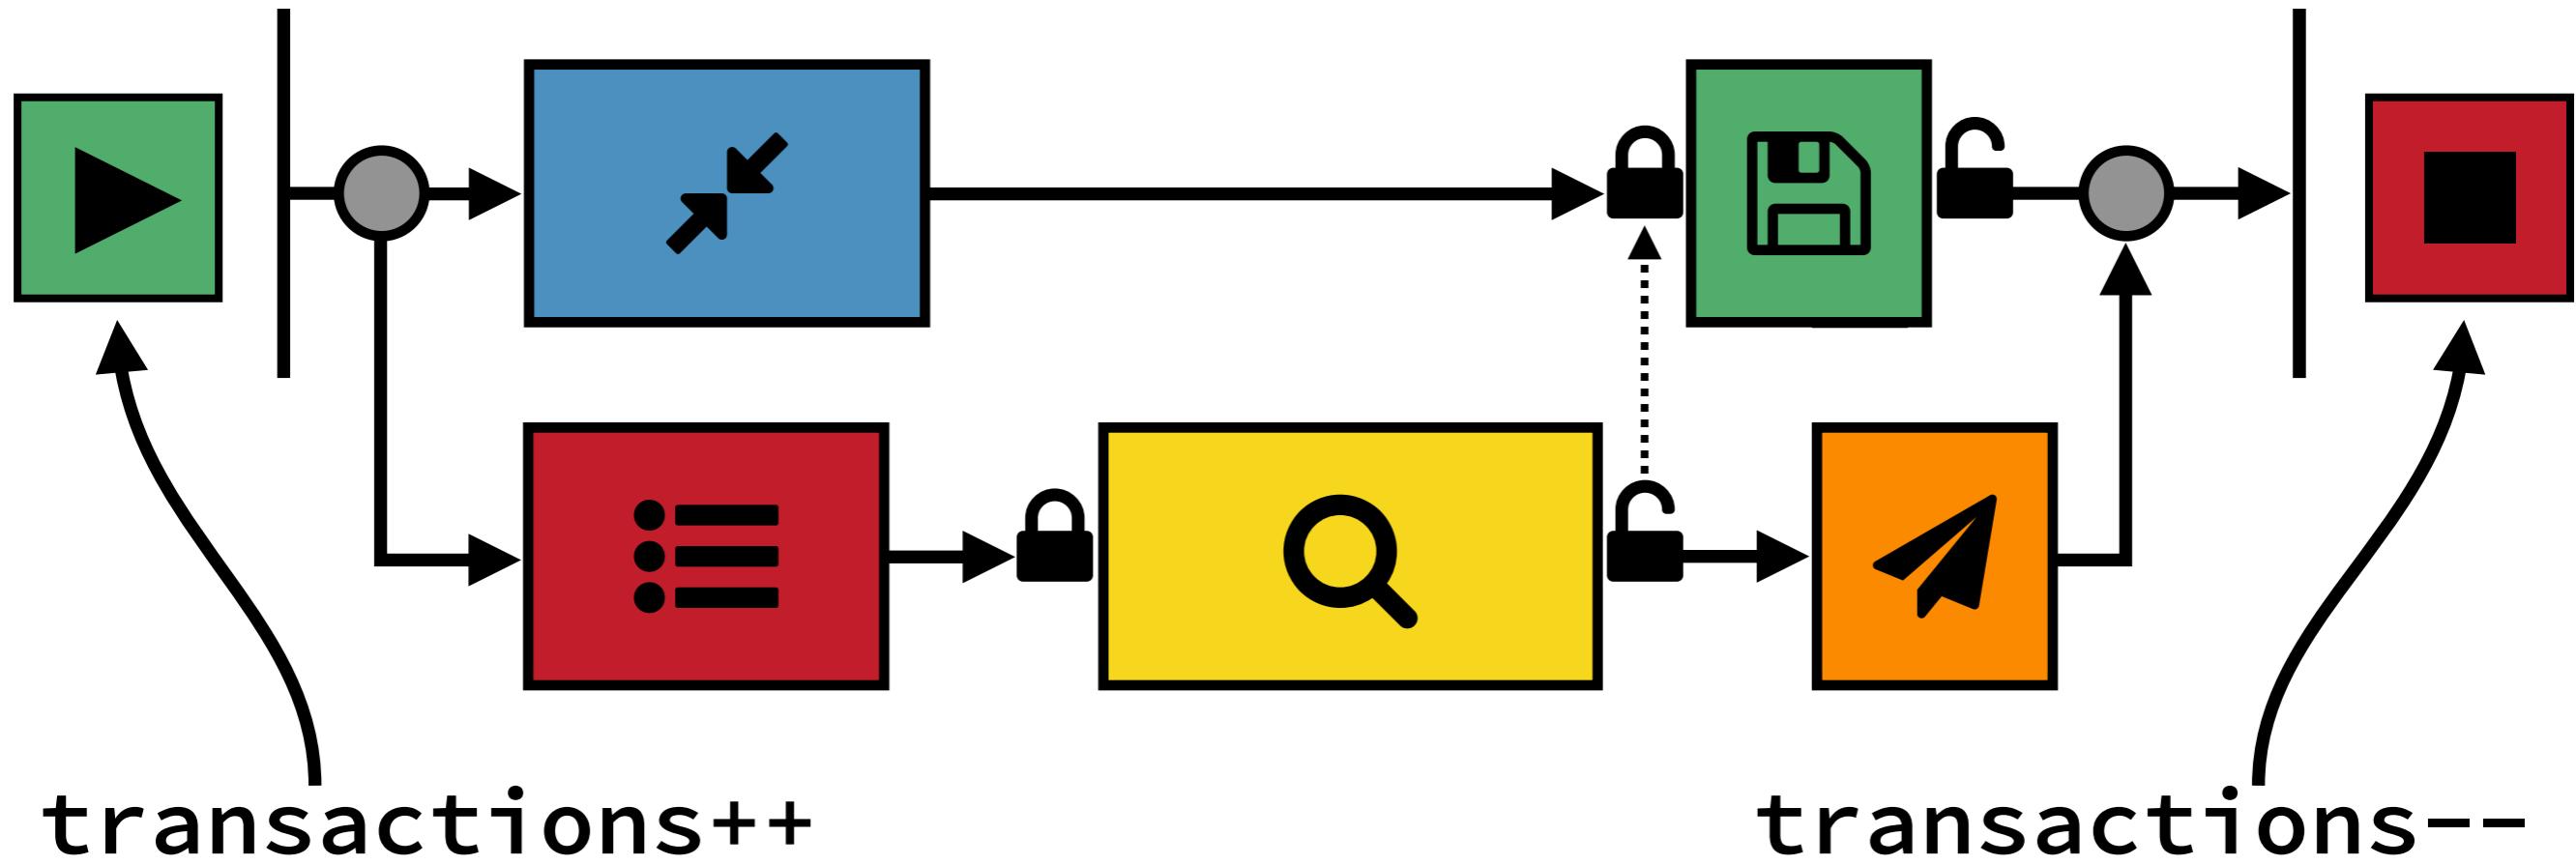
Progress Points

One progress point measures throughput.

If I speed up , how much faster do I run ?



Progress Points



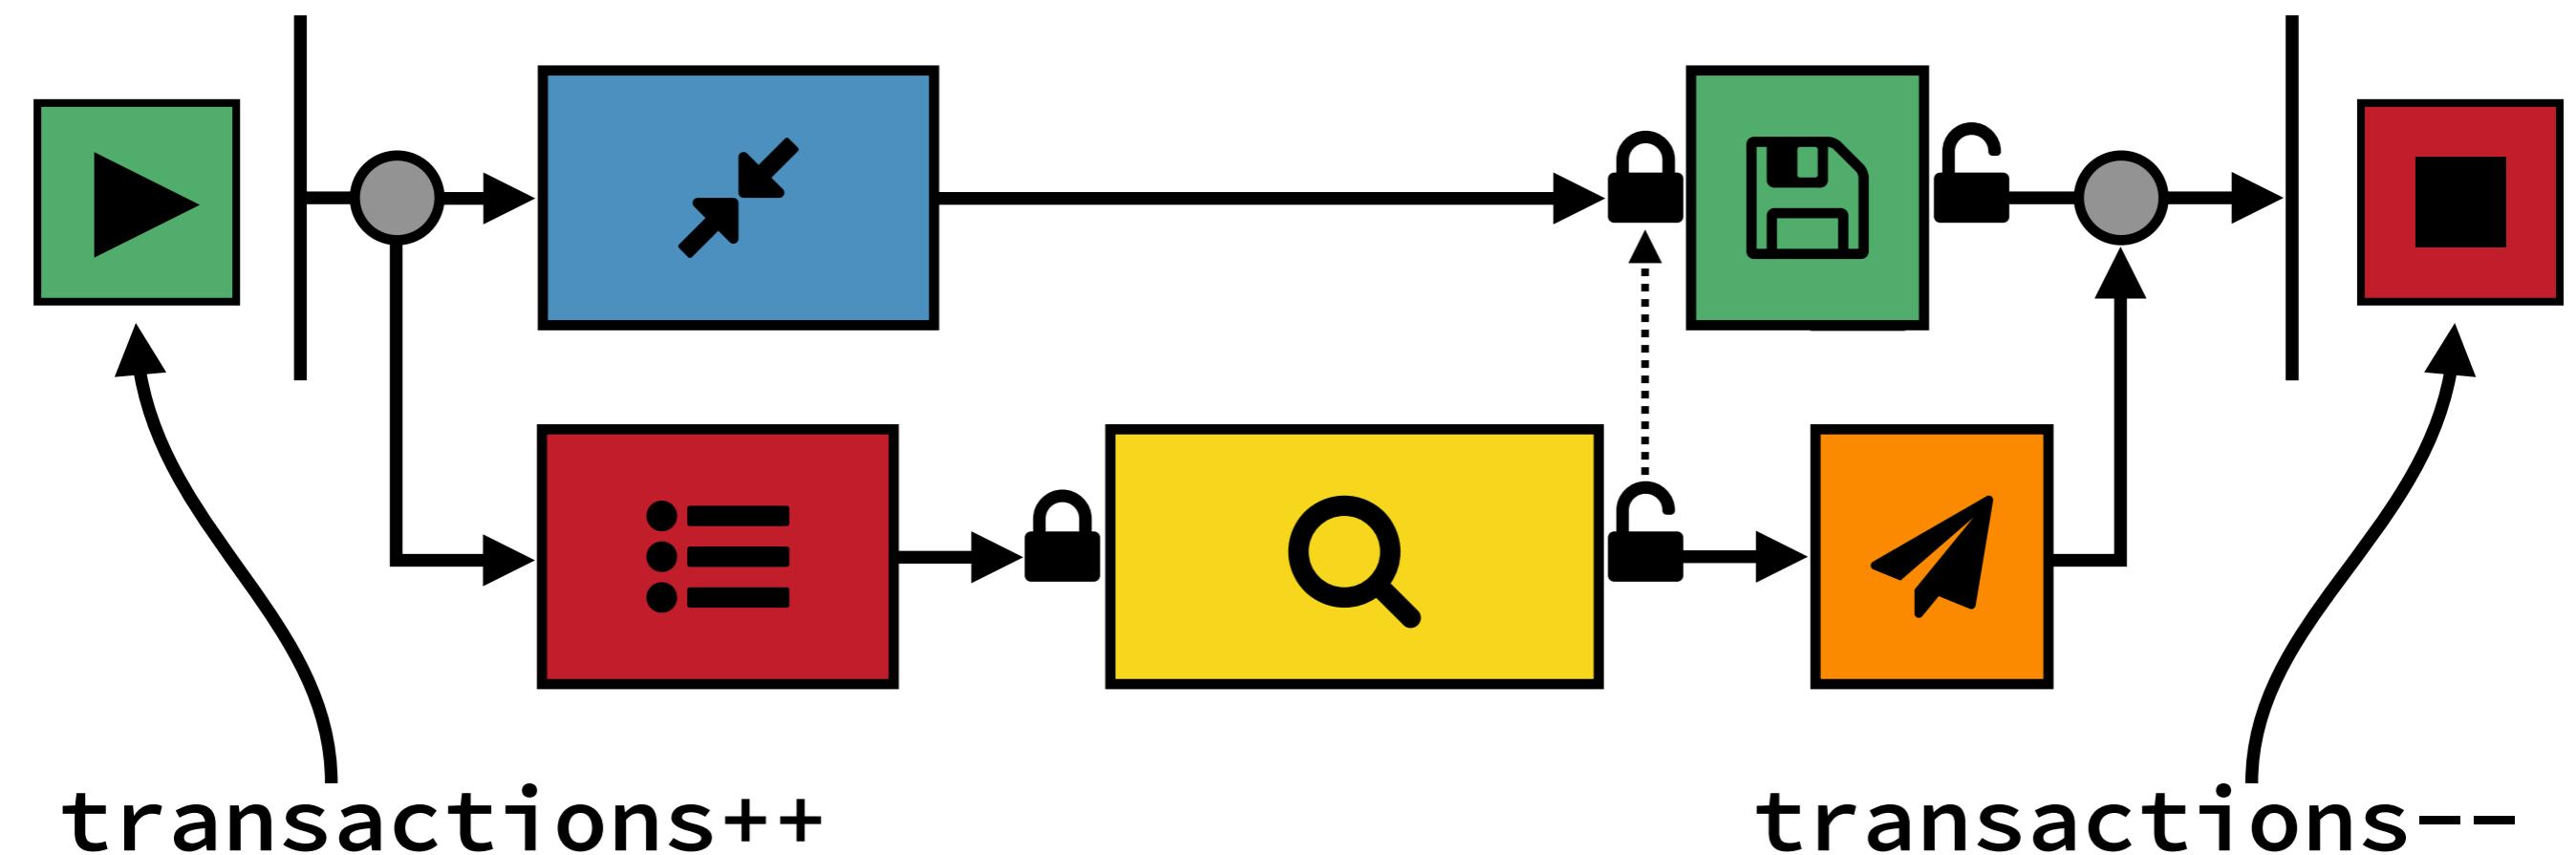
Luke wants to minimize response time.

He adds *latency progress points*.

Progress Points

Little's Law: $W = L / \lambda$

latency = transactions / throughput



Luke wants to minimize response time.

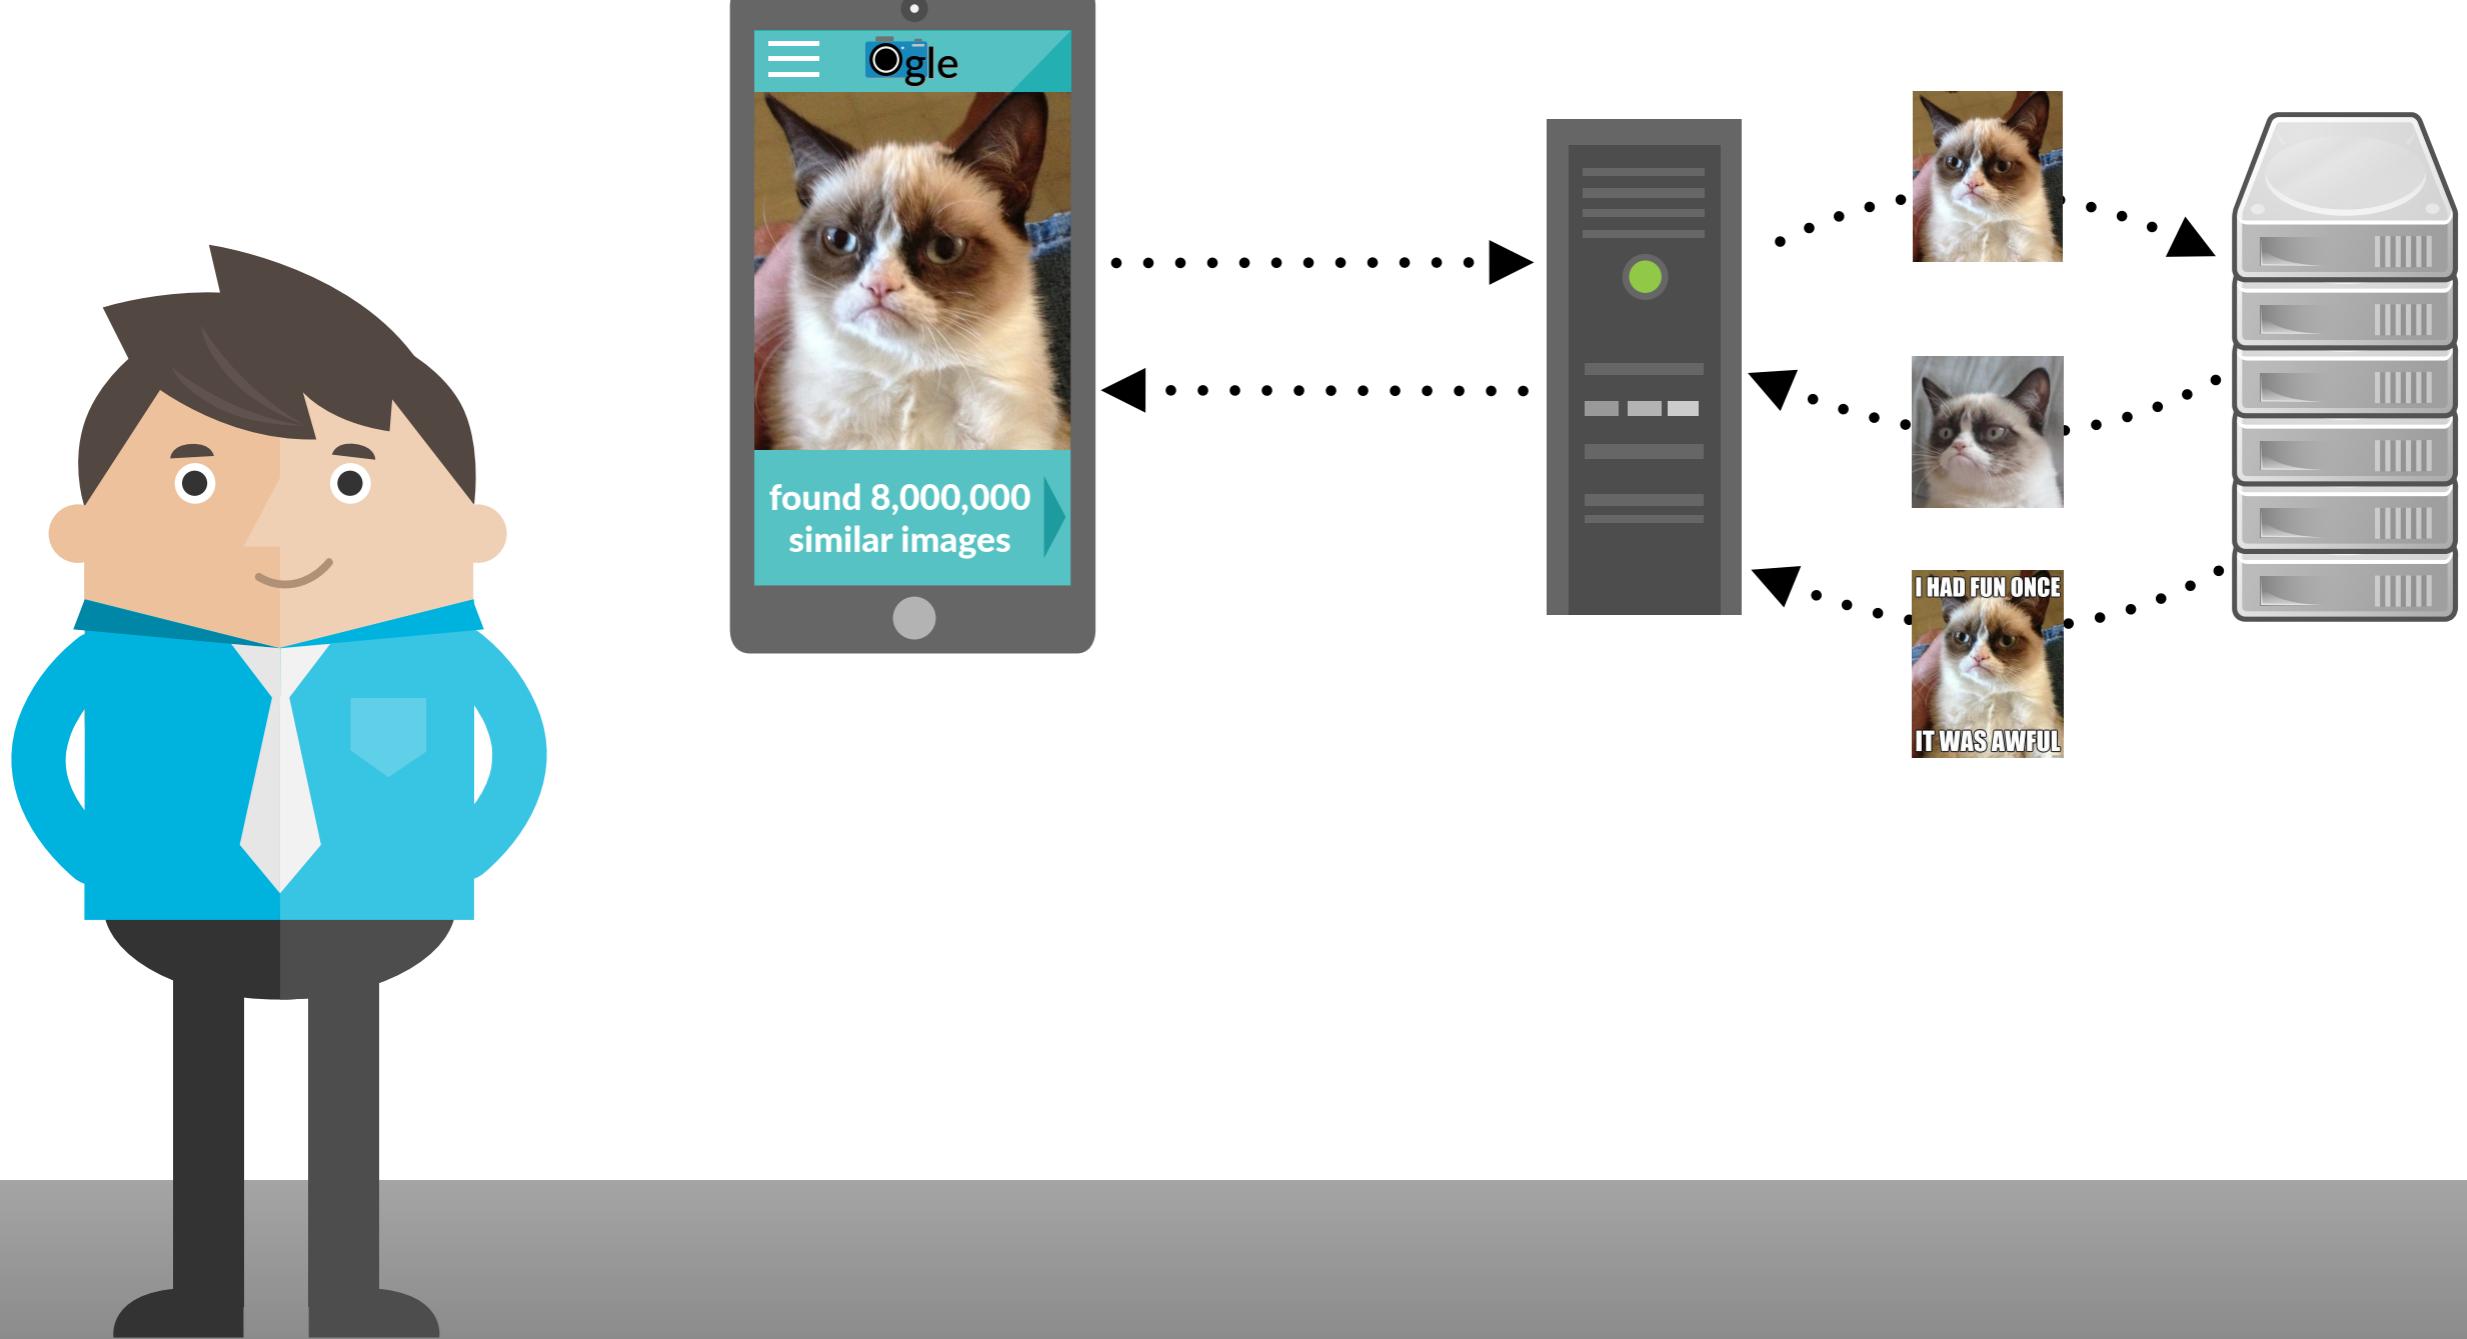
Coz: a Causal Profiler for Linux

(ships with Debian / Ubuntu)

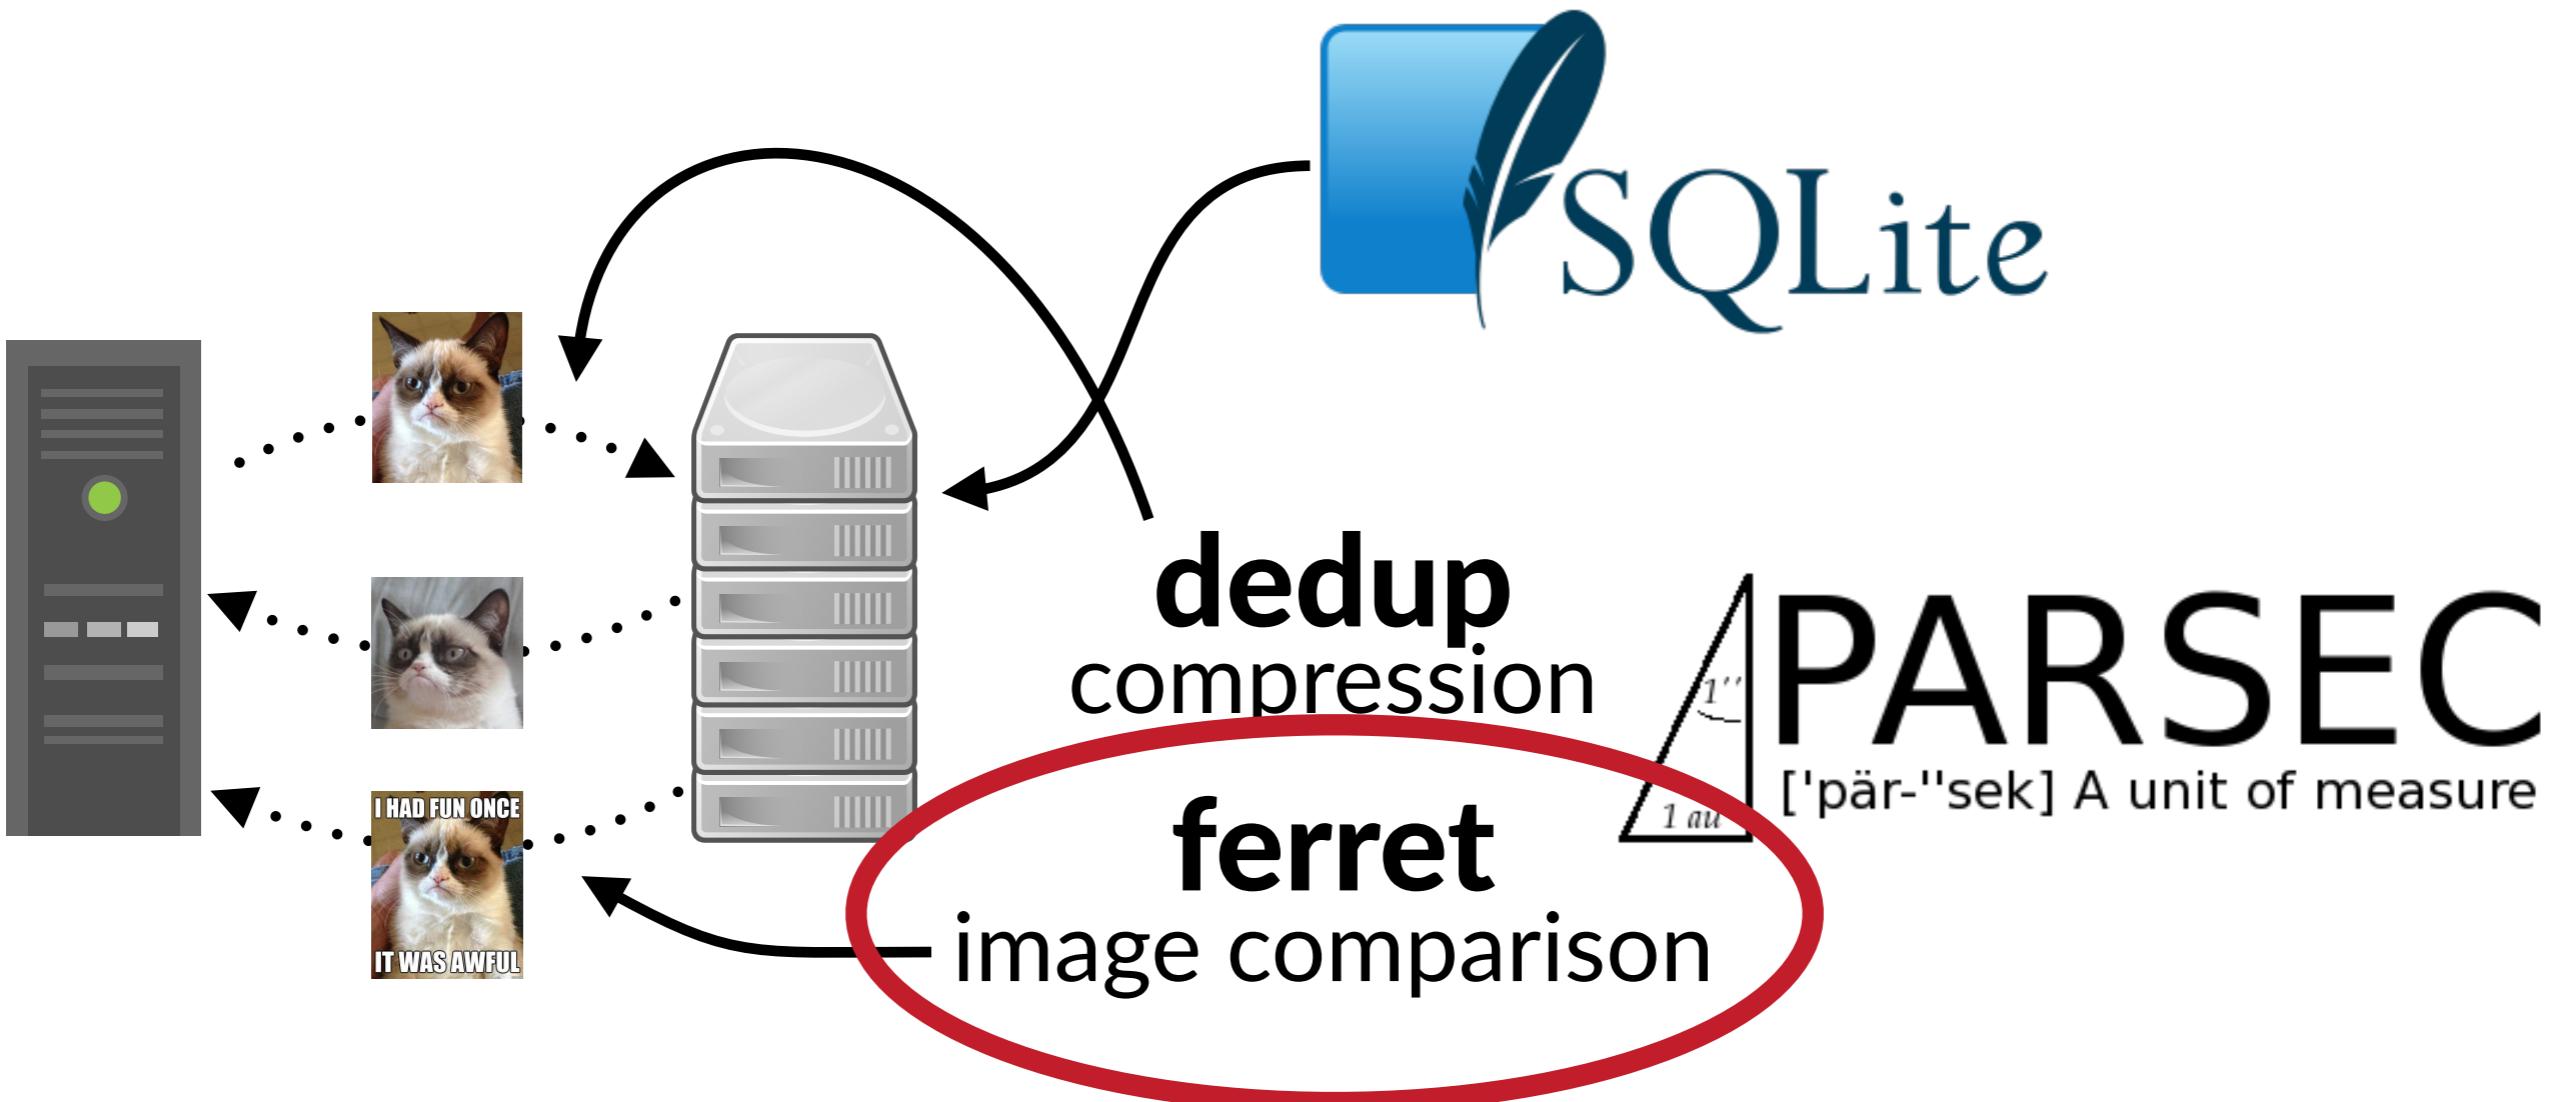
```
> sudo apt install coz-profiler  
> coz run --- ./some_program args
```

Random performance experiments

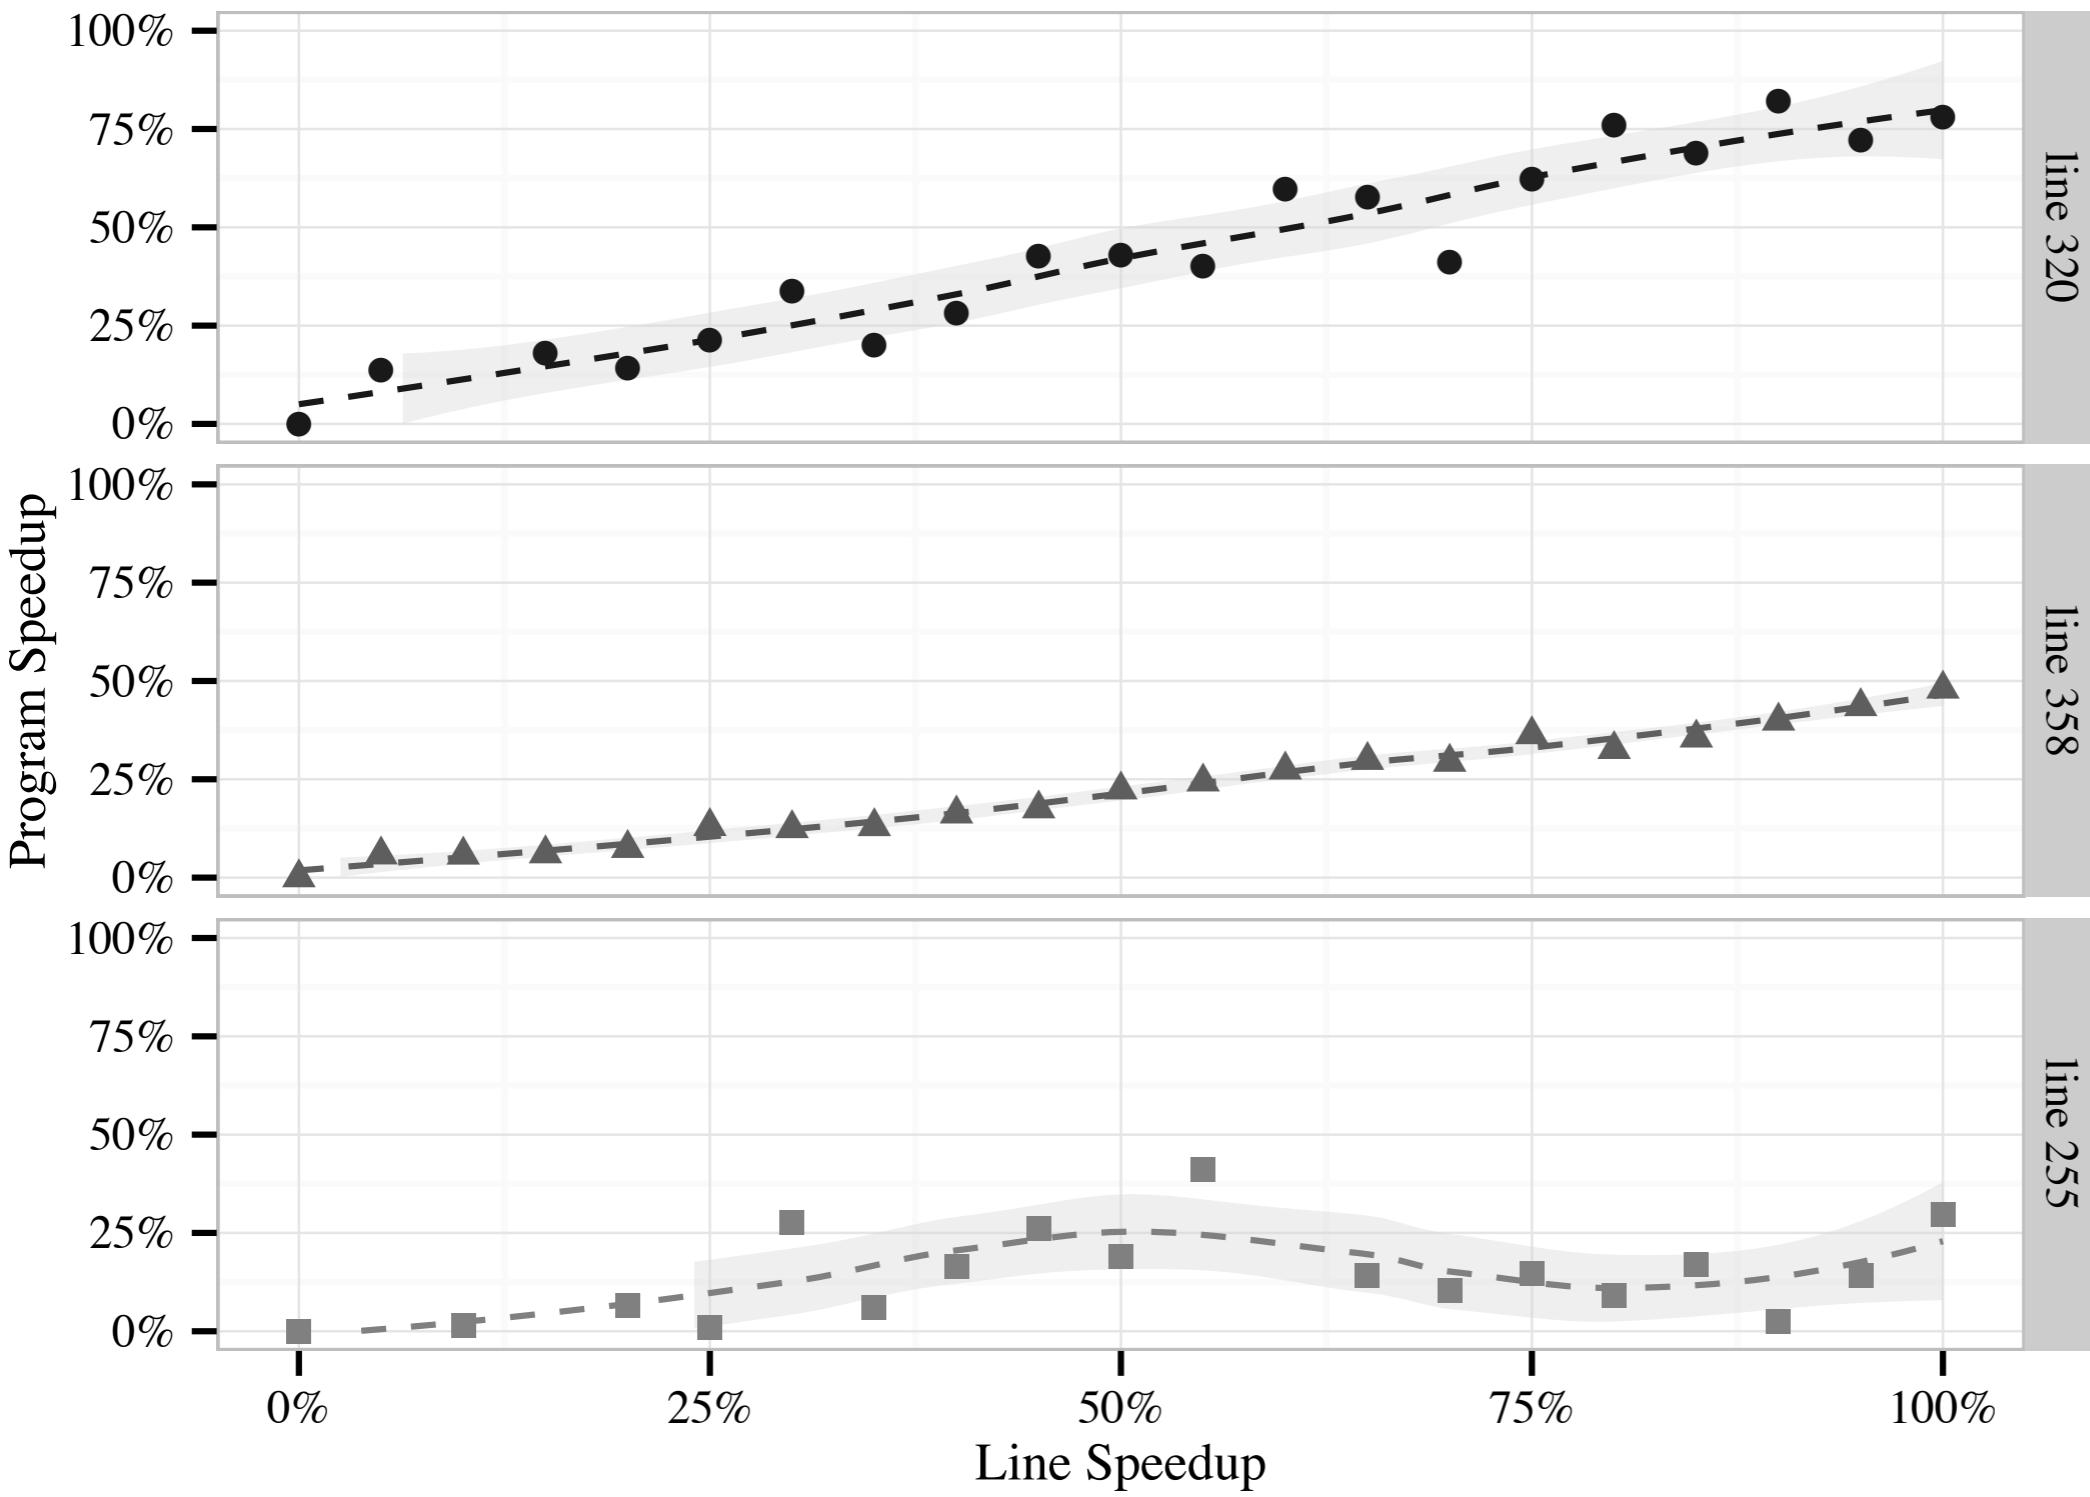
Using Causal Profiling on Ogle



Using Causal Profiling on Ogle



Ferret

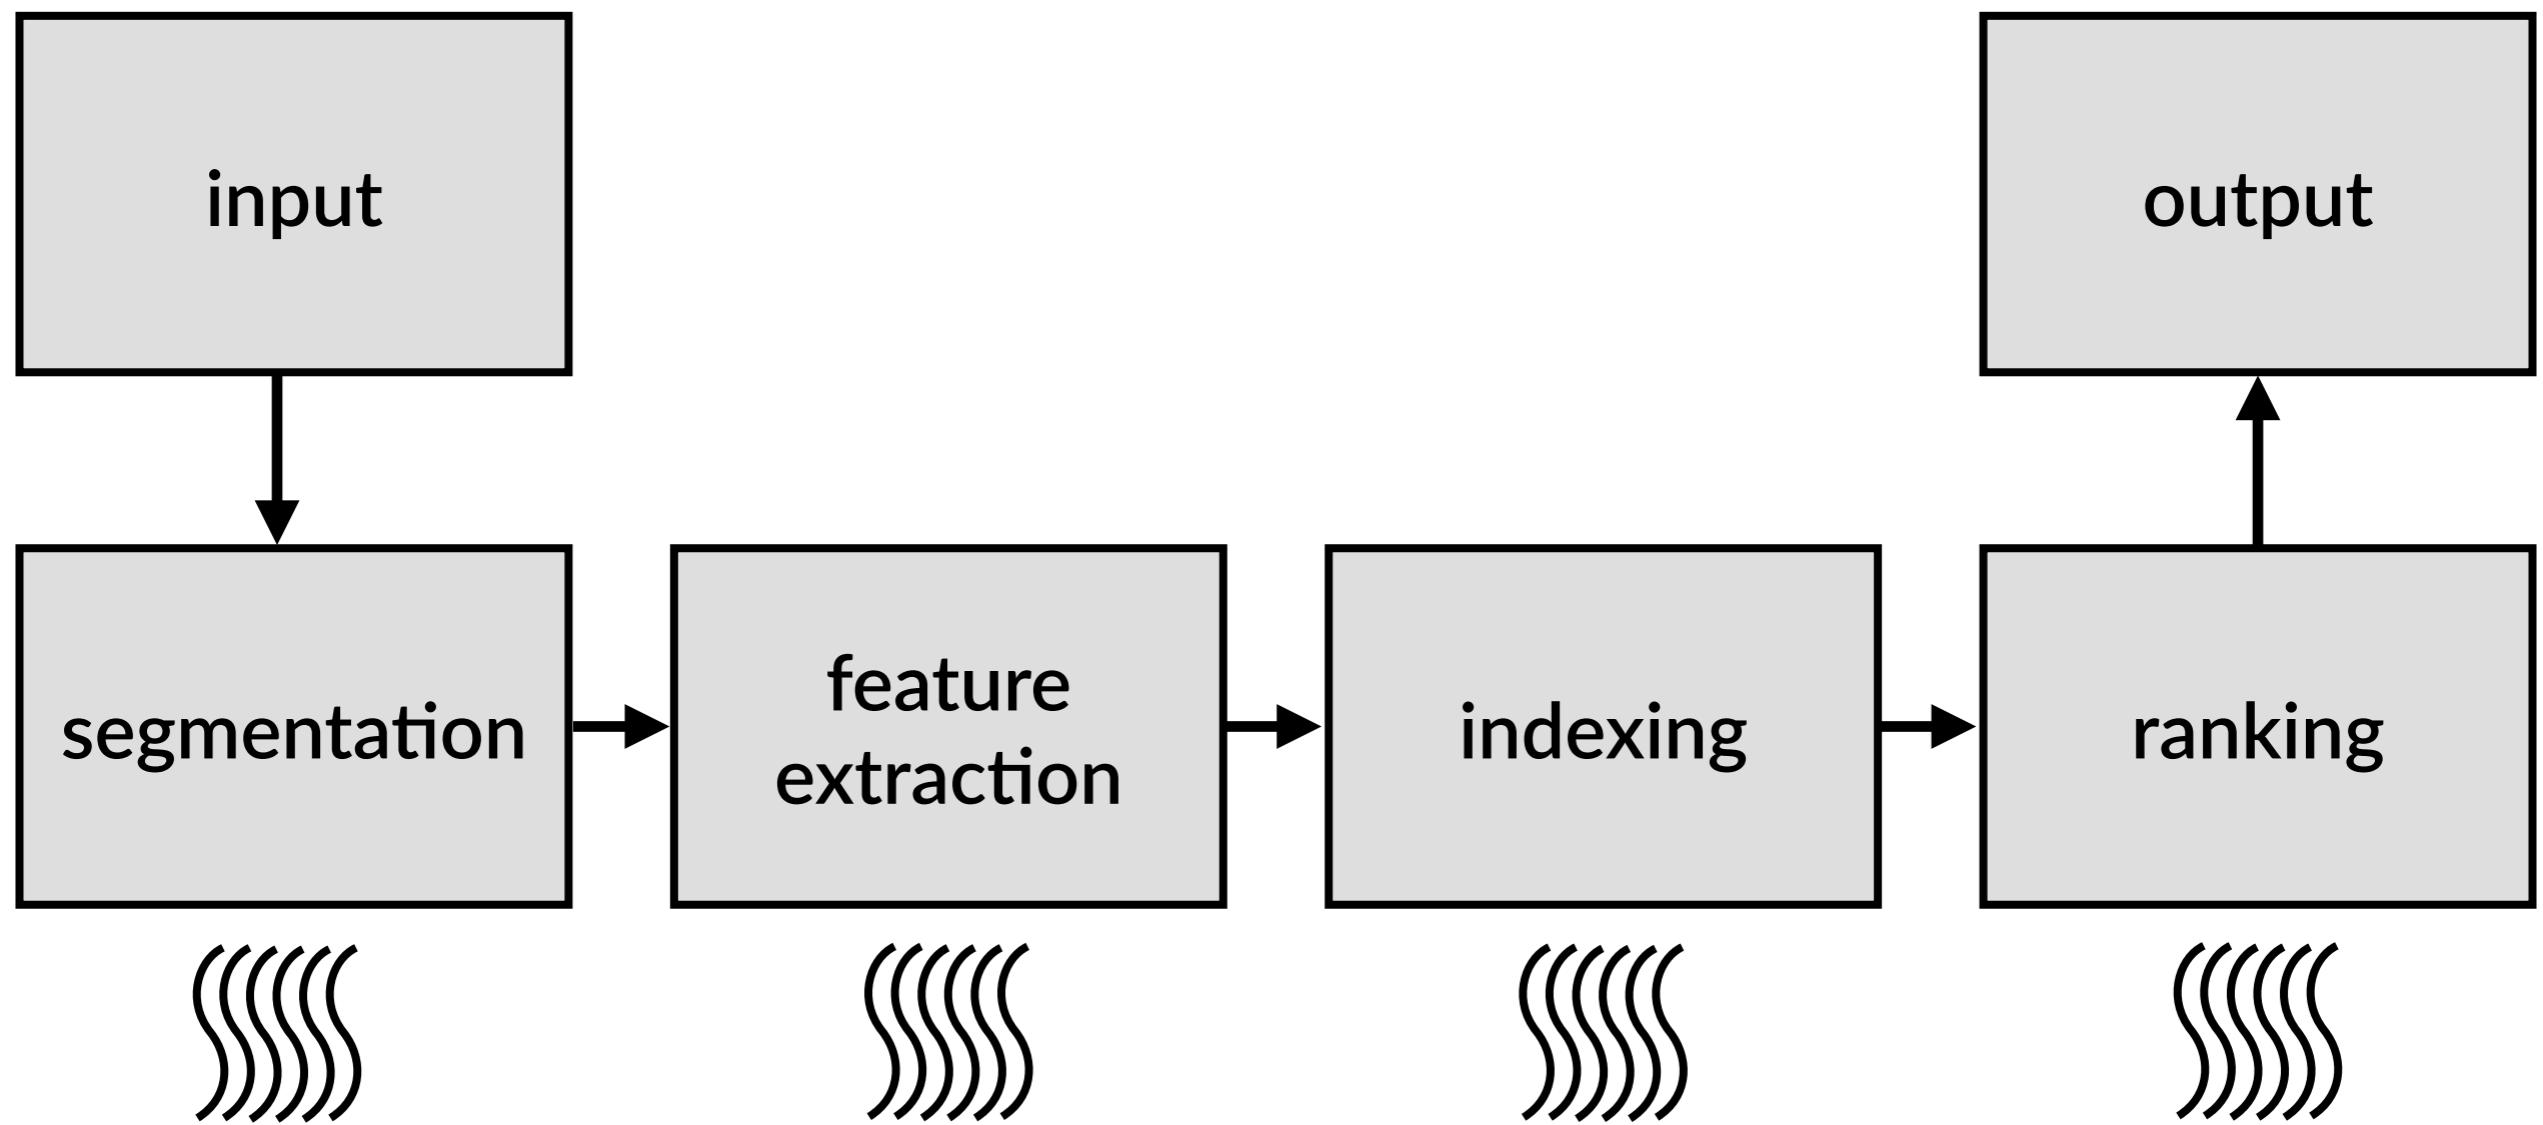


ranking

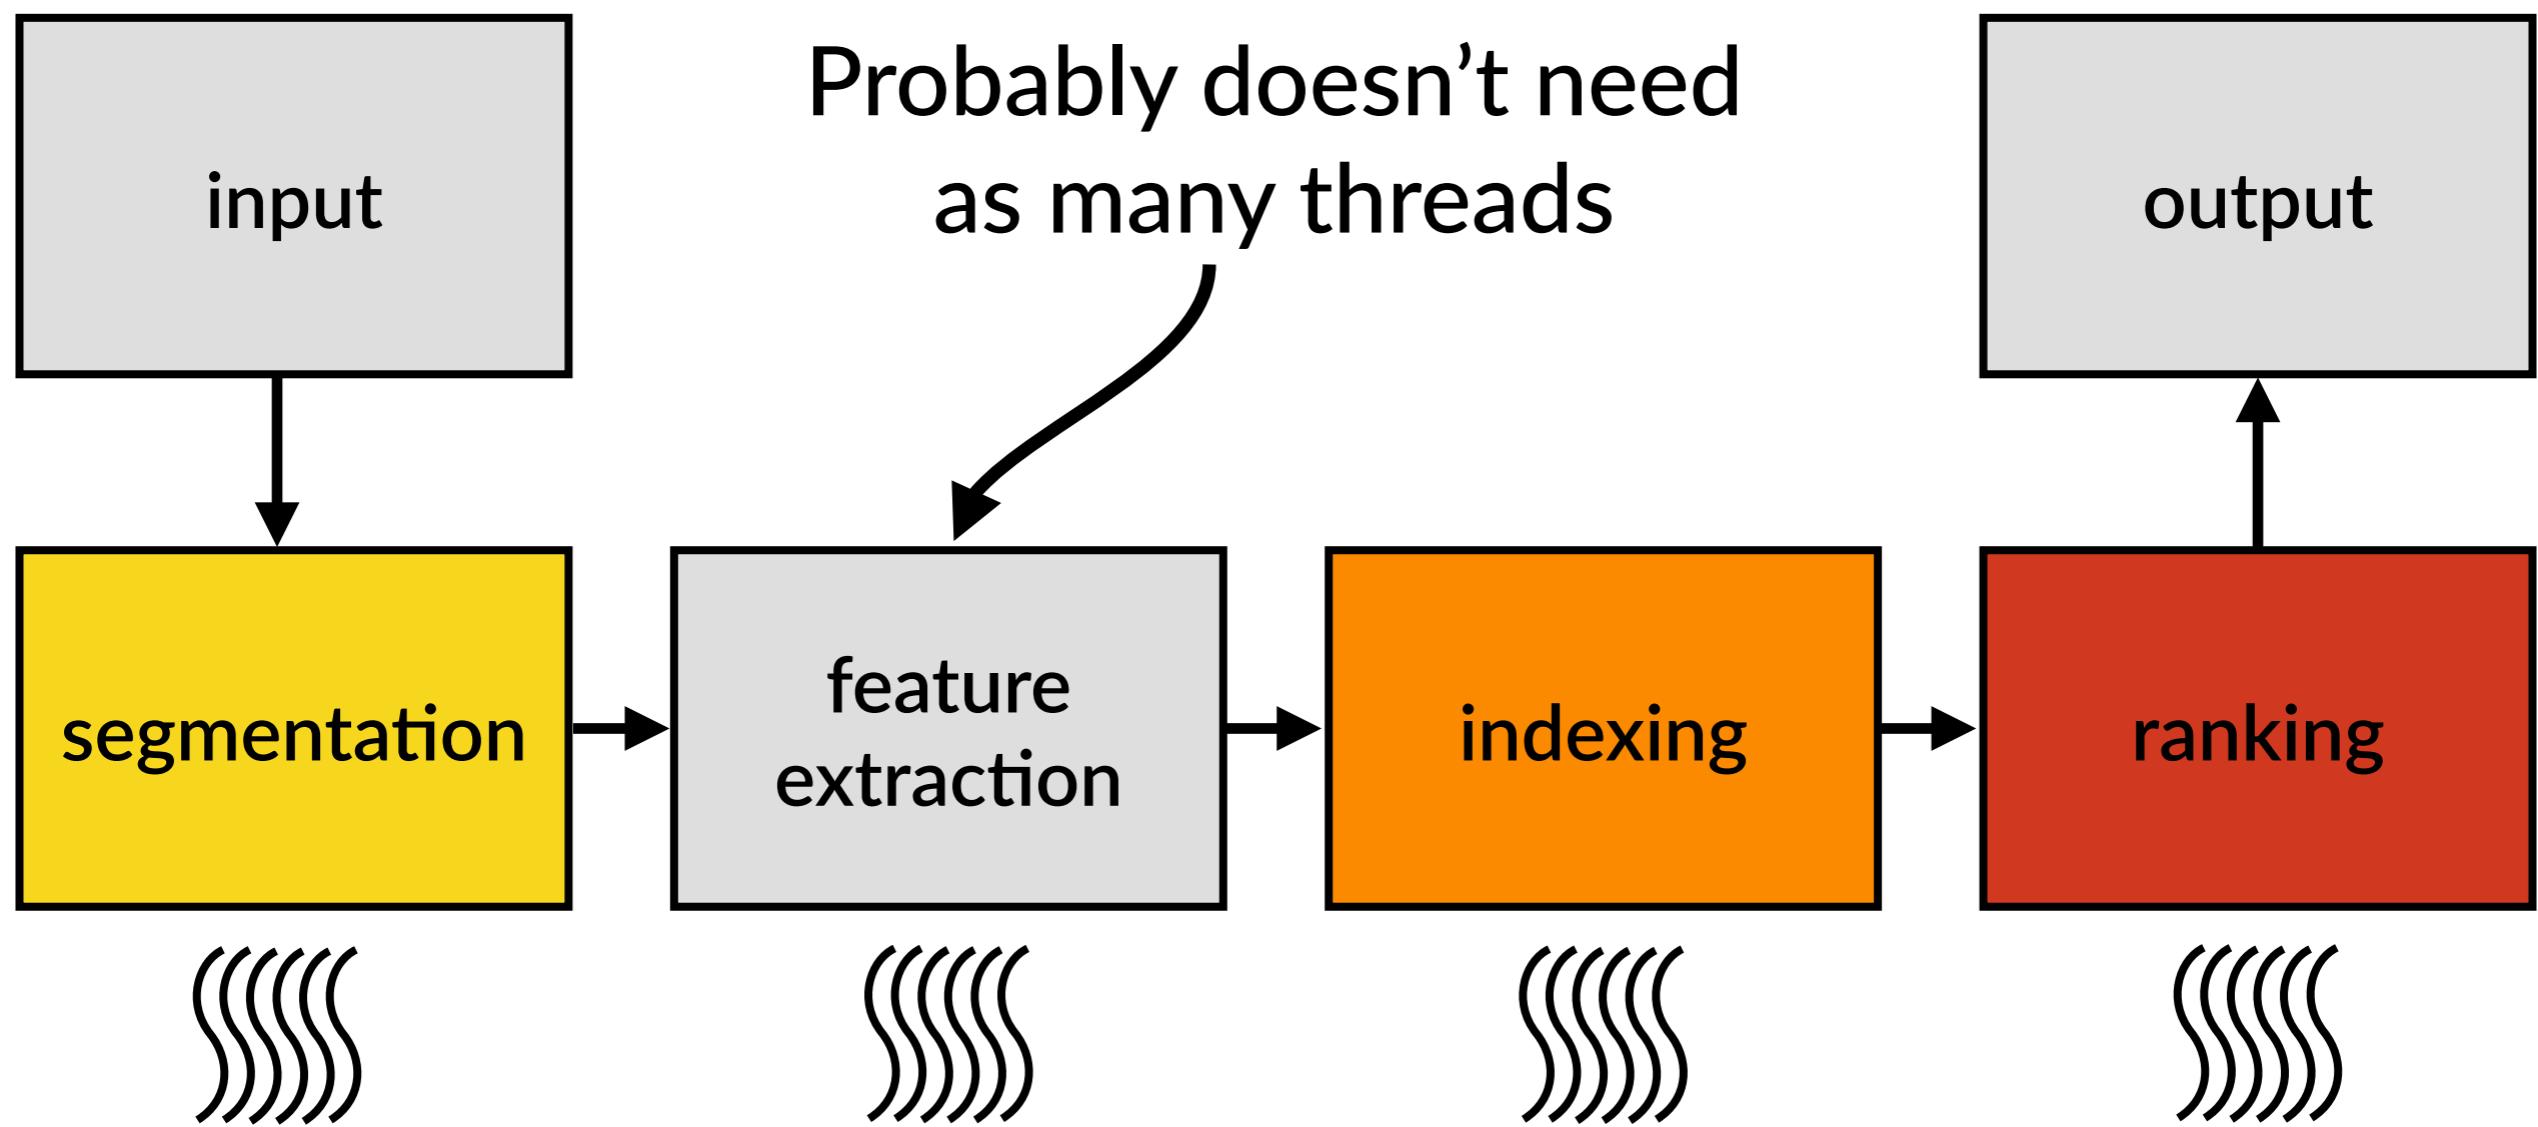
indexing

segmentation

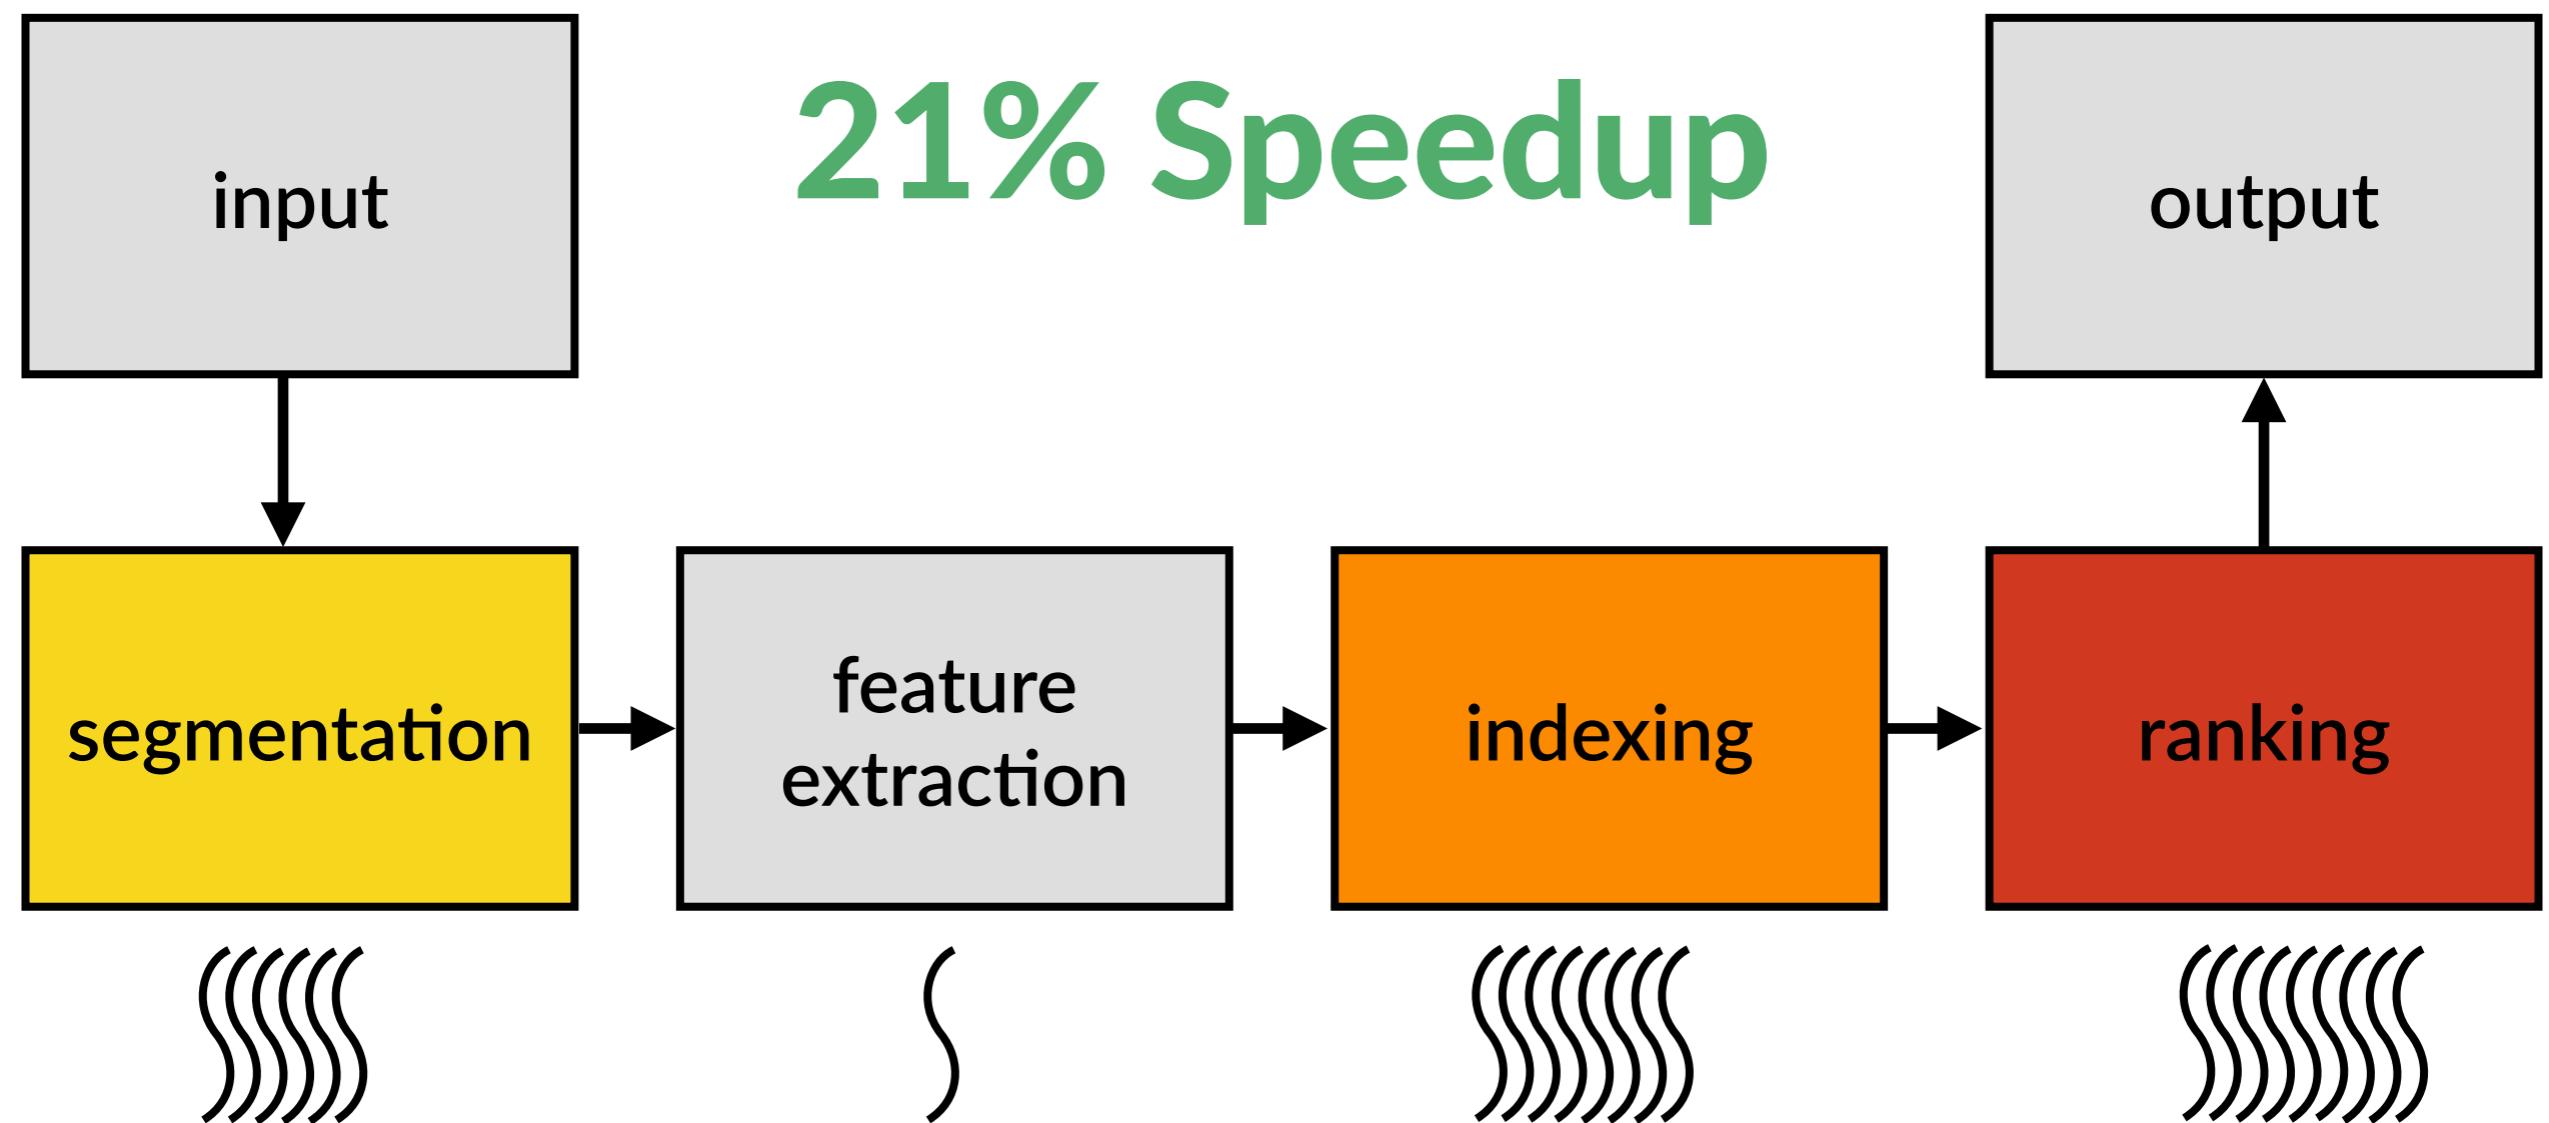
Ferret



Ferret



Ferret



What did Coz predict?

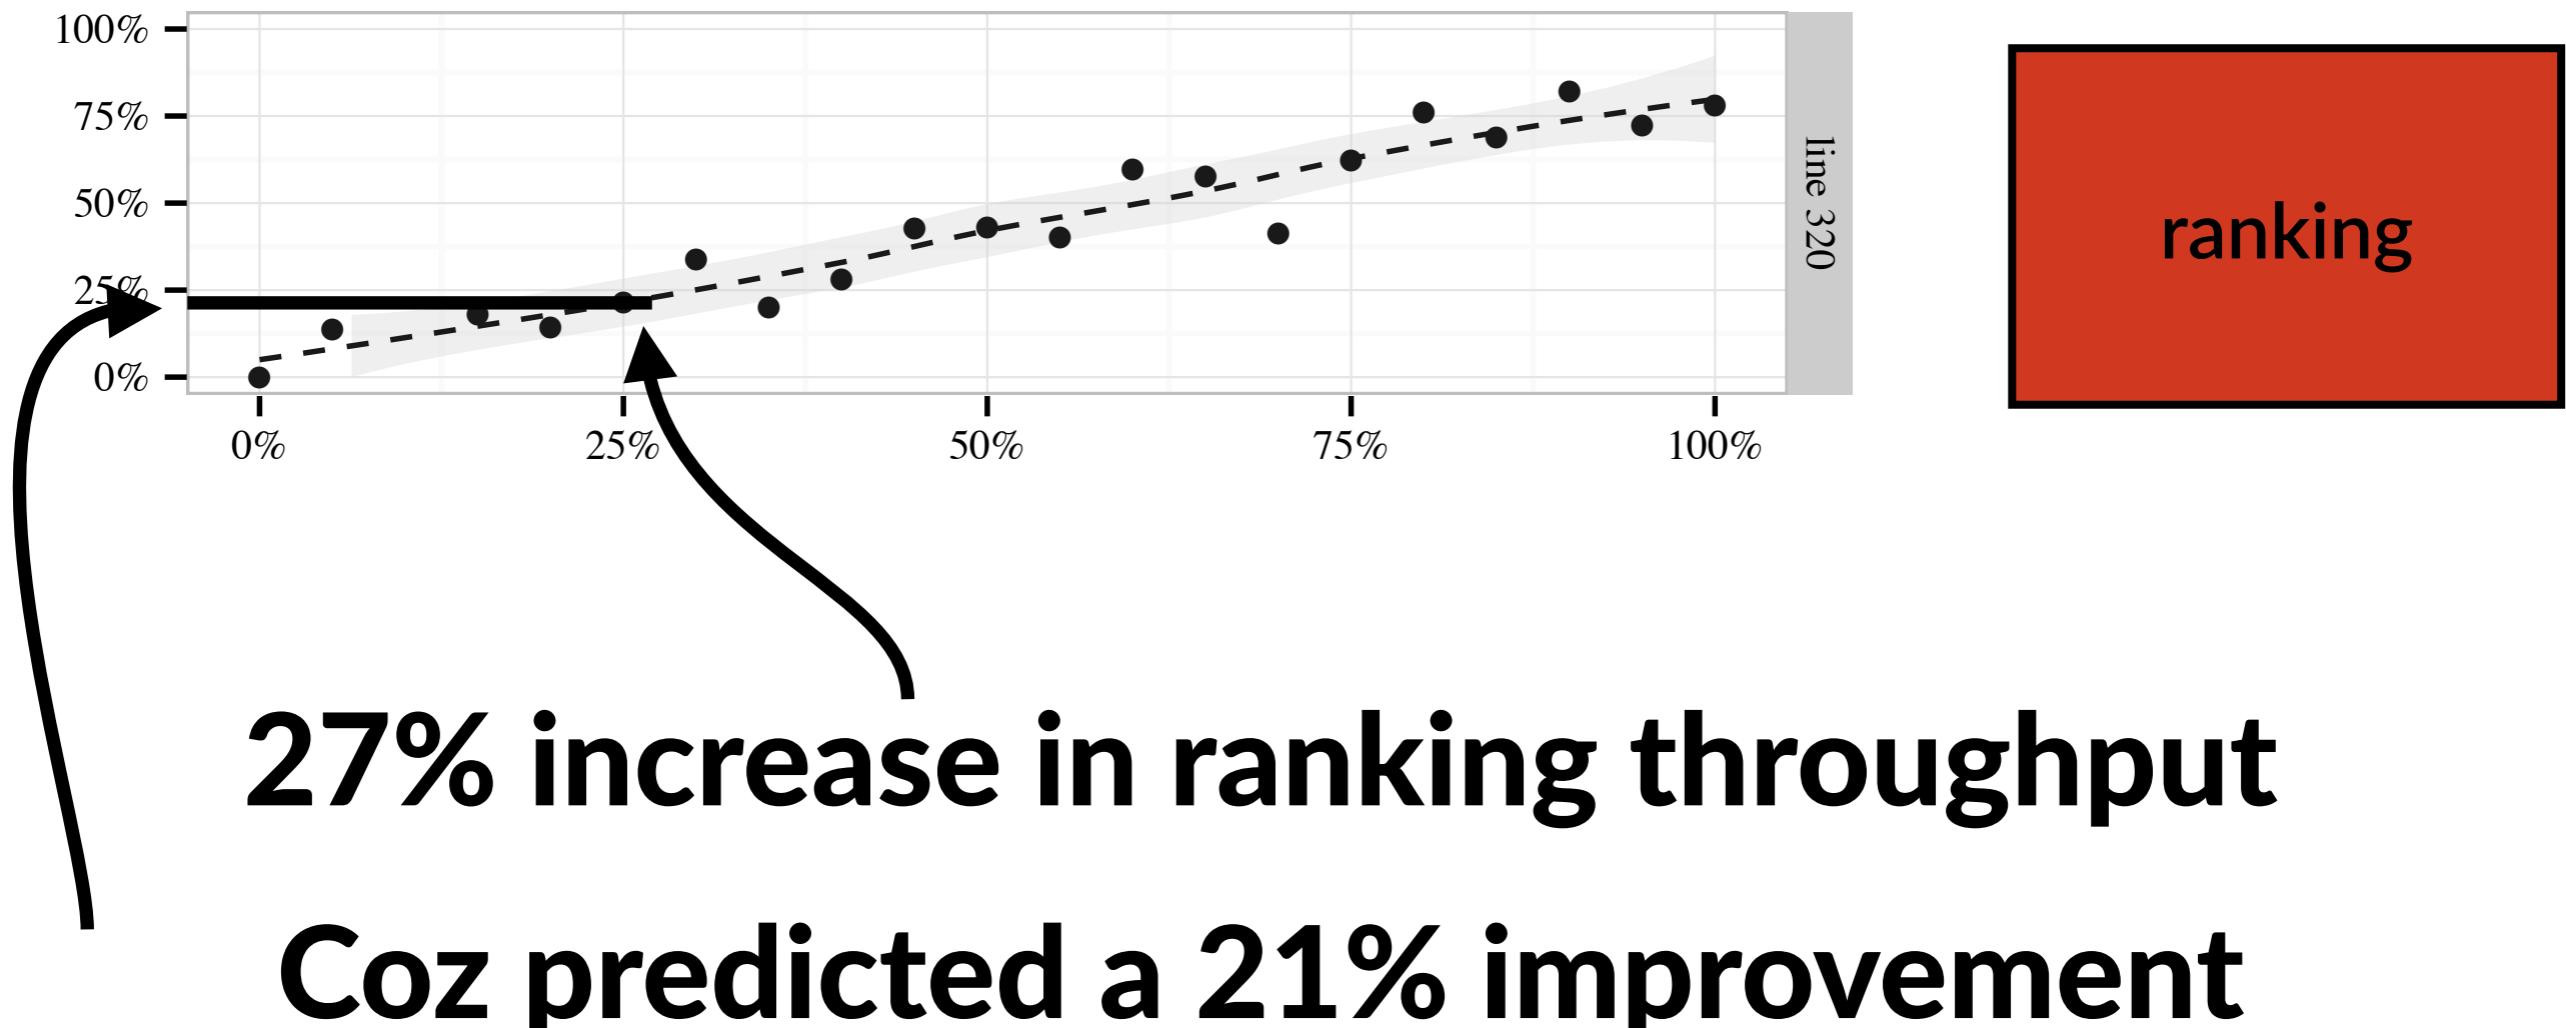


ranking

Increased from 16 to 22 threads

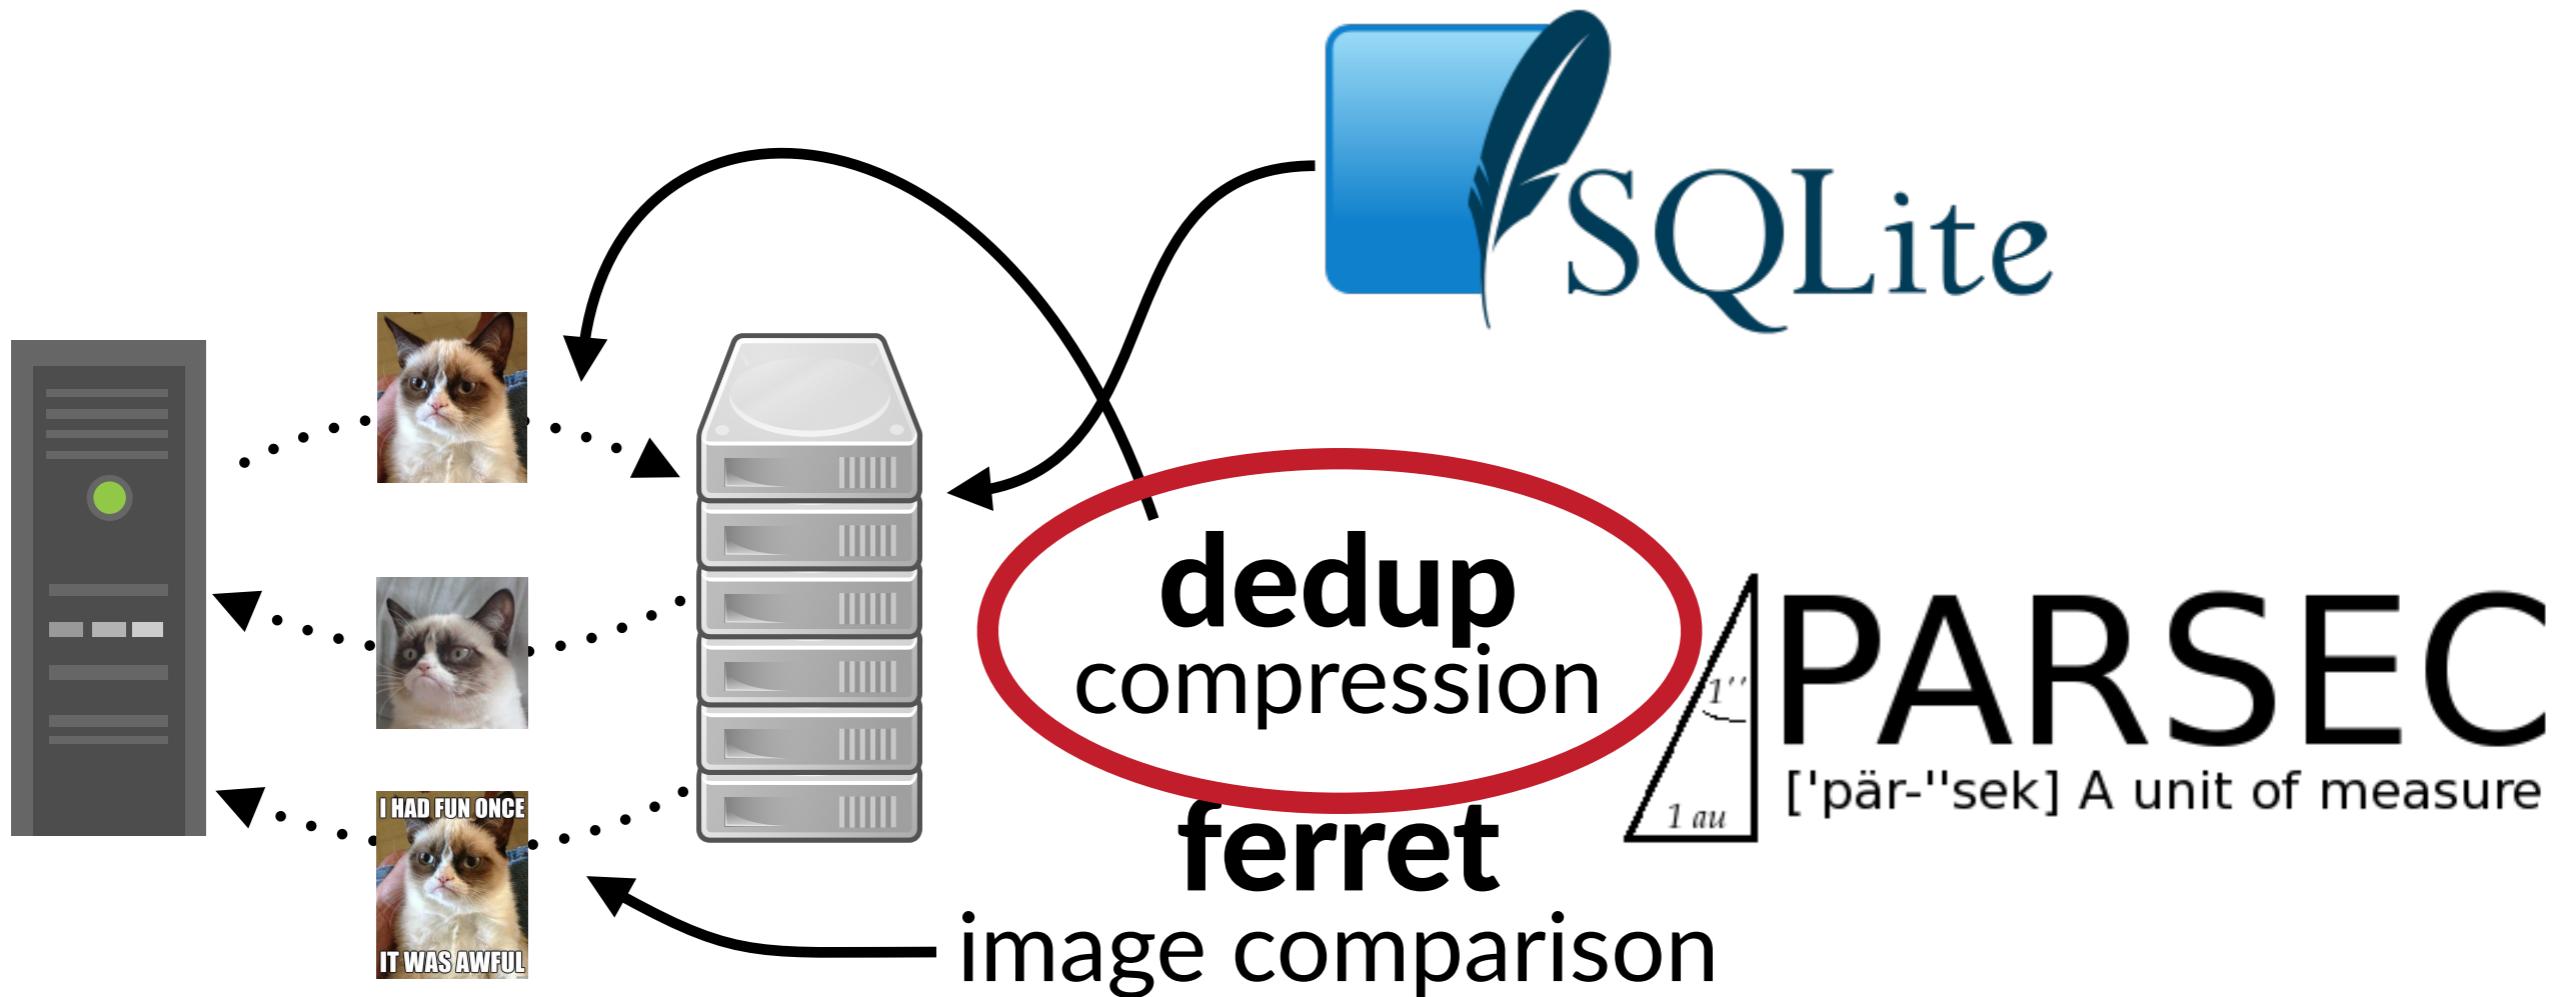
27% increase in ranking throughput

What did Coz predict?



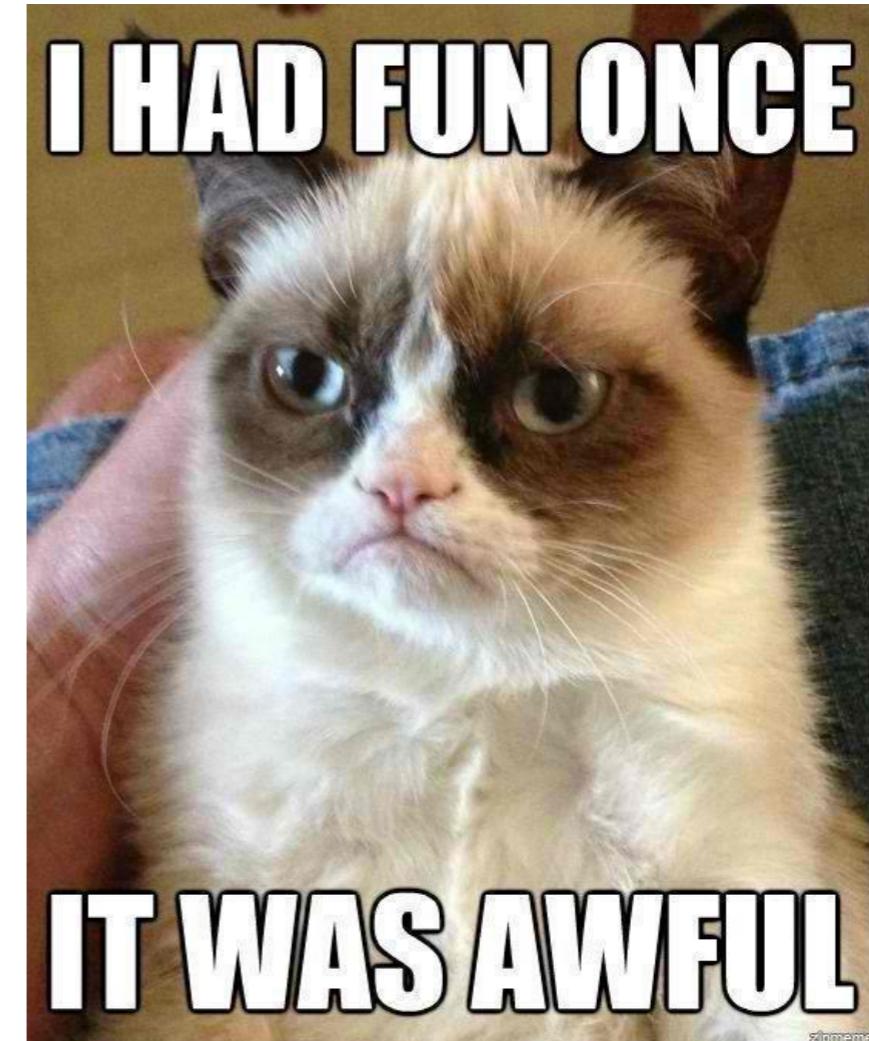
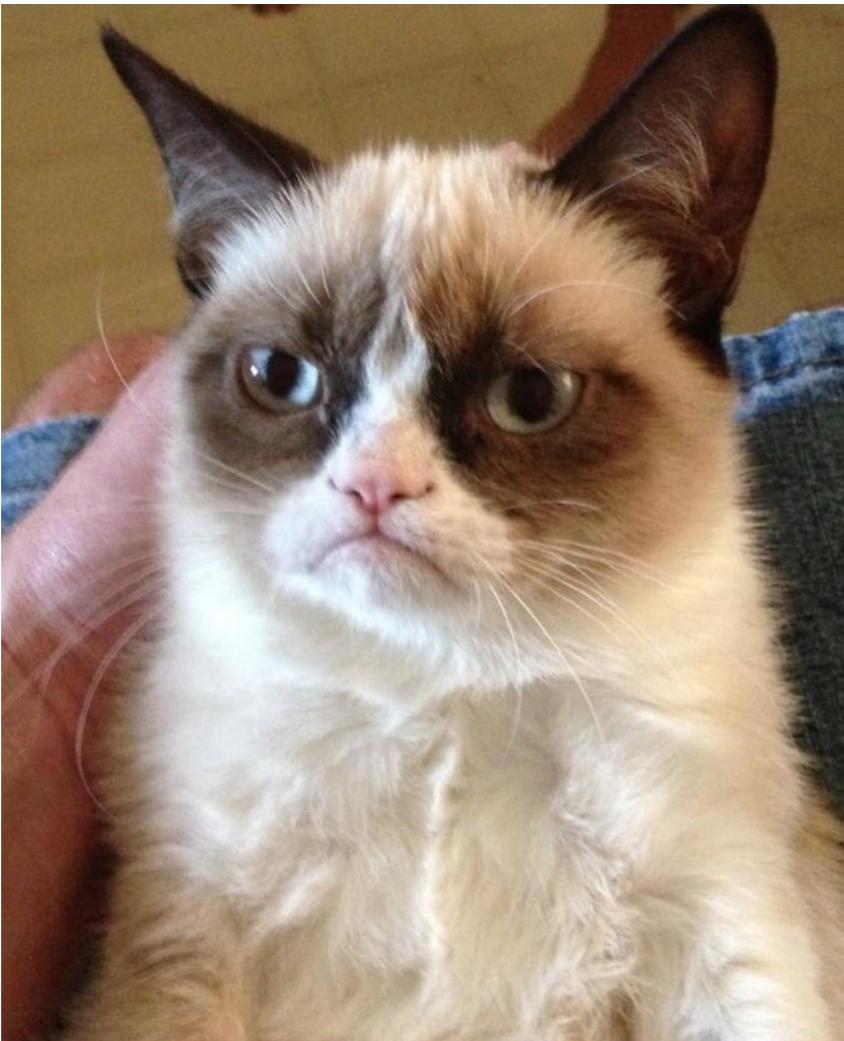
Exactly what we observed

Using Causal Profiling on Ogle



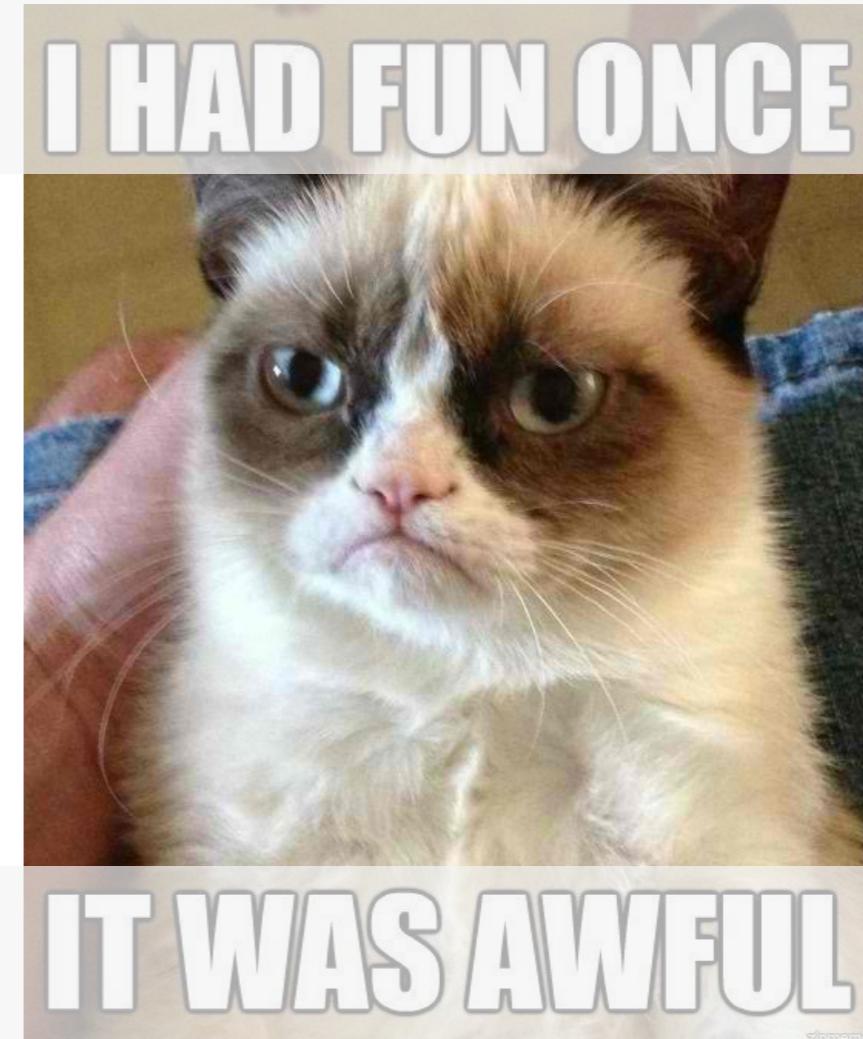
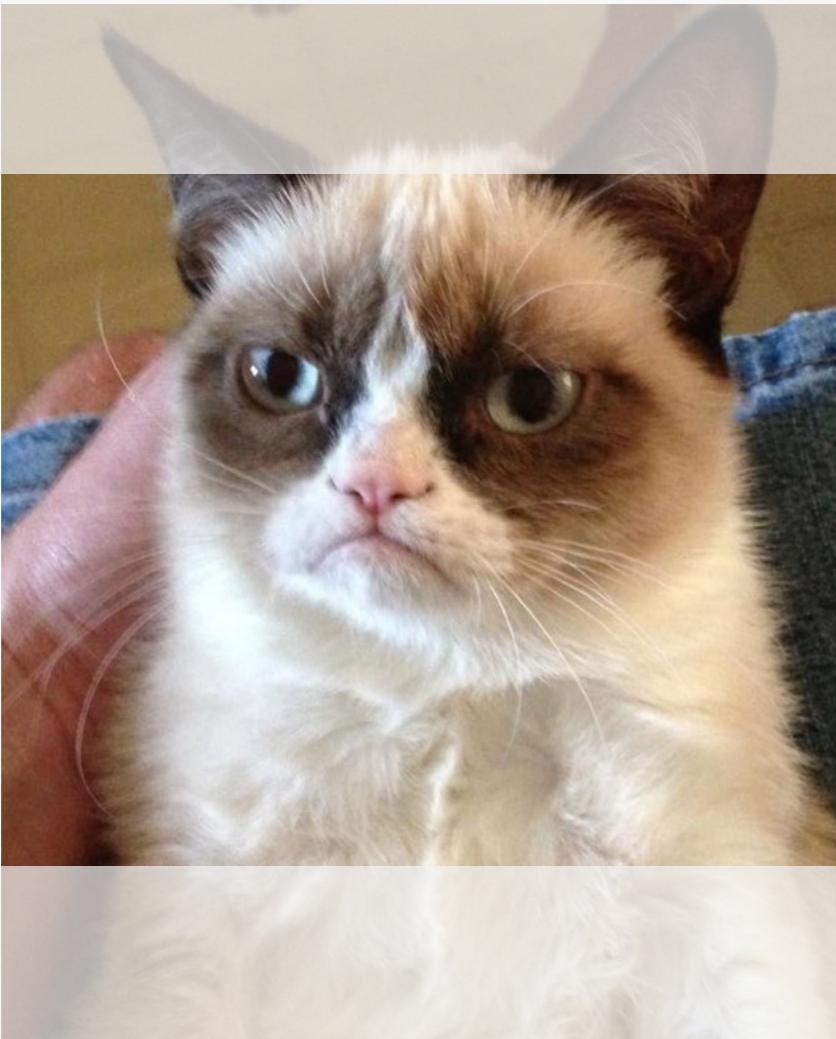
Dedup

Compression via deduplication



Dedup

Compression via deduplication

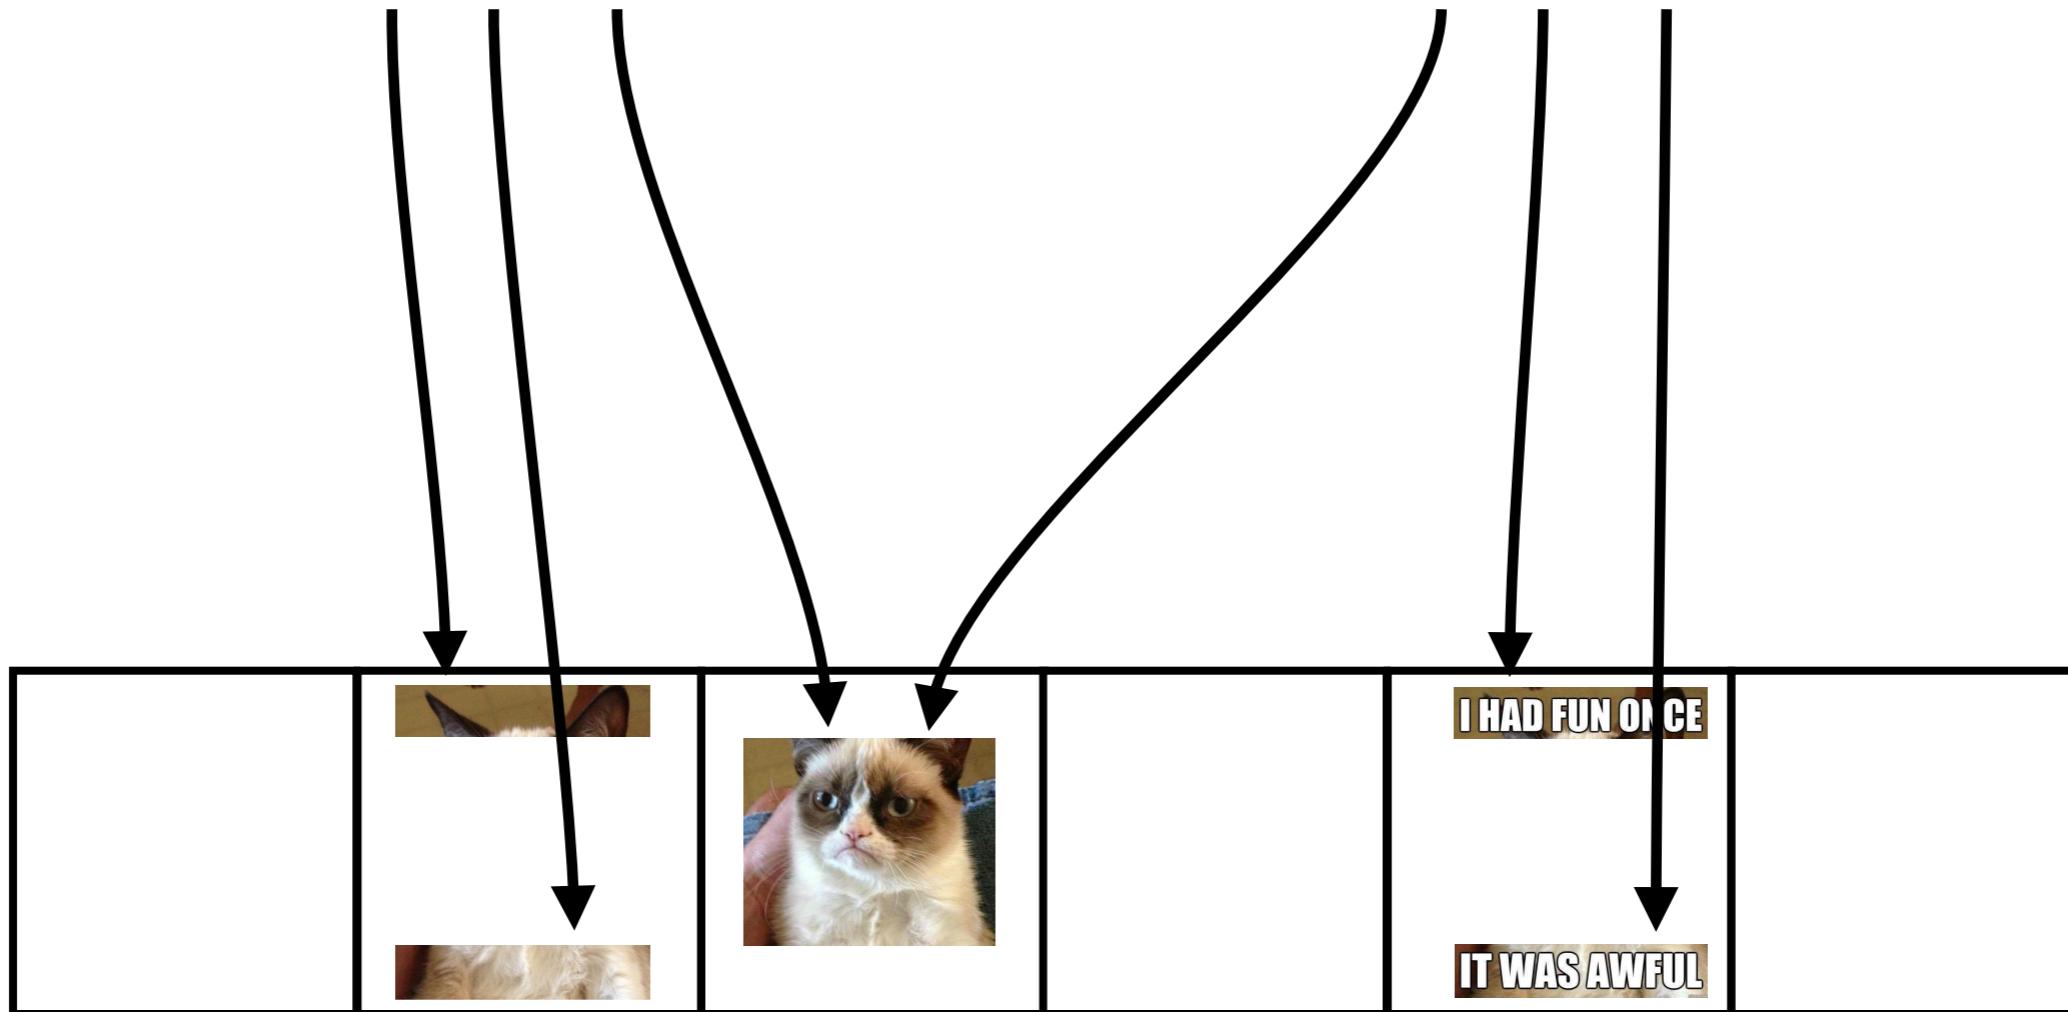


Dedup

Compression via deduplication

grumpycat1.jpg

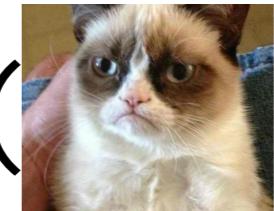
funisawful.jpg



Dedup

Compression via deduplication

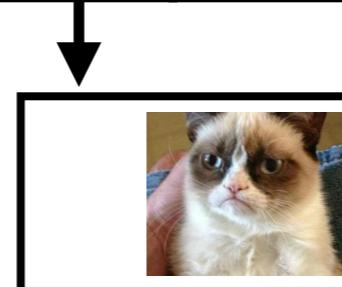
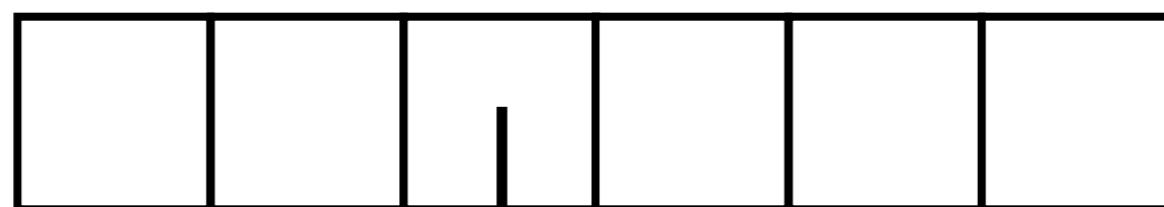
i = hash_function()



Dedup

Compression via deduplication

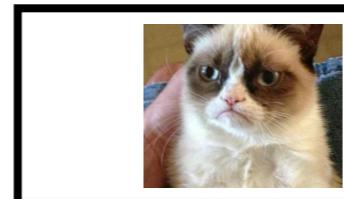
```
i = hash_function()
```



Dedup

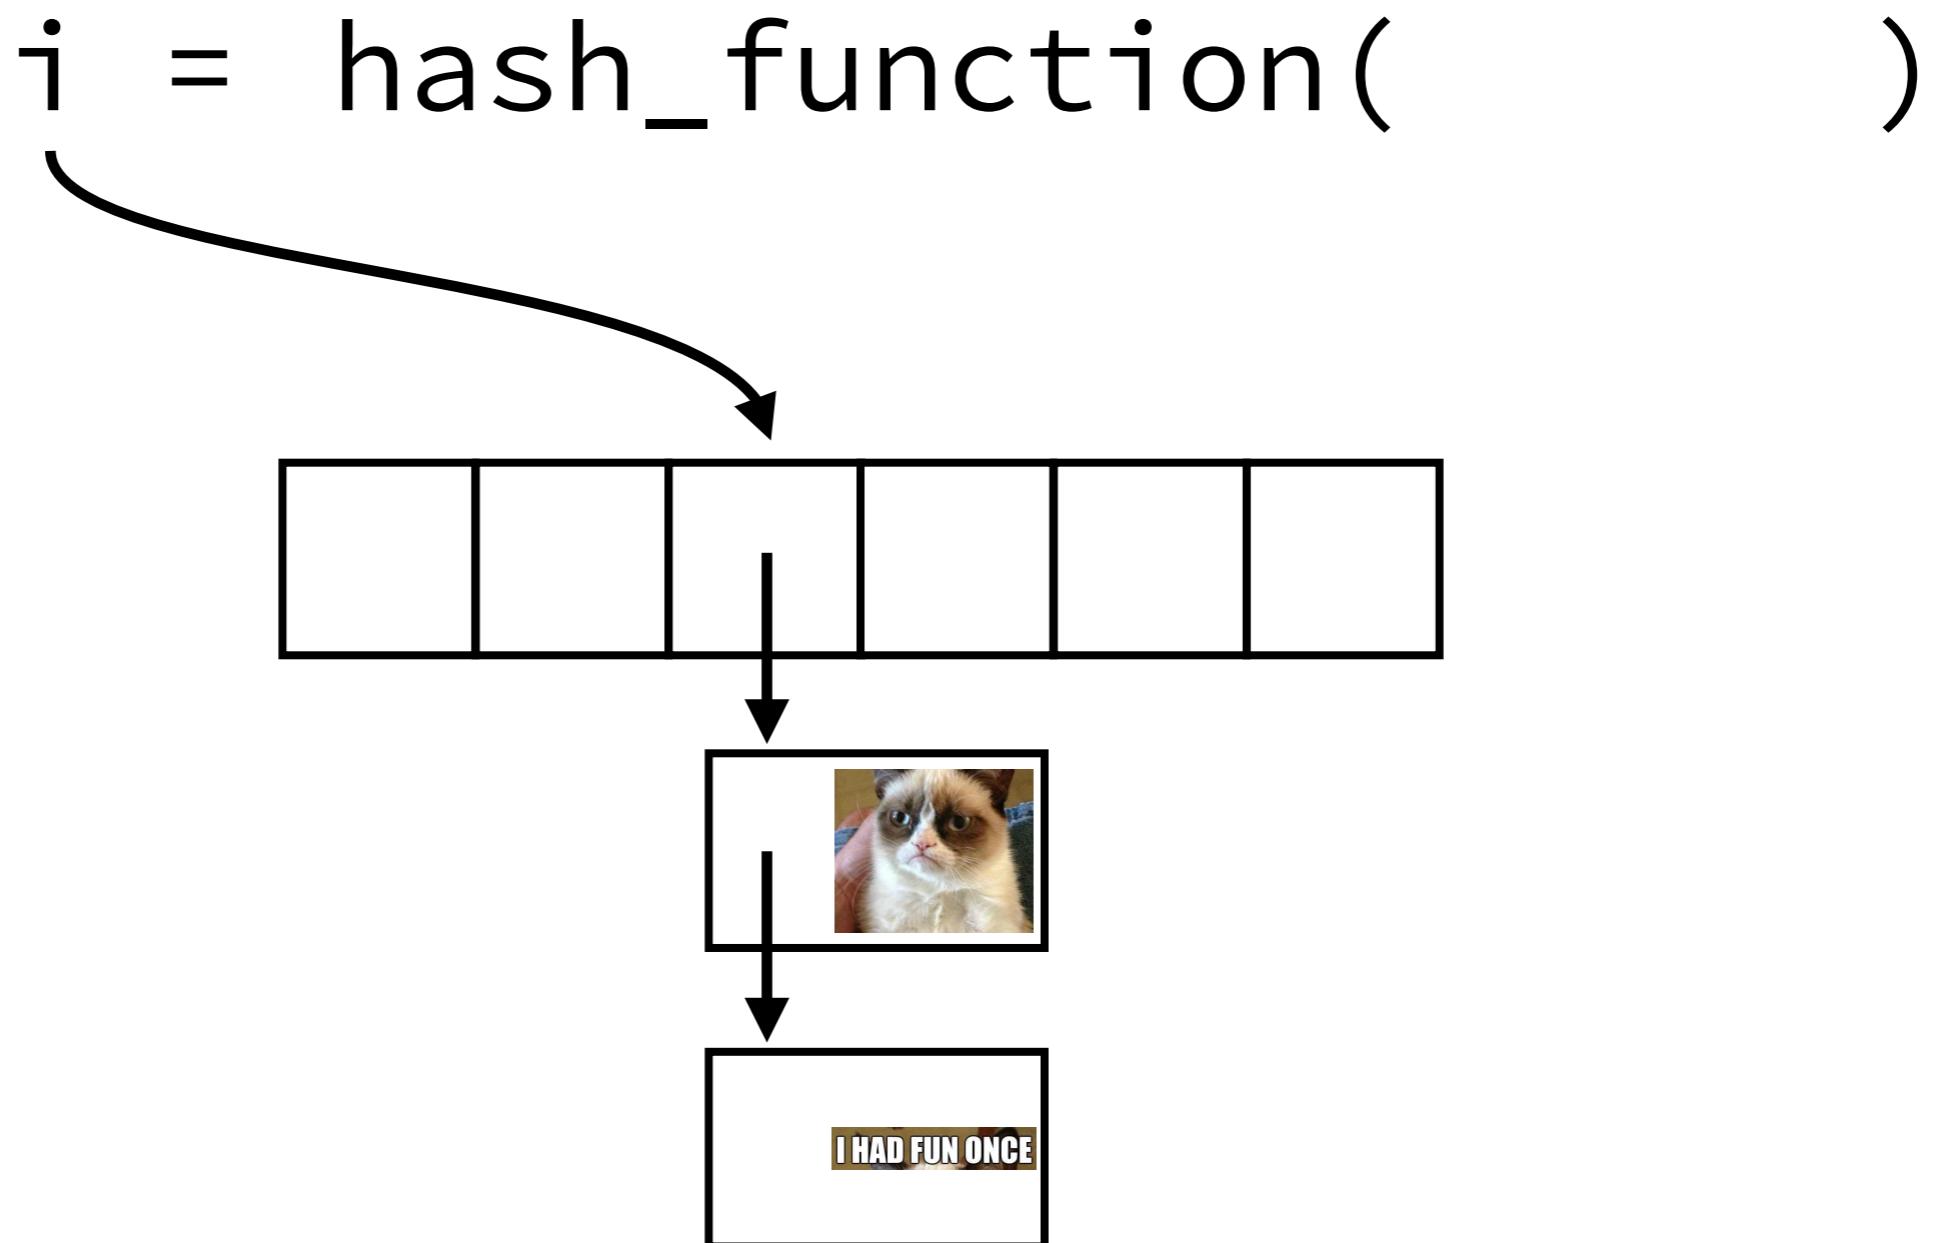
Compression via deduplication

$i = \text{hash_function}(\text{I HAD FUN ONCE})$



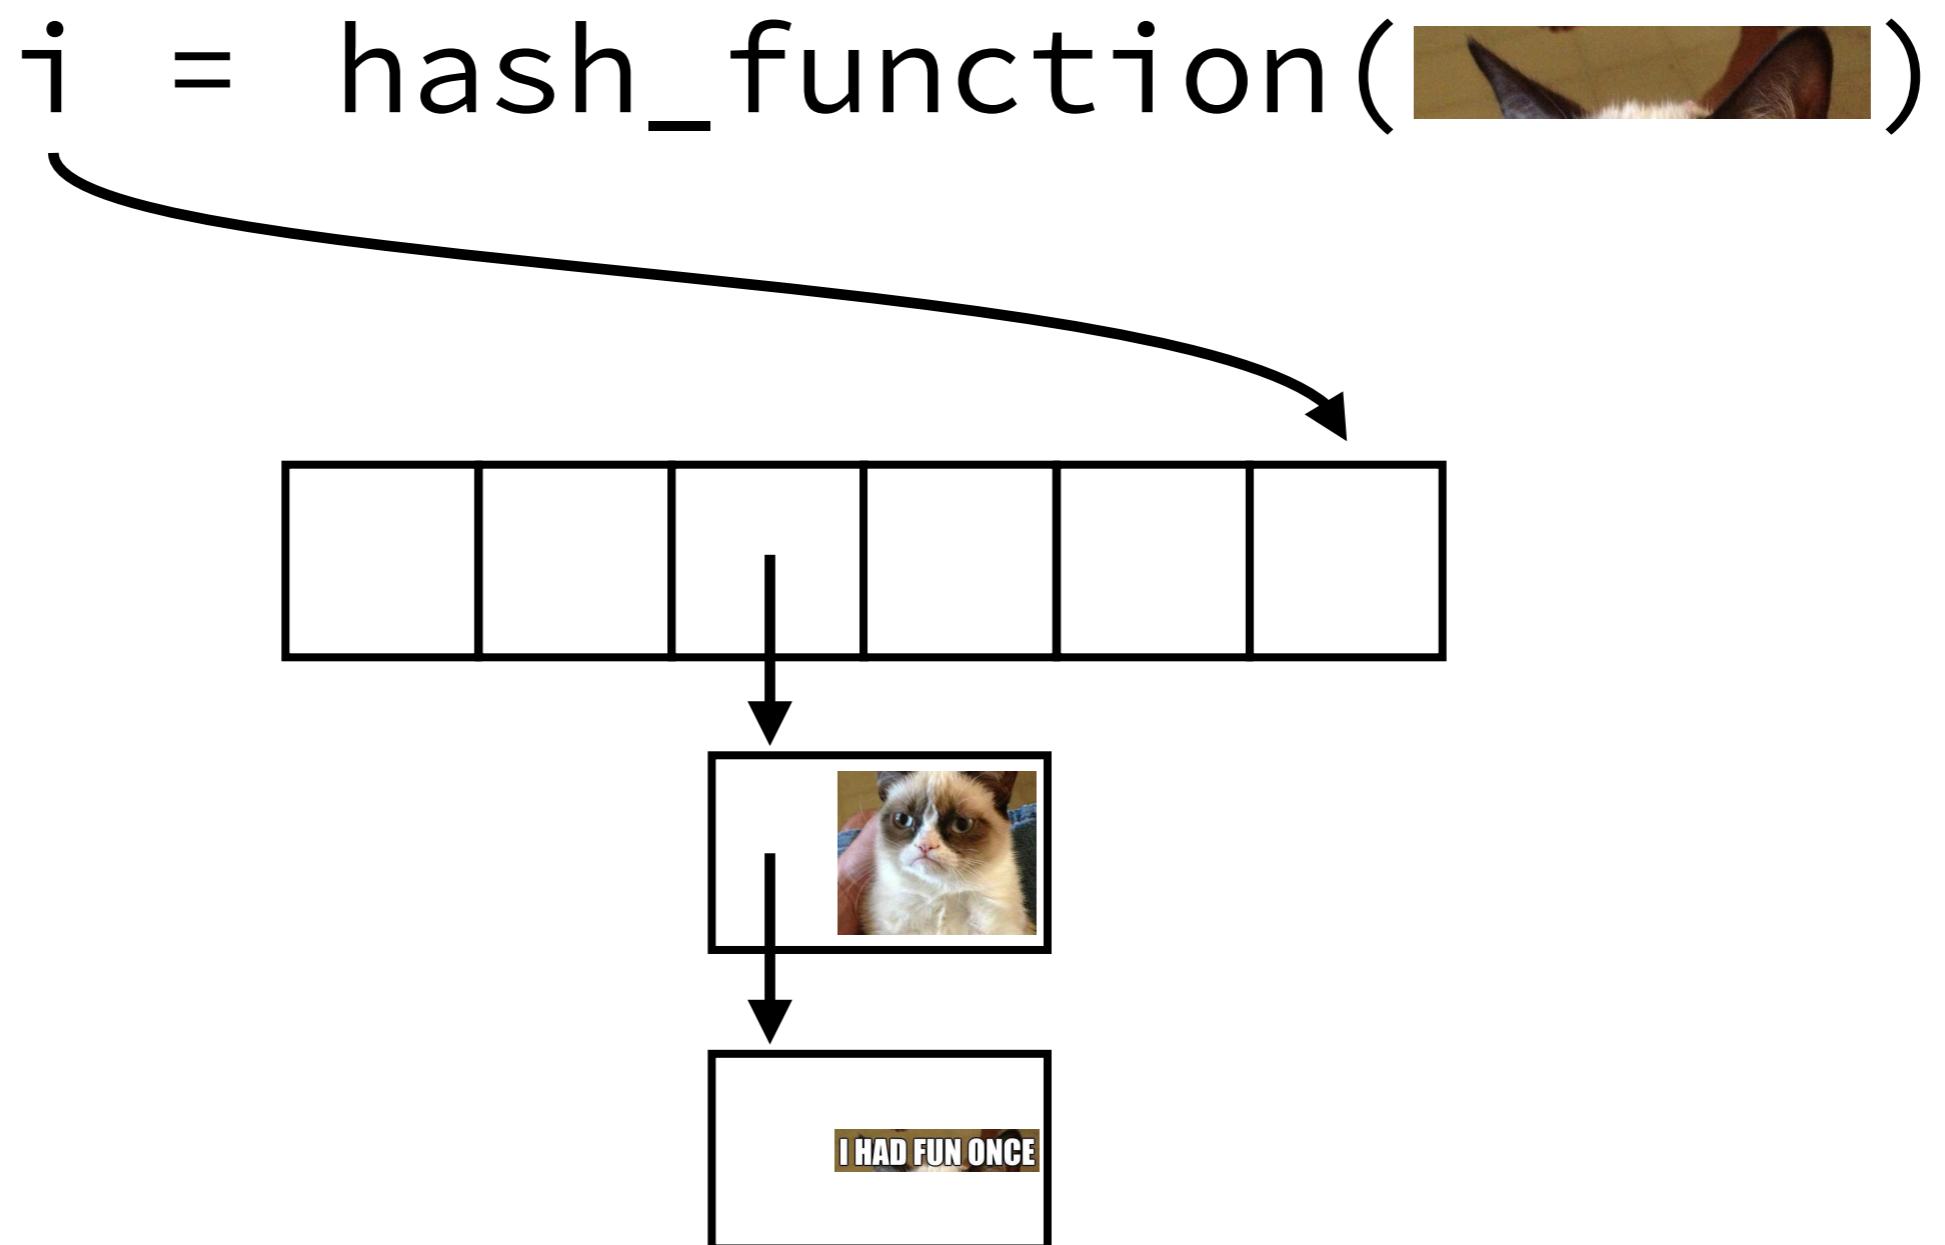
Dedup

Compression via deduplication



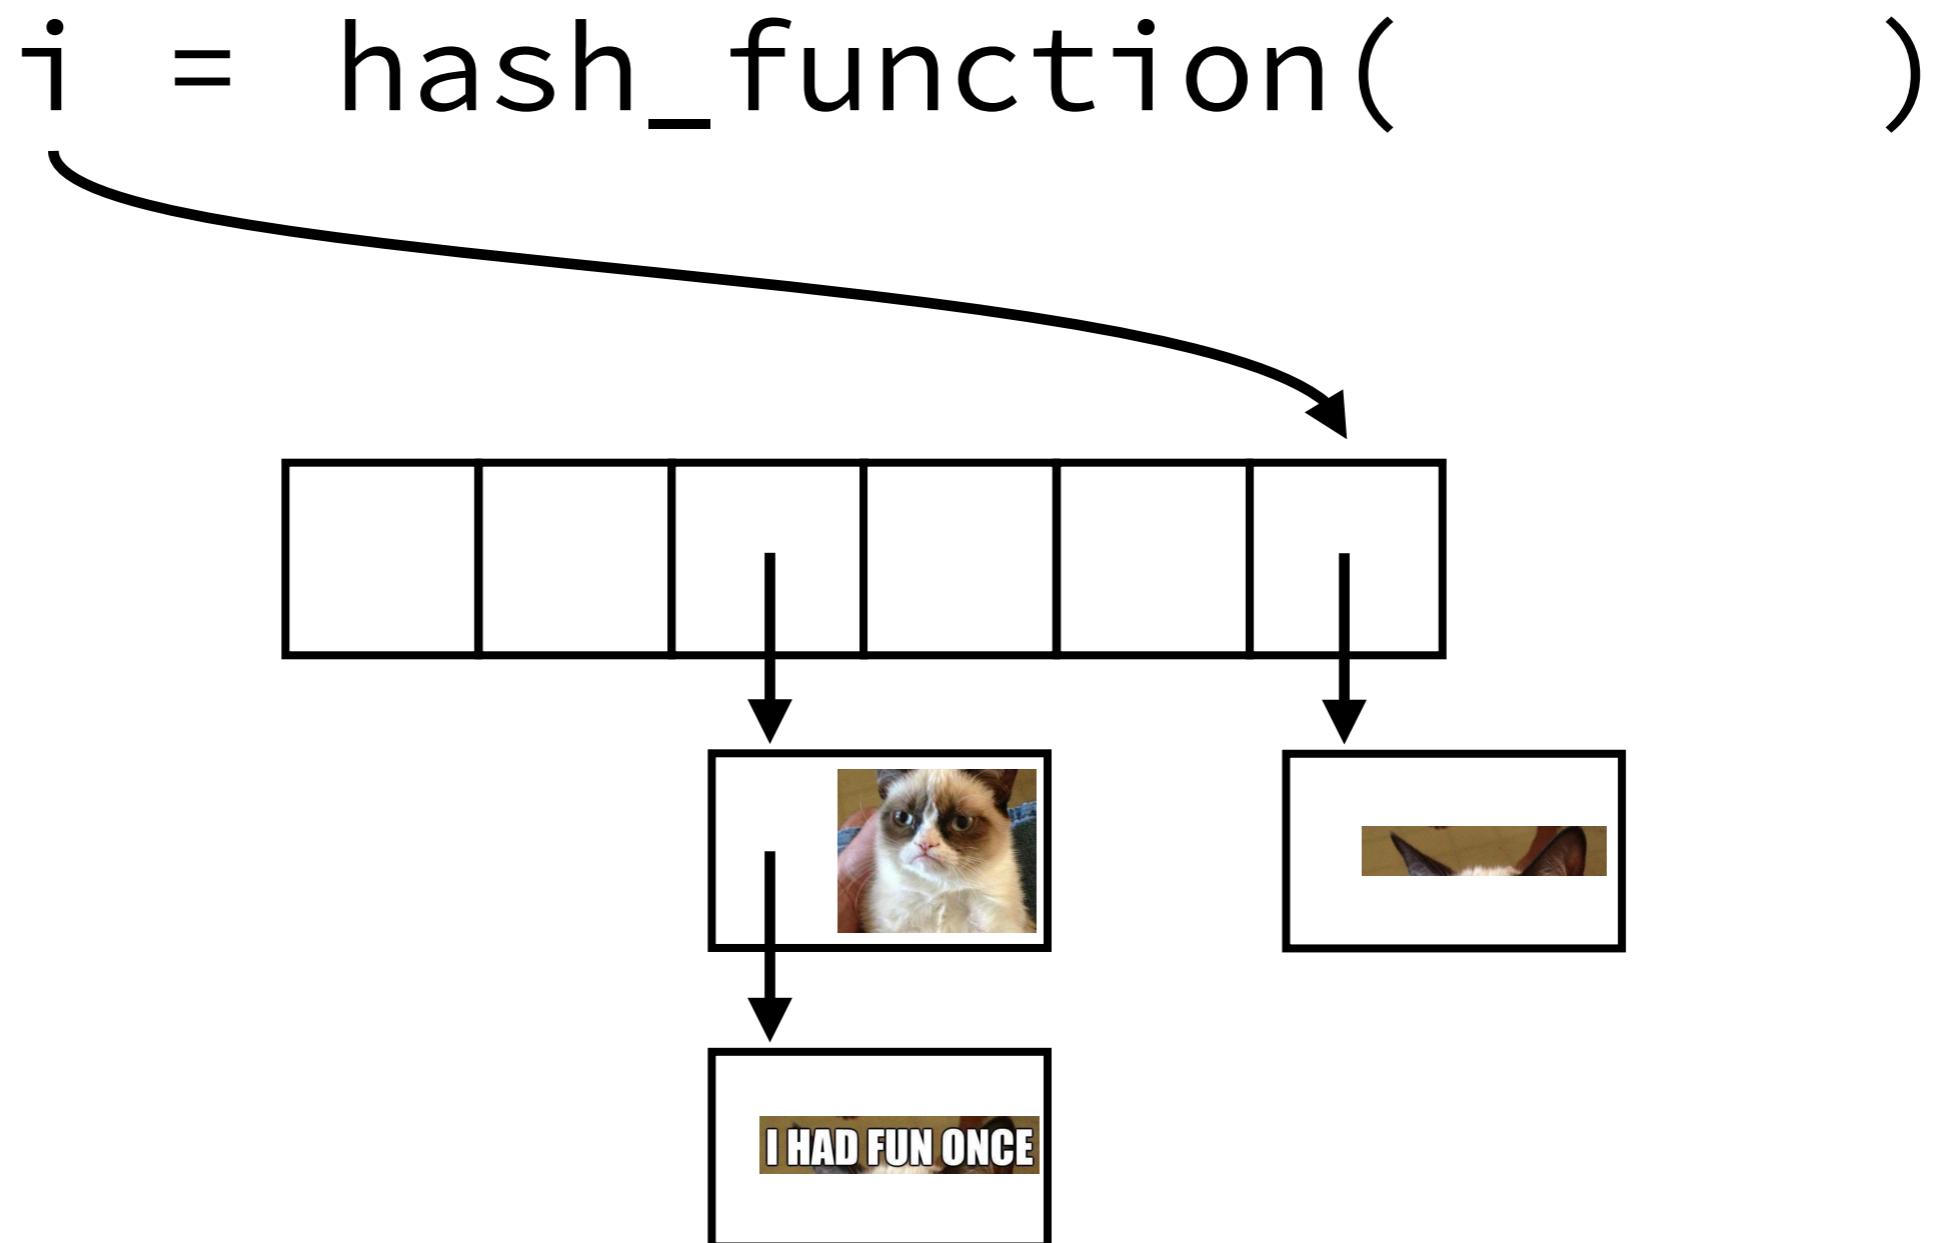
Dedup

Compression via deduplication



Dedup

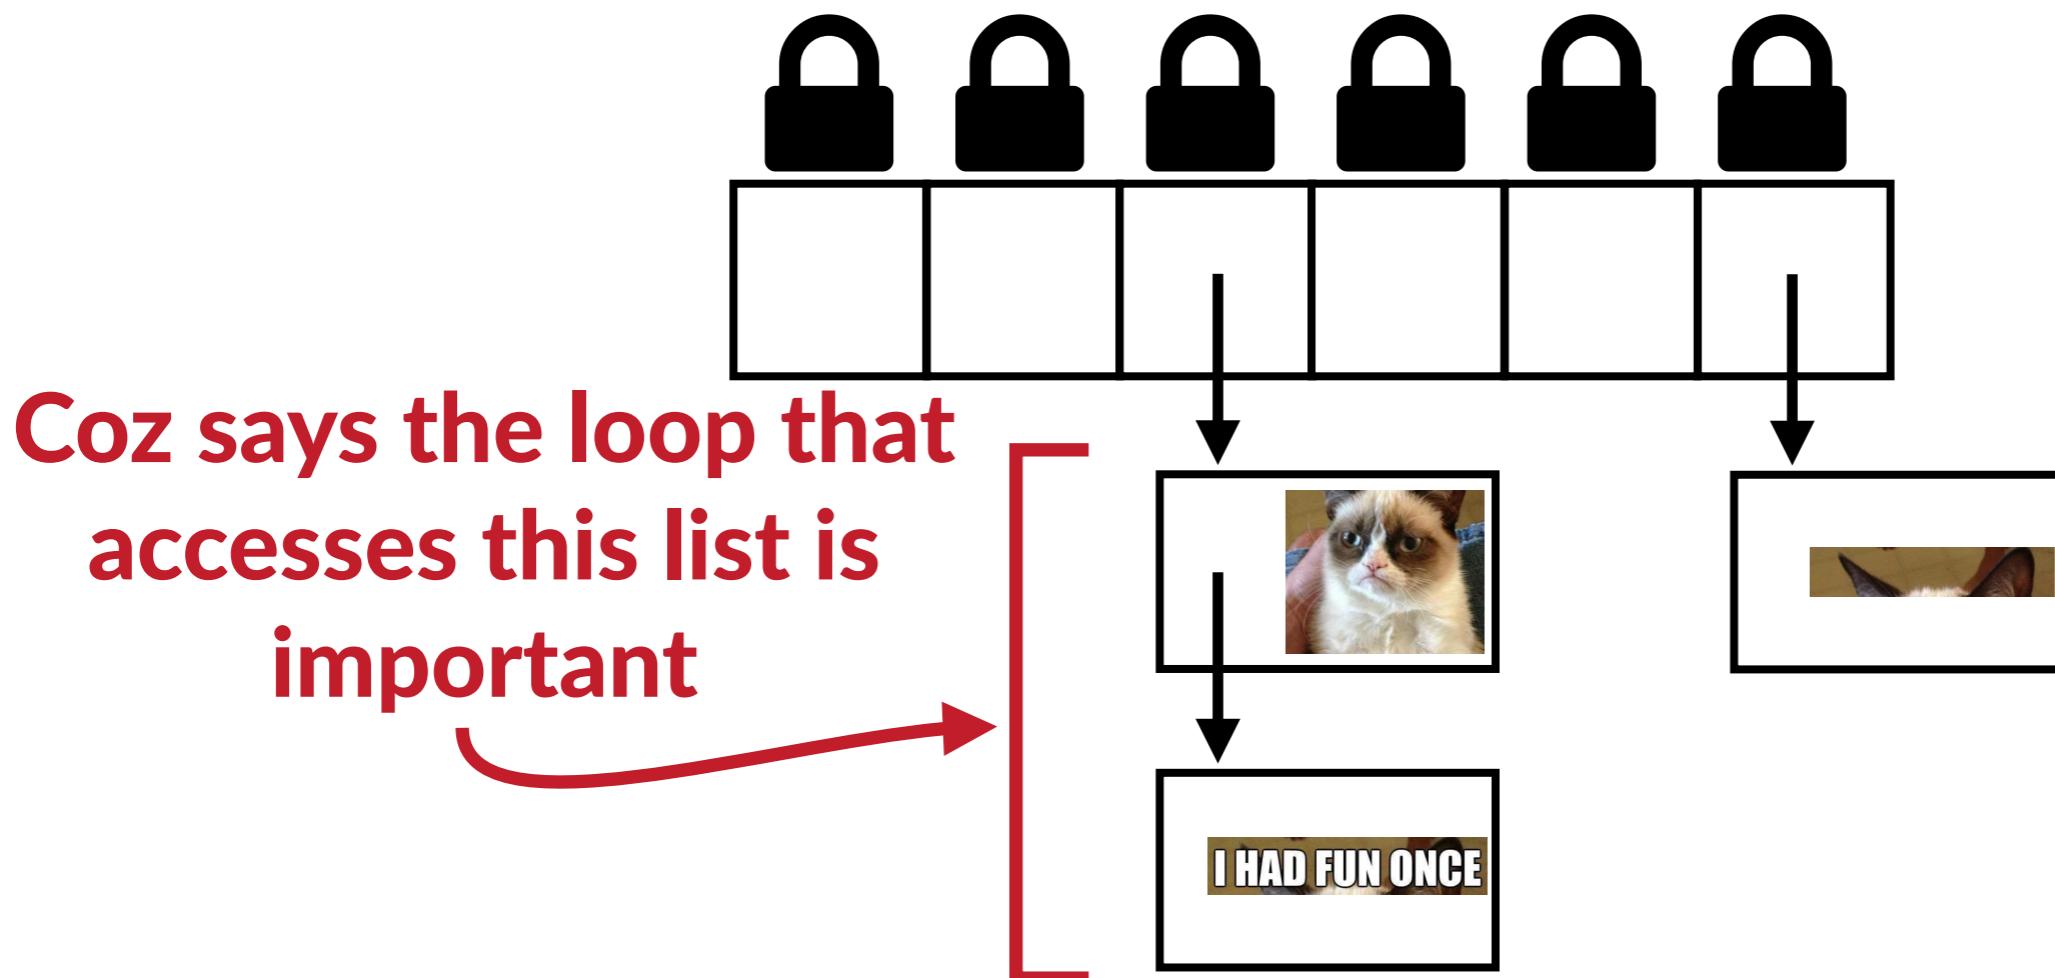
Compression via deduplication



Dedup

Compression via deduplication

Hash table is accessed concurrently by many threads

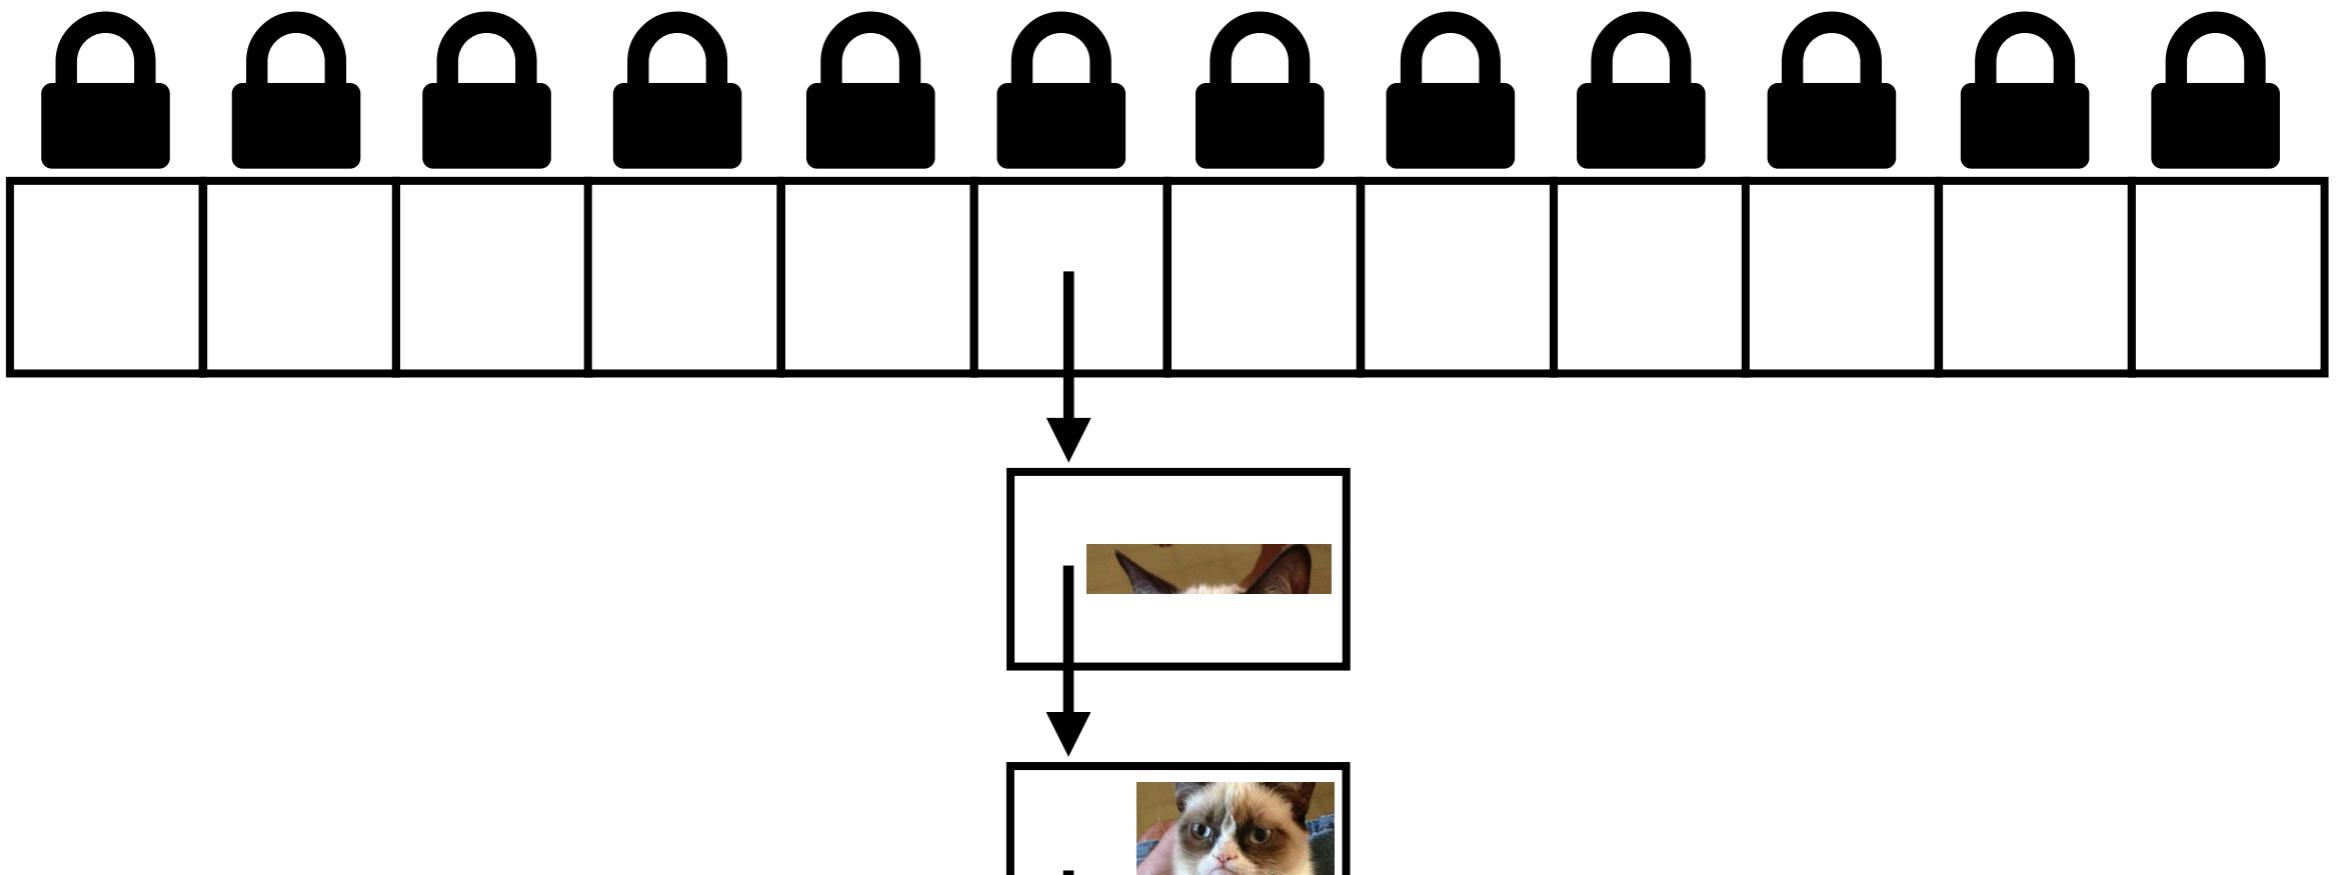


Dedup

Compression via deduplication

More hash buckets should lead to fewer collisions

No performance improvement

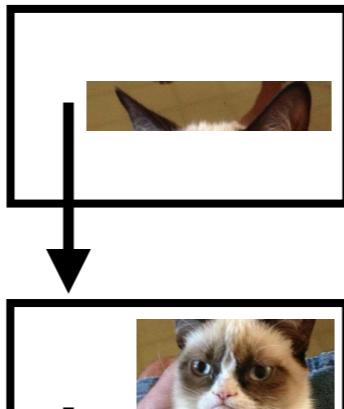
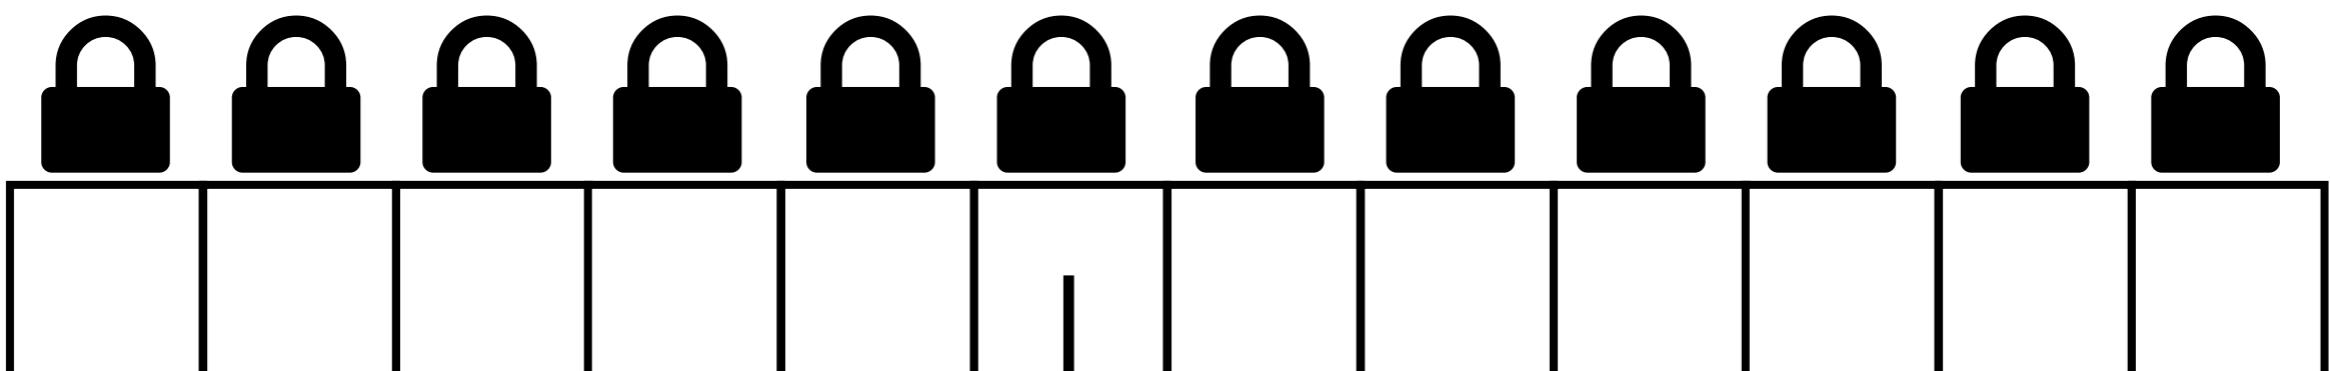


Dedup

Compression via deduplication

What else could be causing collisions?

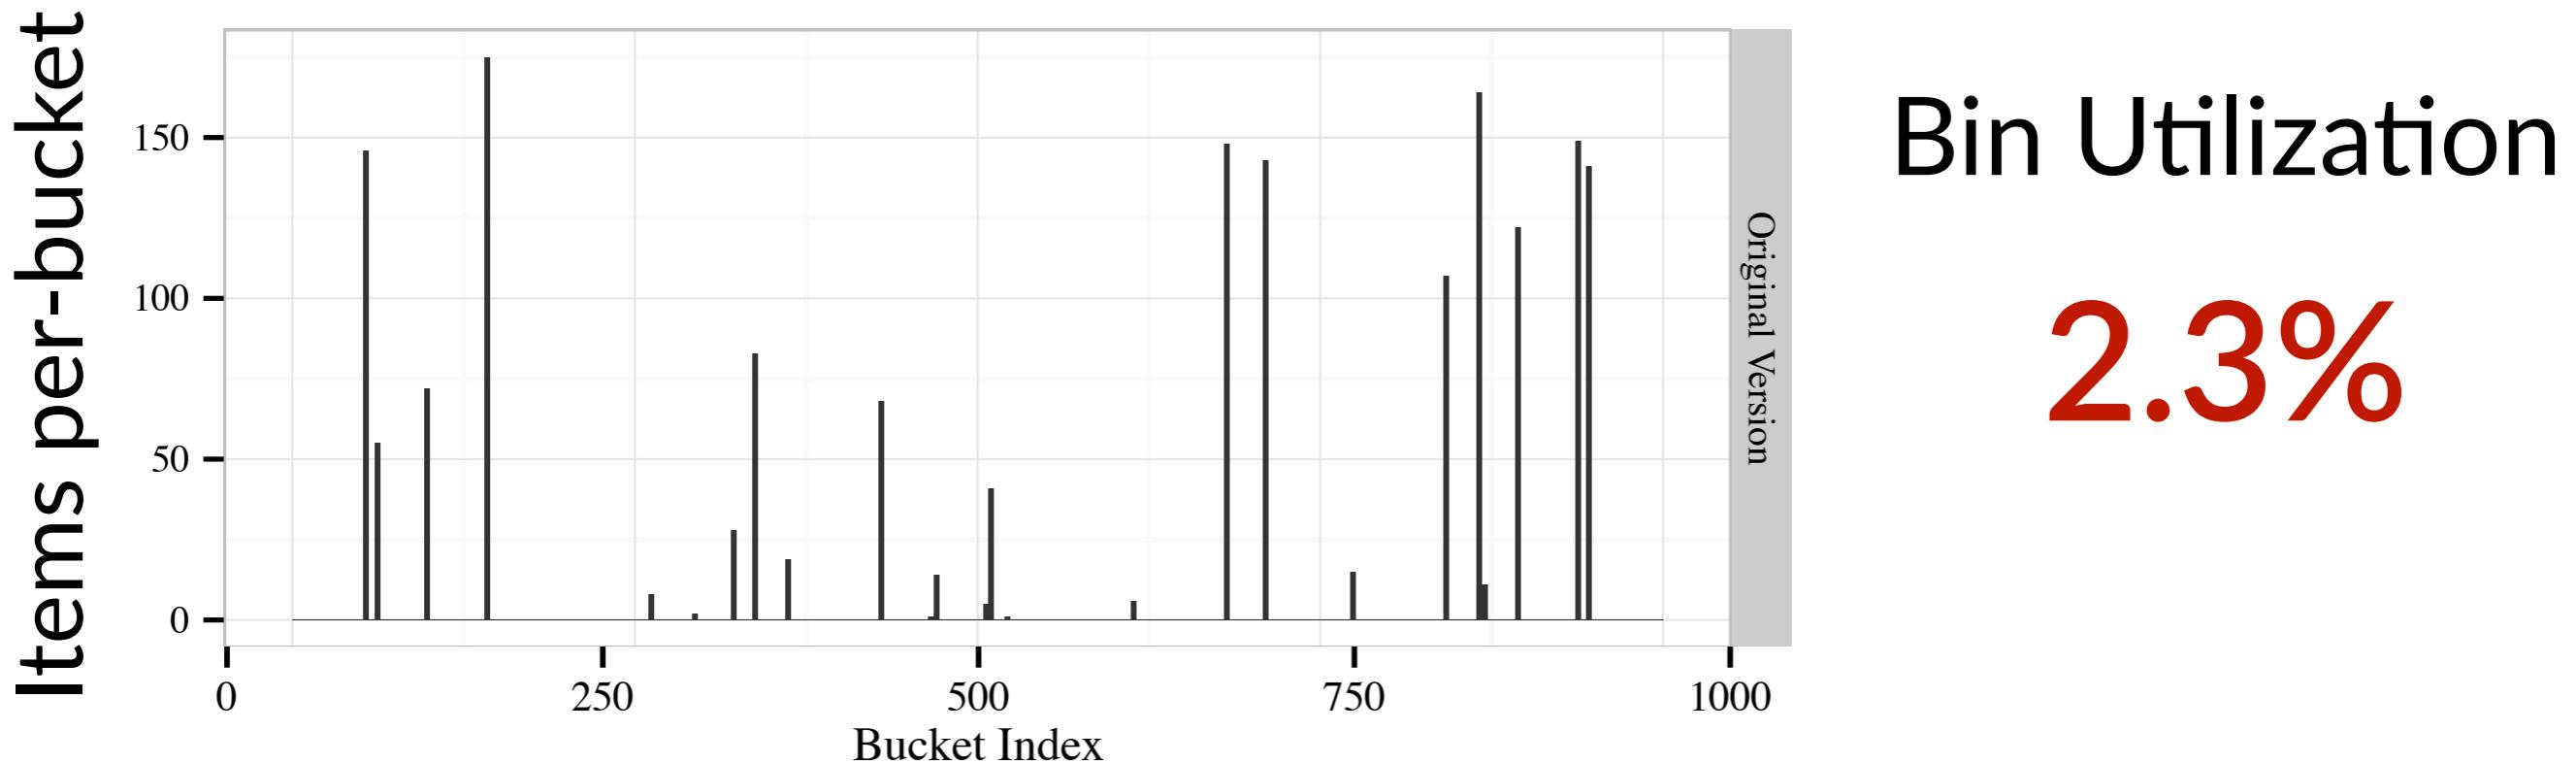
$i = \text{hash_function}(\text{ }$ 



Dedup

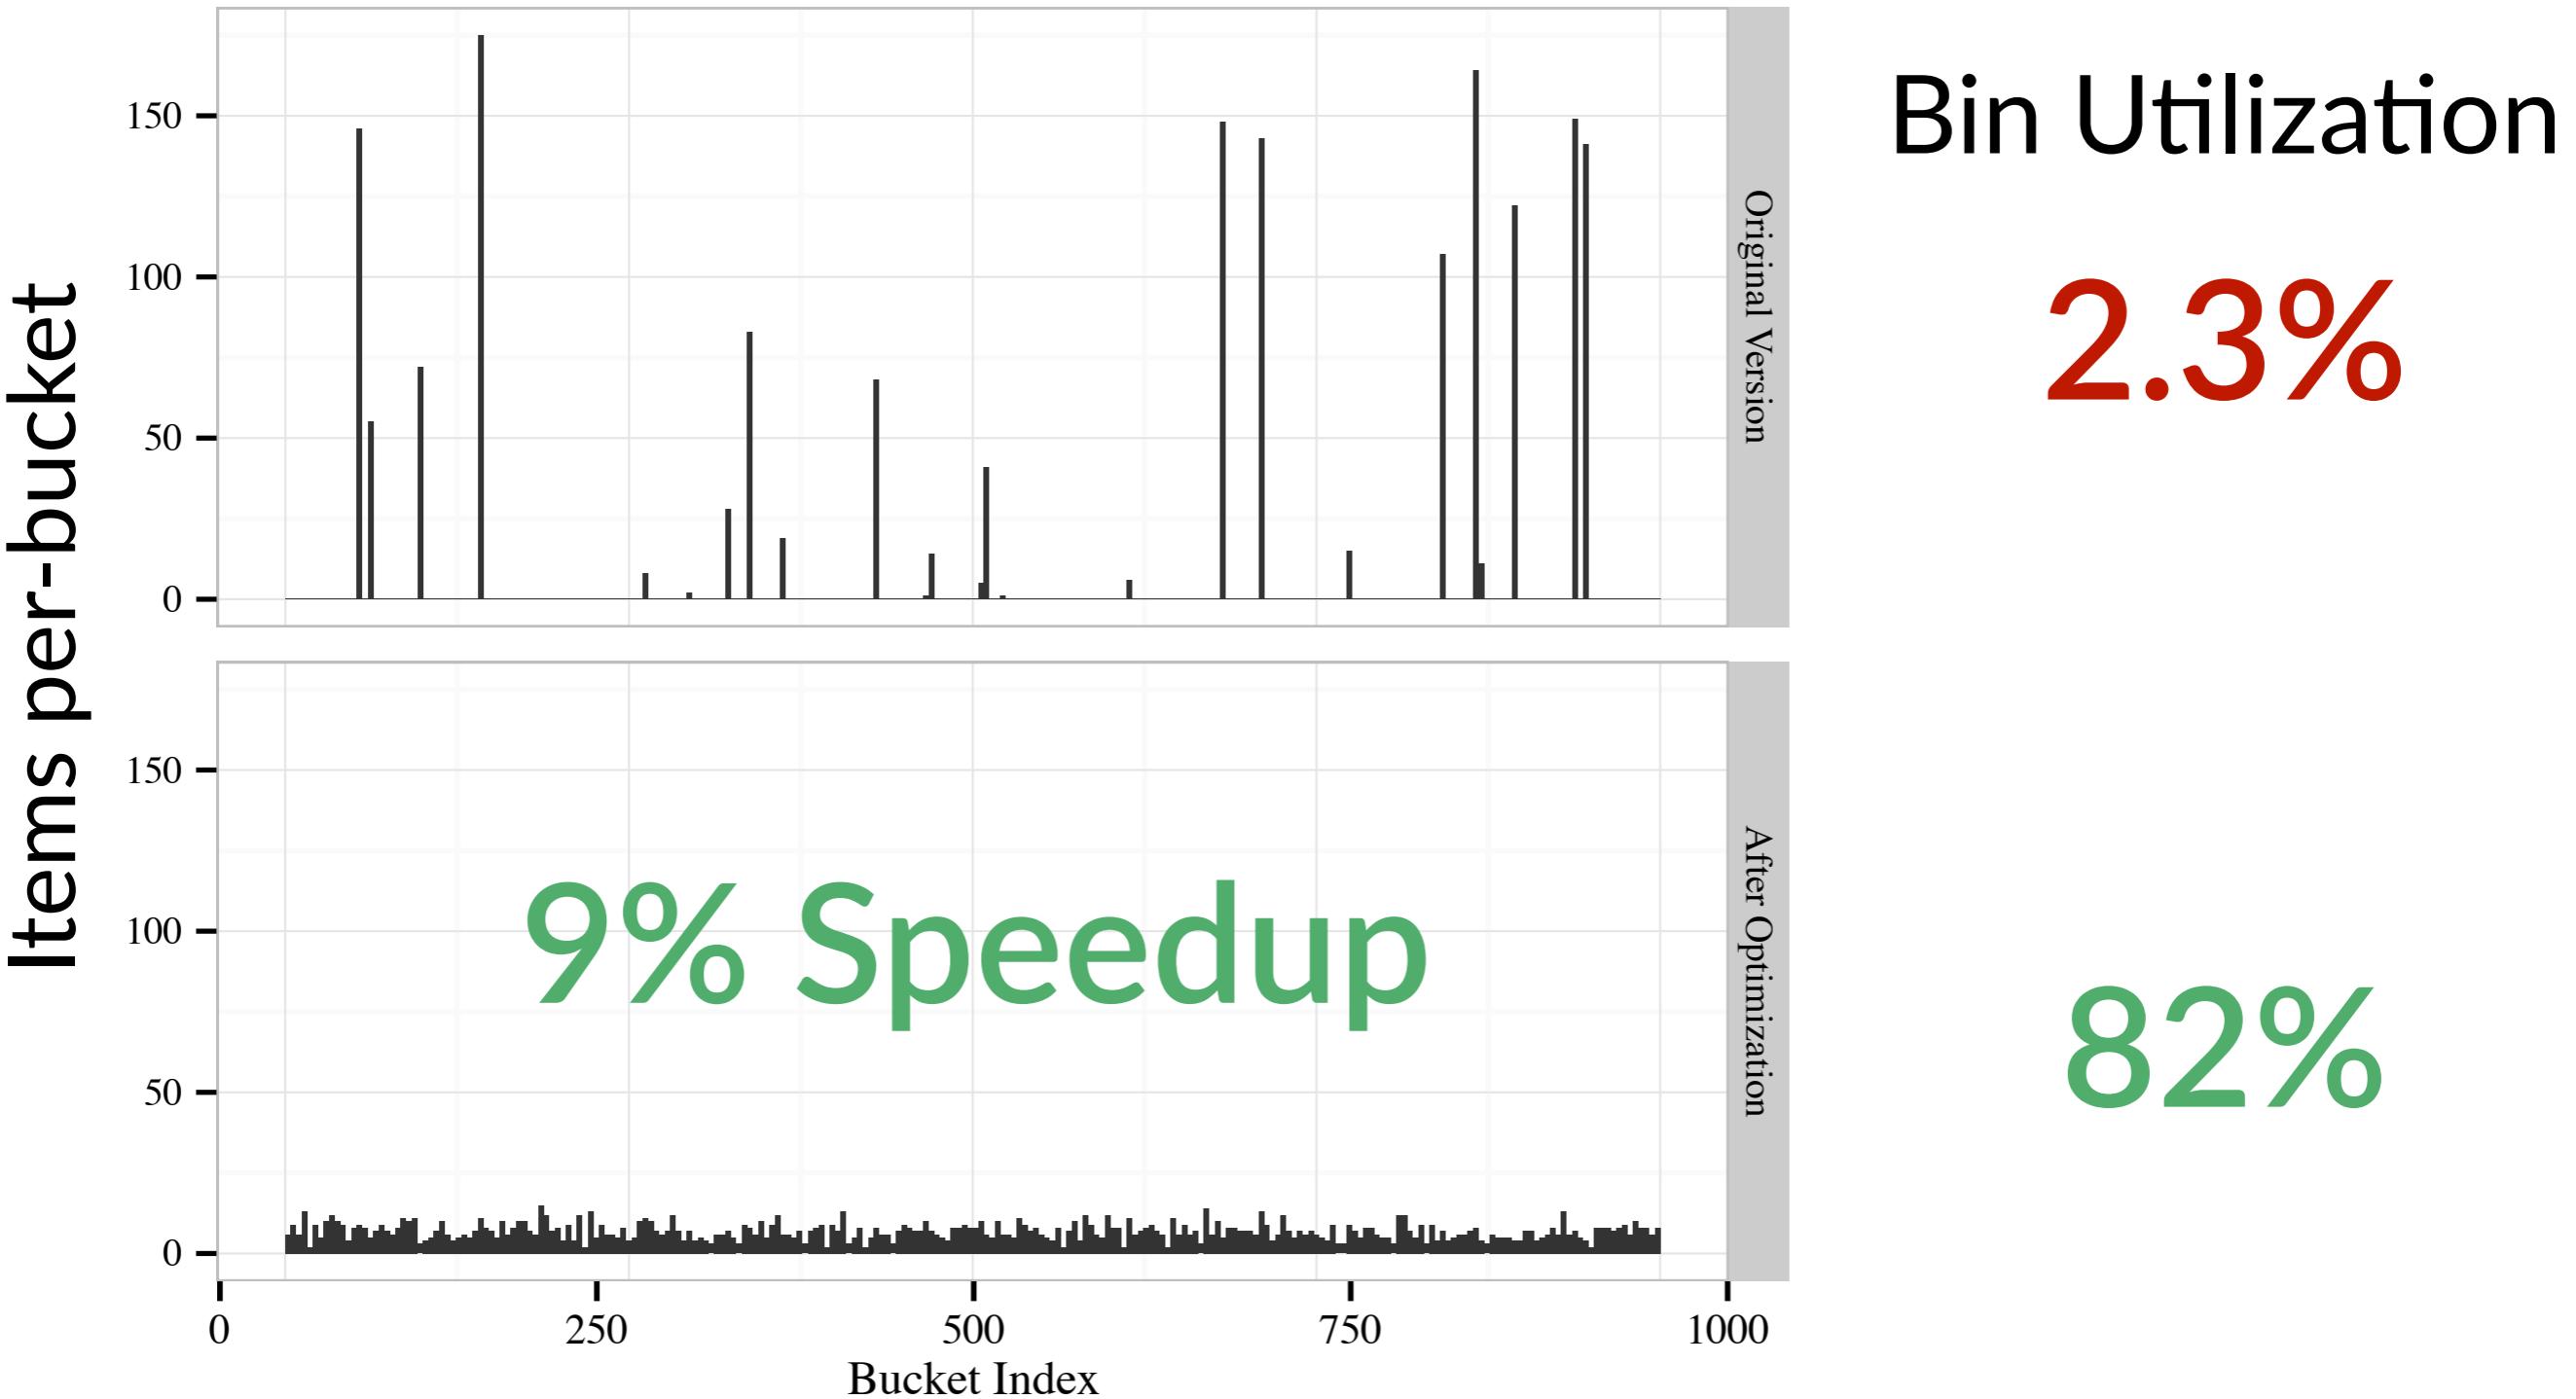
Compression via deduplication

Horrible hash function!



Dedup

Compression via deduplication



Dedup

Compression via deduplication

What did Coz predict?

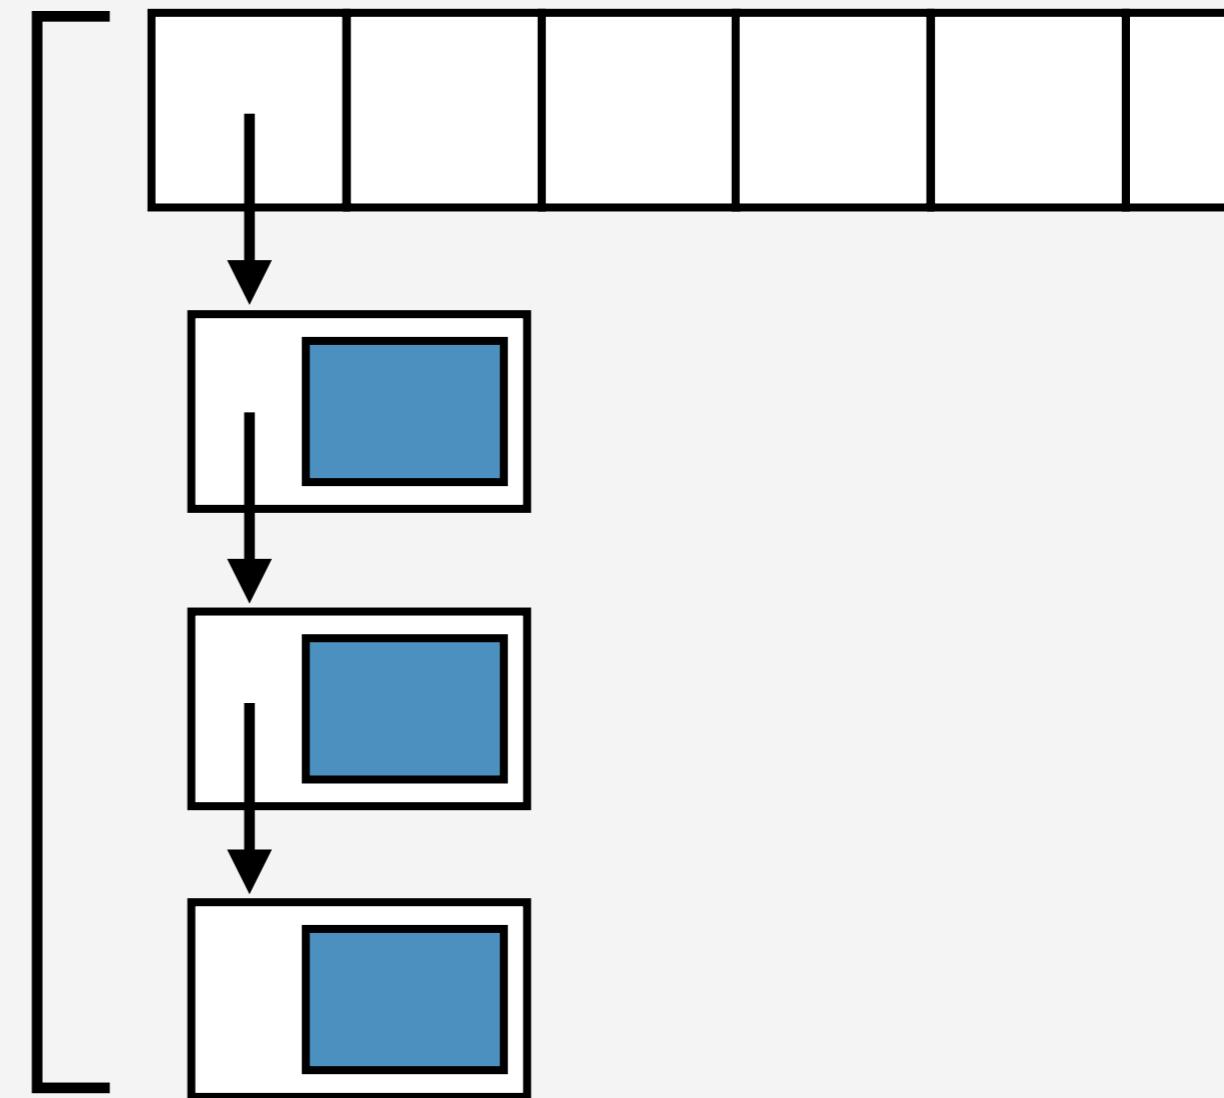
Blocks per-bucket

Before: 76.7

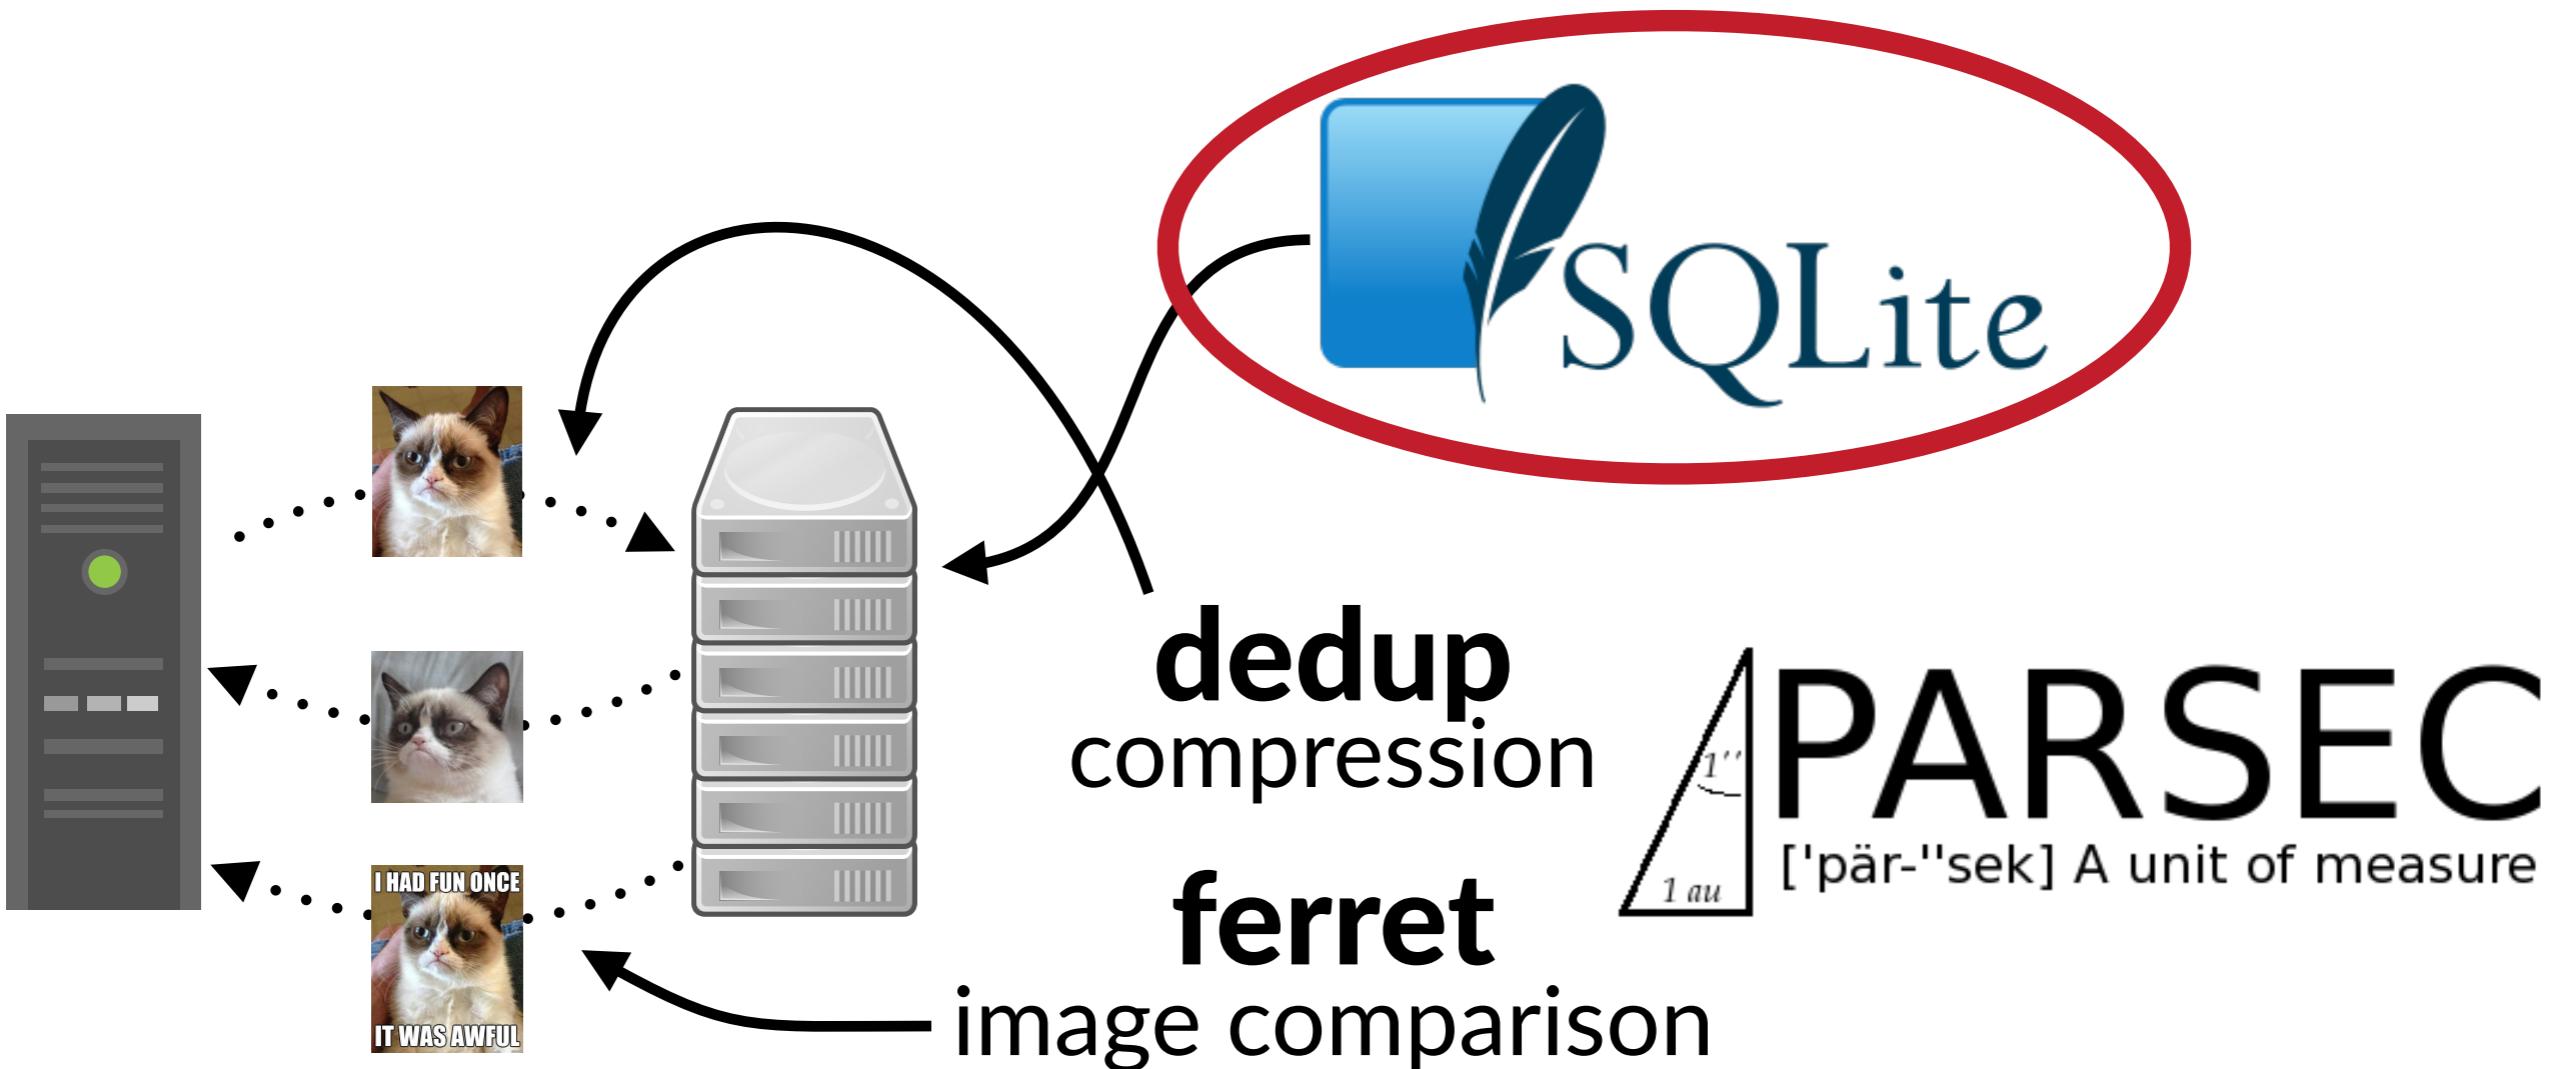
After: 2.09

96% traversal speedup

**9% predicted speedup,
exactly what we observed**

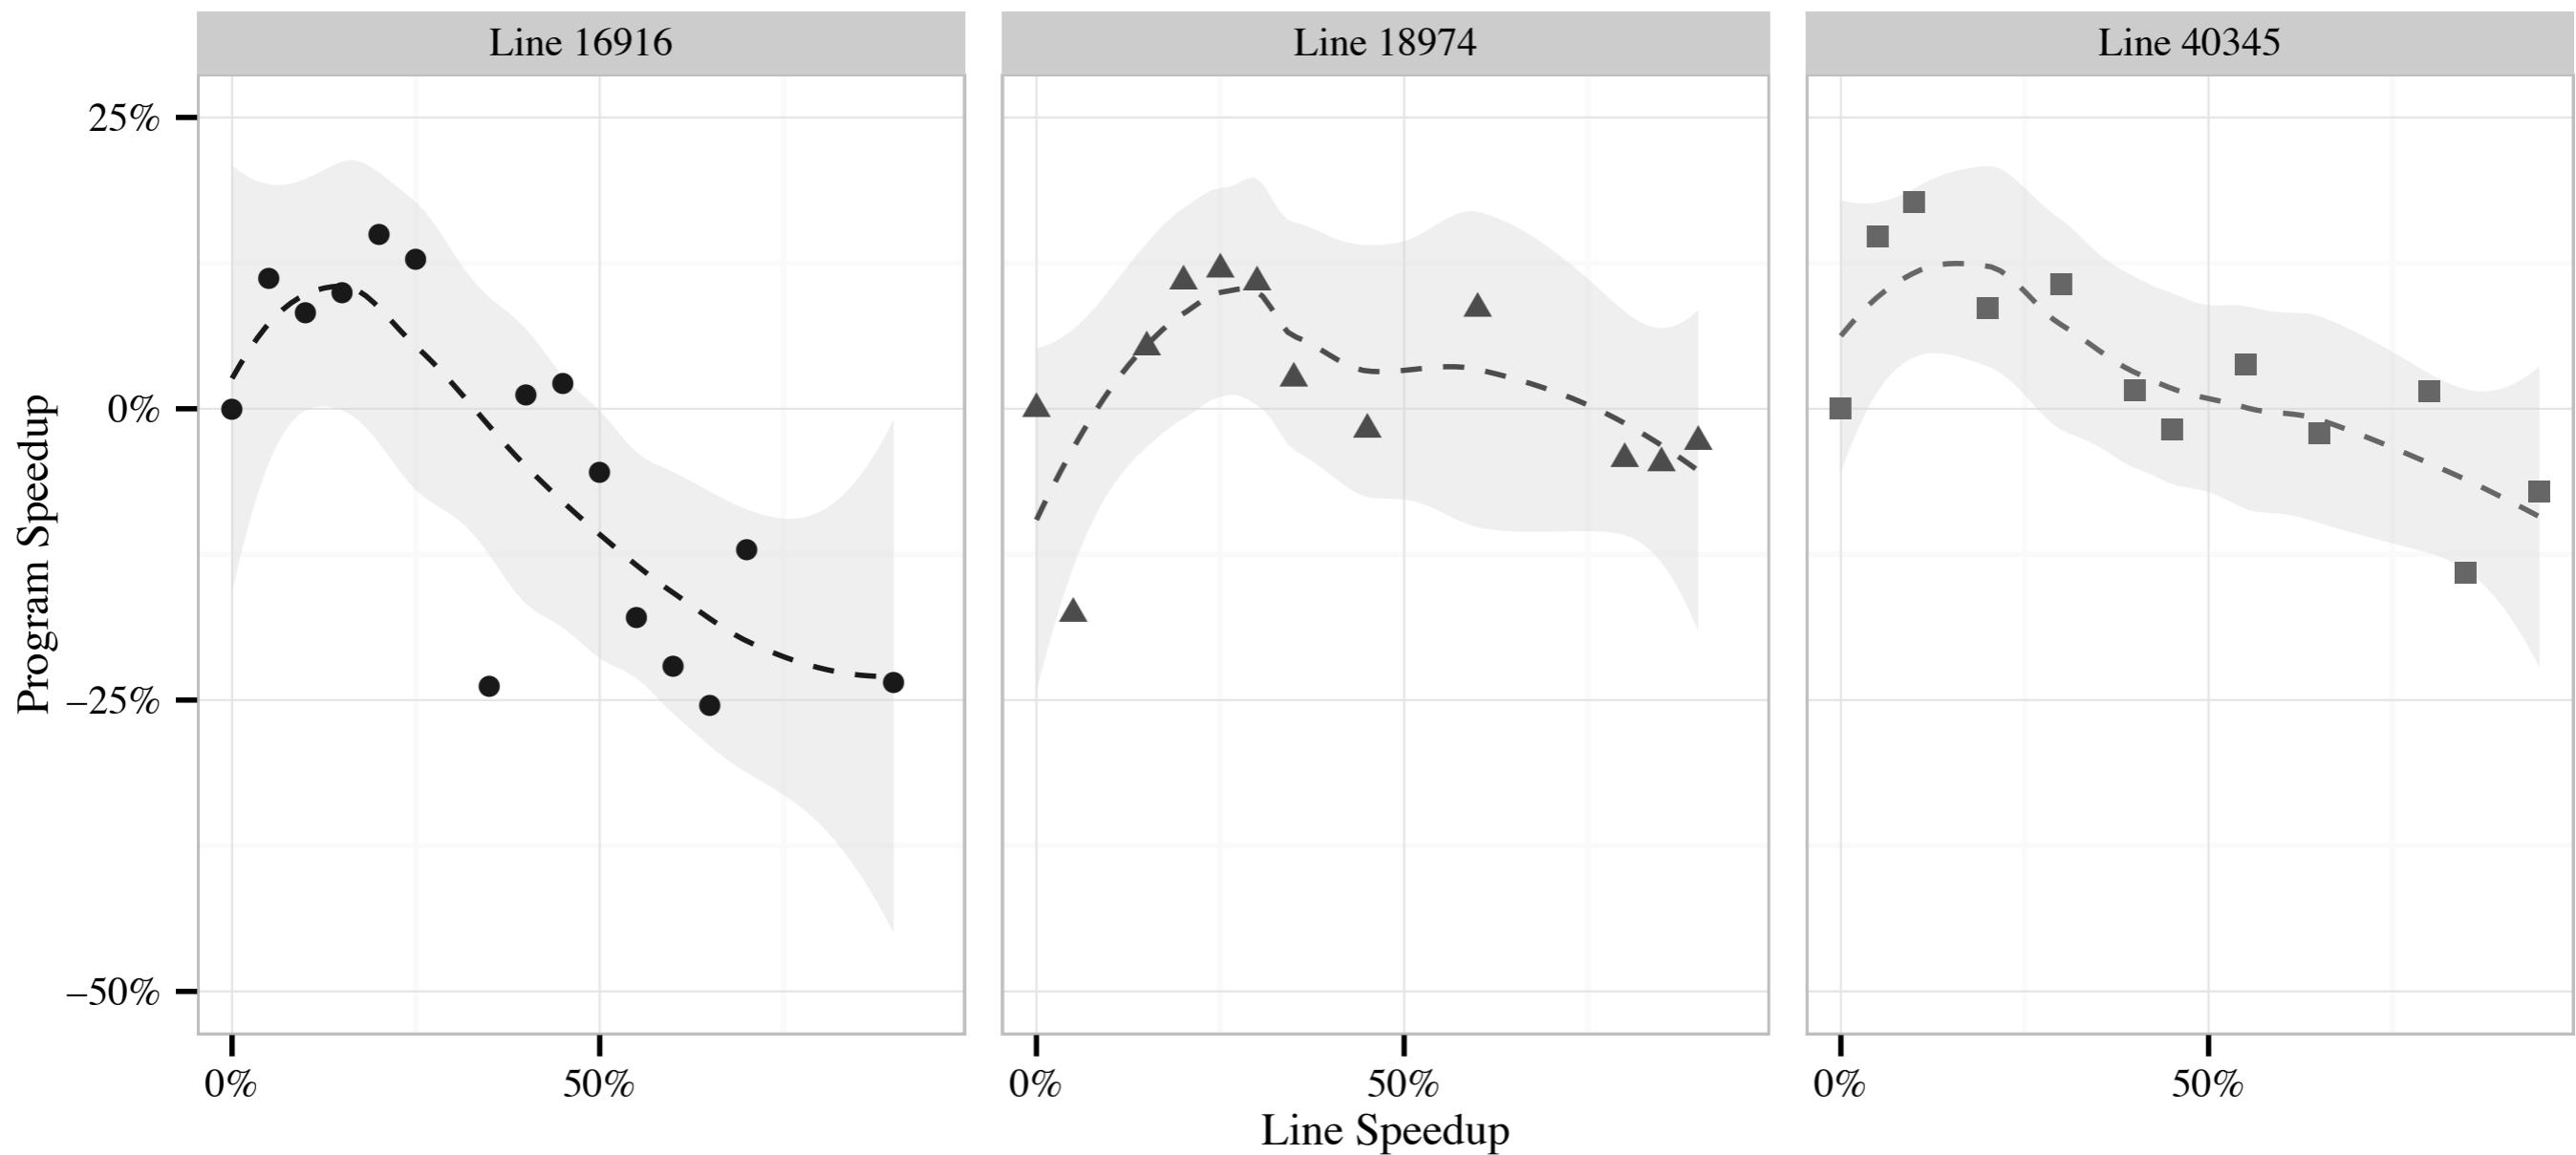


Using Causal Profiling on Ogle





Simple SQL Database





Simple SQL Database

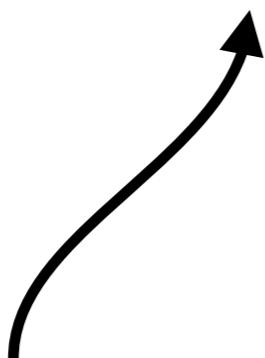
```
#if THREAD_SAFE
config_t global_config = {
...
    .unlock = pthread_mutex_unlock,
    .getsize = sqlite_usable_size,
    .nextitem = sqlite_pagecache_next,
...
};

#endif
```



Simple SQL Database

```
void sqlite_unlock(lock* l) {  
    global_config.unlock(l);  
}
```



Indirect Call

Cheap, but almost the same cost
as pthread_mutex_unlock



Simple SQL Database

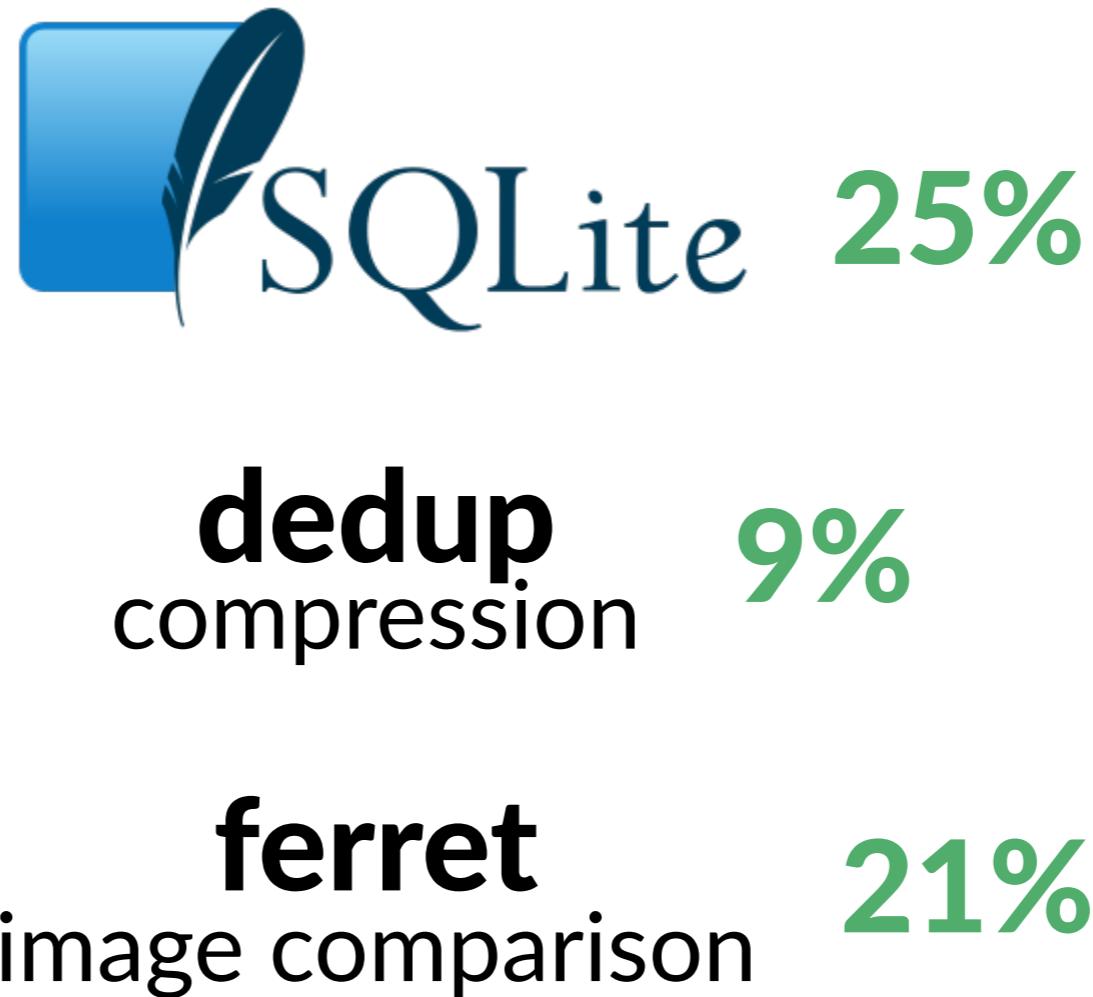
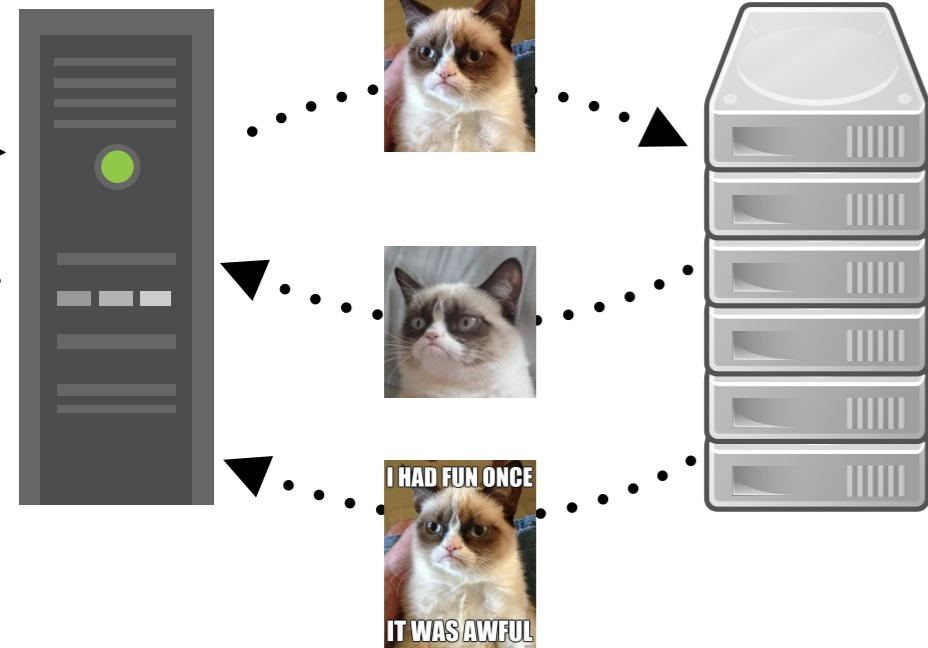
Coz highlights
these lines

```
void sqlite_unlock(lock* l) {  
    global_config.unlock(l); ←  
}
```

```
void sqlite_getsize(void* p) {  
    global_config.getsize(p); ←  
}
```

```
void sqlite_nextitem(item* i) {  
    global_config.nextitem(i); ←  
}
```

Using Causal Profiling on Ogle

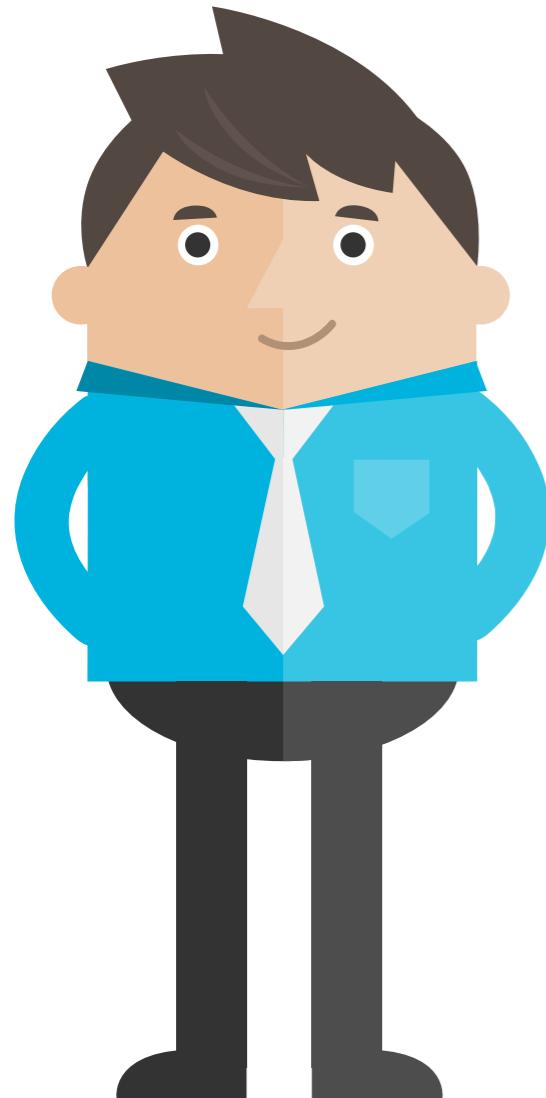


Summary of Optimizations

Benchmark	Speedup	Diff Size	Change Summary
memcached	9.39%	-6, +2	removed unnecessary locks
sqlite	25.60%	-3, +3	removed DIY vtable implementation
blackscholes	2.56%	-61, +4	manual common subexpression elimination
dedup	8.95%	-3, +3	fixed degenerate hash function
ferret	21.27%	-4, +4	rebalanced pipeline thread allocation
fluidanimate	37.50%	-1, +0	removed custom barrier with high contention
streamcluster	68.40%	-1, +0	removed custom barrier with high contention
swaptions	15.80%	-10, +16	reordered loop nests

Sound Performance Analysis

STABILIZER



Effective Performance Profiling

USE THE COZ

% sudo apt install coz-profiler

